

## Week 11: Model Efficiency & Optimization

### From 700GB to 40GB: Making AI Deployable

BSc Natural Language Processing

Discovery-Based Learning Approach

2025

# The 350GB Problem

## The Scenario:

You want to run GPT-3 locally for privacy

Your laptop has 16GB RAM


GPT-3 model size: 350GB

## The Impossibility:

Model is 22x larger than your RAM

Loading would require 175GB of disk swap

Inference: 1 token per minute (unusable)



## The Discovery:

"The 350GB problem has a 40GB solution"

4-bit quantization: 75% size reduction

Accuracy loss: only 3%

Discovery Question: How would YOU make a huge model fit on a small device?

# Paradigm Shift: From Smaller Models to Compressed Models

## OLD Approach (2015):

**Problem:** Large model won't fit

**Solution:** Train a smaller model

### Example:

- GPT-2: 1.5B params → 117M params
- Size: 6GB → 500MB
- Accuracy: 85% → 67%
- Loss: 18 percentage points

### Trade-off:

Smaller size, much worse performance

## NEW Approach (2024):

**Problem:** Large model won't fit

**Solution:** Compress the large model

### Example:

- GPT-3: 175B params (same capability)
- Size: 700GB → 87GB (INT4)
- Accuracy: 92% → 89%
- Loss: 3 percentage points

### Trade-off:

Much smaller size, minimal performance loss

**Key Insight:** Compress post-training preserves learned knowledge better than training smaller

## On-Device LLMs:

### 1. LLaMA-2 7B on Phone

- Original: 28GB (FP32)
- Compressed: 3.5GB (4-bit)
- Method: Quantization
- Performance: 15 tokens/sec

### 2. Whisper in Browser

- Original: 3GB (large model)
- Compressed: 150MB (distilled)
- Method: Knowledge distillation
- Performance: Real-time transcription

## Edge Computing:

### 3. BERT on Arduino

- Original: 440MB (base)
- Compressed: 2MB (pruned + quantized)
- Method: 95% pruning + INT8
- Performance: 200ms inference

### 4. GPT-4 API Efficiency


- Latency: 800ms → 150ms
- Cost: \$0.03/1K → \$0.006/1K
- Method: Mixed precision + distillation
- Scale: Billions of requests/day

---

Deployment Reality: Compression enables AI everywhere (phones, browsers, microcontrollers)

# Foundation 1: Model Size Problem (Visual)

The 350GB Problem: Most Models Won't Fit on Your Laptop  
Compression makes the impossible possible



# Foundation 1: Model Size Problem (Detailed)

## Memory Hierarchy:

Level	Size	Speed
L1 Cache	256KB	1ns
L2 Cache	8MB	5ns
L3 Cache	32MB	20ns
RAM	16GB	100ns
SSD	1TB	100 $\mu$ s

## Fundamental Constraint:

*Inference requires entire model in fast memory*


## Model Size Evolution:

Model	Params	Size (FP32)
BERT-Base	110M	440MB
BERT-Large	340M	1.4GB
GPT-2	1.5B	6GB
GPT-3	175B	700GB
PaLM	540B	2.1TB

**Trend:** Models grow 10 $\times$  every 2 years  
Hardware grows 2 $\times$  every 2 years  
Gap widens without compression

Mathematical Reality: 175B params  $\times$  4 bytes/param = 700GB minimum memory

## Foundation 2: Compression Spectrum (Visual)



### The Spectrum:

Lossless (perfect accuracy, small gains) to Lossy (large gains, small accuracy loss)

### Lossless Methods:

#### 1. Weight Sharing

- Technique: Cluster similar weights
- Reduction: 10-20%
- Accuracy: 100% preserved
- Use case: When zero loss required

#### 2. Low-Rank Factorization

- Technique:  $W = UV^T$  decomposition
- Reduction: 30-40%
- Accuracy: 99-100%
- Use case: Dense layers

### Lossy Methods:

#### 3. Quantization (INT8)

- Technique: FP32  $\rightarrow$  8-bit integers
- Reduction: 75%
- Accuracy: 95-99%
- Use case: Most deployments

#### 4. Quantization (INT4)

- Technique: FP32  $\rightarrow$  4-bit integers
- Reduction: 87.5%
- Accuracy: 90-97%
- Use case: Mobile/edge devices

---

Design Decision: Choose method based on accuracy tolerance and size requirements



## Server (80-512GB RAM):

### Compression:

- GPT-3: 700GB → 350GB (FP16)
- Method: Mixed precision
- Latency: <100ms

## Edge (4-16GB RAM):

### Compression:

- GPT-3: 700GB → 87GB (INT4)
- Method: Quantization
- Latency: <500ms

## Mobile (2-4GB RAM):

### Compression:

- LLaMA-7B: 28GB → 3.5GB
- Method: 4-bit + pruning
- Battery critical

## Microcontroller (256KB-2MB):

### Compression:


- BERT: 440MB → 2MB
- Method: Prune + distill + INT8
- 200× reduction

---

Deployment Reality: Platform constraints drive compression method selection

# Method 1: Quantization (Visual)

Quantization Tradeoffs: Size vs Accuracy  
Fewer bits = Smaller model but slight accuracy loss



# Method 1: Quantization (Detailed Mathematics)

## Quantization Formula:

### Forward (FP32 $\rightarrow$ INT8):

$$q = \text{round} \left( \frac{x - x_{\min}}{s} \right)$$

where  $s = \frac{x_{\max} - x_{\min}}{255}$  (scale)

### Inverse (INT8 $\rightarrow$ FP32):

$$\hat{x} = q \times s + x_{\min}$$

## Numerical Example:

- Weight:  $x = 0.374$  (FP32)
- Range:  $[-1.0, 1.0]$
- Scale:  $s = 2.0/255 = 0.00784$
- Zero-point: 127
- Quantized:  $q = 175$  (INT8)
- Recovered:  $\hat{x} = 0.376$

## Precision Comparison:

Type	Bits	Range	Precision
FP32	32	$\pm 3.4 \times 10^{38}$	7 digits
FP16	16	$\pm 6.5 \times 10^4$	3 digits
INT8	8	-128 to 127	256 values
INT4	4	-8 to 7	16 values

## Real Results:


- BERT-Base FP32: 440MB, 89.5%
- BERT-Base INT8: 110MB, 89.1%
- BERT-Base INT4: 55MB, 87.8%

## When to Use:

Default choice for most deployments  
Hardware support widely available

## Method 2: Knowledge Distillation (Visual)

Knowledge Distillation: Teacher Trains Student  
Transfer knowledge from large model to small model



## Method 2: Knowledge Distillation (Detailed Process)

### Distillation Loss:

$$\mathcal{L} = \alpha \mathcal{L}_{\text{hard}} + (1 - \alpha) \mathcal{L}_{\text{soft}}$$

### Hard Loss (ground truth):

$$\mathcal{L}_{\text{hard}} = - \sum_i y_i \log p_i^{\text{student}}$$

### Soft Loss (teacher knowledge):

$$\mathcal{L}_{\text{soft}} = - \sum_i p_i^{\text{teacher}} \log p_i^{\text{student}}$$

Temperature scaling:  $p_i = \frac{\exp(z_i/T)}{\sum_j \exp(z_j/T)}$

### Typical Values:

- $\alpha = 0.5$  (equal weighting)
- $T = 3 - 5$  (temperature)

Training Process: Student model trained on teacher's logits (soft targets) + true labels

### Concrete Example:

#### Teacher: BERT-Large

- Parameters: 340M
- Size: 1.4GB (FP32)
- Accuracy: 94.0% (GLUE)
- Inference: 120ms

#### Student: DistilBERT


- Parameters: 66M ( $5\times$  smaller)
- Size: 260MB (FP32)
- Accuracy: 92.5% (1.5% loss)
- Inference: 60ms ( $2\times$  faster)

### When to Use:

- When you need  $>5\times$  compression
- When you can retrain the model
- For production deployment at scale

## Method 3: Pruning (Visual)

Pruning Strategies: Random vs Structured  
Both reduce parameters, but structured is hardware-friendly



**Core Idea:** Remove unimportant weights or neurons from the network

**Size Reduction:** 0% sparsity — 10 $\times$  fewer weights

## Method 3: Pruning (Detailed Strategies)

### Unstructured Pruning:

#### Algorithm:

1. Train full model
2. Compute weight magnitudes  $|w_i|$
3. Remove smallest  $p\%$  weights
4. Fine-tune remaining weights

#### Advantages:

- Highest compression (90-95%)
- Minimal accuracy loss
- Flexible per-layer pruning

#### Disadvantages:

- Irregular sparsity patterns
- Requires sparse matrix support
- Limited hardware acceleration

### Structured Pruning:

#### Algorithm:

1. Train full model
2. Compute neuron/channel importance
3. Remove entire neurons/channels
4. Fine-tune remaining network

#### Advantages:

- Direct hardware speedup
- No special sparse libraries
- Smaller actual model size

#### Disadvantages:

- Lower compression (40-60%)
- More accuracy loss
- Coarser granularity

### When to Use:

## Method 4: Low-Rank Factorization (Visual)

### Matrix Decomposition:

Original weight matrix:

$$W \in \mathbb{R}^{m \times n}$$

Decomposed form:

$$W \approx UV^T$$

where  $U \in \mathbb{R}^{m \times r}$ ,  $V \in \mathbb{R}^{n \times r}$ ,  $r \ll \min(m, n)$

### Parameter Count:

- Original:  $m \times n$
- Factorized:  $m \times r + n \times r = r(m + n)$
- Reduction:  $\frac{mn}{r(m+n)}$

### Numerical Example:

Dense layer:  $1024 \times 1024$

#### Original:

- Parameters:  $1024^2 = 1,048,576$
- Size: 4MB (FP32)

#### Factorized ( $r = 64$ ):

- Parameters:  $64(1024 + 1024) = 131,072$
- Size: 512KB (FP32)
- Reduction: 8x smaller
- Accuracy loss: <1%

### SVD Insight:

Most variance captured by first  $r$  singular values  
Remaining  $(n - r)$  dimensions contribute little

**Mathematical Foundation:** Singular Value Decomposition (SVD) provides optimal low-rank approximation

## Method 4: Low-Rank Factorization (Detailed Analysis)

### SVD Algorithm:

#### Step 1: Compute SVD

$$W = U\Sigma V^T$$

where  $\Sigma = \text{diag}(\sigma_1, \dots, \sigma_n)$  with  $\sigma_1 \geq \sigma_2 \geq \dots$

#### Step 2: Choose rank $r$

Energy threshold:  $\frac{\sum_{i=1}^r \sigma_i^2}{\sum_{i=1}^n \sigma_i^2} \geq 0.95$

#### Step 3: Truncate

$$W_r = U_r \Sigma_r V_r^T$$

where  $U_r \in \mathbb{R}^{m \times r}$ ,  $\Sigma_r \in \mathbb{R}^{r \times r}$ ,  $V_r \in \mathbb{R}^{n \times r}$

#### Step 4: Absorb $\Sigma_r$

$$W_r = (U_r \sqrt{\Sigma_r})(\sqrt{\Sigma_r} V_r^T)$$

**Compression Sweet Spot:**  $r \approx 10 - 20\%$  of original dimension balances size and accuracy

### Real Results:

#### BERT Embedding Layer:

- Original:  $30K \times 768 = 23M$  params
- Factorized ( $r = 128$ ):  $128(30K + 768) = 4M$
- Reduction:  $5.8\times$  smaller
- Accuracy:  $89.5\% \rightarrow 89.2\%$

#### GPT-2 Attention:

- Original:  $768 \times 768 = 590K$  params/layer
- Factorized ( $r = 64$ ):  $64 \times 1536 = 98K$
- Reduction:  $6\times$  smaller
- Accuracy: Minimal loss (<0.5%)

#### When to Use:

Dense linear layers (embeddings, attention)

When weight matrix has low intrinsic rank

Combined with quantization for best results

# Method 5: Weight Sharing (Visual)

## Clustering Approach:

**Before:** Each weight is unique

- 175B unique floating-point values
- Full precision per weight
- High memory requirement

**After:** Weights share codebook

- 256 unique cluster centers
- Indices point to codebook
- 2-4 bits per weight (index)

## Storage:

Codebook:  $k$  values (float)

Indices:  $n$  values (2-4 bits)

Total: Much smaller than  $n$  floats

## K-Means Clustering:

### Algorithm:

1. Collect all  $n$  weights
2. Run k-means with  $k$  clusters
3. Replace each weight with nearest cluster center
4. Store: cluster centers + indices

## Numerical Example:

- Weights:  $[0.72, 0.69, -0.31, -0.28, \dots]$
- Clusters ( $k = 4$ ):  $[0.7, -0.3, 0.0, 1.2]$
- Indices:  $[0, 0, 1, 1, \dots]$  (2 bits each)
- Original: 4 bytes/weight
- Compressed: 0.25 bytes/weight
- Reduction: 16 $\times$  smaller

Weight Sharing: Lossless-to-lossy spectrum depending on number of clusters

## Method 5: Weight Sharing (Detailed Implementation)

### Compression Analysis:

#### Storage Requirements:

Codebook size:  $k$  clusters  $\times$  4 bytes

Index size:  $n$  weights  $\times \lceil \log_2 k \rceil$  bits

Total:  $4k + n\lceil \log_2 k \rceil / 8$  bytes

#### Compression Ratio:

$$\text{Ratio} = \frac{4n}{4k + n\lceil \log_2 k \rceil / 8}$$

#### Example ( $n = 1M$ , $k = 256$ ):

- Original:  $1M \times 4 = 4\text{MB}$
- Codebook:  $256 \times 4 = 1\text{KB}$
- Indices:  $1M \times 1 = 1\text{MB}$  (8 bits)
- Total:  $1\text{MB} + 1\text{KB} \approx 1\text{MB}$
- Ratio:  $4\times$  compression

Hybrid Approach: Weight sharing + quantization achieves  $10\text{-}20\times$  compression

### Accuracy Trade-offs:

Clusters	Compression	Accuracy
$k = 2$	$32\times$	60-70%
$k = 16$	$8\times$	85-90%
$k = 256$	$4\times$	95-99%
$k = 4096$	$2.7\times$	99-100%

#### Real Results:

- BERT ( $k = 256$ ):  $440\text{MB} \rightarrow 110\text{MB}$
- Accuracy:  $89.5\% \rightarrow 89.3\%$
- Combined with pruning:  $10\times$  total

#### When to Use:

When you need lossless compression  
Combined with quantization/pruning  
For weight-heavy models

# Method 6: Mixed Precision Training (Visual)

## Precision Strategy:

### FP32 (Master Weights):

- High precision for gradients
- Prevents underflow
- Kept in optimizer state

### FP16 (Forward/Backward):

- Fast computation ( $2\times$ )
- 50% memory reduction
- Hardware acceleration (Tensor Cores)

### INT8 (Inference):

- Minimal memory
- $4\times$  faster than FP32
- Quantized after training

## Training Loop:

1. **Forward:** FP16 computation
2. **Loss:** FP16 calculation
3. **Loss Scaling:** Multiply by  $2^{14}$
4. **Backward:** FP16 gradients
5. **Unscale:** Divide by  $2^{14}$
6. **Update:** FP32 master weights
7. **Copy:** FP32  $\rightarrow$  FP16 for next iteration

### Loss Scaling:

Prevents gradient underflow in FP16  
Typical scale:  $2^{14}$  to  $2^{16}$

Mixed Precision: Best of both worlds (FP32 stability + FP16 speed)

## Method 6: Mixed Precision Training (Detailed Benefits)

### Speed Improvements:

Model	FP32	Mixed
BERT-Base	280 samples/s	560 samples/s
GPT-2	120 samples/s	240 samples/s
ResNet-50	340 images/s	680 images/s

**Speedup:** Consistent 2× across models

### Memory Savings:

Component	FP32	Mixed
Activations	100%	50%
Gradients	100%	50%
Weights	100%	100%
Optimizer	200%	200%
<b>Total</b>	400%	350%

**Industry Standard:** All large model training uses mixed precision (2020+)

### Hardware Support:

#### NVIDIA Tensor Cores:

- FP16: 125 TFLOPS (V100)
- FP32: 15 TFLOPS (V100)
- Speedup: 8× theoretical
- Real speedup: 2-3× (memory bound)

#### TPU v4:

- BF16: 275 TFLOPS
- FP32: 68 TFLOPS
- Speedup: 4×

#### When to Use:

Training large models (GPT-3, BERT)  
When you have Tensor Core GPUs  
Default for modern PyTorch/TensorFlow

## Method 7: Dynamic & Adaptive Computation (Visual)

### Early Exit Strategy:

**Idea:** Not all inputs need full network

Easy examples: Exit after layer 3

Medium examples: Exit after layer 6

Hard examples: Use all 12 layers

### Mechanism:

- Add classifier at each layer
- Compute confidence score
- If confidence  $>$  threshold, exit
- Otherwise, continue to next layer

### Average Speedup:

- Easy:  $4 \times$  (3 layers vs 12)
- Medium:  $2 \times$  (6 layers vs 12)
- Hard:  $1 \times$  (all 12 layers)
- Overall:  $2.5 \times$  average

Adaptive Computation: Allocate resources based on input complexity

### Adaptive Attention:

**Idea:** Not all tokens need full attention

Important tokens: Full attention

Filler words: Sparse attention

### Example (12-word sentence):

- “The”: 20% attention (2 heads)
- “cat”: 100% attention (8 heads)
- “sat”: 100% attention (8 heads)
- “on”: 20% attention (2 heads)
- “the”: 20% attention (2 heads)
- “mat”: 100% attention (8 heads)

### Computation:

- Full:  $12 \times 8 = 96$  head computations
- Adaptive: 48 head computations
- Reduction: 50%

## Method 7: Dynamic & Adaptive Computation (Detailed Results)

### Early Exit Networks:

#### BERT with 3 exits:

- Exit 1 (Layer 4): 35% of samples
- Exit 2 (Layer 8): 45% of samples
- Exit 3 (Layer 12): 20% of samples

#### Performance:

- Average layers: 6.8 vs 12
- Speedup: 1.76 $\times$
- Accuracy: 89.5%  $\rightarrow$  89.1%
- Loss: 0.4 percentage points

#### Confidence Threshold:

- High (0.95): Safe, slower (1.3 $\times$ )
- Medium (0.85): Balanced (1.76 $\times$ )
- Low (0.75): Risky, faster (2.2 $\times$ )

Research Frontier: Adaptive methods are active area of research (2023-2025)

### Adaptive Attention:

#### GPT-2 with Adaptive Heads:

- Content words: 8 heads (100%)
- Function words: 2 heads (25%)
- Punctuation: 1 head (12.5%)

#### Results:

- Computation: 60% of full model
- Speedup: 1.67 $\times$
- Perplexity: 18.2  $\rightarrow$  18.5
- Quality: Minimal degradation

#### When to Use:

Production with varied input complexity  
When average-case matters more than worst-case  
Combined with other compression methods

## GPT-3 Deployment Impossibility

### The Numbers:

- Parameters: 175 billion
- Precision: FP32 (4 bytes each)
- Total size:  $175B \times 4 = 700\text{GB}$
- Typical server RAM: 64-256GB
- Your laptop RAM: 16GB

### Impossibility Ratio:

Model size / Laptop RAM = 44×

Even high-end servers struggle (3-11× over capacity)

### Consequences:

#### Without Compression:

- Must use disk swap
- Inference: 60 seconds per token
- Unusable for production
- Energy: 500W continuous
- Cost: \$10-50 per query

### Business Impact:

- Cannot deploy locally
- Must use cloud APIs
- Privacy concerns
- Latency issues
- Ongoing costs

**Root Problem:** Model capacity requirements exceed deployment hardware by orders of magnitude

## The Naive Solution:

"If GPT-3 is too big, train GPT-2 instead"

### GPT-3 175B:

- Size: 700GB (FP32)
- Parameters: 175B
- Accuracy: 92% (few-shot)
- Training: \$4.6M

↓ Reduce size 100×

### GPT-2 1.5B:

- Size: 6GB (FP32)
- Parameters: 1.5B
- Accuracy: 67% (few-shot)
- Training: \$50K

**Lesson:** Model capacity matters - smaller models cannot simply be trained to match larger ones

## The Problem:

**Accuracy Drop: 25 Percentage Points**

### Capability Loss:

- GPT-3: Complex reasoning, analogies
- GPT-2: Simple pattern matching
- Emergence: Lost at smaller scale


### Scaling Laws:

Performance  $\propto$  (parameters)<sup>0.3</sup>

To match GPT-3 at 1.5B params:

Need 1000× more data (impossible)

# Performance Analysis: Why Smaller Models Fail



# Root Cause: The Capacity Hypothesis

## Theoretical Framework:

### Model Capacity:

$$C = f(\text{parameters, architecture})$$

### Knowledge Stored:

$$K \leq C$$

### Performance:

$$P \propto K$$

### Implications:

- Smaller model  $\Rightarrow$  Less capacity
- Less capacity  $\Rightarrow$  Less knowledge
- Less knowledge  $\Rightarrow$  Worse performance

Diagnosis: Performance tied to parameter count - compression must preserve parameters

## Numerical Evidence:

Model	Params	Accuracy
BERT-Tiny	14M	78%
BERT-Small	28M	83%
BERT-Medium	66M	86%
BERT-Base	110M	89.5%
BERT-Large	340M	94%

### Solution Requirement:

Preserve model capacity (parameters)  
Reduce storage/memory footprint  
 $\Rightarrow$  Compression, not replacement

# Solution Insight: Compress Post-Training

## The Breakthrough Idea:

**OLD:** Train small model (loses knowledge)

↓

**NEW:** Train large, then compress

## Why This Works:

1. Train full-capacity model
2. Model learns all knowledge
3. Compress learned weights
4. Knowledge preserved (mostly)
5. Fit in deployment memory

## Key Observation:

Learned weights have structure

Structure enables compression

Random weights don't compress well

Critical Insight: Trained weights have exploitable structure that random weights lack

## Compression Opportunity:

### Trained Weights Properties:

- Clustered values (weight sharing)
- Low effective rank (factorization)
- Many near-zero (pruning)
- Narrow range (quantization)

## Concrete Example:

- BERT attention weights
- 95% of weights in [-0.5, 0.5]
- Can use 8 bits instead of 32
- 4× compression with 0.4% loss

## Contrast with Random:

- Random weights: Uniform distribution
- No structure to exploit
- Compression hurts accuracy severely

# Quantization Mechanism: FP32 $\rightarrow$ INT8

## The Math:

### Quantization Function:

$$q = \text{round} \left( \frac{x - x_{\min}}{s} \right)$$

where scale  $s = \frac{x_{\max} - x_{\min}}{255}$

### Dequantization Function:

$$\hat{x} = q \times s + x_{\min}$$

### Error:

$$\epsilon = |\hat{x} - x| \leq \frac{s}{2}$$

### Key Idea:

- Map range  $[x_{\min}, x_{\max}]$  to  $[0, 255]$
- Store integer index (1 byte)
- Recover approximate value

Quantization Error: Bounded by half the quantization step (0.00392 in this example)

## Numerical Walkthrough:

### Weight Layer Statistics:

- Min: -1.2
- Max: +0.8
- Range: 2.0
- Scale:  $s = 2.0 / 255 = 0.00784$

### Quantize $x = 0.374$ :

1. Shift:  $0.374 - (-1.2) = 1.574$
2. Scale:  $1.574 / 0.00784 = 200.76$
3. Round:  $q = 201$  (INT8)

### Dequantize $q = 201$ :

1. Unscale:  $201 \times 0.00784 = 1.576$
2. Unshift:  $1.576 + (-1.2) = 0.376$
3. Error:  $|0.376 - 0.374| = 0.002$

# Worked Example: BERT Layer Quantization

## Layer: BERT Attention Weights

### Original (FP32):

- Shape:  $768 \times 768$
- Weights: 590,592
- Min:  $-0.487$
- Max:  $+0.512$
- Mean: 0.003
- Std: 0.124
- Size:  $590K \times 4 = 2.36\text{MB}$

### Quantization Parameters:

- Range:  $[-0.487, 0.512]$
- Scale:  $(0.512 - (-0.487))/255 = 0.00392$
- Zero-point: 127 (symmetric)

Real Result: 4× compression with  $\pm 0.5\%$  accuracy loss on BERT-Base

### Quantized (INT8):

- Shape:  $768 \times 768$  (unchanged)
- Values: INT8 in  $[0, 255]$
- Size:  $590K \times 1 = 590\text{KB}$
- Reduction: 4× smaller

### Sample Weights:

FP32	INT8	Recovered
0.374	201	0.376
-0.251	67	-0.249
0.089	150	0.090
-0.412	25	-0.413
0.501	255	0.512

### Accuracy:

- Original BERT: 89.5%
- Quantized INT8: 89.1%
- Loss: 0.4 percentage points

## BERT-Base Compression:

Method	Size	Accuracy
FP32 Baseline	440MB	89.5%
FP16	220MB	89.5%
INT8	110MB	89.1%
INT4	55MB	87.8%
INT8 + Pruning	22MB	87.5%

## Best Trade-off:

- INT8: 4× smaller, 0.4% loss
- Production standard (2024)
- Hardware accelerated

## GPT-3 Compression:

Precision	Size	Quality
FP32	700GB	100%
FP16	350GB	100%
INT8	175GB	98%
INT4	87GB	95%


## Deployment Reality:

- OpenAI API: INT8 (likely)
- 4× memory reduction
- 2× throughput increase
- Enables profitable deployment

---

Industry Adoption: All major LLM APIs use INT8 quantization (2024)

# Decision Tree: Choosing Compression Method



## Quantization:

### Avoid when:

- Model has high dynamic range
- Batch norm layers (unstable)
- Small models (<100M params)
- Research/debugging phase

## Distillation:

### Avoid when:

- No budget to retrain
- Teacher model unavailable
- Task requires all model capacity
- Target is <5× compression

## Pruning:

### Avoid when:

- No sparse matrix libraries
- Model already small
- All weights are important
- Cannot fine-tune after pruning

## Low-Rank:

### Avoid when:

- Weights are full-rank
- Convolutional layers (better methods)
- Recurrent connections
- Model has few dense layers

---

Anti-Patterns: Know when NOT to use each method to avoid wasted effort

## Pitfall 1: Calibration

### Problem:

- Quantize with training data ranges
- Deploy on different distribution
- Activation ranges differ
- Severe accuracy drop

### Solution:

- Calibrate on representative data
- 1000+ diverse examples
- Measure activation ranges
- Use percentile (99%) not max

## Pitfall 2: INT4 Overflow

### Problem:

- INT4 range:  $[-8, 7]$
- Outlier weights cause clipping

## Pitfall 3: Distillation Failure

### Problem:

- Student too small ( $>20\times$  smaller)
- Cannot learn teacher's knowledge
- Converges to random baseline

### Solution:

- Limit compression to 5-10 $\times$
- Use intermediate layers
- Progressive distillation

## Pitfall 4: Compound Methods

### Problem:

- Prune + quantize + distill = fail
- Errors compound
- $10\% + 5\% + 3\% \neq 18\%$
- Actual: 25% degradation

# Success Metrics: How to Measure

## Primary Metrics:

### 1. Size Reduction

$$R = \frac{\text{Original Size}}{\text{Compressed Size}}$$

Target: 4-10× for deployment

### 2. Accuracy Preservation

$$A = \frac{\text{Compressed Accuracy}}{\text{Original Accuracy}}$$

Target: >95% (absolute <3% loss)

### 3. Latency Improvement

$$L = \frac{\text{Original Latency}}{\text{Compressed Latency}}$$

Target: 2-4× speedup

### 4. Energy Efficiency


## Real Benchmark (BERT):

Method	R	A	L	E
Baseline	1×	100%	1×	1×
FP16	2×	100%	1.5×	1.8×
INT8	4×	99%	2.5×	3.2×
INT4	8×	96%	3.5×	5.1×
Pruned	10×	95%	1.2×	1.5×

## Trade-off Analysis:

- INT8: Best balance (4×, 99%, 2.5×)
- INT4: Maximum compression
- Pruning: Size without speedup (need sparse support)

## The Compression Pareto Frontier Best Accuracy-Size Tradeoffs for Different Methods



## Smartphone LLMs (2024):

### Apple Intelligence (iPhone 15):

- Model: 3B parameter LLM
- Original: 12GB (FP32)
- Compressed: 1.5GB (4-bit + pruning)
- Methods: INT4 + 50% pruning
- Performance: 30 tokens/sec
- Privacy: 100% on-device

### Google Gemini Nano:

- Model: 1.8B parameters
- Size: 900MB (INT8)
- Latency: 40 tokens/sec
- Battery: 1% per 1000 tokens

## Edge Computing:

### Raspberry Pi 4 (8GB):

- LLaMA-2 7B quantized (INT4)
- Size: 3.5GB
- Speed: 2 tokens/sec
- Use case: Local assistant

### NVIDIA Jetson (16GB):

- GPT-J 6B (INT8)
- Size: 6GB
- Speed: 15 tokens/sec
- Use case: Robotics, drones

## Impact:

Compression enables privacy-preserving AI  
Zero cloud dependency  
Millisecond latency

2024 Reality: Compression makes AI ubiquitous (phones, cars, appliances)

# Implementation: PyTorch Quantization in 15 Lines

## Dynamic Quantization:

```
import torch

# Load pre-trained model
model = BertForSequenceClassification
    .from_pretrained('bert-base')

# Quantize to INT8
quantized_model = torch.quantization
    .quantize_dynamic(
        model,
        {torch.nn.Linear},
        dtype=torch.qint8
    )

# Save compressed model
torch.save(quantized_model,
    'bert_int8.pt')
```

**Result:** 440MB → 110MB (4×)

Production Code: PyTorch provides built-in quantization (torch.quantization module)

## Static Quantization:

```
# Prepare model
model.qconfig = torch.quantization
    .get_default_qconfig('fbgemm')
torch.quantization.prepare(model)

# Calibrate with data
for batch in calibration_data:
    model(batch)

# Convert to INT8
quantized_model = torch.quantization
    .convert(model)

# Inference
with torch.no_grad():
    output = quantized_model(input)
```

**Advantage:** Better accuracy (calibrated ranges)

## Model Efficiency Fundamentals

### 1. Compression Preserves Knowledge

Train large, compress post-training beats training small

Example: GPT-3 INT4 (87GB) outperforms GPT-2 (6GB)

### 2. Quantization is the Default

4× reduction, <1% accuracy loss, hardware accelerated

Use INT8 unless you have specific constraints

### 3. Platform Drives Strategy

Server: FP16/INT8 — Edge: INT4 — Mobile: INT4+Pruning — MCU: Distillation+INT8

Deployment memory determines compression needs

### 4. Combine Methods Carefully

Quantization + (Pruning OR Distillation) works

All three together compounds errors

### 5. Measure Four Metrics

Size reduction, accuracy, latency, energy

Optimize for the bottleneck

---

**Summary:** Compression makes modern AI deployable everywhere

### Week 11 Lab:

#### Hands-On Activities:

1. Quantize BERT (FP32 → INT8)
2. Measure size, accuracy, latency
3. Distill GPT-2 (1.5B → 300M)
4. Prune ResNet (90% sparsity)
5. Deploy quantized model

#### Tools:

- PyTorch quantization API
- Hugging Face transformers
- ONNX Runtime

#### Deliverable:

Compress a model 10× with <3% accuracy loss

Bridge to Ethics: Efficiency enables sustainable, accessible, democratized AI

### Week 12: Ethics & Fairness

#### Efficiency → Ethics Link:

#### Sustainability:

- GPT-3 training: 1287 MWh
- Carbon: 550 tons CO<sub>2</sub>
- Compression reduces deployment energy 5×

#### Accessibility:

- On-device AI: No cloud required
- Privacy-preserving inference
- Works in low-connectivity regions

#### Democratization:

- Run LLMs on \$200 hardware
- No API costs
- Open access to AI