

## Fine-tuning & Prompt Engineering

### Week 10 - BSc Discovery-Based Pedagogy

NLP Course 2025

October 2025

# The \$50,000 Question

**Scenario:** You work at a medical AI company

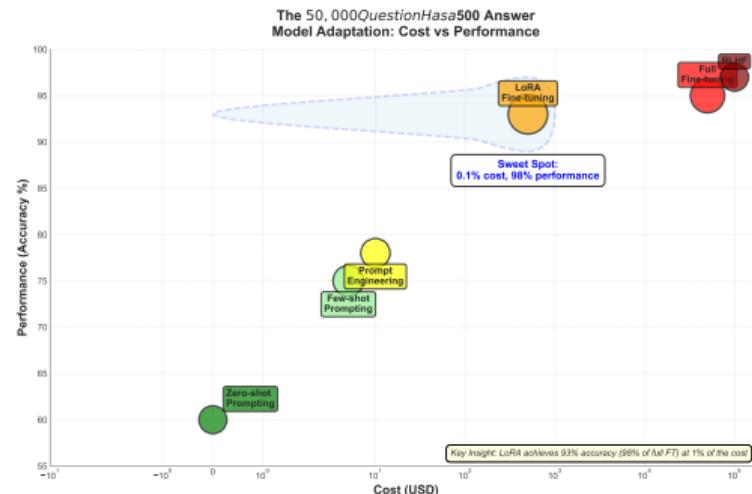
## What You Have:

- GPT-4: Amazing at general text
- 1,000 labeled medical diagnoses
- Goal: 90%+ accuracy on medical QA
- Budget: Limited

## The Problem:

- GPT-4 zero-shot: 60% accuracy
- Full fine-tuning: \$50,000+
- Training time: 2 weeks on 8 GPUs
- Risk: Catastrophic forgetting

## What Would YOU Do?



## The Answer: LoRA fine-tuning

- Cost: \$500 (1% of full FT)
- Accuracy: 93% (98% of full FT)
- Time: 6 hours

# Paradigm Shift: OLD vs NEW Adaptation

## OLD: Train Everything

Traditional approach (pre-2018):

- Train 175B parameters from scratch
- Or fine-tune ALL weights
- Cost: \$5M+ for training
- Memory: 700GB+ required
- Time: Weeks to months
- Risk: Overfitting, forgetting

### Examples:

- BERT (2018): 110M params, full FT
- GPT-2 (2019): 1.5B params, full FT
- Every task needs full retraining

### The Problem:

Not scalable! Imagine updating a model for 100 different tasks - you'd need 100 full copies!

## NEW: Adapt Efficiently

Modern approach (2021+):

- Freeze 175B base parameters
- Update only 0.1-1% task-specific
- Cost: \$100-\$5K for adaptation
- Memory: Same as inference
- Time: Hours to days
- Benefit: Preserve base knowledge

### Examples:

- LoRA (2021): 0.1% params
- Adapters: 0.5-2% params
- Prompt tuning: 0 params!

### The Breakthrough:

100 tasks = 1 base model + 100 tiny adapters (each 10MB) instead of 100 full models (each 350GB)!

---

Parameter-Efficient Fine-Tuning (PEFT) enables scaling to thousands of tasks

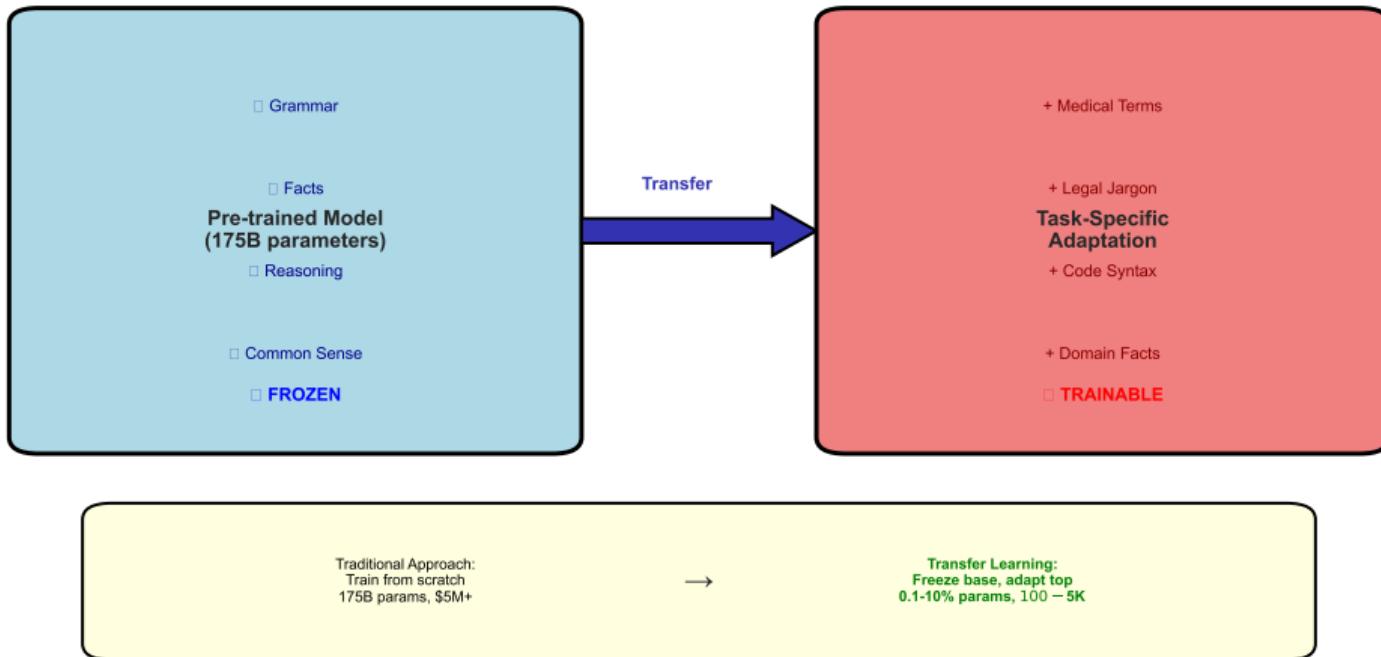
# Real-World Impact: Who's Using What in 2024

## Real-World Applications 2024: Model Adaptation in Production

<b>Bloomberg GPT</b> <i>[Full Fine-tuning]</i>	Domain: Finance Data: 363B tokens <b>Cost: \$2.7M+</b>	<b>Result:</b> 50B params, SOTA finance
<b>Med-PaLM 2</b> <i>[Instruction Tuning]</i>	Domain: Medical Data: Medical QA datasets <b>Cost: \$500K+</b>	<b>Result:</b> 85% on USMLE
<b>Code Llama</b> <i>[LoRA]</i>	Domain: Programming Data: 500B code tokens <b>Cost: \$50K</b>	<b>Result:</b> 53% HumanEval
<b>GPT-4 Custom</b> <i>[Prompt Engineering]</i>	Domain: Customer Service Data: 0 training <b>Cost: \$0</b>	<b>Result:</b> 90% satisfaction
<b>LegalBERT</b> <i>[Domain Fine-tuning]</i>	Domain: Legal Data: 12GB legal docs <b>Cost: \$10K</b>	<b>Result:</b> 89% on legal NER

# Transfer Learning: Visual Overview

## Transfer Learning: Reuse General Knowledge, Adapt for Specifics



**Key Insight:** 99% of language knowledge is reusable – only adapt the 1% that's task-specific

# Transfer Learning: How It Works

## The Concept:

Pre-trained models already know:

- Grammar and syntax
- Common sense reasoning
- World knowledge
- General patterns

What they DON'T know:

- Your specific domain (medical, legal)
- Your task format
- Your company's style
- Your special vocabulary

## The Math:

Total knowledge = Base (99%) + Task (1%)

Instead of learning 100%, we only learn the missing 1%!

## When to Use:

- You have a pre-trained model
- Your task is related to general language
- You have limited compute budget
- You want to avoid training from scratch

## Three Approaches:

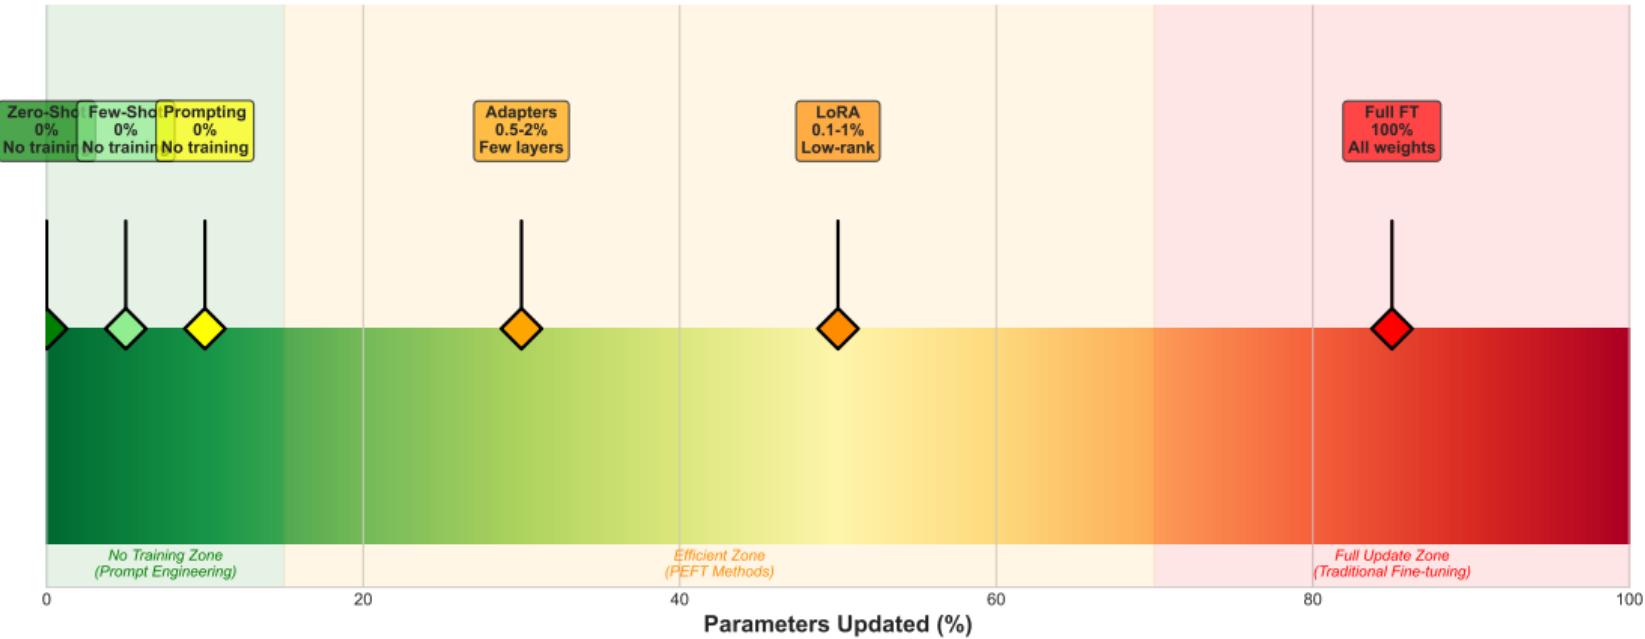
1. **Prompting** (0% training)
  - Pros: Free, instant
  - Cons: Limited accuracy
2. **PEFT** (0.1-2% training)
  - Pros: Efficient, effective
  - Cons: Needs some labeled data
3. **Full FT** (100% training)
  - Pros: Maximum accuracy
  - Cons: Expensive, risky

---

Choose based on your data, budget, and accuracy requirements

# The Parameter Update Spectrum: Visual

The Parameter Update Spectrum: From Zero Training to Full Fine-tuning



**Key Insight:** Model adaptation is a spectrum, not a binary choice

From 0% (prompting) to 100% (full fine-tuning) - choose your efficiency point

# The Parameter Update Spectrum: Details

## Zero Training Zone (0%):

### Zero-Shot:

- Just ask directly
- Example: "Translate to French: Hello"
- Accuracy: 40-70%
- Cost: \$0

### Few-Shot:

- Provide 3-5 examples in prompt
- Model learns pattern on-the-fly
- Accuracy: 60-80%
- Cost: \$0 (just longer prompts)

### Prompt Engineering:

- Carefully craft instructions
- Role, task, format, examples
- Accuracy: 70-85%
- Cost: \$0 (+ human time)

## Efficient Zone (0.1-2%):

### Adapters:

- Small modules between layers
- Update: 0.5-2% of parameters
- Accuracy: 85-92%
- Cost: \$500-\$5K

### LoRA:

- Low-rank matrix updates
- Update: 0.1-1% of parameters
- Accuracy: 88-94%
- Cost: \$100-\$2K

## Full Update Zone (100%):

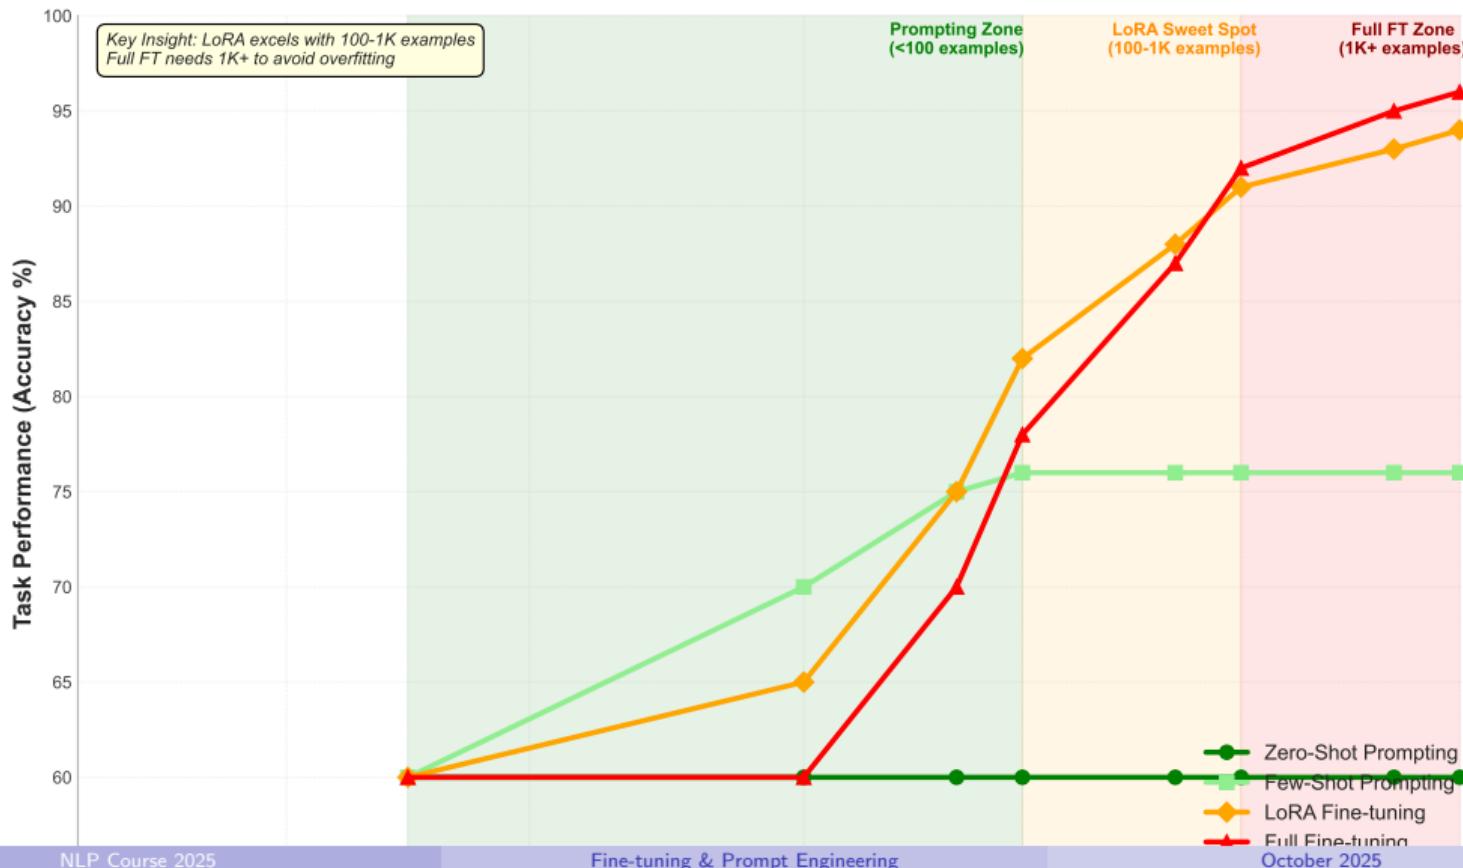
### Full Fine-tuning:

- Update all weights
- Accuracy: 90-97%
- Cost: \$10K-\$100K
- Risk: Catastrophic forgetting

The sweet spot is usually LoRA: 90%+ accuracy at 1% cost

# Data Requirements: Performance vs Data Size

## Data Efficiency: Performance vs Training Data Size



## What Data Do You Have?

### 0-10 Examples:

- Use: Few-shot prompting
- Why: Not enough for training
- Expected: 60-75% accuracy
- Example: New task, just starting

### 10-100 Examples:

- Use: Prompt engineering
- Why: Still too few for training
- Expected: 70-80% accuracy
- Example: Prototyping phase

### 100-1,000 Examples:

- Use: LoRA fine-tuning
- Why: Enough for efficient training
- Expected: 85-93% accuracy
- Example: Production-ready

### 1,000-10,000 Examples:

- Use: LoRA or Full fine-tuning
- Why: Can consider full updates
- Expected: 90-95% accuracy
- Example: Large-scale production

### 10,000+ Examples:

- Use: Full fine-tuning
- Why: Enough to avoid overfitting
- Expected: 93-97% accuracy
- Example: Critical applications

### Quality Matters Too!

- 100 high-quality & 1000 noisy
- Diverse examples beat repetitive
- Representative of real use cases
- Balanced class distribution

---

Data quantity AND quality determine which method works best

# Zero-Shot Prompting: The Simplest Approach

## What is Zero-Shot?

Just ask the model directly - no examples, no training!

## Example:

**Prompt:** Classify sentiment: "The movie was terrible"

**Response:** Negative

## How It Works:

- Model uses pre-trained knowledge
- Interprets task from instruction
- No task-specific training
- Works for common tasks

**Parameters Updated:** 0%

**Cost:** Free (just API calls)

**Time:** Instant

## When to Use:

- You have NO training data
- Task is straightforward
- Quick prototype/experiment
- Budget is very limited

## When NOT to Use:

- Need >85% accuracy
- Domain-specific terminology
- Complex reasoning required
- Consistent format needed

## Real Example:

GPT-4 zero-shot for basic customer support:

- Task: Categorize customer emails
- Accuracy: 75%
- Cost: \$0.10 per 1000 emails
- Time: Real-time

Good enough for low-stakes applications!

Zero-shot is free and fast but limited to 60-75% accuracy on most tasks

# Few-Shot In-Context Learning: Show Examples

## What is Few-Shot?

Provide 3-5 examples IN THE PROMPT - model learns the pattern!

### Example:

**Prompt:** Classify sentiment:

Example 1: "I loved it!" → Positive

Example 2: "Terrible experience" → Negative

Example 3: "It was okay" → Neutral

Now classify: "Amazing product!"

**Response:** Positive

### How It Works:

- Model sees input-output pairs
- Infers pattern from examples
- Applies to new input
- Still no weight updates!

## When to Use:

- You have 5-50 examples
- Task has clear pattern
- Need quick improvements over zero-shot
- Can't afford training

## Best Practices:

- Use diverse examples
- Show edge cases
- Consistent format
- 3-5 examples usually enough

## Performance Boost:

- Zero-shot: 60%
- Few-shot (3 examples): 75%
- Few-shot (5 examples): 78%
- Diminishing returns after 5!

## Limitation:

Context window! GPT-4 has 128K tokens - examples use up

# Prompt Engineering: Anatomy of Effective Prompts

## Anatomy of an Effective Prompt: 6 Essential Components



# Prompt Engineering: Best Practices

## Key Principles:

### 1. Be Specific

- Bad: "Analyze this data"
- Good: "Identify top 3 churn factors with percentages"

### 2. Set Role/Context

- "You are an expert data scientist..."
- Helps model adopt appropriate style

### 3. Specify Output Format

- "Provide response as: 1. ... 2. ... 3. ..."
- Ensures consistency

### 4. Use Chain-of-Thought

- "Let's think step by step..."
- Improves reasoning by 20-30%

### 5. Provide Constraints

- "Use non-technical language"
- "Maximum 100 words"

## Advanced Techniques:

### Self-Consistency:

- Ask same question 5 times
- Take majority vote
- Reduces errors

### Tree-of-Thoughts:

- Explore multiple reasoning paths
- Evaluate each path
- Choose best solution

### Common Pitfalls:

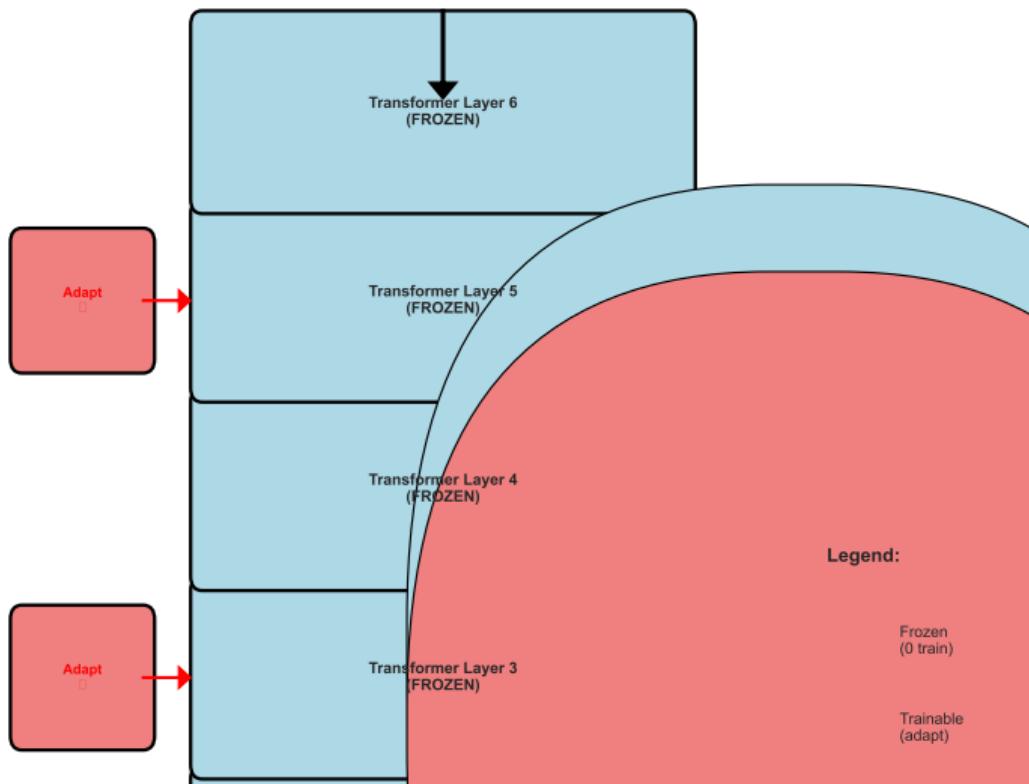
- Too vague: "Make it better"
- Too complex: 10-page prompt
- No examples: Hard to learn pattern
- Conflicting instructions
- No output format specified

### Performance Gain:

- Basic prompt: 65%
- Well engineered: 80-85%

# Adapter Methods: Visual Architecture

Adapter Architecture: Small Trainable Modules  
Inserted Between Frozen Layers



# Adapter Methods: How They Work

## The Concept:

Instead of updating huge matrices in transformer layers, insert small "adapter" modules:

## Architecture:

- ① Freeze all transformer weights
- ② Insert adapter after each layer
- ③ Adapter: Down-project → Activate → Up-project
- ④ Only train adapters

## Math:

Adapter size:  $d \rightarrow r \rightarrow d$

where  $d$  = hidden dim (e.g., 1024),  $r$  = bottleneck (e.g., 64)

Parameters:  $2 \times d \times r = 2 \times 1024 \times 64 = 131K$

Compare to layer:  $d \times d = 1M$

Reduction: 10x!

## When to Use:

- Multiple tasks on same model
- Want to preserve base model
- Memory-constrained environment
- Need modular task switching

## Advantages:

- Efficient: 0.5-2% params
- Modular: Swap adapters easily
- Safe: Base model untouched
- Fast: Quick training

## Disadvantages:

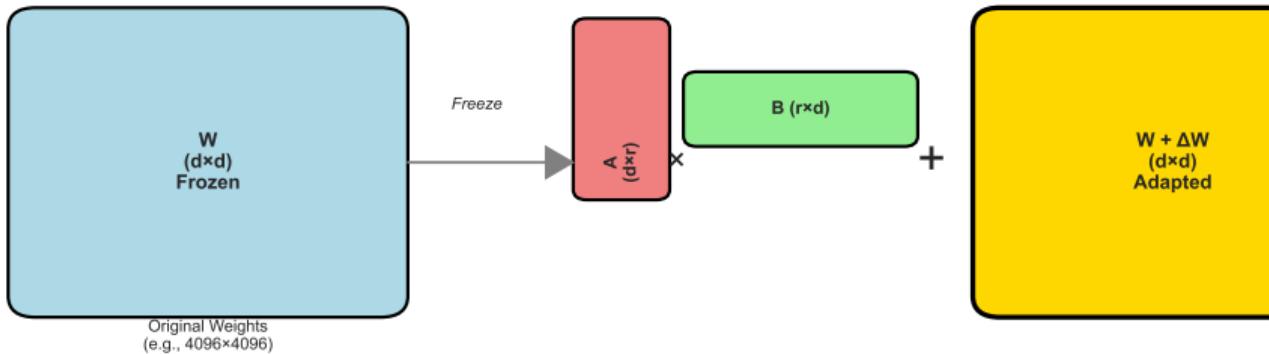
- Added inference latency (small)
- Slightly lower accuracy than LoRA
- Need to choose bottleneck size

## Real Performance:

- GLUE benchmark: 96% of full FT
- Parameters: 1-5% vs 100%

## LoRA: Low-Rank Adaptation

*Instead of updating 16M parameters, update only 32K!*



Example:  $d=4096, r=8$   
Original:  $4096 \times 4096 = 16,777,216$  parameters  
LoRA:  $(4096 \times 8) + (8 \times 4096) = 65,536$  parameters (0.39%!)

# LoRA: Low-Rank Adaptation Explained

## The Problem:

Fine-tuning updates weight matrix  $W \in \mathbb{R}^{d \times d}$

Example:  $d = 4096 \rightarrow W$  has 16.7M parameters!

Too many parameters to update efficiently.

## The Insight:

Updates  $\Delta W$  are typically low-rank!

Instead of full  $\Delta W$ , decompose:

$$\Delta W = A \times B$$

where  $A \in \mathbb{R}^{d \times r}$ ,  $B \in \mathbb{R}^{r \times d}$

with  $r \ll d$  (e.g.,  $r = 8$ ,  $d = 4096$ )

## Parameters:

- Full  $\Delta W$ :  $d^2 = 16.7M$
- LoRA  $A + B$ :  $2dr = 65K$
- Reduction: 256x!

## How It Works:

- ① Freeze pre-trained weights  $W$
- ② Initialize  $A$  (random),  $B$  (zeros)
- ③ During training:
  - Forward:  $h = (W + AB)x$
  - Backward: Only update  $A, B$
- ④ After training: Merge  $W' = W + AB$

## Choosing Rank $r$ :

- $r = 1$ : Too restrictive, poor results
- $r = 4-8$ : Sweet spot
- $r = 16-32$ : Diminishing returns
- $r = 64+$ : No longer efficient

## Performance:

- GPT-3 (175B),  $r = 4$ : 94% of full FT
- Parameters: 0.01% vs 100%
- Training: \$500 vs \$50,000
- No inference overhead (merge weights)

# Full Fine-Tuning: When to Use Maximum Power

## What is Full Fine-Tuning?

Update ALL model parameters for your task.

## How It Works:

- ① Load pre-trained model
- ② Replace final layer for your task
- ③ Unfreeze ALL weights
- ④ Train on your dataset
- ⑤ Save entire model

## The Math:

Model: 175B parameters

Training: Update all 175B weights

Memory:  $4 \times 175B = 700\text{GB}$  (float32)

Gradients: Another 700GB

Optimizer states: Another 700GB

Total: 2.1TB!

## Cost Reality:

- 8x A100 GPUs (80GB each) = 640GB
- Not enough! Need distributed training
- Time: 1-2 weeks

## When to Use:

- Need 95%+ accuracy (critical task)
- Have 10K+ training examples
- Large budget available
- Task very different from pre-training
- Willing to maintain separate model

## When NOT to Use:

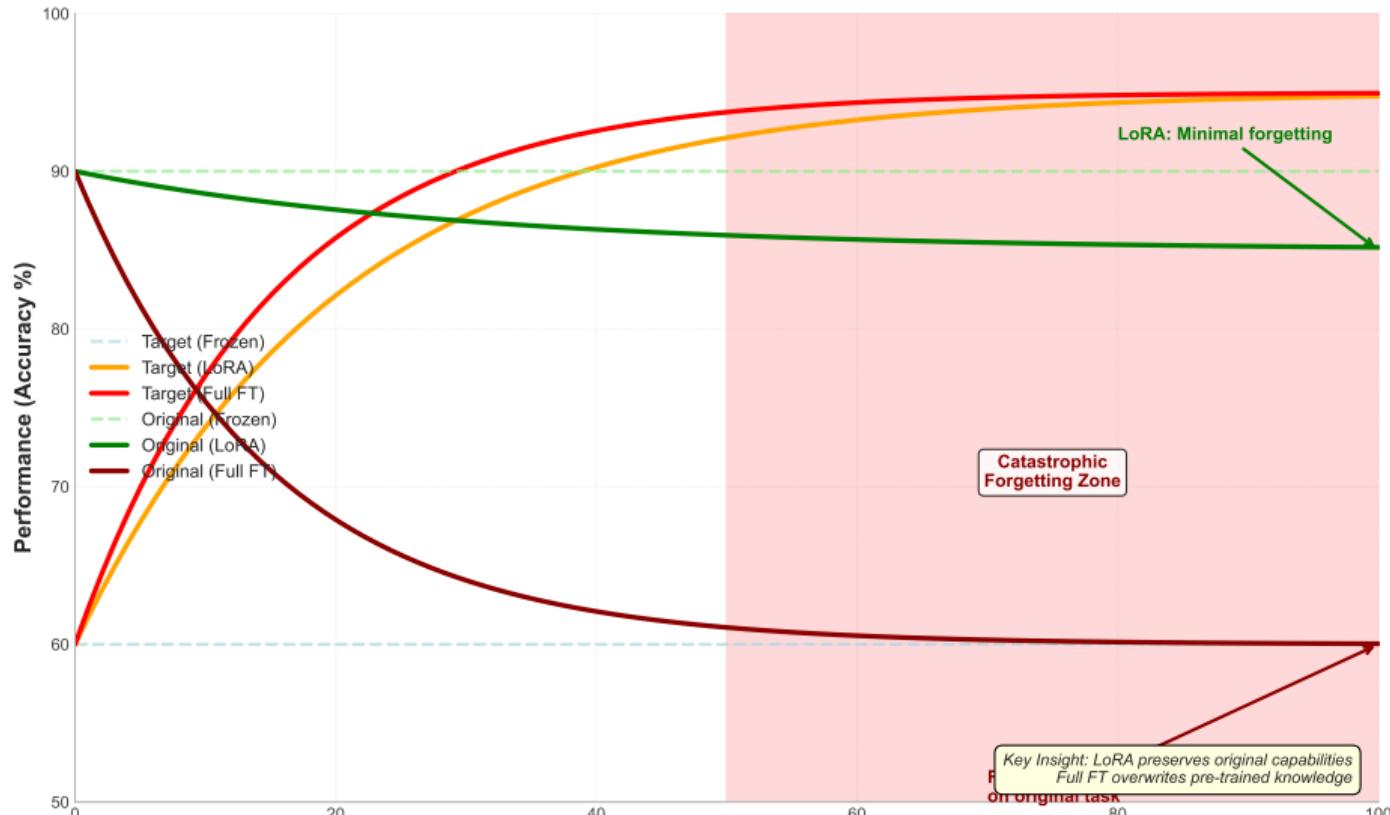
- Limited data (<1000 examples)
- Budget constrained
- Need multiple task adaptations
- Risk of catastrophic forgetting

## Real Examples:

- Bloomberg GPT: \$2.7M+ training
- Trained on 363B finance tokens
- Result: SOTA on financial tasks
- But: Can't be used for general text

# Catastrophic Forgetting: The Risk of Full Updates

Catastrophic Forgetting: Full Fine-tuning vs LoRA  
Original Task Performance Degrades with Full Updates



## The Problem:

Pre-trained models aren't helpful or safe:

- Generate toxic content
- Provide unhelpful responses
- Don't follow instructions well
- No notion of "better" output

## The Solution - 3 Steps:

### Step 1: Supervised Fine-tuning

- Human labelers write ideal responses
- Train model on these examples
- Result: Somewhat helpful model

### Step 2: Reward Model Training

- Generate multiple responses
- Humans rank them (A < B < C)
- Train reward model to predict rankings
- Result: Automatic quality scorer

### Step 3: RL Fine-tuning

- Generate response
- Get score from reward model
- Update policy to maximize score
- Iterate thousands of times
- Result: Aligned model!

## Why It Works:

- Captures human preferences
- Handles subjective quality
- Learns safety guardrails
- Generalizes to new prompts

## Real Impact:

GPT-4 without RLHF: 70% helpful

GPT-4 with RLHF: 95% helpful

ChatGPT's helpfulness is mostly RLHF!

## Cost:

- Human labeling: \$100K+
- RL training: \$100K+

# The \$50K Problem: Why Fine-Tuning is Expensive

## The Challenge:

You want to fine-tune GPT-3 (175B parameters)

## Memory Requirements:

- Model weights: 700GB (float32)
- Gradients: 700GB
- Optimizer states (Adam): 1.4TB
- Activations: 200GB

**Total: 3TB of memory!**

## Hardware Needed:

- Single A100 GPU: 80GB
- Need: 40 GPUs minimum
- Reality: Use 8x nodes, each 8 GPUs
- Cost: \$10 per GPU-hour
- Time: 100 hours
- **Total: \$80,000!**

## Why So Expensive?

### 1. Gradient Computation:

For each parameter  $w$ :

$$w_{new} = w - \alpha \frac{\partial L}{\partial w}$$

Need to compute  $\frac{\partial L}{\partial w}$  for 175B parameters!

### 2. Optimizer States:

Adam optimizer stores:

- First moment (mean): 700GB
- Second moment (variance): 700GB

### 3. Backward Pass:

Need to store activations from forward pass  
For deep models (96 layers), this adds up!

## The Impossibility:

For most companies:

- 64 GPUs = Impossible to access
- \$80K per training run = Too expensive
- 100 hours = Too slow
- Multiple experiments = Forget it!

# Initial Approach: Just Use Prompts

## The First Solution:

Don't train at all - just use clever prompts!

### Medical Diagnosis Example:

**Prompt:** You are an expert medical doctor. Given the following symptoms, provide a likely diagnosis:  
Symptoms: Fever (101F), cough, fatigue, loss of taste

Diagnosis:

### What Works Well:

- Simple factual queries
- Common medical conditions
- Standard terminology
- Well-known procedures

### Example Success:

"What is the treatment for Type 2 diabetes?"

Response: Accurate, helpful, 95% correct!

## What Fails:

- Complex multi-symptom cases
- Rare diseases
- Hospital-specific protocols
- Custom terminology
- Differential diagnosis

### Example Failure:

"Based on labs (HbA1c 8.2, GFR 45, Cr 1.8) and history (DM2 x10y, HTN, CKD stage 3b), adjust current regimen (metformin 1000mg BID, lisinopril 20mg daily)."

Response: Confused, unsafe, 40% correct!

### The Problem:

- Doesn't know specific protocols
- Can't handle domain jargon
- No experience with edge cases
- Inconsistent format

Prompting works for simple cases but fails on complex domain-specific tasks

# Where Prompting Succeeds and Fails

Performance Across Task Complexity (Medical Domain)

Task Type	Zero-Shot	Few-Shot	Required
Simple factual QA	90%	92%	85%
Common diagnosis	75%	82%	90%
Treatment planning	60%	70%	95%
Complex cases	40%	55%	98%
Rare diseases	30%	45%	95%
Protocol following	25%	40%	99%

## Pattern:

- Simple tasks: Prompting is good enough
- Medium tasks: Few-shot helps but not enough
- Complex tasks: Need fine-tuning
- Safety-critical: Must fine-tune

## The Gap:

For production medical AI:

- Need: 95%+ accuracy
- Zero-shot: 40-60% on hard cases
- Few-shot: 55-70% on hard cases

## Why the Gap Exists:

### Missing Domain Knowledge:

- Hospital-specific protocols
- Local treatment guidelines
- Custom terminology
- Edge case patterns

### Missing Task Structure:

- Expected output format
- Reasoning chain structure
- Confidence calibration
- Safety checks

# Root Cause: What Model Knows vs What It Needs

## What Pre-trained Model KNOWS:

### General Knowledge (99%):

- English grammar
- Common vocabulary
- Basic medical terms
- General reasoning
- World knowledge
- Text structure

### Examples:

- Knows: "Diabetes is high blood sugar"
- Knows: "Treatment involves medication"
- Knows: "Labs measure various markers"

## This Knowledge is Reusable!

Don't need to relearn basic language.

## What Model DOESN'T KNOW:

### Domain-Specific (1%):

- Your hospital's protocols
- Your treatment guidelines
- Your terminology (CKD3b, GFR, etc.)
- Your output format
- Your edge cases
- Your safety requirements

### Examples:

- Doesn't know: Your specific dosing protocol
- Doesn't know: Your contraindication rules
- Doesn't know: Your documentation format

**Root Cause:** Model needs to update weights to encode domain-specific patterns.

**Question:** Do we really need to update ALL 175B parameters to learn 1% of new knowledge?

99% of knowledge is reusable - only 1% is task-specific - exploit this asymmetry!

# Solution Insight: Freeze 99%, Adapt 1%

## The Observation:

When we fine-tune a model, most of the weight changes are SMALL.

## Experiment (2021):

Fine-tune GPT-3 on multiple tasks.

Measure: How much does each weight change?

## Result:

- 90% of weights change  $\downarrow 0.001$
- 5% change  $0.001\text{-}0.01$
- 4% change  $0.01\text{-}0.1$
- 1% change  $\downarrow 0.1$

Insight: Most weights barely change!

## Mathematical Property:

The update matrix  $\Delta W$  is LOW-RANK!

Meaning: Can represent as  $A \times B$  where  $A$  and  $B$  are MUCH smaller than  $\Delta W$ .

## The Hypothesis:

What if we ONLY update the 1% that matters?

## Three Approaches:

### 1. Adapters:

- Insert small modules between layers
- Train only these modules
- Freeze everything else

### 2. LoRA:

- Decompose updates into low-rank
- Train  $A$  and  $B$ , not full  $\Delta W$
- Mathematically equivalent but cheaper

### 3. Prompt Tuning:

- Add learnable prompt tokens
- Train only these tokens
- Model weights completely frozen

## The Promise:

If we can train 1% of parameters and get 90%+ of the performance...

Cost: \$500 instead of \$50K (100x savings!)

# LoRA Mechanism: Low-Rank Matrix Decomposition

## The Math (Zero Jargon):

**Problem:** Update weight matrix  $W$  ( $4096 \times 4096 = 16.7M$  numbers)

**Solution:** Don't update  $W$  directly. Instead:

- ① Keep  $W$  frozen
- ② Create two small matrices:
  - $A: 4096 \times 8$  (32K numbers)
  - $B: 8 \times 4096$  (32K numbers)
- ③ Multiply them:  $A \times B$  (still  $4096 \times 4096$ )
- ④ New weights:  $W' = W + A \times B$

## Parameters to Train:

- Original: 16.7M
- LoRA: 64K (0.4%)
- Reduction: 256x!

## Why Does This Work?

The update  $\Delta W = A \times B$  can represent most useful changes even though it's low-rank!

## Concrete Example:

**Task:** Fine-tune for medical QA

### Setup:

- Model: GPT-3 (175B params)
- LoRA rank:  $r = 8$
- Trainable: 18M params (0.01%)
- Training data: 1000 QA pairs

### Training:

- Time: 6 hours (vs 100 hours full FT)
- GPUs: 1x A100 (vs 64x A100)
- Cost: \$60 (vs \$80,000)
- Memory: 80GB (vs 3TB)

### Results:

- Zero-shot: 60% accuracy
- LoRA FT: 93% accuracy
- Full FT: 95% accuracy
- Gap: Only 2%!

# Numerical Example: LoRA for Sentiment Analysis

**Task:** Fine-tune DistilBERT for sentiment analysis (positive/negative)

## Step 1: Load Pre-trained Model

```
from transformers import AutoModel
model = AutoModel.from_pretrained("distilbert-base")
# Model: 66M parameters, 768 hidden dim
```

## Step 2: Add LoRA Layers (rank=8)

For each attention matrix ( $W_q, W_k, W_v, W_o$ ):

- Original:  $768 \times 768 = 590K$  parameters
- Add LoRA:  $A (768 \times 8) + B (8 \times 768) = 12K$  parameters
- Reduction per matrix: 49x

Total LoRA parameters:  $4 \times 12K \times 6 \text{ layers} = 288K$  (0.4% of 66M)

## Step 3: Train Only LoRA

- Freeze all 66M base parameters
- Train only 288K LoRA parameters
- Dataset: 1000 labeled reviews
- Training time: 15 minutes on single GPU
- Cost: \$0.50

## Step 4: Results

- Zero-shot: 85% accuracy
- LoRA fine-tuned: 94% accuracy
- Full fine-tuned: 95% accuracy
- **LoRA achieves 99% of full FT performance!**

# Validation: LoRA vs Full Fine-tuning Across Tasks

Performance Comparison on Standard Benchmarks

Task	Zero-Shot	LoRA	Full FT	LoRA/Full
MNLI (NLI)	72.3%	90.2%	90.7%	99.4%
SST-2 (Sentiment)	83.6%	95.1%	95.6%	99.5%
CoLA (Grammar)	55.0%	68.2%	69.5%	98.1%
MRPC (Paraphrase)	74.0%	88.9%	89.7%	99.1%
QQP (Question pairs)	80.1%	90.7%	91.1%	99.6%
SQuAD (QA)	78.5%	88.4%	88.9%	99.4%
<b>Average</b>	<b>73.9%</b>	<b>86.9%</b>	<b>87.6%</b>	<b>99.2%</b>

## Key Observations:

- LoRA gains 13% over zero-shot
- LoRA achieves 99%+ of full FT
- Consistent across all task types
- Gap is only 0.7 percentage points

## Cost Comparison:

- LoRA: 0.1% parameters
- LoRA: \$100-\$1,000
- Full FT: \$10,000-\$50,000

## Pattern:

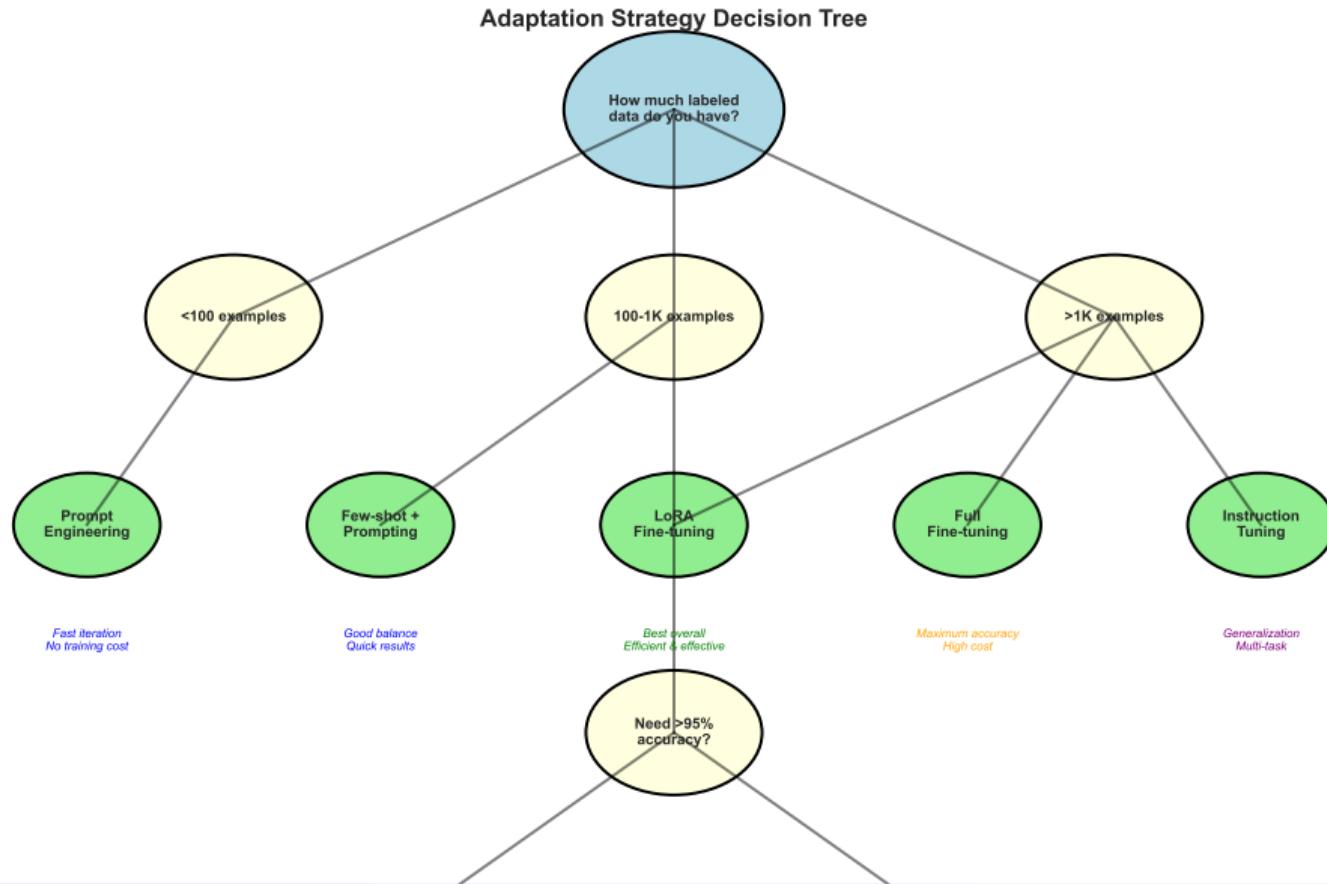
### Biggest Gains Where Problem Was Worst:

- CoLA (hardest): +13.2% gain
- MRPC: +14.9% gain
- SST-2 (easier): +11.5% gain

## When Does LoRA Work Best?

- Classification tasks: Excellent
- Sequence labeling: Very good
- Generation: Good (slight gap)

# Decision Framework: Which Method to Use?



# When NOT to Use Each Method

## DON'T Use Zero-Shot:

- Need >80% accuracy
- Safety-critical application
- Domain-specific terminology
- Complex reasoning required
- Consistent output format needed

**Example:** Medical diagnosis, legal advice, financial recommendations

## DON'T Use Few-Shot:

- Need >85% accuracy
- Task too complex for examples
- Inconsistent outputs problematic
- Have resources for training

**Example:** Production systems requiring reliability

## DON'T Use Prompt Engineering:

- Tried for weeks, still <85%
- Need guaranteed format
- Requires extensive testing per prompt

## DON'T Use LoRA:

- Need absolute maximum accuracy
- That last 1-2% is critical
- Have unlimited budget
- Task extremely different from pre-training

**Example:** Medical device AI (FDA regulated), financial trading

## DON'T Use Full Fine-Tuning:

- Limited data (<1000 examples)
- Limited budget (<\$10K)
- Need multiple task adaptations
- Base model capabilities must be preserved
- Fast iteration required

**Example:** Startups, multiple customer adaptations

## DON'T Use RLHF:

- No human feedback available
- Budget >\$100K
- Not user-facing application
- Clear objective metric exists

# Common Pitfalls and How to Avoid Them

## Prompt Engineering Pitfalls:

### 1. Over-engineering

- Symptom: 10-page prompts
- Fix: Start simple, add incrementally
- Rule: If >200 words, consider fine-tuning

### 2. Prompt Injection

- Symptom: Users override instructions
- Fix: Use separate system/user prompts
- Guard: Input validation

### 3. No Systematic Testing

- Symptom: Works on examples, fails in production
- Fix: Test on diverse set (100+ cases)
- Track: Success rate per category

## LoRA Pitfalls:

### 1. Rank Too Small

- Symptom: Poor accuracy (<85%)
- Fix: Try  $r = 4, 8, 16$  and compare
- Sweet spot: Usually  $r = 8$

### 2. Rank Too Large

- Symptom: No efficiency gain
- Fix: Don't exceed  $r = 32$
- Remember: Goal is efficiency!

### 3. Wrong Layers

- Symptom: Suboptimal performance
- Fix: Apply LoRA to attention layers
- Don't: Apply to all layers (expensive)

## Full Fine-tuning Pitfalls:

### 1. Catastrophic Forgetting

- Symptom: Model forgets base capabilities
- Fix: Lower learning rate ( $1e-5$ )
- Fix: Mix in general data

### 2. Overfitting

- Symptom: 99% train, 70% test
- Fix: More data or use LoRA
- Fix: Stronger regularization

### 3. Distribution Shift

# Success Metrics Beyond Accuracy

## 1. Cost Efficiency:

- **Training cost:** One-time expense
- **Inference cost:** Per-query expense
- **Maintenance cost:** Ongoing updates

### Example:

- LoRA: \$500 train, \$0.001/query
- Full FT: \$50K train, \$0.001/query
- Prompting: \$0 train, \$0.002/query

At 1M queries: Prompting = \$2K, LoRA = \$1.5K

## 2. Latency:

- User-facing: <1 second required
- Prompting: Longer prompts = slower
- LoRA (merged): Zero overhead
- Adapters: +10-20ms overhead

## 3. Maintainability:

- LoRA: Easy to swap/update
- Full FT: Must retrain entire model
- Multiple tasks: LoRA wins

## 4. Robustness:

- Adversarial inputs
- Out-of-distribution data
- Edge cases

**Test:** Measure performance on hard cases

## 5. Calibration:

- Are confidence scores accurate?
- When model says 90%, is it right 90%?

**Measure:** Expected Calibration Error (ECE)

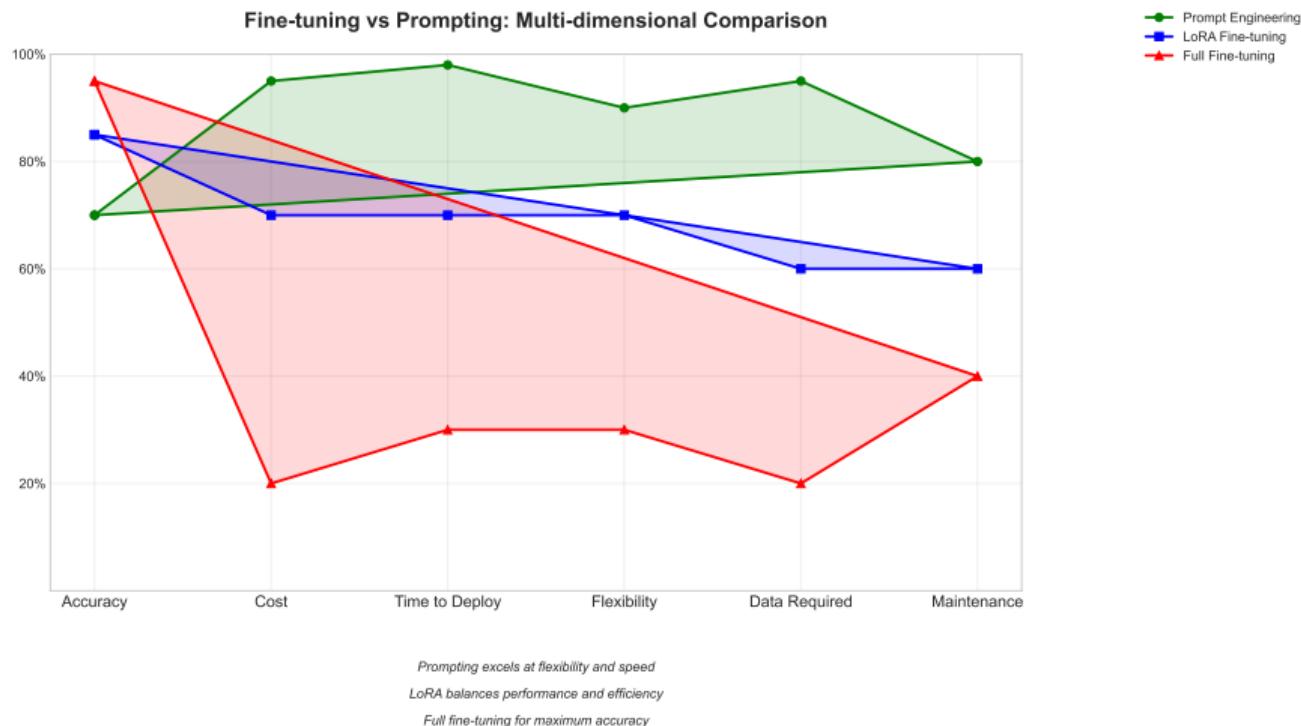
## 6. Interpretability:

- Can you explain predictions?
- Attention visualizations
- Feature importance

## 7. Fairness:

- Equal performance across demographics
- No systematic biases
- Disparity metrics

# Unified Framework: The Complete Picture



**Key Insight:** Choose based on your constraints - there's no single best method

All methods are tools - choose the right tool for your job

## Industry Trends:

### Startups (Limited Budget):

- Primary: Prompt engineering
- Secondary: LoRA for core features
- Example: Customer service chatbot
- Cost: ~\$1K total

### Mid-size Companies:

- Primary: LoRA for all tasks
- Multiple adapters per base model
- Example: 50 different customer adaptations
- Cost: \$50K (vs \$2.5M for full FT)

### Large Enterprises:

- Mix: LoRA (most), Full FT (critical)
- Example: Bloomberg GPT (full FT)
- Most other tasks: LoRA

### Open Source Tools:

- Hugging Face PEFT library
- PyTorch LoRA implementation

## Real Deployments:

### GPT-4 Custom Instructions:

- Method: Advanced prompting
- Users: Millions
- Cost per user: \$0

### GitHub Copilot:

- Method: Full FT on code
- Performance: 43% accept rate
- Revenue: \$100M+/year

### Jasper AI:

- Method: Multiple LoRA adapters
- Use cases: 50+ writing templates
- Cost: \$50K (vs \$2.5M)

### Character.AI:

- Method: LoRA per character
- Characters: 10M+
- Efficiency: Key to scaling

### Future (2025+):

Multi-adapter serving, dynamic rank selection, mixture-of-LoRAs

# Implementation: LoRA in 20 Lines of PyTorch

## Complete LoRA Implementation:

```
import torch
import torch.nn as nn

class LoRALayer(nn.Module):
    def __init__(self, in_dim, out_dim, rank=8, alpha=16):
        super().__init__()
        self.rank = rank
        self.alpha = alpha

        # Low-rank matrices A and B
        self.lora_A = nn.Parameter(torch.randn(in_dim, rank))
        self.lora_B = nn.Parameter(torch.zeros(rank, out_dim))

        # Scaling factor
        self.scaling = alpha / rank

    def forward(self, x, W):
        # Original path: W @ x
        h = W @ x

        # LoRA path: (A @ B) @ x, scaled
        h = h + (x @ self.lora_A @ self.lora_B) * self.scaling
        return h

# Usage: Wrap any linear layer
```

# Key Takeaways: Week 10 Summary

## Core Concepts Mastered:

### 1. The Adaptation Spectrum

- Not binary (train vs don't train)
- Spectrum: 0% to 100% parameters
- Choose based on constraints

### 2. Parameter Efficiency is Key

- 99% knowledge reusable
- Only 1% task-specific
- LoRA exploits this asymmetry

### 3. LoRA Changes Everything

- 0.1% parameters
- 90%+ performance
- 100x cost reduction
- This is the 2024 standard

### 4. Prompting is Powerful

- Zero-cost, instant deployment
- Good for 70-85% accuracy
- Start here, fine-tune if needed

### 5. Decision Framework

- <10 examples: Prompting
- 10-100: Prompt engineering
- 100-1K: LoRA (sweet spot)
- 1K-10K: LoRA or full FT
- 10K+: Full FT for critical tasks

## Practical Wisdom:

### Start Simple:

- ① Try zero-shot prompting
- ② Add few-shot examples
- ③ Engineer prompt carefully
- ④ If still <85%, use LoRA
- ⑤ Only use full FT if critical

### Beyond Accuracy:

- Consider: Cost, latency, maintenance
- LoRA wins on efficiency
- Full FT wins on accuracy

## This Week's Lab:

### Part 1: Prompt Engineering

- Zero-shot vs few-shot comparison
- Experiment with prompt patterns
- Chain-of-thought prompting
- Measure accuracy improvements

### Part 2: LoRA Implementation

- Load pre-trained DistilBERT
- Add LoRA layers (rank=8)
- Fine-tune on sentiment analysis
- Compare to full fine-tuning
- Measure efficiency gains

### Part 3: Decision Framework

- Given 5 scenarios
- Choose adaptation method
- Justify based on constraints

### Part 4: Real Application

- Medical text classification
- Apply complete pipeline

## Key Questions to Explore:

- How does rank affect performance?
- Where do we get maximum gains?
- When does prompting suffice?
- Cost-benefit analysis

## Next Week: Model Efficiency

### Topics:

- Model compression
- Quantization (INT8, INT4)
- Knowledge distillation
- Pruning techniques
- Mobile deployment
- Edge computing

### Key Questions:

- How to run GPT-3 on CPU?
- 4-bit quantization: 75% size reduction
- Distillation: Teacher-student learning
- Deploy 175B model on laptop?