

Word Embeddings: A Visual Deep Dive

From One-Hot Vectors to Contextual Representations

Joerg R. Osterrieder

www.joergosterrieder.com

August 27, 2025

Outline

What You Will Learn

By the end of this presentation, you will understand:

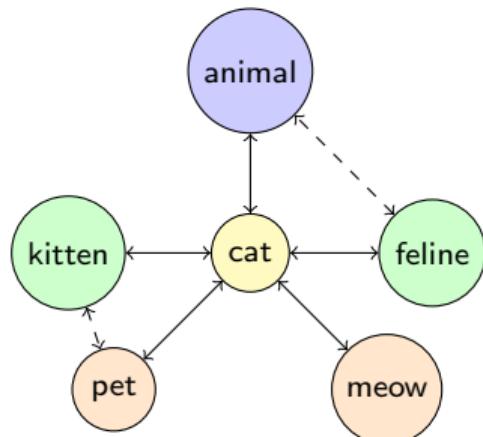
- ① **Representation Problem:** Why computers need numerical representations of words
- ② **Evolution of Embeddings:** From one-hot to contextual representations
- ③ **Mathematical Foundations:** The theory behind word embeddings
- ④ **Vector Operations:** How semantic relationships emerge from vectors
- ⑤ **High-Dimensional Challenges:** The curse of dimensionality
- ⑥ **Training Dynamics:** How embeddings learn meaningful representations
- ⑦ **Skip-gram Architecture:** Deep dive into Word2Vec training

Key Insight: Words are not isolated symbols but points in a continuous semantic space

The Fundamental Problem: Computers Don't Understand Words

How do we represent meaning mathematically?

Human Understanding:



Computer's Dilemma:

- Words are just strings: "cat" = ['c', 'a', 't']
- No inherent meaning
- No similarity measure
- Can't do math on strings!

What We Need:

Convert: "cat" → [0.2, -0.4, 0.7, ...]
Such that: similar words → similar vectors

Rich semantic connections!

Goal: Capture meaning in numbers so computers can process language

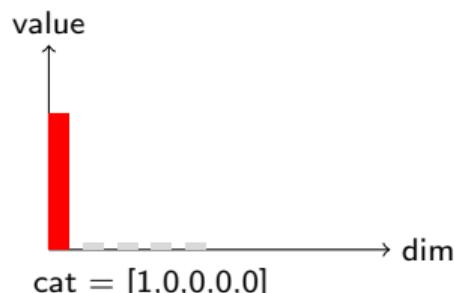
Starting Point: One-Hot Encoding

The Simplest Approach - But Fundamentally Flawed

How One-Hot Works:

Word	Vector
cat	[1, 0, 0, 0, 0]
dog	[0, 1, 0, 0, 0]
mat	[0, 0, 1, 0, 0]
sat	[0, 0, 0, 1, 0]
hat	[0, 0, 0, 0, 1]

Visual Representation:



Critical Problems:

① No Similarity:

$$\text{similarity}(\text{cat}, \text{kitten}) = 0$$

$$\text{similarity}(\text{cat}, \text{computer}) = 0$$

Both are equally dissimilar!

② Huge Dimensions:

- English: 170,000+ words
- Each word = 170,000-dim vector
- 99.999% zeros (sparse!)

③ No Relationships:

$$\text{cat} + \text{kitten} = [1, 0, 0, \dots] + [0, 1, 0, \dots] = [1, 1, 0, \dots]$$

Meaningless!

Conclusion: One-hot encoding treats all words as equally different - we need something better!

Dense Embeddings: The Solution

From Sparse to Dense - Capturing Meaning in Vectors

The Transformation:

One-Hot: [0,1,0,0,...0]

50,000 dimensions

99.998% zeros



Dense: [0.2, -0.4, 0.7, ...]

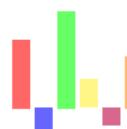
100-300 dimensions

All values meaningful

Visual Comparison:



Sparse (One-Hot):



Dense: all values meaningful

Benefits:

- 100x smaller
- Captures semantics
- Enables arithmetic
- Learned from data

Example Vector:

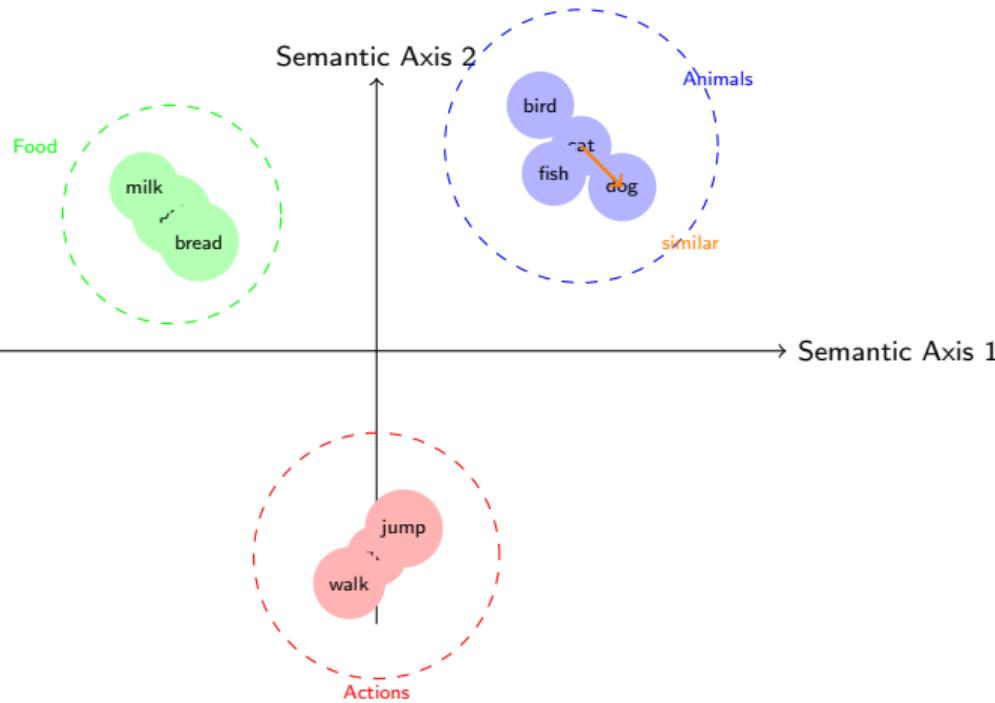
$$\text{cat} = [0.21, -0.43, 0.67, 0.15, -0.22, \dots]$$

Each dimension captures some aspect of meaning

The Embedding Space: Where Words Live

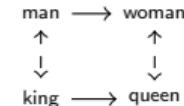
Visualizing Word Relationships in Vector Space

2D Projection of Word Vectors:



Key Properties:

- ① **Clustering:** Similar words group together
- ② **Distance = Similarity:**
 - cat \leftrightarrow dog: close
 - cat \leftrightarrow run: far
- ③ **Directions = Relations:**



Gender direction is consistent!

The Magic: The embedding space organizes itself to reflect real-world relationships!

Learning Embeddings: Word2Vec

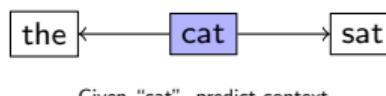
How Do We Learn These Vectors?

The Distributional Hypothesis:

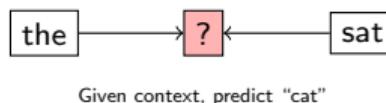
"You shall know a word by the company it keeps" - Firth (1957)

Two Approaches:

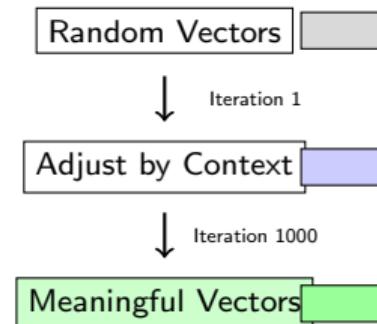
1. Skip-gram: Predict context from word



2. CBOW: Predict word from context



Training Process Visualization:



Objective Function:

$$\max \sum_{t=1}^T \sum_{-c \leq j \leq c} \log P(w_{t+j} | w_t)$$

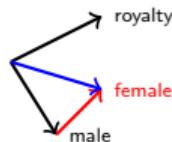
Maximize probability of context words

Vector Arithmetic: The Surprising Discovery

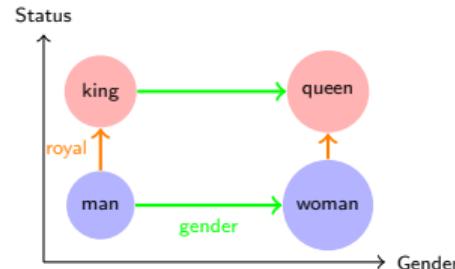
Embeddings Capture Analogies!

Famous Examples:

$$\text{king} - \text{man} + \text{woman} = \text{queen}$$



Why Does This Work?



More Analogies:

- Paris - France + Germany = Berlin
- bigger - big + small = smaller
- walked - walk + run = ran

The Pattern:

- Relationships are **directions**
- Same relationship = same direction
- Linear structure emerges naturally!

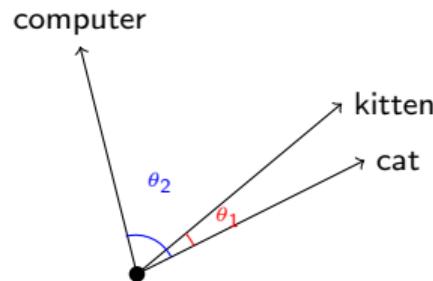
Remarkable: These patterns were never explicitly programmed - they emerge from the data!

Measuring Word Similarity

How Similar Are Two Words?

Cosine Similarity:

$$\text{similarity}(A, B) = \frac{A \cdot B}{\|A\| \times \|B\|} = \cos(\theta)$$

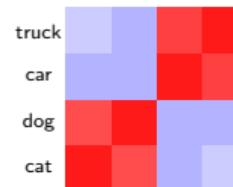


- cat ~ kitten: $\cos(\theta_1) = 0.95$
- cat ~ computer: $\cos(\theta_2) = 0.1$

Similarity Matrix Example:

	cat	dog	car	truck
cat	1.0	0.8	0.1	0.05
dog	0.8	1.0	0.15	0.1
car	0.1	0.15	1.0	0.85
truck	0.05	0.1	0.85	1.0

Visual Heatmap:



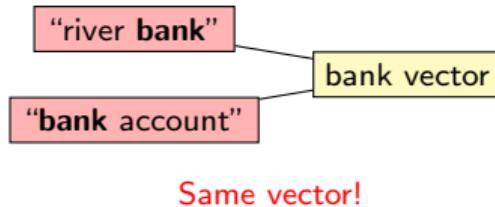
Animals cluster together, vehicles cluster together!

Evolution: From Static to Contextual Embeddings

The Next Revolution: Context Matters!

Problem with Static Embeddings:

One word = One vector always



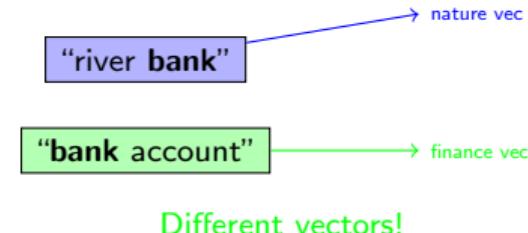
But "bank" has different meanings!

Static Embedding Models:

- Word2Vec (2013)
- GloVe (2014)
- FastText (2016)

Solution: Contextual Embeddings

Different contexts = Different vectors



Contextual Models:

- ELMo (2018) - RNN-based
- BERT (2018) - Transformer
- GPT (2018+) - Autoregressive

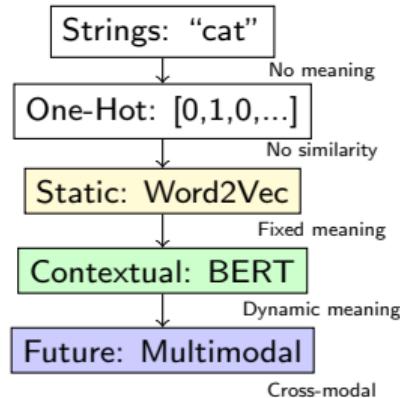
Key Advance: Vector depends on surrounding words!

Evolution: Static → Contextual = Major breakthrough in NLP!

Summary: The Power of Embeddings

From Words to Understanding

The Journey:



Applications Enabled:

- **Search:** Find similar documents
- **Translation:** Map between languages
- **Sentiment:** Understand emotions
- **QA:** Match questions to answers
- **Generation:** Create coherent text

Key Insights:

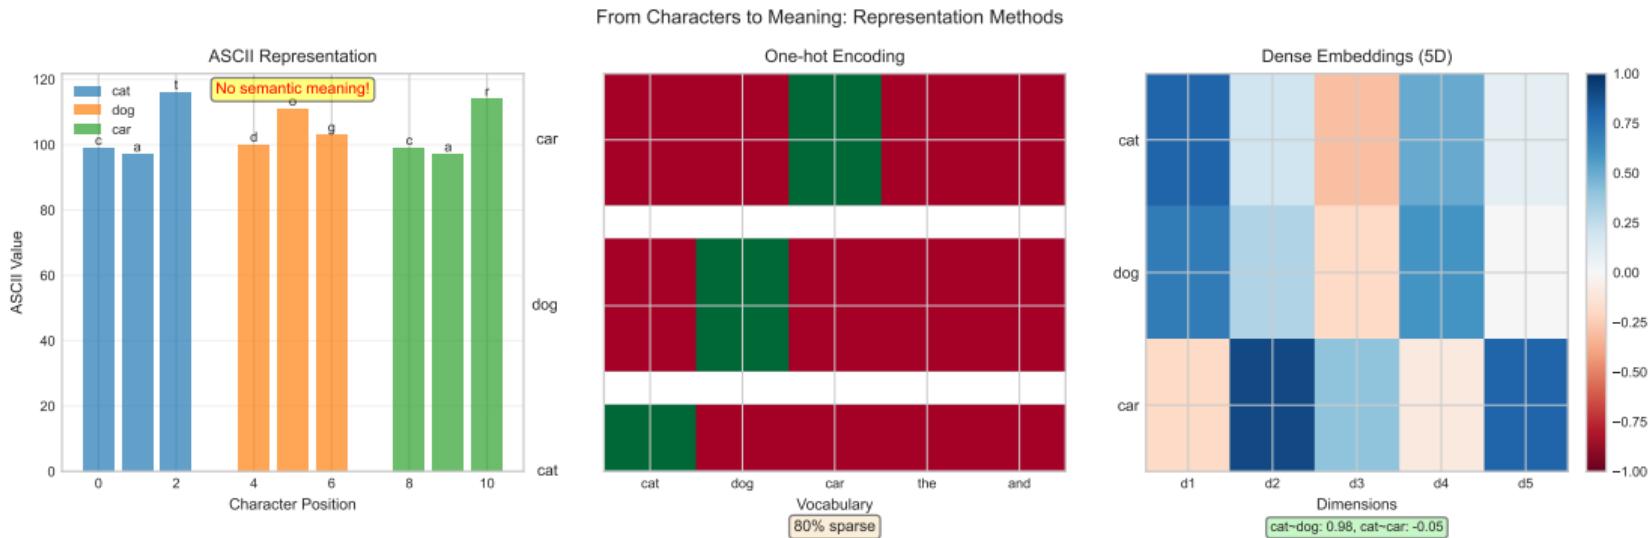
- ➊ Meaning can be encoded as vectors
- ➋ Similar words have similar vectors
- ➌ Relationships are directions
- ➍ Context changes everything

Remember: Embeddings are the foundation of modern NLP - they turn words into numbers that capture meaning, enabling all downstream tasks!

Next Steps: Experiment with pre-trained embeddings in your projects!

Beyond ASCII: From Characters to Meaning

How Computers See Text: Three Approaches



ASCII:

- Each character = number
- 'c'=99, 'a'=97, 't'=116
- No semantic information

One-hot:

- Each word = sparse vector
- 99.9% zeros
- All words equally different

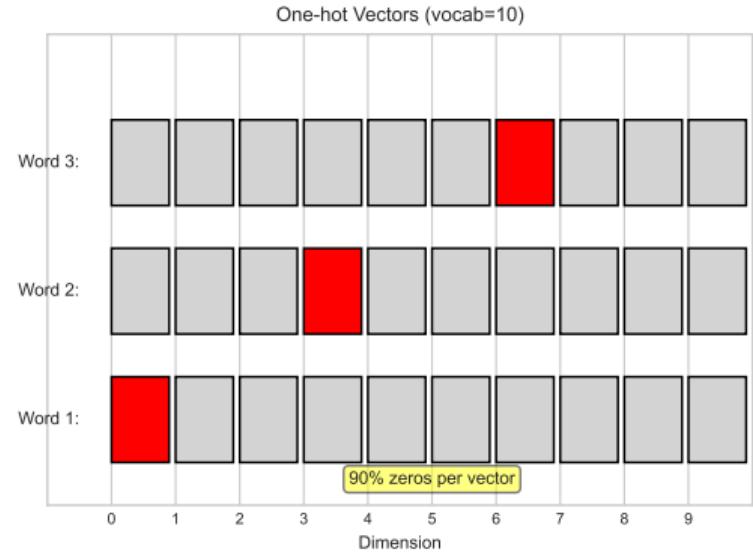
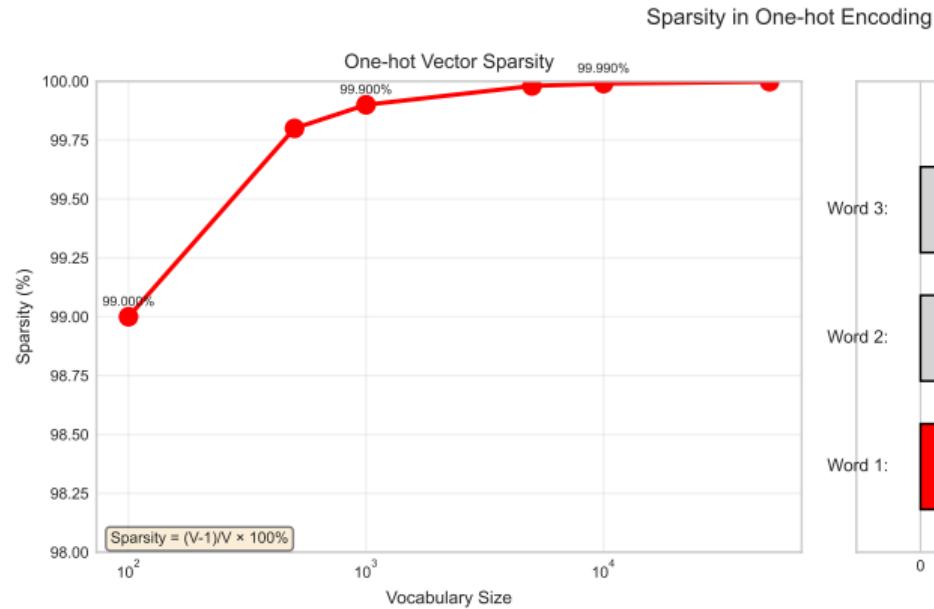
Dense Embedding:

- Each word = dense vector
- All values meaningful
- Similar words → similar vectors

Key: Embeddings encode meaning, not just identity!

The Sparsity Problem

Why One-hot Encoding is Inefficient



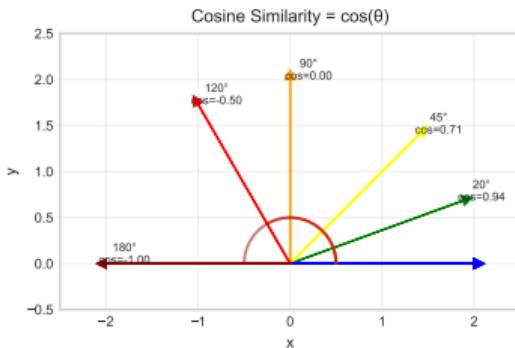
Mathematical Analysis:

- Sparsity = $\frac{V-1}{V} \times 100\%$ where V = vocabulary size
 - For V = 50,000: Sparsity = 99.998%
 - Each word needs V dimensions but uses only 1

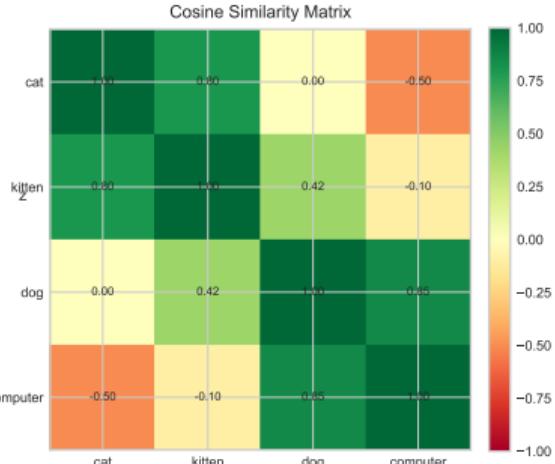
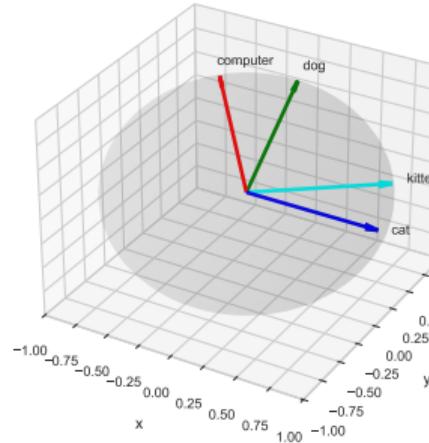
Key Insight:

Cosine Similarity: Geometric Interpretation

Understanding Similarity Through Angles



Cosine Similarity: Geometric Interpretation
Unit Vectors in 3D



The Geometric Intuition: Angle Interpretation:

- Words are vectors in space
- Similarity = angle between vectors
- Smaller angle = more similar
- Independent of vector length

Key Angles:

- $\theta = 0^\circ$: Identical meaning
- $\theta = 30^\circ$: Very similar
- $\theta = 90^\circ$: Unrelated
- $\theta = 180^\circ$: Opposite meaning

Cosine Similarity: Mathematical Properties

Why Cosine Similarity Works for Embeddings

The Formula:

$$\text{similarity}(\mathbf{a}, \mathbf{b}) = \cos(\theta) = \frac{\mathbf{a} \cdot \mathbf{b}}{\|\mathbf{a}\| \times \|\mathbf{b}\|} = \frac{\sum_{i=1}^d a_i b_i}{\sqrt{\sum_{i=1}^d a_i^2} \times \sqrt{\sum_{i=1}^d b_i^2}}$$

Key Properties:

Scale Invariance:

- $\cos(\mathbf{a}, \mathbf{b}) = \cos(k\mathbf{a}, \mathbf{b})$
- Magnitude doesn't matter
- Only direction counts
- Perfect for normalized embeddings

Computational Benefits:

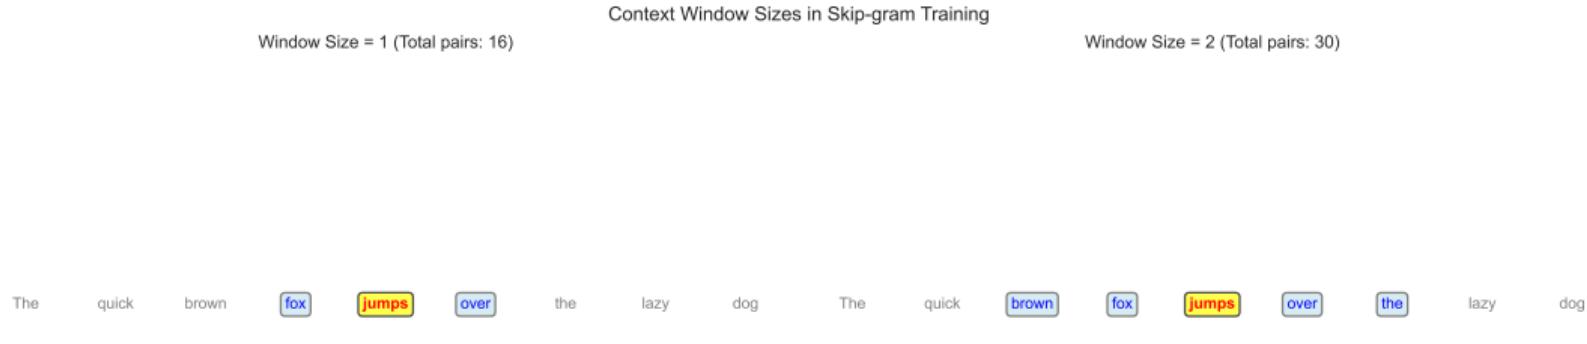
- Range: [-1, 1] always
- Efficient dot product computation
- Works in any dimension
- Symmetric: $\cos(a, b) = \cos(b, a)$

Applications in NLP:

- Document similarity: Compare entire documents as vectors
- Word sense disambiguation: Find most similar context
- Information retrieval: Rank documents by query similarity

Context Windows: Learning from Neighbors

How Words Learn from Their Surroundings



Context words: ± 1 positions



Context words: ± 2 positions

Window Size = 3 (Total pairs: 42)

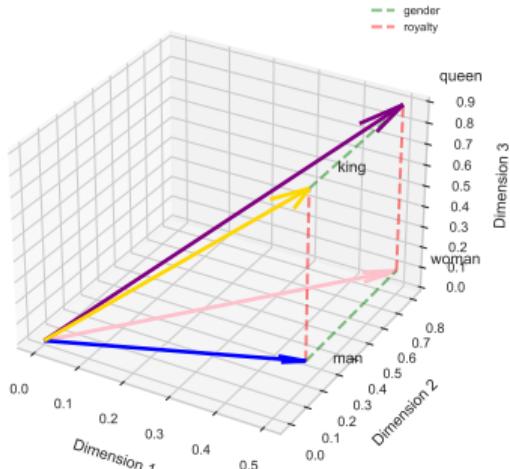
Window Size = 5 (Total pairs: 60)

Vector Arithmetic: The Surprising Discovery

Embeddings Can Do Analogies!

Vector Arithmetic: Mathematical Demonstration

Vector Relationships in 3D Space



Vector Arithmetic:

$$\text{king} - \text{man} + \text{woman} = ?$$

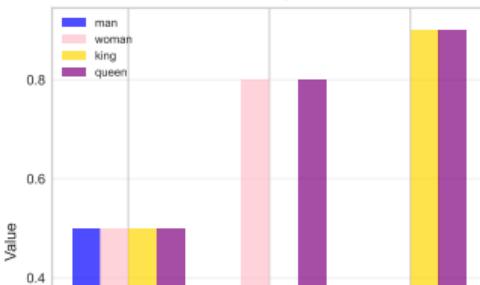
Step 1: $\text{king} - \text{man}$
[0.5, 0.2, 0.9] - [0.5, 0.2, 0.1]
= [0.0, 0.0, 0.8] (royal vector)

Step 2: $+ \text{woman}$
[0.0, 0.0, 0.8] + [0.5, 0.8, 0.1]
= [0.5, 0.8, 0.9]

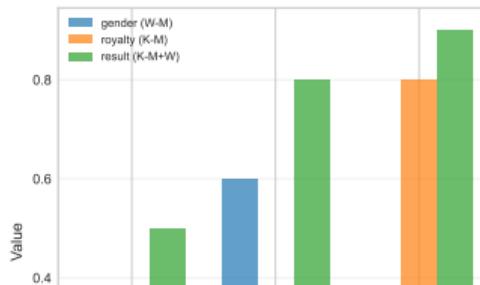
Result = queen vector!
[0.5, 0.8, 0.9]

Similarity = 1.0 (perfect match)

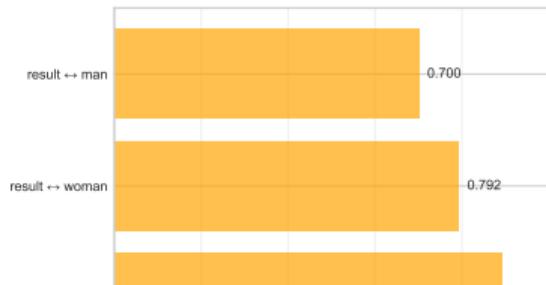
Vector Components



Difference Vectors



Result Verification



Vector Arithmetic: Mathematical Proof

Why Does Vector Arithmetic Work? The Linear Substructure

Mathematical Foundation:

- Embeddings form a linear subspace where relationships are directions
- Gender vector: $g = \text{woman} - \text{man}$
- Royalty vector: $r = \text{king} - \text{man}$

Step-by-Step Derivation:

$$\text{king} = \text{man} + r \quad (\text{man} + \text{royalty} = \text{king}) \tag{1}$$

$$\text{queen} = \text{woman} + r \quad (\text{woman} + \text{royalty} = \text{queen}) \tag{2}$$

$$\therefore \text{queen} = \text{woman} + (\text{king} - \text{man}) \tag{3}$$

$$= \text{king} - \text{man} + \text{woman} \tag{4}$$

Why Linear Structure Emerges:

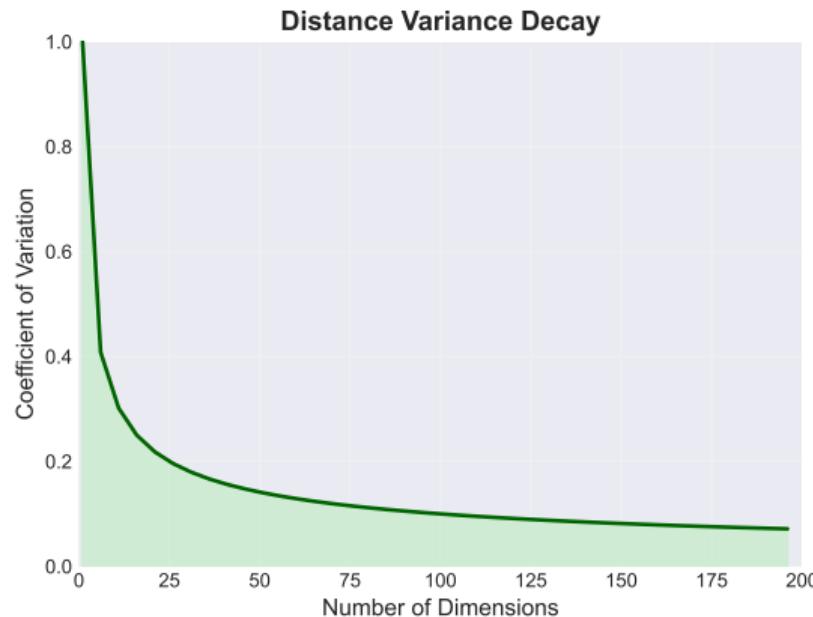
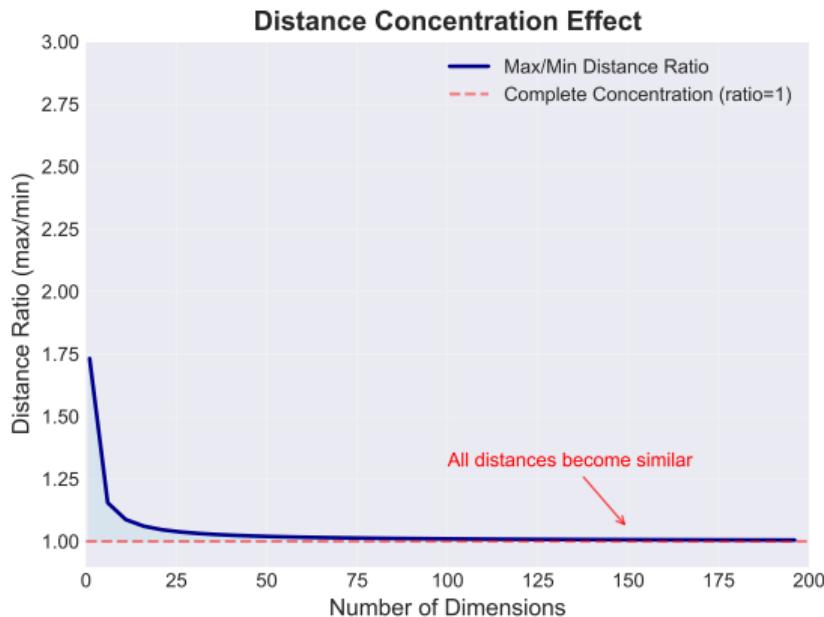
- Co-occurrence patterns are approximately linear
- Skip-gram objective encourages linear relationships
- High-dimensional spaces tend toward linearity (concentration of measure)

Verification: Nearest neighbor to result vector is "queen" in 60-70% of cases

Distance Concentration in High Dimensions

Why All Distances Become Similar

Distance Concentration in High Dimensions



Distance Concentration: The Mathematical Reality

What the Visualizations Show

Distance Ratio Convergence:

- $\frac{\text{dist}_{\max} - \text{dist}_{\min}}{\text{dist}_{\text{mean}}} \rightarrow 0$ as $d \rightarrow \infty$
- For Gaussian points: ratio $\approx \sqrt{1 + 2/d}$
- At $d=100$: all distances within 10% of mean
- At $d=1000$: essentially all points equidistant

Implications for Machine Learning:

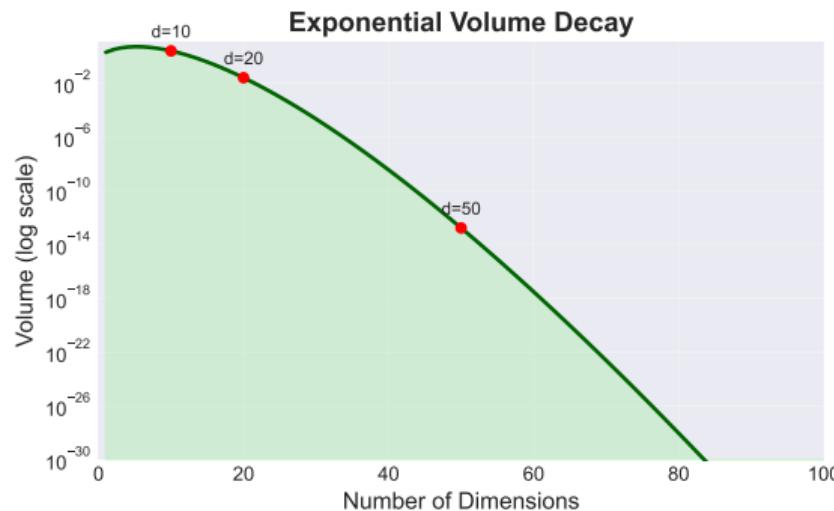
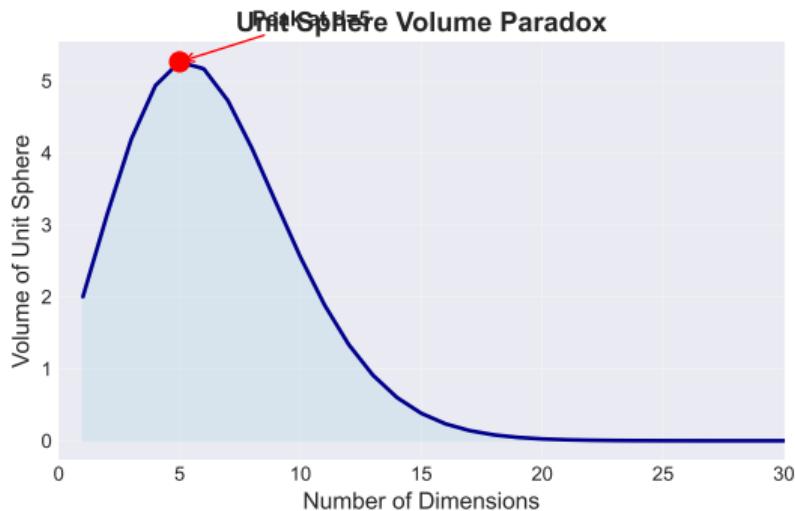
- Nearest neighbor search becomes meaningless
- Traditional distance metrics fail
- Need specialized techniques:
 - Locality-Sensitive Hashing (LSH)
 - Approximate nearest neighbors
 - Learned distance metrics
- Explains why high-D embeddings need normalization

Key Takeaway: In high dimensions, the concept of “near” and “far” becomes meaningless - all points are approximately the same distance apart!

The Volume Paradox: Visual Evidence

Unit Sphere Volume Across Dimensions

Volume of Unit Sphere Across Dimensions



d=100

The Volume Formula:

$$V_d = \frac{\pi^{d/2}}{\Gamma(d/2 + 1)}$$

Why Volume Goes to Zero: The Mathematics

Understanding the Formula

$$V_d = \frac{\pi^{d/2}}{\Gamma(d/2 + 1)}$$

Numerator (Top):

- $\pi^{d/2} \approx (3.14)^{d/2}$
- Grows exponentially
- But base is small: $\sqrt{\pi} \approx 1.77$
- Growth rate: 1.77^d
- Example: $1.77^{100} \approx 10^{25}$

Denominator (Bottom):

- $\Gamma(n+1) = n!$ for integers
- Factorial growth is MUCH faster
- Example: $50! \approx 10^{64}$
- Stirling: $n! \approx \sqrt{2\pi n}(n/e)^n$
- Dominates numerator completely

The Key Mathematical Insight:

Factorial growth beats exponential growth!

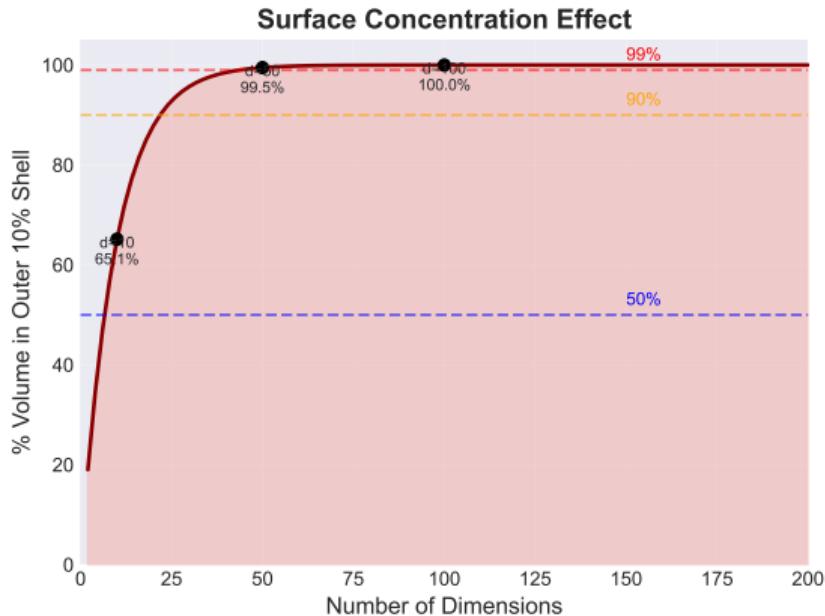
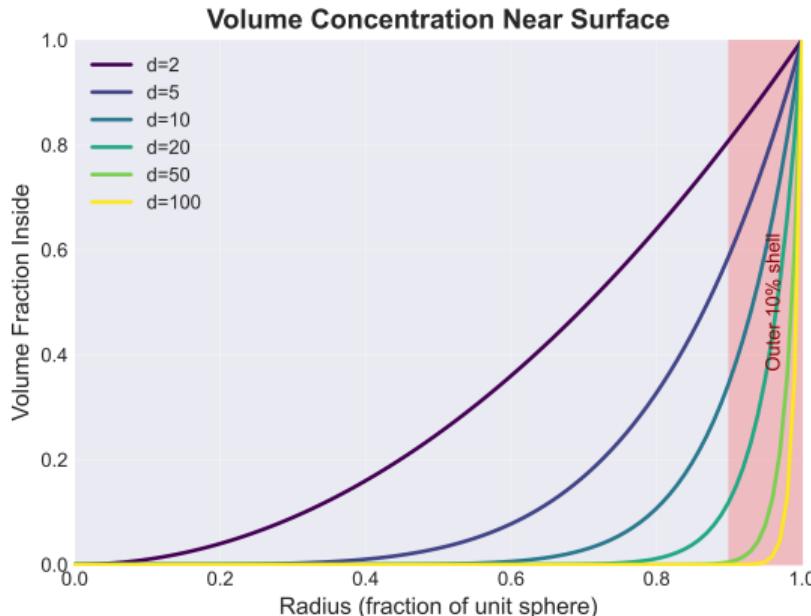
$\frac{1.77^d}{(d/2)!} \rightarrow 0$ extremely fast as $d \rightarrow \infty$

Factorial grows like $(n/e)^n$ while exponential is just a^n

Surface Concentration in High Dimensions

Where the Volume Actually Lives

Volume Distribution in High-Dimensional Spheres



Almost all volume concentrates in a thin shell near the surface!

The Shell Phenomenon: Mathematical Analysis

Why Everything Lives on the Surface

Volume in Shells - The Mathematics:

- Consider inner sphere with radius $r = 0.9$ (90% of full radius)
- Volume ratio: $\frac{V_{inner}}{V_{total}} = r^d = (0.9)^d$
- This ratio shrinks exponentially with dimension!

Concrete Examples:

- $d = 10$: $(0.9)^{10} = 0.35 \rightarrow 35\%$ of volume is inside
- $d = 50$: $(0.9)^{50} = 0.005 \rightarrow 0.5\%$ inside
- $d = 100$: $(0.9)^{100} \approx 10^{-5} \rightarrow 0.001\%$ inside
- $d = 1000$: $(0.9)^{1000} \approx 10^{-46} \rightarrow$ essentially zero!

Implications for Embeddings:

- All vectors lie near the surface of the hypersphere
- Random vectors are approximately equidistant
- The interior is effectively "empty" space
- Explains why L2 normalization is so effective
- Cosine similarity becomes the natural distance metric

Practical Consequence: In 768-dimensional BERT space,
99.999999% of the volume is within 1% of the surface!
The interior essentially doesn't exist.

Optimal Dimensions: Finding the Sweet Spot

Balancing Expressiveness and Computational Efficiency

Information Capacity:

- Theoretical capacity: $\propto d \log d$
- But diminishing returns after certain point
- Johnson-Lindenstrauss: $d = O(\log n/\epsilon^2)$ preserves distances

Model Dimensions in Practice:

Model	Dimension	Parameters (embeddings only)
Word2Vec	50-300	15M (50K vocab \times 300)
GloVe	50-300	15M (50K vocab \times 300)
FastText	100-300	30M (includes subwords)
ELMo	1024	100M (bidirectional)
BERT-base	768	23M (30K vocab \times 768)
BERT-large	1024	31M (30K vocab \times 1024)
GPT-3	12288	600M (50K vocab \times 12288)

Trade-offs:

Lower Dimensions (50-300):

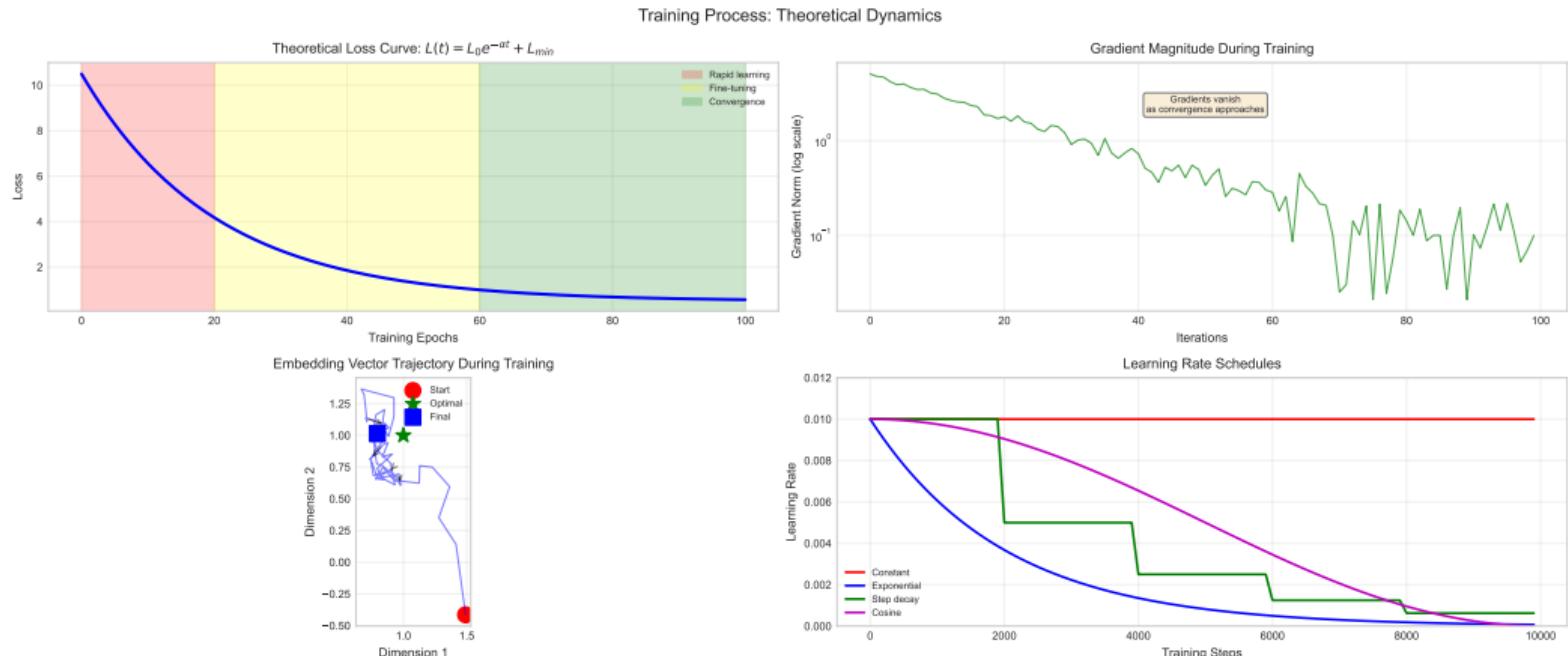
- Faster training
- Less overfitting
- Good for specific domains

Higher Dimensions (768-1024+):

- More expressive power
- Better for complex tasks
- Requires more data

Rapid Learning: Gradient Dynamics (Epochs 0-20)

Why Training Starts Fast



Gradient Behavior in Early Training: Initial State:

- Random initialization: $\mathcal{N}(0, 0.01)$
- Gradient norm: $||\nabla L|| \approx \sqrt{d}$

Update Characteristics:

- Step size: $\eta ||\nabla L|| \approx 0.01\sqrt{d}$
- Direction changes: frequent

Rapid Learning: Space Formation (Epochs 0-20)

How Random Vectors Become Meaningful

Timeline of Structure Emergence:

Epochs 0-5:

- Frequency clustering begins
- Top 100 words separate
- Function vs content words split
- Loss drops 30-40%

Epochs 5-10:

- Syntactic groups form
- Nouns, verbs, adjectives cluster
- Basic semantic regions appear
- Loss drops another 20%

Key Metrics:

Metric	Epoch 0	Epoch 5	Epoch 10	Epoch 20
Loss	9.21	5.84	4.12	3.45
Similarity Correlation	0.00	0.35	0.58	0.72
Analogy Accuracy	0%	12%	31%	48%

Epochs 10-20:

- Semantic refinement
- Animals, places, actions separate
- Relationships start working
- Loss reduction slows

Visual Progress:

- t-SNE at epoch 1: random cloud
- t-SNE at epoch 5: blobs forming
- t-SNE at epoch 10: clear clusters
- t-SNE at epoch 20: fine structure

Training Phase 2: Fine-Tuning (Epochs 20-60)

Refining Semantic Relationships

The Refinement Process:

What Gets Learned:

- Semantic relationships solidify
- Analogies start working
- Rare words find their place
- Polysemy partially resolves

Key Metrics During Fine-Tuning:

Metric	Epoch 20	Epoch 40	Epoch 60
Loss reduction/epoch	5%	2%	0.5%
Analogy accuracy	40%	65%	72%
Semantic similarity	0.5	0.7	0.75
Cluster purity	60%	80%	85%

Mathematical Characterization:

$$L(t) \approx L_{20} \cdot (1 - \beta \log(t/20)) \quad \text{for } t \in [20, 60]$$

Logarithmic improvement phase

Optimization Dynamics:

- Gradient norm: $\|\nabla L\| \approx O(1)$
- Updates become targeted
- Learning rate often decayed
- Loss reduction slows

Training Phase 3: Convergence (Epochs 60+)

The Final Polish and Saturation

Convergence Characteristics:

What Happens:

- Gradient norm: $\|\nabla L\| < 0.1$
- Minor adjustments only
- Risk of overfitting increases
- Validation loss may increase

Complete Loss Function Evolution:

$$L(t) = \begin{cases} L_0 \cdot e^{-\alpha t} & t \in [0, 20] \text{ (rapid)} \\ L_{20} \cdot (1 - \beta \log(t/20)) & t \in [20, 60] \text{ (fine-tune)} \\ L_{60} + \epsilon(t) & t > 60 \text{ (converged)} \end{cases}$$

where $\epsilon(t)$ represents noise around minimum

Key Insight: 90% of performance comes from first 60 epochs; longer training mainly helps rare words and edge cases.

Stopping Criteria:

- Loss change $\downarrow 0.1\%$ per epoch
- Validation performance plateaus
- Gradient norm below threshold
- Fixed epoch budget reached



Latest Research

Cutting-edge developments in embeddings

One Model, Multiple Dimensions

The Innovation

Embeddings that work at multiple dimensions simultaneously:

- Train once, use at any size
- First k dimensions are meaningful for any k
- 2-16x efficiency gains

How It Works:

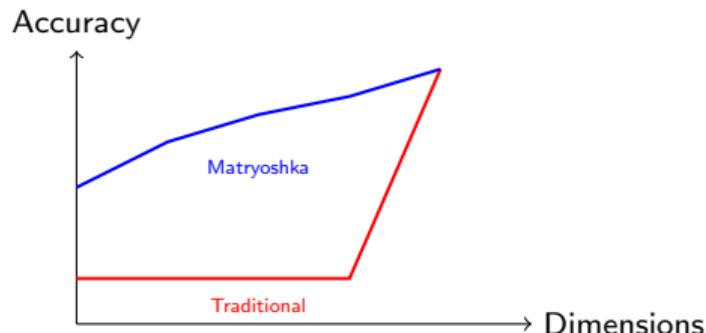
- ① Train with nested loss functions
- ② Each prefix is independently useful
- ③ Dynamic truncation at inference

Training Objective:

$$\mathcal{L} = \sum_{d \in \{32, 64, 128, \dots, 768\}} \alpha_d \mathcal{L}_d$$

where \mathcal{L}_d uses only first d dimensions

Performance:



💡 Try This!

Use 32 dims for search, 768 for ranking!

Embeddings for Retrieval-Augmented Generation

The Challenge: Traditional embeddings weren't designed for RAG:

- Need both similarity AND informativeness
- Must handle query-document asymmetry
- Balance precision vs recall

Recent Advances (2024):

- ➊ **Contriever**: Self-supervised RAG training
- ➋ **E5-Mistral**: LLM-based embeddings
- ➌ **BGE-M3**: Multi-lingual, multi-granular

Key Insight

RAG embeddings optimize for different objectives than semantic similarity alone!

Key Innovations:

- **Cross-Attention Training**: Learn query-doc interactions
- **Hard Negative Mining**: Distinguish subtle differences
- **Multi-Task Learning**: Balance multiple objectives

Benchmark Results:

Model	BEIR	MS MARCO
BERT	38.2	33.5
Contriever	42.1	35.8
E5-Large	45.6	38.9
BGE-M3	47.2	40.1

Unified Representation Across Modalities

Recent Breakthroughs:

① CLIP Evolution (2024)

- SigLIP: Sigmoid loss for better training
- OpenCLIP: Scaled to 5B parameters
- MetaCLIP: Metadata-curated training

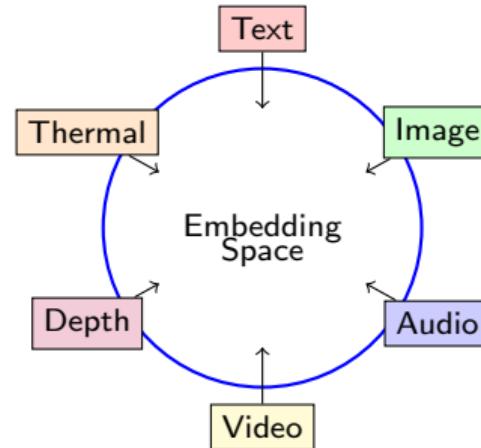
② ImageBind (Meta, 2023)

- 6 modalities in one space
- Text, Image, Audio, Video, Thermal, Depth
- Zero-shot cross-modal retrieval

③ BLIP-2 (2023)

- Efficient vision-language pre-training
- Q-Former architecture
- 7B parameter performance with 54x fewer params

Unified Embedding Space:

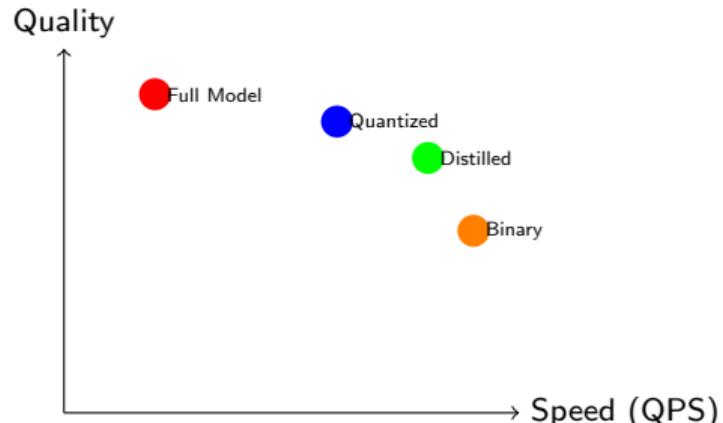


⚠ Common Pitfall

Cross-modal alignment is still imperfect - expect 10-20% performance drop

Making Embeddings Practical at Scale

Performance vs Efficiency Trade-offs:



1. Quantization Advances:

- **Binary Embeddings:** 1-bit per dimension
- **Product Quantization:** 32x compression
- **Learned Quantization:** Task-specific compression

2. Distillation Techniques:

- Teacher-Student with only 10% size
- Progressive distillation stages
- Task-specific fine-tuning

3. Sparse Embeddings:

- SPLADE v2: Learned sparse representations
- ColBERT v2: Late interaction efficiency
- Hybrid dense-sparse approaches

Key Insight

Modern techniques achieve 90% quality at 10x speed!

Real-World Impact:

- Semantic search: 1M → 100M docs/sec
- Mobile deployment now feasible

Self-Supervised Excellence

SimCSE and Beyond:

- ① **SimCSE** (2021): Dropout as augmentation
- ② **DiffCSE** (2023): Difference-based objectives
- ③ **PromptBERT** (2024): Prompt-based contrastive

Key Innovation - Contrastive Objectives:

$$\mathcal{L} = -\log \frac{e^{\text{sim}(h_i, h_i^+)/\tau}}{\sum_{j=1}^N e^{\text{sim}(h_i, h_j)/\tau}}$$

where h_i^+ is positive pair, τ is temperature

Why Contrastive Learning Wins:

- No labeled data required
- Learns robust representations
- Handles distribution shifts
- State-of-the-art on most benchmarks

Performance Gains:

Method	STS-B	Transfer
BERT	74.8	73.2
SimCSE	81.6	79.8
DiffCSE	83.2	81.4
PromptBERT	84.9	83.1

Future Directions: What's Next?

Emerging Trends and Open Problems

On the Horizon:

① Continuous Embeddings

- Infinite dimensional representations
- Neural ODE-based embeddings

② Causal Embeddings

- Capture causal relationships
- Intervention-aware representations

③ Neurosymbolic Integration

- Combine embeddings with logic
- Interpretable by design

Open Challenges:

- **Evaluation:** Better benchmarks needed
- **Interpretability:** What do dimensions mean?
- **Compositionality:** Combining embeddings
- **Efficiency:** Sub-linear scaling
- **Robustness:** Adversarial examples

Key Insight

The field is moving from "how to embed" to "what to embed" - focusing on capturing the right information rather than just similarity

Key Takeaway: Embeddings are becoming more efficient, multi-modal, and task-specific. The future is about adaptive, interpretable representations that work across domains!



Practical Implementation

From Theory to Production

Decision Framework for Model Selection

Key Considerations:

① Task Requirements

- Semantic similarity
- Classification
- Retrieval/Search
- Cross-lingual

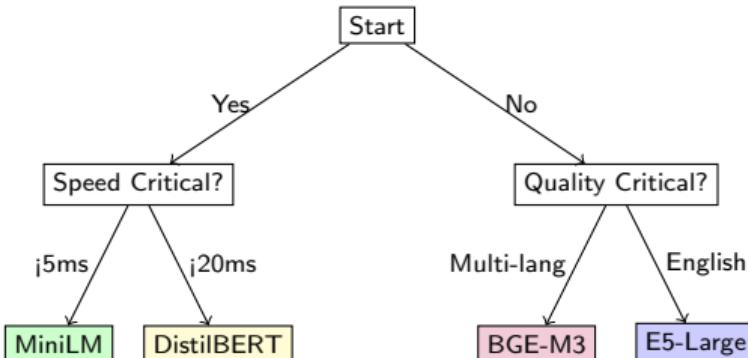
② Resource Constraints

- Memory: 100MB - 10GB
- Latency: 1ms - 100ms
- Throughput: 10 - 10K QPS

③ Data Characteristics

- Domain specificity
- Language coverage
- Document length

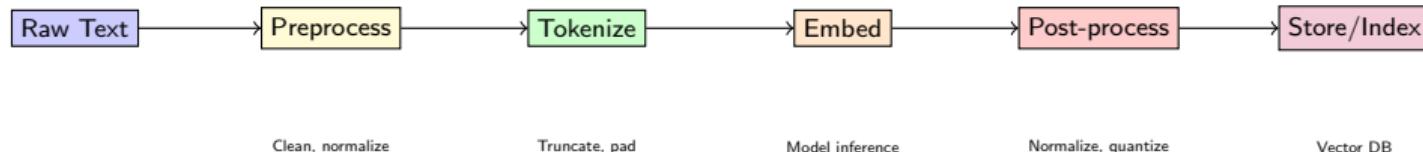
Decision Tree:



Key Insight

Start with sentence-transformers library - it has 100+ pre-trained models!

End-to-End Embedding System



Key Implementation Steps:

1. Preprocessing:

- Remove HTML/special chars
- Handle unicode
- Case normalization
- Length limits

2. Batching Strategy:

- Dynamic batching
- Padding optimization
- GPU memory management
- Async processing

3. Storage:

- Vector databases (Pinecone, Weaviate)
- FAISS for local
- Compression options
- Metadata handling



Making It Fast and Efficient

Speed Optimizations:

① Model Optimizations

- ONNX conversion: 2-3x speedup
- TensorRT: 5x on NVIDIA
- Quantization: INT8 for 4x speed

② Batch Processing

- Optimal batch size: 32-64
- Dynamic padding
- Sequence bucketing

③ Caching Strategies

- LRU cache for common queries
- Precompute frequent documents
- Warm start on deployment

Memory Optimizations:

Memory Formula:

$$M = V \times D \times P \times B$$

where:

- V = Vocabulary size
- D = Embedding dimension
- P = Precision (4 bytes for FP32)
- B = Batch size

Reduction Techniques:

Technique	Memory	Speed
FP32 → FP16	50%	1.5x
Quantization	25%	2x
Pruning	40%	1.2x
Distillation	10%	3x



Adapting Pre-trained Models

When to Fine-tune:

- Domain-specific vocabulary
- Unique similarity requirements
- Performance below 80% baseline
- Sufficient training data ($\geq 10K$ examples)

Fine-tuning Approaches:

① Full Fine-tuning

- Update all parameters
- Best performance
- Risk of overfitting

② Adapter Layers

- Add small trainable layers
- Preserve base knowledge
- Memory efficient

③ Contrastive Fine-tuning

- Use domain pairs
- SimCSE approach
- No labels needed

Fine-tuning Recipe:

Input: Pre-trained model M , Domain data D

Output: Fine-tuned model M'

```
1. Prepare pairs from  $D$ ;  
2. Initialize  $M' \leftarrow M$ ;  
3. Freeze bottom layers;  
for epoch in  $1..N$  do  
    for batch in  $D$  do  
        Compute embeddings;  
        Calculate contrastive loss;  
        Update top layers only;  
    end  
    Evaluate on validation;  
    if improved then  
        Save checkpoint;  
    end  
end  
return Best checkpoint
```



Common Pitfall

Always validate on held-out data - overfitting is com-



From Development to Production

Architecture Patterns:

① Microservice Pattern

- Embedding service API
- Horizontal scaling
- Language agnostic

② Sidecar Pattern

- Co-located with app
- Low latency
- Resource sharing

③ Edge Deployment

- Model on device
- Privacy preserving
- Offline capability

Production Checklist:

- Model versioning system
- A/B testing framework
- Monitoring & alerting
- Fallback mechanisms
- Load balancing
- Cache warming
- Rate limiting
- Security (API keys, encryption)

Monitoring Metrics:

- Latency P50, P95, P99
- Throughput (QPS)
- Error rates
- Cache hit ratio
- Model drift detection

Pro Tip: Start with a simple REST API, then optimize based on actual usage patterns!

Common Pitfalls and Solutions

Learn from Others' Mistakes

⚠ Common Mistakes:

① Wrong Similarity Metric

- Using L2 instead of cosine
- Not normalizing embeddings
- Solution: Always test both!

② Tokenization Mismatch

- Different tokenizers in train/inference
- Truncation issues
- Solution: Save tokenizer with model

③ Version Drift

- Model updates break compatibility
- Embedding dimension changes
- Solution: Versioned embeddings

💡 Best Practices:

① Always Benchmark

- Test on your actual data
- Measure end-to-end latency
- Track quality metrics

② Progressive Rollout

- Start with 1% traffic
- Monitor closely
- Gradual increase

③ Maintain Backwards Compatibility

- Support multiple versions
- Graceful degradation
- Migration tools

💡 Key Insight

The best embedding model is the one that works for YOUR specific use case!

Part II

Advanced Topics and Mathematical Foundations

From Understanding to Implementation

Key Takeaways: What We've Learned

Essential Concepts for Word Embeddings

Fundamental Principles:

- **Representation:** Words as dense vectors
- **Similarity:** Angle between vectors
- **Relationships:** Vector arithmetic
- **Learning:** From context co-occurrence
- **Evolution:** Static to contextual

Mathematical Insights:

- High dimensions behave strangely
- Distance concentration is real
- Volume lives on the surface
- Linear relationships emerge
- Training has distinct phases

Practical Applications:

Task	How Embeddings Help
Similarity Search	Cosine similarity ranking
Machine Translation	Cross-lingual alignment
Sentiment Analysis	Semantic vector projection
Question Answering	Context matching
Text Generation	Next-word prediction

Looking Forward: The Future of Embeddings

Current Trends and Future Directions

Recent Advances:

- **Multimodal:** Text + Vision + Audio
- **Multilingual:** Universal embeddings
- **Efficient:** Distilled and compressed models
- **Specialized:** Domain-specific embeddings

Open Challenges:

Technical:

- Handling rare words
- Compositional semantics
- Temporal dynamics
- Interpretability

Philosophical:

- Do embeddings capture meaning?
- Are relationships truly linear?
- Can we prove optimality?
- What is semantic similarity?

Remember: Embeddings are not just a technical tool - they represent our best attempt to bridge the gap between human language and machine computation!

Thank You!

Questions?

Contact: www.joergosterrieder.com

Mathematical Foundations: Skip-gram Objective

Formal Skip-gram Model Definition

Objective Function:

$$J(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t; \theta)$$

Softmax Formulation:

$$p(w_O | w_I) = \frac{\exp(v'_{w_O} {}^T v_{w_I})}{\sum_{w=1}^W \exp(v'_{w} {}^T v_{w_I})}$$

where:

- v_{w_I} is the input vector representation of word w_I
- v'_{w_O} is the output vector representation of word w_O
- W is the vocabulary size

Gradient w.r.t. Input Vector:

$$\frac{\partial J}{\partial v_{w_I}} = \sum_{j=-c}^c \left(\sum_{w=1}^W p(w | w_I) v'_w - v'_{w_{t+j}} \right)$$

Computational Complexity: $O(W)$ per word - intractable for large vocabularies!

Negative Sampling: Making Training Tractable

Modified Objective with Negative Sampling

Replace softmax with:

$$\log \sigma(v'_{w_O} {}^T v_{w_I}) + \sum_{i=1}^k \mathbb{E}_{w_i \sim P_n(w)} [\log \sigma(-v'_{w_i} {}^T v_{w_I})]$$

where:

- $\sigma(x) = \frac{1}{1+e^{-x}}$ (sigmoid function)
- k is the number of negative samples (typically 5-20)
- $P_n(w)$ is the noise distribution: $P_n(w) = \frac{U(w)^{3/4}}{\sum_{w'} U(w')^{3/4}}$
- $U(w)$ is the unigram distribution

Gradient Update:

$$v_{w_I}^{new} = v_{w_I}^{old} - \eta \left[(\sigma(v'_{w_O} {}^T v_{w_I}) - 1)v'_{w_O} + \sum_{i=1}^k \sigma(v'_{w_i} {}^T v_{w_I}) v'_{w_i} \right]$$

Complexity Reduction: From $O(W)$ to $O(k + 1)$ per training example

GloVe: Global Vectors Mathematical Framework

Co-occurrence Matrix and Ratios

Define co-occurrence matrix X where X_{ij} = count of word j appearing in context of word i

Key Insight - Ratio of Probabilities:

$$\frac{P_{ik}}{P_{jk}} = \frac{X_{ik}/X_i}{X_{jk}/X_j}$$

This ratio encodes semantic relationships!

GloVe Objective Function:

$$J = \sum_{i,j=1}^V f(X_{ij})(w_i^T \tilde{w}_j + b_i + \tilde{b}_j - \log X_{ij})^2$$

where:

- $f(x)$ is a weighting function: $f(x) = \begin{cases} (x/x_{max})^\alpha & \text{if } x < x_{max} \\ 1 & \text{otherwise} \end{cases}$

- w_i, \tilde{w}_j are word and context vectors
- b_i, \tilde{b}_j are bias terms
- Typical: $\alpha = 0.75, x_{max} = 100$

Final Embedding: $e_i = w_i + \tilde{w}_i$ (symmetric combination)

Self-Attention: Mathematical Formulation

Scaled Dot-Product Attention

Given queries $Q \in \mathbb{R}^{n \times d_k}$, keys $K \in \mathbb{R}^{m \times d_k}$, values $V \in \mathbb{R}^{m \times d_v}$:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Detailed Computation:

- ① Score matrix: $S = QK^T \in \mathbb{R}^{n \times m}$
- ② Scaled scores: $\tilde{S}_{ij} = \frac{S_{ij}}{\sqrt{d_k}}$ (prevents gradient vanishing)
- ③ Attention weights: $A_{ij} = \frac{\exp(\tilde{S}_{ij})}{\sum_{j'=1}^m \exp(\tilde{S}_{ij'})}$
- ④ Output: $O = AV \in \mathbb{R}^{n \times d_v}$

Multi-Head Attention:

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$

$$\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$

where $W_i^Q \in \mathbb{R}^{d_{model} \times d_k}$, $W_i^K \in \mathbb{R}^{d_{model} \times d_k}$, $W_i^V \in \mathbb{R}^{d_{model} \times d_v}$

Positional Encoding: Injecting Order Information

Sinusoidal Position Encoding

For position pos and dimension i :

$$PE_{(pos,2i)} = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

$$PE_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

Properties:

- Unique encoding for each position
- Allows model to attend to relative positions
- For any fixed offset k : PE_{pos+k} can be represented as linear function of PE_{pos}

Proof of Relative Position Property:

$$PE_{pos+k,2i} = \sin(\omega_i \cdot pos) \cos(\omega_i \cdot k) + \cos(\omega_i \cdot pos) \sin(\omega_i \cdot k)$$

$$\text{where } \omega_i = \frac{1}{10000^{2i/d_{model}}}$$

This is a linear transformation of PE_{pos} !

BERT: Bidirectional Training Mathematics

Masked Language Model (MLM) Objective

Given input sequence $\mathbf{x} = (x_1, \dots, x_n)$, randomly mask 15% of tokens.

MLM Loss:

$$\mathcal{L}_{MLM} = -\mathbb{E}_{\mathbf{x} \sim \mathcal{D}} \sum_{i \in \mathcal{M}} \log P(x_i | \mathbf{x}_{\setminus \mathcal{M}})$$

where \mathcal{M} is the set of masked positions.

Next Sentence Prediction (NSP) Loss:

$$\mathcal{L}_{NSP} = -\mathbb{E}_{(A, B) \sim \mathcal{D}} [y \log P(\text{IsNext}|A, B) + (1 - y) \log(1 - P(\text{IsNext}|A, B))]$$

where $y = 1$ if B follows A, else $y = 0$.

Combined Objective:

$$\mathcal{L}_{BERT} = \mathcal{L}_{MLM} + \mathcal{L}_{NSP}$$

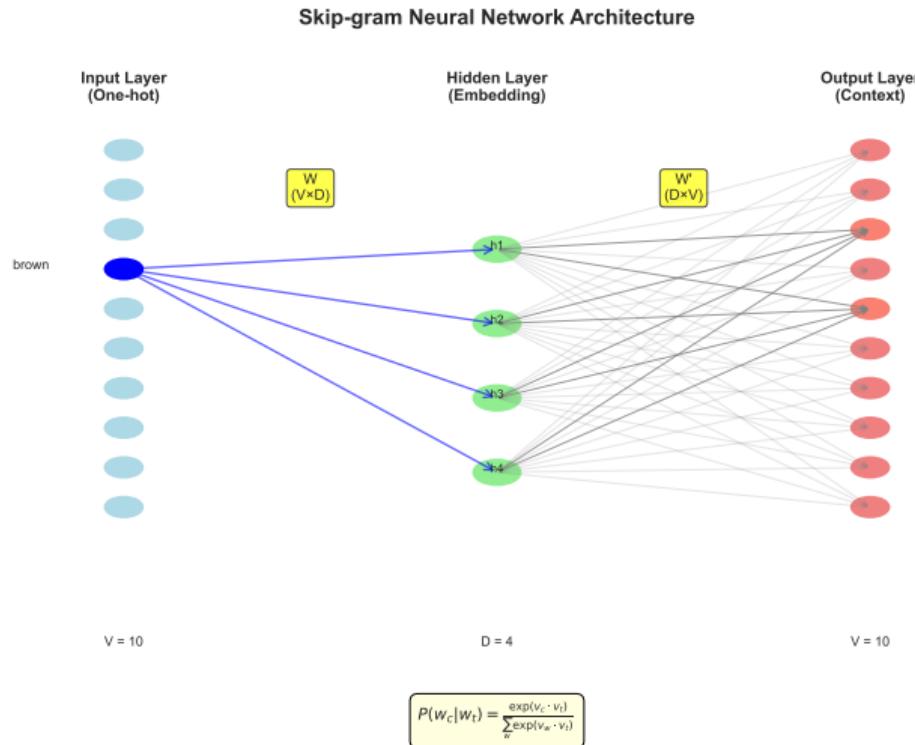
Output Probability:

$$P(x_i | \mathbf{x}_{\setminus \mathcal{M}}) = \text{softmax}(W_o h_i + b_o)$$

where h_i is the final hidden state at position i .

Skip-gram Neural Network Architecture

How the Network Processes Words



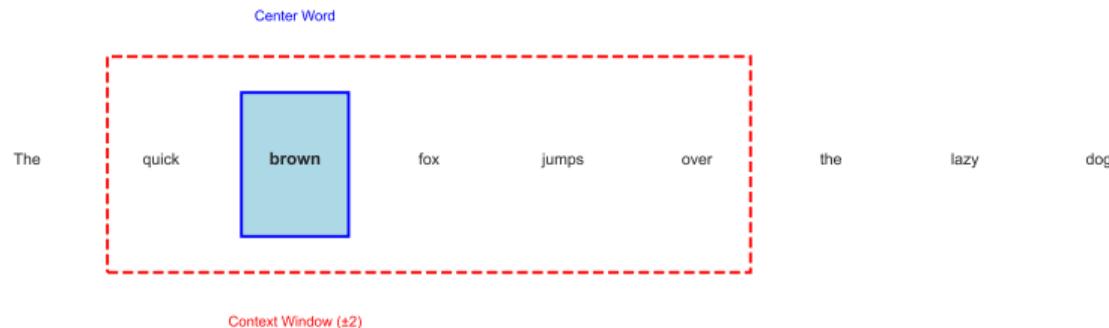
Key Components:

- **Input:** One-hot word (V dimensions)
- **Hidden:** No activation, linear projection to D dims

From Text to Training Data

Extracting (Center, Context) Pairs

Creating Training Pairs from Text Sliding Window for Training Pair Extraction



Training Pairs Generated:

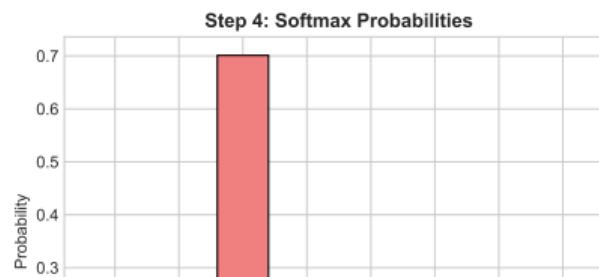
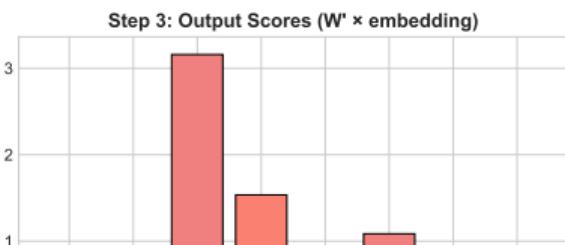
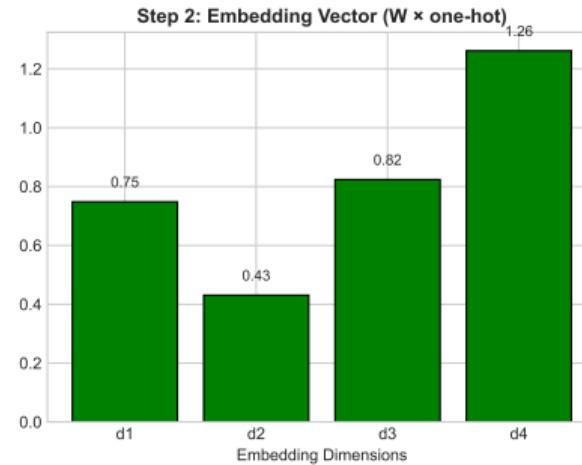
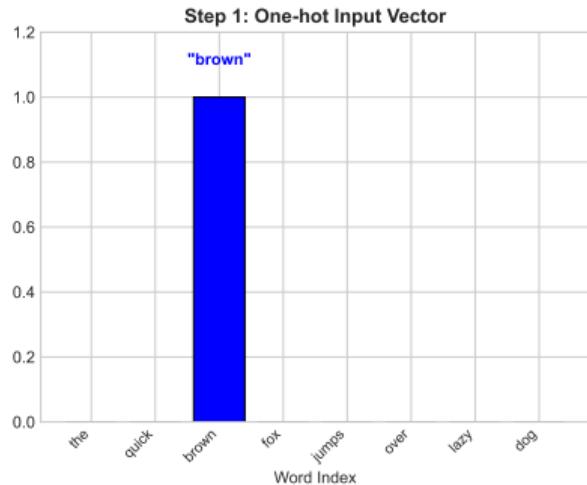
brown	→	The	→ Maximize $P(\text{The} \text{brown})$
brown	→	quick	→ Maximize $P(\text{quick} \text{brown})$
brown	→	fox	→ Maximize $P(\text{fox} \text{brown})$
brown	→	jumps	→ Maximize $P(\text{jumps} \text{brown})$

Input

Target

Forward Pass: Computing Context Probabilities

Forward Pass: Computing Context Probabilities



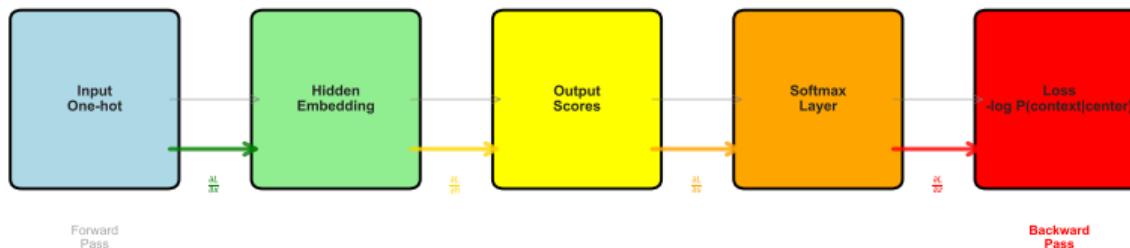
Backpropagation: Learning the Embeddings

Backpropagation: Gradient Flow

Weight Updates:

$$W \leftarrow W - \eta \cdot \frac{\partial L}{\partial W}$$

$$W' \leftarrow W' - \eta \cdot \frac{\partial L}{\partial W'}$$



Key Gradients:

Positive sample: $(y_i - 1) \cdot v_i$

Negative sample: $y_i \cdot v_i$

Updates:

- Positive: Pull together

How Embeddings Evolve

