

# Word Embeddings: A Visual Deep Dive

## From One-Hot Vectors to Contextual Representations

Joerg R. Osterrieder

[www.joergosterrieder.com](http://www.joergosterrieder.com)

August 27, 2025

# Outline

# What You Will Learn

By the end of this presentation, you will understand:

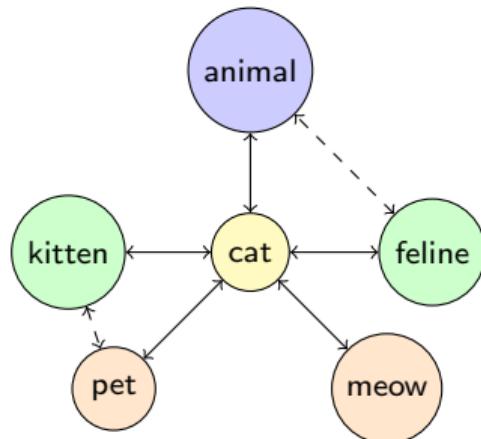
- ① **Representation Problem:** Why computers need numerical representations of words
- ② **Evolution of Embeddings:** From one-hot to contextual representations
- ③ **Mathematical Foundations:** The theory behind word embeddings
- ④ **Vector Operations:** How semantic relationships emerge from vectors
- ⑤ **High-Dimensional Challenges:** The curse of dimensionality
- ⑥ **Training Dynamics:** How embeddings learn meaningful representations
- ⑦ **Skip-gram Architecture:** Deep dive into Word2Vec training

**Key Insight:** Words are not isolated symbols but points in a continuous semantic space

# The Fundamental Problem: Computers Don't Understand Words

How do we represent meaning mathematically?

Human Understanding:



Computer's Dilemma:

- Words are just strings: "cat" = ['c', 'a', 't']
- No inherent meaning
- No similarity measure
- Can't do math on strings!

What We Need:

Convert: "cat"  $\rightarrow$  [0.2, -0.4, 0.7, ...]  
Such that: similar words  $\rightarrow$  similar vectors

Rich semantic connections!

Goal: Capture meaning in numbers so computers can process language

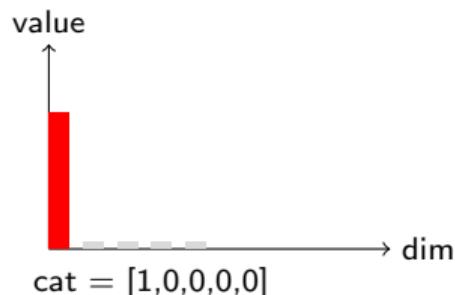
# Starting Point: One-Hot Encoding

## The Simplest Approach - But Fundamentally Flawed

How One-Hot Works:

Word	Vector
cat	[1, 0, 0, 0, 0]
dog	[0, 1, 0, 0, 0]
mat	[0, 0, 1, 0, 0]
sat	[0, 0, 0, 1, 0]
hat	[0, 0, 0, 0, 1]

Visual Representation:



### Critical Problems:

#### ① No Similarity:

$$\text{similarity}(\text{cat}, \text{kitten}) = 0$$

$$\text{similarity}(\text{cat}, \text{computer}) = 0$$

Both are equally dissimilar!

#### ② Huge Dimensions:

- English: 170,000+ words
- Each word = 170,000-dim vector
- 99.999% zeros (sparse!)

#### ③ No Relationships:

$$\text{cat} + \text{kitten} = [1, 0, 0, \dots] + [0, 1, 0, \dots] = [1, 1, 0, \dots]$$

Meaningless!

**Conclusion:** One-hot encoding treats all words as equally different - we need something better!

# Dense Embeddings: The Solution

## From Sparse to Dense - Capturing Meaning in Vectors

The Transformation:

One-Hot: [0,1,0,0,...0]

50,000 dimensions

99.998% zeros



Dense: [0.2, -0.4, 0.7, ...]

100-300 dimensions

All values meaningful

Visual Comparison:



Dense:



Benefits:

- 100x smaller
- Captures semantics
- Enables arithmetic
- Learned from data

Example Vector:

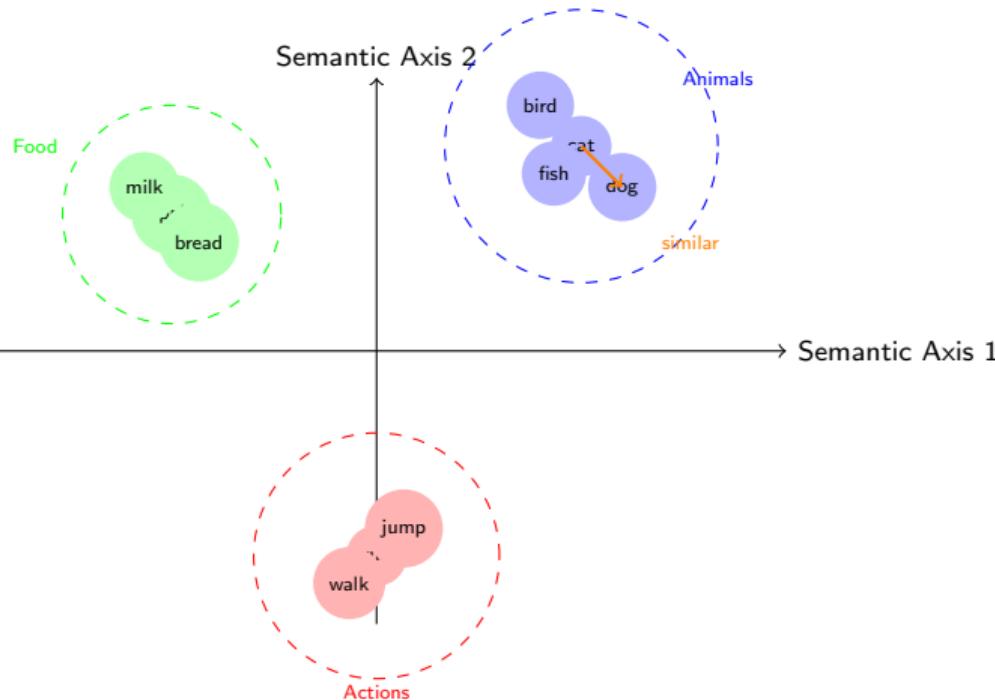
$$\text{cat} = [0.21, -0.43, 0.67, 0.15, -0.22, \dots]$$

Each dimension captures some aspect of meaning

# The Embedding Space: Where Words Live

## Visualizing Word Relationships in Vector Space

2D Projection of Word Vectors:



### Key Properties:

- ① **Clustering:** Similar words group together
- ② **Distance = Similarity:**
  - cat  $\leftrightarrow$  dog: close
  - cat  $\leftrightarrow$  run: far
- ③ **Directions = Relations:**

$$\begin{array}{ccc} \text{man} & \longrightarrow & \text{woman} \\ \uparrow & & \uparrow \\ \downarrow & & \downarrow \\ \text{king} & \longrightarrow & \text{queen} \end{array}$$

Gender direction is consistent!

**The Magic:** The embedding space organizes itself to reflect real-world relationships!

# Learning Embeddings: Word2Vec

## How Do We Learn These Vectors?

### The Distributional Hypothesis:

"You shall know a word by the company it keeps" - Firth (1957)

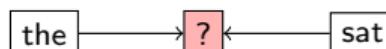
### Two Approaches:

1. Skip-gram: Predict context from word



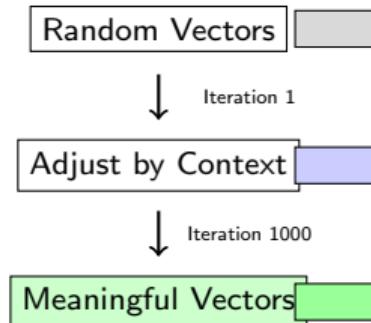
Given "cat", predict context

2. CBOW: Predict word from context



Given context, predict "cat"

### Training Process Visualization:



### Objective Function:

$$\max \sum_{t=1}^T \sum_{-c \leq j \leq c} \log P(w_{t+j} | w_t)$$

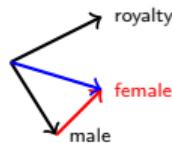
Maximize probability of context words

# Vector Arithmetic: The Surprising Discovery

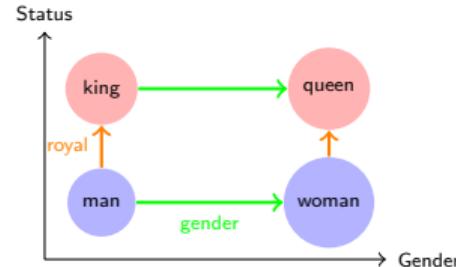
Embeddings Capture Analogies!

Famous Examples:

$$\text{king} - \text{man} + \text{woman} = \text{queen}$$



Why Does This Work?



More Analogies:

- Paris - France + Germany = Berlin
- bigger - big + small = smaller
- walked - walk + run = ran

The Pattern:

- Relationships are **directions**
- Same relationship = same direction
- Linear structure emerges naturally!

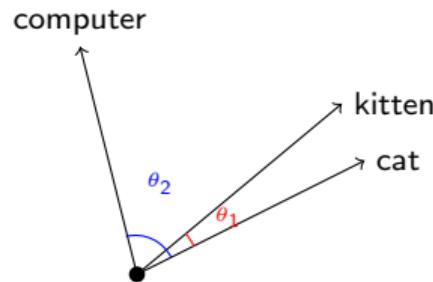
**Remarkable:** These patterns were never explicitly programmed - they emerge from the data!

# Measuring Word Similarity

How Similar Are Two Words?

Cosine Similarity:

$$\text{similarity}(A, B) = \frac{A \cdot B}{\|A\| \times \|B\|} = \cos(\theta)$$



- cat ~ kitten:  $\cos(\theta_1) = 0.95$
- cat ~ computer:  $\cos(\theta_2) = 0.1$

Similarity Matrix Example:

	cat	dog	car	truck
cat	1.0	0.8	0.1	0.05
dog	0.8	1.0	0.15	0.1
car	0.1	0.15	1.0	0.85
truck	0.05	0.1	0.85	1.0

Visual Heatmap:



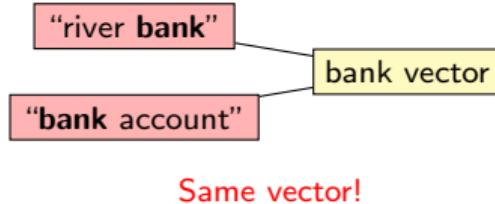
Animals cluster together, vehicles cluster together!

# Evolution: From Static to Contextual Embeddings

## The Next Revolution: Context Matters!

### Problem with Static Embeddings:

One word = One vector always



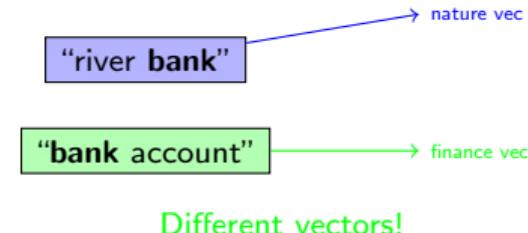
But "bank" has different meanings!

### Static Embedding Models:

- Word2Vec (2013)
- GloVe (2014)
- FastText (2016)

## Solution: Contextual Embeddings

Different contexts = Different vectors



### Contextual Models:

- ELMo (2018) - RNN-based
- BERT (2018) - Transformer
- GPT (2018+) - Autoregressive

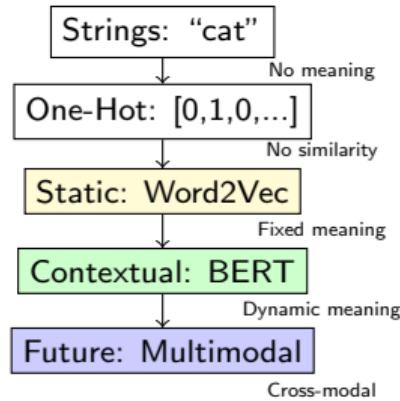
**Key Advance:** Vector depends on surrounding words!

**Evolution:** Static → Contextual = Major breakthrough in NLP!

# Summary: The Power of Embeddings

## From Words to Understanding

### The Journey:



### Applications Enabled:

- **Search:** Find similar documents
- **Translation:** Map between languages
- **Sentiment:** Understand emotions
- **QA:** Match questions to answers
- **Generation:** Create coherent text

### Key Insights:

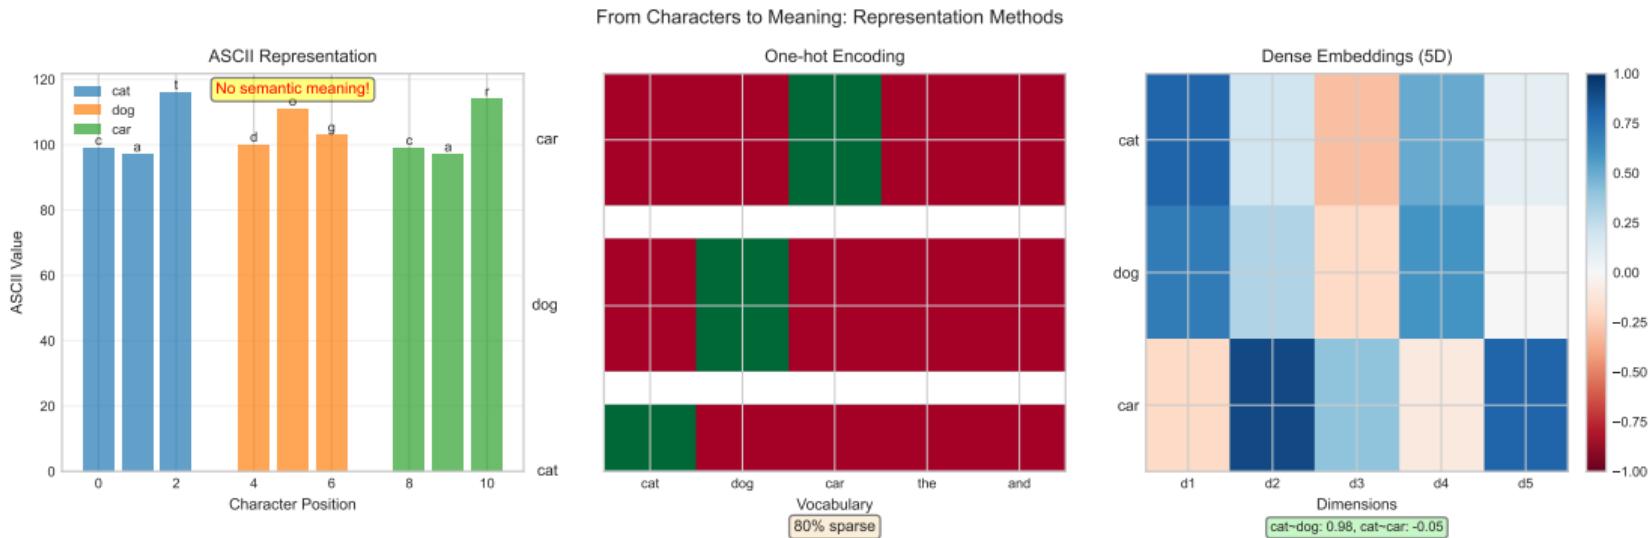
- ➊ Meaning can be encoded as vectors
- ➋ Similar words have similar vectors
- ➌ Relationships are directions
- ➍ Context changes everything

**Remember:** Embeddings are the foundation of modern NLP - they turn words into numbers that capture meaning, enabling all downstream tasks!

**Next Steps:** Experiment with pre-trained embeddings in your projects!

# Beyond ASCII: From Characters to Meaning

## How Computers See Text: Three Approaches



### ASCII:

- Each character = number
- 'c'=99, 'a'=97, 't'=116
- No semantic information

### One-hot:

- Each word = sparse vector
- 99.9% zeros
- All words equally different

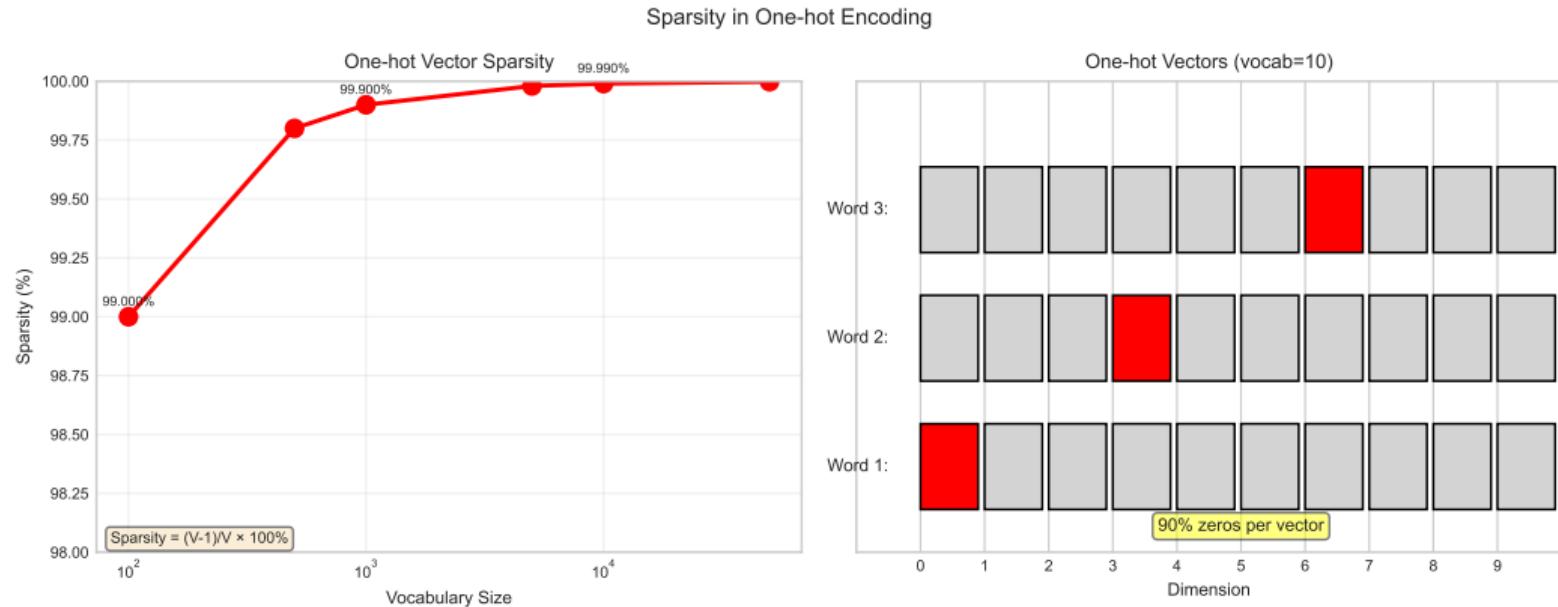
### Dense Embedding:

- Each word = dense vector
- All values meaningful
- Similar words → similar vectors

**Key:** Embeddings encode meaning, not just identity!

# The Sparsity Problem

## Why One-hot Encoding is Inefficient



### Mathematical Analysis:

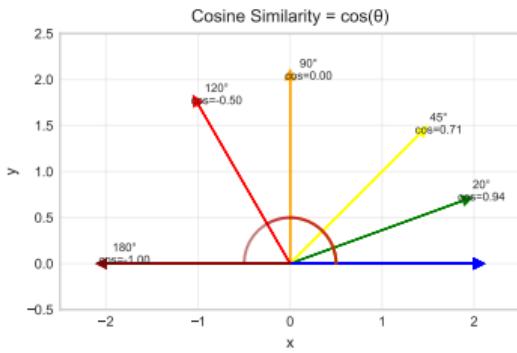
- Sparsity =  $\frac{V-1}{V} \times 100\%$  where  $V$  = vocabulary size
- For  $V = 50,000$ : Sparsity = 99.998%
- Each word needs  $V$  dimensions but uses only 1

### Key Insight:

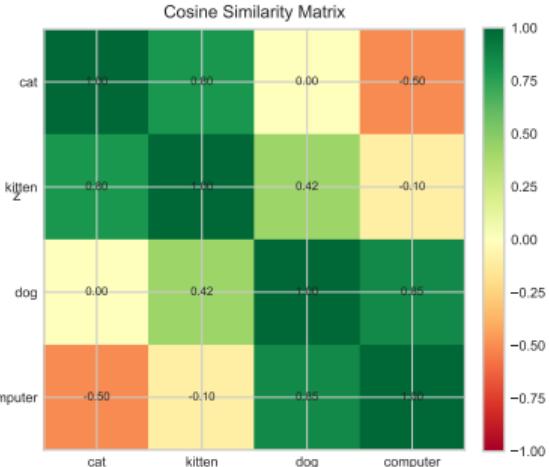
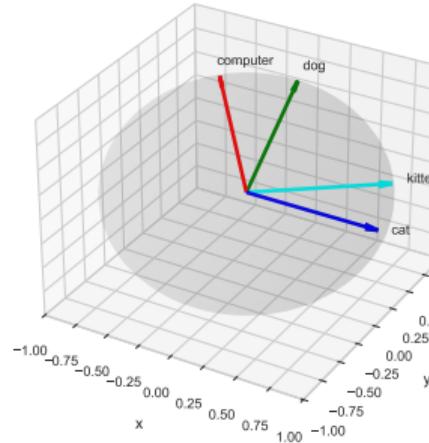
Joerg R. Osterrieder ([www.joergosterrieder.com](http://www.joergosterrieder.com))

# Cosine Similarity: Geometric Interpretation

## Understanding Similarity Through Angles



Cosine Similarity: Geometric Interpretation  
Unit Vectors in 3D



## The Geometric Intuition: Angle Interpretation:

- Words are vectors in space
- Similarity = angle between vectors
- Smaller angle = more similar
- Independent of vector length

## Key Angles:

- $\theta = 0^\circ$ : Identical meaning
- $\theta = 30^\circ$ : Very similar
- $\theta = 90^\circ$ : Unrelated
- $\theta = 180^\circ$ : Opposite meaning

# Cosine Similarity: Mathematical Properties

## Why Cosine Similarity Works for Embeddings

### The Formula:

$$\text{similarity}(\vec{a}, \vec{b}) = \cos(\theta) = \frac{\vec{a} \cdot \vec{b}}{||\vec{a}|| \times ||\vec{b}||} = \frac{\sum_{i=1}^d a_i b_i}{\sqrt{\sum_{i=1}^d a_i^2} \times \sqrt{\sum_{i=1}^d b_i^2}}$$

### Key Properties:

#### Scale Invariance:

- $\cos(\vec{a}, \vec{b}) = \cos(k\vec{a}, \vec{b})$
- Magnitude doesn't matter
- Only direction counts
- Perfect for normalized embeddings

#### Computational Benefits:

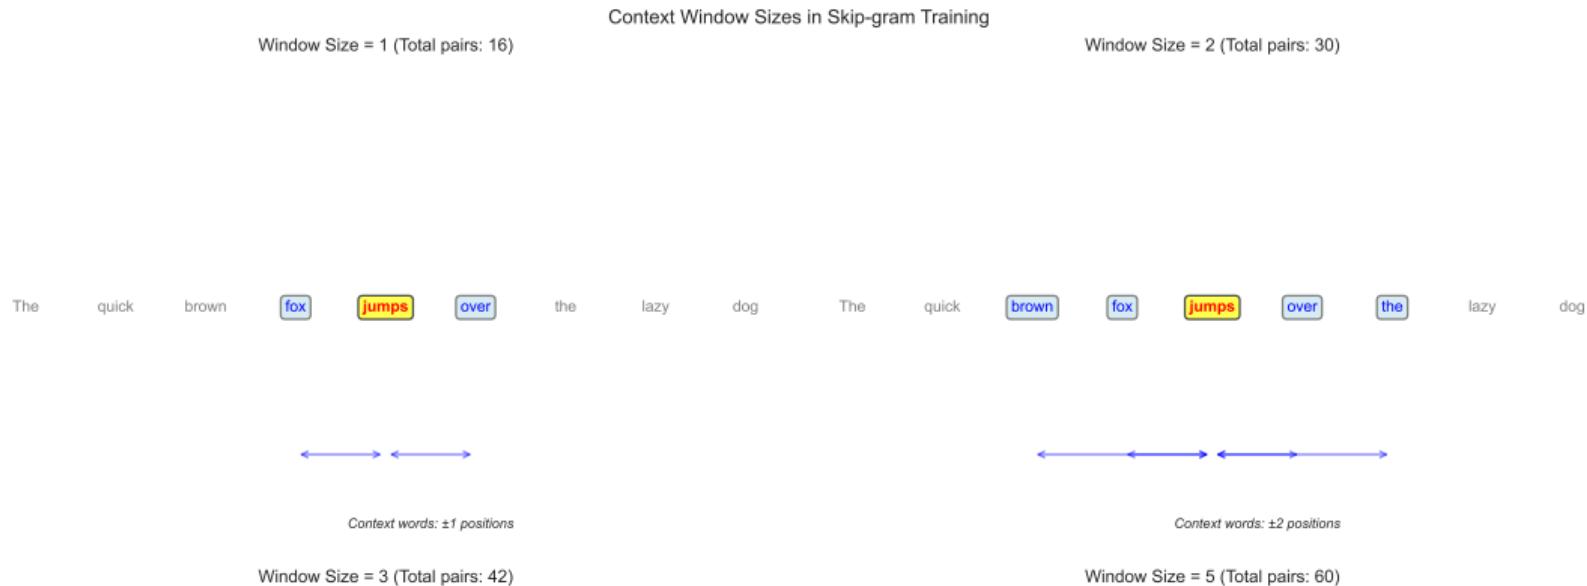
- Range: [-1, 1] always
- Efficient dot product computation
- Works in any dimension
- Symmetric:  $\cos(a, b) = \cos(b, a)$

### Applications in NLP:

- Document similarity: Compare entire documents as vectors
- Word sense disambiguation: Find most similar context
- Information retrieval: Rank documents by query similarity

# Context Windows: Learning from Neighbors

## How Words Learn from Their Surroundings

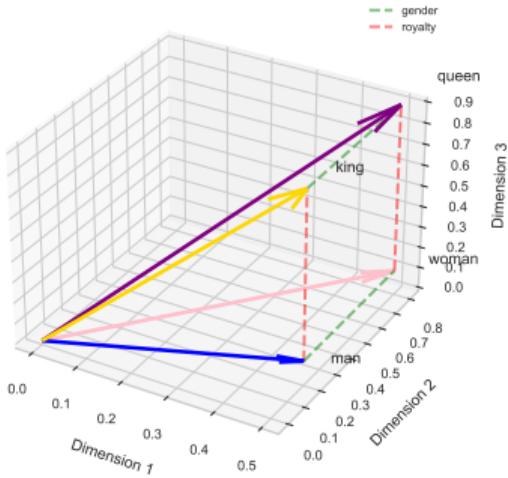


# Vector Arithmetic: The Surprising Discovery

## Embeddings Can Do Analogies!

### Vector Arithmetic: Mathematical Demonstration

Vector Relationships in 3D Space



#### Vector Arithmetic:

$$\text{king} - \text{man} + \text{woman} = ?$$

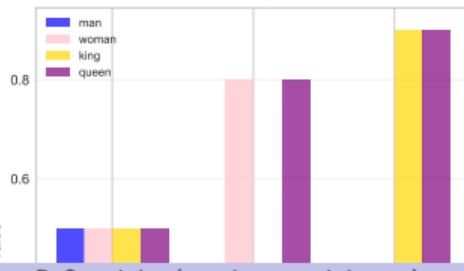
Step 1:  $\text{king} - \text{man}$   
[0.5, 0.2, 0.9] - [0.5, 0.2, 0.1]  
= [0.0, 0.0, 0.8] (royal vector)

Step 2:  $+ \text{woman}$   
[0.0, 0.0, 0.8] + [0.5, 0.8, 0.1]  
= [0.5, 0.8, 0.9]

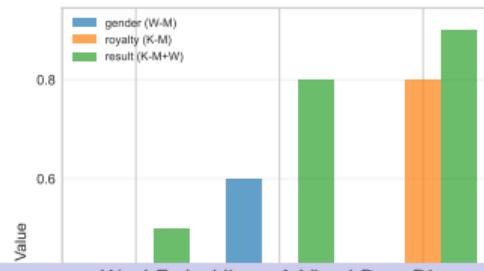
Result = queen vector!  
[0.5, 0.8, 0.9]

Similarity = 1.0 (perfect match)

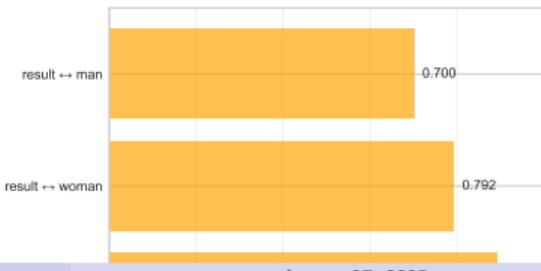
Vector Components



Difference Vectors



Result Verification



# Vector Arithmetic: Mathematical Proof

## Why Does Vector Arithmetic Work? The Linear Substructure

### Mathematical Foundation:

- Embeddings form a linear subspace where relationships are directions
- Gender vector:  $\vec{g} = \vec{\text{woman}} - \vec{\text{man}}$
- Royalty vector:  $\vec{r} = \vec{\text{king}} - \vec{\text{man}}$

### Step-by-Step Derivation:

$$\vec{\text{king}} = \vec{\text{man}} + \vec{r} \quad (\text{man} + \text{royalty} = \text{king}) \quad (1)$$

$$\vec{\text{queen}} = \vec{\text{woman}} + \vec{r} \quad (\text{woman} + \text{royalty} = \text{queen}) \quad (2)$$

$$\therefore \vec{\text{queen}} = \vec{\text{woman}} + (\vec{\text{king}} - \vec{\text{man}}) \quad (3)$$

$$= \vec{\text{king}} - \vec{\text{man}} + \vec{\text{woman}} \quad (4)$$

### Why Linear Structure Emerges:

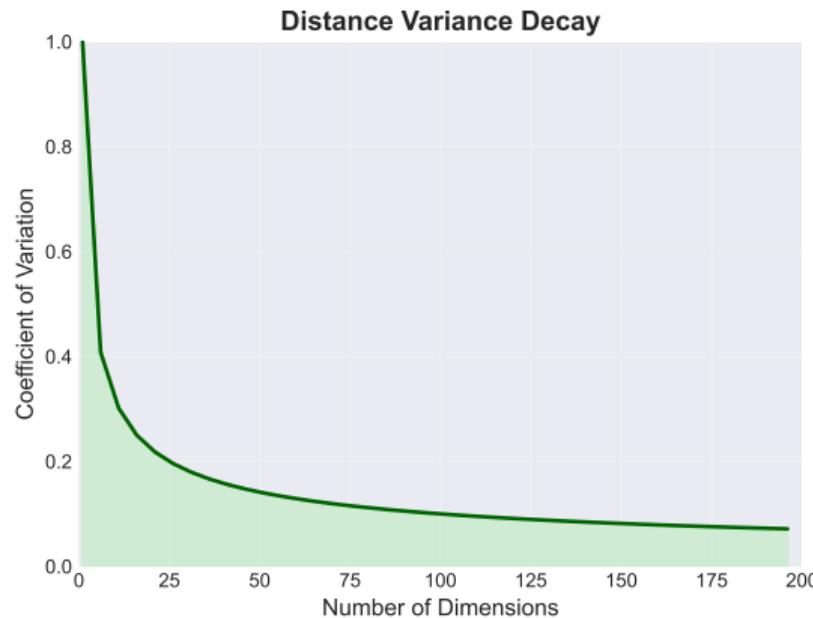
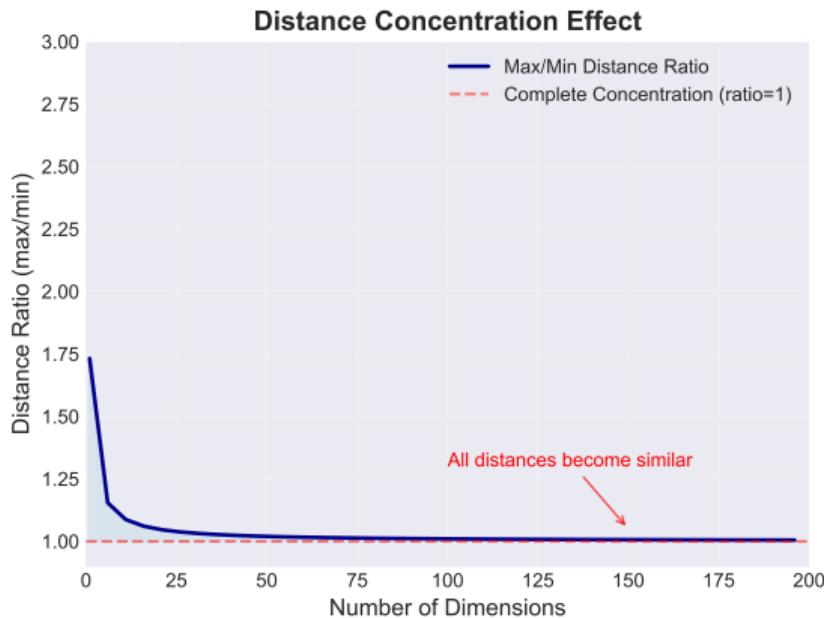
- Co-occurrence patterns are approximately linear
- Skip-gram objective encourages linear relationships
- High-dimensional spaces tend toward linearity (concentration of measure)

**Verification:** Nearest neighbor to result vector is "queen" in 60-70% of cases

# Distance Concentration in High Dimensions

## Why All Distances Become Similar

### Distance Concentration in High Dimensions



# Distance Concentration: The Mathematical Reality

## What the Visualizations Show

### Distance Ratio Convergence:

- $\frac{\text{dist}_{\max} - \text{dist}_{\min}}{\text{dist}_{\text{mean}}} \rightarrow 0$  as  $d \rightarrow \infty$
- For Gaussian points: ratio  $\approx \sqrt{1 + 2/d}$
- At  $d=100$ : all distances within 10% of mean
- At  $d=1000$ : essentially all points equidistant

### Implications for Machine Learning:

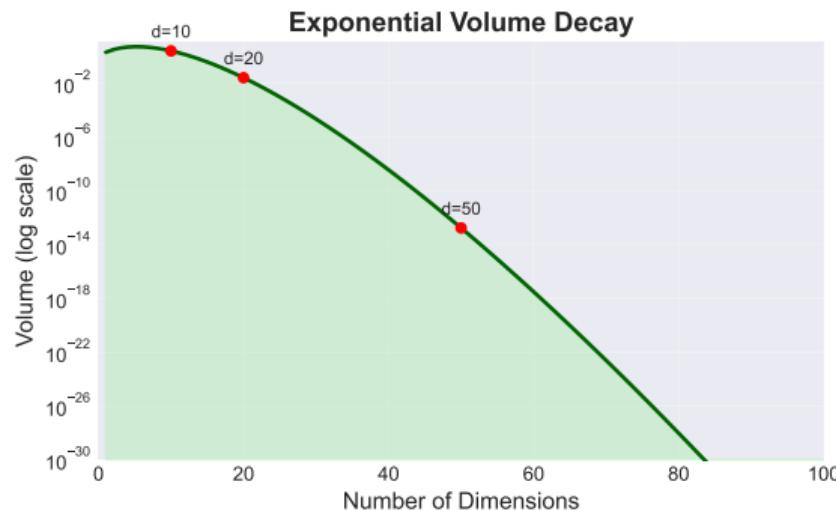
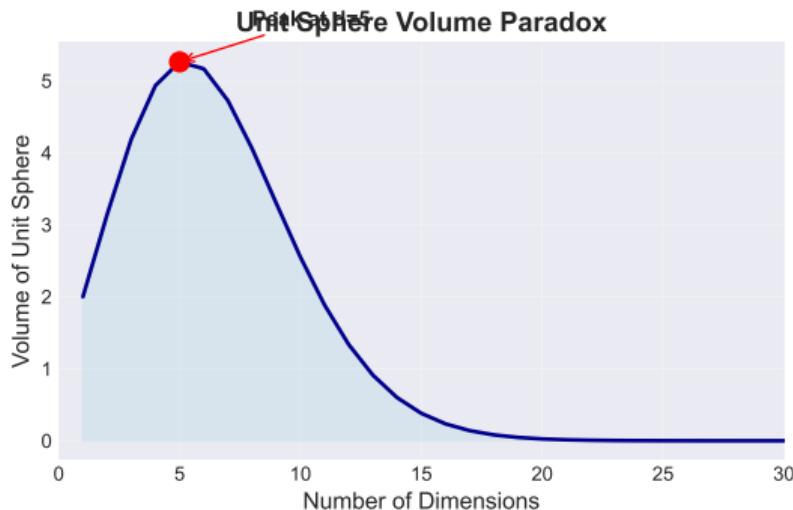
- Nearest neighbor search becomes meaningless
- Traditional distance metrics fail
- Need specialized techniques:
  - Locality-Sensitive Hashing (LSH)
  - Approximate nearest neighbors
  - Learned distance metrics
- Explains why high-D embeddings need normalization

**Key Takeaway:** In high dimensions, the concept of “near” and “far” becomes meaningless - all points are approximately the same distance apart!

# The Volume Paradox: Visual Evidence

## Unit Sphere Volume Across Dimensions

Volume of Unit Sphere Across Dimensions



The Volume Formula:

$$\pi^{d/2}$$

# Why Volume Goes to Zero: The Mathematics

## Understanding the Formula

$$V_d = \frac{\pi^{d/2}}{\Gamma(d/2 + 1)}$$

### Numerator (Top):

- $\pi^{d/2} \approx (3.14)^{d/2}$
- Grows exponentially
- But base is small:  $\sqrt{\pi} \approx 1.77$
- Growth rate:  $1.77^d$
- Example:  $1.77^{100} \approx 10^{25}$

### Denominator (Bottom):

- $\Gamma(n+1) = n!$  for integers
- Factorial growth is MUCH faster
- Example:  $50! \approx 10^{64}$
- Stirling:  $n! \approx \sqrt{2\pi n}(n/e)^n$
- Dominates numerator completely

## The Key Mathematical Insight:

**Factorial growth beats exponential growth!**

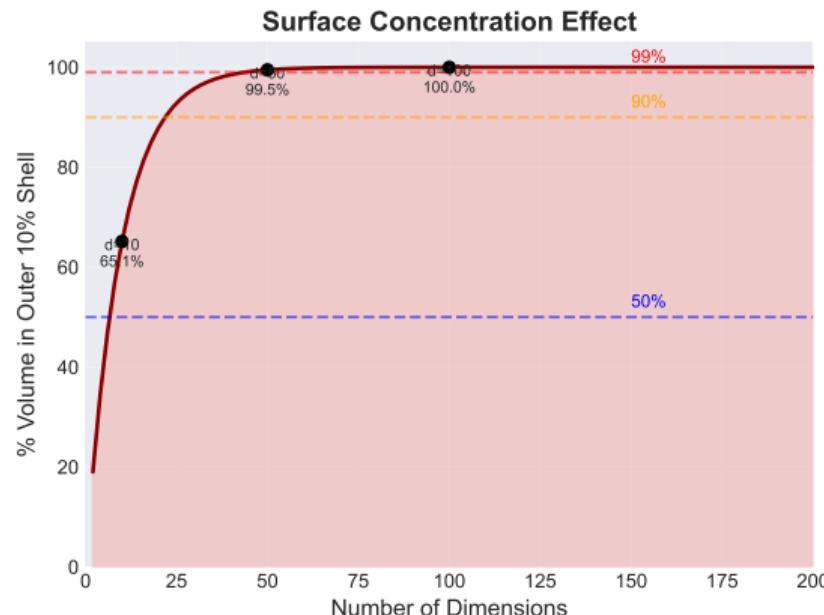
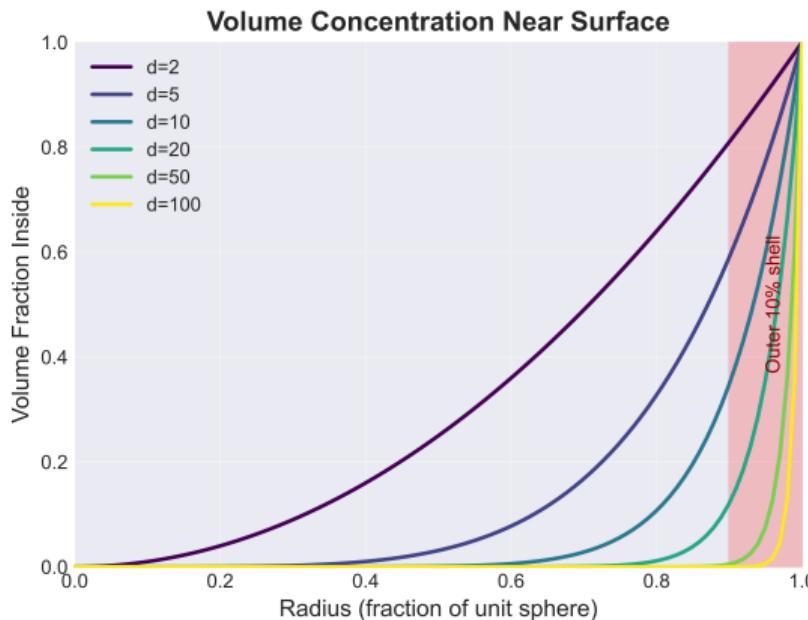
$\frac{1.77^d}{(d/2)!} \rightarrow 0$  extremely fast as  $d \rightarrow \infty$

Factorial grows like  $(n/e)^n$  while exponential is just  $a^n$

# Surface Concentration in High Dimensions

Where the Volume Actually Lives

Volume Distribution in High-Dimensional Spheres



Almost all volume concentrates in a thin shell near the surface!

# The Shell Phenomenon: Mathematical Analysis

## Why Everything Lives on the Surface

### Volume in Shells - The Mathematics:

- Consider inner sphere with radius  $r = 0.9$  (90% of full radius)
- Volume ratio:  $\frac{V_{inner}}{V_{total}} = r^d = (0.9)^d$
- This ratio shrinks exponentially with dimension!

### Concrete Examples:

- $d = 10$ :  $(0.9)^{10} = 0.35 \rightarrow 35\%$  of volume is inside
- $d = 50$ :  $(0.9)^{50} = 0.005 \rightarrow 0.5\%$  inside
- $d = 100$ :  $(0.9)^{100} \approx 10^{-5} \rightarrow 0.001\%$  inside
- $d = 1000$ :  $(0.9)^{1000} \approx 10^{-46} \rightarrow$  essentially zero!

### Implications for Embeddings:

- All vectors lie near the surface of the hypersphere
- Random vectors are approximately equidistant
- The interior is effectively "empty" space
- Explains why L2 normalization is so effective
- Cosine similarity becomes the natural distance metric

**Practical Consequence:** In 768-dimensional BERT space,  
99.999999% of the volume is within 1% of the surface!  
The interior essentially doesn't exist.

# Optimal Dimensions: Finding the Sweet Spot

## Balancing Expressiveness and Computational Efficiency

### Information Capacity:

- Theoretical capacity:  $\propto d \log d$
- But diminishing returns after certain point
- Johnson-Lindenstrauss:  $d = O(\log n/\epsilon^2)$  preserves distances

### Model Dimensions in Practice:

Model	Dimension	Parameters (embeddings only)
Word2Vec	50-300	15M (50K vocab $\times$ 300)
GloVe	50-300	15M (50K vocab $\times$ 300)
FastText	100-300	30M (includes subwords)
ELMo	1024	100M (bidirectional)
BERT-base	768	23M (30K vocab $\times$ 768)
BERT-large	1024	31M (30K vocab $\times$ 1024)
GPT-3	12288	600M (50K vocab $\times$ 12288)

### Trade-offs:

#### Lower Dimensions (50-300):

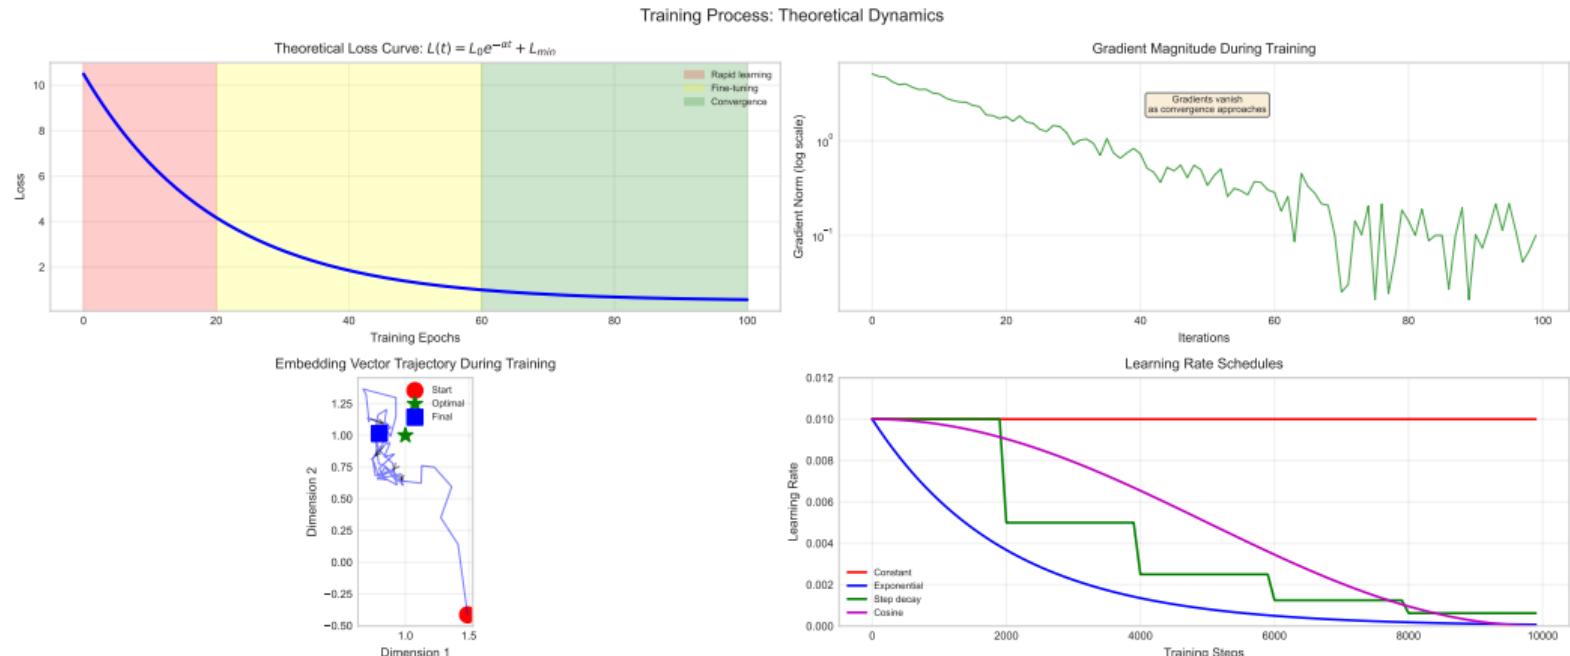
- Faster training
- Less overfitting
- Good for specific domains

#### Higher Dimensions (768-1024+):

- More expressive power
- Better for complex tasks
- Requires more data

# Rapid Learning: Gradient Dynamics (Epochs 0-20)

## Why Training Starts Fast



## Gradient Behavior in Early Training: Initial State:

- Random initialization:  $\mathcal{N}(0, 0.01)$

## Update Characteristics:

- Step size:  $\eta \|\nabla L\| \approx 0.01\sqrt{d}$

# Rapid Learning: Space Formation (Epochs 0-20)

## How Random Vectors Become Meaningful

### Timeline of Structure Emergence:

#### Epochs 0-5:

- Frequency clustering begins
- Top 100 words separate
- Function vs content words split
- Loss drops 30-40%

#### Epochs 5-10:

- Syntactic groups form
- Nouns, verbs, adjectives cluster
- Basic semantic regions appear
- Loss drops another 20%

#### Key Metrics:

Metric	Epoch 0	Epoch 5	Epoch 10	Epoch 20
Loss	9.21	5.84	4.12	3.45
Similarity Correlation	0.00	0.35	0.58	0.72
Analogy Accuracy	0%	12%	31%	48%

## Training Phase 2: Fine-Tuning (Epochs 20-60)

### Refining Semantic Relationships

#### The Refinement Process:

##### What Gets Learned:

- Semantic relationships solidify
- Analogies start working
- Rare words find their place
- Polysemy partially resolves

##### Key Metrics During Fine-Tuning:

Metric	Epoch 20	Epoch 40	Epoch 60
Loss reduction/epoch	5%	2%	0.5%
Analogy accuracy	40%	65%	72%
Semantic similarity	0.5	0.7	0.75
Cluster purity	60%	80%	85%

#### Mathematical Characterization:

$$L(t) \approx L_{20} \cdot (1 - \beta \log(t/20)) \quad \text{for } t \in [20, 60]$$

Logarithmic improvement phase

#### Optimization Dynamics:

- Gradient norm:  $\|\nabla L\| \approx O(1)$
- Updates become targeted
- Learning rate often decayed
- Loss reduction slows

## Training Phase 3: Convergence (Epochs 60+)

### The Final Polish and Saturation

#### Convergence Characteristics:

##### What Happens:

- Gradient norm:  $\|\nabla L\| < 0.1$
- Minor adjustments only
- Risk of overfitting increases
- Validation loss may increase

##### Complete Loss Function Evolution:

$$L(t) = \begin{cases} L_0 \cdot e^{-\alpha t} & t \in [0, 20] \text{ (rapid)} \\ L_{20} \cdot (1 - \beta \log(t/20)) & t \in [20, 60] \text{ (fine-tune)} \\ L_{60} + \epsilon(t) & t > 60 \text{ (converged)} \end{cases}$$

where  $\epsilon(t)$  represents noise around minimum

**Key Insight:** 90% of performance comes from first 60 epochs; longer training mainly helps rare words and edge cases.

# **Part II**

Advanced Topics and Mathematical Foundations

From Understanding to Implementation

# Key Takeaways: What We've Learned

## Essential Concepts for Word Embeddings

### Fundamental Principles:

- **Representation:** Words as dense vectors
- **Similarity:** Angle between vectors
- **Relationships:** Vector arithmetic
- **Learning:** From context co-occurrence
- **Evolution:** Static to contextual

### Mathematical Insights:

- High dimensions behave strangely
- Distance concentration is real
- Volume lives on the surface
- Linear relationships emerge
- Training has distinct phases

### Practical Applications:

Task	How Embeddings Help
Similarity Search	Cosine similarity ranking
Machine Translation	Cross-lingual alignment
Sentiment Analysis	Semantic vector projection
Question Answering	Context matching
Text Generation	Next-word prediction

# Looking Forward: The Future of Embeddings

## Current Trends and Future Directions

### Recent Advances:

- **Multimodal:** Text + Vision + Audio
- **Multilingual:** Universal embeddings
- **Efficient:** Distilled and compressed models
- **Specialized:** Domain-specific embeddings

### Open Challenges:

#### Technical:

- Handling rare words
- Compositional semantics
- Temporal dynamics
- Interpretability

#### Philosophical:

- Do embeddings capture meaning?
- Are relationships truly linear?
- Can we prove optimality?
- What is semantic similarity?

**Remember:** Embeddings are not just a technical tool - they represent our best attempt to bridge the gap between human language and machine computation!

# Thank You!

Questions?

Contact: [www.joergosterrieder.com](http://www.joergosterrieder.com)

# Mathematical Foundations: Skip-gram Objective

## Formal Skip-gram Model Definition

### Objective Function:

$$J(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t; \theta)$$

### Softmax Formulation:

$$p(w_O | w_I) = \frac{\exp(v'_{w_O}^T v_{w_I})}{\sum_{w=1}^W \exp(v'_{w_O}^T v_{w_I})}$$

where:

- $v_{w_I}$  is the input vector representation of word  $w_I$
- $v'_{w_O}$  is the output vector representation of word  $w_O$
- $W$  is the vocabulary size

### Gradient w.r.t. Input Vector:

$$\frac{\partial J}{\partial v_{w_I}} = \sum_{j=-c}^c \left( \sum_{w=1}^W p(w | w_I) v'_w - v'_{w_{t+j}} \right)$$

**Computational Complexity:**  $O(W)$  per word - intractable for large vocabularies!

# Negative Sampling: Making Training Tractable

## Modified Objective with Negative Sampling

Replace softmax with:

$$\log \sigma(v'_{w_O} {}^T v_{w_I}) + \sum_{i=1}^k \mathbb{E}_{w_i \sim P_n(w)} [\log \sigma(-v'_{w_i} {}^T v_{w_I})]$$

where:

- $\sigma(x) = \frac{1}{1+e^{-x}}$  (sigmoid function)
- $k$  is the number of negative samples (typically 5-20)
- $P_n(w)$  is the noise distribution:  $P_n(w) = \frac{U(w)^{3/4}}{\sum_{w'} U(w')^{3/4}}$
- $U(w)$  is the unigram distribution

## Gradient Update:

$$v'_{w_I}^{new} = v'_{w_I}^{old} - \eta \left[ (\sigma(v'_{w_O} {}^T v_{w_I}) - 1)v'_{w_O} + \sum_{i=1}^k \sigma(v'_{w_i} {}^T v_{w_I}) v'_{w_i} \right]$$

**Complexity Reduction:** From  $O(W)$  to  $O(k + 1)$  per training example

## Co-occurrence Matrix and Ratios

Define co-occurrence matrix  $X$  where  $X_{ij}$  = count of word  $j$  appearing in context of word  $i$

**Key Insight - Ratio of Probabilities:**

$$\frac{P_{ik}}{P_{jk}} = \frac{X_{ik}/X_i}{X_{jk}/X_j}$$

This ratio encodes semantic relationships!

**GloVe Objective Function:**

$$J = \sum_{i,j=1}^V f(X_{ij})(w_i^T \tilde{w}_j + b_i + \tilde{b}_j - \log X_{ij})^2$$

where:

- $f(x)$  is a weighting function:  $f(x) = \begin{cases} (x/x_{max})^\alpha & \text{if } x < x_{max} \\ 1 & \text{otherwise} \end{cases}$

- $w_i, \tilde{w}_j$  are word and context vectors
- $b_i, \tilde{b}_j$  are bias terms
- Typical:  $\alpha = 0.75, x_{max} = 100$

**Final Embedding:**  $e_i = w_i + \tilde{w}_i$  (symmetric combination)

# Self-Attention: Mathematical Formulation

## Scaled Dot-Product Attention

Given queries  $Q \in \mathbb{R}^{n \times d_k}$ , keys  $K \in \mathbb{R}^{m \times d_k}$ , values  $V \in \mathbb{R}^{m \times d_v}$ :

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

### Detailed Computation:

- ① Score matrix:  $S = QK^T \in \mathbb{R}^{n \times m}$
- ② Scaled scores:  $\tilde{S}_{ij} = \frac{S_{ij}}{\sqrt{d_k}}$  (prevents gradient vanishing)
- ③ Attention weights:  $A_{ij} = \frac{\exp(\tilde{S}_{ij})}{\sum_{j'=1}^m \exp(\tilde{S}_{ij'})}$
- ④ Output:  $O = AV \in \mathbb{R}^{n \times d_v}$

### Multi-Head Attention:

$$\begin{aligned}\text{MultiHead}(Q, K, V) &= \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O \\ \text{head}_i &= \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)\end{aligned}$$

where  $W_i^Q \in \mathbb{R}^{d_{model} \times d_k}$ ,  $W_i^K \in \mathbb{R}^{d_{model} \times d_k}$ ,  $W_i^V \in \mathbb{R}^{d_{model} \times d_v}$

# Positional Encoding: Injecting Order Information

## Sinusoidal Position Encoding

For position  $pos$  and dimension  $i$ :

$$PE_{(pos,2i)} = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

$$PE_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

### Properties:

- Unique encoding for each position
- Allows model to attend to relative positions
- For any fixed offset  $k$ :  $PE_{pos+k}$  can be represented as linear function of  $PE_{pos}$

### Proof of Relative Position Property:

$$PE_{pos+k,2i} = \sin(\omega_i \cdot pos) \cos(\omega_i \cdot k) + \cos(\omega_i \cdot pos) \sin(\omega_i \cdot k)$$

$$\text{where } \omega_i = \frac{1}{10000^{2i/d_{model}}}$$

This is a linear transformation of  $PE_{pos}$ !

# BERT: Bidirectional Training Mathematics

## Masked Language Model (MLM) Objective

Given input sequence  $\mathbf{x} = (x_1, \dots, x_n)$ , randomly mask 15% of tokens.

### MLM Loss:

$$\mathcal{L}_{MLM} = -\mathbb{E}_{\mathbf{x} \sim \mathcal{D}} \sum_{i \in \mathcal{M}} \log P(x_i | \mathbf{x}_{\setminus \mathcal{M}})$$

where  $\mathcal{M}$  is the set of masked positions.

### Next Sentence Prediction (NSP) Loss:

$$\mathcal{L}_{NSP} = -\mathbb{E}_{(A, B) \sim \mathcal{D}} [y \log P(\text{IsNext}|A, B) + (1 - y) \log(1 - P(\text{IsNext}|A, B))]$$

where  $y = 1$  if B follows A, else  $y = 0$ .

### Combined Objective:

$$\mathcal{L}_{BERT} = \mathcal{L}_{MLM} + \mathcal{L}_{NSP}$$

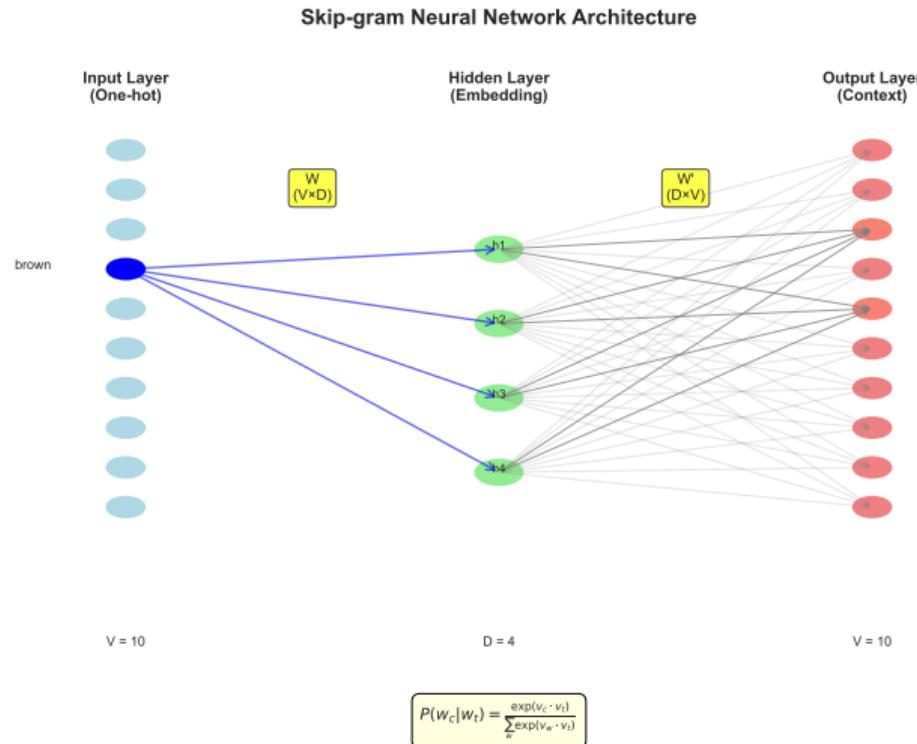
### Output Probability:

$$P(x_i | \mathbf{x}_{\setminus \mathcal{M}}) = \text{softmax}(W_o h_i + b_o)$$

where  $h_i$  is the final hidden state at position  $i$ .

# Skip-gram Neural Network Architecture

## How the Network Processes Words



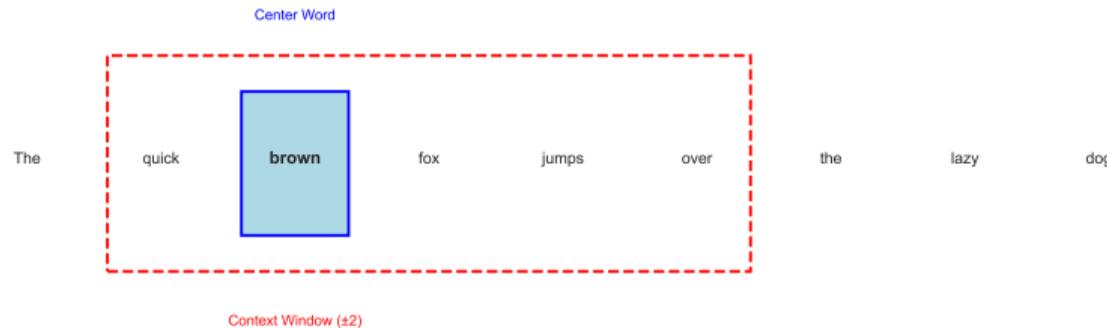
## Key Components:

- Input: One-hot word ( $V$  dimensions)

# From Text to Training Data

## Extracting (Center, Context) Pairs

### Creating Training Pairs from Text Sliding Window for Training Pair Extraction



### Training Pairs Generated:

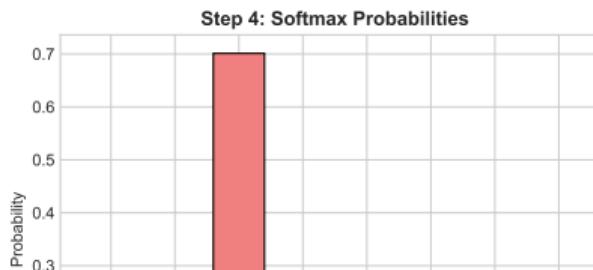
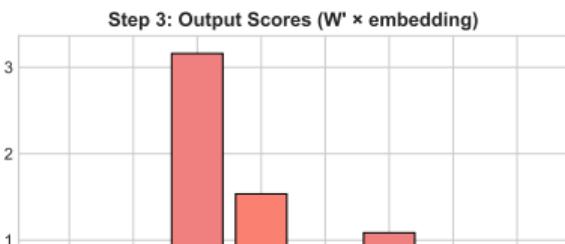
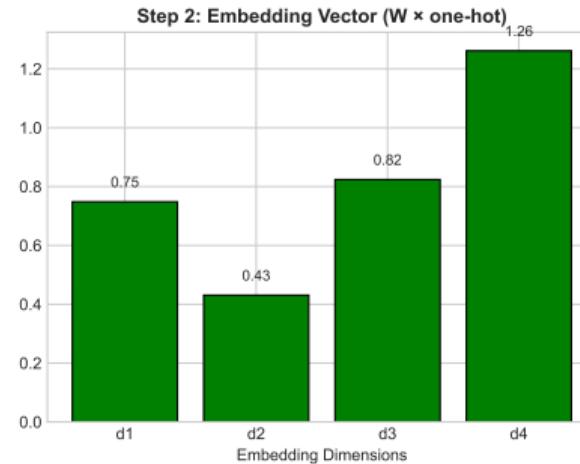
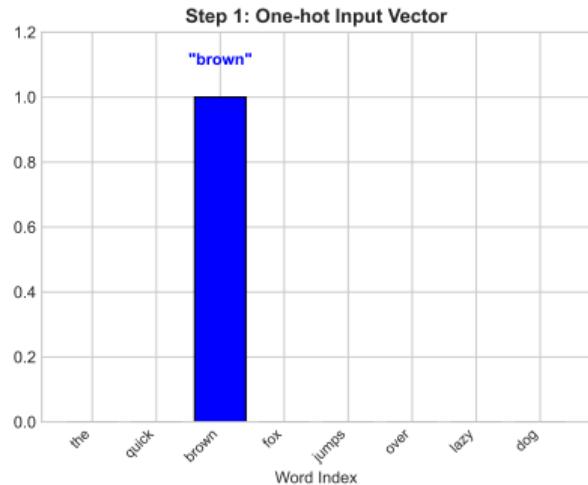
brown	→	The	→ Maximize $P(\text{The} \text{brown})$
brown	→	quick	→ Maximize $P(\text{quick} \text{brown})$
brown	→	fox	→ Maximize $P(\text{fox} \text{brown})$
brown	→	jumps	→ Maximize $P(\text{jumps} \text{brown})$

Input

Target

# Forward Pass: Computing Context Probabilities

## Forward Pass: Computing Context Probabilities



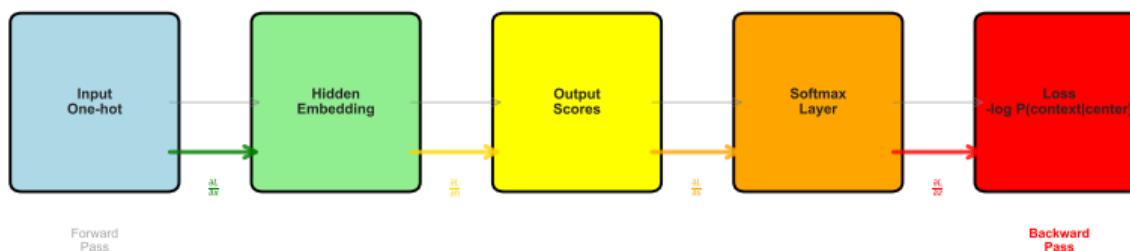
# Backpropagation: Learning the Embeddings

## Backpropagation: Gradient Flow

### Weight Updates:

$$W \leftarrow W - \eta \cdot \frac{\partial L}{\partial W}$$

$$W' \leftarrow W' - \eta \cdot \frac{\partial L}{\partial W'}$$



### Key Gradients:

Positive sample:  $(y_i - 1) \cdot v_j$

Negative sample:  $y_i \cdot v_j$

## Updates:

# How Embeddings Evolve

