

# Word Embeddings: A Visual Deep Dive

## From One-Hot Vectors to Contextual Representations

Joerg R. Osterrieder

[www.joergosterrieder.com](http://www.joergosterrieder.com)

August 28, 2025

# Outline

# What You Will Learn

By the end of this presentation, you will understand:


- ① **Representation Problem:** Why computers need numerical representations of words
- ② **Evolution of Embeddings:** From one-hot to contextual representations
- ③ **Mathematical Foundations:** The theory behind word embeddings
- ④ **Vector Operations:** How semantic relationships emerge from vectors
- ⑤ **High-Dimensional Challenges:** The curse of dimensionality
- ⑥ **Training Dynamics:** How embeddings learn meaningful representations
- ⑦ **Skip-gram Architecture:** Deep dive into Word2Vec training

**Key Insight:** Words are not isolated symbols but points in a continuous semantic space

# The Fundamental Problem: Computers Don't Understand Words

How do we represent meaning mathematically?

Human Understanding:



Computer's Dilemma:

- Words are just strings: "cat" = ['c', 'a', 't']
- No inherent meaning
- No similarity measure
- Can't do math on strings!

What We Need:

Convert: "cat" → [0.2, -0.4, 0.7, ...]  
Such that: similar words → similar vectors

Rich semantic connections!

Goal: Capture meaning in numbers so computers can process language

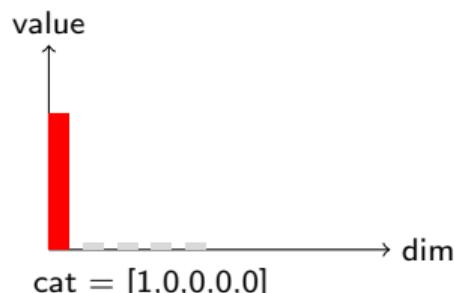
# Starting Point: One-Hot Encoding

## The Simplest Approach - But Fundamentally Flawed

How One-Hot Works:

Word	Vector
cat	[1, 0, 0, 0, 0]
dog	[0, 1, 0, 0, 0]
mat	[0, 0, 1, 0, 0]
sat	[0, 0, 0, 1, 0]
hat	[0, 0, 0, 0, 1]

Visual Representation:



### Critical Problems:

#### ① No Similarity:

$$\text{similarity}(\text{cat}, \text{kitten}) = 0$$

$$\text{similarity}(\text{cat}, \text{computer}) = 0$$

Both are equally dissimilar!

#### ② Huge Dimensions:

- English: 170,000+ words
- Each word = 170,000-dim vector
- 99.999% zeros (sparse!)

#### ③ No Relationships:

$$\text{cat} + \text{kitten} = [1, 0, 0, \dots] + [0, 1, 0, \dots] = [1, 1, 0, \dots]$$

Meaningless!

**Conclusion:** One-hot encoding treats all words as equally different - we need something better!

# Dense Embeddings: The Solution

## From Sparse to Dense - Capturing Meaning in Vectors

The Transformation:

One-Hot: [0,1,0,0,...0]

50,000 dimensions

99.998% zeros



Dense: [0.2, -0.4, 0.7, ...]

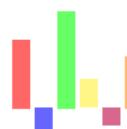
100-300 dimensions

All values meaningful

Visual Comparison:



Sparse (One-Hot):



Dense: all values meaningful

Benefits:

- 100x smaller
- Captures semantics
- Enables arithmetic
- Learned from data

Example Vector:


$$\text{cat} = [0.21, -0.43, 0.67, 0.15, -0.22, \dots]$$

Each dimension captures some aspect of meaning

# The Embedding Space: Where Words Live


## Visualizing Word Relationships in Vector Space

2D Projection of Word Vectors:



### Key Properties:

- ① **Clustering:** Similar words group together
- ② **Distance = Similarity:**
  - cat  $\leftrightarrow$  dog: close
  - cat  $\leftrightarrow$  run: far
- ③ **Directions = Relations:**



Gender direction is consistent!

**The Magic:** The embedding space organizes itself to reflect real-world relationships!

# Learning Embeddings: Word2Vec


## How Do We Learn These Vectors?

### The Distributional Hypothesis:


"You shall know a word by the company it keeps" - Firth (1957)

### Two Approaches:

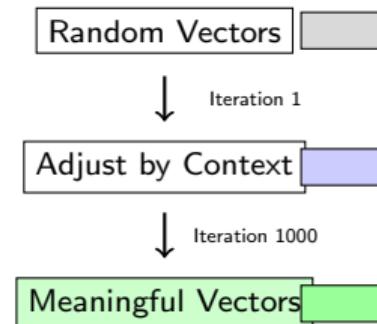
1. Skip-gram: Predict context from word



2. CBOW: Predict word from context



### Training Process Visualization:



### Objective Function:

$$\max \sum_{t=1}^T \sum_{-c \leq j \leq c} \log P(w_{t+j} | w_t)$$

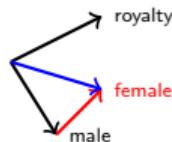
Maximize probability of context words

# Vector Arithmetic: The Surprising Discovery

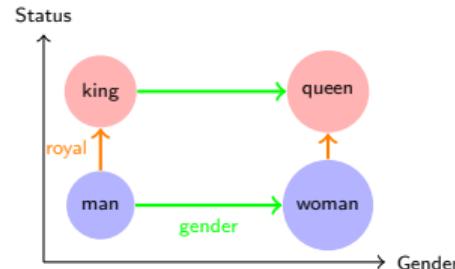
Embeddings Capture Analogies!

Famous Examples:

$$\text{king} - \text{man} + \text{woman} = \text{queen}$$



Why Does This Work?



More Analogies:

- Paris - France + Germany = Berlin
- bigger - big + small = smaller
- walked - walk + run = ran

The Pattern:

- Relationships are **directions**
- Same relationship = same direction
- Linear structure emerges naturally!


**Remarkable:** These patterns were never explicitly programmed - they emerge from the data!

# Measuring Word Similarity

How Similar Are Two Words?

Cosine Similarity:

$$\text{similarity}(A, B) = \frac{A \cdot B}{\|A\| \times \|B\|} = \cos(\theta)$$




- cat ~ kitten:  $\cos(\theta_1) = 0.95$
- cat ~ computer:  $\cos(\theta_2) = 0.1$

Similarity Matrix Example:

	cat	dog	car	truck
cat	1.0	0.8	0.1	0.05
dog	0.8	1.0	0.15	0.1
car	0.1	0.15	1.0	0.85
truck	0.05	0.1	0.85	1.0

Visual Heatmap:




Animals cluster together, vehicles cluster together!

# Evolution: From Static to Contextual Embeddings

## The Next Revolution: Context Matters!

### Problem with Static Embeddings:

One word = One vector always




But "bank" has different meanings!

### Static Embedding Models:

- Word2Vec (2013)
- GloVe (2014)
- FastText (2016)

## Solution: Contextual Embeddings

Different contexts = Different vectors



### Contextual Models:

- ELMo (2018) - RNN-based
- BERT (2018) - Transformer
- GPT (2018+) - Autoregressive

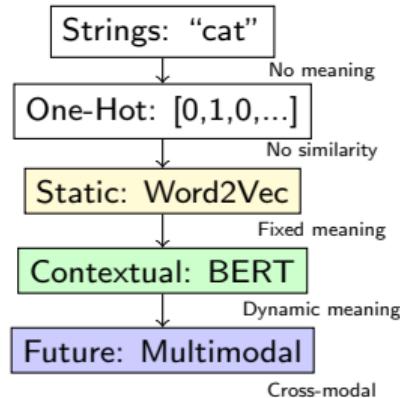
**Key Advance:** Vector depends on surrounding words!

**Evolution:** Static → Contextual = Major breakthrough in NLP!

# Summary: The Power of Embeddings

## From Words to Understanding

### The Journey:



### Applications Enabled:

- **Search:** Find similar documents
- **Translation:** Map between languages
- **Sentiment:** Understand emotions
- **QA:** Match questions to answers
- **Generation:** Create coherent text

### Key Insights:

- ➊ Meaning can be encoded as vectors
- ➋ Similar words have similar vectors
- ➌ Relationships are directions
- ➍ Context changes everything

**Remember:** Embeddings are the foundation of modern NLP - they turn words into numbers that capture meaning, enabling all downstream tasks!

**Next Steps:** Experiment with pre-trained embeddings in your projects!

# Beyond ASCII: From Characters to Meaning

## How Computers See Text: Three Approaches



### ASCII:

- Each character = number
- 'c'=99, 'a'=97, 't'=116
- No semantic information

### One-hot:

- Each word = sparse vector
- 99.9% zeros
- All words equally different

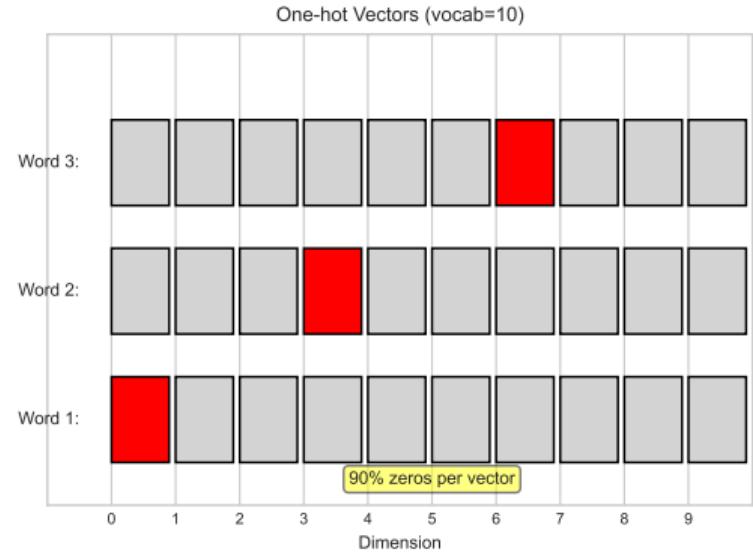

### Dense Embedding:

- Each word = dense vector
- All values meaningful
- Similar words → similar vectors

**Key:** Embeddings encode meaning, not just identity!

## The Sparsity Problem

## Why One-hot Encoding is Inefficient




## **Mathematical Analysis:**

- Sparsity =  $\frac{V-1}{V} \times 100\%$  where  $V$  = vocabulary size
  - For  $V = 50,000$ : Sparsity = 99.998%
  - Each word needs  $V$  dimensions but uses only 1


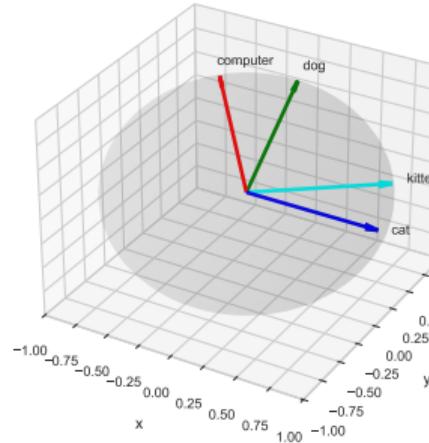
### **Key Insight:**

# Cosine Similarity: Geometric Interpretation

## Understanding Similarity Through Angles



Cosine Similarity: Geometric Interpretation  
Unit Vectors in 3D



## The Geometric Intuition: Angle Interpretation:

- Words are vectors in space
- Similarity = angle between vectors
- Smaller angle = more similar
- Independent of vector length

## Key Angles:

- $\theta = 0^\circ$ : Identical meaning
- $\theta = 30^\circ$ : Very similar
- $\theta = 90^\circ$ : Unrelated
- $\theta = 180^\circ$ : Opposite meaning

# Cosine Similarity: Mathematical Properties

## Why Cosine Similarity Works for Embeddings

The Formula:

$$\text{similarity}(\mathbf{a}, \mathbf{b}) = \cos(\theta) = \frac{\mathbf{a} \cdot \mathbf{b}}{\|\mathbf{a}\| \times \|\mathbf{b}\|} = \frac{\sum_{i=1}^d a_i b_i}{\sqrt{\sum_{i=1}^d a_i^2} \times \sqrt{\sum_{i=1}^d b_i^2}}$$

Key Properties:

Scale Invariance:

- $\cos(\mathbf{a}, \mathbf{b}) = \cos(k\mathbf{a}, \mathbf{b})$
- Magnitude doesn't matter
- Only direction counts
- Perfect for normalized embeddings

Computational Benefits:

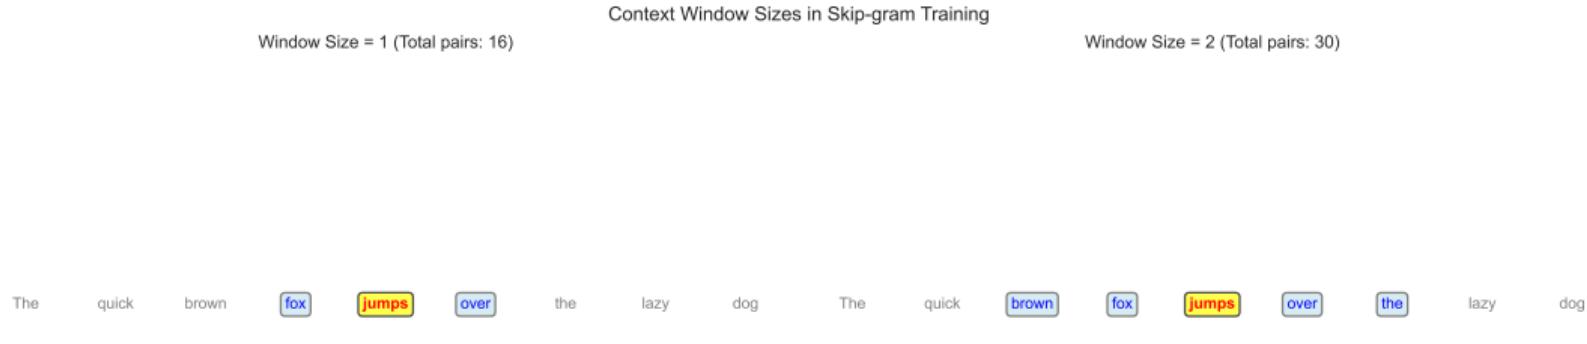
- Range: [-1, 1] always
- Efficient dot product computation
- Works in any dimension
- Symmetric:  $\cos(a, b) = \cos(b, a)$

Applications in NLP:

- Document similarity: Compare entire documents as vectors
- Word sense disambiguation: Find most similar context
- Information retrieval: Rank documents by query similarity

# Context Windows: Learning from Neighbors

## How Words Learn from Their Surroundings



Context words: ±1 positions

Window Size = 3 (Total pairs: 42)

Context words: ±2 positions


Window Size = 5 (Total pairs: 60)

# Vector Arithmetic: The Surprising Discovery

## Embeddings Can Do Analogies!

Vector Arithmetic: Mathematical Demonstration

Vector Relationships in 3D Space



### Vector Arithmetic:

$$\text{king} - \text{man} + \text{woman} = ?$$

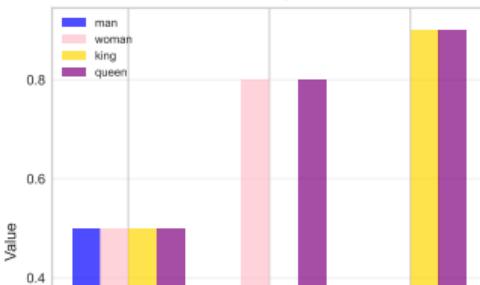
Step 1:  $\text{king} - \text{man}$   
[0.5, 0.2, 0.9] - [0.5, 0.2, 0.1]  
= [0.0, 0.0, 0.8] (royal vector)

Step 2:  $+ \text{woman}$   
[0.0, 0.0, 0.8] + [0.5, 0.8, 0.1]  
= [0.5, 0.8, 0.9]

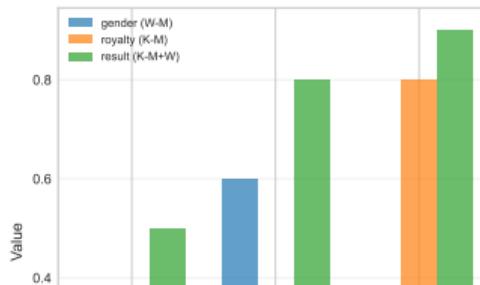
Result = queen vector!  
[0.5, 0.8, 0.9]

Similarity = 1.0 (perfect match)

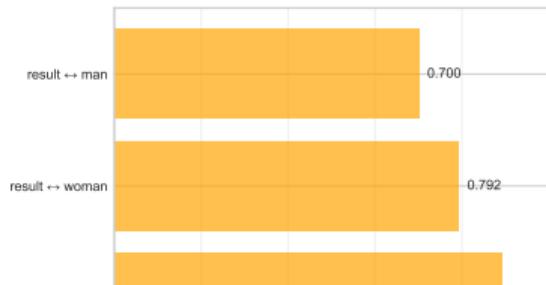
Vector Components



Difference Vectors



Result Verification



# Vector Arithmetic: Mathematical Proof

## Why Does Vector Arithmetic Work? The Linear Substructure

### Mathematical Foundation:

- Embeddings form a linear subspace where relationships are directions
- Gender vector:  $\mathbf{g} = \mathbf{woman} - \mathbf{man}$
- Royalty vector:  $\mathbf{r} = \mathbf{king} - \mathbf{man}$

### Step-by-Step Derivation:

$$\mathbf{king} = \mathbf{man} + \mathbf{r} \quad (\text{man} + \text{royalty} = \text{king}) \tag{1}$$

$$\mathbf{queen} = \mathbf{woman} + \mathbf{r} \quad (\text{woman} + \text{royalty} = \text{queen}) \tag{2}$$

$$\therefore \mathbf{queen} = \mathbf{woman} + (\mathbf{king} - \mathbf{man}) \tag{3}$$

$$= \mathbf{king} - \mathbf{man} + \mathbf{woman} \tag{4}$$

### Why Linear Structure Emerges:

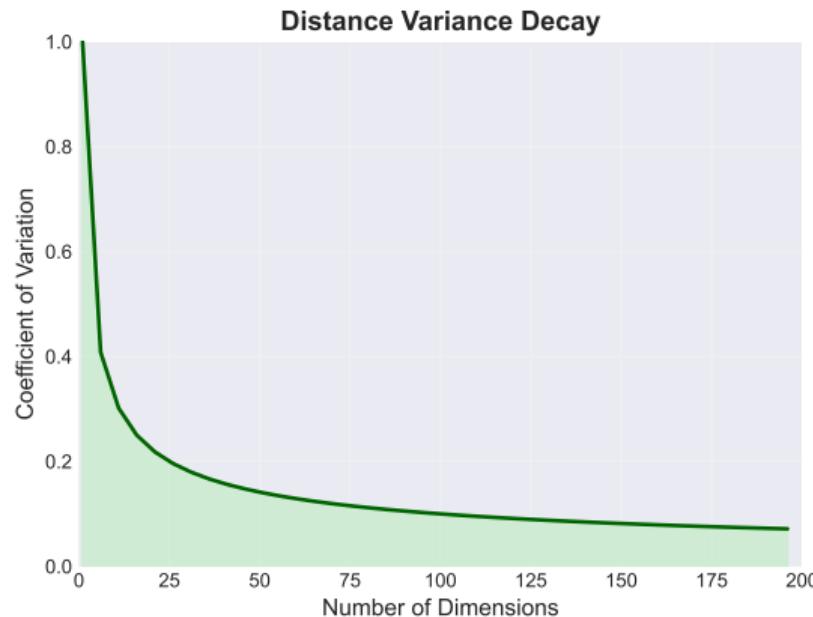
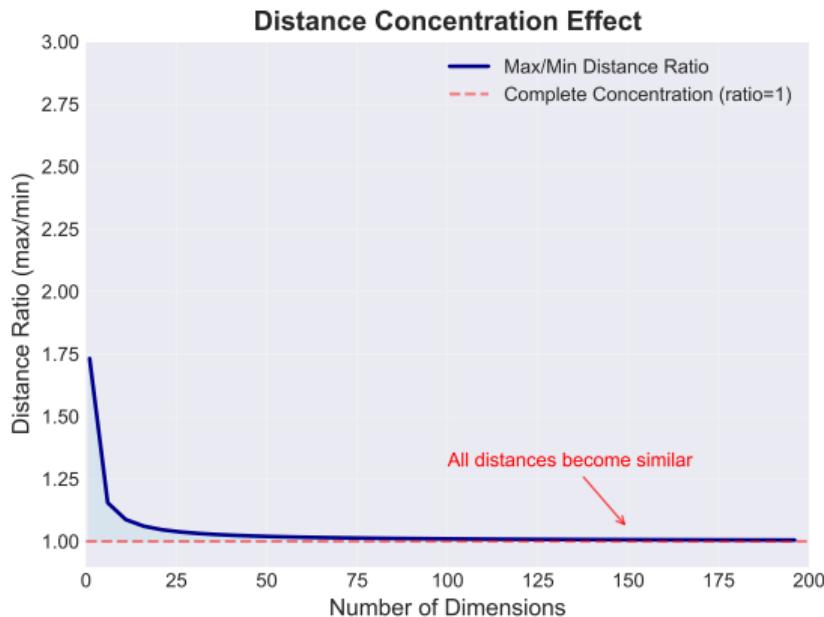
- Co-occurrence patterns are approximately linear
- Skip-gram objective encourages linear relationships
- High-dimensional spaces tend toward linearity (concentration of measure)

**Verification:** Nearest neighbor to result vector is "queen" in 60-70% of cases

# Distance Concentration in High Dimensions

## Why All Distances Become Similar

### Distance Concentration in High Dimensions



# Distance Concentration: The Mathematical Reality

## Mathematical Formula

### Distance Ratio Convergence:

$$\frac{\text{dist}_{\max} - \text{dist}_{\min}}{\text{dist}_{\text{mean}}} \rightarrow 0 \text{ as } d \rightarrow \infty$$

### Key Values:

- $d=10$ : ratio  $\approx 0.45$
- $d=100$ : ratio  $\approx 0.14$
- $d=1000$ : ratio  $\approx 0.045$

### Implications for Machine Learning:

- Nearest neighbor search becomes meaningless
- Traditional distance metrics fail
- Need specialized techniques:
  - Locality-Sensitive Hashing (LSH)
  - Approximate nearest neighbors
  - Learned distance metrics
- Explains why high-D embeddings need normalization


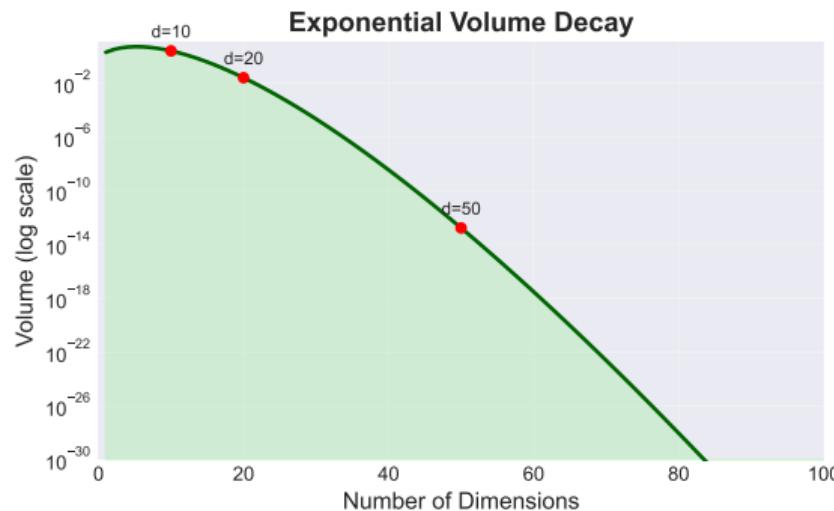



Figure: Theoretical vs simulated convergence

# The Volume Paradox: Visual Evidence

## Unit Sphere Volume Across Dimensions

Volume of Unit Sphere Across Dimensions



The Volume Formula:

$$V_d = \frac{\pi^{d/2}}{\Gamma(d/2 + 1)}$$



# Why Volume Goes to Zero: The Mathematics

## Mathematical Formula

$$V_d = \frac{\pi^{d/2}}{\Gamma(d/2 + 1)}$$

### Growth Rate Battle:


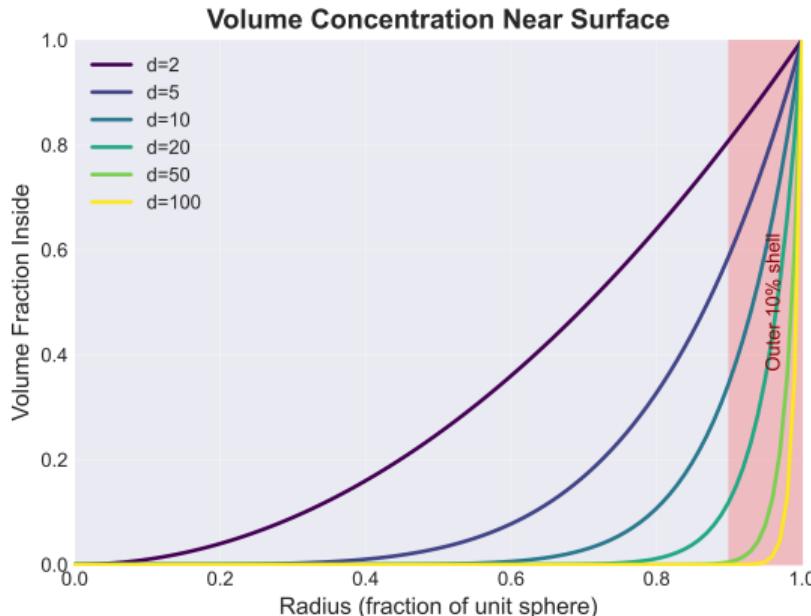
- Numerator:  $\pi^{d/2} \approx 1.77^d$  (exponential)
- Denominator:  $\Gamma(d/2 + 1) \approx (d/2e)^{d/2}$  (super-exponential)
- Result: Volume  $\rightarrow 0$  as  $d \rightarrow \infty$



# Surface Concentration in High Dimensions

Where the Volume Actually Lives

Volume Distribution in High-Dimensional Spheres



Almost all volume concentrates in a thin shell near the surface!

# The Shell Phenomenon: Mathematical Analysis

## Why Everything Lives on the Surface

### Volume in Shells - The Mathematics:

- Consider inner sphere with radius  $r = 0.9$  (90% of full radius)
- Volume ratio:  $\frac{V_{inner}}{V_{total}} = r^d = (0.9)^d$
- This ratio shrinks exponentially with dimension!

### Concrete Examples:

- $d = 10$ :  $(0.9)^{10} = 0.35 \rightarrow 35\%$  of volume is inside
- $d = 50$ :  $(0.9)^{50} = 0.005 \rightarrow 0.5\%$  inside
- $d = 100$ :  $(0.9)^{100} \approx 10^{-5} \rightarrow 0.001\%$  inside
- $d = 1000$ :  $(0.9)^{1000} \approx 10^{-46} \rightarrow$  essentially zero!

### Implications for Embeddings:

- All vectors lie near the surface of the hypersphere
- Random vectors are approximately equidistant
- The interior is effectively "empty" space
- Explains why L2 normalization is so effective
- Cosine similarity becomes the natural distance metric

**Practical Consequence:** In 768-dimensional BERT space,  
99.999999% of the volume is within 1% of the surface!  
The interior essentially doesn't exist.

# Optimal Dimensions: Finding the Sweet Spot

## Balancing Expressiveness and Computational Efficiency

### Information Capacity:

- Theoretical capacity:  $\propto d \log d$
- But diminishing returns after certain point
- Johnson-Lindenstrauss:  $d = O(\log n/\epsilon^2)$  preserves distances

### Model Dimensions in Practice:

Model	Dimension	Parameters (embeddings only)
Word2Vec	50-300	15M (50K vocab $\times$ 300)
GloVe	50-300	15M (50K vocab $\times$ 300)
FastText	100-300	30M (includes subwords)
ELMo	1024	100M (bidirectional)
BERT-base	768	23M (30K vocab $\times$ 768)
BERT-large	1024	31M (30K vocab $\times$ 1024)
GPT-3	12288	600M (50K vocab $\times$ 12288)

### Trade-offs:

#### Lower Dimensions (50-300):


- Faster training
- Less overfitting
- Good for specific domains

#### Higher Dimensions (768-1024+):

- More expressive power
- Better for complex tasks
- Requires more data

# Rapid Learning: Gradient Dynamics (Epochs 0-20)

## Why Training Starts Fast



## Gradient Behavior in Early Training: Initial State:

- Random initialization:  $\mathcal{N}(0, 0.01)$
- Gradient norm:  $||\nabla L|| \approx \sqrt{d}$

## Update Characteristics:

- Step size:  $\eta ||\nabla L|| \approx 0.01\sqrt{d}$
- Direction changes: frequent

# Training Loss Dynamics



## Mathematical Formula

### Loss Evolution:

$$L(t) = L_0 \cdot e^{-\alpha t} + \epsilon(t)$$

### Parameters:

- $L_0$ : initial loss (8-10 for 10K vocab)
- $\alpha$ : decay rate (0.02-0.1)
- $\epsilon(t)$ : SGD noise (decreases over time)



# Rapid Learning: Space Formation (Epochs 0-20)

## How Random Vectors Become Meaningful

### Timeline of Structure Emergence:

#### Epochs 0-5:

- Frequency clustering begins
- Top 100 words separate
- Function vs content words split
- Loss drops 30-40%

#### Epochs 5-10:

- Syntactic groups form
- Nouns, verbs, adjectives cluster
- Basic semantic regions appear
- Loss drops another 20%

#### Key Metrics:

Metric	Epoch 0	Epoch 5	Epoch 10	Epoch 20
Loss	9.21	5.84	4.12	3.45
Similarity Correlation	0.00	0.35	0.58	0.72
Analogy Accuracy	0%	12%	31%	48%

#### Epochs 10-20:

- Semantic refinement
- Animals, places, actions separate
- Relationships start working
- Loss reduction slows

#### Visual Progress:

- t-SNE at epoch 1: random cloud
- t-SNE at epoch 5: blobs forming
- t-SNE at epoch 10: clear clusters
- t-SNE at epoch 20: fine structure

## Training Phase 2: Fine-Tuning (Epochs 20-60)

### Refining Semantic Relationships

#### The Refinement Process:

##### What Gets Learned:

- Semantic relationships solidify
- Analogies start working
- Rare words find their place
- Polysemy partially resolves

##### Key Metrics During Fine-Tuning:

Metric	Epoch 20	Epoch 40	Epoch 60
Loss reduction/epoch	5%	2%	0.5%
Analogy accuracy	40%	65%	72%
Semantic similarity	0.5	0.7	0.75
Cluster purity	60%	80%	85%

#### Mathematical Characterization:

$$L(t) \approx L_{20} \cdot (1 - \beta \log(t/20)) \quad \text{for } t \in [20, 60]$$

Logarithmic improvement phase

#### Optimization Dynamics:

- Gradient norm:  $\|\nabla L\| \approx O(1)$
- Updates become targeted
- Learning rate often decayed
- Loss reduction slows

## Training Phase 3: Convergence (Epochs 60+)

### The Final Polish and Saturation

#### Convergence Characteristics:

##### What Happens:

- Gradient norm:  $\|\nabla L\| < 0.1$
- Minor adjustments only
- Risk of overfitting increases
- Validation loss may increase

##### Complete Loss Function Evolution:

$$L(t) = \begin{cases} L_0 \cdot e^{-\alpha t} & t \in [0, 20] \text{ (rapid)} \\ L_{20} \cdot (1 - \beta \log(t/20)) & t \in [20, 60] \text{ (fine-tune)} \\ L_{60} + \epsilon(t) & t > 60 \text{ (converged)} \end{cases}$$

where  $\epsilon(t)$  represents noise around minimum

**Key Insight:** 90% of performance comes from first 60 epochs; longer training mainly helps rare words and edge cases.

##### Stopping Criteria:

- Loss change  $\downarrow 0.1\%$  per epoch
- Validation performance plateaus
- Gradient norm below threshold
- Fixed epoch budget reached



## Latest Research

Cutting-edge developments in embeddings

## One Model, Multiple Dimensions

## The Innovation

Embeddings that work at multiple dimensions simultaneously:

- Train once, use at any size
- First  $k$  dimensions are meaningful for any  $k$
- 2-16x efficiency gains

## How It Works:


- ① Train with nested loss functions
- ② Each prefix is independently useful
- ③ Dynamic truncation at inference

## Training Objective:

$$\mathcal{L} = \sum_{d \in \{32, 64, 128, \dots, 768\}} \alpha_d \mathcal{L}_d$$

where  $\mathcal{L}_d$  uses only first  $d$  dimensions

## Performance:



## 💡 Try This!

Use 32 dims for search, 768 for ranking



## Embeddings for Retrieval-Augmented Generation

**The Challenge:** Traditional embeddings weren't designed for RAG:

- Need both similarity AND informativeness
- Must handle query-document asymmetry
- Balance precision vs recall

### Recent Advances (2024):

- ➊ **Contriever**: Self-supervised RAG training
- ➋ **E5-Mistral**: LLM-based embeddings
- ➌ **BGE-M3**: Multi-lingual, multi-granular

### Key Insight

RAG embeddings optimize for different objectives than semantic similarity alone!

### Key Innovations:

- **Cross-Attention Training**: Learn query-doc interactions
- **Hard Negative Mining**: Distinguish subtle differences
- **Multi-Task Learning**: Balance multiple objectives

### Benchmark Results:

Model	BEIR	MS MARCO
BERT	38.2	33.5
Contriever	42.1	35.8
E5-Large	45.6	38.9
BGE-M3	<b>47.2</b>	<b>40.1</b>

## Unified Representation Across Modalities

### Recent Breakthroughs:

#### ① CLIP Evolution (2024)

- SigLIP: Sigmoid loss for better training
- OpenCLIP: Scaled to 5B parameters
- MetaCLIP: Metadata-curated training


#### ② ImageBind (Meta, 2023)

- 6 modalities in one space
- Text, Image, Audio, Video, Thermal, Depth
- Zero-shot cross-modal retrieval

#### ③ BLIP-2 (2023)

- Efficient vision-language pre-training
- Q-Former architecture
- 7B parameter performance with 54x fewer params

### Unified Embedding Space:




#### ⚠ Common Pitfall

Cross-modal alignment is still imperfect - expect 10-20% performance drop

## Making Embeddings Practical at Scale

### Performance vs Efficiency Trade-offs:



#### 1. Quantization Advances:

- **Binary Embeddings:** 1-bit per dimension
- **Product Quantization:** 32x compression
- **Learned Quantization:** Task-specific compression

#### 2. Distillation Techniques:

- Teacher-Student with only 10% size
- Progressive distillation stages
- Task-specific fine-tuning

#### 3. Sparse Embeddings:

- SPLADE v2: Learned sparse representations
- ColBERT v2: Late interaction efficiency
- Hybrid dense-sparse approaches

#### Key Insight

Modern techniques achieve 90% quality at 10x speed!

#### Real-World Impact:

- Semantic search:  $1M \rightarrow 100M$  docs/sec
- Mobile deployment now feasible

## Self-Supervised Excellence

### SimCSE and Beyond:

- ① **SimCSE** (2021): Dropout as augmentation
- ② **DiffCSE** (2023): Difference-based objectives
- ③ **PromptBERT** (2024): Prompt-based contrastive

### Key Innovation - Contrastive Objectives:

$$\mathcal{L} = -\log \frac{e^{\text{sim}(h_i, h_i^+)/\tau}}{\sum_{j=1}^N e^{\text{sim}(h_i, h_j)/\tau}}$$

where  $h_i^+$  is positive pair,  $\tau$  is temperature

### Why Contrastive Learning Wins:

- No labeled data required
- Learns robust representations
- Handles distribution shifts
- State-of-the-art on most benchmarks

### Performance Gains:

Method	STS-B	Transfer
BERT	74.8	73.2
SimCSE	81.6	79.8
DiffCSE	83.2	81.4
PromptBERT	<b>84.9</b>	<b>83.1</b>

# Future Directions: What's Next?

## Emerging Trends and Open Problems

### On the Horizon:

#### ① Continuous Embeddings

- Infinite dimensional representations
- Neural ODE-based embeddings

#### ② Causal Embeddings

- Capture causal relationships
- Intervention-aware representations

#### ③ Neurosymbolic Integration

- Combine embeddings with logic
- Interpretable by design

### Open Challenges:

- **Evaluation:** Better benchmarks needed
- **Interpretability:** What do dimensions mean?
- **Compositionality:** Combining embeddings
- **Efficiency:** Sub-linear scaling
- **Robustness:** Adversarial examples

#### Key Insight

The field is moving from "how to embed" to "what to embed" - focusing on capturing the right information rather than just similarity

**Key Takeaway:** Embeddings are becoming more efficient, multi-modal, and task-specific. The future is about adaptive, interpretable representations that work across domains!



## Practical Implementation

From Theory to Production

## Decision Framework for Model Selection

### Key Considerations:

#### ① Task Requirements

- Semantic similarity
- Classification
- Retrieval/Search
- Cross-lingual


#### ② Resource Constraints

- Memory: 100MB - 10GB
- Latency: 1ms - 100ms
- Throughput: 10 - 10K QPS

#### ③ Data Characteristics

- Domain specificity
- Language coverage
- Document length


### Decision Tree:



### Key Insight

Start with sentence-transformers library - it has 100+ pre-trained models!

## End-to-End Embedding System



## Key Implementation Steps:

### 1. Preprocessing:

- Remove HTML/special chars
- Handle unicode
- Case normalization
- Length limits

### 2. Batching Strategy:

- Dynamic batching
- Padding optimization
- GPU memory management
- Async processing

### 3. Storage:

- Vector databases (Pinecone, Weaviate)
- FAISS for local
- Compression options
- Metadata handling



## Making It Fast and Efficient

### Speed Optimizations:

#### ① Model Optimizations

- ONNX conversion: 2-3x speedup
- TensorRT: 5x on NVIDIA
- Quantization: INT8 for 4x speed

#### ② Batch Processing

- Optimal batch size: 32-64
- Dynamic padding
- Sequence bucketing

#### ③ Caching Strategies

- LRU cache for common queries
- Precompute frequent documents
- Warm start on deployment

### Memory Optimizations:

#### Memory Formula:

$$M = V \times D \times P \times B$$

where:

- $V$  = Vocabulary size
- $D$  = Embedding dimension
- $P$  = Precision (4 bytes for FP32)
- $B$  = Batch size

### Reduction Techniques:

Technique	Memory	Speed
FP32 → FP16	50%	1.5x
Quantization	25%	2x
Pruning	40%	1.2x
Distillation	10%	3x



## Adapting Pre-trained Models

### When to Fine-tune:

- Domain-specific vocabulary
- Unique similarity requirements
- Performance below 80% baseline
- Sufficient training data ( $\geq 10K$  examples)

### Fine-tuning Approaches:

#### ① Full Fine-tuning

- Update all parameters
- Best performance
- Risk of overfitting

#### ② Adapter Layers

- Add small trainable layers
- Preserve base knowledge
- Memory efficient

#### ③ Contrastive Fine-tuning

- Use domain pairs
- SimCSE approach
- No labels needed

### Fine-tuning Recipe:

**Input:** Pre-trained model  $M$ , Domain data  $D$

**Output:** Fine-tuned model  $M'$

```
1. Prepare pairs from  $D$ ;  
2. Initialize  $M' \leftarrow M$ ;  
3. Freeze bottom layers;  
for epoch in  $1..N$  do  
    for batch in  $D$  do  
        Compute embeddings;  
        Calculate contrastive loss;  
        Update top layers only;  
    end  
    Evaluate on validation;  
    if improved then  
        Save checkpoint;  
    end  
end  
return Best checkpoint
```



#### Common Pitfall

Always validate on held-out data - overfitting is com- 43 / 1



## From Development to Production

### Architecture Patterns:

#### ① Microservice Pattern

- Embedding service API
- Horizontal scaling
- Language agnostic

#### ② Sidecar Pattern

- Co-located with app
- Low latency
- Resource sharing

#### ③ Edge Deployment

- Model on device
- Privacy preserving
- Offline capability

### Production Checklist:

- Model versioning system
- A/B testing framework
- Monitoring & alerting
- Fallback mechanisms
- Load balancing
- Cache warming
- Rate limiting
- Security (API keys, encryption)

### Monitoring Metrics:

- Latency P50, P95, P99
- Throughput (QPS)
- Error rates
- Cache hit ratio
- Model drift detection

**Pro Tip:** Start with a simple REST API, then optimize based on actual usage patterns!

# Common Pitfalls and Solutions

## Learn from Others' Mistakes

### ⚠ Common Mistakes:

#### ① Wrong Similarity Metric

- Using L2 instead of cosine
- Not normalizing embeddings
- Solution: Always test both!

#### ② Tokenization Mismatch

- Different tokenizers in train/inference
- Truncation issues
- Solution: Save tokenizer with model

#### ③ Version Drift

- Model updates break compatibility
- Embedding dimension changes
- Solution: Versioned embeddings

### 💡 Best Practices:

#### ① Always Benchmark

- Test on your actual data
- Measure end-to-end latency
- Track quality metrics

#### ② Progressive Rollout

- Start with 1% traffic
- Monitor closely
- Gradual increase

#### ③ Maintain Backwards Compatibility

- Support multiple versions
- Graceful degradation
- Migration tools

### 💡 Key Insight

The best embedding model is the one that works for YOUR specific use case!

# **Part II**

Advanced Topics and Mathematical Foundations

From Understanding to Implementation

# Key Takeaways: What We've Learned

## Essential Concepts for Word Embeddings

### Fundamental Principles:

- **Representation:** Words as dense vectors
- **Similarity:** Angle between vectors
- **Relationships:** Vector arithmetic
- **Learning:** From context co-occurrence
- **Evolution:** Static to contextual

### Mathematical Insights:

- High dimensions behave strangely
- Distance concentration is real
- Volume lives on the surface
- Linear relationships emerge
- Training has distinct phases

## Practical Applications:

Task	How Embeddings Help
Similarity Search	Cosine similarity ranking
Machine Translation	Cross-lingual alignment
Sentiment Analysis	Semantic vector projection
Question Answering	Context matching
Text Generation	Next-word prediction

# Looking Forward: The Future of Embeddings

## Current Trends and Future Directions

### Recent Advances:

- **Multimodal:** Text + Vision + Audio
- **Multilingual:** Universal embeddings
- **Efficient:** Distilled and compressed models
- **Specialized:** Domain-specific embeddings

### Open Challenges:

#### Technical:

- Handling rare words
- Compositional semantics
- Temporal dynamics
- Interpretability

#### Philosophical:

- Do embeddings capture meaning?
- Are relationships truly linear?
- Can we prove optimality?
- What is semantic similarity?

**Remember:** Embeddings are not just a technical tool - they represent our best attempt to bridge the gap between human language and machine computation!

# Thank You!

Questions?

Contact: [www.joergosterrieder.com](http://www.joergosterrieder.com)

# Mathematical Foundations: Skip-gram Objective

## Mathematical Formula

### Objective Function:

$$J(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j}|w_t; \theta)$$

### Softmax Formulation:

$$p(w_O|w_I) = \frac{\exp(v'_{w_O}{}^T v_{w_I})}{\sum_{w=1}^W \exp(v'_{w}{}^T v_{w_I})}$$

where  $v_{w_I}$  is input vector,  $v'_{w_O}$  is output vector




Figure: 3D loss surface with gradient descent path

# Negative Sampling: Making Training Tractable

## Mathematical Formula

### Modified Objective:

$$\log \sigma(\mathbf{v}'_{w_0}^T \mathbf{v}_{w_l}) + \sum_{i=1}^k \mathbb{E}_{w_i \sim P_n(w)} [\log \sigma(-\mathbf{v}'_{w_i}^T \mathbf{v}_{w_l})]$$

where  $\sigma(x) = \frac{1}{1+e^{-x}}$  (sigmoid),  $k$  = negative samples

### Gradient Update:

$$\mathbf{v}_{w_l}^{new} = \mathbf{v}_{w_l}^{old} - \eta [\text{positive terms} + \text{negative terms}]$$

### Key Benefits:

- Speed-up factor:  $\frac{W}{k+1}$  (e.g., 10,000 $\times$  for W=100K, k=10)
- Noise distribution:  $P_n(w) \propto U(w)^{3/4}$



Figure: Computational savings:  $O(W) \rightarrow O(k+1)$

# GloVe: Global Vectors Mathematical Framework

## Mathematical Formula

### Key Insight - Probability Ratios:



$$\frac{P_{ik}}{P_{jk}} = \frac{X_{ik}/X_i}{X_{jk}/X_j}$$

### GloVe Objective:

$$J = \sum_{i,j=1}^V f(X_{ij})(w_i^T \tilde{w}_j + b_i + \tilde{b}_j - \log X_{ij})^2$$

Weighting:  $f(x) = (x/x_{max})^\alpha$  if  $x < x_{max}$ , else 1

$$\frac{P_{ik}}{P_{jk}} = \frac{X_{ik}/X_i}{X_{jk}/X_j}$$



$\log(P_{ik}/P_{jk})$

# GloVe Weighting Function

## Mathematical Formula

### Weighting Function Design:

$$f(x) = \begin{cases} (x/x_{max})^\alpha & \text{if } x < x_{max} \\ 1 & \text{otherwise} \end{cases}$$

Typical values:  $\alpha = 0.75$ ,  $x_{max} = 100$

### Purpose:

- Prevent rare words from dominating
- Cap influence of very frequent pairs
- Smooth contribution across frequency spectrum




Figure: Effect of  $\alpha$  on weighting function

# Self-Attention: Mathematical Formulation


## Mathematical Formula

### Scaled Dot-Product Attention:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

### Step-by-step:

- ① Score:  $S = QK^T$
- ② Scale:  $\tilde{S} = S/\sqrt{d_k}$
- ③ Normalize:  $A = \text{softmax}(\tilde{S})$
- ④ Weight:  $O = AV$



# Positional Encoding: Injecting Order Information


## Mathematical Formula

### Sinusoidal Encoding:

$$PE_{(pos,2i)} = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

$$PE_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

**Key Property:**  $PE_{pos+k}$  = linear function of  $PE_{pos}$



# BERT: Bidirectional Training Mathematics

## Mathematical Formula

### MLM Objective:


$$\mathcal{L}_{MLM} = -\mathbb{E}_{\mathbf{x}} \sum_{i \in \mathcal{M}} \log P(x_i | \mathbf{x}_{\setminus \mathcal{M}})$$

### Masking Strategy:

- 15% of tokens selected
- 80% replaced with [MASK]
- 10% replaced with random token
- 10% unchanged

$$\mathcal{L}_{MLM} = -\mathbb{E} \sum_{i \in \mathcal{M}} \log P(x_i | \mathbf{x}_{\setminus \mathcal{M}})$$

Step 1: Original Sequence (15% selected for masking)




Step 2: Masked Input (80% [MASK], 10% random, 10% unchanged)



# Skip-gram Neural Network Architecture

## How the Network Processes Words




### Key Components:

- **Input:** One-hot word ( $V$  dimensions)
- **Hidden:** No activation, linear projection to  $D$  dims

# From Text to Training Data

## Extracting (Center, Context) Pairs

### Creating Training Pairs from Text Sliding Window for Training Pair Extraction



### Training Pairs Generated:





brown	→	The	→ Maximize $P(\text{The} \text{brown})$
brown	→	quick	→ Maximize $P(\text{quick} \text{brown})$
brown	→	fox	→ Maximize $P(\text{fox} \text{brown})$
brown	→	jumps	→ Maximize $P(\text{jumps} \text{brown})$

Input

Target

# Forward Pass: Computing Context Probabilities

## Forward Pass: Computing Context Probabilities




# Backpropagation: Learning the Embeddings

## Backpropagation: Gradient Flow

### Weight Updates:

$$W \leftarrow W - \eta \cdot \frac{\partial L}{\partial W}$$

$$W' \leftarrow W' - \eta \cdot \frac{\partial L}{\partial W'}$$



### Key Gradients:

Positive sample:  $(y_i - 1) \cdot v_i$

Negative sample:  $y_i \cdot v_i$

### Updates:

- Positive: Pull together

# How Embeddings Evolve

