

Generating Shakespeare Sonnets with N-Gram Language Models

Natural Language Processing - Week 1 Extension

BSc Computer Science

2025

Session Outline

Theory (25 min)

- Shakespeare's sonnets: Structure & style
- N-gram models for poetry
- Special challenges in poetry generation
- Evaluation metrics for generated text

Practice (25 min)

- Text preprocessing for sonnets
- Building n-gram models
- Generation techniques
- Hands-on: Generate your own sonnets

Learning Goals: Apply n-gram models to creative text generation, understand domain-specific challenges

Shakespeare's Sonnets: Literary Context

Historical Background

- 154 sonnets published in 1609
- Written over 20 year period
- Themes: love, beauty, mortality, time
- Revolutionary use of English

Why Study with NLP?

- Rich, structured language patterns
- Fixed poetic form (constraints)
- Historical importance
- Challenging generation task

Sonnet 18

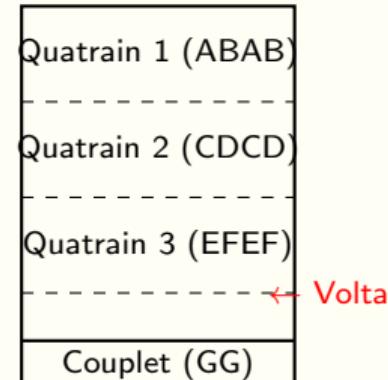
Shall I compare thee to a summer's day?
Thou art more lovely and more temperate:
Rough winds do shake the darling buds of May,
And summer's lease hath all too short a date:
Sometime too hot the eye of heaven shines,
And often is his gold complexion dimm'd;
And every fair from fair sometime declines,
By chance, or nature's changing course untrimm'd;
But thy eternal summer shall not fade
Nor lose possession of that fair thou owest;
Nor shall Death brag thou wander'st in his shade,
When in eternal lines to time thou growest:
So long as men can breathe, or eyes can see,
So long lives this, and this gives life to thee.

Sonnet Structure: The Shakespearean Form

Formal Requirements

- ① **14 lines** exactly
- ② **Iambic pentameter:** 10 syllables per line
 - Pattern: da-DUM da-DUM da-DUM da-DUM da-DUM
 - Example: "Shall I / com-PARE / thee TO / a SUM / mer's DAY?"
- ③ **Rhyme scheme:** ABAB CDCD EFEF GG
- ④ **Volta:** Thematic turn at line 9
- ⑤ **Couplet:** Final two lines summarize/resolve

Structure Visualization



Computational Challenges:

- Maintaining meter while being coherent
- Enforcing rhyme without repetition
- Balancing structure with creativity

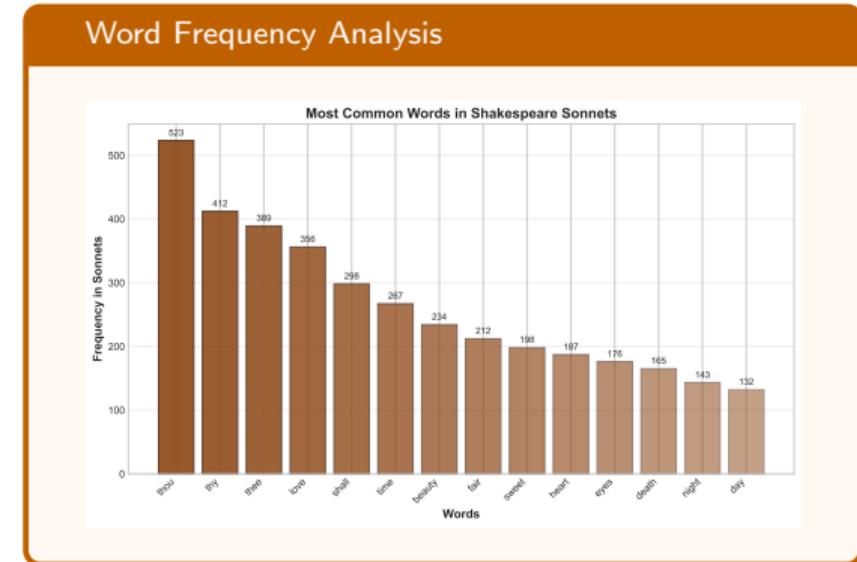
Shakespearean Language Characteristics

Vocabulary Features

- **Archaic forms:** thee, thou, thy, thine, hath, doth
- **Contractions:** o'er, e'er, 'tis, ne'er
- **Inversions:** “Rough winds do shake” vs “Rough winds shake”
- **Rich vocabulary:** 31,000 unique words in complete works

Statistical Properties

- Average sentence length: 12-15 words
- Vocabulary richness: Type-Token Ratio ≈ 0.45
- Most frequent words: thou, thy, thee, love, time



Implications for N-grams:

- Need larger context (3-grams or higher)
- Special handling of archaic forms
- Preserve poetic inversions

N-Gram Models for Poetry Generation

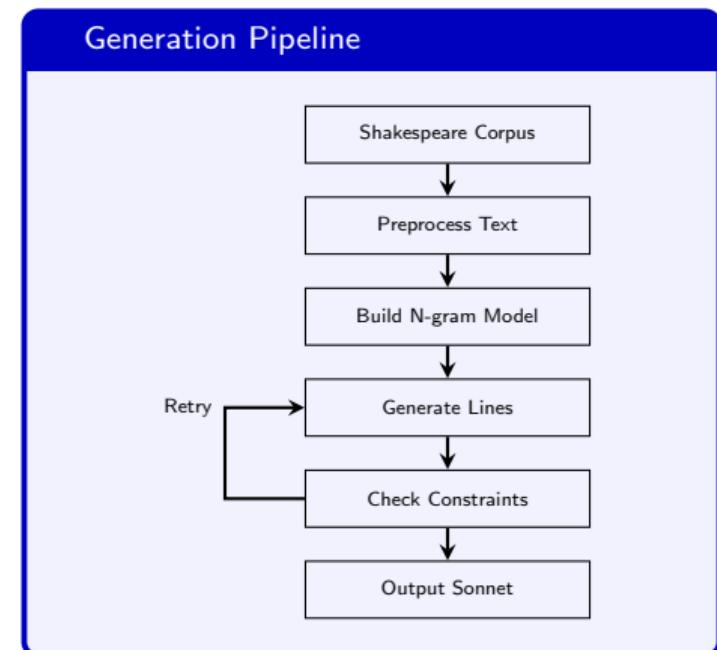
Why N-grams Work for Poetry

- Capture local word dependencies
- Learn stylistic patterns
- Preserve author's "voice"
- Computationally simple

Model Selection	Model	Pros	Cons
	Unigram	Fast	No context
	Bigram	Some context	Limited memory
	Trigram	Good balance	Sparse data
	4-gram+	Rich context	Very sparse

Poetry-Specific Modifications

- Line-aware generation
- Rhyme dictionary integration
- Syllable counting
- Meter checking



Text Preprocessing for Sonnets

Preprocessing Pipeline

```
1 def preprocess_sonnet(text):
2     # 1. Basic cleaning
3     text = text.lower()
4     text = re.sub(r'[\w\s\']+', ' ', text)
5
6     # 2. Handle contractions
7     contractions = {
8         "o'er": "over",
9         "'tis": "it is",
0         "ne'er": "never"
1     }
2     for old, new in contractions.items():
3         text = text.replace(old, new)
4
5     # 3. Tokenization
6     tokens = text.split()
7
8     # 4. Add line markers
9     lines = text.split('\n')
0     processed = []
1     for line in lines:
2         tokens = ['<START>'] + line.split() + ['<END>']
3         processed.extend(tokens)
4
5     return processed
```

Special Considerations

- **Preserve line structure:** Add START/END tokens
- **Handle punctuation:** Keep apostrophes, remove others
- **Case sensitivity:** Lowercase for matching
- **Archaic forms:** Standardize or preserve?

Common Issues

- **Over-preprocessing:** Losing poetic structure
- **Under-preprocessing:** Too many unique tokens
- **Balance:** Preserve style while ensuring generation quality

Building the N-gram Model

Implementation

```
1 from collections import defaultdict, Counter
2
3 def build_ngram_model(tokens, n=3):
4     model = defaultdict(Counter)
5
6     # Create n-grams
7     for i in range(len(tokens) - n):
8         context = tuple(tokens[i:i+n-1])
9         next_word = tokens[i+n-1]
10        model[context][next_word] += 1
11
12    # Convert to probabilities
13    for context in model:
14        total = sum(model[context].values())
15        for word in model[context]:
16            model[context][word] /= total
17
18    return model
19
20
21 # Example usage
22 tokens = preprocess_sonnets(shakespeare_text)
23 bigram_model = build_ngram_model(tokens, n=2)
24 trigram_model = build_ngram_model(tokens, n=3)
```

Model Statistics

- **Corpus size:** 154 sonnets
- **Total tokens:** 17,000
- **Unique tokens:** 3,000
- **Unique bigrams:** 8,000
- **Unique trigrams:** 12,000

Smoothing Techniques

For unseen n-grams:

- **Add-one:** Add 1 to all counts
- **Good-Turing:** Redistribute probability mass
- **Backoff:** Fall back to (n-1)-gram

Generation Techniques

Basic Generation

```
1 def generate_line(model, n=3, max_words=10):
2     # Select good starting context
3     starters = [ctx for ctx in model.keys()
4                  if ctx[0] in ['when', 'shall', 'thy']]
5
6     context = random.choice(starters)
7     line = list(context)
8
9     for _ in range(max_words - n + 1):
10        if context not in model:
11            break
12
13        # Weighted random selection
14        next_words = model[context]
15        words = list(next_words.keys())
16        probs = list(next_words.values())
17
18        next_word = np.random.choice(words, p=probs)
19        line.append(next_word)
20
21        # Update context
22        context = tuple(line[-(n-1):])
23
24        if next_word == '<END>':
25            break
26
27    return ' '.join(line)
```

Advanced Techniques

- **Temperature sampling:**
 - Adjust randomness
 - $P'(w) = \frac{P(w)^{1/T}}{\sum P(w_i)^{1/T}}$
- **Beam search:**
 - Keep top-k candidates
 - Select best complete line
- **Constrained generation:**
 - Force rhyme words
 - Count syllables
 - Check meter

Tip: Start with simple generation, then add constraints gradually

Implementing Rhyme Constraints

Rhyme Detection

```
1 def get_rhyme_sound(word):
2     """Extract rhyme sound (simplified)"""
3     # Real implementation would use CMU Pronouncing Dict
4     endings = {
5         'ay': ['day', 'may', 'say', 'way'],
6         'ight': ['night', 'light', 'sight', 'might'],
7         'ove': ['love', 'dove', 'above'],
8         'ime': ['time', 'rhyme', 'chime']
9     }
10
11    for sound, words in endings.items():
12        if word in words:
13            return sound
14    return None
15
16
17 def generate_rhyming_line(model, rhyme_with):
18     """Generate line that rhymes with given word"""
19     rhyme_sound = get_rhyme_sound(rhyme_with)
20     candidates = []
21
22     for _ in range(100): # Try multiple times
23         line = generate_line(model)
24         last_word = line.split()[-1]
25         if get_rhyme_sound(last_word) == rhyme_sound:
26             candidates.append(line)
27
28
29 return random.choice(candidates) if candidates else None
```

ABAB CDCD EFEF GG Pattern

- ➊ Generate line 1 freely (A)
- ➋ Generate line 2 freely (B)
- ➌ Generate line 3 rhyming with 1 (A)
- ➍ Generate line 4 rhyming with 2 (B)
- ➎ Repeat for next quatrains
- ➏ Final couplet: two rhyming lines

Challenges

- Limited rhyming vocabulary
- Maintaining coherence
- Avoiding forced rhymes
- Balancing rhyme with meaning

Evaluating Generated Sonnets

Quantitative Metrics

- **Perplexity:** How well model predicts next word

$$PPL = 2^{-\frac{1}{N} \sum_{i=1}^N \log_2 P(w_i | context)}$$

- **BLEU Score:** Compare with real sonnets
- **Rhyme accuracy:** % correct rhymes
- **Meter compliance:** % correct rhythm
- **Vocabulary diversity:** Type-Token Ratio

Qualitative Assessment

- Coherence and meaning
- Poetic quality
- Style consistency
- Creativity/originality



Human Evaluation: Best judge of poetic quality

- Turing test: Can readers distinguish from real?

Common Problems and Solutions

Problem 1: Repetitive Output

- **Symptom:** Same phrases repeatedly
- **Cause:** High-probability sequences dominate
- **Solutions:**
 - Temperature sampling
 - Penalty for recently used words
 - Diverse starting contexts

Problem 2: Incoherent Lines

- **Symptom:** Grammatically wrong or nonsensical
- **Cause:** Limited context in bigrams
- **Solutions:**
 - Use trigrams or higher
 - Add syntax checking
 - Post-process filtering

Problem 3: Poor Rhyming

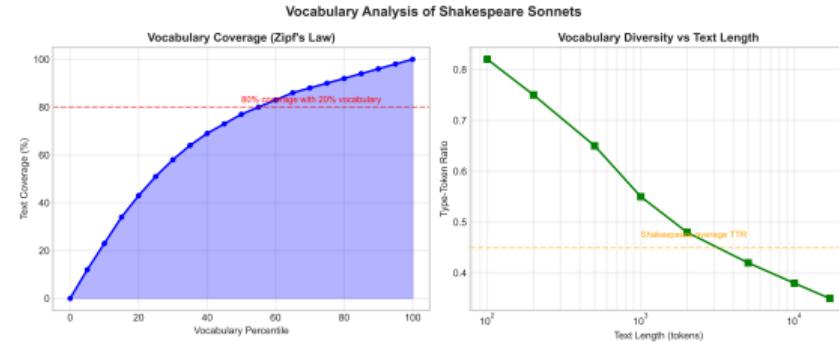
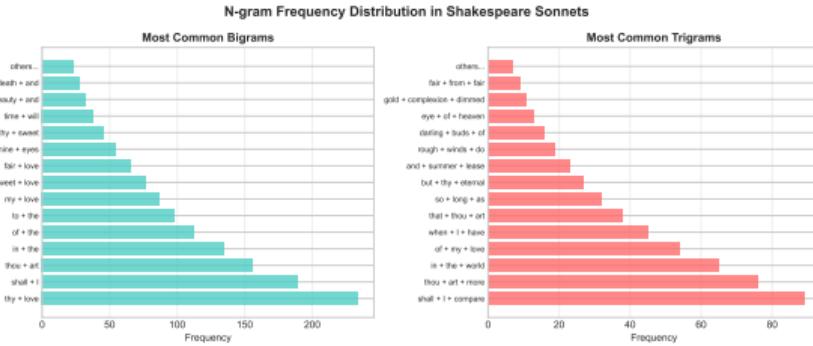
- **Symptom:** Forced or missing rhymes
- **Cause:** Limited rhyming vocabulary
- **Solutions:**
 - Pre-compute rhyme sets
 - Generate multiple candidates
 - Relax exact rhyme requirement

Problem 4: Wrong Meter

- **Symptom:** Lines too long/short
- **Cause:** No syllable counting
- **Solutions:**
 - Add syllable dictionary
 - Generate with constraints
 - Post-generation filtering

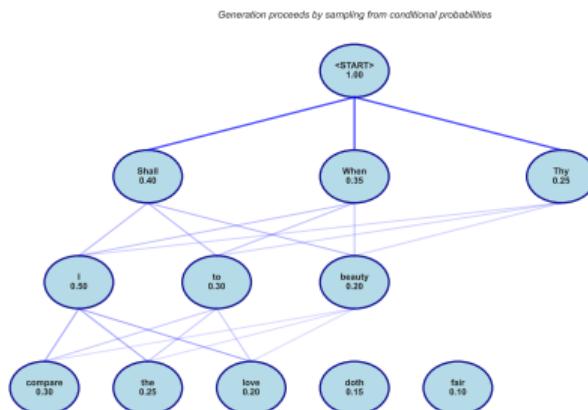
Key Insight: Balance between structure and creativity - too many constraints kill creativity, too few produce nonsense

Visualizing Model Behavior



Key Observations

- Power law distribution of n-grams
- Most contexts have few continuations
- 20% of vocabulary covers 80% of usage
- Trigrams capture more style than bigrams



Hands-On Exercise: Generate Your First Sonnet

Step-by-Step Implementation

```
1 # 1. Load and preprocess sonnets
2 sonnets = load_shakespeare_sonnets()
3 tokens = preprocess_sonnets(sonnets)
4
5 # 2. Build models
6 bigram = build_ngram_model(tokens, n=2)
7 trigram = build_ngram_model(tokens, n=3)
8
9 # 3. Generate a quatrain
10 def generate_quatrain(model):
11     lines = []
12     # Generate ABAB pattern
13     lines.append(generate_line(model)) # A
14     lines.append(generate_line(model)) # B
15     lines.append(generate_rhyming_line(model,
16         get_last_word(lines[0]))) # A
17     lines.append(generate_rhyming_line(model,
18         get_last_word(lines[1]))) # B
19     return lines
20
21 # 4. Complete sonnet
22 quatrain1 = generate_quatrain(trigram)
23 quatrain2 = generate_quatrain(trigram)
24 quatrain3 = generate_quatrain(trigram)
25 couplet = generate_couplets(trigram)
26
27 sonnet = quatrain1 + quatrain2 + quatrain3 + couplet
```

Try These Experiments

① Compare n-gram orders:

- Generate with bigram vs trigram
- Which sounds more Shakespearean?

② Temperature variation:

- T=0.5 (conservative)
- T=1.0 (balanced)
- T=2.0 (creative)

③ Different seeds:

- Start with "When"
- Start with "Love"
- Start with "Time"

④ Hybrid approaches:

- Mix bigram and trigram
- Combine multiple sonnets

Challenge: Generate a sonnet about "artificial intelligence" in Shakespeare's style!

Technical Improvements

- **Character-level models:**
 - Generate at character level
 - Better for archaic forms
 - Harder to train
- **Neural approaches:**
 - RNN/LSTM for better context
 - GPT-style transformers
 - Fine-tuning on sonnets
- **Hybrid models:**
 - N-grams + neural networks
 - Rule-based + statistical
 - Multiple n-gram orders

Creative Applications

- **Style transfer:**
 - Modern text → Shakespearean
 - Mix different poets' styles
- **Interactive generation:**
 - User provides first line
 - Collaborative writing
 - Real-time constraints
- **Educational tools:**
 - Teach sonnet structure
 - Explore vocabulary
 - Understand poetic devices

Research Direction: Controllable generation - specify theme, emotion, or style while maintaining structure

Real vs Generated: Can You Tell?

Sonnet A

When in the chronicle of wasted time
I see descriptions of the fairest wights,
And beauty making beautiful old rhyme
In praise of ladies dead and lovely knights,
Then, in the blazon of sweet beauty's best,
Of hand, of foot, of lip, of eye, of brow,
I see their antique pen would have express'd
Even such a beauty as you master now.

Sonnet B

When to the sessions of sweet silent thought
I summon up remembrance of things past,
Thy beauty's form in table of my heart
And all the world of love was ever lost,
Then can I drown an eye with precious friends
And weep afresh love's long since cancell'd woe,
But if the while I think on thee, dear friends,
All losses are restored and sorrows go.

Which is real Shakespeare?

Real vs Generated: Can You Tell?

Sonnet A

When in the chronicle of wasted time
I see descriptions of the fairest wights,
And beauty making beautiful old rhyme
In praise of ladies dead and lovely knights,
Then, in the blazon of sweet beauty's best,
Of hand, of foot, of lip, of eye, of brow,
I see their antique pen would have express'd
Even such a beauty as you master now.

Sonnet B

When to the sessions of sweet silent thought
I summon up remembrance of things past,
Thy beauty's form in table of my heart
And all the world of love was ever lost,
Then can I drown an eye with precious friends
And weep afresh love's long since cancell'd woe,
But if the while I think on thee, dear friends,
All losses are restored and sorrows go.

Which is real Shakespeare? Answer: A is real (Sonnet 106), B is trigram-generated!

Summary: N-grams Meet the Bard

What We Learned

- N-gram models can capture writing style
- Poetry has unique challenges:
 - Structure (rhyme, meter)
 - Vocabulary (archaic forms)
 - Coherence vs creativity
- Higher-order n-grams better for style
- Constraints improve output quality
- Evaluation is multi-dimensional

Key Takeaways

- ➊ Simple models can produce impressive results
- ➋ Domain knowledge improves generation
- ➌ Balance constraints with creativity
- ➍ Preprocessing matters enormously
- ➎ Multiple evaluation metrics needed

Next Steps: Try the notebook exercises, experiment with different parameters, and generate your own sonnets!

“The code, dear Brutus, is not in our stars, but in our n-grams”

Further Reading

Academic Papers

- Jurafsky & Martin (2024). "Speech and Language Processing", Chapter 3: N-gram Language Models
- Kao & Jurafsky (2012). "A Computational Analysis of Style, Affect, and Imagery in Contemporary Poetry"
- Ghazvininejad et al. (2016). "Generating Topical Poetry"

Online Resources

- <https://www.shakespeareswords.com/> - Shakespeare language database
- <http://www.speech.cs.cmu.edu/cgi-bin/cmudict> - CMU Pronouncing Dictionary
- <https://github.com/topics/poetry-generation> - Code examples

Related Notebooks

- `shakespeare_sonnets_simple_bsc.ipynb` - Basic implementation
- `week01_ngrams.bsc.ipynb` - General n-gram models