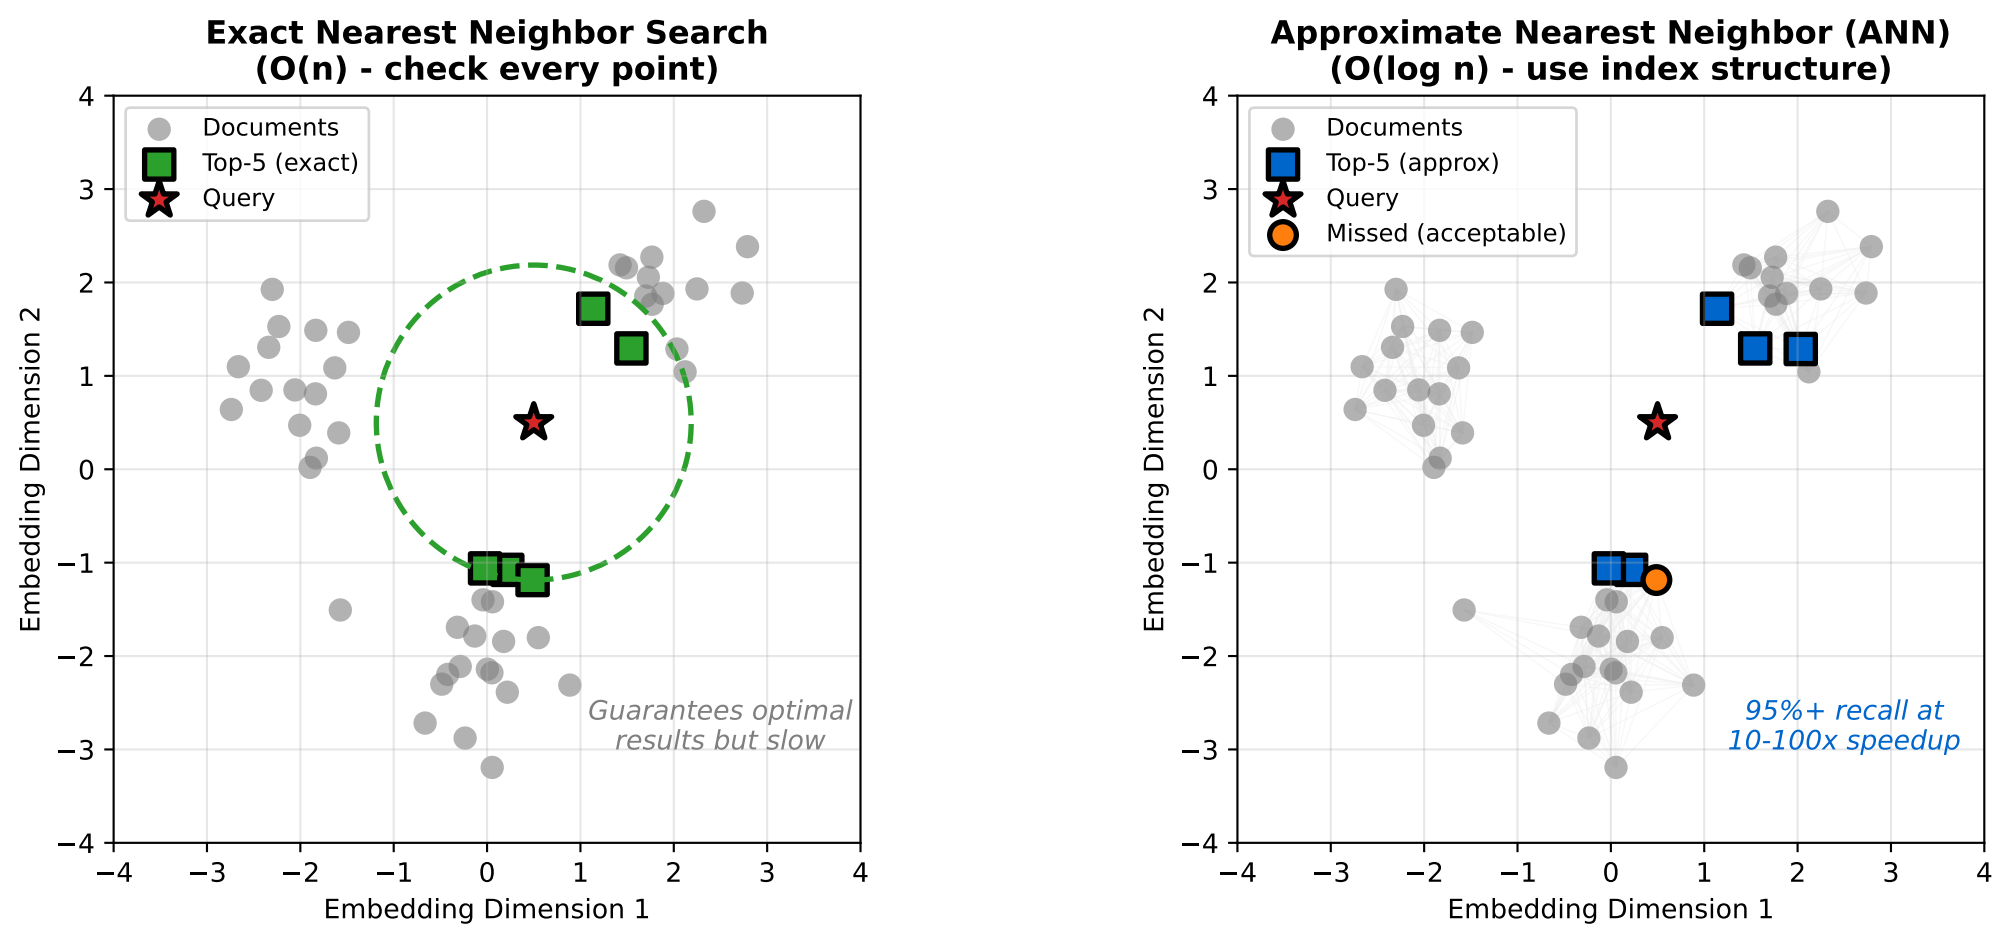
Exact Nearest Neighbor Search
(O(n) - check every point)
Documents
Top-5 (exact)
Query
Guarantees optimal
results but slow
Embedding Dimension 1
Embedding Dimension 2
Approximate Nearest Neighbor (ANN)
(O(log n) - use index structure)
Documents
Top-5 (approx)
Query
Missed (acceptable)
95%+ recall at
10-100x speedup
Embedding Dimension 1
Embedding Dimension 2