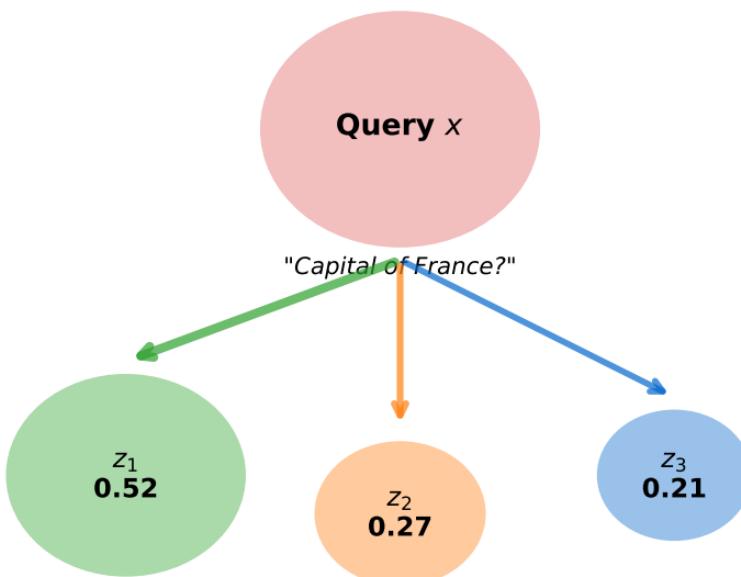


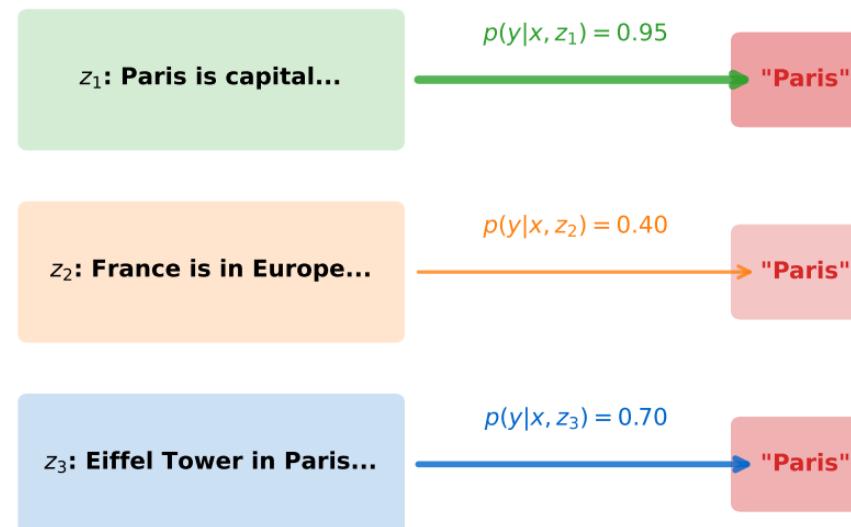
RAG Conditional Probabilities: Visual Intuition

$p(z_i|x)$: Retrieval Probability



Size/arrow = retrieval probability
(how relevant is doc to query?)

$p(y|x, z_i)$: Generation Probability



Arrow thickness = generation probability
(given doc, how likely is answer?)

$$p(y|x) = \sum_i p(z_i|x) \cdot p(y|x, z_i)$$

$$p(y|x) = 0.75$$

$$\begin{aligned} z_3: & 0.21 \times 0.70 \\ & = 0.147 \\ z_2: & 0.27 \times 0.40 \\ & = 0.108 \\ z_1: & 0.52 \times 0.95 \\ & = 0.494 \end{aligned}$$

$$\begin{aligned} & 0.52 \times 0.95 \\ & + 0.27 \times 0.40 \\ & + 0.21 \times 0.70 \\ & = 0.75 \end{aligned}$$

Each document contributes to final answer
weighted by its retrieval probability