Adaptation Strategy Decision Tree

How much labeled data do you have?

- <100 examples
- 100-1K examples
- >1K examples

Prompt Engineering
Fast iteration
No training cost

Few-shot + Prompting
Good balance
Quick results

LoRA Fine-tuning
Best overall
Efficient & effective

Full Fine-tuning
Maximum accuracy
High cost

Instruction Tuning
Generalization
Multi-task

Need >95% accuracy?

- Consider full fine-tuning
- LoRA is sufficient