

# LSTM Primer: Next Word Prediction

## From Autocomplete to Modern Language Models

BSc Level Introduction


September 27, 2025

# The Autocomplete Challenge

You're typing on your phone...

Your Phone Predicts the Next Word

"I love chocolate ice cream but I prefer..."



## Why It's Hard:

- Context can be long
- Words far back still matter
- Grammar rules complex
- Meaning changes with context

## Example Challenges:

- "I love chocolate. But I prefer \_\_\_"
- Need to remember "love/prefer" pattern
- Forget specific "chocolate"
- Context = 7 words back!

## The Goal:

- Predict what word comes next

This is a sequence modeling problem

# What We've Tried So Far: N-grams

## Simple Baseline: Count Word Pairs

N-Gram: Fixed 2-Word Window (Forgets "cat")



### How N-grams Work:

- Look at previous 1-2 words only
- Count what usually comes next
- Pick most common continuation

### Example:

- Saw "I love" 100 times
- Followed by "you": 60 times
- Followed by "chocolate": 30 times
- Followed by "pizza": 10 times
- Predict: "you" (most common)

LSTM: Selective Memory (Remembers "cat")



### Fatal Limitations:

- Only sees 1-2 words back
- Can't handle long context
- No real understanding
- Millions of word combinations

# The Memory Problem

## Reading a Book: What Do You Remember?

### Human Memory:

- You read 200 pages
- Don't remember every word
- Remember: Main plot
- Remember: Key characters
- Forget: Minor details
- Forget: Exact sentences

### The Insight:

- Memory is **selective**
- Keep important information
- Discard irrelevant details
- Update as story progresses

### What We Need in AI:

- Remember relevant past words
- Forget irrelevant words
- Decide what's important NOW
- Update memory as we read

### Example:

*"I grew up in Paris. I learned to speak fluent ..."*

- Remember: "Paris" (8 words back)
- Forget: "grew up" (not needed now)
- Predict: "French"

We need controllable memory

## Recurrent Neural Networks: A Loop of Memory

### The RNN Idea:

- Hidden state = memory
- Update memory at each word
- Use memory to predict next word
- Memory flows through time

### The Math:

$$h_t = \tanh(W_h h_{t-1} + W_x x_t)$$

$$y_t = \text{softmax}(W_y h_t)$$

- $h_t$  = memory at time  $t$
- $x_t$  = current word
- $h_{t-1}$  = previous memory
- $y_t$  = prediction

### The Problem:

When we train (backpropagation):

- Gradient flows backward through time
- Multiplied by same weight matrix
- After 10 steps:  $0.9^{10} = 0.35$
- After 20 steps:  $0.9^{20} = 0.12$
- After 50 steps:  $0.9^{50} = 0.005$

### Vanishing Gradient:

- Signal gets weaker each step
- Can't learn from distant past
- Memory effectively only 5-10 words
- Same problem as n-grams!

We need a gradient highway

## Why RNNs Forget Long-Term Context

### What Happens During Training:

#### Step 1: Make prediction

- Forward pass through network
- Get prediction error

#### Step 2: Compute gradient

- How much to adjust weights?
- Flow gradient backward in time

#### Step 3: Problem appears

- Gradient multiplied at each step
- $\frac{\partial h_t}{\partial h_{t-1}} \approx 0.9$
- After  $n$  steps:  $0.9^n$

### The Numbers:

Steps Back	Gradient Strength
1	0.90 (90%)
5	0.59 (59%)
10	0.35 (35%)
20	0.12 (12%)
50	0.005 (0.5%)


### The Impact:

- Word 50 steps back: Almost no learning
- Network can't learn long-term patterns
- Effectively limited to 5-10 words
- Need 50-100+ word context!

Solution: Create a direct path for gradients

# LSTM Solution: Selective Memory with Gates

## Long Short-Term Memory: Three Gates Control Information Flow



### Traffic Light Analogy:

- **Red Gate (Forget):** What to remove from memory

BSc Level Introduction

### How It Works:

#### Forget Gate: (0 to 1)

- 0.0 = Completely forget
- 1.0 = Keep everything
- Example: 0.9 at period = forget old topic

#### Input Gate: (0 to 1)

- 0.0 = Ignore new word
- 1.0 = Fully store
- Example: 0.95 on "Paris" = remember!

#### Output Gate: (0 to 1)

- 0.0 = Hide memory
- 1.0 = Reveal everything
- Example: 0.8 when predicting = use memory

# Inside the LSTM: Gate Equations

## From Intuition to Mathematics

### The Four Equations:

#### 1. Forget Gate:

$$f_t = \sigma(W_f[h_{t-1}, x_t] + b_f)$$

#### 2. Input Gate:

$$i_t = \sigma(W_i[h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C[h_{t-1}, x_t] + b_C)$$

#### 3. Update Cell State:

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$

#### 4. Output Gate:

$$o_t = \sigma(W_o[h_{t-1}, x_t] + b_o)$$

$$h_t = o_t \odot \tanh(C_t)$$

### Concrete Example:

Input: "I love"

### Values at "love":


- $f_t = 0.3$  (forget 70% of "I")
- $i_t = 0.9$  (strongly store "love")
- $\tilde{C}_t = [0.8, -0.3, 0.5, \dots]$  (new info)
- $C_t = 0.3 \cdot C_{t-1} + 0.9 \cdot \tilde{C}_t$
- $o_t = 0.7$  (reveal 70% of memory)
- $h_t = 0.7 \cdot \tanh(C_t)$

### Key Insight:

- $\sigma$  = sigmoid (0 to 1)
- $\odot$  = element-wise multiply
- $C_t$  = cell state (long-term memory)
- $h_t$  = hidden state (short-term output)

# Training Progression: Watching It Learn

## How LSTM Improves Over Time



### What's Happening:

#### Epoch 1 (Random):

- Gates untrained
- Random predictions
- No pattern learning
- Loss: 8.5

#### Epoch 10 (Bigrams):

- Learns immediate pairs
- "I" → "love" pattern
- Still struggles with context
- Loss: 4.2

#### Epoch 50 (Context):

- Gates start working
- Remembers further back

# Why LSTMs Work: The Gradient Highway

## Comparing RNN vs LSTM Gradient Flow

### RNN Problem:

- Gradient:  $0.9^{10} = 0.35$
- Multiplied at every step
- Exponential decay
- Can't learn long-term



### LSTM Solution:

- Gradient:  $1.0^{10} = 1.0$
- Addition in cell state
- No multiplication!
- Learns 50+ steps back

### The Key Equation:

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$

- Addition (+) not multiplication

## From Theory to Real-World Impact

### Key Takeaways:

#### 1. The Problem:

- Need to remember long context
- RNNs couldn't learn beyond 5-10 words
- Vanishing gradient problem

#### 2. The Solution:

- Three gates control memory
- Cell state = gradient highway
- Additive updates preserve gradients

#### 3. The Impact:

- 50-100+ word context
- Breakthrough in NLP (2015-2018)
- Foundation for modern transformers

### Where LSTMs Are Used:

- Autocomplete: Phone keyboards
- Translation: Google Translate (2016-2019)
- Speech Recognition: Siri, Alexa
- Sentiment Analysis: Product reviews
- Text Generation: Creative writing
- Music Generation: Compose melodies
- Video Captioning: Describe videos

### PyTorch Implementation:

```
import torch.nn as nn
lstm = nn.LSTM(input_size=100,
    hidden_size=256, num_layers=2)
output (h_n, c_n) = lstm(input)
```