

LSTM Networks: A Visual Journey

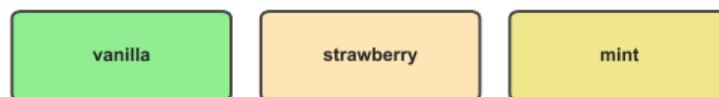
Understanding Long Short-Term Memory Through Charts

Neural Language Processing

The Challenge: Long-Distance Dependencies

Your Phone Predicts the Next Word

"I love chocolate ice cream but I prefer..."



vanilla

strawberry

mint

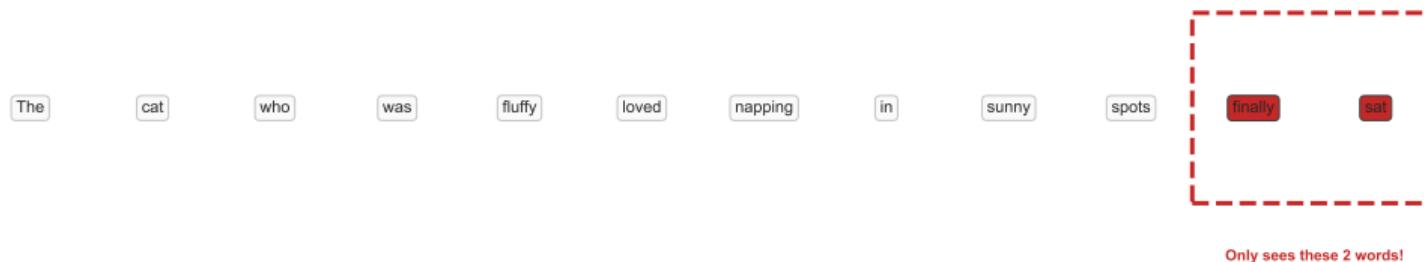
How does it know?
It remembered "chocolate" context!



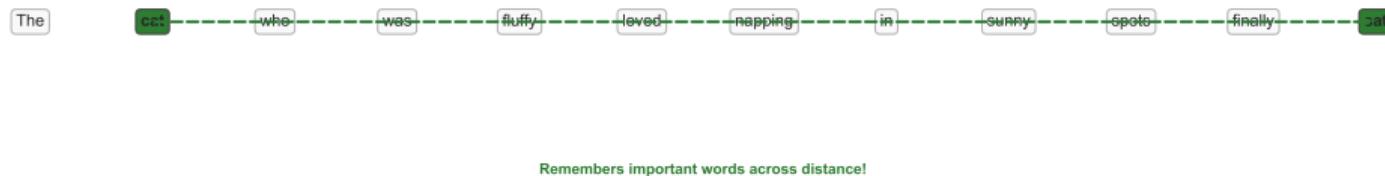
The Problem:

N-grams Can't See Far Enough

N-Gram: Fixed 2-Word Window (Forgets "cat"!)



LSTM: Selective Memory (Remembers "cat"!)



What Humans Remember vs What N-grams Remember

Human Memory:

- + Remembers Paris mentioned earlier
- + Connects Paris → French
- + Forgets irrelevant details
- + Updates memory with new info

N-gram Memory:

- Only sees last 1-3 words
- No connection to Paris
- Can't distinguish important from irrelevant
- Fixed window, can't adapt

Key Insight:

Humans have **selective memory**:

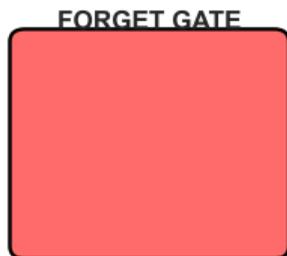
- Keep important info
- Forget irrelevant details
- Update with new context

What We Need:

A memory system with:

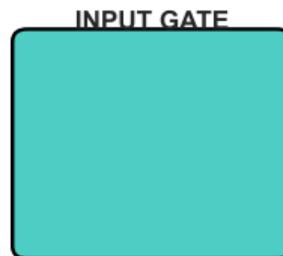
1. Forget unimportant info
2. Store new important info
3. Retrieve when needed

Three Gates Control Memory Flow



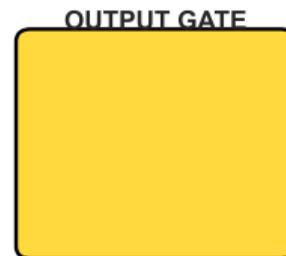
What to remove
from memory

0.0 = Erase all
1.0 = Keep all



What new info
to store

0.0 = Ignore new
1.0 = Store all

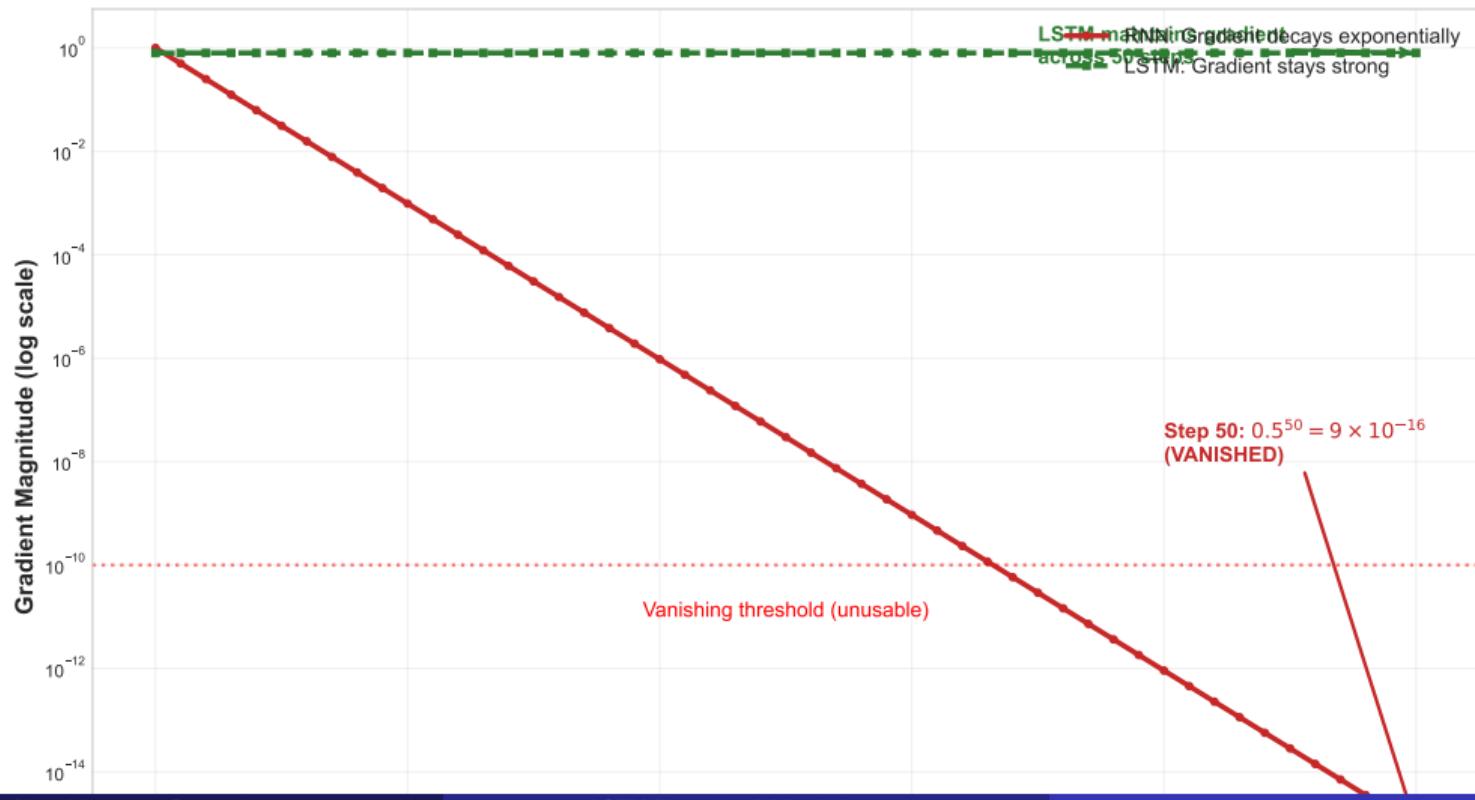


What to reveal
from memory

0.0 = Hide all
1.0 = Show all

All gates use Sigmoid function: output between 0 and 1

Why RNNs Fail: The Vanishing Gradient Problem



Checkpoint 1: Understanding the Problem

Q1: Why can't N-grams solve the Paris problem?

- A) Too slow
- B) Fixed 1-3 word window
- C) Too much memory
- D) Don't understand French

Q2: What causes RNN gradient vanishing?

- A) Too many layers
- B) Exponential decay over time
- C) Learning rate too high
- D) Wrong activation function

A1: B - Fixed window

N-grams use fixed 1-3 word context.
Paris is 18 words back - completely invisible!

A2: B - Exponential decay

Gradients multiply by <1 at each step. 0.5^{50} becomes vanishingly small!

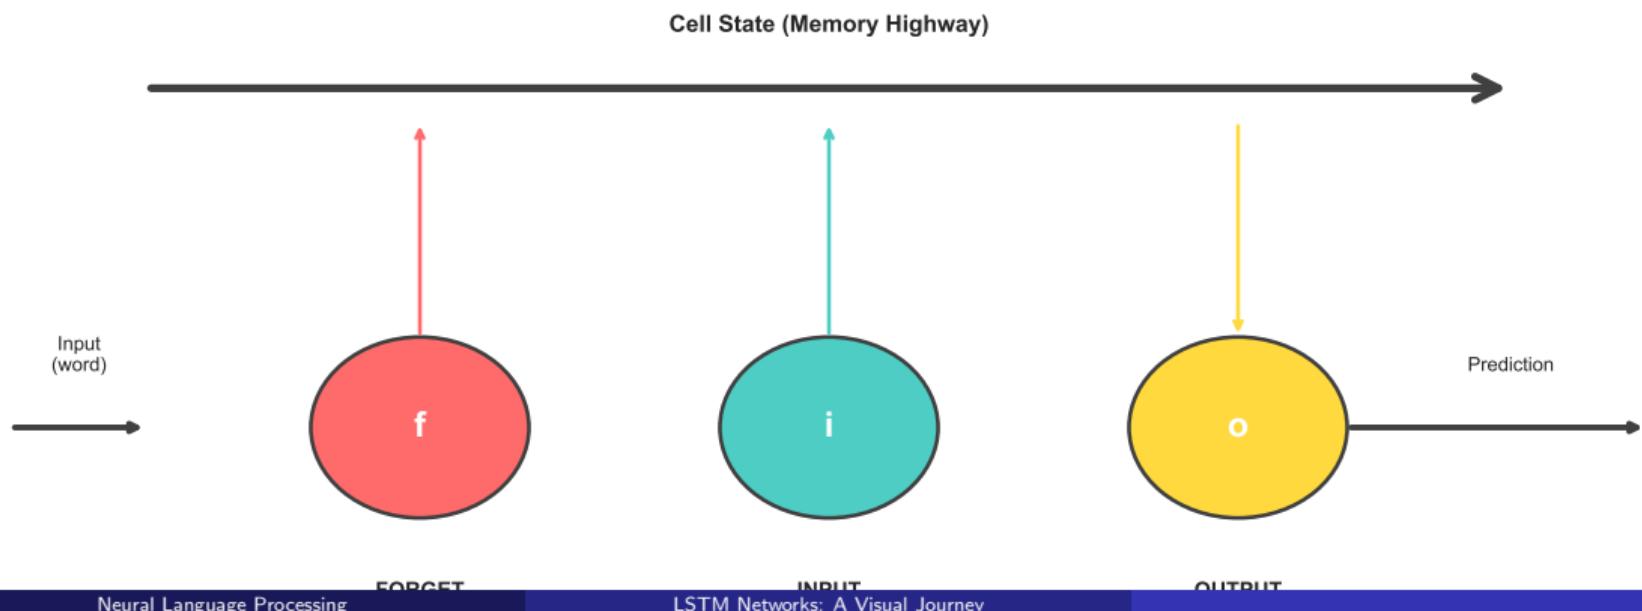
Q3: How many gate mechanisms does LSTM need?

- A) 1: Memory
- B) 2: Input + Output
- C) 3: Forget + Input + Output
- D) 4: All gates + Cell

A3: C - Three gates

Forget (remove), Input (add), Output (reveal)

LSTM Cell: Three Gates Control Memory



Notation Guide: Understanding the Symbols

States and Inputs:

- x_t - Input at time t (current word)
- h_t - Hidden state (output) at time t
- h_{t-1} - Previous hidden state
- C_t - Cell state (memory) at time t
- C_{t-1} - Previous cell state

Gates (all 0 to 1):

- f_t - Forget gate (what to erase)
- i_t - Input gate (what to store)
- o_t - Output gate (what to reveal)
- \tilde{C}_t - Candidate memory (-1 to 1)

Operations:

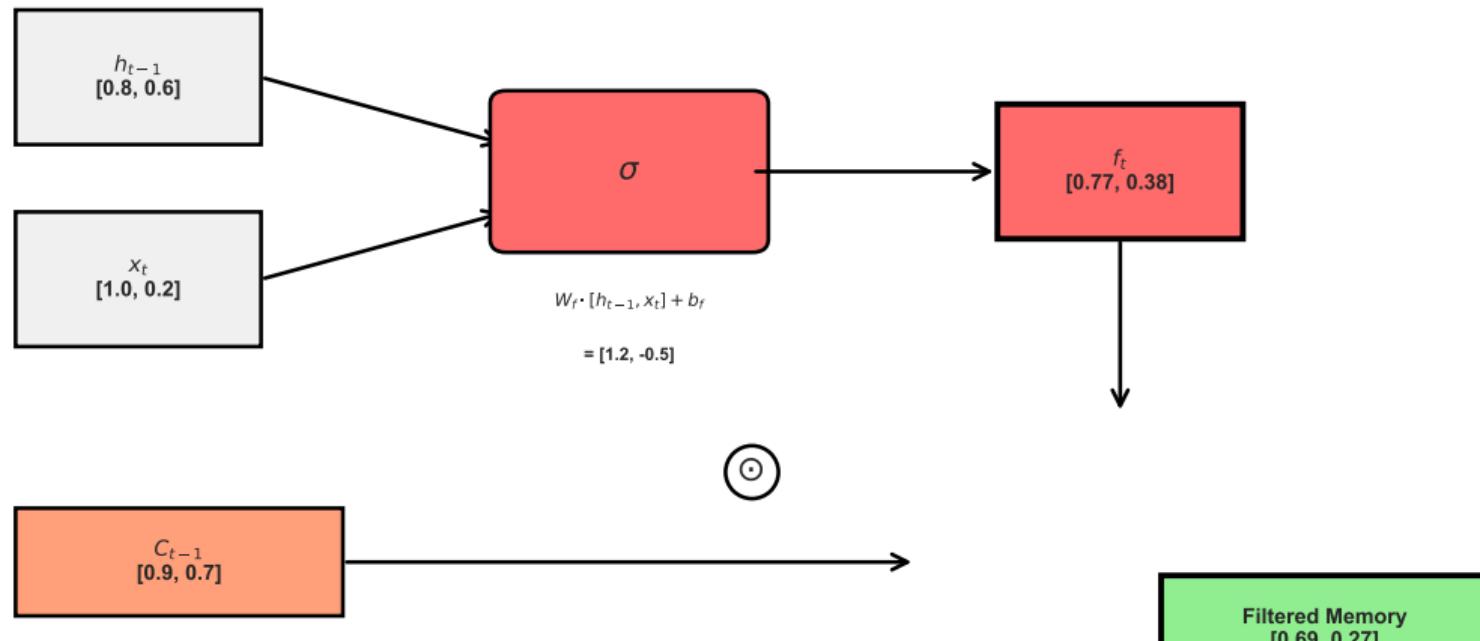
- σ - Sigmoid (0 to 1) for gates
- \tanh - Tanh (-1 to 1) for memory
- \odot - Element-wise multiplication
- $[a, b]$ - Concatenation (stick together)

Learned Parameters:

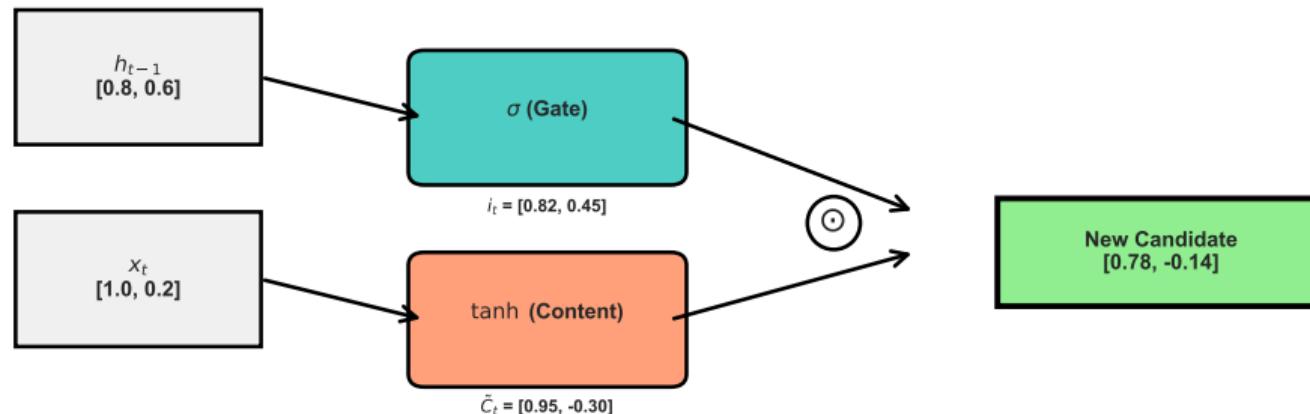
- W_f, W_i, W_o, W_C - Weight matrices
- b_f, b_i, b_o, b_C - Bias vectors

Time step t : Current position in sequence

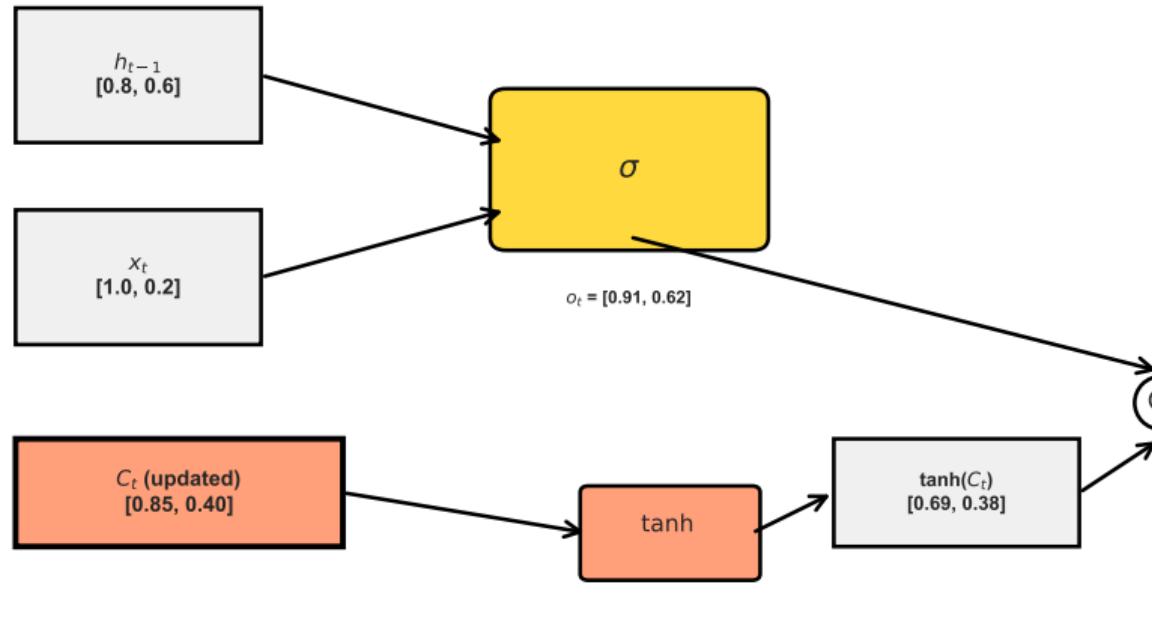
Forget Gate: Step-by-Step



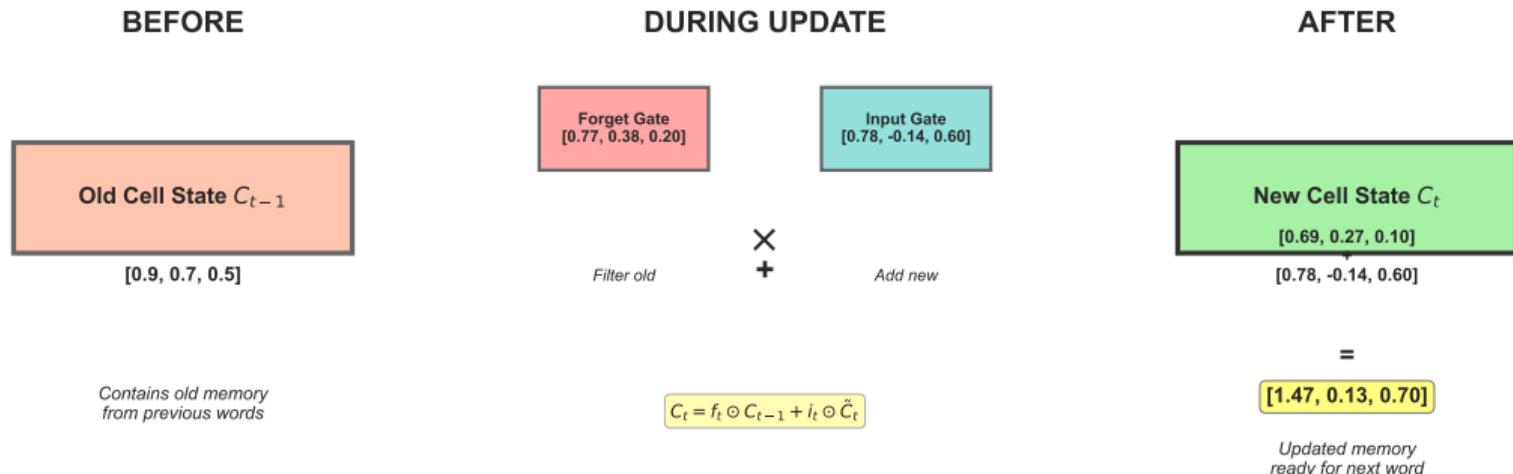
Input Gate: Step-by-Step



Output Gate: Step-by-Step



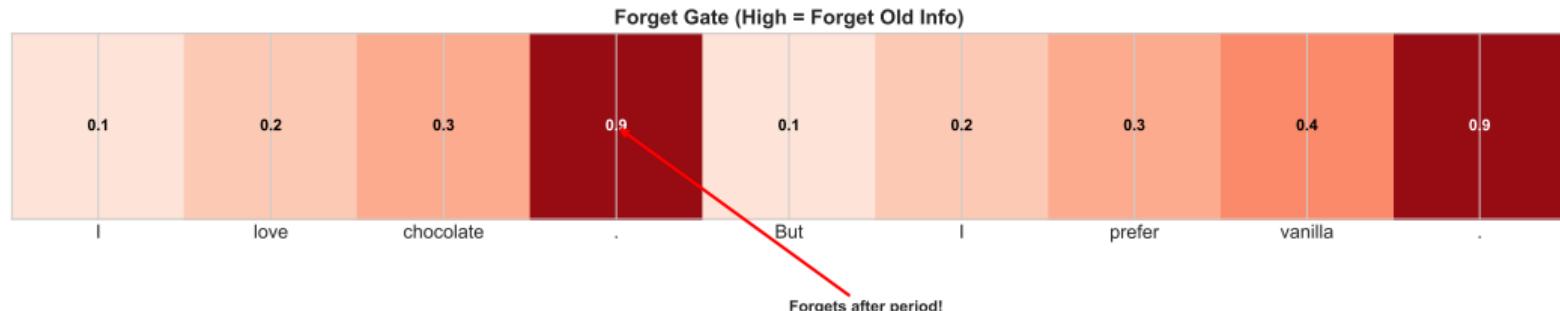
Cell State Update: The Memory Highway



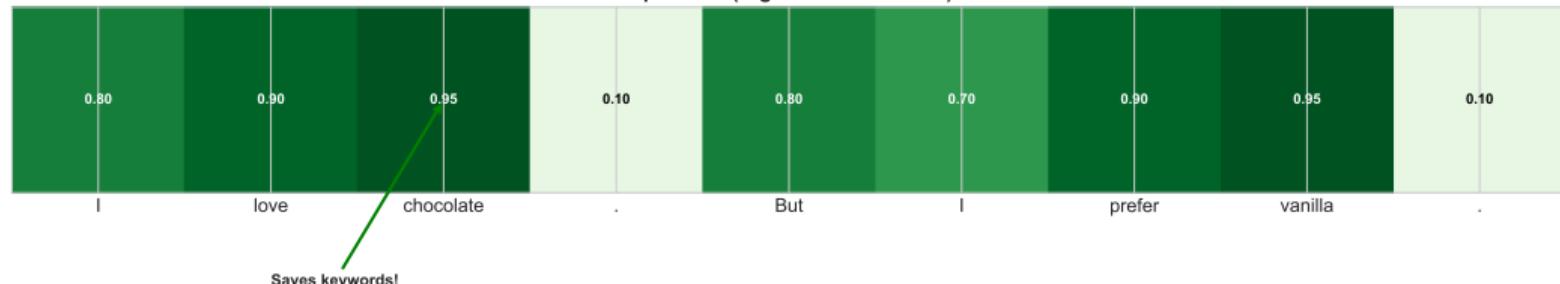
In Plain English: Old memory C_{t-1} flows through: (1) Forget gate filters it, (2) Input gate adds new info, (3) Result is updated memory C_t . This is the ADDITIVE update that prevents gradient vanishing!

Gate Activations: Real Numbers in Action

LSTM Gate Activations: "I love chocolate. But I prefer vanilla."



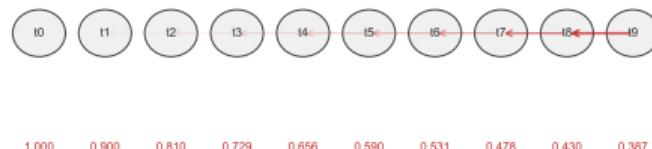
Input Gate (High = Save New Info)



Output Gate (High = Use Memory for Prediction)

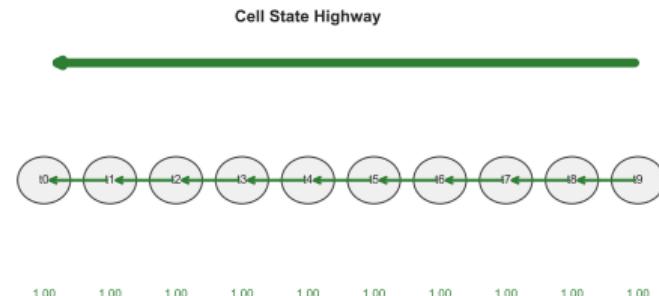
Why LSTM Solves Gradient Vanishing

RNN: Vanishing Gradient



Gradient shrinks exponentially: $0.9^{10} \approx 0.35$

LSTM: Gradient Highway



Gradient preserved: $1.0^{10} = 1.0$

In Plain English: RNN: Gradients MULTIPLY through time ($\times 0.5$ each step = exponential decay). LSTM: Gradients ADD through cell state ($+f_t \odot$ = linear flow). Addition preserves gradients!

Checkpoint 2: Understanding Gates and Cell State

Q1: What does $f_t = 0.2$ mean?

- A) Keep 20% of old memory
- B) Erase 20% of old memory
- C) Add 20% new memory
- D) Output 20%

Q2: Why does LSTM have TWO highways (C_t and h_t)?

- A) Redundancy
- B) Speed
- C) C_t = long-term memory, h_t = short-term output
- D) Prevent overfitting

A1: A - Keep 20%

Forget gate f_t controls what to KEEP, not erase. $f_t = 0.2$ means keep 20%, erase 80%.

A2: C - Different roles

C_t stores unfiltered long-term memory (highway). h_t is filtered output for predictions.

Q3: What's the key operation that prevents vanishing gradients?

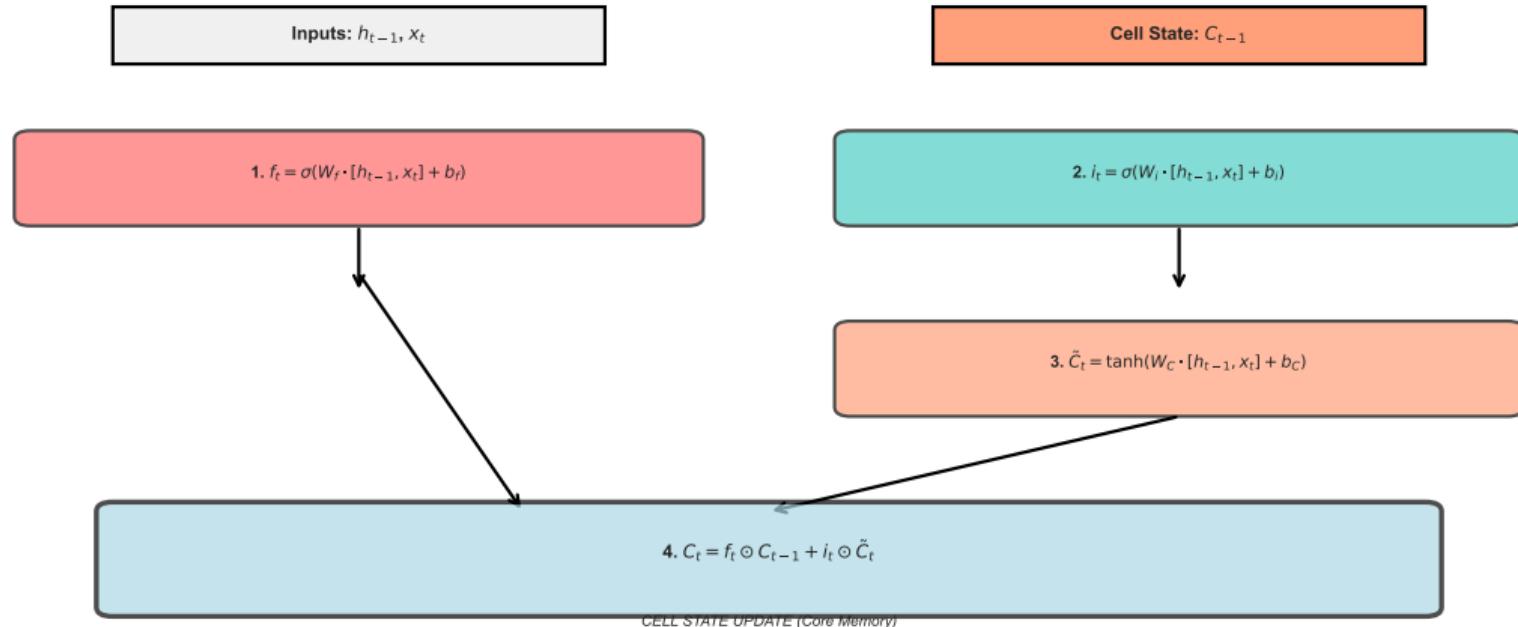
- A) Sigmoid activation
- B) Additive cell state update
- C) Element-wise multiplication
- D) Tanh normalization

A3: B - Addition

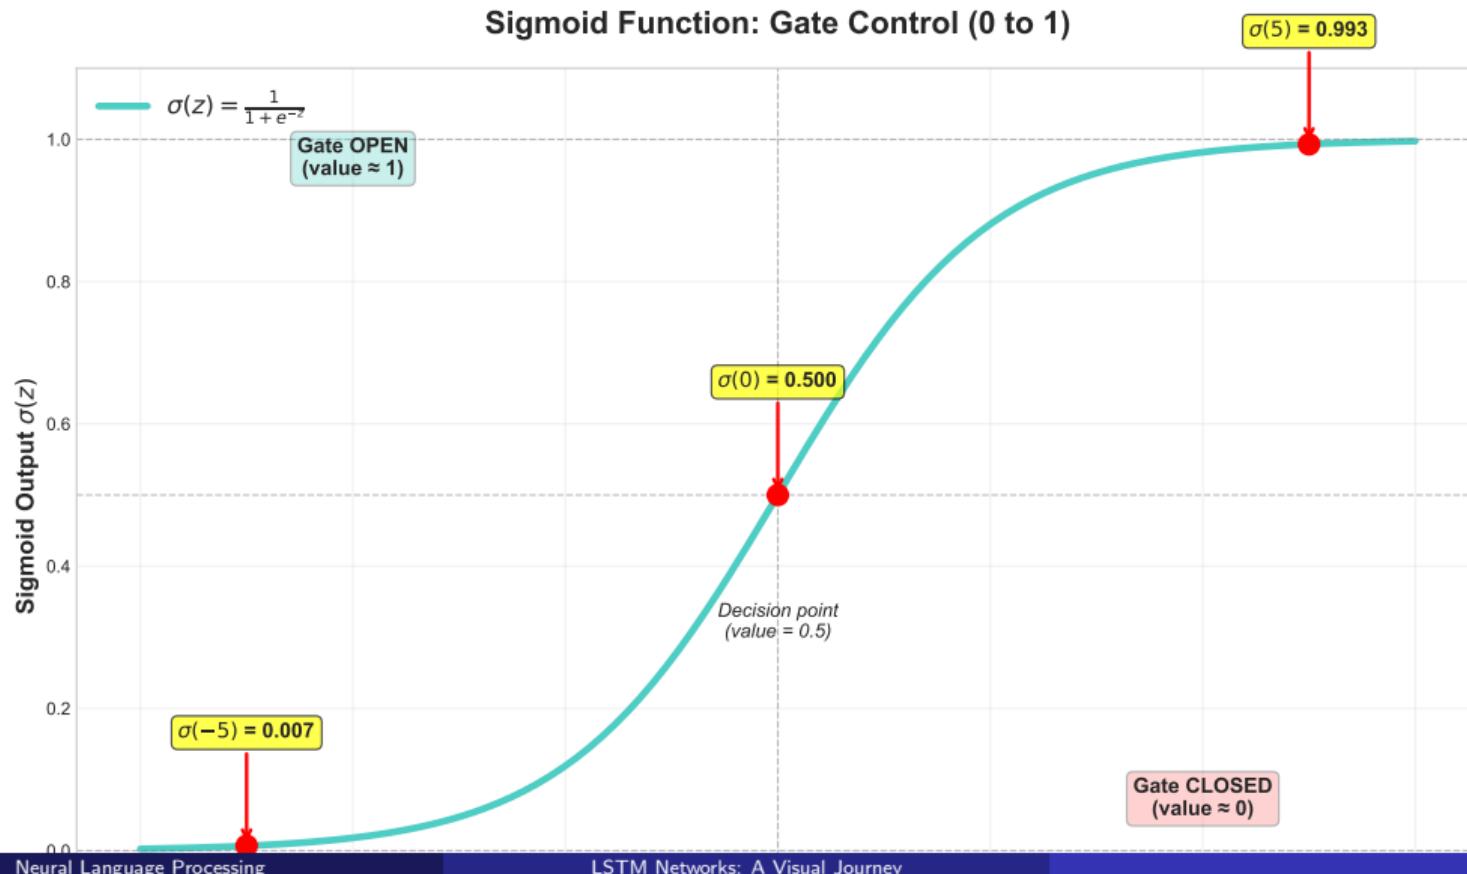
$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$ uses addition, creating gradient highway!

Complete Forward Pass: All 6 Equations as Flowchart

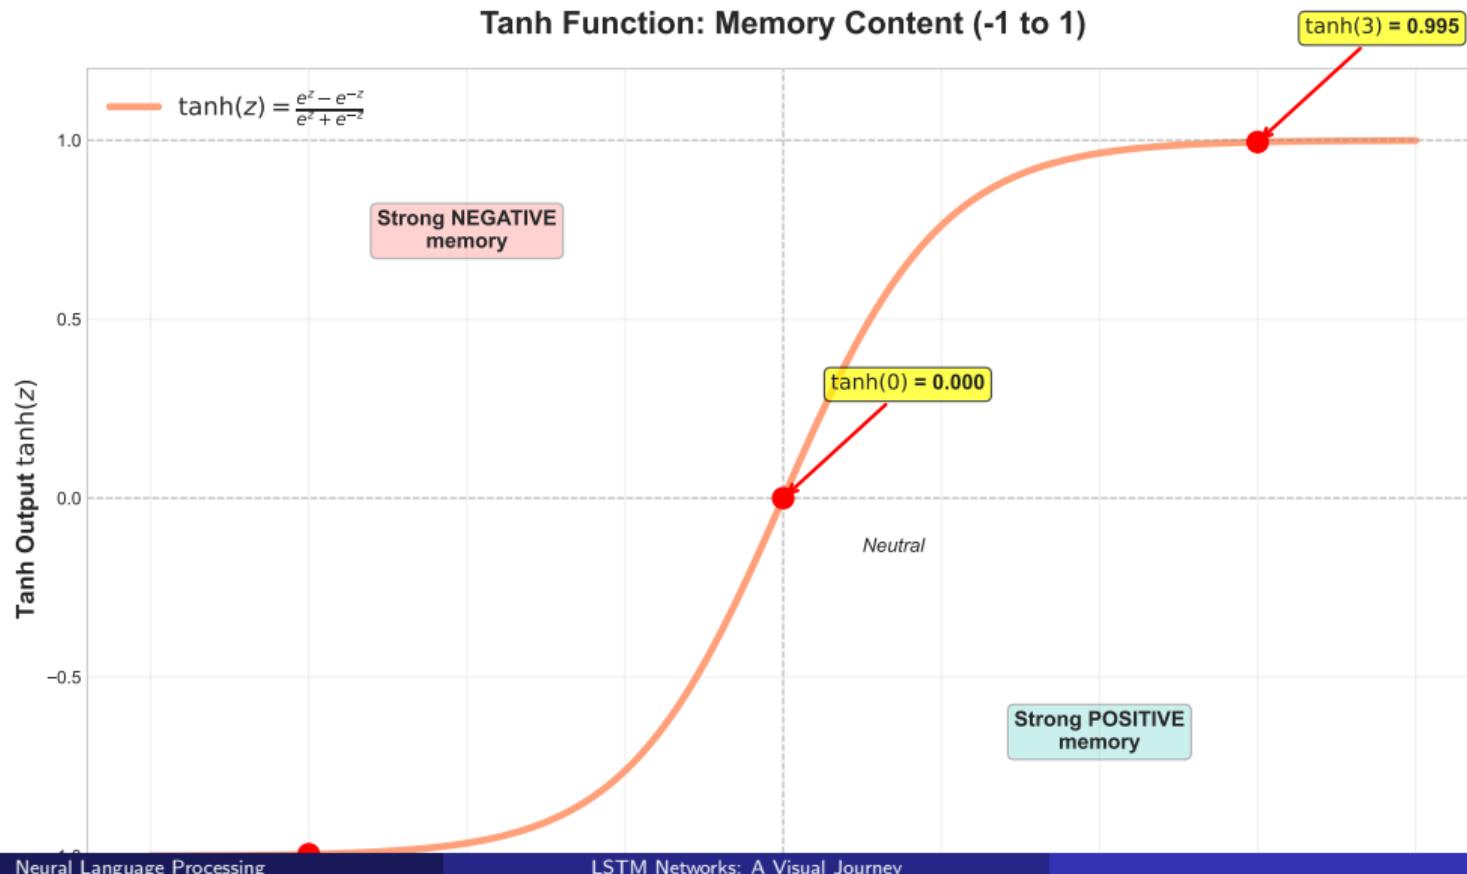
Complete LSTM Forward Pass: All 6 Equations



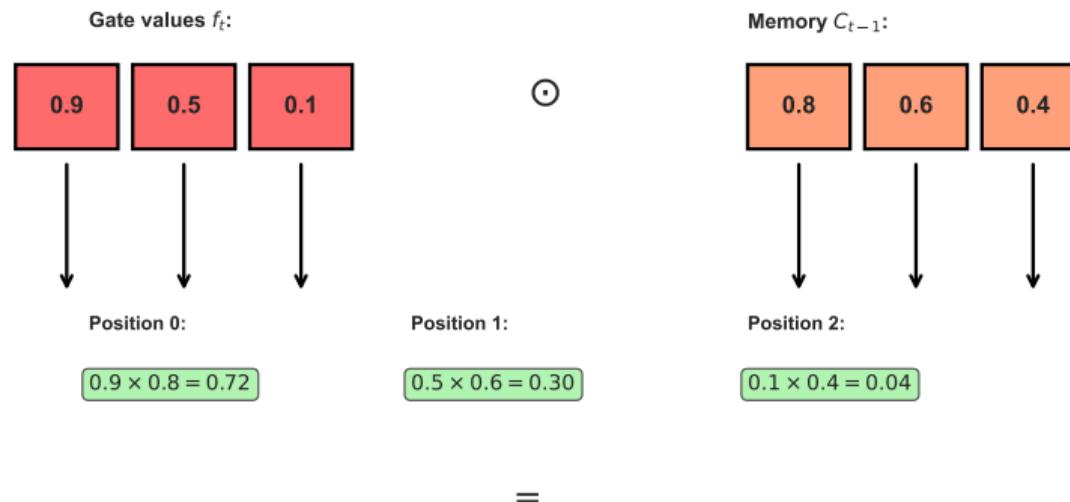
Sigmoid Function: Gate Control Mechanism



Tanh Function: Memory Content Normalization

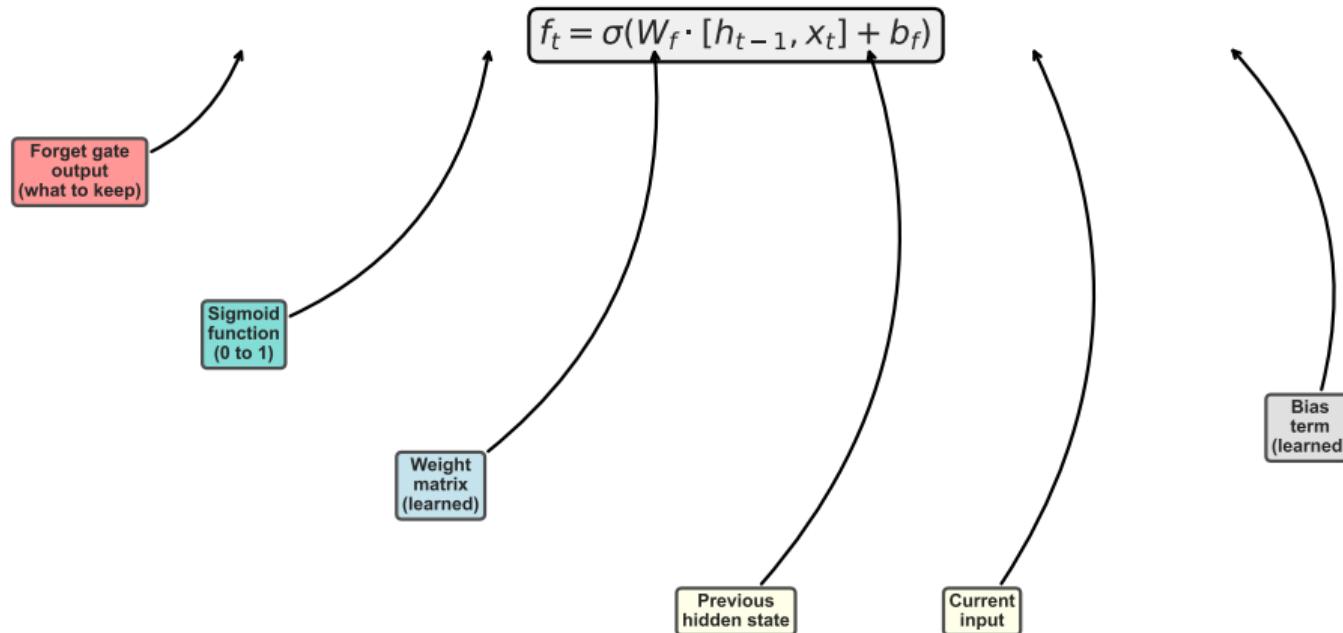


Element-wise Multiplication: Position by Position



Equation Anatomy: Reading LSTM Formulas

Equation Anatomy: Reading LSTM Formulas



Step 1: Inputs

- $h_{t-1} = [0.8, 0.6]$
- $x_t = [1.0, 0.2]$
- $C_{t-1} = [0.9, 0.7]$

Step 2: Compute Gates

(Assume weights trained)

- $f_t = \sigma([1.2, -0.5]) = [0.77, 0.38]$
- $i_t = \sigma([1.5, -0.2]) = [0.82, 0.45]$
- $o_t = \sigma([2.0, 0.5]) = [0.88, 0.62]$
- $\tilde{C}_t = \tanh([2.0, -0.3]) = [0.96, -0.29]$

Step 3: Update Cell State

$$\begin{aligned} C_t &= f_t \odot C_{t-1} + i_t \odot \tilde{C}_t \\ &= [0.77, 0.38] \odot [0.9, 0.7] \\ &\quad + [0.82, 0.45] \odot [0.96, -0.29] \\ &= [0.69, 0.27] + [0.79, -0.13] \\ &= [1.48, 0.14] \end{aligned}$$

Step 4: Compute Output

$$\begin{aligned} h_t &= o_t \odot \tanh(C_t) \\ &= [0.88, 0.62] \odot \tanh([1.48, 0.14]) \\ &= [0.88, 0.62] \odot [0.90, 0.14] \\ &= [0.79, 0.09] \end{aligned}$$

In Plain English: Watch numbers flow: Forget gate keeps 77% of first memory value, Input gate adds strong new info (0.79), Output gate reveals 88% of normalized cell state to output.

Checkpoint 3: Understanding the Math

Q1: What does $\sigma(-3)$ output?

- A) -3
- B) Close to 0
- C) 0.5
- D) Close to 1

Q2: What is $[0.5, 0.8] \odot [2.0, 3.0]$?

- A) [2.5, 3.8]
- B) [1.0, 2.4]
- C) [1.5, 2.2]
- D) [0.25, 0.27]

A1: B - Close to 0

$\sigma(-3) = 0.047 \approx 0$. Large negative input → gate closed!

A2: B - [1.0, 2.4]

Element-wise: $[0.5 \times 2.0, 0.8 \times 3.0] = [1.0, 2.4]$

Q3: Why use Tanh for \tilde{C}_t but Sigmoid for gates?

- A) Tanh faster
- B) Tanh allows negative values,
- Sigmoid for 0-1 control
- C) Historical reasons
- D) Random choice

A3: B - Different purposes

Gates need 0-1 (off/on). Memory needs positive AND negative evidence (-1 to +1).

Training Progression: Watching LSTM Learn

LSTM Training: Watching It Learn

Epoch 1: Random Initialization

Input: "I love chocolate"

Prediction: "xjwkq"

Loss: 8.5 (Gibberish!)

Epoch 10: Learning Letters

Input: "I love chocolate"

Prediction: "cream"

Loss: 2.1 (Better!)

Epoch 50: Understanding Context

Input: "I love chocolate"

Prediction: "ice cream"

Epoch 200: Fluent Generation

Input: "I love chocolate"

Prediction: "ice cream and strawberry cake"

Training Recipe: Step-by-Step LSTM Training

1. Data Preparation

- Tokenize text into sequences
- Create input-target pairs
- Batch sequences of same length
- Pad if needed

2. Model Setup

- Initialize weight matrices W_f, W_i, W_o, W_C
- Initialize bias vectors b_f, b_i, b_o, b_C
- Choose hidden size (e.g., 128, 256, 512)
- Choose number of layers (1-3 typical)

3. Forward Pass

- Process sequence word by word
- Update cell state C_t at each step
- Collect outputs h_t for predictions

4. Loss Computation

- Compare predictions to targets
- Use cross-entropy loss
- Average over sequence

5. Backward Pass (BPTT)

- Compute gradients backward through time
- Cell state acts as gradient highway
- No vanishing gradient problem!

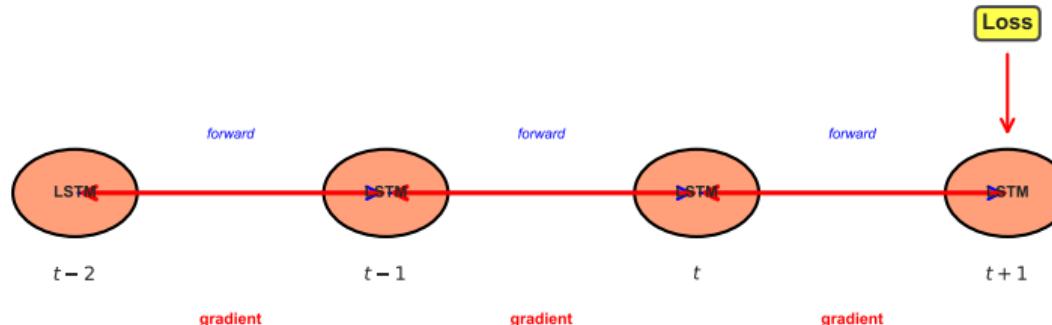
6. Weight Update

- Apply optimizer (Adam recommended)
- Learning rate: 10^{-3} to 10^{-4}
- Gradient clipping: max norm 5.0

7. Iterate

- Repeat for multiple epochs (10-50)
- Monitor validation loss

Backpropagation Through Time (BPTT)



Forward Pass (blue): Compute outputs from left to right

Backward Pass (red): Propagate gradients from right to left

LSTM cell state acts as gradient highway - prevents vanishing!



Model Comparison: When to Use LSTM

Comparison: N-gram vs RNN vs LSTM

Feature	N-gram	RNN	LSTM
Memory Type	Fixed window	Fading	Selective
Long Context	□	□	□
Parameters	Few	Moderate	Many
Training Speed	Fast	Medium	Slow
Vanishing Gradient	N/A	Yes □	Solved □
Best For	Short (2-3 words)	Medium (10 words)	Long (50+ words)
Example	"I love..."	"The cat sat..."	"The cat, who was...finally..."

Applications: Where LSTMs Excel

Natural Language Processing:

- Machine translation
- Text generation
- Sentiment analysis
- Named entity recognition
- Question answering

Time Series:

- Stock price prediction
- Weather forecasting
- Energy demand prediction
- Anomaly detection

Speech and Audio:

- Speech recognition
- Music generation
- Voice synthesis
- Audio classification

Video Analysis:

- Action recognition
- Video captioning
- Motion prediction
- Event detection

Key Advantage:

LSTM excels when **context from distant past** matters for current prediction!

Checkpoint 4: Training and Applications

Q1: What is BPTT?

- A) Backward Pass Through Training
- B) Backpropagation Through Time
- C) Batch Processing Training Technique
- D) Bi-directional Processing Through Time

Q2: Why use gradient clipping in LSTM training?

- A) Speed up training
- B) Prevent exploding gradients
- C) Reduce memory usage
- D) Improve accuracy

A1: B - Backpropagation Through Time

BPTT unrolls sequence through time and backpropagates errors from future to past.

A2: B - Prevent explosion

While LSTM solves vanishing, gradients can still explode. Clipping caps maximum gradient norm.

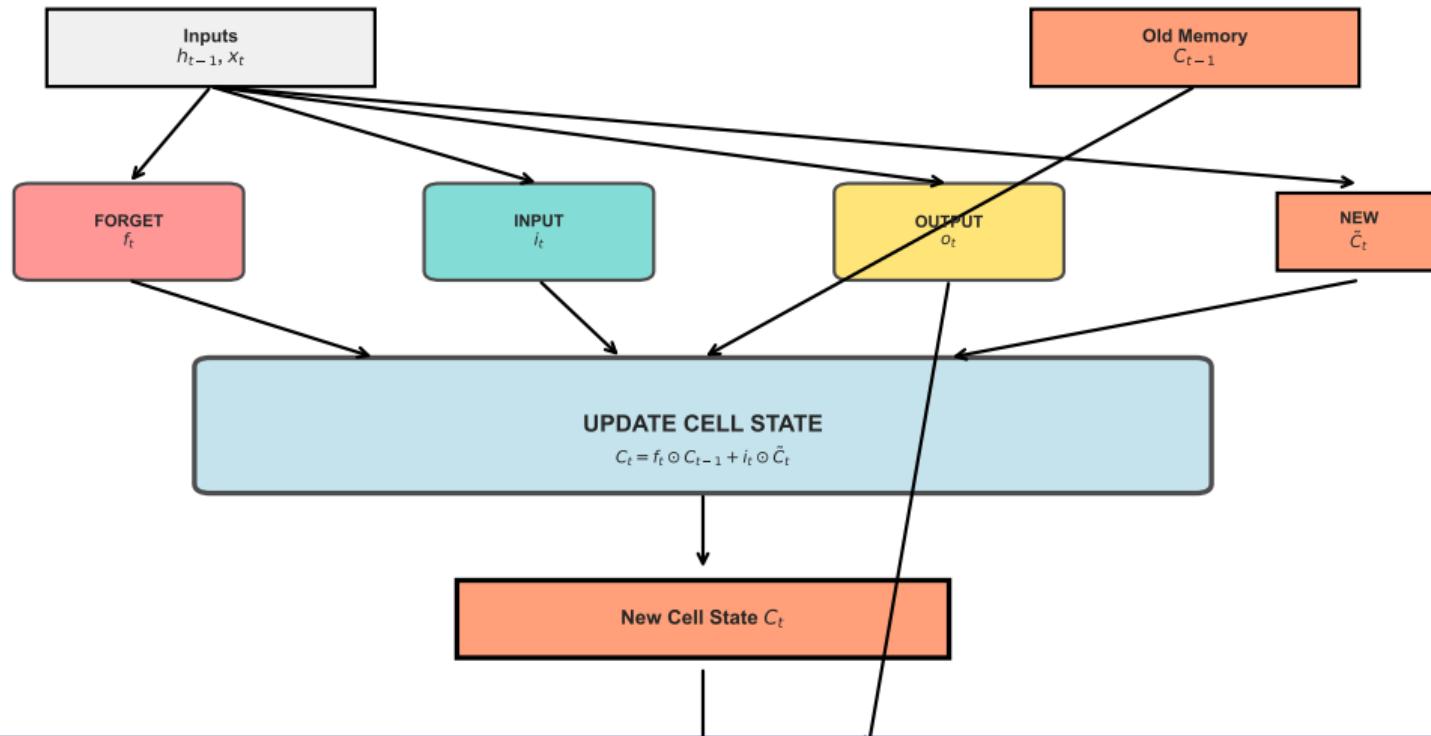
Q3: When should you NOT use LSTM?

- A) Machine translation
- B) Short 2-3 word context
- C) Speech recognition
- D) Time series prediction

A3: B - Short context

For 2-3 word context, simpler models (N -gram, simple RNN) are faster.

Complete LSTM Flow: End-to-End



Quick Reference: Essential LSTM Equations

The Six Core Equations:

1. Forget Gate:

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

2. Input Gate:

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

3. Candidate Memory:

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

4. Cell State Update:

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$

5. Output Gate:

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

6. Hidden State:

$$h_t = o_t \odot \tanh(C_t)$$

Key Takeaways:

- ➊ **Problem:** N-grams limited window, RNNs vanishing gradient
- ➋ **Solution:** Three gates + cell state highway
- ➌ **Forget gate:** Decides what to erase (0-1)
- ➍ **Input gate:** Decides what to store (0-1)
- ➎ **Output gate:** Decides what to reveal (0-1)
- ➏ **Cell state:** Long-term memory highway (additive update)
- ➐ **Hidden state:** Short-term filtered output
- ➑ **Why it works:** Addition in C_t update creates gradient highway
- ➒ **Use when:** Long-distance dependencies matter
- ➓ **Avoid when:** Short context or parallel processing needed

Remember: Cell state is the **memory highway**, gates control the **traffic**!