

## Week 4: Sequence-to-Sequence Models From Variable Lengths to Attention Mechanisms

BSc-Level Enhanced Version with Layout Templates

# Today's Journey

## Where We Are:

- Week 3: RNNs handle sequences
- Week 3: Fixed input → fixed output
- **Today: Variable input → variable output**

## Today's Learning Objectives:

- ① Understand why translation needs special architecture
- ② Master encoder-decoder framework
- ③ Discover attention mechanism
- ④ Apply seq2seq in practice

### Prerequisite

You should know:

- RNN basics
- Hidden states
- Backpropagation

# Why Can't We Just Use RNNs?



## Build Your Intuition

RNNs expect: one input → one output


Translation needs: any length → any length

## The Fundamental Problem:

- English: “I love you” (3 words)
- French: “Je t'aime” (2 words)
- Japanese: “Aishiteru” (1 word)

## RNN Limitation:

- Fixed mapping between positions
- Can't handle length mismatch
- No way to “wait” or “generate multiple”



## Common Misconception

“Just pad with zeros!”  
No - that changes meaning and wastes computation.

# Concrete Example: “Hello World” Translation

Let's trace through actual numbers:

**Input:** “Hello world” (English)

- Word 1: “Hello” → embed = [0.2, -0.1, 0.5]
- Word 2: “world” → embed = [0.3, 0.4, -0.2]

**Target:** “Bonjour monde” (French)

- Word 1: “Bonjour”
- Word 2: “monde”

👍 Try It

With standard RNN:

- Position 1: Hello → Bonjour ✓
- Position 2: world → monde ✓
- Position 3: ??? → Nothing needed

But what if French had 3 words?  
The system breaks!

✓ Checkpoint

Can you explain why padding doesn't solve this?

Answer: Because we don't know output length in advance!


# The Brilliant Insight: Two-Stage Process

## 🧠 Build Your Intuition

Think like a human translator:

- ① READ and understand the whole sentence
- ② THINK about the meaning
- ③ WRITE the translation

Encoder-Decoder Flow with Actual Values



# Encoder: Step-by-Step with Numbers

Processing “The cat sat”:

## Step 1: Embedding

- “The” →  $x_1 = [0.1, -0.2, 0.3]$
- “cat” →  $x_2 = [0.5, 0.1, -0.1]$
- “sat” →  $x_3 = [0.3, 0.6, 0.2]$

## Step 2: LSTM Processing

$$h_0 = [0, 0, 0] \text{ (initial)}$$

$$h_1 = \text{LSTM}(x_1, h_0) = [0.2, -0.1, 0.1]$$

$$h_2 = \text{LSTM}(x_2, h_1) = [0.4, 0.3, -0.2]$$

$$h_3 = \text{LSTM}(x_3, h_2) = [0.6, 0.5, 0.3]$$

## Step 3: Context Vector

$$c = h_3 = [0.6, 0.5, 0.3]$$

### Try It

Notice how each hidden state builds on the previous:

- $h_1$ : knows “The”
- $h_2$ : knows “The cat”
- $h_3$ : knows “The cat sat”

Final  $h_3$  contains the complete meaning!

# Encoder Implementation

```
1 class Encoder(nn.Module):
2     def __init__(self, vocab_size,
3                  embed_dim=100,
4                  hidden_dim=256):
5         super().__init__()
6         self.embedding = nn.Embedding(
7             vocab_size, embed_dim)
8         self.lstm = nn.LSTM(
9             embed_dim, hidden_dim)
10
11    def forward(self, source):
12        # source: [seq_len, batch]
13        embedded = self.embedding(source)
14        # embedded: [seq_len, batch, embed]
15
16        outputs, (hidden, cell) = \
17            self.lstm(embedded)
18        # hidden: [1, batch, hidden_dim]
19
20        return hidden, cell # Context!
```

## Key Implementation Points:

- **vocab\_size**: How many words we know
- **embed\_dim**: Word vector size (100)
- **hidden\_dim**: Context size (256)

## Dimensions Example:

- Input: “Hello world” = [2, 1]
- After embedding: [2, 1, 100]
- Context output: [1, 1, 256]

### Checkpoint

The context vector is always the same size (256) regardless of input length!

# Decoder: Generating Step by Step

Generating “Le chat” from context  $c = [0.6, 0.5, 0.3]$ :

## Step 1: Initialize with context

- $h_0^{dec} = c = [0.6, 0.5, 0.3]$
- Input:  $\langle \text{START} \rangle_i$  token

## Step 2: Generate first word


- $h_1^{dec} = \text{LSTM}(\langle \text{START} \rangle, h_0^{dec})$
- Output probabilities: {Le: 0.7, La: 0.2, Un: 0.1}
- Select: “Le”

## Step 3: Generate second word

- $h_2^{dec} = \text{LSTM}(\text{“Le”}, h_1^{dec})$
- Output probabilities: {chat: 0.8, chien: 0.2}
- Select: “chat”

```
1 # Decoder forward pass
2 def forward(self, context):
3     hidden = context
4     outputs = []
5
6     word = "<START>"
7     for _ in range(max_len):
8         embed = embedding(word)
9         hidden = lstm(embed, hidden)
10        probs = softmax(hidden)
11        word = sample(probs)
12
13        if word == "<END>":
14            break
15        outputs.append(word)
16
17    return outputs
```

# The Information Bottleneck



## The Problem:

- 10-word sentence:  $\approx 100$  bits of info
- 256-dim vector: 256 numbers capacity
- 50-word sentence:  $\approx 500$  bits of info
- Same 256-dim vector! **Overflow!**



### Common Misconception

“Just make the vector bigger!”  
Problems with that:

- More parameters to learn
- Slower training
- Still fixed size

# Attention: The Solution

**Key Insight:** Look back at ALL encoder states, not just the last one!

**Generating “chat” (cat in French):**

**Step 1: Score each encoder state**


- $e_1 = \text{score}(h_{\text{chat}}^{\text{dec}}, h_{\text{The}}^{\text{enc}}) = 0.1$
- $e_2 = \text{score}(h_{\text{chat}}^{\text{dec}}, h_{\text{cat}}^{\text{enc}}) = 0.9$
- $e_3 = \text{score}(h_{\text{chat}}^{\text{dec}}, h_{\text{sat}}^{\text{enc}}) = 0.3$

**Step 2: Normalize to probabilities**

- $\alpha_1 = \text{softmax}(0.1) = 0.15$
- $\alpha_2 = \text{softmax}(0.9) = 0.70$
- $\alpha_3 = \text{softmax}(0.3) = 0.15$

**Step 3: Weighted combination**

$$c_{\text{chat}} = 0.15 \cdot h_1 + 0.70 \cdot h_2 + 0.15 \cdot h_3$$



Try It

Attention weights tell us:

- 70% focus on “cat”
- 15% on “The”
- 15% on “sat”

# Attention Implementation

```
1 class BahdanauAttention(nn.Module):
2     def __init__(self, hidden_dim):
3         super().__init__()
4         self.W1 = nn.Linear(hidden_dim, hidden_dim)
5         self.W2 = nn.Linear(hidden_dim, hidden_dim)
6         self.V = nn.Linear(hidden_dim, 1)
7
8     def forward(self, decoder_hidden,
9                 encoder_outputs):
10        # decoder_hidden: [batch, hidden]
11        # encoder_outputs: [seq_len, batch, hidden]
12
13        # Score each encoder output
14        scores = self.V(torch.tanh(
15            self.W1(decoder_hidden) +
16            self.W2(encoder_outputs)))
17        # scores: [seq_len, batch, 1]
18
19        # Convert to probabilities
20        weights = F.softmax(scores, dim=0)
21
22        # Weighted sum
23        context = torch.sum(
24            weights * encoder_outputs, dim=0)
25
26        return context, weights
```

## Numerical Example:

Given:

- Decoder hidden: [0.5, 0.3]
- Encoder outputs:
  - $h_1$ : [0.1, 0.2]
  - $h_2$ : [0.7, 0.4]
  - $h_3$ : [0.3, 0.6]

Computation:

- Scores: [0.2, 0.8, 0.3]
- Weights: [0.15, 0.70, 0.15]
- Context:  $0.15 \times [0.1, 0.2] +$   
 $0.70 \times [0.7, 0.4] +$   
 $0.15 \times [0.3, 0.6]$   
 $= [0.55, 0.46]$

### Checkpoint

Attention gives different context for each output word!

# Training: Teacher Forcing

**Problem:** Early in training, model makes mistakes that compound.

## Without Teacher Forcing:

- Target: “Le chat dort”
- Step 1: Generate “La” (wrong!)
- Step 2: Based on “La”, generate “maison”
- Step 3: Based on “maison”, generate “est”
- Result: “La maison est” (completely wrong!)

## With Teacher Forcing:

- Target: “Le chat dort”
- Step 1: Generate “La”, but feed “Le” to next
- Step 2: Based on “Le”, generate “chat” ✓
- Step 3: Based on “chat”, generate “dort” ✓
- Learning signal much stronger!

### Try It

Teacher forcing schedule:

- Epoch 1-10: 100% forcing
- Epoch 11-20: 50% forcing
- Epoch 21+: 0% forcing

Gradually let the model learn to trust itself!




### Common Misconception

“Always use teacher forcing!”  
No - creates mismatch between training and inference.

## Inference: Beam Search

### Finding the best translation with beam size = 2:

## Beam Search: Exploring Multiple Translation Paths



# Evaluation: BLEU Score Calculation

## Evaluating: “The cat sits on mat”

Reference: “Le chat est sur le tapis”

Generated: “Le chat sur tapis”

### Step 1: Count n-grams

- 1-grams: Le(1/1), chat(1/1), sur(1/1), tapis(1/1)
- 2-grams: “Le chat”(1/1), “chat sur”(0/1), “sur tapis”(0/1)
- 3-grams: None match

### Step 2: Calculate precision

- $P_1 = 4/4 = 1.00$
- $P_2 = 1/3 = 0.33$
- $P_3 = 0/2 = 0.00$

### Step 3: Brevity penalty

- Generated length: 4
- Reference length: 6
- $BP = e^{1-6/4} = e^{-0.5} = 0.61$

### Step 4: Final BLEU

$$\begin{aligned}BLEU &= BP \times \sqrt[3]{P_1 \times P_2 \times P_3} \\&= 0.61 \times \sqrt[3]{1.0 \times 0.33 \times 0.01} \\&= 0.61 \times 0.22 = 0.13\end{aligned}$$

### Try It

#### BLEU scores:

- < 0.1: Useless
- 0.1 – 0.3: Understandable
- 0.3 – 0.5: Good
- > 0.5: Very good

# Summary: Key Takeaways

## What We Learned:

- ① **Problem:** Variable length sequences
- ② **Solution:** Encoder-Decoder architecture
- ③ **Enhancement:** Attention mechanism
- ④ **Training:** Teacher forcing
- ⑤ **Inference:** Beam search
- ⑥ **Evaluation:** BLEU score

## Key Insights:

- Context vector = compressed understanding
- Attention = dynamic focus on relevant parts
- Trade-off between exploration and exploitation

## Next Week:

Transformers - Attention is all you need!

- No more recurrence
- Parallel processing
- Self-attention
- Much faster training

### ✓ Checkpoint

Final check: Can you explain why we need TWO networks?  
Answer: Input length  $\neq$  output length!



# Complete Mathematical Formulation

**Encoder:**

$$h_t^{enc} = f_{enc}(x_t, h_{t-1}^{enc}) \quad (1)$$

$$c = h_T^{enc} \quad (2)$$

**Decoder (without attention):**

$$h_t^{dec} = f_{dec}(y_{t-1}, h_{t-1}^{dec}, c) \quad (3)$$

$$p(y_t | y_{<t}, x) = \text{softmax}(W_o h_t^{dec}) \quad (4)$$

**Decoder (with attention):**

$$e_{ti} = a(h_{t-1}^{dec}, h_i^{enc}) \quad (5)$$

$$\alpha_{ti} = \frac{\exp(e_{ti})}{\sum_j \exp(e_{tj})} \quad (6)$$

$$c_t = \sum_i \alpha_{ti} h_i^{enc} \quad (7)$$

$$h_t^{dec} = f_{dec}(y_{t-1}, h_{t-1}^{dec}, c_t) \quad (8)$$

Where  $a(\cdot)$  is the attention scoring function.



# Seq2Seq in 2024

## Still Using Seq2Seq:

- Machine Translation (Google Translate)
- Speech Recognition (Whisper)
- Summarization
- Code Generation (GitHub Copilot)

## Evolution to Transformers:

- BERT: Encoder-only
- GPT: Decoder-only
- T5: Full encoder-decoder

## Key Improvements:

- Self-attention (next week!)
- Multi-head attention
- Positional encoding
- Layer normalization

## Performance Gains:

- 2014 Seq2Seq: BLEU 20
- 2017 Transformer: BLEU 28
- 2024 GPT-4: BLEU 45+