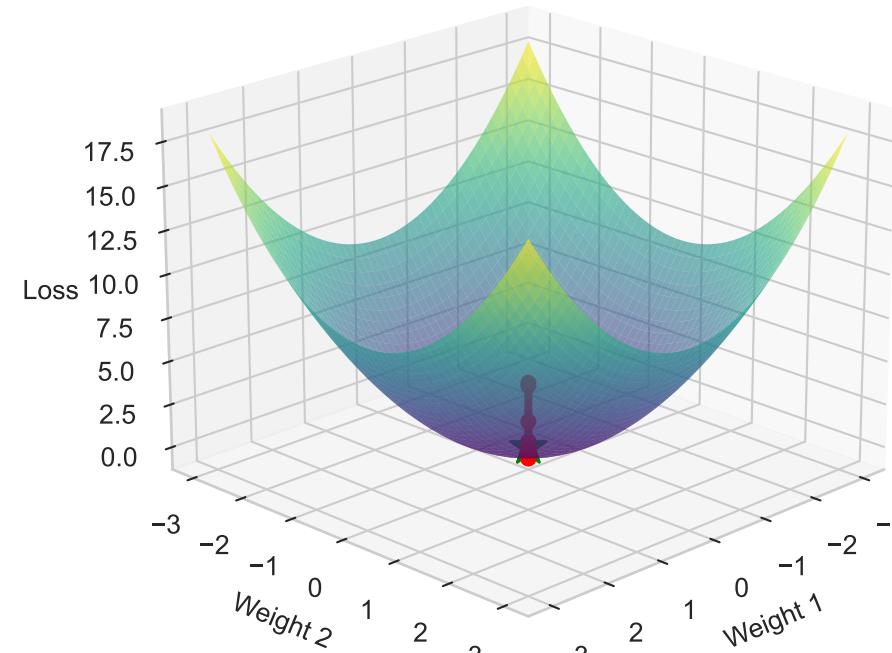
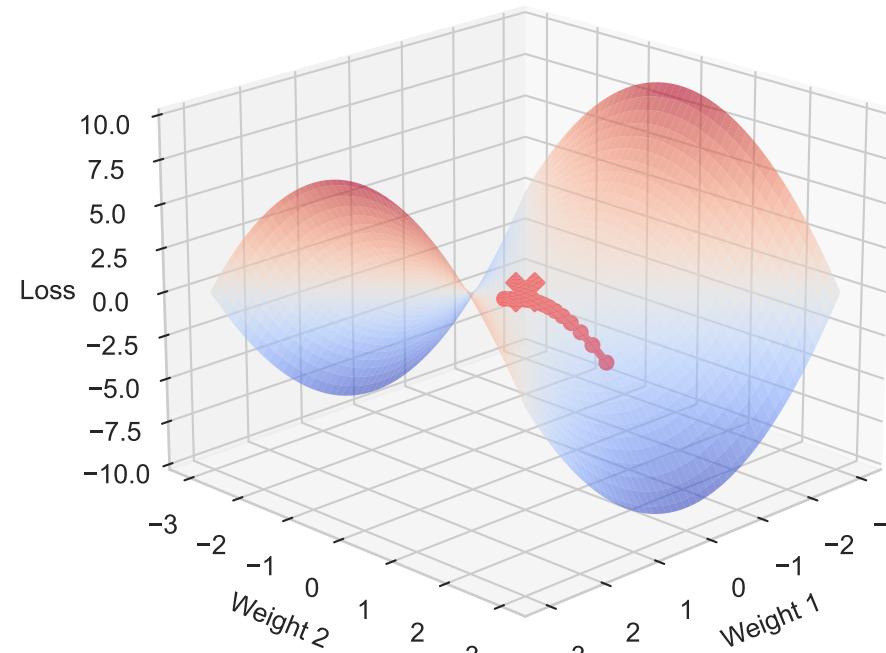


Gradient Descent: Navigating the Loss Landscape

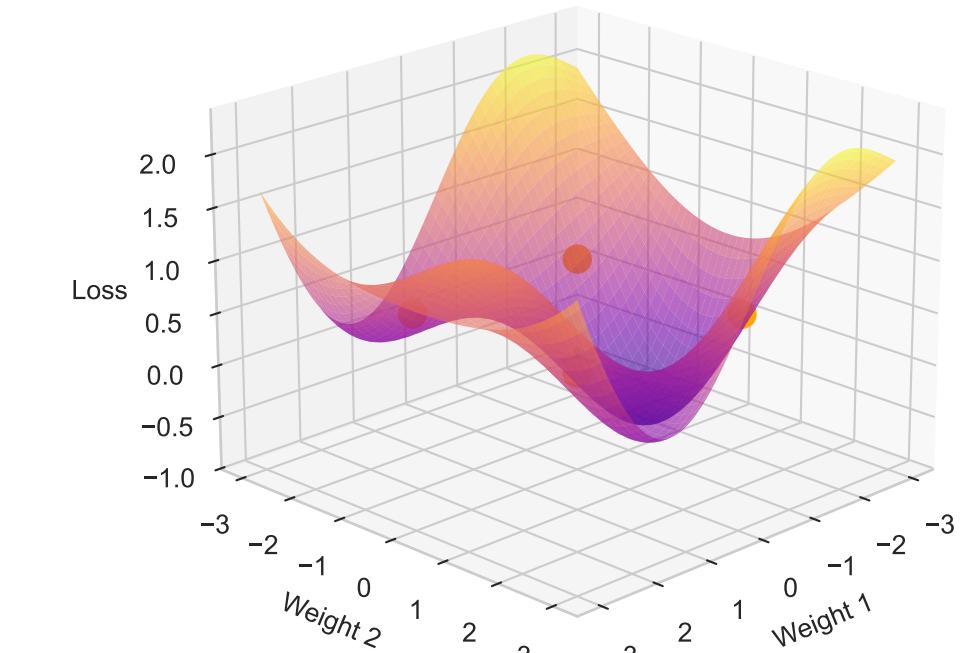
Ideal Case: Convex Loss
(Easy optimization)



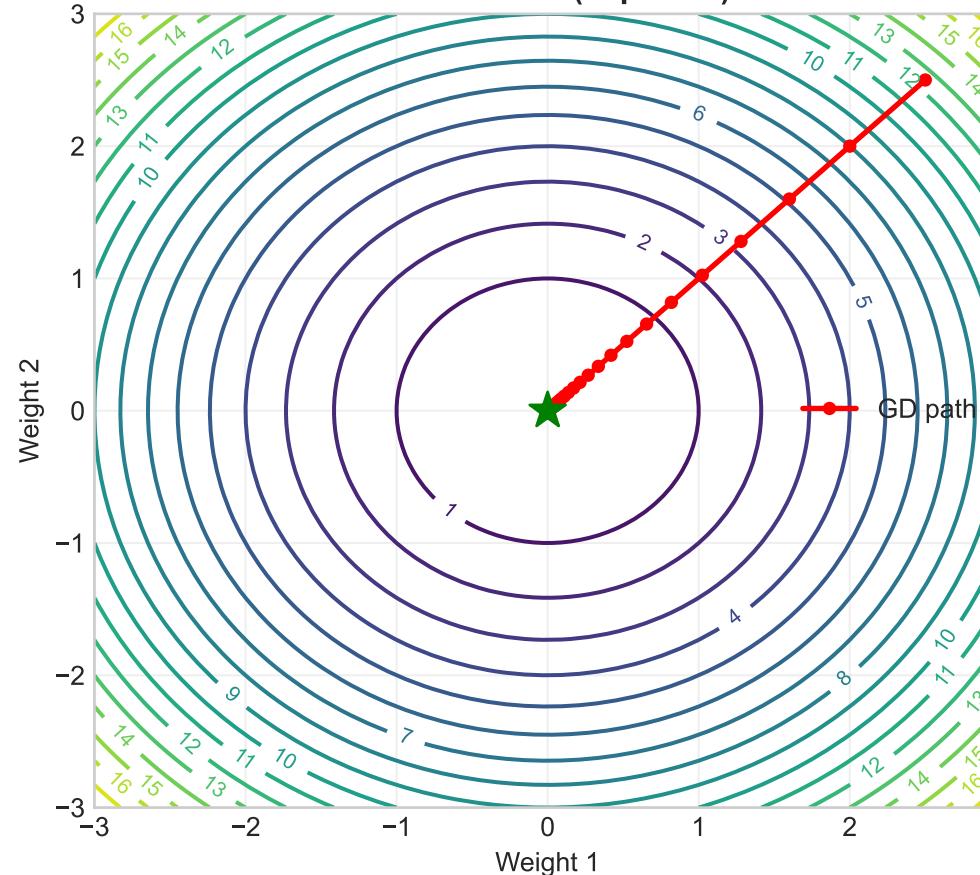
Saddle Point Problem
(Gradient = 0, not minimum)



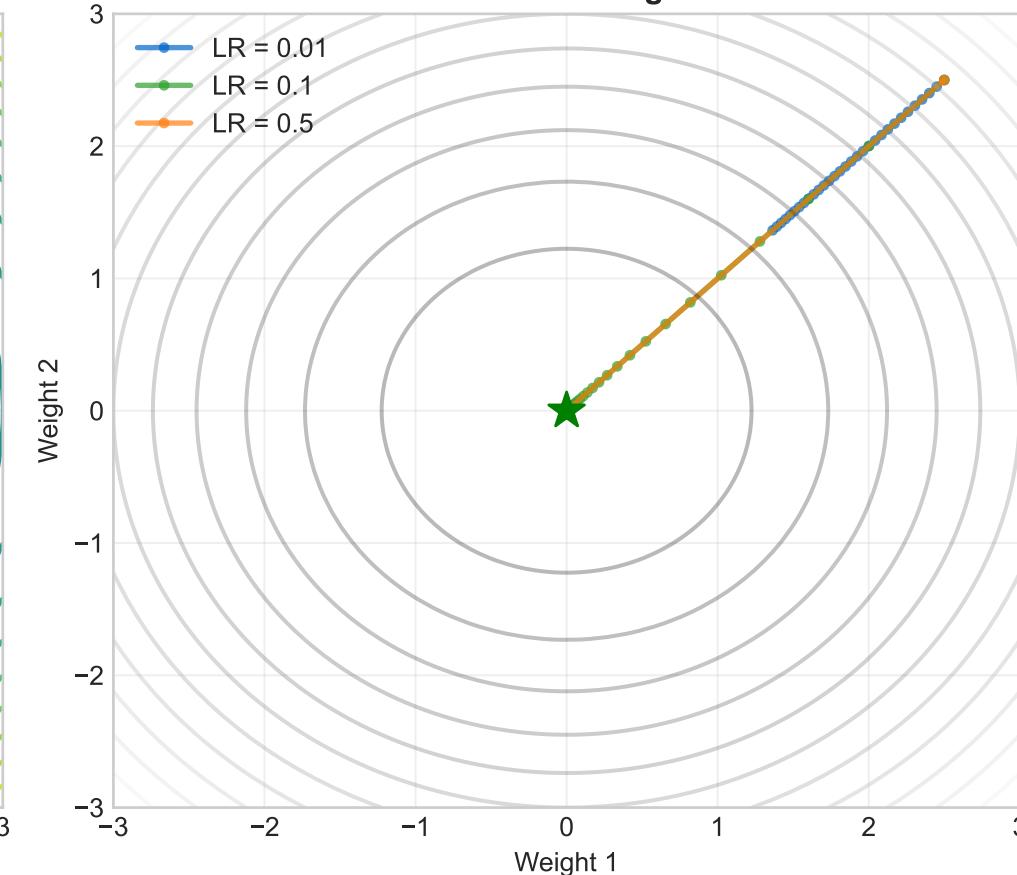
Multiple Local Minima
(Can get stuck)



Convex Loss (Top View)



Effect of Learning Rate



Loss Decrease Over Iterations

