

# Neural Networks: Complete Summary

## From Zero to Understanding in 10 Slides

NLP Course 2025

# 1. The Problem & Motivation

## Why Traditional Programming Fails

Traditional code: IF-THEN-ELSE rules

But how to code:

- Handwritten digit recognition?
- Spam detection?
- Chess at grandmaster level?

## 1959 Mail Sorting Crisis

- U.S. Postal Service: millions of ZIP codes
- Every person writes differently
- Traditional OCR failed

## Paradigm Shift

Instead of programming rules, let computers *learn patterns from examples!*

## Historical Timeline

- 1943: McCulloch-Pitts neuron
- 1958: Perceptron (first learning)
- 1969: Limitations proved (AI Winter)
- 1986: Backpropagation rediscovered
- 1998: LeNet reads bank checks
- 2012: AlexNet (deep learning revolution)

**Key:** Neural networks excel at pattern recognition where rules are unclear

## 2. The Neuron: Building Block

### Mathematical Definition

$$z = \sum_{i=1}^n w_i x_i + b$$

### Components:

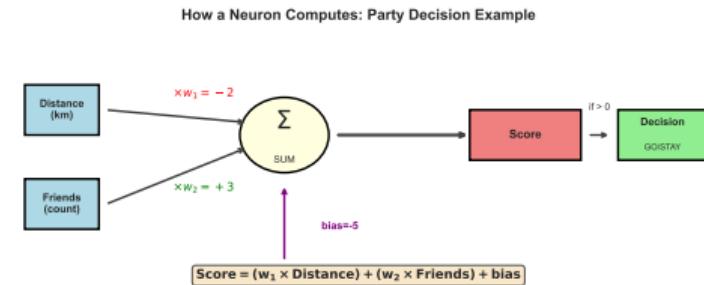
- $x_i$  = inputs (data)
- $w_i$  = weights (importance)
- $b$  = bias (baseline)
- $z$  = output (score)

### Party Decision Example

Alex's formula:

$$\text{Score} = -2 \cdot \text{Distance} + 3 \cdot \text{Friends} - 5$$

If Score > 0 → GO, else STAY



### Geometric Interpretation

Decision boundary: line where Score = 0

$$-2d + 3f - 5 = 0 \Rightarrow f = \frac{2d + 5}{3}$$



### 3. Activation Functions: The Secret to Power

#### The Linearity Problem

Without activation: multiple neurons = another linear function!

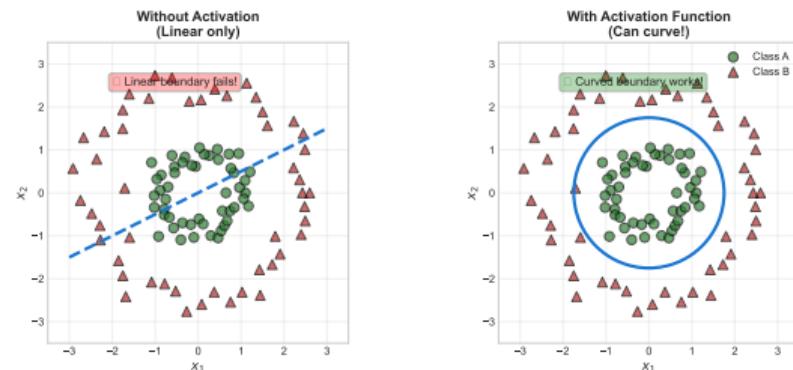
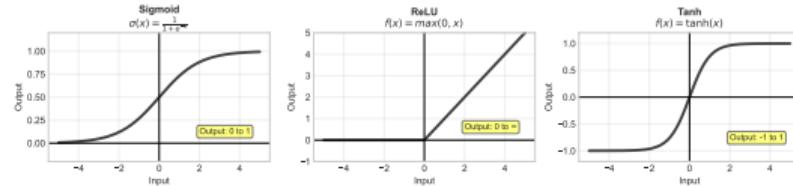
$$z_2 = w_1(w_2x + b_2) + b_1 = (w_1w_2)x + (w_1b_2 + b_1)$$

**Solution:** Add non-linear activation

$$a = f(z) = f\left(\sum_i w_i x_i + b\right)$$

#### Common Activations

- **Sigmoid:**  $\sigma(z) = \frac{1}{1+e^{-z}}$  (0 to 1)
- **ReLU:**  $\max(0, z)$  (modern standard)
- **Tanh:**  $\frac{e^z - e^{-z}}{e^z + e^{-z}}$  (-1 to 1)
- **Leaky ReLU:**  $\max(0.01z, z)$  (prevents dying)



**Critical:** Without activation, 100 layers = 1 neuron!

## 4. The XOR Crisis

### XOR Problem

Output 1 if inputs different, 0 if same

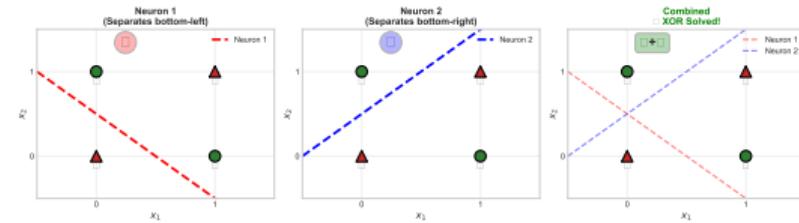
$x_1$	$x_2$	Output
0	0	0
0	1	1
1	0	1
1	1	0

**Challenge:** Draw ONE line separating 1's from 0's

**Result:** *Impossible!*

### Geometric Proof

- (0,1) and (1,0) must be on one side
- (0,0) and (1,1) on the other
- No straight line separates opposite corners!



### Historical Impact (1969)

Minsky & Papert proved single-layer networks cannot solve XOR

⇒ First AI Winter

Funding dried up for decades

**Limitation:** Single neurons only solve *linearly separable* problems

## 5. Hidden Layers: The Breakthrough

### The Solution

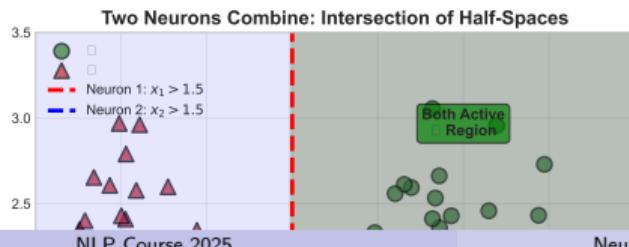
Use TWO neurons in hidden layer, combine outputs!

### Architecture

- Input layer: 2 neurons ( $x_1, x_2$ )
- Hidden layer: 2 neurons (two boundaries)
- Output layer: 1 neuron (combines)

### Geometric Intuition

- Hidden 1: Separates (0,0) from others
- Hidden 2: Separates (1,1) from others
- Output: Finds *intersection*
- Only (0,1) and (1,0) satisfy both!



### Forward Pass Example

Given weights:

- Hidden 1:  $w = [1, 1], b = -0.5$
- Hidden 2:  $w = [1, 1], b = -1.5$
- Output:  $w = [1, -1], b = 0$

For input (1, 0):

$$\begin{aligned} h_1 &= \sigma(1 \cdot 1 + 1 \cdot 0 - 0.5) \\ &= \sigma(0.5) \approx 0.62 \\ h_2 &= \sigma(1 \cdot 1 + 1 \cdot 0 - 1.5) \\ &= \sigma(-0.5) \approx 0.38 \\ y &= \sigma(1 \cdot 0.62 - 1 \cdot 0.38) \\ &= \sigma(0.24) \approx 0.56 \text{ (close to 1!)} \end{aligned}$$

**Why it works:** Each neuron learns different feature. Enough neurons → any boundary!

# 6. Backpropagation: How Networks Learn

## Credit Assignment Problem

Given output error, which weights to adjust by how much?

## The Algorithm (4 steps)

### 1. Forward Pass

$$z^{[l]} = W^{[l]} a^{[l-1]} + b^{[l]}, \quad a^{[l]} = f(z^{[l]})$$

### 2. Compute Error

$$L = \frac{1}{2} (y_{\text{pred}} - y_{\text{true}})^2$$

### 3. Backward Pass (chain rule)

$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial a} \cdot \frac{\partial a}{\partial z} \cdot \frac{\partial z}{\partial w}$$

### 4. Update Weights

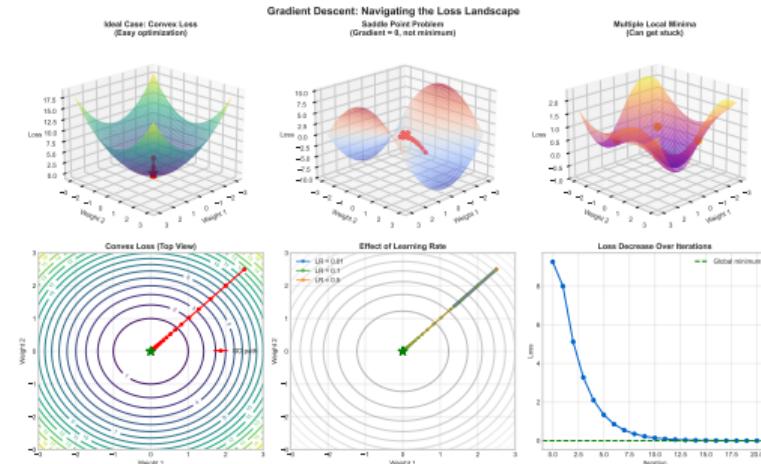
$$w \leftarrow w - \eta \frac{\partial L}{\partial w}$$

where  $\eta$  = learning rate (step size)

## Gradient Descent Intuition

Hiking in fog to reach valley:

- ➊ Feel slope under feet (gradient)
- ➋ Take step downhill (update)
- ➌ Repeat until can't go lower (converge)



## Historical Note

Invented 1970s, famous 1986  
(Rumelhart / Hinton / Williams)

## 7. Universal Approximation Theorem

### Cybenko's Theorem (1989)

Network with:

- One hidden layer
- Finite neurons
- Sigmoid activation

can approximate *any* continuous function to *any* accuracy!

### How It Works

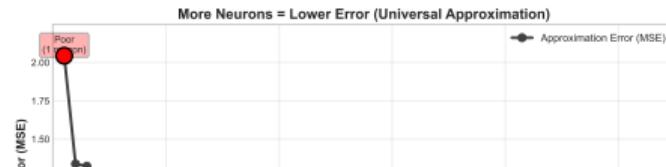
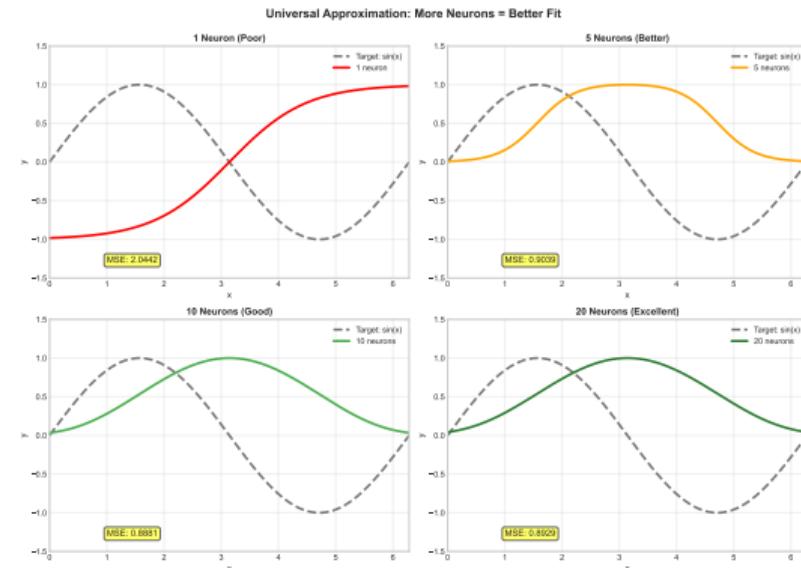
Each sigmoid = smooth step function

Position steps at different locations/heights → build any curve:

$$f(x) \approx \sum_{i=1}^n a_i \sigma(w_i x + b_i)$$

### Caveats

- Guarantees existence, not efficient learning
- May need exponential neurons
- Deep networks often better than wide



## 8. From Theory to Modern Practice

### Key Breakthroughs Timeline

- 1998 LeNet-5: First CNN, read checks
- 2012 AlexNet: ImageNet 26% → 16%
- 2015 ResNet: Skip connections, 152 layers
- 2017 Transformers: Attention revolutionized NLP
- 2022 ChatGPT: LLMs mainstream

### Why 2012 Was Different

- ① **ReLU**: Solved vanishing gradients
- ② **Dropout**: Prevents overfitting
- ③ **GPUs**: 50x faster training
- ④ **Big Data**: ImageNet 14M images
- ⑤ **Batch Norm**: Stabilizes layers

### Modern Architectures

- **CNNs**: Images (ResNet, EfficientNet)
- **RNNs**: Sequences (LSTM, GRU)
- **Transformers**: Everything (BERT, GPT, ViT)

### Real-World Applications Today

- Medical diagnosis
- Autonomous vehicles
- Drug discovery
- Language translation
- Code generation
- Art generation
- Speech recognition
- Recommendation systems
- Protein folding
- Climate modeling
- Fraud detection
- Robotics

### Key Innovation: Transfer Learning

Pre-train large model → Fine-tune for specific task  
Examples: BERT, GPT-3, DALL-E, AlphaFold

**Today:** Neural networks power most modern AI systems

# 9. Building Your First Network

## 7-Step Process

### 1. Define Problem

- Classification vs Regression?
- Input/output sizes?
- Target accuracy?

### 2. Prepare Data

- Split: 70% train, 15% val, 15% test
- Normalize:  $[0,1]$  or  $\text{mean}=0, \text{std}=1$
- Augment: flip, rotate, paraphrase

### 3. Design Architecture

- Start simple: 1-2 hidden, 32-128 neurons
- ReLU hidden, sigmoid/softmax output
- Add dropout (0.2-0.5)

### 4. Hyperparameters

- Learning rate: 0.001 (most critical!)
- Batch size: 32-256
- Optimizer: Adam
- Loss: CrossEntropy/MSE

### 5. Train

Forward  $\rightarrow$  Loss  $\rightarrow$  Backward  $\rightarrow$  Update

### 6. Debug Common Issues

Symptom	Solution
Loss not decreasing	Try 10x higher/lower LR
Train good, val bad	Dropout, more data
Loss = NaN	Lower LR, clip gradients

### 7. Evaluate

- Never touch test until final!
- Multiple metrics (Accuracy, F1, Precision)
- Visualize: confusion matrix, learning curves

## Best Practices

- Start simple, add complexity
- Log everything
- Save checkpoints
- Monitor training (TensorBoard)

# 10. Summary: The Complete Picture

## Essential Formulas

Neuron	$z = \sum_i w_i x_i + b$
Sigmoid	$\sigma(z) = 1/(1 + e^{-z})$
ReLU	$\max(0, z)$
Forward	$a^{[l]} = f(W^{[l]} a^{[l-1]} + b^{[l]})$
Loss	$L = \frac{1}{n} \sum (y_{pred} - y_{true})^2$
Gradient	$w \leftarrow w - \eta \frac{\partial L}{\partial w}$

## Logical Flow

Problem → Neuron → Activation → XOR Crisis → Hidden Layers  
→ Backprop → Theory → Practice → Applications

## Key Concepts Checklist

- Neuron = weighted sum
- Weights control importance
- Activation adds non-linearity
- Single neuron = linear
- Hidden layers = non-linear
- XOR impossible alone
- Backprop assigns credit
- Gradient descent optimizes

## What's Next

### Implement:

- Code from scratch (NumPy)
- Use frameworks (PyTorch/TensorFlow)

### Learn:

- Fast.ai, CS231n, Coursera
- Papers: LeNet, AlexNet, ResNet, Attention

### Build:

- Image classifier
- Text generator
- Game AI

## Resources

- Book: Deep Learning (Goodfellow et al.)
- Viz: playground.tensorflow.org
- Code: github.com/pytorch/examples
- Papers: arxiv-sanity.com