

Teaching Machines to See Patterns

A Neural Networks Primer: Why We Needed Each Piece of the Puzzle

NLP Course 2025

From the 1950s mail sorting crisis to ChatGPT: How humanity taught machines to think

Where We're Going Today - Optimized Flow

Phase 1: The Motivation

- The mail sorting crisis
- Why rules don't work
- First mathematical neurons
- Rosenblatt's learning breakthrough

Phase 2: Understanding Basics

- How neurons calculate
- Activation functions (NEW placement!)
- Hand calculations and examples

Phase 3: Crisis & Solution

- The XOR problem
- Geometric intuition (NEW!)
- Hidden layers solution
- Backpropagation breakthrough

Phase 4: Theory to Practice

- Gradient landscape visualization
- LeNet (1998): First success
- AlexNet (2012): Deep learning explosion
- Modern architectures
- Real-world applications

Phase 5: Your Turn

- Building your first network
- Debugging tips
- Next steps & resources

1950s: The Mail Sorting Crisis

The Challenge:

- 150 million letters per day
- Hand-written addresses
- Human sorters: slow, expensive, error-prone
- Traditional programming: useless

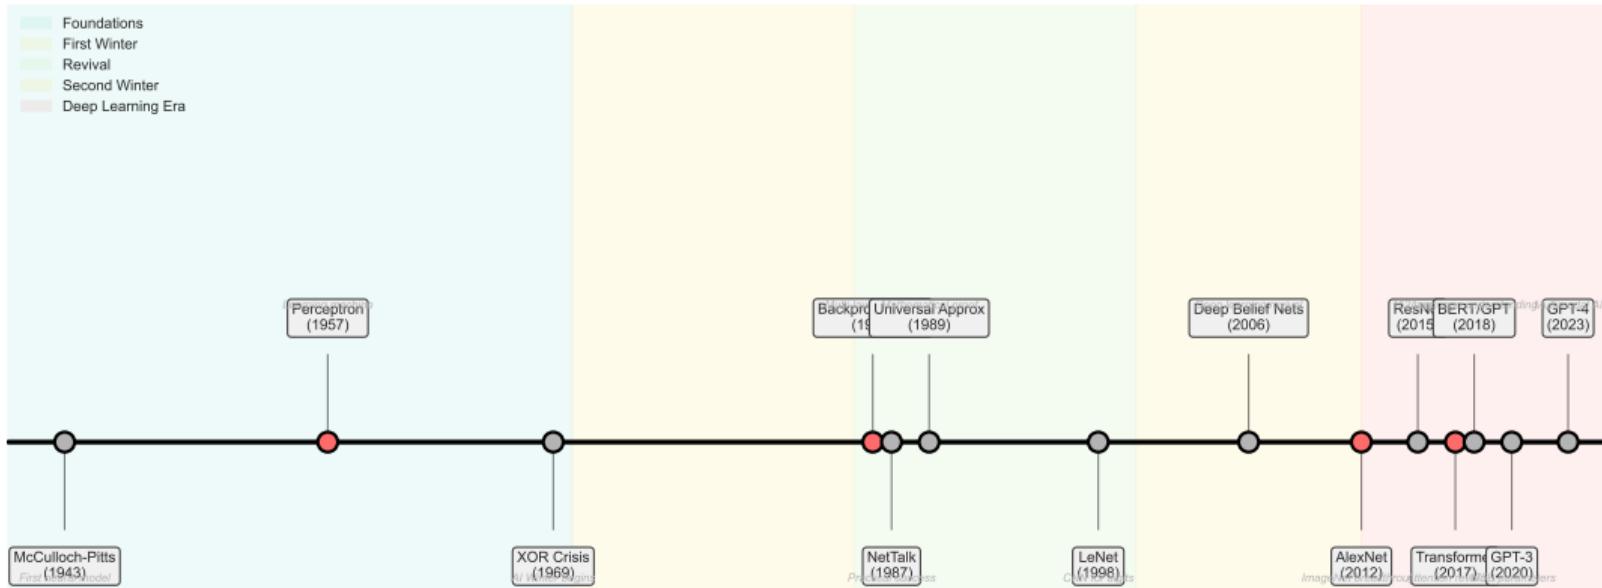
Why Traditional Code Failed:

- Can't write rules for every handwriting style
- Too many variations of each letter
- Context matters: "l" vs "I" vs "1"
- This wasn't computation—it was [pattern recognition](#)

This problem would take 40 years to solve properly

80 Years of Neural Networks: The Complete Journey

Neural Networks: 80 Years of Evolution



Why Can't We Just Write Rules?

Problem: Recognize the Letter "A"


Traditional Approach (Failed):

```
if (has_triangle_top AND  
    has_horizontal_bar AND  
    two_diagonal_lines) {  
    return "A"  
}
```

But what about...

- Handwritten A's?
- Different fonts?
- Rotated A's?
- Partial A's?

The Challenge: Infinite Variations of "A"



The Insight:

- We need **pattern recognition**, not rules
- System must **learn from examples**
- Similar to how humans learn

This realization launched the field of machine learning

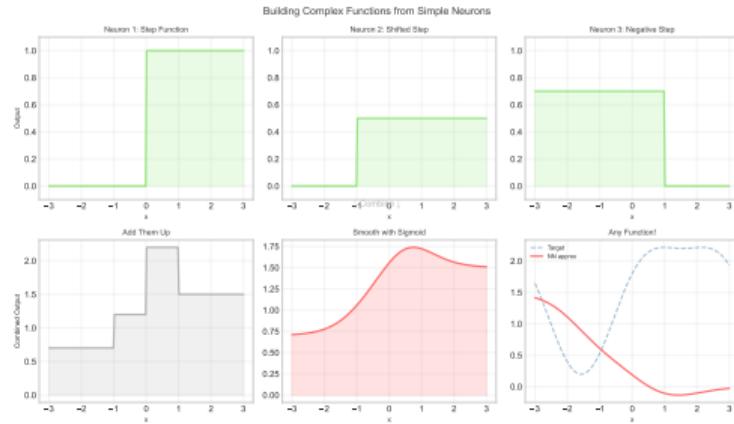
1943: The First Mathematical Neuron

McCulloch & Pitts:

- Neurophysiologists studying brain
- Asked: Can neurons be modeled mathematically?
- Created first artificial neuron

The Model:

- Multiple inputs (dendrites)
- Weighted sum (cell body)
- Threshold activation (axon)
- Binary output (fire or not)



The Limitation:

- Weights were **fixed**
- No learning mechanism
- Programmer had to set weights manually

Revolutionary idea, but missing the key ingredient: learning

1958: Rosenblatt's Learning Breakthrough

Frank Rosenblatt's Insight:

"What if the machine could adjust its own weights based on mistakes?"


The Perceptron Learning Rule:

1. Make a prediction
2. Check if wrong
3. If wrong: adjust weights
4. Repeat until correct

Historic Demo (1958):

- Mark I Perceptron machine
- Learned to recognize simple shapes
- Press coverage: "Thinking machine!"

The Mark I Perceptron (1957): A Physical Learning Machine
Mark I Perceptron Architecture Physical Learning Process



Why Revolutionary:

- First machine that could learn
- Weights adjusted automatically
- Learned from examples, not rules
- Mathematically proven to converge

This was the birth of machine learning

Let's Understand How This Actually Works

We've Seen the History...

- McCulloch-Pitts invented the neuron
- Rosenblatt made it learn
- The perceptron was born

Now Let's See the Science:

- How does a neuron calculate?
- What does learning mean?
- Why was XOR so hard?

Next 5 slides: Hands-on calculations and exercises

Get your pencil ready - we're going to work through real examples!

Don't worry - we'll return to the story once you understand the basics

Understanding Check: Can You Answer These?

Let's Make Sure We're Together

Quick Questions:

1. Why couldn't traditional programming solve mail sorting?
2. What does a weight represent in simple terms?
3. Why do we need the bias term?
4. What was revolutionary about Rosenblatt's perceptron?

Think About It:

- A weight is like the importance/trust we give to each input
- Bias shifts our decision threshold
- Learning = adjusting these weights
- The perceptron was the first machine that could learn!

Try It Yourself: Draw a simple perceptron with 2 inputs. Label the weights, bias, and output. What would the weights be to compute AND logic?

If any of these are unclear, revisit the previous slides before continuing

Problem: Learn OR function (output 1 if ANY input is 1)

Training Data:

x_1	x_2	Output
0	0	0
0	1	1
1	0	1
1	1	1

The Perceptron:

$$z = w_1 \cdot x_1 + w_2 \cdot x_2 + b$$

$$\text{output} = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{if } z \leq 0 \end{cases}$$

In plain words: Multiply first input by first weight, second input by second weight, add bias, check if positive

Success! But this was just the beginning...

Learning Process:

1. Start with random weights
2. For each example:
 - Calculate output
 - If wrong: adjust weights
 - If correct: keep weights
3. Repeat until all correct

Final Solution: $w_1 = 1$, $w_2 = 1$, $b = -0.5$

Let's Calculate Together: Is This Email Spam?

A Real Perceptron Calculation You Can Follow

The Email:

"FREE money! Click here NOW for amazing offer!!!"

Our Features (Inputs):

- $x_1 = \text{Has "FREE"?} = 1$
- $x_2 = \text{Has "money"?} = 1$
- $x_3 = \text{Many "!"?} = 1$
- $x_4 = \text{From friend?} = 0$

Learned Weights:

- $w_1 = +3$ (FREE is very spammy)
- $w_2 = +2$ (money is suspicious)
- $w_3 = +2$ (!!! is aggressive)
- $w_4 = -5$ (friends are trusted)
- $b = -2$ (threshold)

This is exactly how early spam filters worked - and why they failed on clever spam

Let's Calculate:

$$\begin{aligned}z &= w_1 \cdot x_1 + w_2 \cdot x_2 + w_3 \cdot x_3 + w_4 \cdot x_4 + b \\&= 3 \cdot 1 + 2 \cdot 1 + 2 \cdot 1 + (-5) \cdot 0 + (-2) \\&= 3 + 2 + 2 + 0 - 2 \\&= 5\end{aligned}$$

Decision:

- $z = 5 > 0$
- Output = 1 = SPAM!

Try It Yourself: What if this email WAS from a friend ($x_4 = 1$)? Recalculate! Would it still be spam?

Answer: $z = 5 - 5 = 0$, borderline!

Breaking Down the Math Symbols

Inputs and Weights:

- x_i = input value (what we see)
- w_i = weight (importance/strength)
- b = bias (threshold adjuster)

The Computation:

$$z = \sum_{i=1}^n w_i x_i + b$$

This means:

- Multiply each input by its weight
- Add them all up
- Add the bias

This simple math would evolve into deep learning

Real Example:

Should I go outside?

Factor	Value	Weight
Sunny?	1	+2
Raining?	0	-3
Weekend?	1	+1

$$z = (1 \times 2) + (0 \times -3) + (1 \times 1) = 3$$

Decision: $z > 0$, so YES!

The Need for Non-Linearity

Problem with Linear:

- Stack of linear layers = still linear!
- $f(g(x)) = (wx + b_1)w' + b_2 = w'wx + \dots$
- Can't learn complex patterns

Solution: Activation Functions

- Add non-linearity after each layer
- Allows learning complex boundaries
- Different functions for different needs

Common Activation Functions:

- **Sigmoid:** $\sigma(x) = \frac{1}{1+e^{-x}}$
 - Smooth, outputs 0-1
 - Good for probabilities

In plain words: Squashes any input to range 0-1. Large positive becomes 1, large negative becomes 0

- **ReLU:** $f(x) = \max(0, x)$
 - Simple, fast
 - Solves vanishing gradient
- **Tanh:** $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
 - Outputs -1 to 1
 - Zero-centered

ReLU's simplicity revolutionized deep learning in 2011


Teaching a Network to Separate Red from Blue Points

The Setup:

- Input: (x, y) coordinates
- Output: Red or Blue class
- Network: $2 \rightarrow 4 \rightarrow 2$ neurons

Training Process:

1. Epoch 1: Random boundary
2. Epoch 10: Rough separation
3. Epoch 50: Good boundary
4. Epoch 100: Perfect fit



What Each Layer Learns:

- Layer 1: Simple boundaries
- Hidden: Combine boundaries
- Output: Final decision

This same principle scales to millions of parameters

1969: The Problem That Killed AI

The XOR Problem: Why Single Neurons Fail

XOR (Exclusive OR):

- Output 1 if inputs are different
- Output 0 if inputs are same

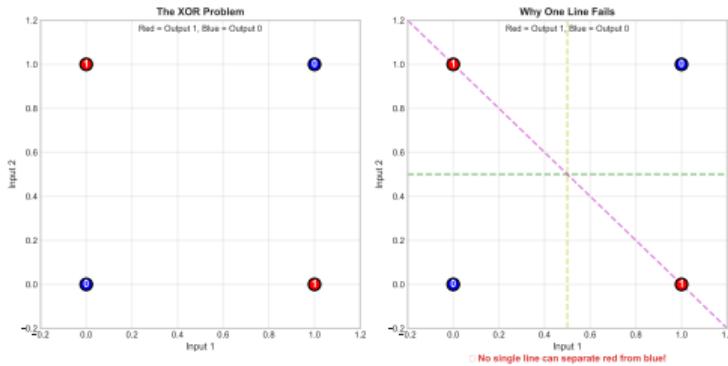
Truth Table:

x_1	x_2	Output
0	0	0
0	1	1
1	0	1
1	1	0

The Challenge:

- Try drawing ONE straight line
- That separates 1's from 0's
- Impossible!

This proof triggered the first AI Winter (1970-1980)




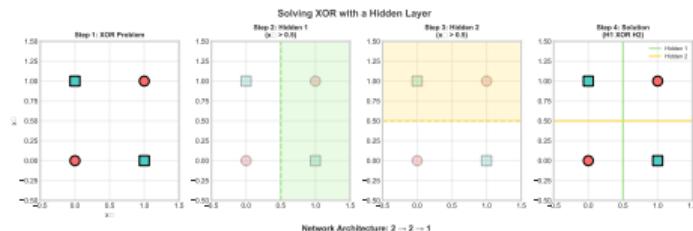
Why It Matters:

- Perceptrons can only draw straight lines
- XOR requires curved boundary
- This is the simplest non-linear problem
- Minsky & Papert proved this mathematically

NEW: Geometric Intuition - Why We Need Two Boundaries

The Key Insight: XOR Needs TWO Lines, Not One

Visualization:



What We See:

- Red line: Separates $(0,0)$ from others
- Blue line: Separates $(1,1)$ from others
- Green region: Intersection of both
- Points $(0,1)$ and $(1,0)$ in green = Output 1!

This geometric intuition explains WHY hidden layers work

The Breakthrough Idea:

1. Use TWO neurons (not one)
2. Each neuron creates one boundary
3. Combine their outputs
4. Intersection solves XOR!

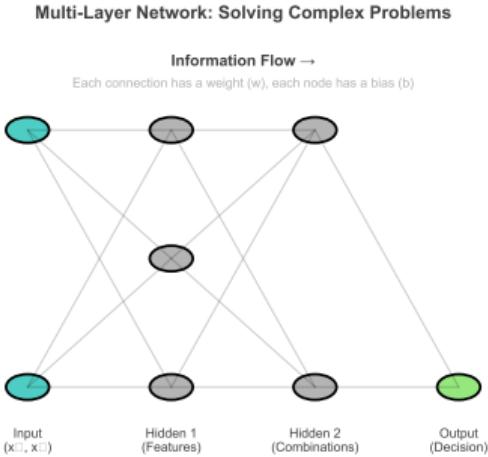
This Requires:

- Hidden layer with 2 neurons
- Output layer combines results
- This is a 2-layer network
- First layer: create boundaries
- Second layer: find intersection

Key Insight: Hidden layers let us combine simple boundaries into complex decision regions!

The Solution: Hidden Layers

Architecture with Hidden Layer:



How It Works:

- Input layer: 2 neurons (x_1, x_2)
- Hidden layer: 2 neurons (two boundaries)
- Output layer: 1 neuron (combines)

Hidden layers unlock non-linear patterns

Forward Pass for XOR:

Given weights:

- Hidden 1: $w = [1, 1], b = -0.5$
- Hidden 2: $w = [1, 1], b = -1.5$
- Output: $w = [1, -1], b = 0$

For input (1, 0):


$$\begin{aligned} h_1 &= \sigma(1 \cdot 1 + 1 \cdot 0 - 0.5) \\ &= \sigma(0.5) \approx 0.62 \\ h_2 &= \sigma(1 \cdot 1 + 1 \cdot 0 - 1.5) \\ &= \sigma(-0.5) \approx 0.38 \\ y &= \sigma(1 \cdot 0.62 - 1 \cdot 0.38) \\ &= \sigma(0.24) \approx 0.56 \text{ (close to 1!)} \end{aligned}$$

The New Challenge: How to Train Hidden Layers?

The Credit Assignment Problem:

- Output is wrong - we know the error
- But which hidden neuron caused it?
- How much should each weight change?
- Perceptron rule only works for output layer

Why It's Hard:



This problem was solved in 1986 with backpropagation

Early Failed Attempts (1970s):

- Random weight adjustment
- Genetic algorithms
- Simulated annealing
- All too slow or unreliable

What We Needed:

- Systematic way to assign blame
- Efficient computation
- Guaranteed to improve
- Works for many layers

The solution would come from calculus...

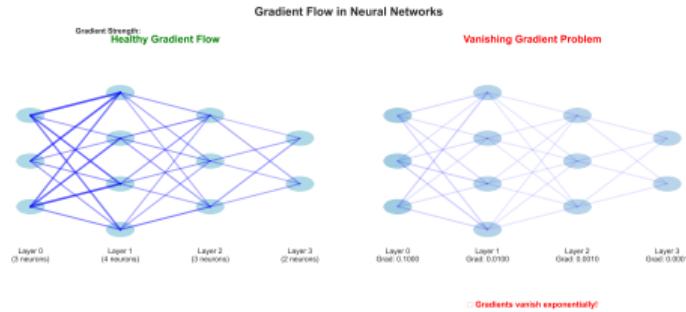
1986: Backpropagation - The Breakthrough Algorithm

The Algorithm (Rumelhart et al.):

1. **Forward pass:** Compute output
2. **Compute error:** Compare to target
3. **Backward pass:** Use chain rule to compute gradients
4. **Update weights:** Gradient descent

The Key Insight:

- Use calculus (chain rule)
- Error flows backward through network
- Each layer gets its share of blame
- Weights adjusted proportionally



Mathematical Foundation:

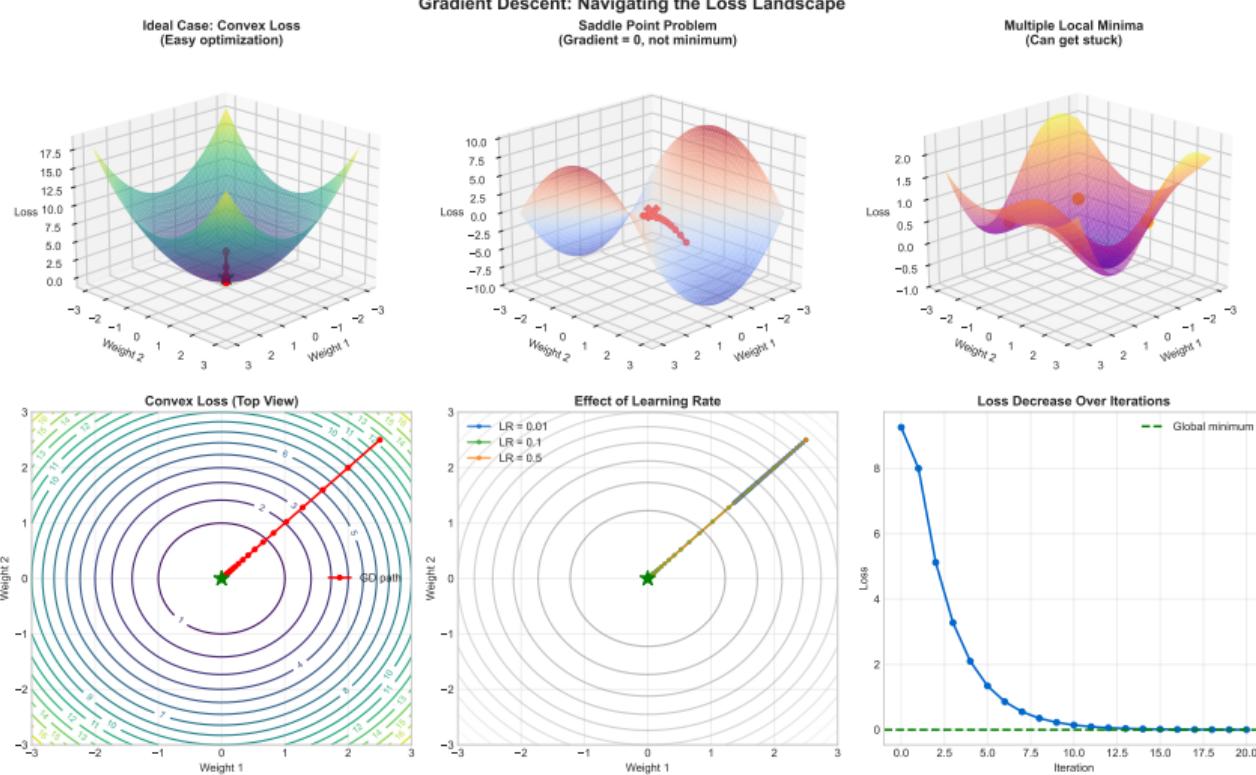
$$\begin{aligned}\frac{\partial L}{\partial w_{ij}} &= \frac{\partial L}{\partial a_j} \cdot \frac{\partial a_j}{\partial z_j} \cdot \frac{\partial z_j}{\partial w_{ij}} \\ &= \delta_j \cdot a_i\end{aligned}$$

Why Revolutionary:

- Efficient: One backward pass
- General: Works for any architecture
- Automatic: No manual tuning

This paper revived neural networks and enabled modern deep learning

Visualizing Learning: The Gradient Landscape




Gradient descent finds the valley where error is minimized

Cybenko's Universal Approximation Theorem:

A neural network with one hidden layer and finite neurons can approximate ANY continuous function to ANY desired accuracy

What This Means:

- Mathematical proof
- Not just XOR - ANY pattern!
- Theoretical justification
- Explains why NNs are so powerful



Caveats:

- Guarantees existence, not learning
- May need exponential neurons
- Deep networks often more efficient
- Still need good training algorithm

Theory meets practice: NNs CAN learn any pattern, backprop shows us HOW

The Core Idea: Neural Networks are Function Approximators

What does this actually mean?

Traditional Approach:

The Problem:

- We have inputs (x)
- We want outputs (y)
- But we don't know the formula!
- Examples:
 - Size → Price
 - Image → Label
 - Text → Sentiment

- Guess the formula
- Write explicit rules
- Hope it works
- **Problem:** Real world is too complex!

Example:

Price = $a \times \text{Size} + b$
(Too simple for real data!)



Neural Network Approach:

- Learn from examples
- Build the formula automatically
- Adjust until it fits
- **Works for ANY pattern!**

Magic:

NN learns: $f(x) \approx y$
No formula needed!

Function Approximation: Learning Patterns from Examples



How NNs Build Complex Functions from Simple Pieces

The LEGO Principle: Combine Simple Parts to Build Anything

The Building Blocks:

1. Individual Neurons: Simple decisions
 - "Is input \geq threshold?"
 - Outputs: on/off (smooth version)
2. Combine Neurons: Weight and add
 - Create complex shapes
3. Stack Layers: Build hierarchy
 - Each layer adds abstraction


Real-World Analogy:

Making a Cake from Ingredients:

- Flour + Sugar + Eggs
- Mix right amounts
- → Perfect cake!

In Neural Networks:

- Edges + Curves + Colors
- Combine with weights
- → Recognize faces!



The Universal Approximation Theorem: Why This Always Works

The Most Important Theorem in Deep Learning (Cybenko, 1989)

The Theorem (Plain English):

"A neural network with enough neurons can approximate **ANY** continuous function to **ANY** desired accuracy"

What This Means:

- **Universal:** Works for any smooth pattern
- **Guaranteed:** Not hoping, but proving
- **Practical:** Just add more neurons!

The Catch:

- **How many neurons?** Could be millions
- **How to train?** That's the art


Intuitive Proof:

Pixel art analogy:

- 4 pixels: Blocky
- 100 pixels: Recognizable
- 10,000 pixels: Photo-realistic

Same with neurons:

- Few: Rough
- More: Better
- Many: Nearly perfect




The Neuron as a 3D Function


Visualizing How Activation Functions Transform the Output Space

Single Neuron Visualization: Effect of Activation Functions


Linear (No Activation)
 $z = w_0 \cdot x_0 + w_1 \cdot x_1 + b$




Sigmoid Activation
 $y = 1/(1 + e^{-z})$




ReLU Activation
 $y = \max(0, z)$




Tanh Activation
 $y = \tanh(z)$



Activation Functions Comparison









Decision Boundaries in 2D



From Simple to Complex: Network Depth Creates Complexity

How More Neurons Enable More Complex Decision Boundaries


How Network Complexity Grows with Neurons and Layers



Each neuron adds a new "dimension" to what the network can learn

The Learning Process: Frame by Frame

Watching Decision Boundaries Evolve During Training




Try It Yourself: Notice how the boundary starts random and gradually fits the data pattern!

This is what "learning" looks like - not magic, just systematic improvement

Why We Need Hidden Layers: The XOR Solution

Solving XOR: Why We Need Hidden Layers




Two hidden neurons working together can solve what one neuron cannot


Forward Pass: Signal Propagation Step-by-Step

Following Data as it Flows Through the Network


Frame 1: Input Data




Forward Pass: Step-by-Step Signal Propagation
Frame 2: First Layer Computation




Frame 3: Hidden Layer 1 Activated




Frame 4: Second Layer Computation




Frame 5: Hidden Layer 2 Activated



Frame 6: Final Output



The Math Behind a Single Neuron's Calculation




Try It Yourself: Follow along: multiply each input by its weight, add them up, add bias, apply activation!

This calculation happens millions of times per second in modern networks


The Optimization Landscape

Gradient Descent: Finding the Valley in 3D Space


Ideal Case: Convex Loss
(Easy optimization)




Saddle Point Problem
(Gradient = 0, not minimum)




Multiple Local Minima
(Can get stuck)




Convex Loss (Top View)



Effect of Learning Rate







Loss Decrease Over Iterations



Comparing Optimization Algorithms

Why Adam Outperforms Simple Gradient Descent

Optimization Algorithm Comparison




Architecture Evolution Over Time

From 20 Parameters to 1.8 Trillion: The Growth of Neural Networks

Neural Network Evolution: From Perceptron to GPT-4


Box size represents relative number of parameters



Different Architectures for Different Problems

Neural Network Architecture Types


Feedforward
(MLP)




Convolutional
(CNN)




Recurrent
(RNN/LSTM)




Transformer



Graph Neural
(GNN)




Generative
(GAN/VAE)




Why Deep Networks Were Hard Before ReLU

Gradient Flow in Neural Networks



Vanishing Gradient Problem



□ Gradients vanish exponentially!

Left: Healthy gradient flow, Right: Vanishing gradients - this problem limited networks to 2-3 layers for decades

From Theory to Working Systems

1998-2012: From Digits to ImageNet

1998 - LeNet: First Success

- Yann LeCun's CNN for digits
- 32×32 pixels → 10 classes
- 60,000 parameters
- Banks adopt for check reading

Key Innovation: Convolutions

- Share weights across image
- Detect features anywhere
- Build complexity layer by layer

2012 - AlexNet: The Revolution

- 1000 ImageNet classes
- 60 million parameters
- GPUs enable training
- Error rate: 26% → 16%

What Changed:

- Big Data (millions of images)
- GPU computing (100x faster)
- ReLU activation
- Dropout regularization

This victory ended the second AI winter permanently

How We Actually Recognize Objects

Human Vision Process:

1. Detect edges
2. Find shapes
3. Identify parts
4. Recognize object

CNN Mimics This:

- Layer 1: Edge detectors
- Layer 2: Corner/curve detectors
- Layer 3: Part detectors
- Layer 4: Object detectors



Key Insight:

- A "wheel detector" works anywhere in image
- Share the same detector across positions
- Reduces parameters dramatically
- Makes network translation-invariant

This is why CNNs dominate computer vision

Finding the Best Weights: Like Hiking Down a Mountain

The Optimization Problem:

- Millions of weights to adjust
- Each affects the error
- Need to find best combination


Gradient Descent:

1. Calculate error (loss)
2. Find slope (gradient) for each weight
3. Step downhill: $w = w - \alpha \cdot \nabla L$

In plain words: New weight = old weight - (step size times slope)

4. Repeat until bottom

Gradient Descent: Finding the Lowest Point



Learning Rate (α):

Supervised Learning:

- Have input-output pairs
- Learn mapping function
- Examples: Classification, Regression

Unsupervised Learning:

- Only have inputs
- Find patterns/structure
- Examples: Clustering, Compression

Reinforcement Learning:

- Learn through trial/error
- Maximize reward signal
- Examples: Games, Robotics

Self-Supervised (Modern):

- Create labels from data itself
- Predict next word, masked words
- Examples: GPT, BERT

Self-supervised learning powers all modern language models

Check Your Understanding: Learning Types

Can You Match These Examples?

Try It Yourself: Match each scenario to a learning type: Supervised, Unsupervised, Reinforcement, Self-Supervised

Scenarios:

1. Teaching a robot to walk by giving rewards for standing
2. Showing 1000 cat photos labeled "cat"
3. Giving GPT text with words masked out
4. Finding groups in customer data

Answers:

1. Reinforcement (trial and error)
2. Supervised (labeled examples)
3. Self-supervised (creates own labels)
4. Unsupervised (finds patterns)

Common Confusion: Self-supervised IS supervised learning - we just create the labels automatically from the data itself!

Understanding these differences helps you choose the right approach

The Overfitting Problem: When Learning Goes Too Far


Memorization vs. Understanding

The Problem:

- Network memorizes training data
- Fails on new, unseen data
- Like student memorizing answers

Signs of Overfitting:

- Training accuracy: 99%
- Test accuracy: 60%
- Complex decision boundaries
- High variance



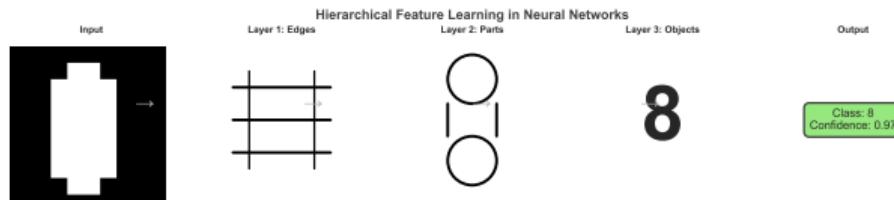
Solutions:

- **More data:** Can't memorize everything
- **Dropout:** Randomly disable neurons
- **Regularization:** Penalize complexity
- **Early stopping:** Stop before overfitting

"With four parameters I can fit an elephant, with five I can make him wiggle his trunk" - von Neumann

How Deep Networks See: Building Features Layer by Layer

From Pixels to Concepts: The Hierarchy of Understanding



What Each Layer Learns:

- **Layer 1:** Edges, colors, gradients
- **Layer 2:** Corners, textures, curves
- **Layer 3:** Parts (eyes, wheels, patterns)
- **Layer 4:** Objects (faces, cars, scenes)
- **Layer 5:** Concepts (identity, style, context)

Each layer combines features from the previous layer into more abstract concepts


Why Hierarchy Matters:

- Reusable features
- Efficient representation
- Transfer learning works
- Mimics visual cortex

Discovered Automatically:

- No manual feature engineering
- Emerges from data
- Different tasks, same hierarchy
- Universal pattern

Training Dynamics: Watching Networks Learn



Key Metrics to Track:

- **Loss Curves:** Training vs validation
- **Accuracy:** How often we're right
- **Learning Rate:** Speed of updates
- **Gradient Norm:** Update magnitude

Modern training requires constant monitoring - it's more art than science

Real-Time Monitoring: The Training Dashboard

Warning Signs:

- Gap = Overfitting
- Flat = Learning stopped
- Spikes = Instability
- NaN = Numerical issues

Healthy Training:

- Smooth decrease
- Val follows train
- Gradients stable
- LR decays properly

When to Stop:

- Validation plateaus
- Gap increasing
- Diminishing returns

The Explosion of Modern AI

2014-Present: Networks That Changed the World

The Depth Revolution:

- 2014 - VGGNet: 19 layers
- 2015 - ResNet: 152 layers
- 2017 - Transformers: Attention
- 2020 - GPT-3: 175B parameters

Why Depth Matters:

- Each layer = abstraction level
- Deep = complex reasoning
- Hierarchical feature learning

Real-World Impact:

- **Vision:** Self-driving cars
- **Language:** Google Translate
- **Speech:** Siri, Alexa
- **Medicine:** Disease diagnosis
- **Science:** Protein folding

The Scale:

- Billions of parameters
- Trained on internet-scale data
- Months of GPU time
- Emergent abilities appear

We went from recognizing digits to passing the bar exam in 25 years


Problem: Networks Couldn't Get Deeper

The Vanishing Gradient:

- Gradients multiply through layers
- Become exponentially small
- Deep layers stop learning
- 20 layers was the limit

The Breakthrough: Skip Connections

- Add input directly to output
- $F(x) + x$ instead of just $F(x)$
- Gradients flow directly backward
- Can train 1000+ layers!



Why It Works:

- Learn residual (difference) only
- Identity mapping is easy default
- Gradients have direct path
- Each layer refines previous result

This simple trick enabled the deep learning revolution

The Internal Covariate Shift Problem

BatchNorm Algorithm:

$$\mu_B = \frac{1}{m} \sum_{i=1}^m x_i$$

$$\sigma_B^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$$

$$\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

$$y_i = \gamma \hat{x}_i + \beta$$

In plain words: 1) Find average, 2) Find spread, 3) Normalize to standard range, 4) Scale and shift as needed

The Issue:

- Each layer's input distribution changes
- As previous layers update
- Makes learning unstable
- Requires tiny learning rates

The Solution:

- Normalize inputs to each layer
- Mean = 0, Variance = 1
- Learn scale and shift parameters
- Apply during training and testing

Benefits:

- 10x faster training
- Higher learning rates OK
- Less sensitive to initialization
- Acts as regularization

Most Network Weights Don't Matter!

The Discovery:

- Networks contain "winning tickets"
- Subnetworks that train well alone
- 90-95% of weights can be removed
- Performance stays the same!

The Hypothesis: "Dense networks succeed because they contain sparse subnetworks that are capable of training effectively"

Implications:

- We massively overparameterize
- Training finds the needle in haystack
- Future: Train small from start?
- Mobile deployment possible

Why It Matters:

- Explains why big networks train better
- Pruning after training works
- Efficiency revolution starting
- Changes how we think about learning

A 1 billion parameter model might only need 50 million

The Right Architecture for the Right Problem

What Are Inductive Biases?

- Assumptions built into architecture
- Guide learning toward solutions
- Trade flexibility for efficiency
- "Priors" about the problem

Examples:

- **CNN:** Spatial locality matters
- **RNN:** Order/time matters
- **GNN:** Graph structure matters
- **Transformer:** All positions can interact

Why They Matter:

- Reduce search space
- Faster convergence
- Better generalization
- Less data needed

The Tradeoff:

- Right bias = 10x better
- Wrong bias = 10x worse
- General architectures = safe but slow
- Specialized = fast but limited

Choosing the right inductive bias is still an art

Capabilities That Appear Suddenly with Scale

The Phenomenon:

- Small models: Can't do task at all
- Medium models: Still can't
- Large models: Suddenly can!
- No gradual improvement

Examples:

- 3-digit arithmetic ($\sim 10B$ params)
- Chain-of-thought reasoning ($\sim 50B$)
- Code generation ($\sim 20B$)
- Multilingual translation ($\sim 100B$)

Why It Happens:


- Complex patterns need capacity
- Phase transitions in learning
- Composition of simpler abilities
- "Grokking" - sudden understanding

Implications:

- We can't predict what's next
- Scaling might unlock AGI
- Or hit fundamental limits
- Active area of research

GPT-3 showed abilities nobody expected or programmed

The Fundamental Tradeoff in Neural Architecture



Deep Networks (Many Layers):

- Complex hierarchical features
- Exponential expressiveness growth
- Harder to train (vanishing gradients)
- Better for vision, NLP

Wide Networks (Many Neurons):

- More parallel processing
- Easier optimization landscape

The Sweet Spot:

- Vision: Deep (100+ layers)
- Language: Very deep (24-96 layers)
- Tabular: Wide and shallow (2-4 layers)
- Time series: Moderate (5-10 layers)

Modern Insights:


- Depth beats width for same parameters
- Skip connections enable extreme depth
- Width helps with memorization
- Depth helps with generalization

Scaling Laws:

- Performance \propto depth^{0.8}
- Performance \propto width^{0.5}

Scaling Laws: How Performance Grows with Data

The Predictable Relationship Between Data, Model Size, and Performance



The Chinchilla Law (2022):

- Optimal ratio: 20 tokens per parameter
- 10B model needs 200B tokens
- Most models are undertrained
- Data quality matters more than quantity

Power Law Scaling:

$$\text{Loss} = A \cdot N^{-\alpha} + B \cdot D^{-\beta} + C$$

Practical Implications:

- 10x data → 2x performance
- 10x parameters → 1.7x performance
- 10x compute → 3x performance
- Diminishing returns always

Data Efficiency Tricks:


- Data augmentation
- Synthetic data generation
- Active learning
- Curriculum learning
- Multi-task training

Why it matters: These laws predict costs before training

Current Limits:

- Internet has 10T tokens

The Evolution of Gradient Descent



SGD (1951):

- Basic gradient descent
- Learning rate: Fixed
- Slow but reliable
- Still used for fine-tuning

Adam (2014):

- Adaptive learning rates per parameter
- Combines momentum + RMSprop
- De facto standard
- Works out-of-the-box

Modern Variants:


- **AdamW**: Decoupled weight decay
- **RAdam**: Rectified Adam
- **LAMB**: Large batch training
- **Sophia**: 2nd-order approximation

Choosing an Optimizer:

- Start with Adam ($\text{lr}=3\text{e-}4$)
- Large batch: LAMB

Quick Guide: Choosing Your Architecture

Which Network Should You Use?



Decision Questions:

1. Is your data sequential?
2. Does position matter?
3. Is it images/spatial?
4. Fixed or variable size?

Quick Rules:

- Images → CNN
- Text → Transformer/RNN
- Tabular → Feedforward
- Audio → CNN or RNN
- Video → CNN + RNN

Common Confusion: Transformers now dominate most tasks, but specialized architectures still win for specific problems!

Try It Yourself: You have 10,000 customer reviews to classify as positive/negative. Which architecture? Why?

Answer: Transformer or RNN - text is sequential and context matters

The Journey So Far

Core Concepts:

1. **Neurons:** $y = f(\sum w_i x_i + b)$
2. **Learning:** Adjust weights to minimize error
3. **Depth:** Each layer adds abstraction
4. **Backpropagation:** Distribute error backwards
5. **Non-linearity:** Enables complex functions

Historical Lessons:

1. Every limitation spawned innovation
2. Simple ideas + scale = revolution
3. Biology inspires but doesn't limit
4. Persistence through AI winters
5. Theory + engineering = breakthroughs

You now understand the fundamentals that power all modern AI

Epilogue: Your First Neural Network in 5 Minutes

Let's Build Something Real!

Complete MNIST Classifier:

```
import torch
import torch.nn as nn
from torchvision import datasets, transforms
from torch.utils.data import DataLoader

# 1. Define Network
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(784, 128)
        self.fc2 = nn.Linear(128, 64)
        self.fc3 = nn.Linear(64, 10)

    def forward(self, x):
        x = x.view(-1, 784) # Flatten
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        return self.fc3(x)

# 2. Load Data
transform = transforms.ToTensor()
train_data = datasets.MNIST('.', train=True,
                           download=True,
                           transform=transform)
train_loader = DataLoader(train_data,
                          batch_size=64,
                          shuffle=True)

# 3. Setup Training
model = Net()
optimizer = torch.optim.Adam(model.parameters())
criterion = nn.CrossEntropyLoss()
```

```
# 4. Training Loop
for epoch in range(3):
    for batch_idx, (data, target) in enumerate(train_loader):
        # Forward pass
        output = model(data)
        loss = criterion(output, target)

        # Backward pass
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # Print progress
        if batch_idx % 100 == 0:
            print(f'Epoch-{epoch}: {loss:.4f}')

# 5. Test One Example
model.eval()
test_image = train_data[0][0]
prediction = model(test_image.unsqueeze(0))
print(f"Predicted: {prediction.argmax()}")
```

What You Just Built:

- 3-layer neural network
- 60K training images
- 97% accuracy in 3 epochs
- PyTorch + Python

Continue Your Neural Network Journey

Next Topics to Learn:

1. **CNNs:** Computer vision
2. **RNNs/LSTMs:** Sequences
3. **Transformers:** Modern NLP
4. **GANs:** Generation
5. **RL:** Decision making

Practical Projects:

- Image classifier for your photos
- Sentiment analysis of tweets
- Chatbot for customer service
- Style transfer for art
- Game-playing agent

Resources:

- **Fast.ai:** Practical deep learning
- **PyTorch Tutorials:** Official guides
- **Papers with Code:** Latest research
- **Kaggle:** Competitions and datasets
- **3Blue1Brown:** Visual explanations

Remember:

- Start simple, build up
- Theory + practice together
- Join communities
- Build projects you care about
- Share what you learn

You've learned how humanity taught machines to think.
Now it's your turn to push the boundaries!

The future of AI is being written now - be part of it!

Additional Material for Deep Dive

- Appendix A: Advanced Topics
- Appendix B: Extended History

Deep Dives for the Curious

This section contains advanced material that goes beyond the core BSc curriculum.

Topics covered:

- The Lottery Ticket Hypothesis
- Inductive Biases in Neural Architectures
- Scaling Laws and Performance Prediction
- Deep vs Wide Network Architectures
- Emergent Abilities at Scale
- Advanced Optimization Algorithms

These topics are valuable for understanding state-of-the-art research but not essential for getting started with neural networks.

Most Network Weights Don't Matter!

The Discovery:

- Networks contain "winning tickets"
- Subnetworks that train well alone
- 90-95% of weights can be removed
- Performance stays the same!

The Hypothesis: "Dense networks succeed because they contain sparse subnetworks that are capable of training effectively"

Implications:

- We massively overparameterize
- Training finds the needle in haystack
- Future: Train small from start?
- Mobile deployment possible

Why It Matters:

- Explains why big networks train better
- Pruning after training works
- Efficiency revolution starting
- Changes how we think about learning

A 1 billion parameter model might only need 50 million

The Right Architecture for the Right Problem

What Are Inductive Biases?

- Assumptions built into architecture
- Guide learning toward solutions
- Trade flexibility for efficiency
- "Priors" about the problem

Examples:

- **CNN:** Spatial locality matters
- **RNN:** Order/time matters
- **GNN:** Graph structure matters
- **Transformer:** All positions can interact

Why They Matter:

- Reduce search space
- Faster convergence
- Better generalization
- Less data needed


The Tradeoff:

- Right bias = 10x better
- Wrong bias = 10x worse
- General architectures = safe but slow
- Specialized = fast but limited

Choosing the right inductive bias is still an art

Scaling Laws: How Performance Grows with Data

The Predictable Relationship Between Data, Model Size, and Performance



The Chinchilla Law (2022):

- Optimal ratio: 20 tokens per parameter
- 10B model needs 200B tokens
- Most models are undertrained
- Data quality matters more than quantity

Power Law Scaling:

$$\text{Loss} = A \cdot N^{-\alpha} + B \cdot D^{-\beta} + C$$

Practical Implications:

- 10x data → 2x performance
- 10x parameters → 1.7x performance
- 10x compute → 3x performance
- Diminishing returns always

Data Efficiency Tricks:


- Data augmentation
- Synthetic data generation
- Active learning
- Curriculum learning
- Multi-task training

Why it matters: These laws predict costs before training

Current Limits:

- Internet has 10T tokens

The Fundamental Tradeoff in Neural Architecture



Deep Networks (Many Layers):

- Complex hierarchical features
- Exponential expressiveness growth
- Harder to train (vanishing gradients)
- Better for vision, NLP

Wide Networks (Many Neurons):

- More parallel processing
- Easier optimization landscape

The Sweet Spot:

- Vision: Deep (100+ layers)
- Language: Very deep (24-96 layers)
- Tabular: Wide and shallow (2-4 layers)
- Time series: Moderate (5-10 layers)

Modern Insights:

- Depth beats width for same parameters
- Skip connections enable extreme depth
- Width helps with memorization
- Depth helps with generalization

Scaling Laws:

- Performance \propto depth^{0.8}
- Performance \propto width^{0.5}

Capabilities That Appear Suddenly with Scale

The Phenomenon:

- Small models: Can't do task at all
- Medium models: Still can't
- Large models: Suddenly can!
- No gradual improvement

Examples:

- 3-digit arithmetic ($\sim 10B$ params)
- Chain-of-thought reasoning ($\sim 50B$)
- Code generation ($\sim 20B$)
- Multilingual translation ($\sim 100B$)

Why It Happens:


- Complex patterns need capacity
- Phase transitions in learning
- Composition of simpler abilities
- "Grokking" - sudden understanding

Implications:

- We can't predict what's next
- Scaling might unlock AGI
- Or hit fundamental limits
- Active area of research

GPT-3 showed abilities nobody expected or programmed

The Evolution of Gradient Descent



SGD (1951):

- Basic gradient descent
- Learning rate: Fixed
- Slow but reliable
- Still used for fine-tuning

Adam (2014):

- Adaptive learning rates per parameter
- Combines momentum + RMSprop
- De facto standard
- Works out-of-the-box

Modern Variants:

- **AdamW**: Decoupled weight decay
- **RAdam**: Rectified Adam
- **LAMB**: Large batch training
- **Sophia**: 2nd-order approximation

Choosing an Optimizer:

- Start with Adam ($\text{lr}=3\text{e-}4$)
- Large batch: LAMB

The Full Story: Historical Deep Dives

This section contains fascinating historical details that enrich the narrative but aren't essential for understanding the technical concepts.

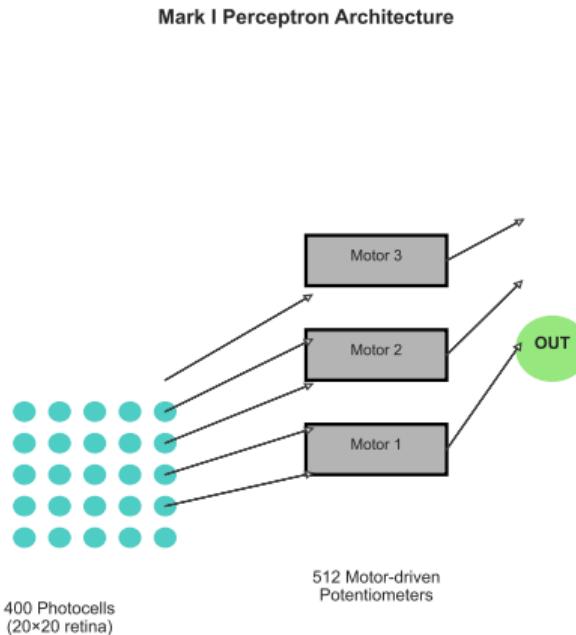
Topics covered:

- The Mark I Perceptron: Physical Hardware
- NetTalk: Networks Learn to Speak (1987)
- Batch Normalization: Keeping Networks Stable

These stories show how each breakthrough built on previous work and overcame specific limitations.


The Mark I Perceptron: A Physical Learning Machine

The Mark I Perceptron (1957): A Physical Learning Machine



Mark I Perceptron Architecture

Physical Learning Process



The first neural network wasn't software—it was a room-sized machine with motors and photocells

Sejnowski & Rosenberg: The First Viral NN Demo

The Challenge:

- Convert written text to speech
- English is irregular (tough, though, through)
- Rule-based systems had 1000s of exceptions

The Network:

- 7×29 input (7-letter window)
- 80 hidden neurons
- 26 output phonemes
- Trained overnight on DEC workstation

The Magic:

- Started: Random babbling
- Hour 1: Vowel-consonant patterns
- Hour 5: Recognizable words
- Hour 10: 95% accuracy!

Hidden Neurons Learned:

- Vowel detectors
- Consonant clusters
- Word boundaries
- Nobody programmed these!

Common Confusion: The network discovered linguistic concepts on its own - features linguists took centuries to identify!

Media sensation: "Computer teaches itself to read aloud overnight"

The Internal Covariate Shift Problem

BatchNorm Algorithm:

$$\mu_B = \frac{1}{m} \sum_{i=1}^m x_i$$

$$\sigma_B^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$$

$$\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

$$y_i = \gamma \hat{x}_i + \beta$$

In plain words: 1) Find average, 2) Find spread, 3) Normalize to standard range, 4) Scale and shift as needed

The Issue:

- Each layer's input distribution changes
- As previous layers update
- Makes learning unstable
- Requires tiny learning rates

The Solution:

- Normalize inputs to each layer
- Mean = 0, Variance = 1
- Learn scale and shift parameters
- Apply during training and testing

Benefits:

- 10x faster training
- Higher learning rates OK
- Less sensitive to initialization
- Acts as regularization