

Neural Language Models and Word Embeddings

Week 2: From Discrete Symbols to Continuous Representations

NLP Course 2025

Professional Template Edition

September 29, 2025

Learning Journey: From one-hot encodings to dense semantic vectors. Master Word2Vec, understand negative sampling, explore semantic arithmetic, and build the foundation for modern NLP.

By the end of this lecture, you will:

Understand

- Why one-hot encodings fail
- Distributional hypothesis principle
- Word2Vec architectures (CBOW vs Skip-gram)
- Negative sampling optimization
- Semantic vector arithmetic

Be Able To

- Train word embeddings
- Evaluate embedding quality
- Visualize semantic spaces
- Apply embeddings to NLP tasks
- Debug common issues

Core Message: Words are not just symbols - they carry meaning in their vector geometry

Prerequisites: Basic linear algebra, probability theory, neural network fundamentals

Part 1: Motivation & Problem Setting

Why do we need word embeddings?

Interactive Exercise: Word Association

Complete these sentences - What word naturally comes next?

1. The cat sat on the _____

Interactive Exercise: Word Association

Complete these sentences - What word naturally comes next?

1. The cat sat on the _____
mat, floor, chair

2. Coffee with milk and _____

Interactive Exercise: Word Association

Complete these sentences - What word naturally comes next?

1. The cat sat on the _____
mat, floor, chair

2. Coffee with milk and _____
sugar, cream

3. The capital of France is _____

Interactive Exercise: Word Association

Complete these sentences - What word naturally comes next?

1. The cat sat on the _____
mat, floor, chair

2. Coffee with milk and _____
sugar, cream

3. The capital of France is _____
Paris

4. She was happy but also _____

Interactive Exercise: Word Association

Complete these sentences - What word naturally comes next?

1. The cat sat on the _____
mat, floor, chair

2. Coffee with milk and _____
sugar, cream

3. The capital of France is _____
Paris

4. She was happy but also _____
sad, anxious, tired

5. King is to queen as man is to _____

Interactive Exercise: Word Association

Complete these sentences - What word naturally comes next?

1. The cat sat on the _____
mat, floor, chair

2. Coffee with milk and _____
sugar, cream

3. The capital of France is _____
Paris

4. She was happy but also _____
sad, anxious, tired

5. King is to queen as man is to _____
woman

6. Python is a programming _____

Interactive Exercise: Word Association

Complete these sentences - What word naturally comes next?

1. The cat sat on the _____
mat, floor, chair

2. Coffee with milk and _____
sugar, cream

3. The capital of France is _____
Paris

4. She was happy but also _____
sad, anxious, tired

5. King is to queen as man is to _____
woman

6. Python is a programming _____
language

Interactive Exercise: Word Association

Complete these sentences - What word naturally comes next?

1. The cat sat on the _____
mat, floor, chair

2. Coffee with milk and _____
sugar, cream

3. The capital of France is _____
Paris

4. She was happy but also _____
sad, anxious, tired


5. King is to queen as man is to _____
woman

6. Python is a programming _____
language

Key Insight: You use semantic understanding to predict words!

Humans naturally understand word relationships - how can we teach this to machines?

Historical Context: The Evolution Journey



1980s-2000s

- N-gram models

NLP Course 2025 (Professional Template Edition)

2003-2013

- Neural LMs

Neural Language Models and Word Embeddings

2013-2017

- Word2Vec era

2017-Present

- Transformers

September 29, 2025

The Fundamental Problem: One-Hot Encoding

Traditional Approach

- Vocabulary size: 50,000 words
- Each word = unique ID
- One-hot vectors:

cat =	[0,0,1,0,0,...,0]
dog =	[0,0,0,1,0,...,0]
car =	[0,0,0,0,1,...,0]
democracy =	[0,0,0,0,0,...,1]

Mathematical Issues

Cosine similarity between any two words:

$$\cos(\vec{w}_i, \vec{w}_j) = \frac{\vec{w}_i \cdot \vec{w}_j}{|\vec{w}_i| \cdot |\vec{w}_j|} = 0$$

All words are orthogonal!

Critical Problems:

- No similarity information
- $\text{distance}(\text{cat}, \text{dog}) = \text{distance}(\text{cat}, \text{car})$
- 50,000-dimensional vectors!
- Can't generalize patterns

Storage nightmare:

- 1M vocabulary = 1M dimensions
- Neural network input: $1M \times \text{hidden size}$
- Impossible to scale

One-hot encoding treats all words as equally different - destroying semantic information

Why This Problem Matters

Example Task: Sentiment Analysis

Training data:

- “This movie is great” → Positive
- “This film is excellent” → ?

With one-hot encoding:

- “great” and “excellent” unrelated
- Can’t transfer knowledge
- Need examples for every word

With embeddings:

- Similar words → similar vectors
- Knowledge transfers automatically
- Generalizes to unseen combinations

Real-World Impact

Search Engines:

- Query: “cheap cars”
- Finds: “affordable vehicles”

Translation:

- Unseen word combinations
- Cross-lingual transfer

Recommendations:

- Similar content discovery
- User preference modeling

Chatbots:







- Intent understanding
- Paraphrase detection

Dense embeddings enable knowledge transfer - the foundation of modern NLP success

The Solution: Dense Distributed Representations

..../figures/sparse_to_dense.pdf

Word Embeddings: Real-World Impact Across Industries



Part 2: Core Theory & Architecture

How Word2Vec learns meaning from context

“You shall know a word by the company it keeps”

- J.R. Firth (1957)

Core Principle: Words appearing in similar contexts have similar meanings

Example Sentences:

- The **cat** sat on the mat
- The **dog** sat on the floor
- A **cat** chased the mouse
- A **dog** chased the ball
- The **cat** is sleeping
- The **dog** is sleeping

Shared Contexts:

- Both follow “The” and “A”
- Both precede “sat”, “chased”, “is”
- Similar syntactic positions

Context Windows:

Window size = 2:

The	cat	mllavender3sat	on	the
←context		target	context→	

Mathematical View:

$$\text{context}(\text{cat}) = \{\text{the}, \text{sat}, \text{on}, \text{chased}, \text{is}, \dots\}$$

$$\text{context}(\text{dog}) = \{\text{the}, \text{sat}, \text{on}, \text{chased}, \text{is}, \dots\}$$




similar contexts \Rightarrow similar vectors


Word2Vec: Two Complementary Architectures

Word2Vec Architecture Comparison

CBOW: Context → Center



Skip-gram: Center → Context



CBOW (Continuous Bag-of-Words)

- Don't predict context from center

Objective: Maximize probability of context words given center word

Formal Objective Function:

$$J(\theta) = \frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log P(w_{t+j} | w_t)$$

Where:

- T = Total words in corpus
- c = Context window size
- w_t = Center word at position t
- w_{t+j} = Context word at offset j

Probability using Softmax:

$$P(w_O | w_I) = \frac{\exp(v_{w_O}^T \cdot v_{w_I})}{\sum_{w=1}^{|V|} \exp(v_w^T \cdot v_{w_I})}$$

Two Embedding Matrices:

- $W_{in} \in \mathbb{R}^{|V| \times d}$ (input embeddings)
- $W_{out} \in \mathbb{R}^{|V| \times d}$ (output embeddings)

Computational Problem:

- Denominator sums over entire vocabulary
- 50,000 words = 50,000 exponentials!
- Too expensive per training step

The softmax normalization is beautiful mathematically but computationally prohibitive

Negative Sampling: Making Training Feasible

Key Idea: Convert to binary classification problem

Instead of Full Softmax:

- Classify real vs fake pairs
- Positive: (center, actual context)
- Negative: (center, random word)
- Much faster computation

Example:

- Center word: "sat"
- Positive: (sat, cat) $\rightarrow 1$
- Negative samples:
 - (sat, democracy) $\rightarrow 0$
 - (sat, quantum) $\rightarrow 0$
 - (sat, purple) $\rightarrow 0$

New Objective Function:

$$\log \sigma(v_{w_0}^T \cdot v_{w_l}) + \sum_{i=1}^k \mathbb{E}_{w_i \sim P_n(w)} [\log \sigma(-v_{w_i}^T \cdot v_{w_l})]$$

Where:




- $\sigma(x) = \frac{1}{1+e^{-x}}$
- $k =$ number of negative samples (typically 5-20)
- $P_n(w) =$ noise distribution $\propto f(w)^{3/4}$

Computational Speedup: From $O(\text{---V---})$ to $O(k)$ where $k \ll \text{---V---}$

Negative sampling reduces 50,000 computations to just 5-20 - enabling practical training

Training Dynamics: How Embeddings Evolve

Word2Vec Training Dynamics



Implementation: Critical Details

Preprocessing:

- Lowercase text
- Remove punctuation (optional)
- Build vocabulary
- Subsampling threshold: 10^{-3}

Initialization:

```
1 # Random init [-0.5, 0.5] / dim
2 W_in = (np.random.rand(V, d) - 0.5) / d
3 W_out = np.zeros((V, d))
```

Hyperparameters:

- Embedding dim: 100-300
- Window size: 5-10
- Negative samples: 5-20
- Min word frequency: 5

These implementation details make the difference between mediocre and excellent embeddings

Optimization Tricks:

1. Hierarchical Softmax:

- Binary tree over vocabulary
- $O(\log V)$ instead of $O(V)$
- Good for large vocabularies

2. Subsampling Frequent Words:

- "the", "a", "is" less informative
- Probability of keeping: $P(w_i) = \sqrt{\frac{t}{f(w_i)}}$
- Speeds training, improves quality

3. Dynamic Window:

- Actual window: `random(1, window_size)`
- Gives more weight to closer words

Checkpoint Quiz: Understanding Check

Test your understanding so far:

Questions:

Q1: Why do we need dense embeddings?

Q2: What's the key insight of distributional hypothesis?

Q3: Skip-gram vs CBOW - which to use?

Q4: Why negative sampling?

Checkpoint Quiz: Understanding Check

Test your understanding so far:

Questions:

Q1: Why do we need dense embeddings?

Q2: What's the key insight of distributional hypothesis?

Q3: Skip-gram vs CBOW - which to use?

Q4: Why negative sampling?

Answers:

A1: One-hot is sparse, no similarity info, can't generalize

A2: Words with similar contexts have similar meanings

A3: Skip-gram for quality, CBOW for speed/frequent words

A4: Avoid expensive softmax over entire vocabulary

Checkpoint Quiz: Understanding Check

Test your understanding so far:

Questions:

Q1: Why do we need dense embeddings?

Q2: What's the key insight of distributional hypothesis?

Q3: Skip-gram vs CBOW - which to use?

Q4: Why negative sampling?

Answers:

A1: One-hot is sparse, no similarity info, can't generalize

A2: Words with similar contexts have similar meanings

A3: Skip-gram for quality, CBOW for speed/frequent words

A4: Avoid expensive softmax over entire vocabulary

If you understand these, you grasp the core of Word2Vec!

These four concepts form the foundation of all modern word embedding methods

Part 3: Properties & Evaluation

The surprising geometry of meaning

The Magic: Semantic Vector Arithmetic

Semantic Arithmetic: Mathematical Operations on Meaning

Gender Relationship

$$\begin{matrix} \text{King} \\ \boxed{\text{Blue}} \end{matrix} - \begin{matrix} \text{Man} \\ \boxed{\text{Orange}} \end{matrix} + \begin{matrix} \text{Woman} \\ \boxed{\text{Teal}} \end{matrix} = \begin{matrix} \text{Queen} \\ \boxed{\text{Green}} \end{matrix}$$

Capital Cities

$$\begin{matrix} \text{Paris} \\ \boxed{\text{Blue}} \end{matrix} - \begin{matrix} \text{France} \\ \boxed{\text{Orange}} \end{matrix} + \begin{matrix} \text{Italy} \\ \boxed{\text{Teal}} \end{matrix} = \begin{matrix} \text{Rome} \\ \boxed{\text{Green}} \end{matrix}$$


Verb Conjugation

$$\begin{matrix} \text{Walking} \\ \boxed{\text{Blue}} \end{matrix} - \begin{matrix} \text{Walk} \\ \boxed{\text{Orange}} \end{matrix} + \begin{matrix} \text{Swim} \\ \boxed{\text{Teal}} \end{matrix} = \begin{matrix} \text{Swimming} \\ \boxed{\text{Green}} \end{matrix}$$

Comparative Forms

$$\begin{matrix} \text{Bigger} \\ \boxed{\text{Blue}} \end{matrix} - \begin{matrix} \text{Big} \\ \boxed{\text{Orange}} \end{matrix} + \begin{matrix} \text{Small} \\ \boxed{\text{Teal}} \end{matrix} = \begin{matrix} \text{Smaller} \\ \boxed{\text{Green}} \end{matrix}$$

Word Embeddings in 3D Space



Exploring Neighborhoods: Semantic Similarity

Finding nearest neighbors reveals semantic understanding:

Query: “king”

1. queen (0.72)
2. prince (0.70)
3. monarch (0.68)
4. emperor (0.65)
5. royal (0.64)

Query: “beautiful”

1. gorgeous (0.83)
2. lovely (0.81)
3. stunning (0.78)
4. pretty (0.76)
5. wonderful (0.72)

Query: “computer”

1. laptop (0.78)
2. PC (0.75)
3. software (0.71)
4. desktop (0.70)
5. machine (0.68)

Query: “Microsoft”

1. Google (0.75)
2. Apple (0.73)
3. IBM (0.71)
4. Amazon (0.70)
5. Intel (0.68)

Similarity = Cosine Distance in Embedding Space

Evaluation: How Good Are Your Embeddings?

1. Intrinsic Evaluation

Word Similarity:

- WordSim-353
- SimLex-999
- MEN dataset

Analogies:

- Google analogy test
- BATS dataset
- Semantic/syntactic

Metrics:

- Spearman correlation
- Accuracy@1, @5

2. Extrinsic Evaluation

Downstream Tasks:

- Sentiment analysis
- Named entity recognition
- POS tagging

Improvements:

- +3-5% accuracy typical
- Faster convergence
- Better generalization

Transfer Learning:

- Cross-domain
- Cross-lingual

3. Visualization

Techniques:

- t-SNE projection
- PCA analysis
- UMAP

Quality Checks:

- Cluster coherence
- Outlier detection
- Coverage analysis

Tools:

- TensorBoard
- Embedding Projector
- Custom notebooks

Good embeddings: 0.6+ correlation on similarity, 80%+ on analogies, clear visual clusters

Analogy Categories: Testing Semantic Understanding

Google Analogy Test Set: 19,544 questions across categories

Semantic (8,869 questions)

Capital-Country (506):

- Athens : Greece :: Tokyo : Japan
- Paris : France :: Rome : Italy

Currency (866):

- USA : dollar :: Japan : yen
- Europe : euro :: UK : pound

Family (506):

- brother : sister :: father : mother
- uncle : aunt :: nephew : niece

Typical Performance: Word2Vec on 1B word corpus

- Semantic analogies: 75-85% accuracy
- Syntactic analogies: 70-80% accuracy

Syntactic (10,675 questions)

Plural (1,332):

- cat : cats :: dog : dogs
- child : children :: person : people

Past Tense (1,560):

- walk : walked :: run : ran
- see : saw :: go : went

Comparative (1,332):

- good : better :: bad : worse
- big : bigger :: small : smaller

Fundamental Limitations:

1. Single Vector per Word

- “bank” (river) = “bank” (financial)
- No polysemy handling
- Context-independent

2. Out-of-Vocabulary

- Can't handle new words
- Misspellings break
- Rare words poorly represented

3. Bias in Data

- “doctor” closer to “he”
- “nurse” closer to “she”
- Amplifies stereotypes

Solutions & Evolution:

FastText (2016):

- Character n-grams
- Handles OOV words
- Better for morphology

ELMo (2018):

- Context-dependent
- Different vectors per use
- Bidirectional LSTM

BERT (2018):

- Transformer-based
- Deeply contextual
- Masked language modeling

Word2Vec's limitations directly inspired the transformer revolution

Part 4: Practical Implementation

From theory to practice

Implementation: Training Word2Vec

```
1 import gensim.downloader as api
2 from gensim.models import Word2Vec
3
4 # Load corpus (or use your own)
5 corpus = api.load('text8') # Wikipedia sample
6
7 # Train Word2Vec model
8 model = Word2Vec(
9     sentences=corpus,
10    vector_size=300,           # Embedding dimension
11    window=5,                 # Context window
12    min_count=5,              # Ignore rare words
13    workers=4,                # Parallel training
14    sg=1,                     # 1=skip-gram, 0=CBOW
15    negative=15,              # Negative samples
16    epochs=5,                 # Training iterations
17    seed=42                   # Reproducibility
18 )
19
20 # Save model
21 model.save("word2vec.model")
22
23 # Get vocabulary size
24 print(f"Vocabulary: {len(model.wv)} words")
25 print(f"Embedding shape: {model.wv.vectors.shape}")
```

Gensim makes Word2Vec incredibly accessible - production-quality embeddings in minutes

Using Pre-trained Embeddings

```
1 # Load pre-trained model
2 model = Word2Vec.load("word2vec.model")
3
4 # 1. Get word vector
5 king_vector = model.wv['king']
6 print(f"Vector for 'king': {king_vector[:5]}...")
7
8 # 2. Find similar words
9 similar = model.wv.most_similar('computer', topn=5)
10 for word, score in similar:
11     print(f"{word}: {score:.3f}")
12
13 # 3. Word analogies
14 result = model.wv.most_similar(
15     positive=['king', 'woman'],
16     negative=['man'],
17     topn=3
18 )
19 print(f"king - man + woman = {result[0][0]}")
20
21 # 4. Similarity scores
22 sim = model.wv.similarity('cat', 'dog')
23 print(f"Similarity(cat, dog) = {sim:.3f}")
24
25 # 5. Odd one out
26 odd = model.wv.doesnt_match(['breakfast', 'lunch', 'dinner', 'car'])
27 print(f"Odd one out: {odd}")
```

Pre-trained embeddings can be used immediately for various NLP tasks

Data Preparation:

- **Corpus Size:** Minimum 1M words
- **Preprocessing:**
 - Lowercase consistently
 - Keep punctuation (context!)
 - Handle numbers: "123" → "NUM"
- **Sentence Boundaries:** Respect them
- **Domain-Specific:** Add your data

Hyperparameter Tuning:

- **Dimension:** 100 (small), 300 (standard)
- **Window:** 5 (syntax), 10 (semantics)
- **Min Count:** 5-10 typical
- **Negative:** 5-20 (more is better)

Training Strategy:

- **Epochs:** 5-10 for large corpus
- **Learning Rate:** Start 0.025, decay
- **Subsampling:** $t = 10^{-5}$ typical
- **Threads:** Use multiple cores

Quality Checks:

- Monitor loss curve
- Check known analogies
- Visualize samples with t-SNE
- Test on downstream task early
- Compare against baselines

Pro Tip: Start with pre-trained, fine-tune on your domain

Good embeddings require careful attention to both data and hyperparameters

Problem 1: Poor Analogies

- **Symptom:** Random results
- **Causes:**
 - Too small corpus
 - Too few epochs
 - Window too small
- **Fix:** More data, longer training

Problem 2: Memory Issues

- **Symptom:** Out of memory
- **Causes:**
 - Vocabulary too large
 - Dimension too high
- **Fix:** Increase min_count, reduce dim

Problem 3: Slow Training

- **Symptom:** Takes days
- **Causes:**
 - Single thread
 - No negative sampling
 - Full softmax
- **Fix:** Use workers, negative sampling

Problem 4: Domain Mismatch

- **Symptom:** Works poorly on your data
- **Causes:**
 - Pre-trained on different domain
 - Technical jargon missing
- **Fix:** Fine-tune or train from scratch

Golden Rule: Always validate on your specific use case

Most embedding issues stem from data quality or hyperparameter mismatches

Integration: Embeddings in Neural Networks

```
1 import torch
2 import torch.nn as nn
3 import numpy as np
4
5 class SentimentClassifier(nn.Module):
6     def __init__(self, word2vec_model, num_classes=2):
7         super().__init__()
8
8         # Load pre-trained embeddings
9         weights = word2vec_model.wv.vectors
10        vocab_size, embed_dim = weights.shape
11
12        # Embedding layer (frozen initially)
13        self.embedding = nn.Embedding(vocab_size, embed_dim)
14        self.embedding.weight = nn.Parameter(torch.FloatTensor(weights))
15        self.embedding.weight.requires_grad = False # Freeze
16
17        # Classification layers
18        self.lstm = nn.LSTM(embed_dim, 128, batch_first=True)
19        self.dropout = nn.Dropout(0.3)
20        self.fc = nn.Linear(128, num_classes)
21
22    def forward(self, x):
23        embeds = self.embedding(x) # Use pre-trained
24        lstm_out, _ = self.lstm(embeds)
25        pooled = lstm_out[:, -1, :] # Last hidden state
26        return self.fc(self.dropout(pooled))
```

Pre-trained embeddings provide excellent initialization for downstream tasks

`.../figures/applications_gallery.pdf`

Timeline of Impact:

2013: Word2Vec Published

- Mikolov et al. papers
- Open-source release
- Immediate adoption

2014-2015: Rapid Adoption

- GloVe competition (Stanford)
- Industry integration
- 10,000+ citations

2016-2017: Refinements

- FastText (Facebook)
- Multilingual embeddings
- Domain-specific models

2018: Contextual Era

- ELMo introduces context
- BERT revolution begins
- GPT demonstrates scale

2019-Present: Foundation

- Still used in production
- Initialization for transformers
- 50,000+ citations
- Taught in every NLP course

Final Quiz: Mastery Check

Can you answer these advanced questions?

1. Why does Word2Vec use two embedding matrices (input and output)?

Final Quiz: Mastery Check

Can you answer these advanced questions?

1. Why does Word2Vec use two embedding matrices (input and output)?
→ Asymmetric relationship: predicting context from center
2. How does negative sampling approximate the softmax?

Final Quiz: Mastery Check

Can you answer these advanced questions?

1. Why does Word2Vec use two embedding matrices (input and output)?
→ Asymmetric relationship: predicting context from center
2. How does negative sampling approximate the softmax?
→ Binary classification: real pairs vs random pairs
3. Why does “king - man + woman = queen” work mathematically?

Final Quiz: Mastery Check

Can you answer these advanced questions?

1. Why does Word2Vec use two embedding matrices (input and output)?
→ Asymmetric relationship: predicting context from center
2. How does negative sampling approximate the softmax?
→ Binary classification: real pairs vs random pairs
3. Why does “king - man + woman = queen” work mathematically?
→ Relationships encoded as consistent vector offsets
4. What's the computational complexity: full softmax vs negative sampling?

Final Quiz: Mastery Check

Can you answer these advanced questions?

1. Why does Word2Vec use two embedding matrices (input and output)?
→ Asymmetric relationship: predicting context from center
2. How does negative sampling approximate the softmax?
→ Binary classification: real pairs vs random pairs
3. Why does “king - man + woman = queen” work mathematically?
→ Relationships encoded as consistent vector offsets
4. What's the computational complexity: full softmax vs negative sampling?
→ $O(V^2)$ vs $O(k)$, where $k \ll V$
5. Why can't Word2Vec handle “apple” (fruit) vs “Apple” (company)?

Final Quiz: Mastery Check

Can you answer these advanced questions?

1. Why does Word2Vec use two embedding matrices (input and output)?
→ Asymmetric relationship: predicting context from center
2. How does negative sampling approximate the softmax?
→ Binary classification: real pairs vs random pairs
3. Why does “king - man + woman = queen” work mathematically?
→ Relationships encoded as consistent vector offsets
4. What's the computational complexity: full softmax vs negative sampling?
→ $O(V^2)$ vs $O(k)$, where $k \ll V$
5. Why can't Word2Vec handle “apple” (fruit) vs “Apple” (company)?
→ Single vector per word type, no context dependence

If you can answer these, you truly understand Word2Vec's design and limitations

Core Concepts

- One-hot → Dense embeddings
- Distributional hypothesis
- Skip-gram architecture
- Negative sampling trick
- Semantic arithmetic
- Evaluation methods

Technical Skills

- Training Word2Vec
- Hyperparameter tuning
- Using pre-trained models
- Integration with neural nets
- Debugging embeddings

Key Insights

- Context defines meaning
- Geometry encodes semantics
- Relationships are vectors
- Efficiency enables scale
- Limitations inspire progress

Practical Impact

- Powers search engines
- Enables recommendations
- Foundation for transformers
- \$100B+ market impact
- 50,000+ citations

Word2Vec: Simple idea, profound impact, lasting legacy

Hands-On Exercises: week02_word_embeddings_lab.ipynb

Part 1: Training (45 min)

- Load and preprocess corpus
- Train Word2Vec from scratch
- Experiment with hyperparameters
- Compare CBOW vs Skip-gram
- Visualize training progress

Part 2: Exploration (30 min)

- Find nearest neighbors
- Test word analogies
- Visualize with t-SNE
- Explore semantic clusters
- Identify interesting patterns

Part 3: Application (45 min)

- Build similarity search engine
- Create analogy solver
- Simple sentiment classifier
- Document similarity system
- Performance evaluation

Bonus Challenges:

- Train on your own corpus
- Implement negative sampling
- Cross-lingual embeddings
- Bias analysis and debiasing
- Compare with GloVe/FastText

Goal: Hands-on mastery of word embeddings

Resources for Deep Dive

Essential Papers

- Mikolov et al. (2013a): Efficient Estimation
- Mikolov et al. (2013b): Distributed Representations
- Goldberg & Levy (2014): word2vec Explained
- Pennington et al. (2014): GloVe
- Bojanowski et al. (2016): FastText

Implementations

- Gensim (Python)
- TensorFlow/Keras
- PyTorch
- FastText library
- spaCy integration

Tutorials

- TensorFlow tutorials
- PyTorch examples
- Gensim documentation

Pre-trained Models

- Google News (3M words)
- GloVe (6B tokens)
- FastText (157 languages)
- Domain-specific models

Visualization

- TensorBoard Projector
- Embedding Explorer
- Custom t-SNE tools

Continue Learning: <https://github.com/your-course/week02-resources>

These resources will deepen your understanding and practical skills

What's Next in Our Journey:

Next Week: RNNs & LSTMs

- Sequential processing
- Hidden state dynamics
- Vanishing gradient problem
- LSTM architecture
- Text generation

Building On Word2Vec:

- Embeddings as RNN input
- Sequence modeling
- Context evolution
- Memory mechanisms

Future Topics:

- Week 4: Seq2Seq & Attention
- Week 5: Transformers
- Week 6: BERT & GPT
- Week 7: Advanced architectures
- Week 8-12: Modern NLP

The Big Picture:

- Word2Vec → RNN → Transformer
- Static → Contextual embeddings
- Understanding → Generation

You now understand the foundation of all modern NLP!