

Neural Language Models

Week 2: Word Embeddings and Word2Vec

NLP Course 2025

Professional Template Edition

September 29, 2025

Week 2: Journey Through Word Embeddings

Learning Path: From discrete word IDs to continuous semantic vectors. Master how neural networks learn word meaning through context, leading to the Word2Vec revolution that powers modern NLP.

Part 1: Introduction & Motivation

Interactive: Word Association Game

Fill in the blank - What word naturally comes next?

1. The cat sat on the _____

Interactive: Word Association Game

Fill in the blank - What word naturally comes next?

1. The cat sat on the _____ → mat, floor, chair (physical objects)

2. I drink my coffee with milk and _____

Interactive: Word Association Game

Fill in the blank - What word naturally comes next?

1. The cat sat on the _____ → mat, floor, chair (physical objects)

2. I drink my coffee with milk and _____ → sugar, cream, honey (additives)

3. The capital of France is _____

Interactive: Word Association Game

Fill in the blank - What word naturally comes next?

1. The cat sat on the _____ → mat, floor, chair (physical objects)
2. I drink my coffee with milk and _____ → sugar, cream, honey (additives)
3. The capital of France is _____ → Paris (factual knowledge)
4. She was happy but also felt _____


Interactive: Word Association Game

Fill in the blank - What word naturally comes next?

1. The cat sat on the _____ → mat, floor, chair (physical objects)
2. I drink my coffee with milk and _____ → sugar, cream, honey (additives)
3. The capital of France is _____ → Paris (factual knowledge)
4. She was happy but also felt _____ → sad, anxious, confused (emotions)

Humans predict words using semantic understanding - how can computers learn this?

The Evolution of Language Modeling



Four Major Eras in Next-Word Prediction:

Traditional Approach: One-Hot Encoding

- Words as discrete IDs
- Vocabulary size: 10,000 words
- “cat” = [0,0,1,0,...,0] (position 3)
- “dog” = [0,0,0,0,1,...,0] (position 5)

Problems:

- No notion of similarity
- $\text{distance}(\text{cat}, \text{dog}) = \text{distance}(\text{cat}, \text{democracy})$
- Can't generalize knowledge
- Huge, sparse vectors

Solution: Dense Embeddings

- Words as dense vectors
- Dimension: 100-300 (not 10,000!)
- “cat” = [0.2, -0.4, 0.7, ...]
- “dog” = [0.3, -0.3, 0.8, ...]

Benefits:

- Similar words have similar vectors
- $\text{distance}(\text{cat}, \text{dog}) \neq \text{distance}(\text{cat}, \text{democracy})$
- Knowledge transfers between similar words
- Compact, meaningful representation

Key Insight: Learn representations where geometric distance = semantic distance

Search Engines

- Semantic search
- Query understanding
- “car” finds “automobile”
- Intent matching

Used by:

- Google Search
- Bing
- DuckDuckGo

Recommendations

- Content similarity
- User preferences
- Cross-lingual matching
- Cold-start solutions

Used by:

- Netflix
- Spotify
- Amazon

Language AI

- Machine translation
- Sentiment analysis
- Chatbots
- Foundation for LLMs

Used by:

- ChatGPT
- Google Translate
- Grammarly

Market Impact: Word embeddings power \$100B+ in NLP applications worldwide

Word2Vec papers cited 50,000+ times - one of the most influential ML innovations

Part 2: Core Concepts

“You shall know a word by the company it keeps”

- J.R. Firth (1957)

Example Context Windows:

- The **cat** sat on the mat
- The **dog** sat on the floor
- A **cat** chased the mouse
- A **dog** chased the ball

Shared contexts:

- Both appear after “The” and “A”
- Both appear before “sat”, “chased”
- Both are subjects of similar actions

This simple idea - words with similar contexts have similar meanings - drives all embeddings

Mathematical Formulation:

Context window of size 2:

$$\text{context(cat)} = \{\text{The, sat, on, the}\}$$

$$\text{context(dog)} = \{\text{The, sat, on, the}\}$$


Key Insight:

- Similar contexts \Rightarrow Similar meanings
- Learn vectors to predict context
- Vectors capture semantic similarity


Word2Vec: Two Revolutionary Architectures

Word2Vec Architecture Comparison

CBOW: Context → Center



Skip-gram: Center → Context



CBOW (Continuous Bag-of-Words)

- Predict center word from context

Skip-gram

- Predict context from center word

Skip-gram Objective Function:

Maximize the probability of context words given center word:

$$J(\theta) = \frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log P(w_{t+j} | w_t)$$

Probability Calculation using Softmax:

$$P(w_O | w_I) = \frac{\exp(v_{w_O}^T \cdot v_{w_I})}{\sum_{w=1}^V \exp(v_w^T \cdot v_{w_I})}$$

Where:

- v_{w_I} : Input vector for center word
- v_{w_O} : Output vector for context word
- V : Vocabulary size
- c : Context window size


Problem: Denominator sums over entire vocabulary (expensive!)

Computing softmax over 50,000 words for every training example is computationally prohibitive


Negative Sampling: Making Training Feasible

Negative Sampling: The Optimization That Made Word2Vec Practical

Full Softmax: Compute All 50,000 Words



Negative Sampling: Only 5-20 Words



Convert to Binary Classification:
New Objective:





Instead of softmax over all words:

- Positive sample: $(\text{cat}, \text{sat}) \rightarrow 1$

Part 3: Training & Solutions

Training Dynamics: How Embeddings Evolve

Word2Vec Training Dynamics



Semantic Arithmetic: Mathematical Operations on Meaning

Gender Relationship

$$\begin{matrix} \text{King} \\ \boxed{\text{King}} \end{matrix} - \begin{matrix} \text{Man} \\ \boxed{\text{Man}} \end{matrix} + \begin{matrix} \text{Woman} \\ \boxed{\text{Woman}} \end{matrix} = \begin{matrix} \text{Queen} \\ \boxed{\text{Queen}} \end{matrix}$$

Capital Cities

$$\begin{matrix} \text{Paris} \\ \boxed{\text{Paris}} \end{matrix} - \begin{matrix} \text{France} \\ \boxed{\text{France}} \end{matrix} + \begin{matrix} \text{Italy} \\ \boxed{\text{Italy}} \end{matrix} = \begin{matrix} \text{Rome} \\ \boxed{\text{Rome}} \end{matrix}$$

Verb Conjugation

$$\begin{matrix} \text{Walking} \\ \boxed{\text{Walking}} \end{matrix} - \begin{matrix} \text{Walk} \\ \boxed{\text{Walk}} \end{matrix} + \begin{matrix} \text{Swim} \\ \boxed{\text{Swim}} \end{matrix} = \begin{matrix} \text{Swimming} \\ \boxed{\text{Swimming}} \end{matrix}$$

Comparative Forms

$$\begin{matrix} \text{Bigger} \\ \boxed{\text{Bigger}} \end{matrix} - \begin{matrix} \text{Big} \\ \boxed{\text{Big}} \end{matrix} + \begin{matrix} \text{Small} \\ \boxed{\text{Small}} \end{matrix} = \begin{matrix} \text{Smaller} \\ \boxed{\text{Smaller}} \end{matrix}$$

Evaluating Word Embeddings

Intrinsic Evaluation

- Word similarity tasks
- Analogy completion
- Clustering quality

Benchmarks:

- WordSim-353
- Google Analogy Test
- SimLex-999

Metrics:

- Spearman correlation
- Accuracy@1, @5
- Silhouette score

Extrinsic Evaluation

- Downstream task performance
- NER improvement
- Sentiment accuracy

Tasks:

- Text classification
- Machine translation
- Question answering

Metrics:

- F1 score improvement
- BLEU score gain
- Task-specific metrics

Visualization

- t-SNE projections
- PCA analysis
- Nearest neighbors

Qualitative:

- Semantic coherence
- Cluster separation
- Outlier detection

Tools:

- TensorBoard
- Embedding Projector
- Custom visualizations

Best Practice: Combine all three - numbers alone don't tell the whole story

Good embeddings show 0.6+ correlation on similarity tasks and 3-5% improvement on downstream tasks

Fundamental Limitations:

- Out-of-vocabulary words
→ FastText with subword units
- Single vector per word
→ Contextual embeddings (ELMo, BERT)
- No word order information
→ Position encodings
- Bias in training data
→ Debiasing techniques
- Fixed after training
→ Fine-tunable embeddings


Advanced Techniques:

- GloVe (2014):
Combines global statistics + local context
- FastText (2016):
Character n-grams for OOV handling
- ELMo (2018):
Context-dependent embeddings
- BERT (2018):
Bidirectional contextual representations
- GPT (2018+):
Autoregressive language modeling







Word2Vec's limitations led directly to the transformer revolution in NLP

Part 4: Applications & Future

Word Embeddings in 3D Space



Word Embeddings: Real-World Impact Across Industries




Hands-On: Using Word2Vec in Practice

```
1 from gensim.models import Word2Vec
2 import numpy as np
3
4 # Train Word2Vec model
5 sentences = [["the", "cat", "sat", "on", "the", "mat"],
6               ["the", "dog", "sat", "on", "the", "floor"]]
7
8 model = Word2Vec(sentences, vector_size=100, window=5,
9                   min_count=1, sg=1) # sg=1 for skip-gram
10
11 # Get word vectors
12 cat_vector = model.wv['cat']
13 dog_vector = model.wv['dog']
14
15 # Compute similarity
16 similarity = model.wv.similarity('cat', 'dog')
17 print(f"Similarity(cat, dog)={similarity:.3f}")
18
19 # Find similar words
20 similar_words = model.wv.most_similar('cat', topn=3)
21 print(f"Words similar to 'cat': {similar_words}")
22
23 # Word arithmetic
24 result = model.wv.most_similar(positive=['king', 'woman'],
25                                negative=['man'], topn=1)
26 print(f"king - man + woman = {result[0][0]}")
```

Gensim makes Word2Vec incredibly easy to use - training takes just minutes

From Word2Vec to Modern Transformers



Word2Vec's Legacy:

- Proved semantic learning possible

NLP Course 2025 (Professional Template Edition)

Modern Evolution:

- BERT: Contextual embeddings

Neural Language Models

September 29, 2025

22 / 1

What We Learned:

Key Concepts

- Distributional hypothesis
- Dense vector representations
- Skip-gram vs CBOW
- Negative sampling optimization
- Semantic arithmetic

Technical Skills

- Training Word2Vec models
- Evaluating embedding quality
- Visualizing semantic spaces
- Applying to downstream tasks

Practical Impact

- Powers modern search engines
- Enables machine translation
- Foundation for ChatGPT/Claude
- \$100B+ market impact

Historical Significance

- 50,000+ citations
- Revolutionized NLP (2013)
- Led to transformer era
- Still widely used today

Core Insight: Words are not just symbols - they carry meaning in their

Quick Quiz: Test Your Understanding

Answer these questions to check your understanding:

1. What is the key insight of the distributional hypothesis?

Quick Quiz: Test Your Understanding

Answer these questions to check your understanding:

1. What is the key insight of the distributional hypothesis? → Words with similar contexts have similar meanings
2. Why is negative sampling needed in Word2Vec?

Quick Quiz: Test Your Understanding

Answer these questions to check your understanding:

1. What is the key insight of the distributional hypothesis? → Words with similar contexts have similar meanings
2. Why is negative sampling needed in Word2Vec? → To avoid expensive softmax over entire vocabulary
3. What's the difference between Skip-gram and CBOW?

Quick Quiz: Test Your Understanding

Answer these questions to check your understanding:

1. What is the key insight of the distributional hypothesis? → Words with similar contexts have similar meanings
2. Why is negative sampling needed in Word2Vec? → To avoid expensive softmax over entire vocabulary
3. What's the difference between Skip-gram and CBOW? → Skip-gram: word→context, CBOW: context→word
4. Why does “king - man + woman = queen” work?

Quick Quiz: Test Your Understanding

Answer these questions to check your understanding:

1. What is the key insight of the distributional hypothesis? → Words with similar contexts have similar meanings
2. Why is negative sampling needed in Word2Vec? → To avoid expensive softmax over entire vocabulary
3. What's the difference between Skip-gram and CBOW? → Skip-gram: word→context, CBOW: context→word
4. Why does “king - man + woman = queen” work? → Relationships are encoded as vector directions
5. What's the main limitation of Word2Vec?

Quick Quiz: Test Your Understanding

Answer these questions to check your understanding:

1. What is the key insight of the distributional hypothesis? → Words with similar contexts have similar meanings
2. Why is negative sampling needed in Word2Vec? → To avoid expensive softmax over entire vocabulary
3. What's the difference between Skip-gram and CBOW? → Skip-gram: word→context, CBOW: context→word
4. Why does “king - man + woman = queen” work? → Relationships are encoded as vector directions
5. What's the main limitation of Word2Vec? → Single vector per word (no context dependence)

If you can answer these, you understand the core of word embeddings!

Essential Papers

- Mikolov et al. (2013a): Efficient Estimation
- Mikolov et al. (2013b): Distributed Representations
- Goldberg & Levy (2014): word2vec Explained
- Pennington et al. (2014): GloVe

Implementations

- Gensim (Python)
- TensorFlow Embeddings
- PyTorch nn.Embedding
- FastText library

Datasets & Tools

- Google News vectors
- GloVe pre-trained
- Embedding Projector
- Word2Vec demos

Lab Session Preview:

- Train Word2Vec on real corpus
- Explore semantic relationships
- Build a similarity search engine
- Visualize your own embeddings

Lab notebook: week02_word_embeddings.lab.ipynb