

Natural Language Processing Course

Week 2: Neural Language Models and Word Embeddings

Restructured with 4-Part Format

2024

Week 2: Neural Language Models - Overview

Part 1: Introduction & Motivation

- Interactive word association
- The semantic understanding problem
- Real-world impact and applications
- Historical journey to Word2Vec

Part 2: Core Concepts

- Distributional hypothesis
- From discrete to continuous
- Word2Vec architecture
- Implementation deep dive

Part 3: Challenges & Solutions

- Training at scale
- Evaluation methodologies
- Fundamental limitations
- Advanced techniques

Part 4: Applications & Future

- Hands-on applications
- Modern evolution (BERT, GPT)
- Industry state-of-the-art
- Looking forward

Goal: Master how computers learn word meaning through context

Part 1

Introduction and Motivation

Why Computers Need to Understand Word Meaning

Interactive Exercise: Word Association Game

When you see this word, what comes to mind?

OCEAN

Interactive Exercise: Word Association Game

When you see this word, what comes to mind?

OCEAN

water

35% of you

sea

25% of you

beach

20% of you

waves

20% of you

You naturally understand semantic relationships!

But until 2003, computers saw:

- ocean = ID 7849
- water = ID 2341
- No connection whatsoever!

The Semantic Gap: Computers vs Humans

How Humans See Words:

- cat ≈ kitten (similar animals)
- Paris ↔ France (location relation)
- running ~ ran (same verb, different tense)
- doctor ↔ hospital (association)

Rich semantic network with relationships, similarities, and associations

How Computers Saw Words (Pre-2003):

- cat = 1247
- kitten = 8923
- Paris = 4567
- France = 2109

Arbitrary IDs with no notion of meaning or relationships

The Challenge: Bridge this semantic gap!

Real System Failures Without Semantic Understanding

Early Google Search (2000):

- Search: "car" → Missed: "automobile", "vehicle"
- Search: "running shoes" → Missed: "jogging sneakers"

Machine Translation Disasters:

- "The spirit is willing but the flesh is weak"
- → Russian → English:
- "The vodka is good but the meat is rotten"

Customer Service Chatbots (2005):

- Customer: "I want to return my purchase"
- Bot: "I don't understand. Did you mean 'buy'?"
- → Couldn't link "return" with "refund", "exchange"

Economic Impact: Billions lost due to poor search and translation

Where Word Embeddings Power Your Life (2024)

Entertainment:

- **Spotify**: 256-dim song embeddings
- **Netflix**: Show similarity vectors
- **TikTok**: Video understanding
- **YouTube**: Related videos

Productivity:

- **Gmail**: Smart compose (BERT)
- **Grammarly**: Context awareness
- **Notion AI**: Semantic search
- **Slack**: Message threading

Commerce:

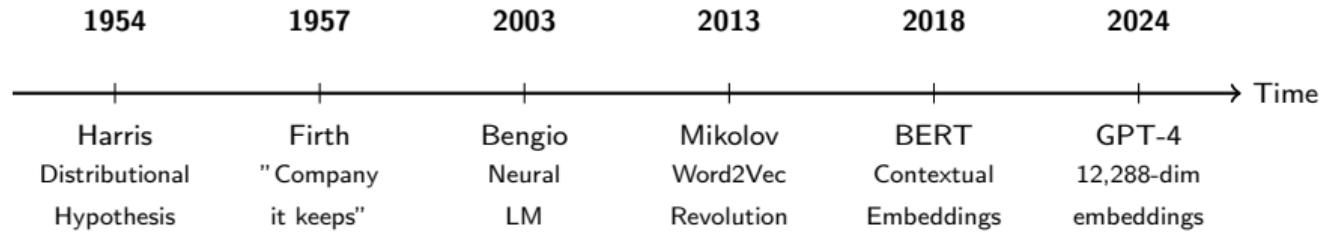
- **Amazon**: Product similarity
- **Google Ads**: Ad matching
- **Airbnb**: Listing embeddings
- **Uber**: Location understanding

Market Size:

- Embedding API Market: \$2.7B by 2025
- OpenAI Embeddings: 1M+ developers
- Vector Database Market: \$4.3B by 2028

Every AI application today relies on word embeddings!

The Journey to Understanding: Timeline



Key Breakthroughs:

- 1954-1957: Theoretical foundation - words defined by context
- 2003: First neural language model with continuous representations
- 2013: Word2Vec makes embeddings practical and scalable
- 2018: Contextualized embeddings (same word, different contexts)
- 2024: Massive embeddings powering GPT-4, Claude, Gemini

The 2013 Breakthrough: King - Man + Woman = ?

The demo that shocked the NLP world:¹

king - man + woman =

¹Mikolov et al. (2013). "Linguistic regularities in continuous space word representations"

The 2013 Breakthrough: King - Man + Woman = ?

The demo that shocked the NLP world:¹

$$\text{king} - \text{man} + \text{woman} = \text{queen}$$

Why this was revolutionary:

- Computer discovered gender relationships automatically
- No one programmed these rules
- Learned purely from reading text
- Worked across many relationship types

More examples that work:

- | | |
|---|---------------------------------------|
| • Paris - France + Italy = Rome | • bigger - big + small = smaller |
| • sushi - Japan + Mexico = tacos | • walking - walk + swim = swimming |
| • Einstein - scientist + artist = Picasso | • CEO - company + country = president |

¹Mikolov et al. (2013). "Linguistic regularities in continuous space word representations"

Part 1 Summary: Why This Matters

Key Insights:

- ① **The Problem:** Computers treating words as meaningless IDs
- ② **The Impact:** Billions in losses, poor user experiences
- ③ **The Solution:** Learn meaning from context (distributional hypothesis)
- ④ **The Breakthrough:** Word2Vec made it practical (2013)

What's Next:

- Part 2: How do we actually create these word vectors?
- Understanding the mathematics and algorithms
- Building Word2Vec from scratch

Remember: Every modern AI system (ChatGPT, Claude, Gemini) started here!

Part 2

Core Concepts

How Computers Learn Word Meaning from Context

The Distributional Hypothesis: Foundation

Core Principle (Firth, 1957):

"You shall know a word by the company it keeps"

Example: What is a "zorb"?

- The zorb ate the cheese
- I saw a zorb in my garden
- The zorb ran under the couch
- My cat chased the zorb

The Distributional Hypothesis: Foundation

Core Principle (Firth, 1957):

"You shall know a word by the company it keeps"

Example: What is a "zorb"?

- The zorb ate the cheese
- I saw a zorb in my garden
- The zorb ran under the couch
- My cat chased the zorb

You probably guessed: zorb = mouse (or similar small animal)

Mathematical Formulation:

- Words with similar distributions have similar meanings
- $\text{similarity}(w_1, w_2) \propto P(\text{context}|w_1) \cdot P(\text{context}|w_2)$
- Context defines meaning!

Interactive: Guess the Word from Context

Mystery word = [BLANK]. What is it?

- ① The [BLANK] was delicious
- ② I ordered [BLANK] with extra cheese
- ③ The [BLANK] delivery arrived in 30 minutes
- ④ We shared a large [BLANK] at the party
- ⑤ My favorite [BLANK] topping is pepperoni

Interactive: Guess the Word from Context

Mystery word = [BLANK]. What is it?

- ① The [BLANK] was delicious
- ② I ordered [BLANK] with extra cheese
- ③ The [BLANK] delivery arrived in 30 minutes
- ④ We shared a large [BLANK] at the party
- ⑤ My favorite [BLANK] topping is pepperoni

Answer: pizza

This is exactly how Word2Vec learns:

- Sees millions of sentences
- Learns what words appear in similar contexts
- Groups them close together in vector space
- No dictionary needed!

From Discrete IDs to Continuous Vectors

One-Hot Encoding (Old Way):

- Vocabulary size: 50,000 words
- $\text{cat} = [0,0,1,0,0,\dots,0]$ (50K dimensions!)
- $\text{dog} = [0,0,0,1,0,\dots,0]$

Problems:

- **No similarity**: $\text{cat} \cdot \text{dog} = 0$
- **Huge vectors**: 50K dimensions
- **Sparse**: 49,999 zeros
- **No learning**: Fixed representation

Dense Embeddings (Word2Vec):

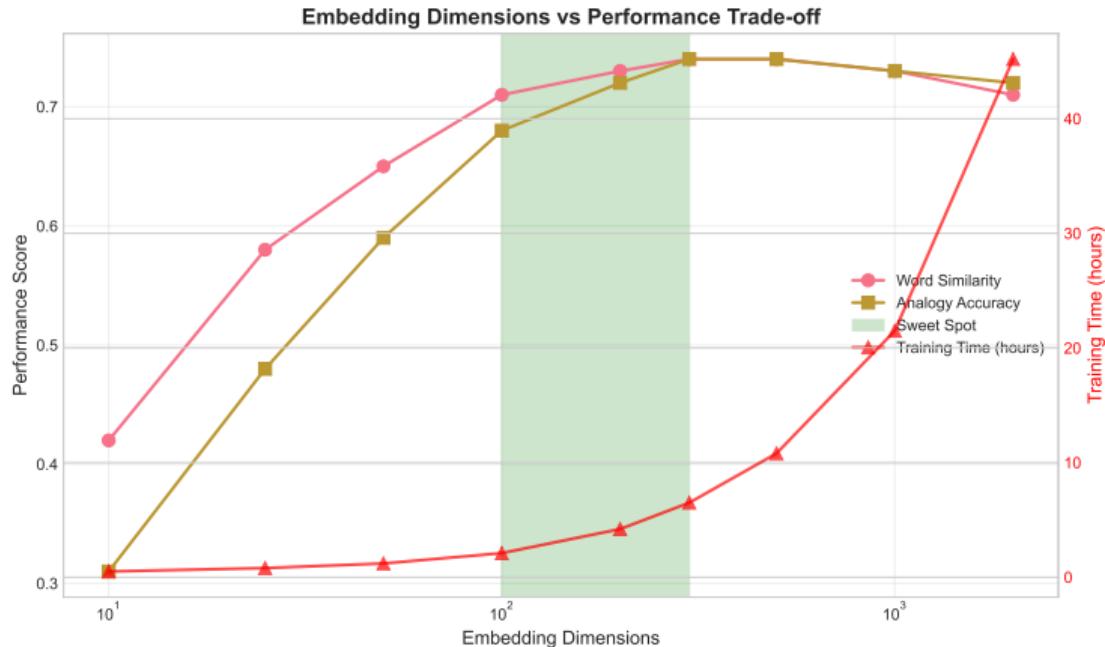
- Typical size: 100-300 dimensions
- $\text{cat} = [0.2, -0.4, 0.7, \dots, 0.1]$
- $\text{dog} = [0.3, -0.3, 0.6, \dots, 0.2]$

Benefits:

- **Similarity**: $\text{cat} \cdot \text{dog} = 0.89$
- **Compact**: 100-300 dims
- **Dense**: All values meaningful
- **Learnable**: Updated during training

Key: Every dimension captures some semantic property

The Goldilocks Zone: Why 100-300 Dimensions?



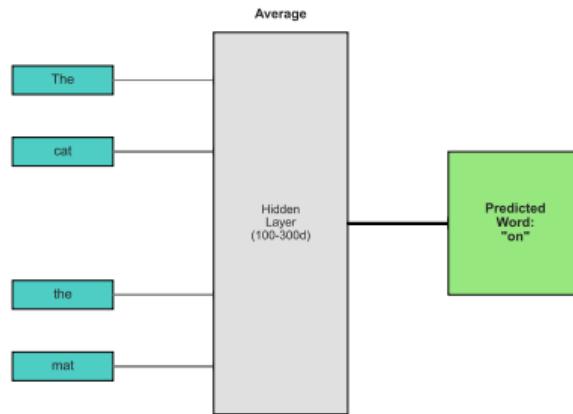
Empirical Findings:

- ↳ 50 dims: Too compressed, loses nuances
- 100-300 dims: Sweet spot for most tasks
- ↳ 500 dims: Diminishing returns, overfitting risk

Word2Vec: Two Architectures

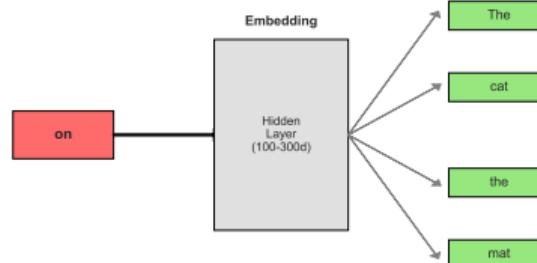
CBOW (Continuous Bag of Words):

CBOW Architecture: Predict Center from Context



Skip-gram:

Skip-gram Architecture: Predict Context from Center



- Predict center from context
- Input: [the, cat, on, mat]
- Output: sat
- Faster to train
- Better for frequent words

- Predict context from center
- Input: sat
- Output: [the, cat, on, mat]
- Slower but more accurate
- Better for rare words

Skip-gram Training: Step by Step

Sentence: "The quick brown fox jumps"

Window size = 2 (look 2 words left and right)

Step	Center Word	Context to Predict
1	quick	[the, brown]
2	brown	[the, quick, fox, jumps]
3	fox	[quick, brown, jumps]

Training Process:

- ① Take center word embedding
- ② Try to predict context words
- ③ Measure prediction error
- ④ Update embeddings to reduce error
- ⑤ Repeat millions of times

Result: Words appearing in similar contexts get similar embeddings

Implementing Word2Vec in PyTorch

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4
5 class Word2Vec(nn.Module):
6     def __init__(self, vocab_size, embed_dim=100):
7         super().__init__()
8         # Two embedding matrices
9         self.center_embeddings = nn.Embedding(
10             vocab_size, embed_dim
11         )
12         self.context_embeddings = nn.Embedding(
13             vocab_size, embed_dim
14         )
15
16     def forward(self, center, context, neg_samples):
17         # Get embeddings
18         center_emb = self.center_embeddings(center)
19         context_emb = self.context_embeddings(context)
20         neg_emb = self.context_embeddings(neg_samples)
21
22         # Positive samples (should be similar)
23         pos_score = torch.sum(
24             center_emb * context_emb, dim=1
25         )
26         pos_loss = F.logsigmoid(pos_score)
27
28         # Negative samples (should be different)
29         neg_score = torch.bmm(
30             neg_emb, center_emb.unsqueeze(2)
31         ).squeeze()
32         neg_loss = F.logsigmoid(-neg_score).sum(1)
33
34         return -(pos_loss + neg_loss).mean()
```

Key Components:

- **Two matrices:** Center and context embeddings
- **Positive samples:** Real context words
- **Negative samples:** Random words (not in context)

Training Trick:

- Full softmax over 50K words is expensive
- Solution: Negative sampling
- Only update a few random words
- 5-20 negative samples typical

Loss Function:

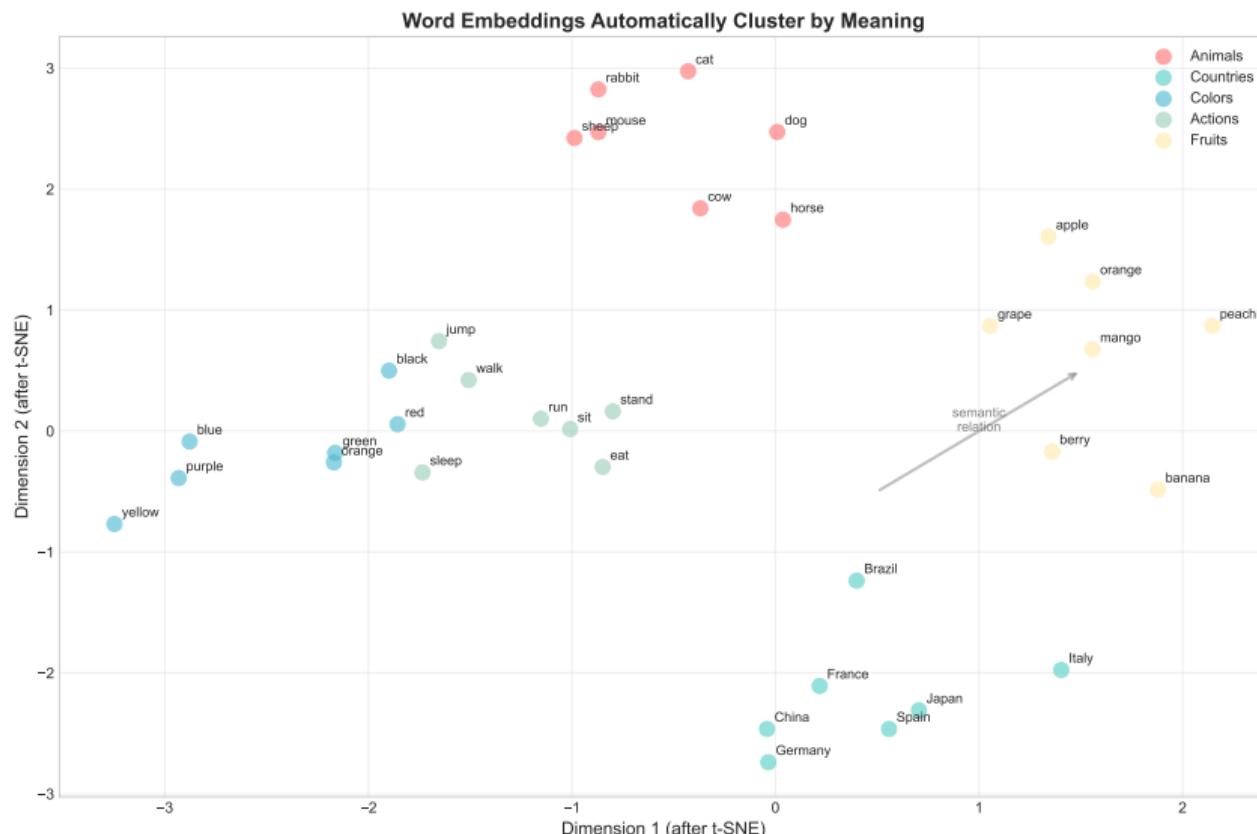
- Maximize similarity with real context
- Minimize similarity with random words

Training Word2Vec: The Complete Loop

```
1 def train_word2vec(corpus, vocab_size, embed_dim=100, epochs=5, window=2):
2     model = Word2Vec(vocab_size, embed_dim)
3     optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
4
5     for epoch in range(epochs):
6         total_loss = 0
7         for sentence in corpus:
8             # Generate training samples from sentence
9             for i, center_word in enumerate(sentence):
10                 # Get context words within window
11                 context_words = []
12                 for j in range(max(0, i-window), min(len(sentence), i+window+1)):
13                     if i != j:
14                         context_words.append(sentence[j])
15
16                 # Get negative samples (5 random words not in context)
17                 neg_samples = get_negative_samples(vocab_size, 5, avoid=context_words)
18
19                 # Forward pass
20                 loss = model(center_word, context_words, neg_samples)
21
22                 # Backward pass
23                 optimizer.zero_grad()
24                 loss.backward()
25                 optimizer.step()
26
27                 total_loss += loss.item()
28
29             print(f"Epoch {epoch}: Loss = {total_loss:.4f}")
30
31     return model.center_embeddings.weight.data # Final embeddings
```

Result: After training on millions of sentences, similar words cluster together!

Visualizing What Word2Vec Learns



Mathematical Intuition: Why Dot Product = Similarity

The Skip-gram Objective:

$$\max \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log P(w_{t+j} | w_t)$$

Where probability is defined as:

$$P(w_O | w_I) = \frac{\exp(v_{w_O}^T v_{w_I})}{\sum_{w=1}^W \exp(v_w^T v_{w_I})}$$

Key Insight:

- Dot product $v_{w_O}^T v_{w_I}$ measures similarity
- Higher dot product \rightarrow higher probability of co-occurrence
- Training maximizes dot product for words that appear together
- Result: Similar words have high dot product (cosine similarity)

Geometry emerges from statistics: Similar contexts \rightarrow Similar vectors

Part 2 Summary: Core Concepts Mastered

What We Learned:

- ① **Distributional Hypothesis:** Context defines meaning
- ② **Dense Vectors:** 100-300 dimensions capture semantics
- ③ **Skip-gram Model:** Predict context from center word
- ④ **Training Process:** Maximize co-occurrence probability
- ⑤ **Implementation:** Two embedding matrices + negative sampling

Key Takeaways:

- Word meaning emerges from statistical patterns
- No linguistic knowledge required
- Scalable to millions of words
- Foundation for all modern NLP

Next: Part 3 - Challenges and Solutions

- How to train on billions of words efficiently?
- How to evaluate embedding quality?
- What are the limitations?

Part 3

Challenges and Solutions

Scaling, Evaluation, and Limitations

Challenge 1: Computational Complexity

The Softmax Bottleneck:

Original formulation requires normalizing over entire vocabulary:

$$P(w_O|w_I) = \frac{\exp(v_{w_O}^T v_{w_I})}{\sum_{w=1}^W \exp(v_w^T v_{w_I})}$$

Problem:

- Vocabulary size $W = 50,000+$ words
- Must compute 50,000 dot products per training step
- Billions of training steps needed
- Computationally infeasible!

Solutions:

1. Hierarchical Softmax:

- Binary tree of words
- $O(\log W)$ instead of $O(W)$
- Path through tree to each word

2. Negative Sampling:

- Only update k random words
- Typically $k = 5-20$
- Dramatic speedup
- Better performance!

Solution: Negative Sampling Explained

Instead of: Predicting the right word from 50,000 options

We ask: Is this word the right context word? (Binary classification)

Center	Word	Label
cat	sits (real context)	1
cat	on (real context)	1
cat	elephant (random)	0
cat	democracy (random)	0
cat	quantum (random)	0

Sampling Strategy:

- Sample negative words by frequency: $P(w) \propto f(w)^{3/4}$
- The 3/4 power reduces dominance of very common words
- Gives rare words more chance to be negative samples

Result: 1000x speedup with better quality embeddings!

Challenge 2: How Do We Evaluate Embeddings?

The Problem: How do we know if our embeddings are "good"?

Intrinsic Evaluation:

- Word Similarity:
 - Dataset: WordSim-353
 - Human ratings vs cosine similarity
 - Correlation: 0.6-0.7 typical
- Word Analogies:
 - king - man + woman = ?
 - Google analogy dataset
 - Accuracy: 60-75% typical

Extrinsic Evaluation:

- Downstream Tasks:
 - Sentiment analysis
 - Named entity recognition
 - Machine translation
- Key Finding:
 - Good intrinsic \neq Good extrinsic
 - Task-specific fine-tuning helps

Modern Approach (2024):

- Skip intrinsic evaluation
- Directly evaluate on downstream tasks
- Use pre-trained embeddings as starting point

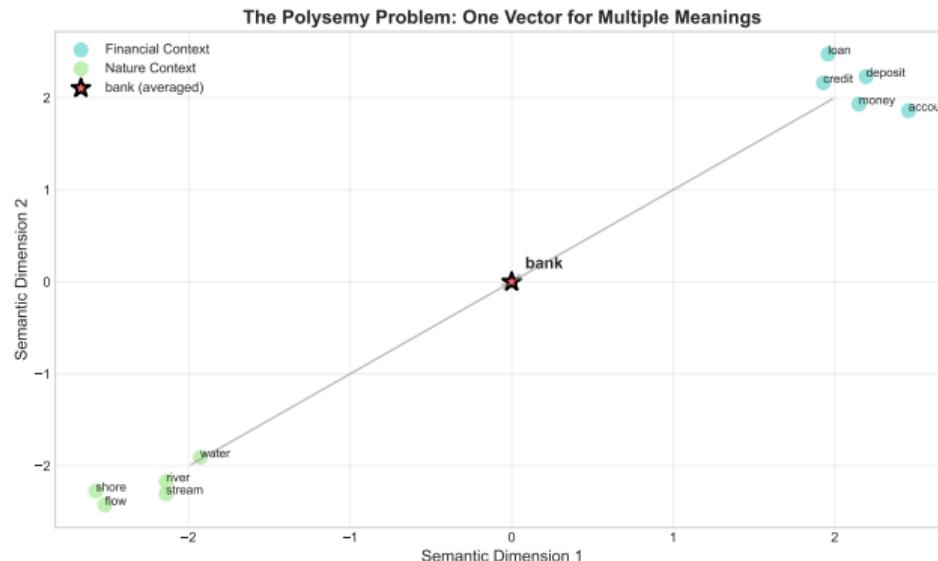
Challenge 3: The Polysemy Problem

One Vector Per Word... But Words Have Multiple Meanings!

Example: "bank"

- "I deposited money at the **bank**" (financial institution)
- "We sat by the river **bank**" (edge of river)

Word2Vec gives one vector that averages both meanings:

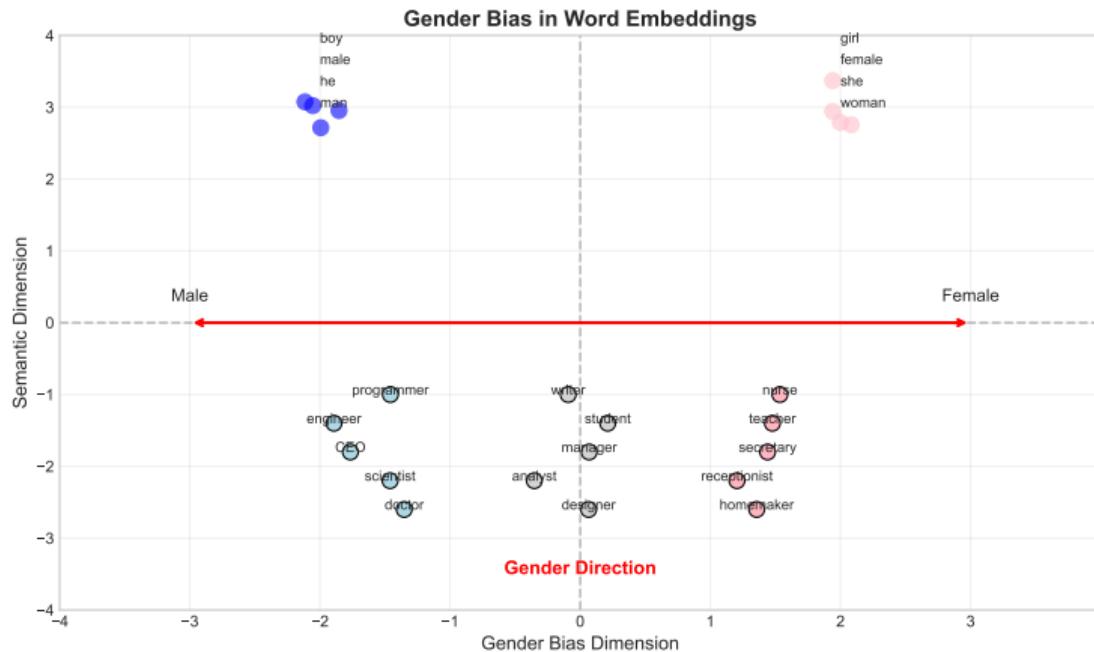


Challenge 4: Bias in Embeddings

Embeddings Learn Societal Biases from Text:

Problematic Analogies Found:

- man : computer programmer :: woman : homemaker
- man : doctor :: woman : nurse



Advanced Techniques: Beyond Basic Word2Vec

FastText (2016):

- Uses character n-grams
- "where" = "wh", "whe", "her", "ere", "re"
- Handles unseen words
- Better for morphologically rich languages

GloVe (2014):

- Global matrix factorization
- Combines count-based and predictive
- Often better for word analogies

ELMo (2018):

- Contextualized embeddings
- Different vector per context
- Bi-directional LSTM
- Solves polysemy problem

Modern (2024):

- BERT/GPT embeddings
- Learned during pre-training
- Task-specific fine-tuning
- 768-12,288 dimensions

Word2Vec pioneered the field, but modern methods build on its foundation

Part 3 Summary: Challenges Addressed

Challenges We Explored:

- ① **Computational:** Softmax over 50K words → Negative sampling
- ② **Evaluation:** Intrinsic vs extrinsic metrics
- ③ **Polysemy:** One vector per word limitation
- ④ **Bias:** Embeddings reflect societal biases

Solutions and Evolution:

- Negative sampling: 1000x speedup
- Task-specific evaluation
- Contextualized embeddings (BERT/GPT)
- Debiasing techniques

Next: Part 4 - Applications and Future

- Build real applications with embeddings
- See modern systems in action
- Understand the path forward

Part 4

Applications and Future

From Word2Vec to Modern AI Systems

Build It: Semantic Search Engine

Let's build a search engine that understands meaning!

Traditional Search:

- Query: "car"
- Finds: Only documents with "car"
- Misses: "automobile", "vehicle"

Semantic Search:

- Query: "car"
- Finds: "car", "automobile", "vehicle", "BMW"
- Understands synonyms and related concepts

Implementation Steps:

- ① Load pre-trained Word2Vec embeddings
- ② Convert documents to vectors (average word embeddings)
- ③ Convert query to vector
- ④ Find documents with highest cosine similarity
- ⑤ Return ranked results

This is the foundation of Google Search, Elastic Search, and more!

Semantic Search Implementation

```
1 import numpy as np
2 from sklearn.metrics.pairwise import cosine_similarity
3
4 class SemanticSearch:
5     def __init__(self, word2vec_model):
6         self.w2v = word2vec_model
7         self.documents = []
8         self.doc_vectors = []
9
10    def add_document(self, doc):
11        """Add document to search index"""
12        self.documents.append(doc)
13        # Convert document to vector (average of word vectors)
14        words = doc.lower().split()
15        vectors = [self.w2v[word] for word in words if word in self.w2v]
16        doc_vector = np.mean(vectors, axis=0) if vectors else np.zeros(100)
17        self.doc_vectors.append(doc_vector)
18
19    def search(self, query, top_k=5):
20        """Find most similar documents to query"""
21        # Convert query to vector
22        words = query.lower().split()
23        vectors = [self.w2v[word] for word in words if word in self.w2v]
24        query_vector = np.mean(vectors, axis=0) if vectors else np.zeros(100)
25
26        # Calculate similarities
27        similarities = cosine_similarity([query_vector], self.doc_vectors)[0]
28
29        # Return top k results
30        top_indices = np.argsort(similarities)[-1:-top_k:-1]
31        return [(self.documents[i], similarities[i]) for i in top_indices]
```

Real-World Applications (2024)

Content Recommendation:

- Netflix: Show embeddings
- Spotify: Song2Vec
- YouTube: Video embeddings
- Amazon: Product2Vec

Language Understanding:

- ChatGPT: Token embeddings
- Google Translate: Multilingual embeddings
- Grammarly: Context understanding

Search and Retrieval:

- Google: Semantic search
- Bing: Neural matching
- Enterprise search: Document similarity

Novel Applications:

- Code2Vec: Programming embeddings
- Molecule2Vec: Drug discovery
- Graph2Vec: Social networks

Any data with context can use embedding techniques!

Evolution: From Word2Vec to GPT

Word2Vec (2013)	GloVe/FastText	ELMo (2018)	BERT (2018)	GPT-4 (2024)
-----------------	----------------	-------------	-------------	--------------



Static embedding	Better coverage	Contextualized	Transformers	12,288 dims
One vector per word	Subword units	Bi-LSTM based	Bidirectional	Multimodal

Key Progression:

- Static → Contextualized
- Words → Subwords → Tokens
- 300 dims → 12,288 dims
- Single task → Multi-task → General intelligence

Future Directions: What's Next?

Current Research (2024):

- **Efficient Embeddings:** Maintain quality at 64 dimensions
- **Multimodal:** Text + Image + Audio in same space
- **Dynamic:** Embeddings that update with new information
- **Personalized:** User-specific embedding spaces

Challenges Being Solved:

- **Long Context:** Embed entire books (1M+ tokens)
- **Cross-lingual:** Universal embeddings for all languages
- **Interpretability:** Understanding what each dimension means
- **Continual Learning:** Updating without forgetting

Connection to Next Week:

- Week 3: RNNs - Processing sequences with embeddings
- Embeddings are the input to all modern NLP models
- Foundation we'll build on for rest of course

Week 2 Summary: Words Have Meaning!

Journey We Took:

- ① Started with words as meaningless IDs
- ② Learned the distributional hypothesis
- ③ Built Word2Vec from scratch
- ④ Tackled challenges (scale, bias, polysemy)
- ⑤ Applied to real problems

Key Takeaways:

- **Core Insight:** Similar contexts → Similar meanings
- **Technical:** Skip-gram + negative sampling = efficient training
- **Practical:** Embeddings power all modern AI
- **Evolution:** Static → Contextualized → Multimodal

Your Homework:

- Build semantic search engine (notebook provided)
- Explore biases in pre-trained embeddings
- Try word arithmetic with different models

Remember: Every ChatGPT response starts with embeddings!

Appendix A: Skip-gram Objective Derivation

Full Mathematical Formulation:

Given a sequence of words w_1, w_2, \dots, w_T , maximize:

$$\mathcal{L} = \frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log P(w_{t+j} | w_t)$$

Where:

$$P(w_O | w_I) = \frac{\exp(v_{w_O}' v_{w_I})}{\sum_{w=1}^W \exp(v_w' v_{w_I})}$$

Gradient with respect to center word:

$$\frac{\partial \log P(w_O | w_I)}{\partial v_{w_I}} = v_{w_O}' - \sum_{w=1}^W P(w | w_I) \cdot v_w'$$

Interpretation:

- First term: Move toward actual context word
- Second term: Move away from expected context (all words weighted by probability)
- Result: Embeddings organize by co-occurrence patterns

Appendix A: Negative Sampling Mathematics

Original objective (expensive):

$$\log P(w_O|w_I) = \log \frac{\exp(v_{w_O}'^T v_{w_I})}{\sum_{w=1}^W \exp(v_w'^T v_{w_I})}$$

Negative sampling objective (efficient):

$$\log \sigma(v_{w_O}'^T v_{w_I}) + \sum_{i=1}^k \mathbb{E}_{w_i \sim P_n(w)} [\log \sigma(-v_{w_i}'^T v_{w_I})]$$

Where:

- $\sigma(x) = \frac{1}{1+e^{-x}}$ (sigmoid function)
- k = number of negative samples (typically 5-20)
- $P_n(w) \propto f(w)^{3/4}$ (noise distribution)

Why it works:

- Transforms problem to binary classification
- Distinguishing real from noise is sufficient
- Dramatically reduces computation: $O(k)$ instead of $O(W)$

Concrete Example: Computing $P(\text{cat} \mid \text{context})$

Given: "The cat sat on the mat" with window=1

Simplified vocabulary: {the, cat, sat, on, mat} = {0, 1, 2, 3, 4}

Embeddings (2D for visualization):

Word	v (input)	v' (output)
the	[0.5, 0.3]	[0.4, 0.6]
cat	[0.8, 0.2]	[0.7, 0.4]
sat	[0.3, 0.7]	[0.5, 0.8]
on	[0.4, 0.5]	[0.3, 0.6]
mat	[0.6, 0.4]	[0.8, 0.3]

Calculate $P(\text{sat} \mid \text{cat})$:

$$v_{\text{cat}}^T \cdot v'_{\text{sat}} = [0.8, 0.2] \cdot [0.5, 0.8] = 0.4 + 0.16 = 0.56 \quad (1)$$

$$\exp(0.56) = 1.75 \quad (2)$$

$$Z = \sum_w \exp(v_{\text{cat}}^T \cdot v'_w) = 1.73 + 1.75 + 1.61 + 1.87 + 1.94 = 8.9 \quad (3)$$

$$P(\text{sat} \mid \text{cat}) = \frac{1.75}{8.9} = 0.197 \approx 19.7\% \quad (4)$$

Concrete Example: Gradient Update

One gradient step for center word "cat":

Initial: $v_{cat} = [0.8, 0.2]$, learning rate $\alpha = 0.1$

Gradient components:

- Positive: Pull toward "sat" (actual context)
- Negative: Push away from expected distribution

Calculation:

$$\nabla = v'_{sat} - \sum_w P(w|cat) \cdot v'_w \quad (5)$$

$$= [0.5, 0.8] - (0.194 \cdot [0.4, 0.6] + 0.197 \cdot [0.5, 0.8] + \dots) \quad (6)$$

$$= [0.5, 0.8] - [0.48, 0.62] \quad (7)$$

$$= [0.02, 0.18] \quad (8)$$

Update:

$$v_{cat}^{new} = v_{cat} + \alpha \cdot \nabla = [0.8, 0.2] + 0.1 \cdot [0.02, 0.18] = [0.802, 0.218]$$

Result: "cat" moves slightly closer to "sat" in embedding space!

Concrete Example: Negative Sampling Calculation

Training instance: Center="cat", Positive="sat", Negatives={the, mat}

Loss calculation step-by-step:

1. Positive pair (cat, sat):

$$\text{score}_{pos} = v_{cat}^T \cdot v'_{sat} = 0.56 \quad (9)$$

$$\text{loss}_{pos} = -\log \sigma(0.56) = -\log(0.636) = 0.452 \quad (10)$$

2. Negative pairs:

$$\text{score}_{neg1} = v_{cat}^T \cdot v'_{the} = 0.52 \quad (11)$$

$$\text{loss}_{neg1} = -\log \sigma(-0.52) = -\log(0.373) = 0.987 \quad (12)$$

$$\text{score}_{neg2} = v_{cat}^T \cdot v'_{mat} = 0.70 \quad (13)$$

$$\text{loss}_{neg2} = -\log \sigma(-0.70) = -\log(0.332) = 1.103 \quad (14)$$

3. Total loss:

$$\text{Loss} = \text{loss}_{pos} + \text{loss}_{neg1} + \text{loss}_{neg2} = 0.452 + 0.987 + 1.103 = 2.542$$

Much faster: Only 3 calculations instead of 5,000!

Concrete Example: Cosine Similarity Step-by-Step

Question: How similar are "king" and "queen"?

Given embeddings (3D for simplicity):

- $v_{king} = [0.6, 0.8, 0.2]$
- $v_{queen} = [0.5, 0.7, 0.4]$

Step 1: Calculate dot product

$$v_{king} \cdot v_{queen} = (0.6 \times 0.5) + (0.8 \times 0.7) + (0.2 \times 0.4) = 0.3 + 0.56 + 0.08 = 0.94$$

Step 2: Calculate magnitudes

$$\|v_{king}\| = \sqrt{0.6^2 + 0.8^2 + 0.2^2} = \sqrt{0.36 + 0.64 + 0.04} = 1.02 \quad (15)$$

$$\|v_{queen}\| = \sqrt{0.5^2 + 0.7^2 + 0.4^2} = \sqrt{0.25 + 0.49 + 0.16} = 0.95 \quad (16)$$

Step 3: Compute cosine similarity

$$\cos(\theta) = \frac{v_{king} \cdot v_{queen}}{\|v_{king}\| \times \|v_{queen}\|} = \frac{0.94}{1.02 \times 0.95} = \frac{0.94}{0.969} = 0.97$$

Interpretation: 0.97 = Very similar! (1.0 = identical, 0 = orthogonal)

Appendix B: Hands-On Notebook Exercises

Available in: `week02_word2vec_notebook.ipynb`

Exercise 1: Train Word2Vec from Scratch

- Load text corpus (Wikipedia sample)
- Implement skip-gram with negative sampling
- Visualize training progress
- Save trained embeddings

Exercise 2: Word Arithmetic Playground

- Load pre-trained embeddings
- Implement analogy solver
- Test on Google analogy dataset
- Create your own analogies

Exercise 3: Bias Detection

- Measure gender bias in embeddings
- Visualize bias directions
- Implement simple debiasing
- Compare before/after

Exercise 4: Build Applications

- Semantic search engine
- Document clustering
- Simple chatbot with semantic understanding

Appendix C: Domain-Specific Embeddings

Medical Embeddings (BioWordVec):

- Trained on PubMed + MIMIC-III
- Captures: drug-disease relationships
- Application: Clinical decision support

Legal Embeddings (Law2Vec):

- Trained on case law + statutes
- Captures: Legal concept similarity
- Application: Legal document search

Financial Embeddings (FinBERT):

- Trained on financial news + reports
- Captures: Market sentiment
- Application: Trading signals

Code Embeddings (CodeBERT):

- Trained on GitHub repositories
- Captures: Programming patterns
- Application: Code search, bug detection

Lesson: Domain-specific training dramatically improves performance

References and Resources

Essential Papers:

- Mikolov et al. (2013). "Efficient estimation of word representations in vector space"
- Mikolov et al. (2013). "Distributed representations of words and phrases"
- Pennington et al. (2014). "GloVe: Global vectors for word representation"
- Peters et al. (2018). "Deep contextualized word representations" (ELMo)

Implementations:

- Gensim: <https://radimrehurek.com/gensim/>
- FastText: <https://fasttext.cc/>
- Hugging Face: <https://huggingface.co/>

Datasets:

- Google Analogy Test Set
- WordSim-353
- SimLex-999

Next Week: Recurrent Neural Networks - Processing sequences with embeddings