

# Recurrent Neural Networks

## Week 3 - Teaching Networks to Remember

NLP Course 2025

September 22, 2025

From Feedforward to Recurrent: Adding Memory to Neural Networks

## Processing Sequences with State

### The Challenge

- Sequential data everywhere
- Order matters
- Variable length inputs
- Long-term dependencies

### The Solution

- Recurrent connections
- Hidden state memory
- Parameter sharing
- Backprop through time

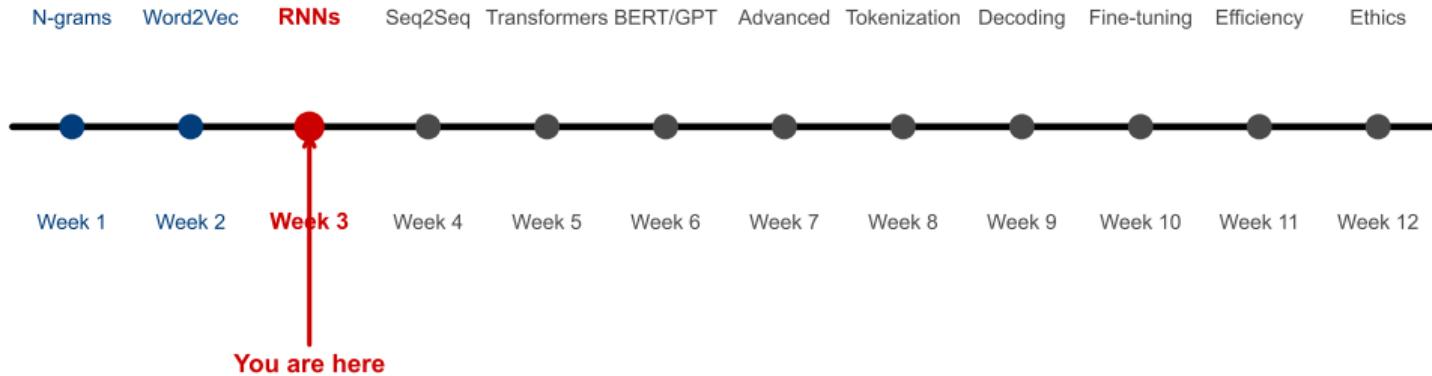
### The Evolution

- Vanilla RNN → LSTM
- Solving gradients
- Gating mechanisms
- Modern variants (GRU)

The foundation of sequence modeling before transformers

# Course Journey: Where We Are

## NLP Course Journey



### Journey So Far:

- Week 1: Statistical language models (n-grams)
- Week 2: Word embeddings (Word2Vec, dense vectors)
- **Week 3: Sequential processing with memory**

### Coming Next:

- Week 4: Encoder-decoder architectures
- Week 5: Attention mechanisms
- Week 6+: Transformers and beyond

# The Importance of Order

### Word Order Changes Meaning

- "Dog bites man" ≠ "Man bites dog"
- "Not bad" ≠ "Bad, not!"
- Context flows through sequence

### Feedforward Limitations

- Fixed input size
- No memory between inputs
- Can't model sequences naturally
- Position information lost

Sequential processing is fundamental to understanding language

### Sequential Tasks in NLP

Task	Type
Language Model	Many-to-many
Translation	Seq-to-seq
Sentiment	Many-to-one
Named Entity	Many-to-many
Speech Rec.	Seq-to-seq

Most NLP is inherently sequential

# The Core Idea: Recurrence

## Mathematical Definition

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t + b_h)$$

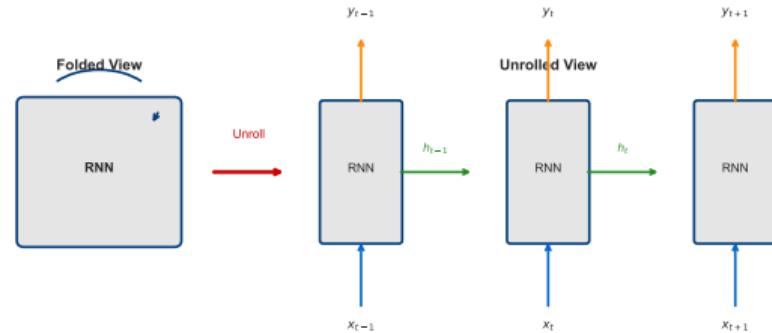
$$y_t = W_{hy}h_t + b_y$$

Where:

- $h_t$  = hidden state at time  $t$
- $x_t$  = input at time  $t$
- $y_t$  = output at time  $t$
- $W_*$  = weight matrices (shared!)

**Key Insight:** Same weights at every timestep

One network, applied repeatedly through time

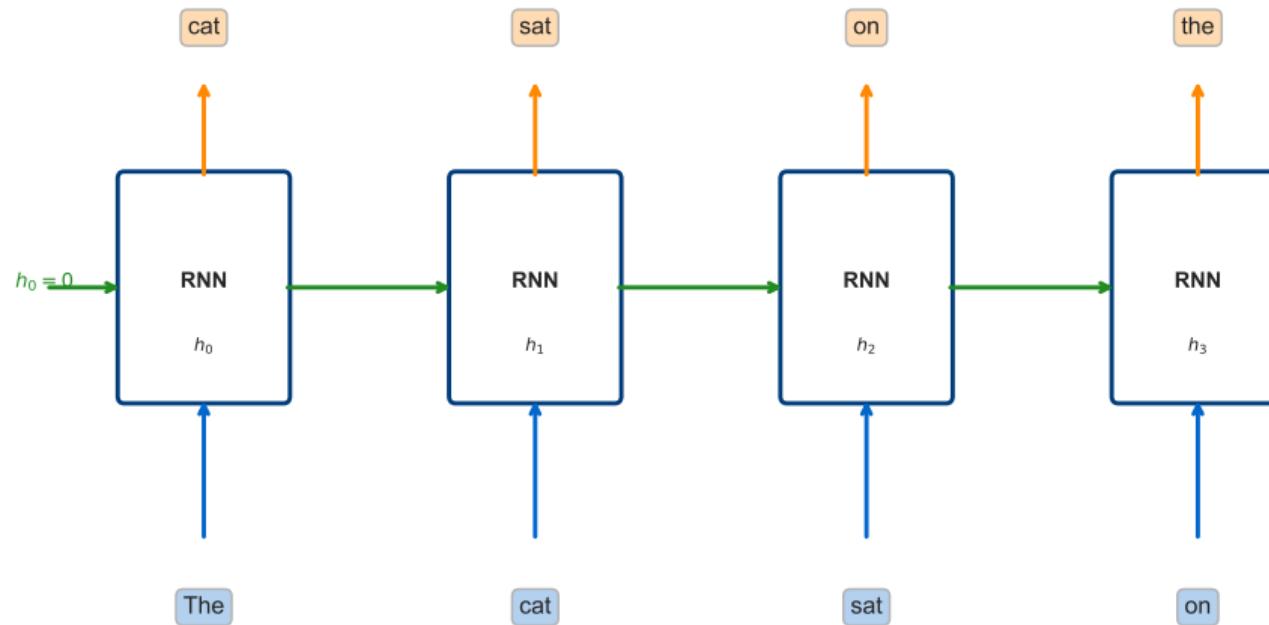


## Unrolled View Shows:

- Information flows left-to-right
- Hidden state carries memory
- Parameters shared across time
- Can handle any sequence length

# Forward Pass: Step by Step

RNN Forward Pass: Processing "The cat sat on"



# Implementation: Simple RNN Cell

```
1 import numpy as np
2
3 class RNNCell:
4     def __init__(self, input_size, hidden_size):
5         # Initialize weights
6         self.Wxh = np.random.randn(input_size,
7             hidden_size) * 0.01
8         self.Whh = np.random.randn(hidden_size,
9             hidden_size) * 0.01
10        self.Why = np.random.randn(hidden_size,
11            output_size) * 0.01
12        self.bh = np.zeros((1, hidden_size))
13        self.by = np.zeros((1, output_size))
14
15    def step(self, x, h_prev):
16        # Single timestep forward
17        h = np.tanh(np.dot(x, self.Wxh) +
18                    np.dot(h_prev, self.Whh) + self.bh)
19        y = np.dot(h, self.Why) + self.by
20        return y, h
21
22    def forward(self, inputs):
23        h = np.zeros((1, self.hidden_size))
24        outputs = []
25
26        for x in inputs:
27            y, h = self.step(x, h)
```

## PyTorch Equivalent:

```
1 import torch.nn as nn
2
3 # Built-in RNN
4 rnn = nn.RNN(
5     input_size=100,
6     hidden_size=256,
7     num_layers=1,
8     batch_first=True
9 )
10
11 # Or use LSTM/GRU
12 lstm = nn.LSTM(
13     input_size=100,
14     hidden_size=256,
15     num_layers=2,
16     dropout=0.2,
17     bidirectional=True
18 )
19
20 # Forward pass
21 output, (hn, cn) = lstm(input_seq)
```

## Part 2: The Vanishing Gradient Problem

# Why Simple RNNs Fail

### The Problem

Gradient through time:

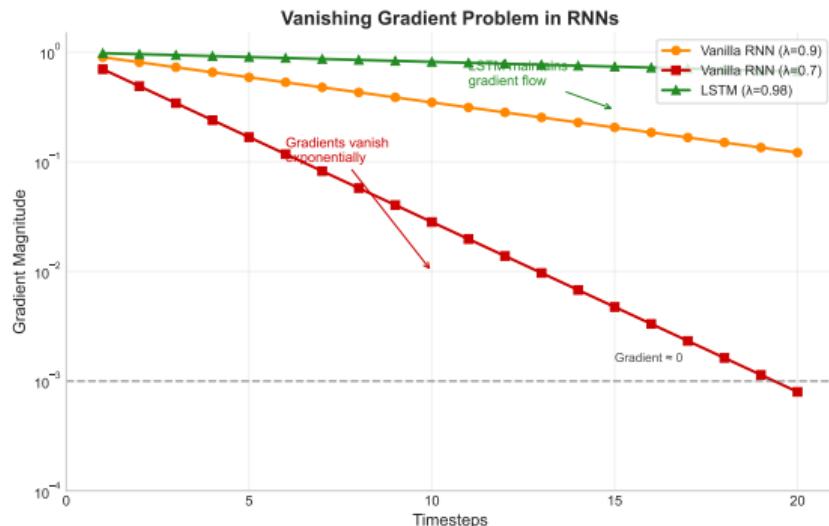
$$\frac{\partial L}{\partial h_0} = \frac{\partial L}{\partial h_T} \prod_{t=1}^T \frac{\partial h_t}{\partial h_{t-1}}$$

Each term:  $\frac{\partial h_t}{\partial h_{t-1}} = W_h^T \cdot \text{diag}(f'(h_{t-1}))$

For tanh:  $|f'(x)| \leq 1$

If  $\|W_h\| < 1$ : gradients  $\rightarrow 0$  (vanish)

If  $\|W_h\| > 1$ : gradients  $\rightarrow$  (explode)

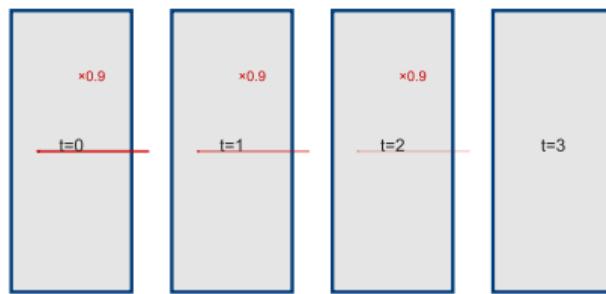


### Consequences:

- Can't learn long dependencies
- Gradient 0 after 10-20 steps
- Network "forgets" early inputs

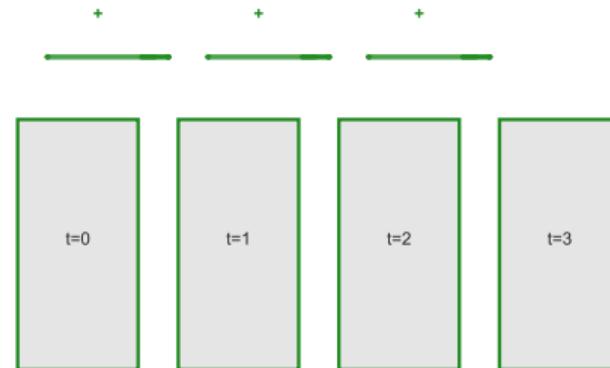
# Visualizing Gradient Flow

Vanilla RNN Gradient Flow



*Gradient vanishes through multiplications*

LSTM Gradient Flow



*Gradient flows through cell state highway*

## Vanilla RNN

- Exponential decay/growth
- Gradient magnitude:  $O(\lambda^T)$
- Effective memory: 5-10 steps
- **Cannot learn long patterns**

## LSTM (Next Section)

- Constant error flow
- Gradient highways
- Effective memory: 100+ steps
- **Learns long dependencies**

## Part 3: Long Short-Term Memory (LSTM)

# Engineering Memory

### The Innovation (1997)

Hochreiter Schmidhuber's insight:

- Add a **memory cell**  $C_t$
- Control flow with **gates**
- Create gradient highways
- Selective reading/writing

### Three Gates:

1. **Forget:** What to discard
2. **Input:** What to store
3. **Output:** What to expose

### LSTM Equations

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

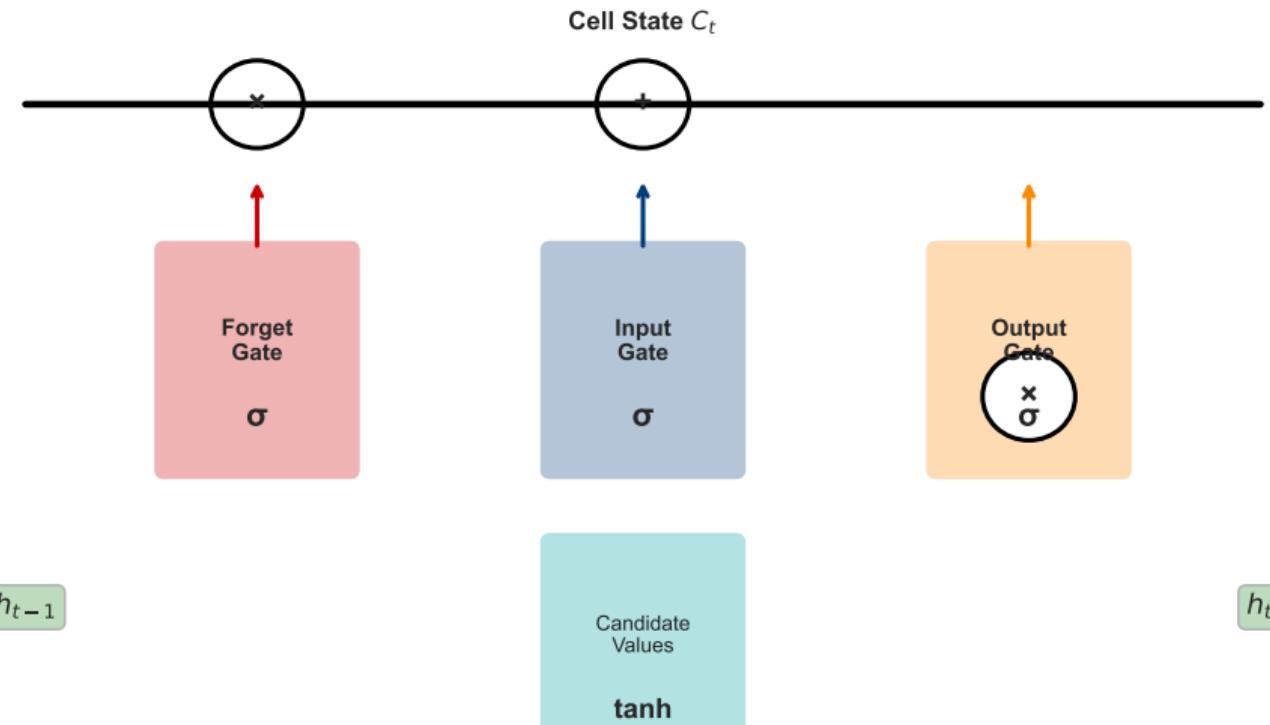
$$h_t = o_t * \tanh(C_t)$$

Gates use sigmoid (0-1) for control

The architecture that made deep sequence modeling possible

# LSTM Architecture: Gate Mechanisms

## LSTM Architecture: Information Flow Through Gates



# RNN vs LSTM: Key Differences

Aspect	Vanilla RNN	LSTM
Parameters	$O(h^2)$	$O(4h^2)$
Memory	Short (5-10 steps)	Long (100+ steps)
Gradient flow	Multiplicative	Additive
Training speed	Fast	Slower (4x params)
Gradient problem	Severe	Largely solved
Use cases	Short sequences	Most applications

## When to Use RNN:

- Very short sequences
- Real-time constraints
- Limited compute
- Simple patterns

## When to Use LSTM:

- Long dependencies
- Complex patterns
- Production systems
- Default choice (pre-2017)

LSTM's complexity is justified by superior performance

# GRU: Gated Recurrent Unit

## Simplification of LSTM (2014)

Cho et al. merged gates:

- Only **2 gates** instead of 3
- No separate cell state
- Fewer parameters (3x vs 4x)
- Similar performance

## GRU Equations:

$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

## GRU Architecture: Simplified Gating

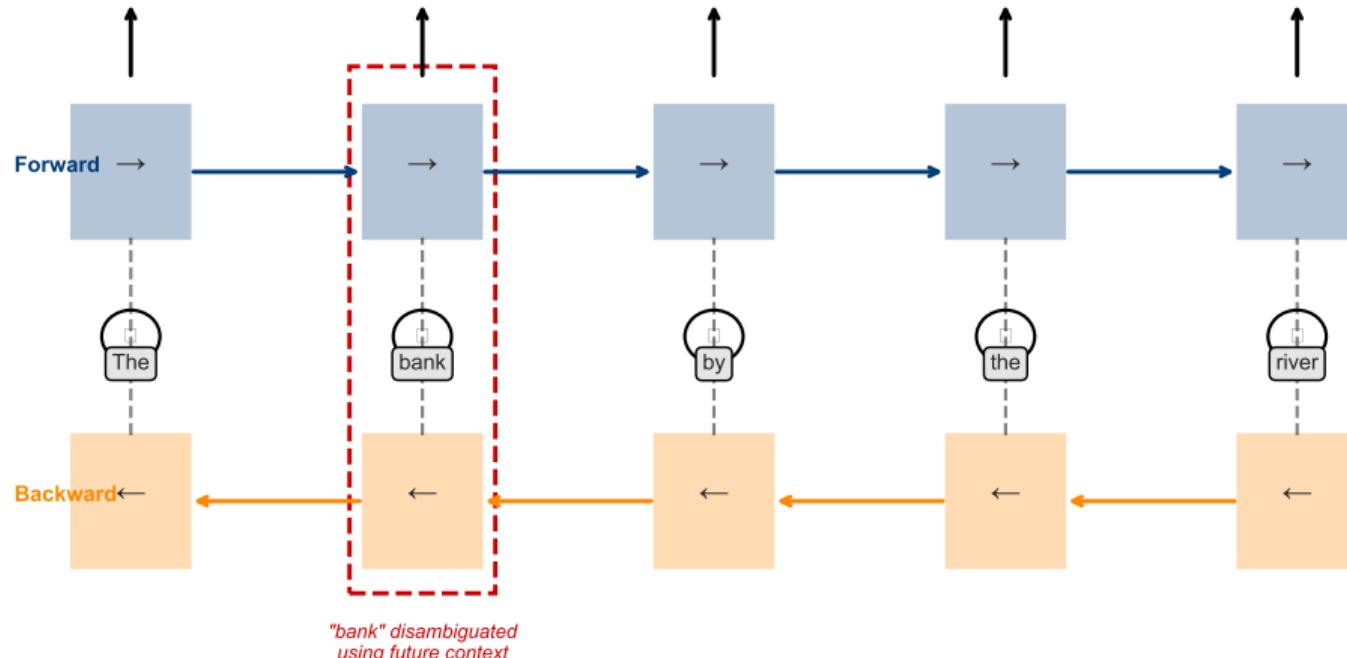


## Gates:

- **Update gate ( $z_t$ ):** How much to update
- **Reset gate ( $r_t$ ):** How much past matters

# Bidirectional RNNs: Using Future Context

Bidirectional RNN: Using Past and Future Context



# Training RNNs: Practical Tips

## Common Issues & Solutions

### 1. Gradient Explosion

- Solution: Gradient clipping
- `'torch.nn.utils.clip_grad_norm_'`
- Typical value: 1.0 - 5.0

### 2. Initialization

- Xavier/He initialization
- Forget gate bias = 1.0 (LSTM)
- Helps gradient flow

### 3. Overfitting

- Dropout (between layers)
- Recurrent dropout (careful!)
- Weight decay

## Hyperparameters

Parameter	Typical Range
Hidden size	128 - 512
Num layers	1 - 3
Learning rate	1e-3 - 1e-2
Batch size	32 - 128
Sequence length	20 - 200
Gradient clip	1.0 - 5.0
Dropout	0.2 - 0.5

## Training Strategy:

- Start with small sequences
- Gradually increase length
- Monitor gradient norms
- Use teacher forcing wisely

RNN training requires careful tuning and monitoring

# Real-World Applications

## Natural Language Processing

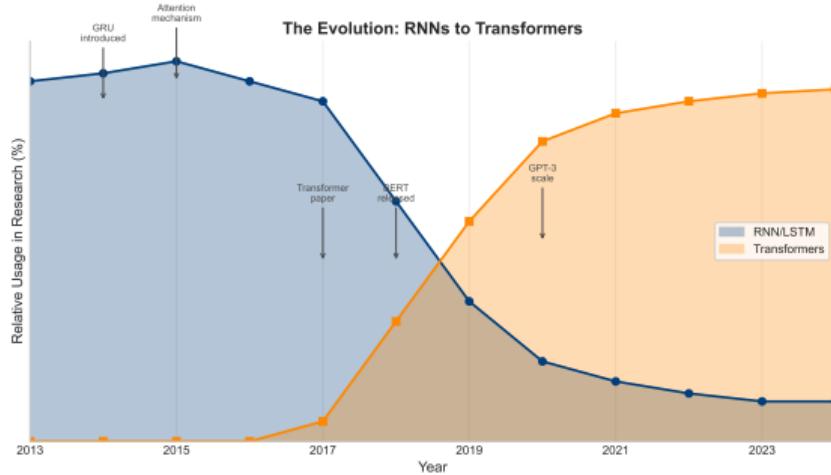
- Language modeling (pre-2018)
- Machine translation (pre-2017)
- Speech recognition (still used)
- Named entity recognition
- Sentiment analysis

## Time Series

- Stock price prediction
- Weather forecasting
- Anomaly detection
- Signal processing

RNNs remain relevant for specific use cases

## Modern Context (2024)

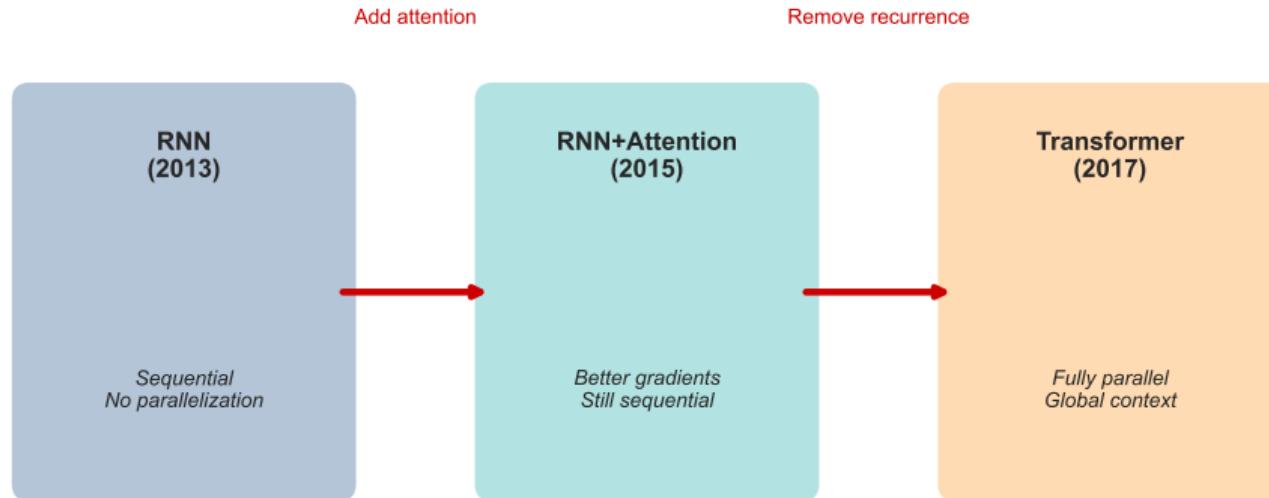


## Where RNNs Still Win:

- Streaming/online processing
- Edge devices (memory constraints)
- Variable-length sequences
- Time series with clear temporal patterns

# From RNNs to Transformers: Evolution

## Evolution of Sequence Models



# Key Takeaways

## What We Learned About RNNs

### Core Concepts

- **Recurrence** for sequences
- Hidden state as memory
- Parameter sharing
- Backprop through time

### Challenges

- Vanishing gradients
- Sequential bottleneck
- Training difficulty
- Limited context window

### Solutions

- LSTM/GRU gates
- Gradient clipping
- Bidirectional processing
- Attention (next week!)

RNNs introduced memory to neural networks - a crucial innovation

### Next Week: Sequence-to-Sequence Models

How to translate, summarize, and generate with encoder-decoder architectures

# References & Further Reading

## Foundational Papers:

- Hochreiter & Schmidhuber (1997). "Long Short-Term Memory"
- Cho et al. (2014). "Learning Phrase Representations using RNN Encoder-Decoder" (GRU)
- Graves (2013). "Generating Sequences With RNNs"
- Karpathy (2015). "The Unreasonable Effectiveness of RNNs" (blog)

## Practical Resources:

- PyTorch RNN Tutorial: [pytorch.org/tutorials/intermediate/char\\_rnn](https://pytorch.org/tutorials/intermediate/char_rnn.html)
- Understanding LSTMs: [colah.github.io/posts/2015-08-Understanding-LSTMs/](https://colah.github.io/posts/2015-08-Understanding-LSTMs/)
- Stanford CS224N Lecture 6: RNNs and Language Models

## Code Examples:

- Week 3 Lab: 'week03\_rnn\_lab.ipynb'
- GitHub: Various char-RNN implementations
- Hugging Face: Modern RNN models