

Recurrent Neural Networks (RNNs)

Understanding Sequential Data Processing

NLP Course - Week 3

BSc Level - Introduction

Why Do We Need RNNs?

Traditional Neural Networks:

- Process fixed-size inputs
- No memory of previous inputs
- Each prediction is independent

Problem: Many tasks involve sequences!

- Text: Words depend on previous words
- Speech: Sounds form words over time
- Video: Frames are connected

Intuition: Key Idea

We need networks that can “remember” previous inputs to understand context.

Examples of Sequential Data:

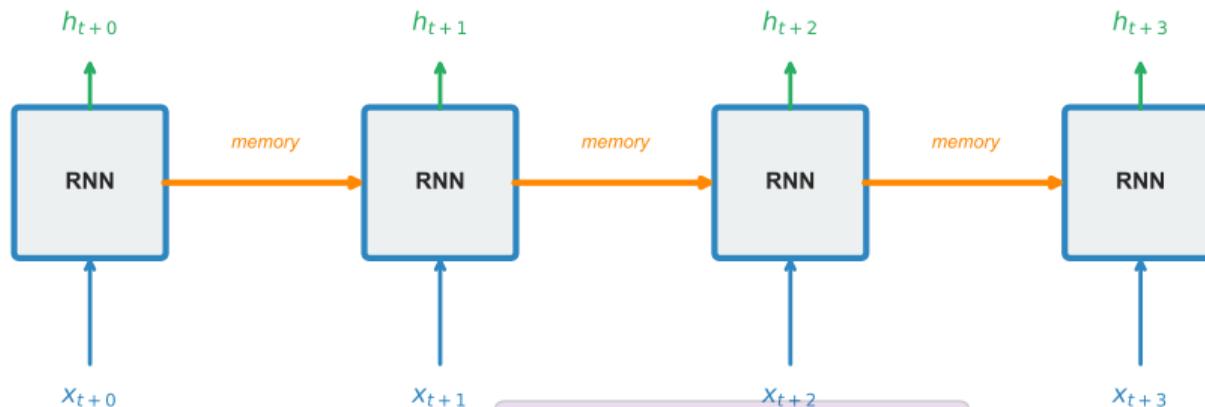
- ➊ **Names:** “Joh” → “n”
- ➋ **Sentences:** “The cat is ...”
- ➌ **Stock prices:** Yesterday → Today
- ➍ **Weather:** Past week → Tomorrow

Where Are RNNs Used? Real-World Applications



The RNN Solution: Adding Memory

RNN Unrolled Through Time



How RNNs Work:

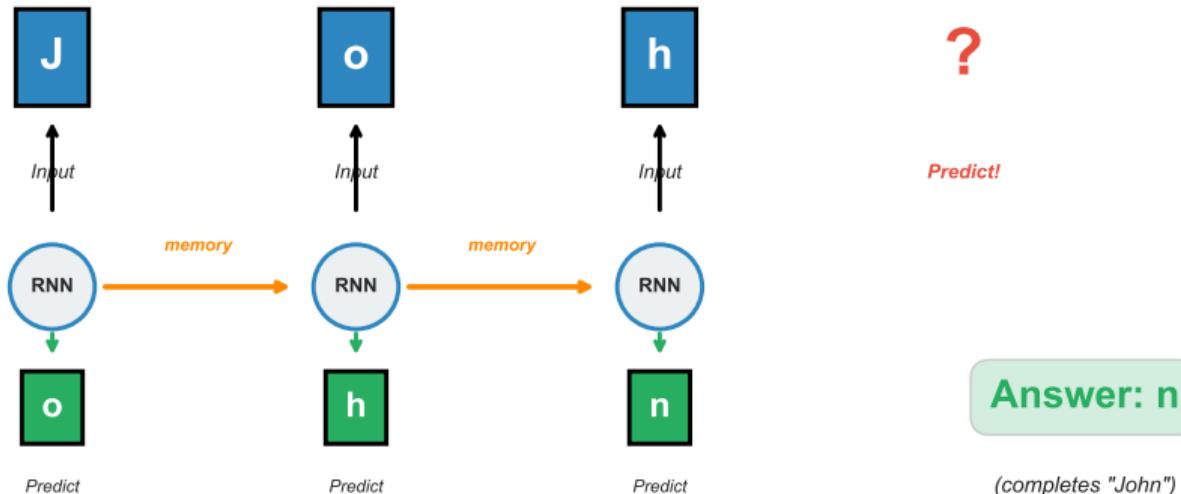
- ① Take current input x_t
- ② Combine with previous memory h_{t-1}
- ③ Produce output h_t
- ④ Pass memory to next step

Key Innovation:

- Same weights at each time step
- Memory flows through time
- Can process any sequence length

Example 1: Predicting Names

Example: Predicting Names Character by Character



Task: Given "Joh", predict "n"

RNN Process:

Checkpoint: Understanding

How RNN Processes Sequences

RNN Equations:

Hidden state update:

$$h_t = \tanh(W_x x_t + W_h h_{t-1} + b)$$

Output:

$$y_t = \text{softmax}(W_y h_t + b_y)$$

What each part means:

- x_t : Current input
- h_{t-1} : Previous memory
- W_x, W_h : Weight matrices
- \tanh : Activation function (-1 to 1)
- y_t : Output prediction

Intuition: Think of it like

A student taking notes:

- ① Read new information (x_t)
- ② Check your notes (h_{t-1})
- ③ Update your notes (h_t)
- ④ Answer question (y_t)

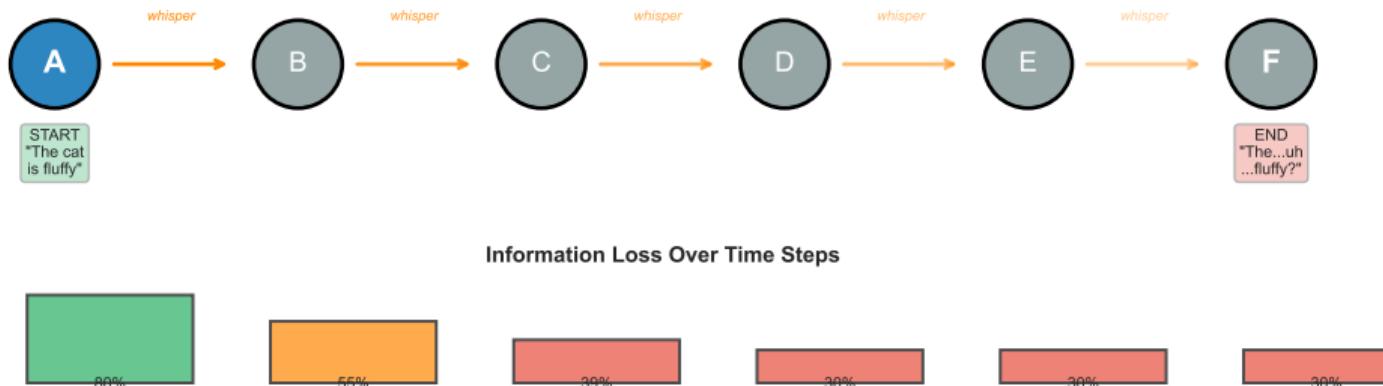
Key Parameters:

- Input size: 100-300
- Hidden size: 128-512
- Layers: 1-3

Problem: Vanishing Gradient (Memory Fading)

The Memory Problem: Telephone Game Analogy

Problem: Earlier information gets WEAKER as it travels through time
Solution: LSTM uses gates to preserve important information



The Telephone Game Problem:

- Message gets weaker as it passes
- Early information gets lost
- RNN forgets long-term context

Why This Happens:

- Gradients get multiplied many times
- Small numbers \times small numbers \rightarrow tiny
- Long sequences = many multiplications

Example 2: Sentiment Analysis

Task: Classify movie review as positive/negative

Input Sentence:

"This movie is absolutely fantastic!"

RNN Processing:

- ① Process "This" → h_1
- ② Process "movie" → h_2
- ③ Process "is" → h_3
- ④ Process "absolutely" → h_4
- ⑤ Process "fantastic" → h_5

Final hidden state h_5 contains sentiment info

→ Classify: **Positive (98%)**

PyTorch Code Example:

```
import torch.nn as nn

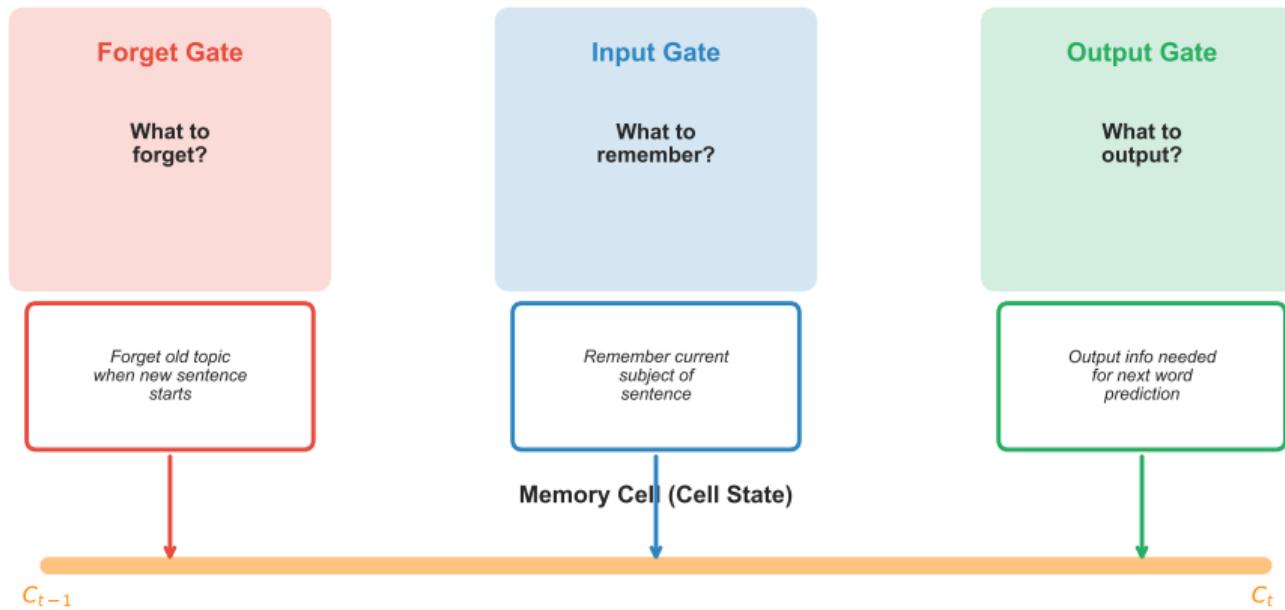
class SentimentRNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.embedding = nn.Embedding(
            num_embeddings=10000,
            embedding_dim=128
        )
        self.rnn = nn.RNN(
            input_size=128,
            hidden_size=256,
            num_layers=2
        )
        self.fc = nn.Linear(256, 2)

    def forward(self, x):
        # x: [seq_len, batch]
        embedded = self.embedding(x)
        output, hidden = self.rnn(embedded)
        # Use final hidden state
        prediction = self.fc(hidden[-1])
        return prediction
```

Output: [Negative, Positive]

Solution: LSTM (Long Short-Term Memory)

LSTM Solution: Three Smart Gates



Key: Gates control information flow to solve vanishing gradient problem

Example 3: Text Generation

Task: Generate text character by character

Training Data: Shakespeare's works

How it works:

- ① Start with seed text: "To be"
- ② LSTM predicts next character: " "
- ③ Add to sequence: "To be "
- ④ Predict next: "o"
- ⑤ Continue: "To be or"
- ⑥ Keep generating...

Generated Output:

"To be or not to be, that is the question"

LSTM learned:

- English grammar
- Shakespeare's style
- Common phrases

Simple Generation Code:

```
import torch
import torch.nn as nn

class CharRNN(nn.Module):
    def __init__(self, vocab_size):
        super().__init__()
        self.lstm = nn.LSTM(
            input_size=vocab_size,
            hidden_size=512,
            num_layers=2
        )
        self.fc = nn.Linear(512, vocab_size)

    def generate(self, start_text, length=100):
        hidden = None
        chars = [start_text]

        for _ in range(length):
            # Encode current char
            x = encode(chars[-1])
            # LSTM prediction
            out, hidden = self.lstm(x, hidden)
            # Sample next char
            next_char = sample(self.fc(out))
            chars.append(next_char)

        return ''.join(chars)

# Generate text
model.generate("To\u00ebl", length=50)
```

Summary: What We Learned

Key Concepts:

- ① **Sequential Data:** Many tasks involve sequences
- ② **RNN:** Networks with memory
- ③ **Hidden State:** Carries information forward
- ④ **Vanishing Gradient:** Memory fades over time
- ⑤ **LSTM:** Uses gates to preserve memory

Three Examples We Saw:

- Name prediction (supervised)
- Sentiment analysis (classification)
- Text generation (autoregressive)

Real World: Where RNNs Are Used

- **ChatGPT:** Text generation
- **Google Translate:** Language translation
- **Siri/Alexa:** Speech recognition
- **YouTube:** Video captioning
- **Gmail:** Smart compose

Next Steps:

- Week 4: Sequence-to-Sequence models
- Week 5: Transformers (attention mechanism)
- Lab: Implement your own LSTM

Questions?