LSTM - Long Short-Term Memory

Understanding Through a Complete Example

Watch LSTM Process a Sentence

Sentence: "The cat was hungry. The dog was sleeping."

Word	Forget	Input	Output	Memory
"The"	0.9	0.3	0.2	[article]
"cat"	0.8	0.9	0.8	[cat]
"was"	0.9	0.7	0.9	[cat, was]
"hungry"	0.8	0.8	0.7	[cat, hungry]
"."	**0.1**	0.4	0.3	[end]
"The"	0.1	0.8	0.2	[article]
"dog"	0.7	**0.9**	0.9	[dog]
"was"	0.9	0.8	**0.9**	[dog, was]

Notice:

- Three mysterious numbers per word
- Memory changes as we read
- "cat" disappears at period
- "dog" appears as new subject

Intuition: The Magic

How does LSTM know to:

- Forget "cat" at the period?
- Remember "dog" as new subject?
- Keep the right information?

Let's find out...

Why Do We Need This?

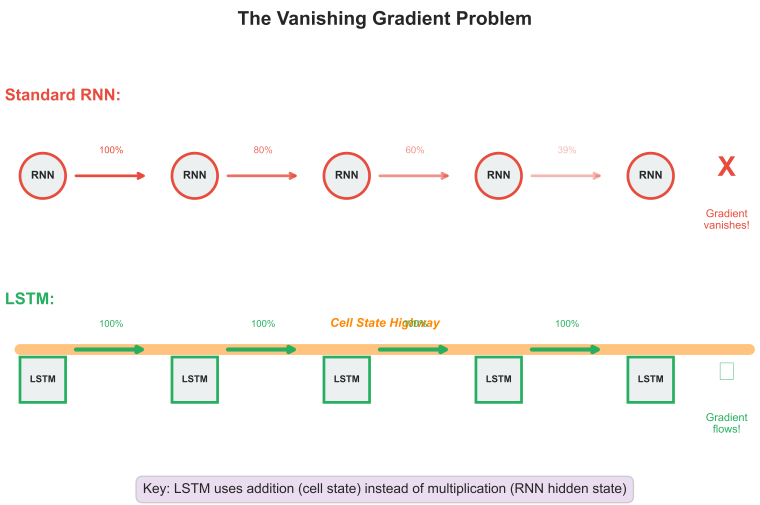

RNN Problem:

LSTM Solution:

Gate 1: Forget - What to Erase?

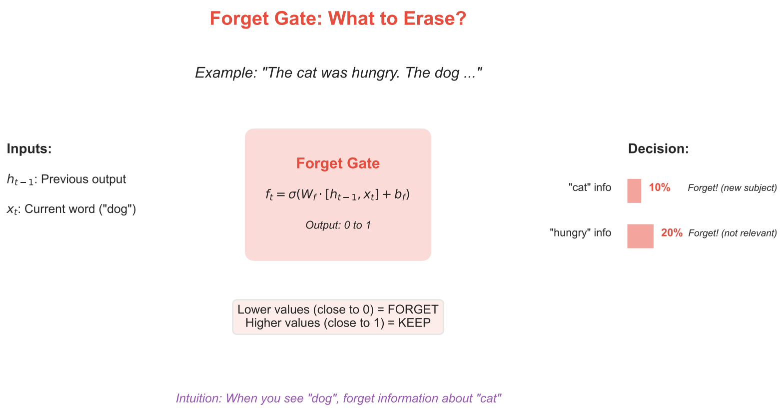

Remember Our Table?
Row 5: "." had **Forget = 0.1**

What This Means:

Formula:

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

How It Decides:

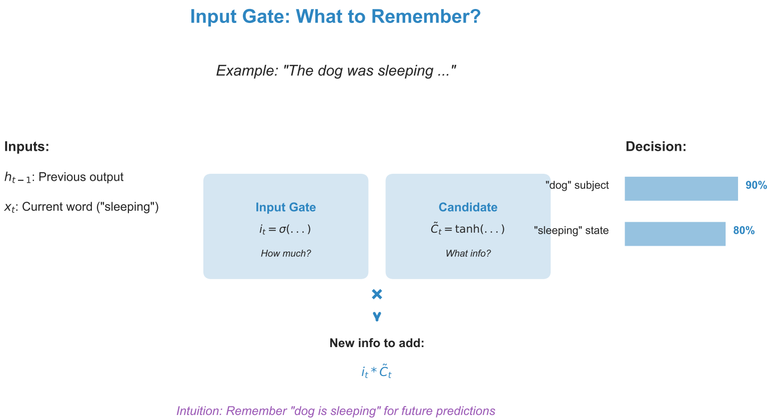

Remember Our Table?
Row 7: "dog" had **Input = 0.9**

What This Means:

Formulas (Two Parts):

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

Gate 3: Output - What to Share?

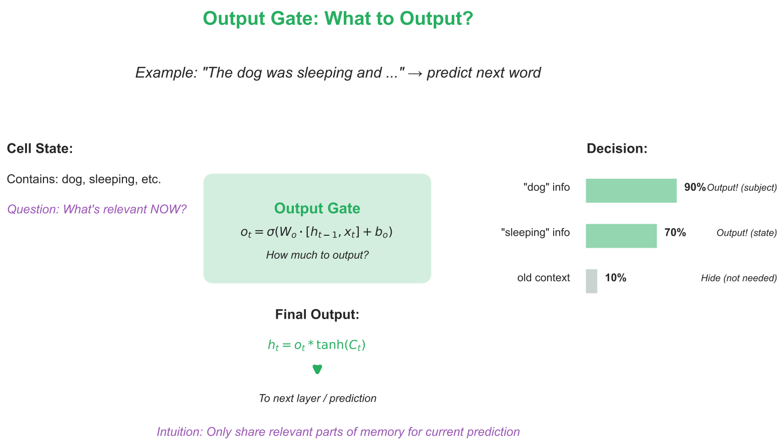

Remember Our Table?
Row 8: "was" had **Output = 0.9**

What This Means:

Formulas:

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t \odot \tanh(C_t)$$

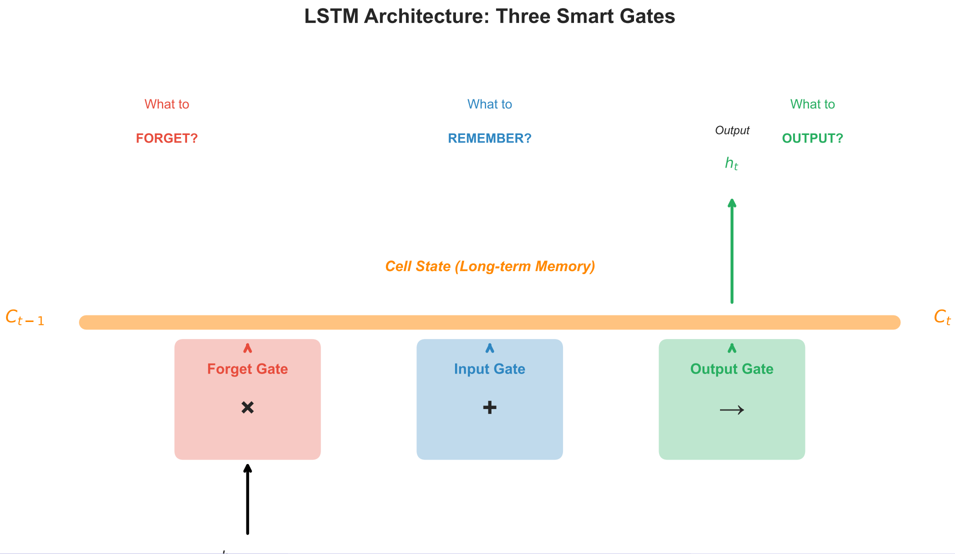
LSTM Architecture: Three Smart Gates
What to
FORGET?
What to
REMEMBER?
What to
OUTPUT?
Output
h_t
Cell State (Long-term Memory)
C_{t-1}
C_t
Forget Gate
×
Input Gate
+
Output Gate
→

Now Let's Look Again - You Understand It!

Sentence: "The cat was hungry. The dog was sleeping."

Word	Forget	Input	Output	Memory
"The"	0.9 (keep)	0.3 (small add)	0.2 (hide)	[article]
"cat"	0.8 (keep)	0.9 (add subject!)	0.8 (show)	[cat]
"was"	0.9 (keep)	0.7 (add verb)	0.9 (need it!)	[cat, was]
"hungry"	0.8 (keep)	0.8 (add state)	0.7 (show)	[cat, hungry]
"."	**0.1 (ERASE!)**	0.4 (ending)	0.3 (hide)	[end]
"The"	0.1 (clear old)	0.8 (new start)	0.2 (hide)	[article]
"dog"	0.7 (keep)	**0.9 (NEW subject!)**	0.9 (show)	[dog]
"was"	0.9 (keep)	0.8 (add verb)	**0.9 (USE dog!)**	[dog, was]

Checkpoint: The Critical Transition

Watch rows 4-7: "hungry" → "." → "The" → "dog"
Memory evolves: [cat, hungry] → FORGET → [end] → ADD → [dog]
This is what RNNs cannot do! LSTM uses gates to control memory intelligently.

The Complete Mathematics

All Three Gates:

Forget Gate:

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

Input Gate:

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Output Gate:

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

Cell State Update:

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$

Output:

$$h_t = o_t \odot \tanh(C_t)$$

In Our Example:

- x_t = current word embedding
- h_{t-1} = previous output
- C_t = cell state (memory)
- σ = sigmoid (0 to 1)
- tanh = tanh (-1 to 1)
- $\odot$ = element-wise multiply

Summary: LSTM in Practice

What We Learned:

1. LSTM uses three gates to control memory
2. Forget gate: what to erase (0.1 at "." = erase "cat")
3. Input gate: what to add (0.9 at "dog" = add subject)
4. Output gate: what to use (0.9 at "was" = use "dog")
5. Cell state flows information forward

When to Use LSTM:

- Long sequences (100+ words)
- Long-term dependencies
- Context matters
- Grammar and structure

Real World: Applications

Where LSTMs Excel:

- Machine Translation (Google Translate)
- Speech Recognition (Siri, Alexa)
- Text Generation (ChatGPT foundations)
- Video Analysis
- Music Generation
- Handwriting Recognition

Key Takeaway:
The sentence example showed you *exactly* how LSTM gates work in practice. That's the real magic!

Questions?