

Sequence-to-Sequence Models

Nature Professional Theme Template

NLP Course - Week 4

2025

Overview

The Nature Professional Theme

Color Palette

- **Forest Green** - Primary (#14532D)
- **Teal** - Secondary (#0D9488)
- **Amber** - Accent (#F59E0B)
- **Slate** - Support (#475569)
- **Mint Cream** - Background

Design Principles

- **Natural harmony**
- **Clear hierarchy**
- **Minimal distraction**
- **Maximum readability**



Why Nature Professional?

- Reduces eye strain
- Professional appearance
- Excellent contrast
- Calming effect

Sequence-to-Sequence Architecture

The Encoder-Decoder Framework


- Encoder: Processes input sequence
 - Maps variable-length input to fixed representation
 - Captures semantic information
- Decoder: Generates output sequence
 - Converts representation to target sequence
 - Autoregressive generation

Mathematical Formulation

$$h_t = \text{RNN}_{\text{enc}}(x_t, h_{t-1}) \quad (1)$$

$$s_t = \text{RNN}_{\text{dec}}(y_{t-1}, s_{t-1}, c) \quad (2)$$

$$p(y_t | y_{<t}, x) = \text{softmax}(W_s s_t) \quad (3)$$



Key Innovation

Variable-length input → Variable-length output

Nature-Themed Attention Heatmap

Attention Mechanism Heatmap
Translation: English → French




Model Complexity Comparison

Seq2Seq Model Evolution: Nature Professional Theme





Evolution of model architectures with nature-inspired visualization

Beam Search Tree Visualization




Training Dynamics Dashboard

Training Dynamics Dashboard - Nature Professional



Architecture Evolution Timeline



Historical progression with nature-themed annotations

Nature Professional Theme

Encoder Implementation

```
class Encoder(nn.Module):
    def __init__(self, vocab_size,
                 hidden_dim):
        super().__init__()
        self.embedding = nn.Embedding(
            vocab_size, hidden_dim)
        self.rnn = nn.LSTM(
            hidden_dim, hidden_dim)

    def forward(self, x):
        embedded = self.embedding(x)
        output, hidden = self.rnn(embedded)
        return output, hidden
```

Attention Mechanism

```
def attention(query, keys, values):
    # Compute attention scores
    scores = torch.matmul(
        query, keys.transpose(-2, -1))
    scores = scores / sqrt(d_k)

    # Apply softmax
    weights = F.softmax(scores, dim=-1)

    # Weighted sum of values
    context = torch.matmul(weights, values)
    return context, weights
```

Performance Metrics

Model	BLEU	Time
Simple RNN	65.2	1.0x
LSTM	78.4	2.5x
GRU	76.1	2.2x
Basic Seq2Seq	82.3	3.0x
+ Attention	89.1	4.0x
Transformer	94.5	2.0x

Key Findings

- **45% improvement** with attention
- **2x faster** training with Transformers
- **Better long-range dependencies**

Translation Examples

Source: The cat sat on the mat

Basic: Le chat assis tapis

Attention: Le chat s'est assis sur le tapis

Source: I love natural language processing

Basic: J'aime traitement langue

Attention: J'adore le traitement du langage naturel

Error Analysis

- **Word order:** 15% errors
- **Grammar:** 8% errors
- **Vocabulary:** 5% errors

Beyond Basic Seq2Seq

Multi-Head Attention


- Parallel attention mechanisms
- Different representation subspaces
- Enhanced expressiveness

Copy Mechanisms

- Direct copying from source
- Hybrid generation-copy approach
- Improved rare word handling

Coverage Models

- Track attention history
- Prevent repetition
- Ensure completeness



Future Directions

- ① *Cross-lingual* models
- ② *Multimodal* seq2seq
- ③ *Few-shot* learning

Key Takeaways

What We Learned


- ① Encoder-Decoder architecture
- ② Attention mechanism importance
- ③ Beam search for decoding
- ④ Training dynamics and optimization
- ⑤ Evolution to Transformers

Practical Applications

- Machine Translation
- Text Summarization
- Dialog Systems
- Code Generation
- Image Captioning

Next Steps

- Week 5: **Transformers**
- Week 6: *Pre-trained Models*
- Week 7: **Advanced Architectures**



Thank You!