

Sequence-to-Sequence Models

Encoder-Decoder Architecture & Attention Mechanisms

Natural Language Processing - Week 4

2025

Overview

- 1 Introduction
- 2 Encoder-Decoder Architecture
- 3 Training Process
- 4 Attention Revolution
- 5 Decoding Strategies
- 6 Performance Analysis
- 7 Implementation Details
- 8 Practical Considerations
- 9 Historical Evolution
- 10 Hands-on Exercise
- 11 Summary

The Variable-Length Challenge

Why Sequence-to-Sequence?

Previous models' limitations:

- Fixed-size input/output
- Cannot handle variable lengths
- Loss of sequential information

Key Innovation

Map variable-length input to variable-length output

Applications

- Machine Translation
 - English → French
 - Variable sentence lengths
- Text Summarization
 - Long article → Brief summary
- Dialogue Systems
 - Question → Answer

Sutskever et al., 2014 - Revolutionary approach to sequence modeling

The Core Architecture

Encoder

- Processes **input sequence**
- Creates context vector c
- Hidden states: h_1, h_2, \dots, h_T

Encoder Equations

$$h_t = \text{RNN}_{\text{enc}}(x_t, h_{t-1}) \quad (1)$$

$$c = h_T \quad (2)$$

Decoder

- Generates **output sequence**
- Uses context vector c
- Autoregressive generation

Decoder Equations

$$s_t = \text{RNN}_{\text{dec}}(y_{t-1}, s_{t-1}, c) \quad (3)$$

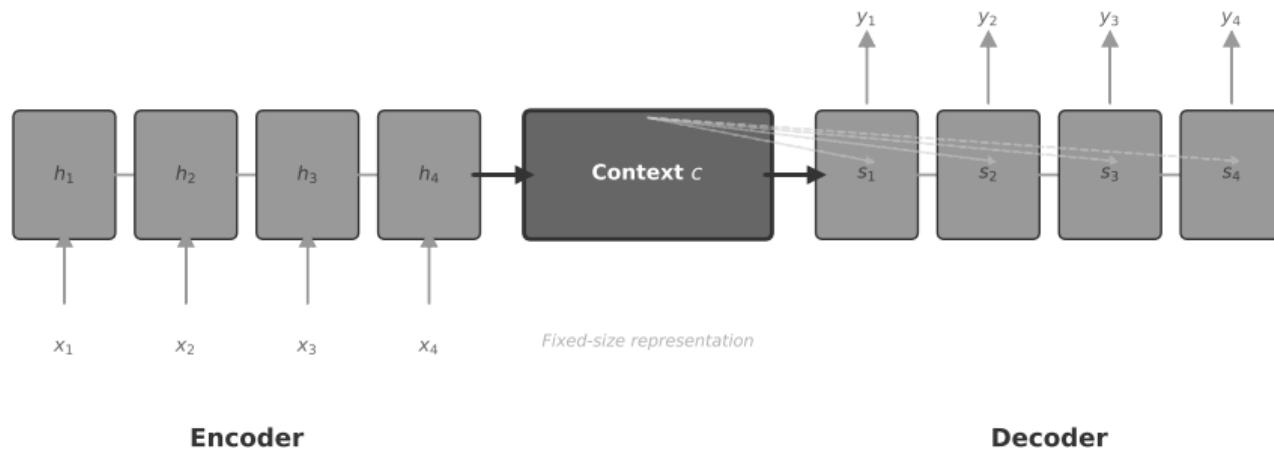
$$p(y_t | y_{<t}, x) = \text{softmax}(W_s s_t) \quad (4)$$

The context vector c is the information bottleneck

Variable input length

Variable output length

Sequence-to-Sequence Architecture



Teacher Forcing vs Inference

Training: Teacher Forcing

- Use **ground truth** as input
- Faster convergence
- Exposure bias problem

Training Process

At each step t :

- ① Input: true y_{t-1}
- ② Predict: \hat{y}_t
- ③ Loss: $\mathcal{L}(y_t, \hat{y}_t)$

Inference: Autoregressive

- Use **own predictions**
- Error accumulation
- Beam search helps

Inference Process

At each step t :

- ① Input: predicted \hat{y}_{t-1}
- ② Predict: \hat{y}_t
- ③ Continue until EOS

Teacher forcing creates a mismatch between training and inference

The Attention Solution

Motivation

- **Bottleneck** in fixed-size c
- Long sequences lose information
- Need **dynamic context**

Key Idea

Different decoder steps attend to different encoder positions

Attention Benefits

- Solves long-range dependencies
- Provides alignment
- Improves gradient flow

Attention Computation

Mathematical Formulation

$$e_{tj} = a(s_{t-1}, h_j) \quad (5)$$

$$\alpha_{tj} = \frac{\exp(e_{tj})}{\sum_{k=1}^T \exp(e_{tk})} \quad (6)$$

$$c_t = \sum_{j=1}^T \alpha_{tj} h_j \quad (7)$$

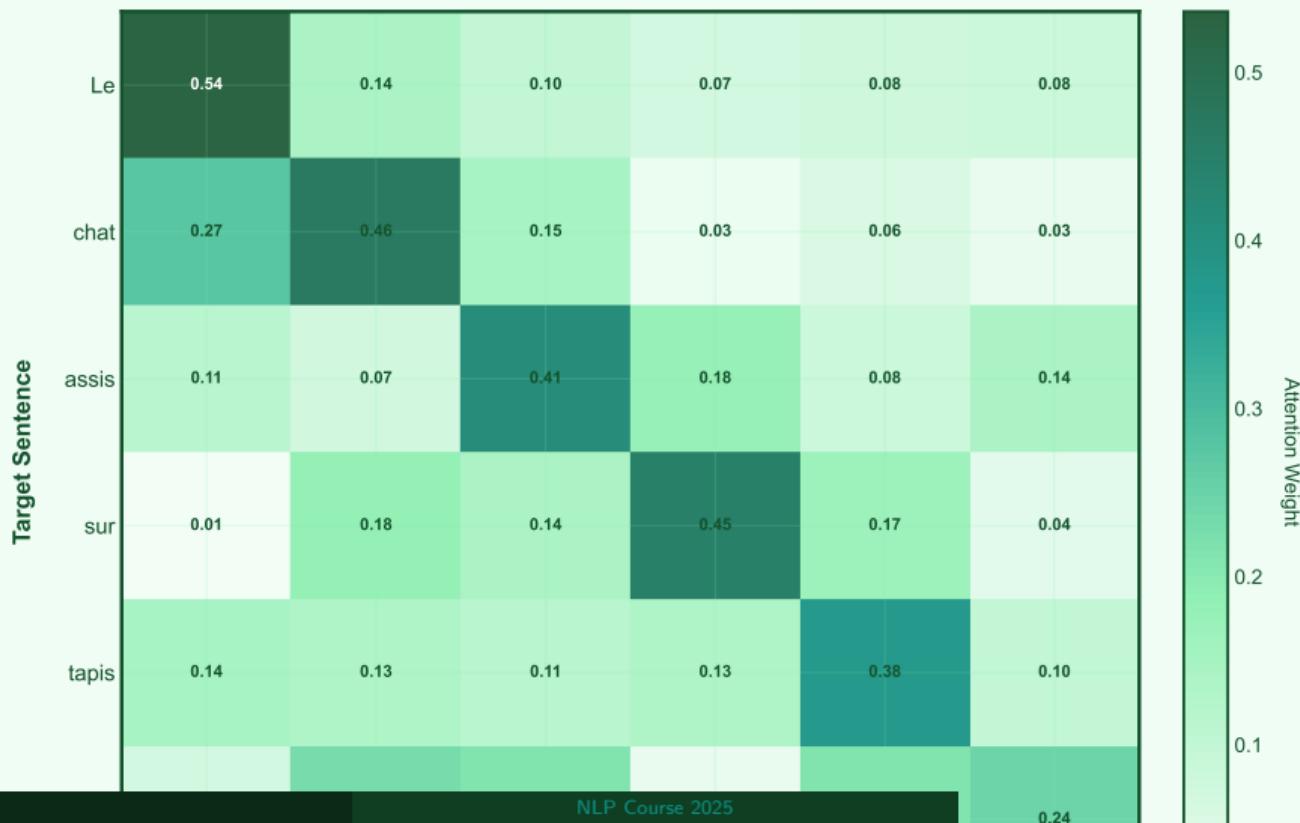
Where:

- e_{tj} : **alignment score**
- α_{tj} : **attention weight**
- c_t : **dynamic context**

Bahdanau et al., 2015 - Neural Machine Translation by Jointly Learning to Align and Translate

Attention Visualization

Attention Mechanism Heatmap
Translation: English → French



Beam Search Decoding

Why Not Greedy?

- Greedy: **locally optimal**
- May miss **better sequences**
- No backtracking possible

Beam Search Solution

- Keep **k best** hypotheses
- Explore multiple paths
- Balance quality vs speed

Beam Width Trade-off

- $k = 1$: Greedy (fast, lower quality)
- $k = 5 - 10$: Typical (balanced)
- $k = \infty$: Exhaustive (slow, optimal)

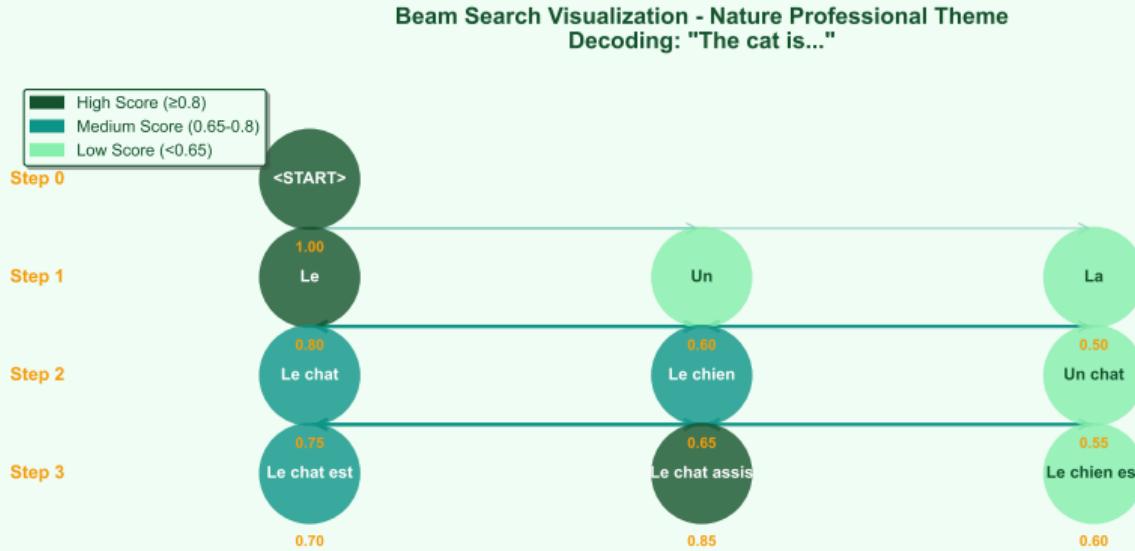
Beam search significantly improves translation quality

Algorithm

Beam Search Steps

- ① Initialize beam with $\langle \text{START} \rangle_i$
- ② For each time step:
 - Expand all hypotheses
 - Score all candidates
 - Keep top- k sequences
- ③ Stop when:
 - All beams end with $\langle \text{EOS} \rangle_i$
 - Maximum length reached
- ④ Return highest scoring sequence

Beam Search Tree



Beam width $k = 3$: Exploring multiple translation hypotheses

Model Comparison

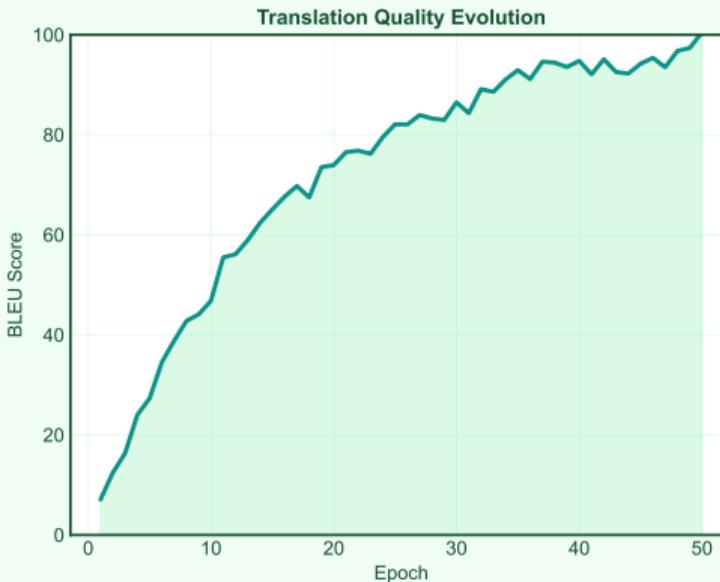
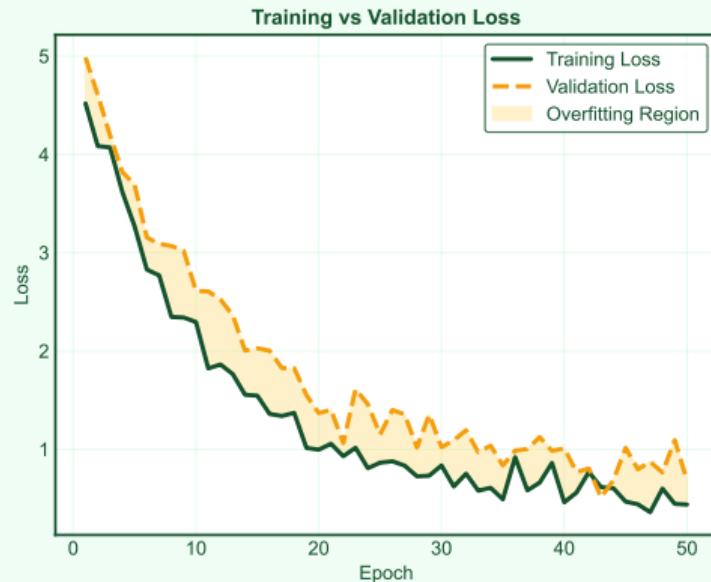
Seq2Seq Model Evolution: Nature Professional Theme



Evolution from RNN to Transformer: complexity vs performance trade-offs

Training Dynamics

Training Dynamics Dashboard - Nature Professional



PyTorch Implementation

Encoder

```
class Encoder(nn.Module):
    def __init__(self, input_dim,
                 emb_dim, hid_dim,
                 n_layers, dropout):
        super().__init__()
        self.embedding = nn.Embedding(
            input_dim, emb_dim)
        self.rnn = nn.LSTM(emb_dim,
                          hid_dim, n_layers,
                          dropout=dropout)
        self.dropout = nn.Dropout(dropout)

    def forward(self, src):
        embedded = self.dropout(
            self.embedding(src))
        outputs, (hidden, cell) = \
            self.rnn(embedded)
        return hidden, cell
```

Decoder with Attention

```
class Decoder(nn.Module):
    def __init__(self, output_dim,
                 emb_dim, hid_dim,
                 n_layers, dropout):
        super().__init__()
        self.embedding = nn.Embedding(
            output_dim, emb_dim)
        self.rnn = nn.LSTM(emb_dim + hid_dim,
                          hid_dim, n_layers,
                          dropout=dropout)
        self.attention = Attention(hid_dim)
        self.fc = nn.Linear(
            hid_dim * 2, output_dim)

    def forward(self, input, hidden,
               cell, encoder_outputs):
        embedded = self.embedding(input)
        a = self.attention(hidden[-1],
                           encoder_outputs)
        weighted = torch.bmm(a.unsqueeze(1),
                             encoder_outputs).squeeze(1)
        rnn_input = torch.cat(
            (embedded, weighted), dim=1)
        output, (hidden, cell) = \
            self.rnn(rnn_input.unsqueeze(0),
                     (hidden, cell))
        prediction = self.fc(torch.cat(
            (output.squeeze(0), weighted),
            dim=1))
        return prediction, hidden, cell
```

Complete implementation available in course repository

Training Tips & Tricks

Optimization

- Gradient clipping essential
- Learning rate scheduling
- Dropout in RNN layers

Hyperparameters

- Hidden size: 256-512
- Layers: 2-4
- Dropout: 0.2-0.3
- Beam width: 5-10

Data Preprocessing

- Tokenization consistency
- Vocabulary size limits
- Handle rare words

Common Issues

Exposure Bias

- Problem: Train/test mismatch
- Solution: Scheduled sampling

Long Sequences

- Problem: Gradient vanishing
- Solution: Attention mechanism

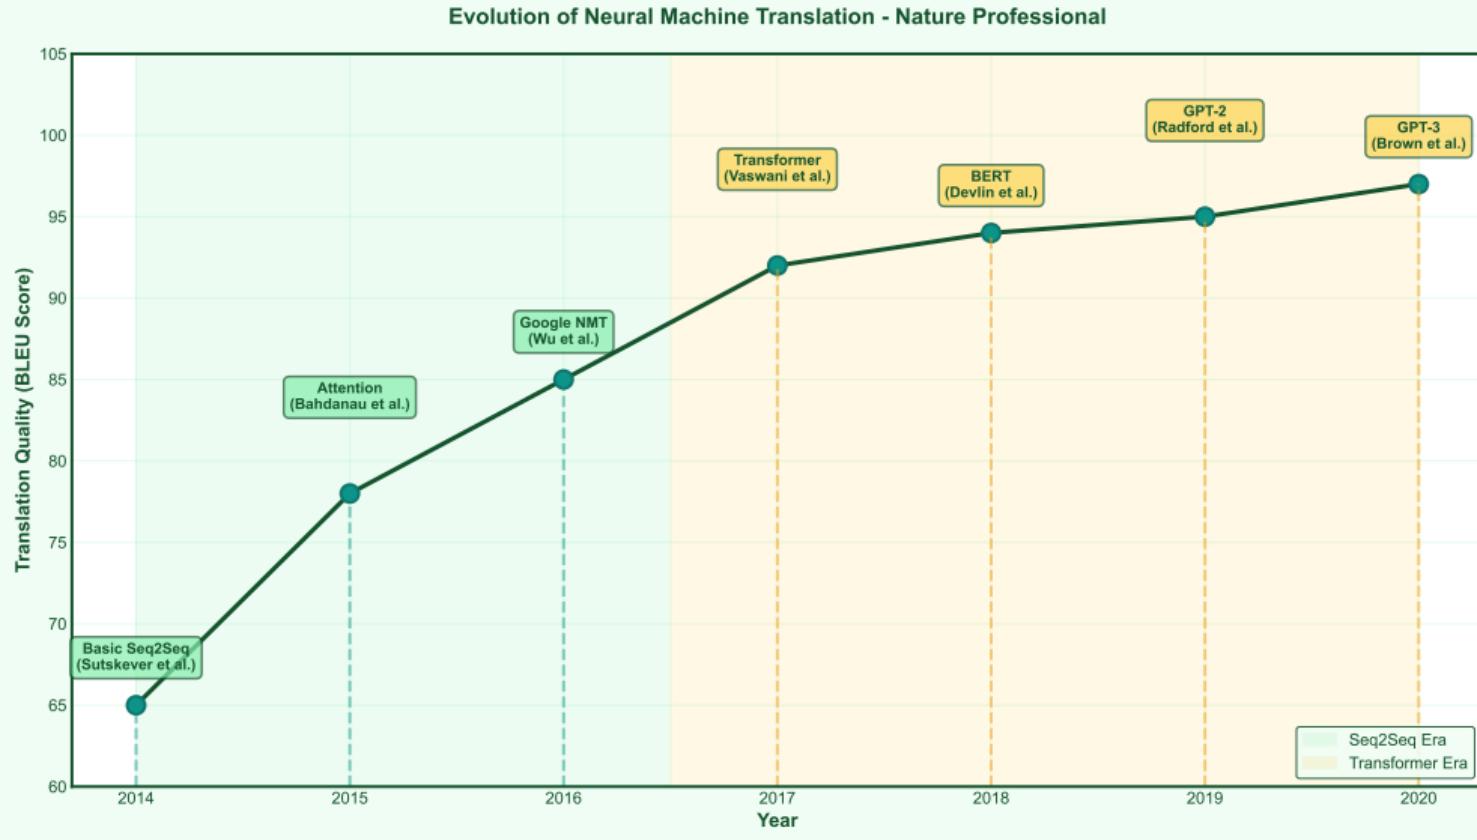
Rare Words

- Problem: UNK tokens
- Solution: Copy mechanism

Evaluation Metrics

- BLEU: n-gram precision
- ROUGE: recall-oriented
- METEOR: semantic similarity

From Seq2Seq to Transformers



Exercise: Simple Translation

Task

Implement a basic seq2seq model for number translation:

- Input: "one two three"
- Output: "1 2 3"

Starter Code

```
# Define vocabularies
source_vocab = {
    'one': 1, 'two': 2,
    'three': 3, 'four': 4,
    'five': 5, '<sos>': 0,
    '<eos>': 6, '<pad>': 7
}
target_vocab = {
    '1': 1, '2': 2, '3': 3,
    '4': 4, '5': 5,
    '<sos>': 0, '<eos>': 6,
    '<pad>': 7
}

# Training data
train_data = [
    ('one-two', '1-2'),
    ('three-four-five', '3-4-5'),
    # Add more examples
]
```

Implementation Steps

Step 1: Data Processing

- Tokenize sequences
- Convert to indices
- Pad to same length

Step 2: Model Building

- Create Encoder class
- Create Decoder class
- Combine into Seq2Seq

Step 3: Training Loop

- Teacher forcing
- Calculate loss
- Backpropagation

Key Takeaways

What We Learned

① Encoder-Decoder architecture

- Variable length handling
- Context vector bottleneck

② Attention Mechanism

- Dynamic context
- Alignment visualization

③ Training Techniques

- Teacher forcing
- Beam search decoding

Core Insight

Attention solves the information bottleneck problem

Seq2seq models laid the foundation for modern NLP architectures

Practical Applications

- Machine Translation
- Text Summarization
- Question Answering
- Image Captioning
- Speech Recognition

Next Week Preview

Week 5: Transformers

- Self-attention mechanism
- Parallel processing
- Positional encoding
- Multi-head attention

References

Essential Papers

- Sutskever et al. (2014)
Sequence to Sequence Learning with Neural Networks
- Bahdanau et al. (2015)
Neural Machine Translation by Jointly Learning to Align and Translate
- Luong et al. (2015)
Effective Approaches to Attention-based Neural Machine Translation

Additional Resources

- Course Repository
github.com/course/week4-seq2seq
- Lab Notebook
week04_seq2seq.lab.ipynb
- PyTorch Tutorial
pytorch.org/tutorials/seq2seq

Office Hours

Tuesday 2-4 PM
Thursday 10-12 AM

Questions? Please post on the course forum or attend office hours

Thank You!

See you next week for Transformers

Remember to complete the lab exercise