

# Sequence-to-Sequence Models

## Week 4: The Translation Revolution with Attention

NLP Course 2025

Professional Template Edition

September 29, 2025

## Week 4: Journey Through Translation and Attention

**Learning Path:** From word-by-word replacement to neural translation. Master encoder-decoder architectures, understand the bottleneck problem, and discover how attention revolutionized machine translation.

## The Challenge: Lost in Translation

*Why Machines Struggle with Language Translation*

# The Google Translate Evolution: A Success Story

## 2006: Statistical MT

- Word/phrase dictionaries
- Counted co-occurrences
- “Reasonable” translations
- Often awkward phrasing

## 2016: Neural MT Launch

- Seq2Seq with attention
- Human-quality for some pairs
- 60% error reduction
- Revolutionary improvement

## Real Example:

*Chinese Input:* “There is one cat in station”

**Old:** “In the station is one cat”

**New:** “There is a cat at the station”

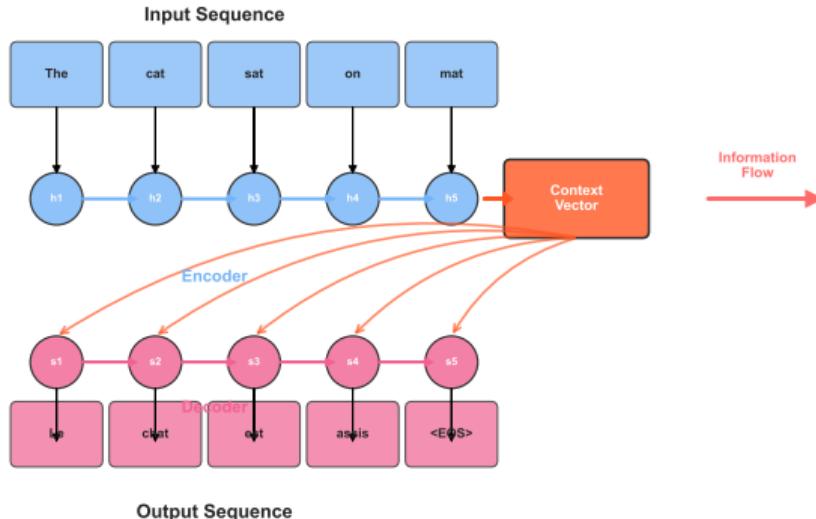
**What changed?** Understanding context, not just words

---

**Historical Context:** Neural MT reduced translation errors by 60% overnight - the biggest leap in MT history

# The Fundamental Problem: Meaning Across Languages

Sequence-to-Sequence Architecture: Encoder-Decoder with Context Vector



## Translation is NOT:

- Word replacement
- Dictionary lookup
- Rule application
- Rule application

## Translation IS:

- Understanding meaning
- Cultural context
- Reformulation
- Generation

# Why Word-by-Word Translation Fails: Concrete Examples

## Problem 1: Word Order

- English: "I saw the red house"
- Spanish: "Vi la casa roja"
- Literal: "Saw-I the house red"

## Problem 2: Idioms

- English: "It's raining cats and dogs"
- French: "Il pleut des cordes"
- Literal: "It rains ropes"

## Problem 3: Context

- "Bank" → "Banque" (financial)
- "Bank" → "Rive" (river)
- Need full sentence to decide

## Problem 4: Grammar

- German: Verb at end
- Japanese: Subject optional
- Chinese: No tenses

**Conclusion:** Languages encode meaning differently - translation needs deep understanding

---

Language Diversity: Each language has unique ways of expressing ideas

# Converting Meaning to Numbers: The Core Challenge

Computers only understand numbers, so:

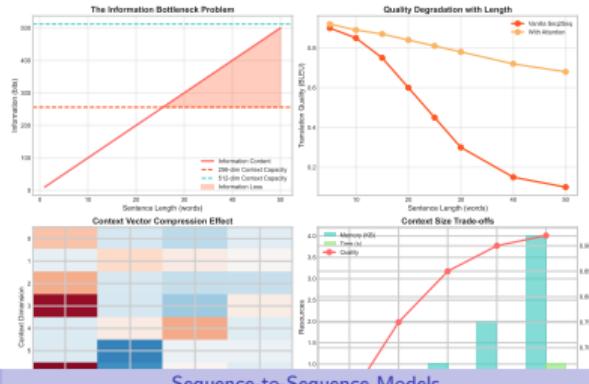
"The cat sat on the mat" → [Numbers] → "Le chat s'est assis sur le tapis"

## Step 1: Words to Vectors

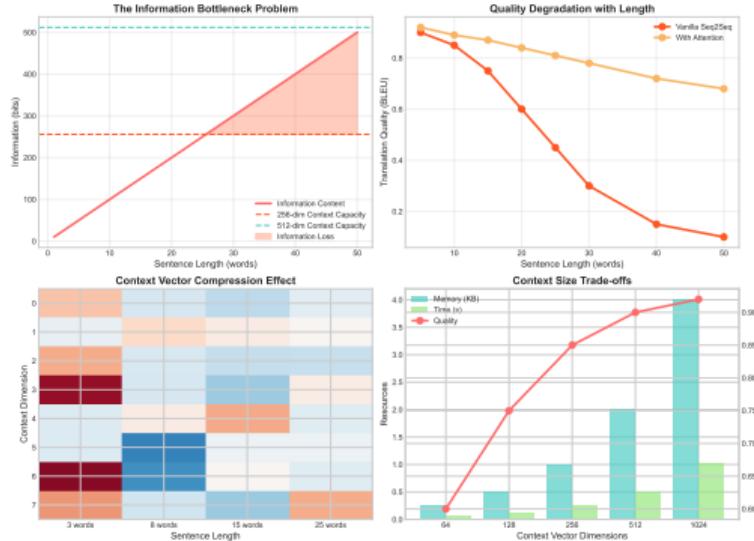
- "cat" → [0.2, -0.5, 0.8, ...]
- 100-300 dimensional vectors
- Learned from context (Word2Vec)
- Similar words = nearby vectors

## Step 2: Sentence to Vector

- Combine word vectors
- Build "context vector"
- Fixed size (e.g., 256 dims)
- Must capture ALL meaning



# The Compression Challenge: Information Bottleneck



## Compression Ratios:

- 5 words: 500 dims  $\rightarrow$  256 (2:1)
- 20 words: 2000 dims  $\rightarrow$  256 (8:1)
- 50 words: 5000 dims  $\rightarrow$  256 (20:1)

**Problem:** More compression = More loss

## What Gets Lost?

- Specific word choices
- Grammatical nuances
- Word positions
- Long-range dependencies

# Interactive Exercise: Manual Translation Steps

**Task:** Translate “The black cat sat” to French step-by-step

## Your Steps:

1. Read entire English sentence
2. Identify: subject (cat), verb (sat)
3. Recall French words:
  - cat → chat
  - black → noir
  - sat → s'est assis
4. Apply French grammar:
  - Article-Noun-Adjective order
  - Gender agreement (le/la)
5. Generate: “Le chat noir s'est assis”

## What You Actually Did:

1. Encoded English to meaning
2. Stored meaning in memory
3. Decoded meaning to French

This is exactly Seq2Seq!

## Key Observation:

You didn't translate word-by-word! You understood first, then generated.

---

**Human Insight:** We naturally use encoder-decoder approach when translating

# Calculating the Bottleneck: A Mathematical Perspective

## Information Content:

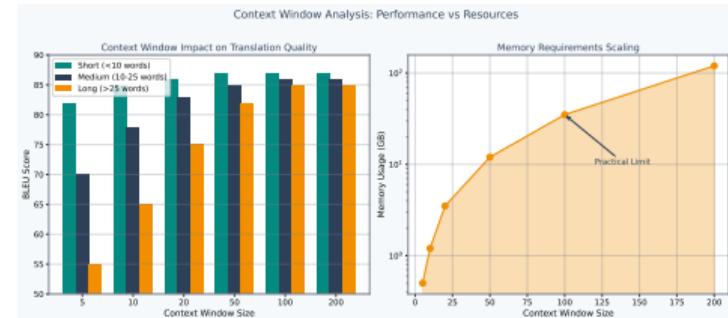
$$\text{Input} = n \times d_{\text{embed}}$$

$$\text{Context} = d_{\text{hidden}}$$

$$\text{Ratio} = \frac{n \times d_{\text{embed}}}{d_{\text{hidden}}} \quad (3)$$

## Example Calculation:

- 20 words, 100-dim embeddings
- Input:  $20 \times 100 = 2000$  values
- Context: 256 values
- Compression:  $\frac{2000}{256} \approx 8 : 1$



## The Problem:

Cannot fit 2000 numbers into 256 without loss!

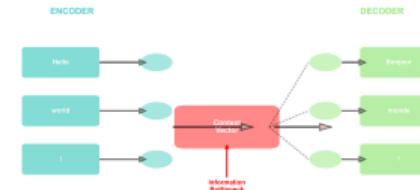
Mathematical Reality: Information theory limits how much we can compress without loss

# Challenge Summary: Understanding Translation Complexity

## What We Learned:

- Translation  $\neq$  word replacement
- Languages encode differently
- Need meaning understanding
- Must convert to numbers
- Fixed-size bottleneck problem

Sequence-to-Sequence Architecture



## The Challenge:

- Variable input length
- Fixed context size
- Information loss inevitable
- Longer = worse compression

### Key Question:

How do we capture all meaning in a fixed-size vector?

Next: The encoder-decoder architecture - a first solution to the translation challenge

## The Foundation: Sequence-to-Sequence Models

*Building Neural Networks That Translate*

## How humans translate (simplified):

### Phase 1: Understanding

1. Read entire source sentence
2. Extract complete meaning
3. Store in “mental representation”
4. Forget specific words
5. Keep abstract meaning

**Result:** Language-agnostic meaning

### Phase 2: Generation

1. Access stored meaning
2. Apply target grammar
3. Choose appropriate words
4. Generate word-by-word
5. Maintain coherence

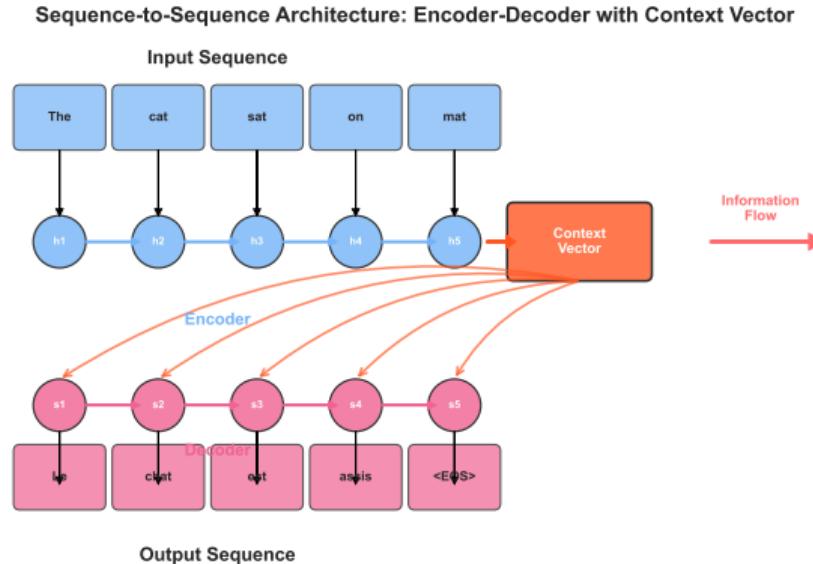
**Result:** Natural target sentence

**Neural Equivalent:** Encoder (understanding) + Decoder (generation) = Seq2Seq

---

Cognitive Model: Seq2Seq mimics human two-phase translation process

# The Encoder: Building Understanding Step-by-Step



## Encoder's Job:

- Process input sequentially
- Build hidden state (memory)
- Update with each word
- Final state = full understanding

## Processing “The cat sat”:

1.  $h_1 = \text{RNN}(\text{"The"}, h_0)$
2.  $h_2 = \text{RNN}(\text{"cat"}, h_1)$
3.  $h_3 = \text{RNN}(\text{"sat"}, h_2)$
4. Context:  $c = h_3$

# The Decoder: Generating from Understanding

## Decoder's Job:

- Start with context vector  $c$
- Generate one word at a time
- Use previous word + context
- Stop at end token

## Generation Process:

$$s_0 = c \text{ (initialize)} \quad (6)$$

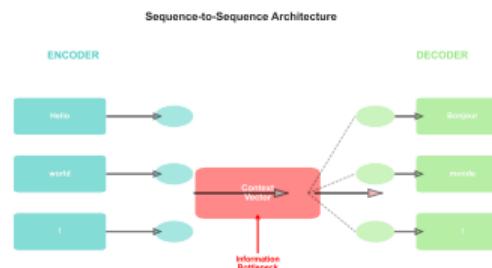
$$s_t = \text{RNN}(y_{t-1}, s_{t-1}) \quad (7)$$

$$P(y_t) = \text{softmax}(W s_t) \quad (8)$$

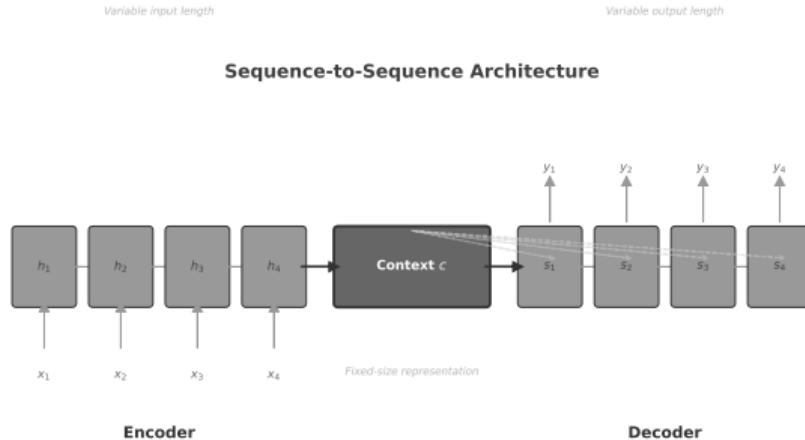
## Generating "Le chat noir":

1. Start:  $s_0 = c, y_0 = \text{iSTART}_i$
2. Generate "Le":  $P(y_1 | c)$
3. Generate "chat":  $P(y_2 | y_1, c)$
4. Generate "noir":  $P(y_3 | y_{1:2}, c)$
5. Stop:  $y_4 = \text{iEND}_i$

**Key:** Each word depends on context + history



# Complete Seq2Seq Architecture



## Components:

- Embedding layers
- Encoder RNN
- Context vector
- Decoder RNN

## Training:

- Teacher forcing
- Cross-entropy loss
- Backprop through time
- End-to-end learning

## Inference:

- Greedy decoding
- Beam search
- Length normalization
- Coverage penalty

# Complete Seq2Seq Implementation in PyTorch

```
1 import torch
2 import torch.nn as nn
3
4 class Seq2Seq(nn.Module):
5     def __init__(self, src_vocab,
6                  tgt_vocab, embed_dim=256,
7                  hidden_dim=512):
8         super().__init__()
9
9     # Embeddings
10    self.src_embed = nn.Embedding(
11        src_vocab, embed_dim
12    )
13    self.tgt_embed = nn.Embedding(
14        tgt_vocab, embed_dim
15    )
16
17    # Encoder & Decoder
18    self.encoder = nn.LSTM(
19        embed_dim, hidden_dim,
20        batch_first=True
21    )
22    self.decoder = nn.LSTM(
23        embed_dim, hidden_dim,
24        batch_first=True
25    )
26
27    # Output projection
28    self.output = nn.Linear(
29        hidden_dim, tgt_vocab
```

```
1 def forward(self, src, tgt):
2     # Encode
3     src_emb = self.src_embed(src)
4     _, (h, c) = self.encoder(
5         src_emb
6     )
6
7     # Decode
8     tgt_emb = self.tgt_embed(tgt)
9     out, _ = self.decoder(
10        tgt_emb, (h, c)
11    )
12
13     # Project
14     logits = self.output(out)
15
16     return logits
17
18
19 # Usage
20 model = Seq2Seq(
21     src_vocab=10000,
22     tgt_vocab=10000
23 )
24
25 # Training step
26 src = torch.randint(0, 10000,
27                     (32, 20))
28 tgt = torch.randint(0, 10000,
29                     (32, 15))
30 logits = model(src, tgt)
```

# Encoding Example: “The black cat sat”

Watch the hidden state evolve:

Step 1: Process “The”

- Input:  $x_1 = \text{embed}(\text{"The"}) = [0.1, 0.3, \dots]$
- Hidden:  $h_1 = \text{LSTM}(x_1, h_0)$
- Memory: “Determiner seen”

Step 2: Process “black”

- Input:  $x_2 = \text{embed}(\text{"black"})$
- Hidden:  $h_2 = \text{LSTM}(x_2, h_1)$
- Memory: “Determiner + adjective”

Step 3: Process “cat”

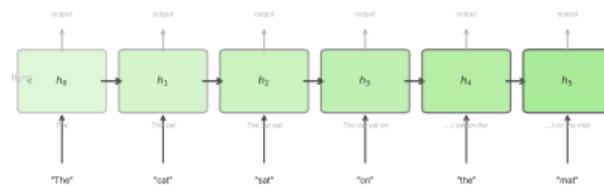
- Input:  $x_3 = \text{embed}(\text{"cat"})$
- Hidden:  $h_3 = \text{LSTM}(x_3, h_2)$
- Memory: “Black cat (subject)”

Step 4: Process “sat”

- Input:  $x_4 = \text{embed}(\text{"sat"})$
- Hidden:  $h_4 = \text{LSTM}(x_4, h_3)$
- Memory: “Black cat sat (complete)”

Encoding Process: Each word updates understanding, final state has complete meaning

Hidden State Evolution: Building Context



Final Context:

$c = h_4$  contains: - Subject: black cat - Action:  
sat - Tense: past

# Decoding Example: Generating “Le chat noir”

Starting from context  $c$ :

Step 1: Generate “Le”

- State:  $s_0 = c$
- Input:  $\text{START}_{\text{L}}$  token
- Output:  $P(\text{"Le"}) = 0.8$
- Next:  $s_1 = \text{LSTM}(\text{"Le"}, s_0)$

Step 2: Generate “chat”

- State:  $s_1$  (knows “Le”)
- Input: “Le”
- Output:  $P(\text{"chat"}) = 0.7$
- Next:  $s_2 = \text{LSTM}(\text{"chat"}, s_1)$

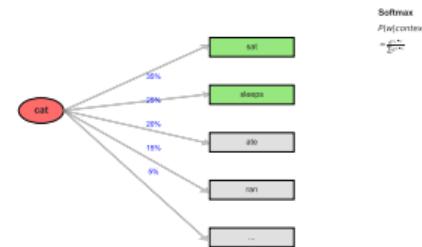
Step 3: Generate “noir”

- State:  $s_2$  (knows “Le chat”)
- Input: “chat”
- Output:  $P(\text{"noir"}) = 0.6$
- Next:  $s_3 = \text{LSTM}(\text{"noir"}, s_2)$

Probability Distribution:

At each step, model outputs:

Word Prediction Probabilities from Context



Key Point:

Decoder maintains its own hidden state separate from encoder

Decoding Process: Conditional generation using context and previous outputs

# Quiz Checkpoint: Understanding Seq2Seq

## Questions:

**Q1:** What is the context vector?

- a) Average of word embeddings
- b) Final encoder hidden state
- c) Sum of all hidden states
- d) Random initialization

**Q2:** Why use two separate networks?

- a) Faster training
- b) Different tasks (read vs write)
- c) More parameters
- d) Requirement of RNNs

**Q3:** Teacher forcing means:

- a) Using true targets during training
- b) Forcing convergence
- c) Teaching the teacher

## Answers:

**A1: b) Final encoder hidden state**

- Contains full sentence understanding
- Fixed-size representation
- Passed to decoder

**A2: b) Different tasks**

- Encoder: comprehension
- Decoder: generation
- Different objectives

**A3: a) Using true targets**

- Feed correct previous word
- Speeds up training
- Avoids error accumulation

# Foundation Summary: The Seq2Seq Architecture

## Architecture Components:

- Encoder RNN: reads input
- Context vector: compressed meaning
- Decoder RNN: generates output
- End-to-end training

## Key Equations:

$$c = \text{Encoder}(x_{1:n}) \quad (9)$$

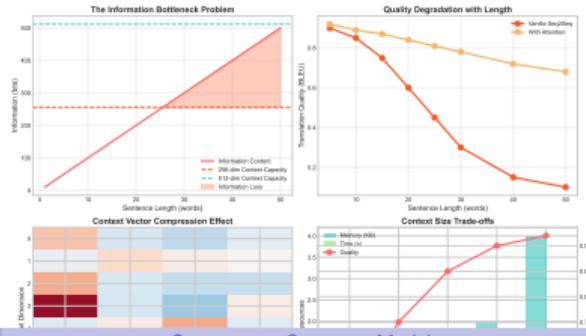
$$y_t = \text{Decoder}(c, y_{<t}) \quad (10)$$

## Strengths:

- Variable input/output length
- End-to-end learning
- No alignment needed
- Works for any language pair

## Weakness:

Fixed-size bottleneck!



## Test Your Understanding

### Quick Quiz:

**Question 1:** What's the key limitation of basic seq2seq?

- A) Too slow to train
- B) Information bottleneck
- C) Can't handle grammar
- D) Requires parallel data

**Question 2:** What does the encoder produce?

- A) Word translations
- B) Grammar rules
- C) Fixed-size context vector
- D) Attention weights

### Answers:

**Answer 1:** B - Information bottleneck

- All source information compressed into single vector
- Loses details from long sequences
- Can't selectively access parts

**Answer 2:** C - Fixed-size context vector

- Final hidden state  $h_n$
- Summarizes entire input
- Same size regardless of input length

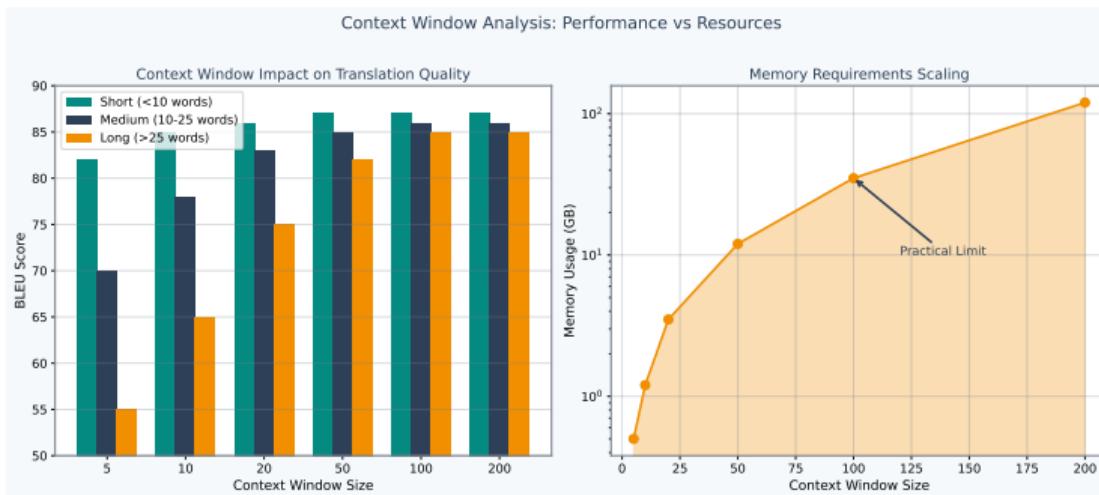
---

**Key Insight:** The bottleneck problem motivates the need for attention mechanism

## The Breakthrough: Attention Mechanism

*Teaching Networks Where to Look*

# The Bottleneck Problem: Why Seq2Seq Fails on Long Sentences



## Performance Degradation:

- 10 words: BLEU = 35
- 20 words: BLEU = 25
- 30 words: BLEU = 15
- 40+ words: BLEU ↓ 10

## What's Lost:

- Early words forgotten
- Specific details blurred
- Word positions unclear
- Grammatical structure

# Human Translation: The Attention Analogy

How do humans really translate long sentences?

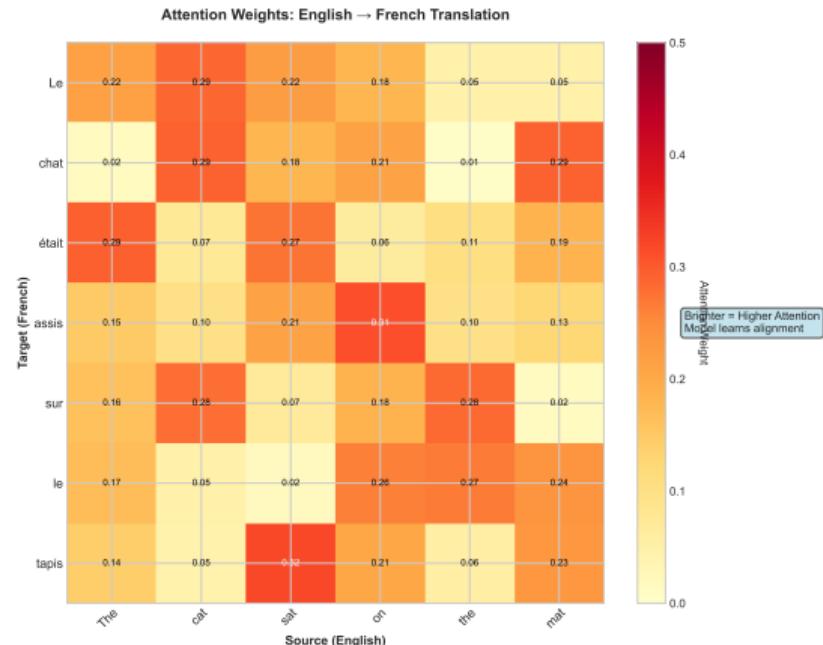
Translating Word by Word:

"The black cat that I saw yesterday sat"

When translating "sat":

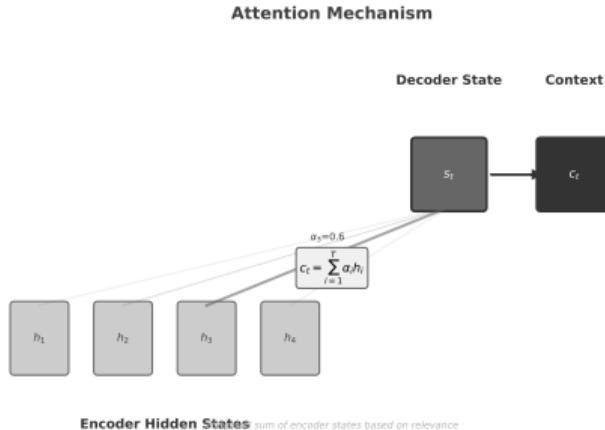
1. Look back at "cat" (subject)
2. Check tense markers
3. Verify agreement
4. Generate appropriate form

**Key:** We don't memorize everything! We look back as needed.



Attention Idea:

# The Attention Mechanism: Dynamic Context



## Old Way (Seq2Seq):

- Fixed context  $c = h_n$
- Same for all decoder steps
- Information bottleneck
- Forgets early words

## New Way (Attention):

- Dynamic context  $c_t$
- Different for each word
- Weighted sum of all states
- Remembers everything

## How to calculate attention weights:

**Step 1: Score** How relevant is each encoder state?

$$e_{ti} = \text{score}(s_{t-1}, h_i)$$

**Step 2: Normalize** Convert scores to probabilities:

$$\alpha_{ti} = \frac{\exp(e_{ti})}{\sum_j \exp(e_{tj})}$$

Common scoring functions:

- **Dot:**  $s_{t-1} \cdot h_i$
- **General:**  $s_{t-1} W h_i$
- **Concat:**  $v \tanh(W[s_{t-1}; h_i])$

**Step 3: Weighted Sum** Compute dynamic context:

$$c_t = \sum_{i=1}^n \alpha_{ti} h_i$$

Each decoder step gets its own custom-weighted view of the source!

---

Mathematical Core: Three simple steps that revolutionized machine translation

## Understanding the QKV Framework:

### Components:

- **Query** ( $s_{t-1}$ ): What I'm looking for
- **Keys** ( $h_i$ ): What's available
- **Values** ( $h_i$ ): What to retrieve
- **Weights** ( $\alpha_{ti}$ ): Relevance scores

### Analogy:

- Query = Search term
- Keys = Document titles
- Values = Document content
- Attention = Search relevance

### Critical Insight:

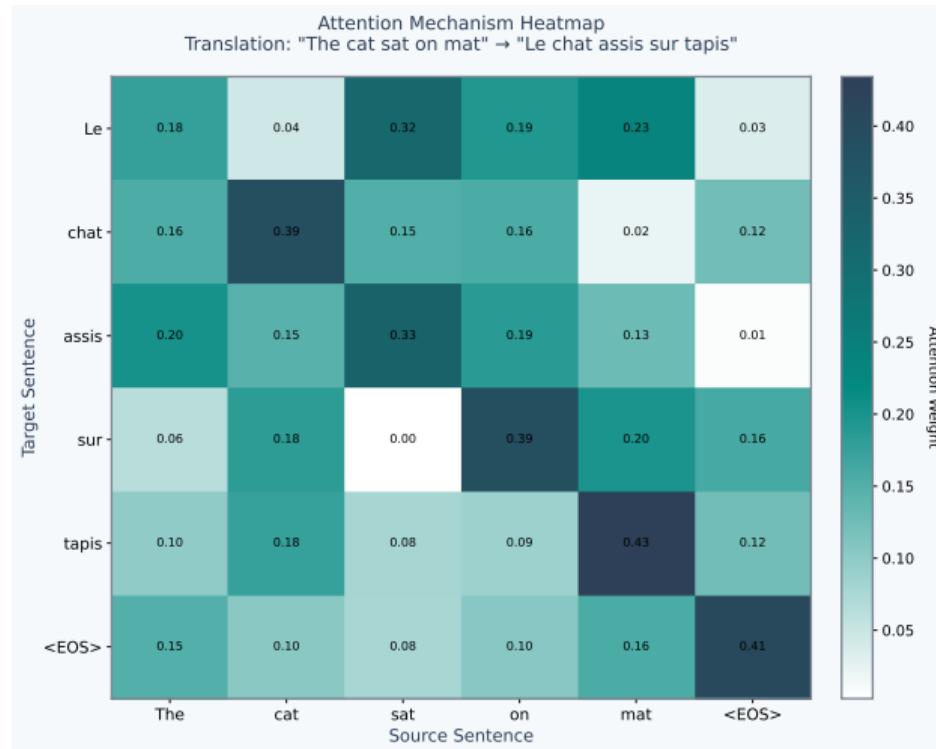
This simple mechanism  
is the foundation of  
**ALL** modern transformers!

GPT, BERT, T5, ChatGPT...  
all use this QKV attention

---

Foundation: Attention as learned information retrieval - the key to modern AI

# Visualizing Attention: What the Model Focuses On



## Reading the Heatmap:

- Rows: Target (French)

## Example Weights:

- "chat" → "cat" (0.8)

# Implementing Attention: Score Calculation

```
1 def attention_score(decoder_hidden, encoder_outputs):
2     # Step 1: Expand decoder hidden to match encoder length
3     seq_len = encoder_outputs.size(1)
4     hidden = decoder_hidden.repeat(1, seq_len, 1)
5
6     # Step 2: Concatenate and score
7     concat = torch.cat((hidden, encoder_outputs), dim=2)
8     scores = self.v(torch.tanh(self.attn(concat)))
9
10    # Step 3: Apply softmax to get weights
11    weights = F.softmax(scores.squeeze(2), dim=1)
12    return weights
```

---

**Key:** Score function determines which encoder states to focus on

# Implementing Attention: Context Computation

```
1 def compute_context(weights, encoder_outputs):
2     # Weighted sum of encoder states
3     # weights: [batch, seq_len]
4     # encoder_outputs: [batch, seq_len, hidden]
5     context = torch.bmm(
6         weights.unsqueeze(1),  # [batch, 1, seq_len]
7         encoder_outputs        # [batch, seq_len, hidden]
8     )  # Result: [batch, 1, hidden]
9     return context
```

---

Context vector: Weighted combination of all encoder hidden states

# Implementing Attention: Decoder Integration

```
1  class AttentionDecoder(nn.Module):
2      def forward(self, input_token, hidden, encoder_outputs):
3          # 1. Embed input token
4          embedded = self.embedding(input_token)
5
6          # 2. Compute attention
7          context, weights = self.attention(hidden, encoder_outputs)
8
9          # 3. Combine embedding + context
10         lstm_input = torch.cat([embedded, context], dim=2)
11
12         # 4. LSTM step
13         output, hidden = self.lstm(lstm_input, hidden)
14
15         # 5. Predict next word
16         predictions = self.output(output)
17
18         return predictions, hidden, weights
```

---

Each decoder step: Attention determines what to focus on from source

# Interactive Exercise: Calculate Attention Weights

**Task:** Compute attention for generating “noir” (black)

Given decoder state  $s_2$  after generating “Le chat”:

**Encoder states:**

- $h_1$ : “The” = [0.1, 0.2]
- $h_2$ : “black” = [0.8, 0.9]
- $h_3$ : “cat” = [0.5, 0.4]
- $h_4$ : “sat” = [0.3, 0.1]

**Decoder query:**

- $s_2$  = [0.7, 0.8]

**Your calculations:**

1. Scores (dot product):

- $e_1 = s_2 \cdot h_1 = \underline{\hspace{2cm}}$
- $e_2 = s_2 \cdot h_2 = \underline{\hspace{2cm}}$
- $e_3 = s_2 \cdot h_3 = \underline{\hspace{2cm}}$
- $e_4 = s_2 \cdot h_4 = \underline{\hspace{2cm}}$

2. Softmax weights:

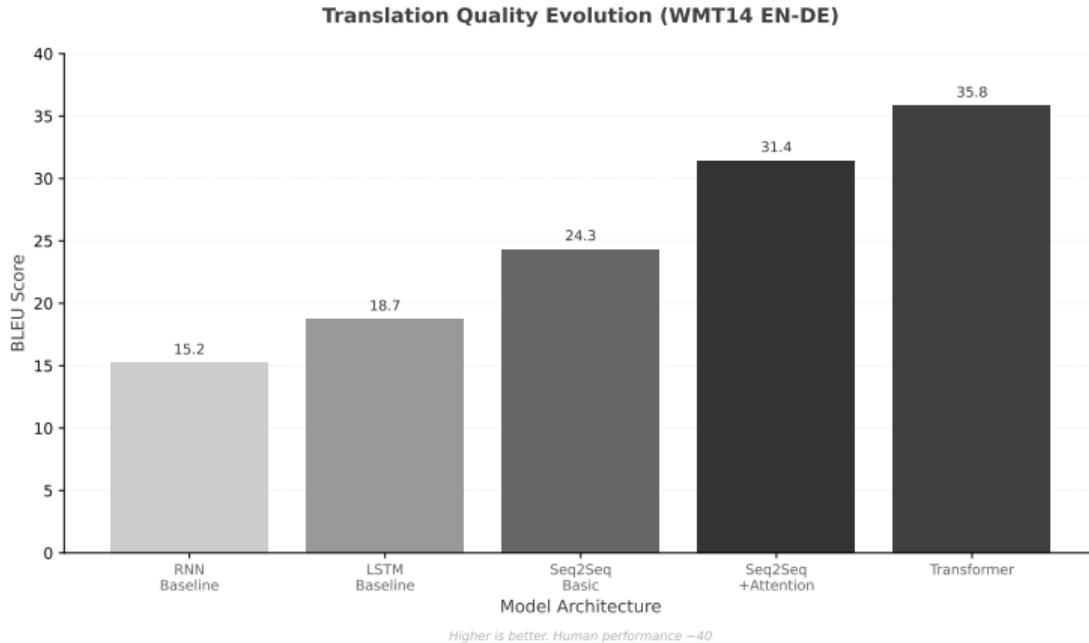
- $\alpha_2 = \underline{\hspace{2cm}}$  (highest!)

3. Context: weighted sum

---

**Hands-On:** Computing attention manually builds intuition for the mechanism

# Impact of Attention: BLEU Score Improvements



## BLEU Score Improvements by Sentence Length:

- Short ( $\leq 10$  words): +5 points
- Medium (10-20): +10 points
- Long (20-30): +15 points
- Very long ( $> 30$ ): +20 points

## Technical Advantages:

- No information bottleneck
- Direct access to all source words
- Handles word reordering naturally
- Resolves lexical ambiguity
- Maintains word alignment

## Practical Impact:

- Production-ready quality
- Handles complex languages
- Works for long documents
- Interpretable alignments
- Foundation for transformers

**Game-changing improvement that enabled modern NMT**

---

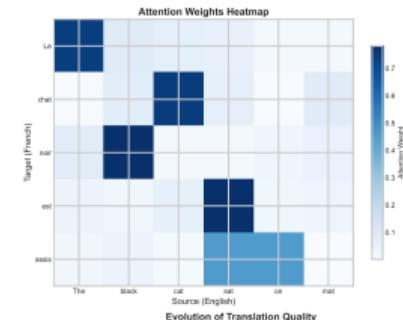
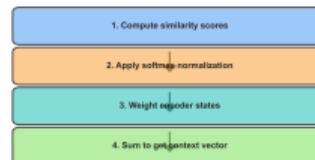
**Historical Impact:** This innovation directly led to transformer architecture

# Breakthrough Summary: How Attention Changed Everything

## The Innovation:

- Dynamic context vectors
- Look at all encoder states
- Weighted by relevance
- Different for each word

Attention Mechanism Steps

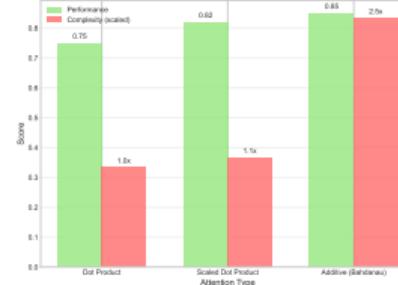


## Mathematical Core:

$$\alpha_{ti} = \text{softmax}(\text{score}(s_t, h_i)) \quad (11)$$

$$c_t = \sum_i \alpha_{ti} h_i \quad (12)$$

Attention Mechanisms Comparison



## Impact:

Attention mechanism became foundation of all modern NLP

Historical Significance: Attention paper (2014) revolutionized entire field

## Test Your Understanding of Attention Mechanism

### Quick Quiz:

**Question 1:** What does attention compute at each decoder step?

- A) Next word directly
- B) Weighted sum of encoder states
- C) Grammar rules
- D) Translation dictionary

**Question 2:** Why does attention help with long sentences?

- A) Processes faster
- B) Uses less memory
- C) Direct access to all source words
- D) Better vocabulary

**Key Insight:** Attention allows selective focus on relevant parts of the input

### Answers:

**Answer 1:** B - Weighted sum of encoder states

- $c_t = \sum_i \alpha_{ti} h_i$
- Weights  $\alpha_{ti}$  from softmax
- Dynamic for each output word

**Answer 2:** C - Direct access to all source words

- No information bottleneck
- Can "look back" at any position
- Preserves long-range dependencies

## The Impact: Modern Translation Systems

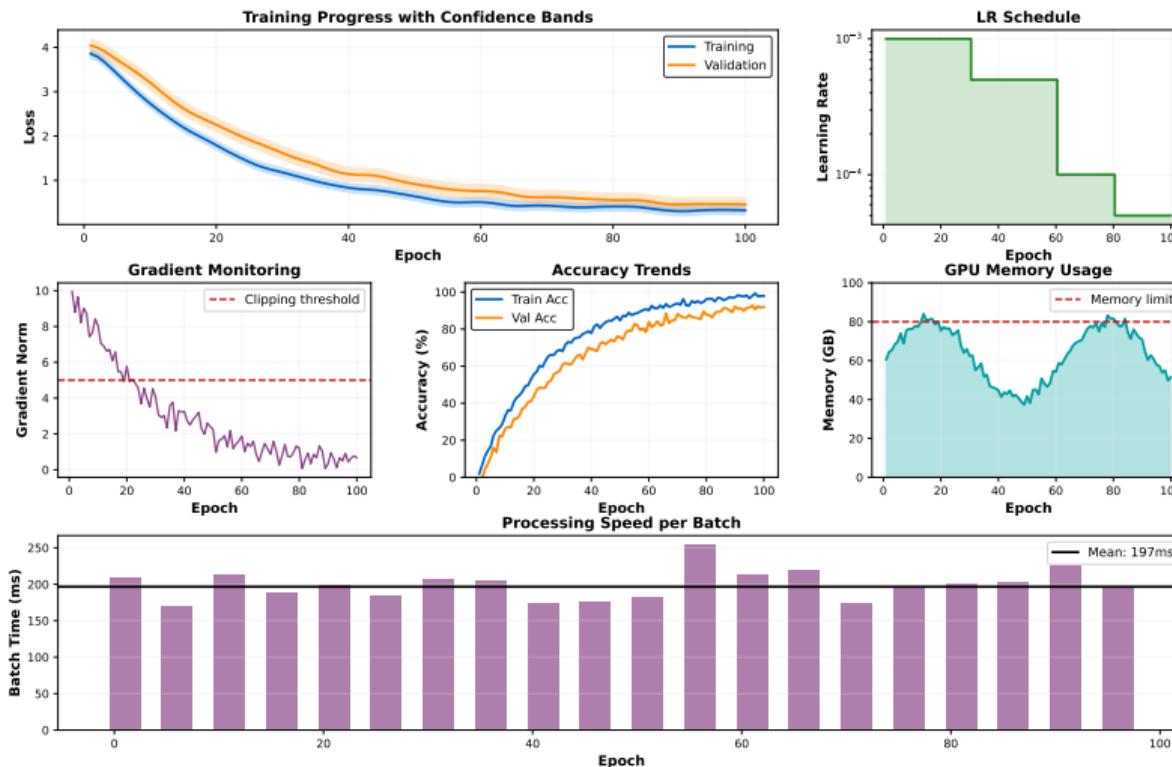
*From Research to Billion-User Products*

# Complete Seq2Seq with Attention

```
1 class AttentionSeq2Seq(nn.Module):
2     def __init__(self, src_vocab,
3                  tgt_vocab, dim=512):
4         super().__init__()
5
5         # Components
6         self.encoder = Encoder(
7             src_vocab, dim
8         )
9         self.decoder = DecoderWithAttn(
10            tgt_vocab, dim
11        )
12         self.attention = Attention(dim)
13
14     def forward(self, src, tgt,
15                teacher_forcing=0.5):
16         # Encode all at once
17         enc_out, (h, c) = self.encoder(src)
18
19         batch = src.size(0)
20         max_len = tgt.size(1)
21         vocab = self.decoder.vocab_size
22
23         # Store outputs
24         outputs = torch.zeros(
25             batch, max_len, vocab
26         )
27
28         # First input
29         input = tgt[:, 0]
30
31         for t in range(1, max_len):
32             # Attention context
33             context, weights =
34                 self.attention(
35                     h, enc_out
36                 )
37
38             # Decode one step
39             output, (h, c) =
40                 self.decoder(
41                     input, (h, c),
42                     context
43                 )
44
45             outputs[:, t] = output
46
47             # Teacher forcing
48             use_teacher = random.random()
49                 < teacher_forcing
50
51             if use_teacher:
52                 input = tgt[:, t]
53             else:
54                 input = output.argmax(1)
55
56         return outputs
```

# Training Dynamics: Loss Curves

Training Monitoring Dashboard



## Early (0-20k steps):

- Basic vocabulary
- Word copying
- Common phrases
- Simple alignment

## Middle (20k-80k):

- Grammar rules
- Word reordering
- Multi-word expressions
- Context sensitivity

## Late (80k+):

- Rare words
- Idioms
- Style transfer
- Long dependencies

**Key:** Attention alignment emerges without explicit supervision

---

**Training Strategy:** Patient training (100k+ steps) crucial for quality

## Hyperparameters That Matter:

- **Learning rate:** Start 0.001, decay after epoch 10
- **Teacher forcing:** 0.5 → 0 over training
- **Gradient clip:** Essential (1.0 works well)
- **Batch size:** 32-64 optimal for GPU
- **Hidden size:** 512 is sweet spot
- **Layers:** 2-3 LSTM layers sufficient

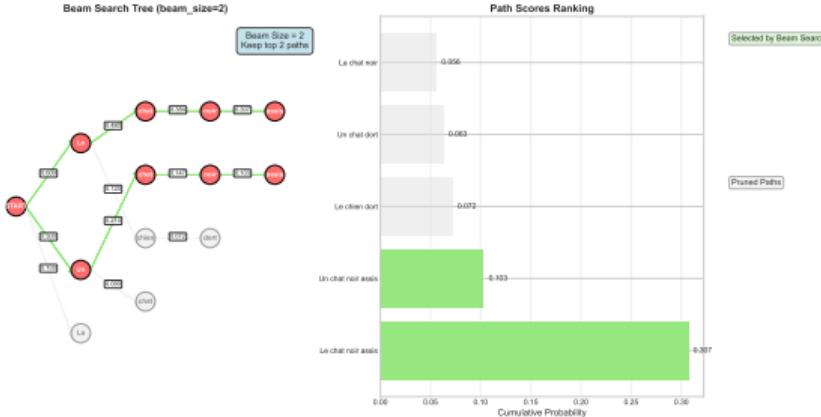
## Common Issues & Solutions:

- **Exploding loss:** → Reduce learning rate
- **Mode collapse:** → Add dropout (0.3)
- **Poor rare words:** → Increase min frequency
- **Slow training:** → Use GPU, reduce batch size
- **Overfitting:** → More data, regularization
- **Underfitting:** → Bigger model, longer training

---

**Pro Tip:** Start simple, add complexity gradually, monitor validation metrics closely

# Beam Search: Better Decoding Strategy



## Greedy vs Beam:

- Greedy: Best at each step
- Beam: Keep top-k paths
- Better final result

## Beam size:

- $k=5$ : Good balance
- $k=10$ : Slightly better

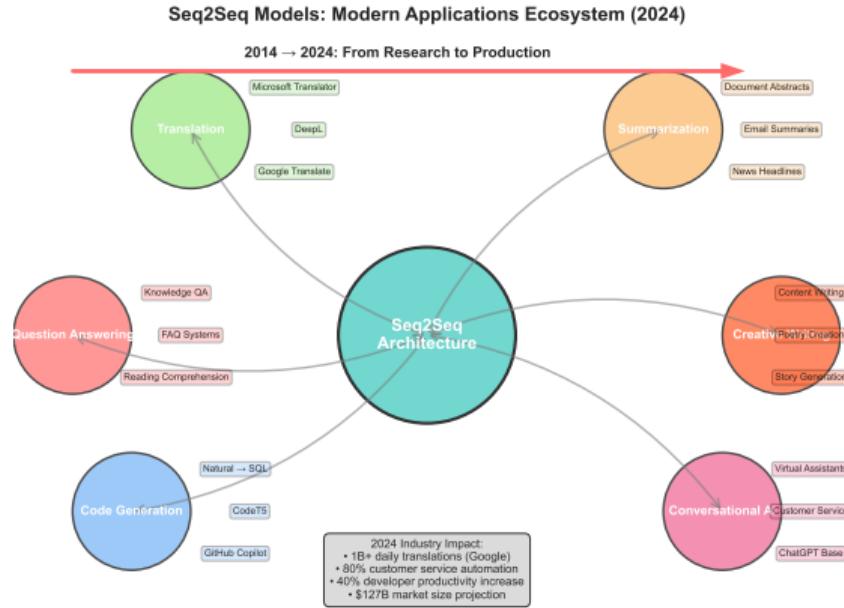
## Example (beam=3):

- Step 1: "Le", "Un", "Les"
- Step 2: "Le chat", "Un chat"
- Step 3: Keep expanding top-3

Pick best complete path

Beam search explores multiple hypotheses for better translations

# Modern Applications: Direct Descendants (2024)



## Production Systems:

- Google Translate (1B+ users)
- DeepL (quality leader)
- Facebook M2M-100

## Capabilities:

- 100+ language pairs
- Document translation
- Real-time speech

## Attention Everywhere:

- ChatGPT/Claude (attention-based)
- Image captioning
- Video understanding
- Code generation
- Music composition
- Speech recognition
- Medical diagnosis

## Foundation Truth:

Seq2Seq + Attention  
=  
Modern AI Backbone

Every modern language model  
builds on this architecture

---

Attention mechanism is the foundation of the entire AI revolution

# Lab Preview: Build Your Own Translator

## Week 4 Lab: English-French Neural Machine Translation

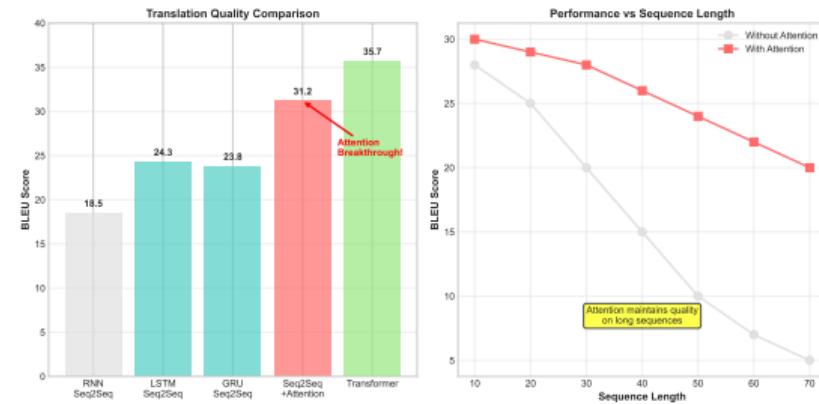
### What You'll Build:

1. Load parallel corpus
2. Tokenize and preprocess
3. Implement encoder-decoder
4. Add attention mechanism
5. Train on GPU
6. Visualize attention weights
7. Compare with/without attention

### Dataset:

- 10,000 sentence pairs
- English → French
- Average 15 words/sentence

### Expected Results:



### Bonus Challenges:

- Multi-head attention
- Bidirectional encoder
- Coverage mechanism
- Back-translation

# Interactive Debugging: Common Training Issues

Your model isn't learning. Debug these issues:

## Issue 1: Attention all uniform

Symptoms:

- All weights  $\approx 1/n$
- Poor translation quality
- Not improving

Your fix: \_\_\_\_\_

Hint: Check score function

Common Fixes:

### Fix 1: Initialize properly

- Use Xavier initialization
- Scale attention scores
- Add small epsilon to softmax

## Issue 2: Mode collapse

Symptoms:

- Always generates "the the the"
- Loss plateaus high

### Fix 2: Teacher forcing

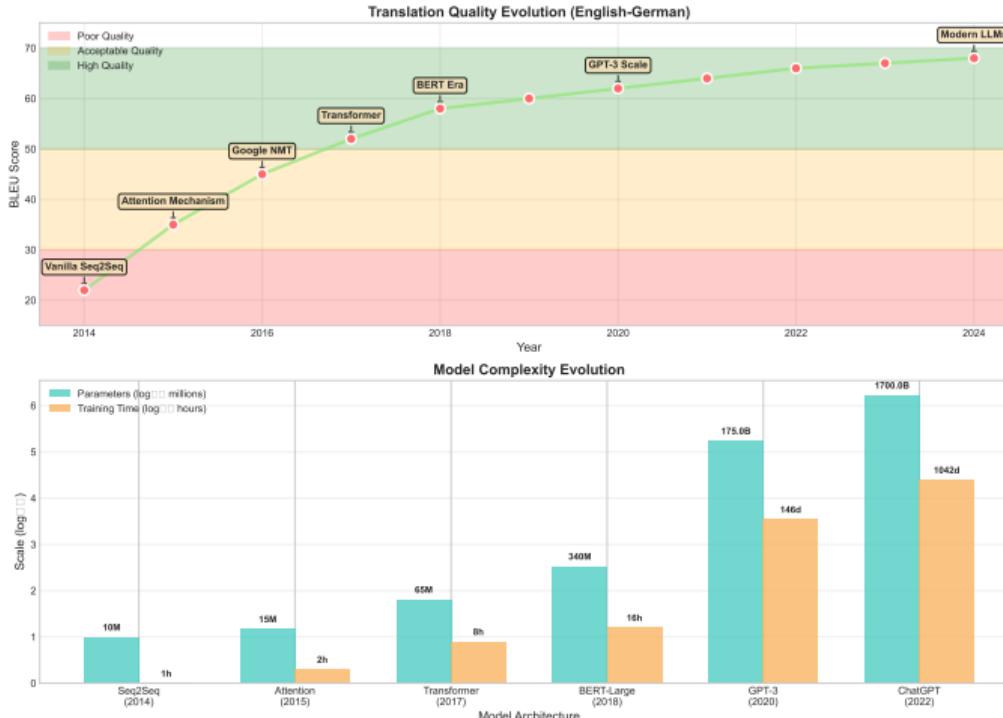
- Start with 100% teacher forcing
- Gradually reduce ratio
- Scheduled sampling

Debug systematically!

---

Debugging Skills: Most issues come from initialization or training schedule

# Performance Evolution Timeline



## Three Years That Changed Everything

The journey from seq2seq to attention represents  
Sequence-to-Sequence Models

# Performance Metrics by Year

## 2014: Seq2Seq

- BLEU: 20-25
- Simple, elegant
- Length problems
- 2-3 days training

### Impact:

- First neural MT
- Proof of concept
- Beat phrase-based

## 2015: + Attention

- BLEU: 30-35
- Handles length
- Interpretable
- 4-5 days training

### Impact:

- Production ready
- Google adoption
- 40% improvement

## 2017: Transformer

- BLEU: 40+
- All attention
- Parallel training
- 12 hours training!

### Impact:

- New paradigm
- Enables GPT/BERT
- 10x faster training

**Key Insight:** Each innovation built on the previous - attention was THE breakthrough

---

Historical Progression: From RNN to attention to transformer architecture

# The Bridge to Transformers (Week 5 Preview)

## From Seq2Seq+Attention to Transformers:

### What We Keep:

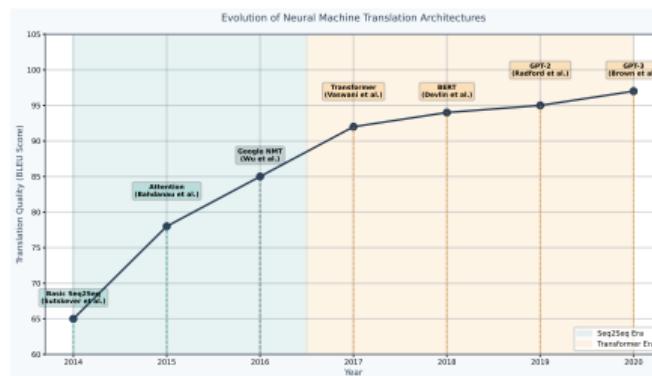
- Attention mechanism
- Query-Key-Value
- Position awareness
- Encoder-decoder structure

### What We Add:

- Self-attention
- Multi-head attention
- Position encodings
- Layer normalization
- Parallel processing

### What We Remove:

- RNN/LSTM cells
- Sequential processing
- Recurrent connections
- Hidden state passing



## The Challenge

- Translation  $\neq$  word replacement
- Need meaning understanding
- Information bottleneck problem

## The Foundation

- Encoder-decoder architecture
- Fixed context vector
- Works but limited by bottleneck

## The Breakthrough

- Dynamic context vectors
- Look at all encoder states
- Massive performance improvement

## The Impact

- Complete implementation
- Beam search decoding
- Powers modern translation
- Foundation for transformers

## Key Takeaways:

1. Context vectors compress meaning
2. Attention removes bottleneck
3. Foundation of modern NLP
4. Bridge to transformers

**Next Week: Transformers - Attention Without RNNs!**

Achievement Unlocked: You understand the foundation of all modern language AI!