

Sequence-to-Sequence Models

Week 4: The Translation Revolution with Attention

NLP Course 2025

Professional Template Edition

September 29, 2025

Week 4: Journey Through Translation and Attention

Learning Path: From word-by-word replacement to neural translation. Master encoder-decoder architectures, understand the bottleneck problem, and discover how attention revolutionized machine translation.

Part 1: Translation Challenge & Motivation

Why Word-by-Word Translation Fails

The Google Translate Evolution: A Success Story

2006: Statistical MT

- Word/phrase dictionaries
- Counted co-occurrences
- “Reasonable” translations
- Often awkward phrasing

2016: Neural MT Launch

- Seq2Seq with attention
- Human-quality for some pairs
- 60% error reduction
- Revolutionary improvement

Real Example:

Chinese Input: “There is one cat in station”

Old: “In the station is one cat”

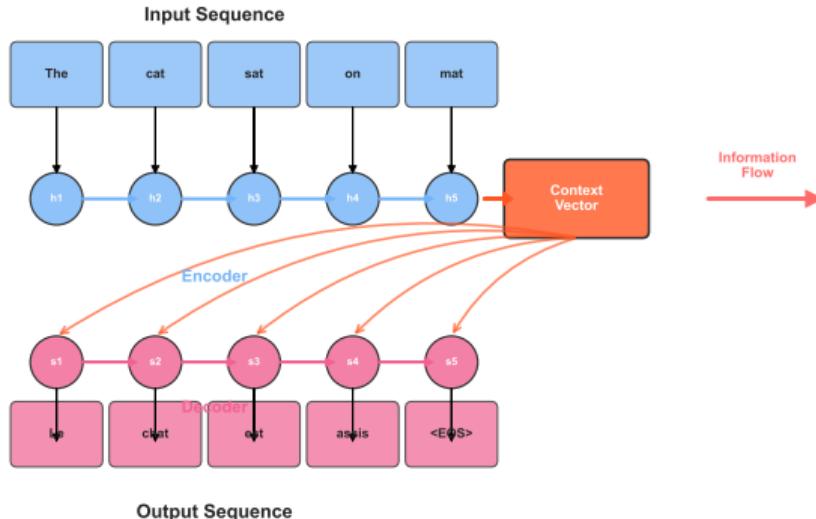
New: “There is a cat at the station”

What changed? Understanding context, not just words

Historical Context: Neural MT reduced translation errors by 60% overnight - the biggest leap in MT history

The Fundamental Problem: Meaning Across Languages

Sequence-to-Sequence Architecture: Encoder-Decoder with Context Vector



Translation is NOT:

- Word replacement
- Dictionary lookup
- Rule application
- Rule application

Translation IS:

- Understanding meaning
- Cultural context
- Reformulation
- Generation

Why Word-by-Word Translation Fails: Concrete Examples

Problem 1: Word Order

- English: "I saw the red house"
- Spanish: "Vi la casa roja"
- Literal: "Saw-I the house red"

Problem 2: Idioms

- English: "It's raining cats and dogs"
- French: "Il pleut des cordes"
- Literal: "It rains ropes"

Problem 3: Context

- "Bank" → "Banque" (financial)
- "Bank" → "Rive" (river)
- Need full sentence to decide

Problem 4: Grammar

- German: Verb at end
- Japanese: Subject optional
- Chinese: No tenses

Conclusion: Languages encode meaning differently - translation needs deep understanding

Language Diversity: Each language has unique ways of expressing ideas

Converting Meaning to Numbers: The Core Challenge

Computers only understand numbers, so:

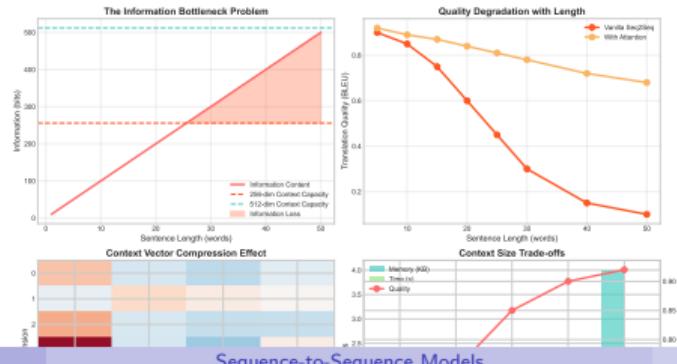
"The cat sat on the mat" → [Numbers] → "Le chat s'est assis sur le tapis"

Step 1: Words to Vectors

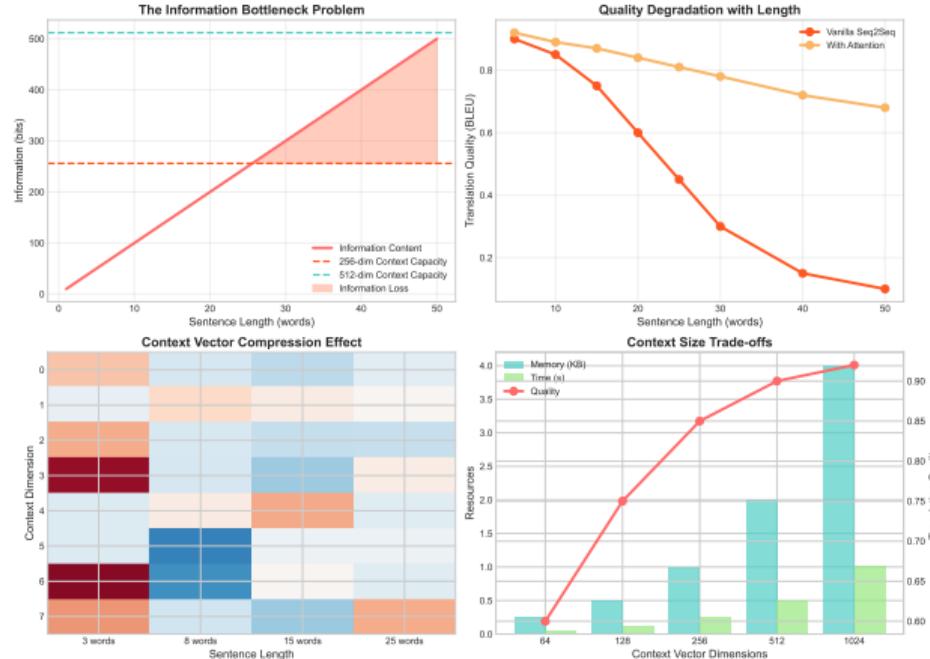
- "cat" → [0.2, -0.5, 0.8, ...]
- 100-300 dimensional vectors
- Learned from context (Word2Vec)
- Similar words = nearby vectors

Step 2: Sentence to Vector

- Combine word vectors
- Build "context vector"
- Fixed size (e.g., 256 dims)
- Must capture ALL meaning



The Compression Challenge: Information Bottleneck



Compression Ratios:

- 5 words: 500 dims \rightarrow 256 (2:1)
- 20 words: 2000 dims \rightarrow 256 (8:1)

What Gets Lost?

- Specific word choices
- Grammatical nuances

Interactive Exercise: Manual Translation Steps

Task: Translate “The black cat sat” to French step-by-step

Your Steps:

1. Read entire English sentence
2. Identify: subject (cat), verb (sat)
3. Recall French words:
 - cat → chat
 - black → noir
 - sat → s'est assis
4. Apply French grammar:
 - Article-Noun-Adjective order
 - Gender agreement (le/la)
5. Generate: “Le chat noir s'est assis”

What You Actually Did:

1. Encoded English to meaning
2. Stored meaning in memory
3. Decoded meaning to French

This is exactly Seq2Seq!

Key Observation:

You didn't translate word-by-word! You understood first, then generated.

Human Insight: We naturally use encoder-decoder approach when translating

Calculating the Bottleneck: A Mathematical Perspective

Information Content:

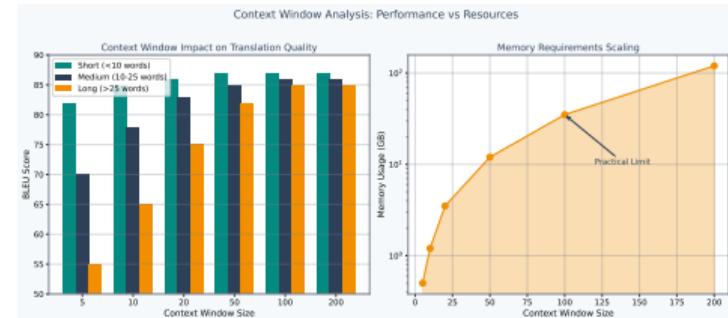
$$\text{Input} = n \times d_{\text{embed}}$$

$$\text{Context} = d_{\text{hidden}}$$

$$\text{Ratio} = \frac{n \times d_{\text{embed}}}{d_{\text{hidden}}} \quad (3)$$

Example Calculation:

- 20 words, 100-dim embeddings
- Input: $20 \times 100 = 2000$ values
- Context: 256 values
- Compression: $\frac{2000}{256} \approx 8 : 1$



The Problem:

Cannot fit 2000 numbers into 256 without loss!

Mathematical Reality: Information theory limits how much we can compress without loss

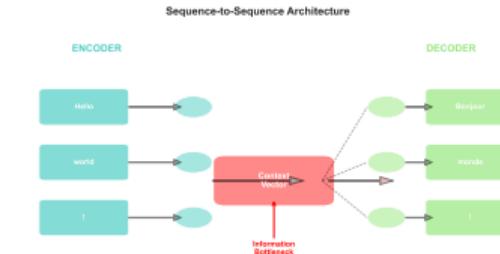
Part 1 Summary: Understanding the Challenge

What We Learned:

- Translation \neq word replacement
- Languages encode differently
- Need meaning understanding
- Must convert to numbers
- Fixed-size bottleneck problem

The Challenge:

- Variable input length
- Fixed context size
- Information loss inevitable
- Longer = worse compression



Key Question:

How do we capture all meaning in a fixed-size vector?

Next: The encoder-decoder architecture - a first solution to the translation challenge

Part 2: Encoder-Decoder Architecture

Building Understanding, Then Generating

How humans translate (simplified):

Phase 1: Understanding

1. Read entire source sentence
2. Extract complete meaning
3. Store in “mental representation”
4. Forget specific words
5. Keep abstract meaning

Result: Language-agnostic meaning

Phase 2: Generation

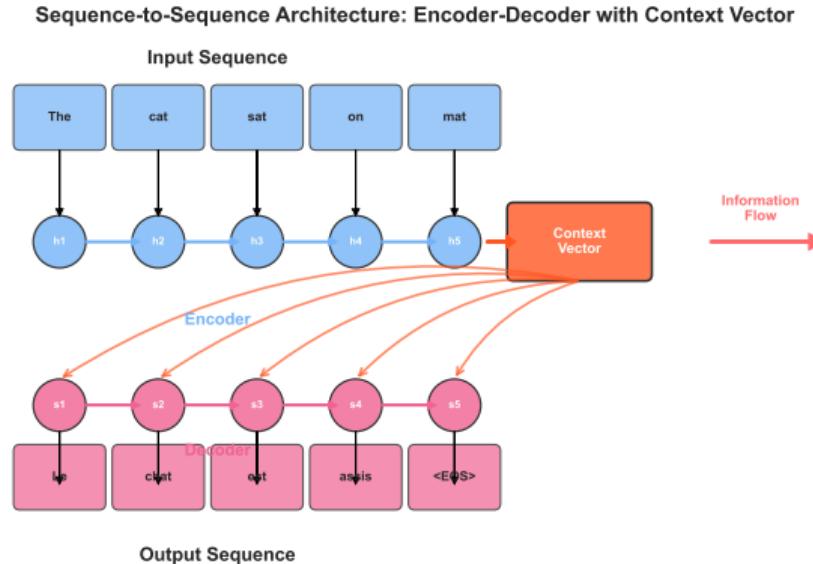
1. Access stored meaning
2. Apply target grammar
3. Choose appropriate words
4. Generate word-by-word
5. Maintain coherence

Result: Natural target sentence

Neural Equivalent: Encoder (understanding) + Decoder (generation) = Seq2Seq

Cognitive Model: Seq2Seq mimics human two-phase translation process

The Encoder: Building Understanding Step-by-Step



Encoder's Job:

- Process input sequentially
- Build hidden state (memory)
- Update with each word
- Final state = full understanding

Processing “The cat sat”:

1. $h_1 = \text{RNN}(\text{"The"}, h_0)$
2. $h_2 = \text{RNN}(\text{"cat"}, h_1)$
3. $h_3 = \text{RNN}(\text{"sat"}, h_2)$
4. Context: $c = h_3$

The Decoder: Generating from Understanding

Decoder's Job:

- Start with context vector c
- Generate one word at a time
- Use previous word + context
- Stop at end token

Generation Process:

$$s_0 = c \text{ (initialize)} \quad (6)$$

$$s_t = \text{RNN}(y_{t-1}, s_{t-1}) \quad (7)$$

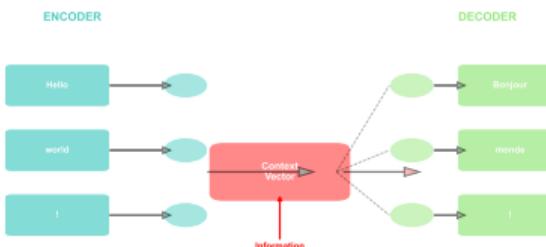
$$P(y_t) = \text{softmax}(W s_t) \quad (8)$$

Generating "Le chat noir":

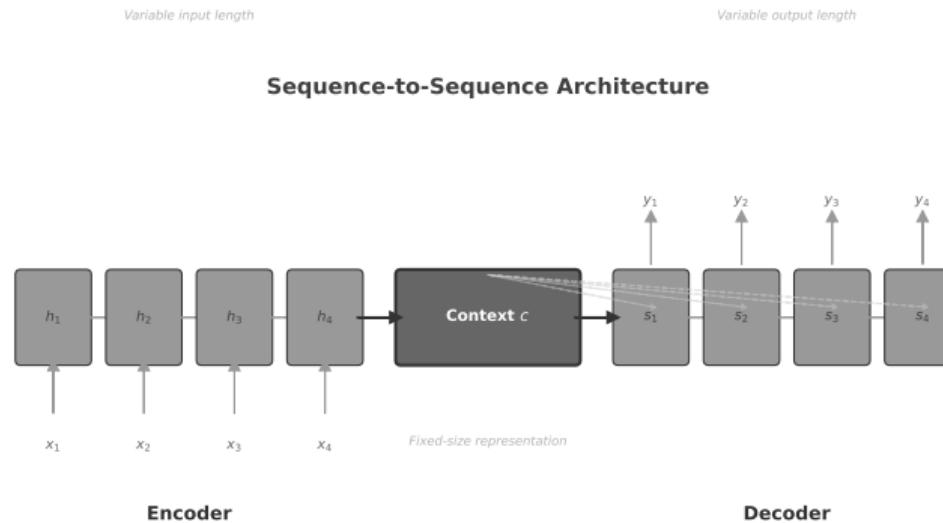
1. Start: $s_0 = c, y_0 = \text{iSTART}_i$
2. Generate "Le": $P(y_1 | c)$
3. Generate "chat": $P(y_2 | y_1, c)$
4. Generate "noir": $P(y_3 | y_{1:2}, c)$
5. Stop: $y_4 = \text{iEND}_i$

Key: Each word depends on context + history

Sequence-to-Sequence Architecture



Complete Seq2Seq Architecture



Components:

- Embedding layers
- Encoder RNN

Training:

- Teacher forcing
- Cross-entropy loss

Inference:

- Greedy decoding
- Beam search

Complete Seq2Seq Implementation in PyTorch

```
1 import torch
2 import torch.nn as nn
3
4 class Seq2Seq(nn.Module):
5     def __init__(self, src_vocab,
6                  tgt_vocab, embed_dim=256,
7                  hidden_dim=512):
8         super().__init__()
9
9     # Embeddings
10    self.src_embed = nn.Embedding(
11        src_vocab, embed_dim
12    )
13    self.tgt_embed = nn.Embedding(
14        tgt_vocab, embed_dim
15    )
16
17    # Encoder & Decoder
18    self.encoder = nn.LSTM(
19        embed_dim, hidden_dim,
20        batch_first=True
21    )
22    self.decoder = nn.LSTM(
23        embed_dim, hidden_dim,
24        batch_first=True
25    )
26
27    # Output projection
28    self.output = nn.Linear(
29        hidden_dim, tgt_vocab
```

```
1 def forward(self, src, tgt):
2     # Encode
3     src_emb = self.src_embed(src)
4     _, (h, c) = self.encoder(
5         src_emb
6     )
6
7     # Decode
8     tgt_emb = self.tgt_embed(tgt)
9     out, _ = self.decoder(
10        tgt_emb, (h, c)
11    )
12
13     # Project
14     logits = self.output(out)
15
16     return logits
17
18
19 # Usage
20 model = Seq2Seq(
21     src_vocab=10000,
22     tgt_vocab=10000
23 )
24
25 # Training step
26 src = torch.randint(0, 10000,
27                     (32, 20))
28 tgt = torch.randint(0, 10000,
29                     (32, 15))
30 logits = model(src, tgt)
```

Encoding Example: “The black cat sat”

Watch the hidden state evolve:

Step 1: Process “The”

- Input: $x_1 = \text{embed}(\text{"The"}) = [0.1, 0.3, \dots]$
- Hidden: $h_1 = \text{LSTM}(x_1, h_0)$
- Memory: “Determiner seen”

Step 2: Process “black”

- Input: $x_2 = \text{embed}(\text{"black"})$
- Hidden: $h_2 = \text{LSTM}(x_2, h_1)$
- Memory: “Determiner + adjective”

Step 3: Process “cat”

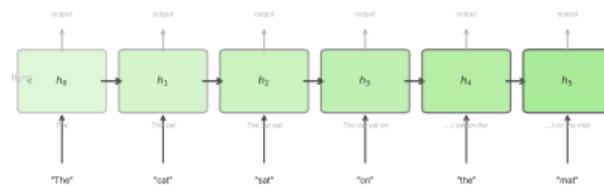
- Input: $x_3 = \text{embed}(\text{"cat"})$
- Hidden: $h_3 = \text{LSTM}(x_3, h_2)$
- Memory: “Black cat (subject)”

Step 4: Process “sat”

- Input: $x_4 = \text{embed}(\text{"sat"})$
- Hidden: $h_4 = \text{LSTM}(x_4, h_3)$
- Memory: “Black cat sat (complete)”

Encoding Process: Each word updates understanding, final state has complete meaning

Hidden State Evolution: Building Context



Final Context:

$c = h_4$ contains: - Subject: black cat - Action: sat - Tense: past

Decoding Example: Generating “Le chat noir”

Starting from context c :

Step 1: Generate “Le”

- State: $s_0 = c$
- Input: START_{L} token
- Output: $P(\text{"Le"}) = 0.8$
- Next: $s_1 = \text{LSTM}(\text{"Le"}, s_0)$

Step 2: Generate “chat”

- State: s_1 (knows “Le”)
- Input: “Le”
- Output: $P(\text{"chat"}) = 0.7$
- Next: $s_2 = \text{LSTM}(\text{"chat"}, s_1)$

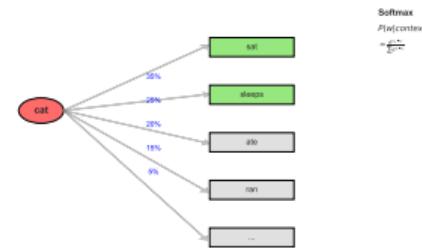
Step 3: Generate “noir”

- State: s_2 (knows “Le chat”)
- Input: “chat”
- Output: $P(\text{"noir"}) = 0.6$
- Next: $s_3 = \text{LSTM}(\text{"noir"}, s_2)$

Probability Distribution:

At each step, model outputs:

Word Prediction Probabilities from Context



Key Point:

Decoder maintains its own hidden state separate from encoder

Decoding Process: Conditional generation using context and previous outputs

Quiz Checkpoint: Understanding Seq2Seq

Questions:

Q1: What is the context vector?

- a) Average of word embeddings
- b) Final encoder hidden state
- c) Sum of all hidden states
- d) Random initialization

Q2: Why use two separate networks?

- a) Faster training
- b) Different tasks (read vs write)
- c) More parameters
- d) Requirement of RNNs

Q3: Teacher forcing means:

- a) Using true targets during training
- b) Forcing convergence
- c) Teaching the teacher

Answers:

A1: b) Final encoder hidden state

- Contains full sentence understanding
- Fixed-size representation
- Passed to decoder

A2: b) Different tasks

- Encoder: comprehension
- Decoder: generation
- Different objectives

A3: a) Using true targets

- Feed correct previous word
- Speeds up training
- Avoids error accumulation

Part 2 Summary: The Encoder-Decoder Solution

Architecture Components:

- Encoder RNN: reads input
- Context vector: compressed meaning
- Decoder RNN: generates output
- End-to-end training

Key Equations:

$$c = \text{Encoder}(x_{1:n}) \quad (9)$$

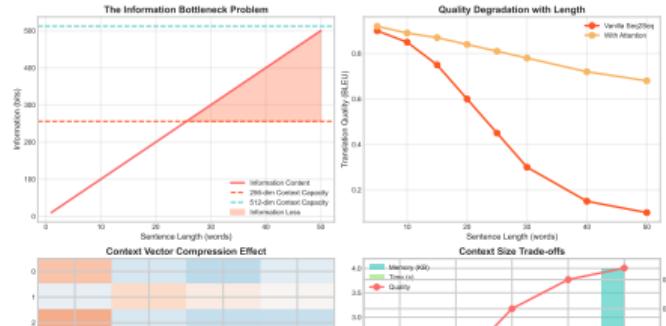
$$y_t = \text{Decoder}(c, y_{<t}) \quad (10)$$

Strengths:

- Variable input/output length
- End-to-end learning
- No alignment needed
- Works for any language pair

Weakness:

Fixed-size bottleneck!

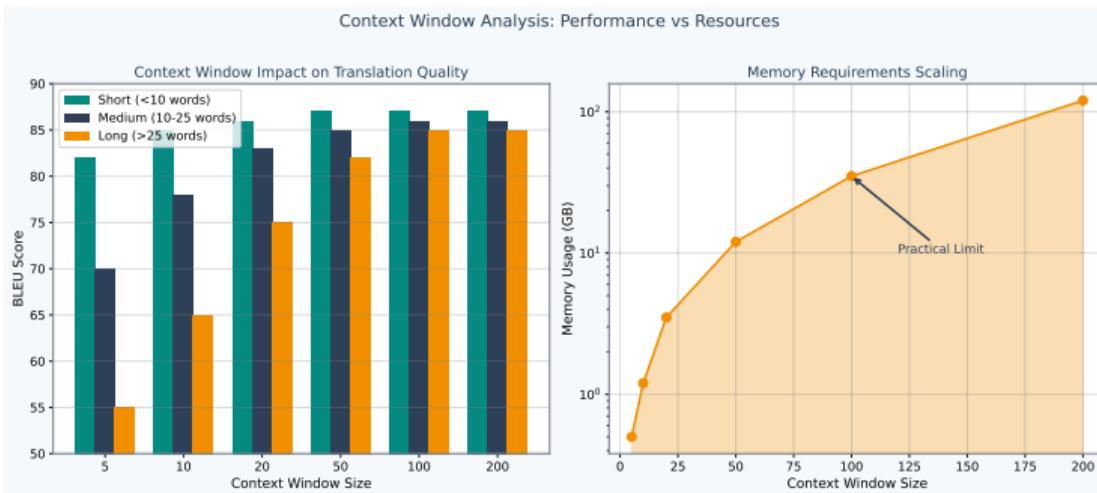


Sequence-to-Sequence Models

Part 3: The Attention Revolution

Looking Back at All Hidden States

The Bottleneck Problem: Why Seq2Seq Fails on Long Sentences



Performance Degradation:

- 10 words: BLEU = 35
- 20 words: BLEU = 25
- 30 words: BLEU = 15
- 40+ words: BLEU ↓ 10

What's Lost:

- Early words forgotten
- Specific details blurred
- Word positions unclear
- Grammatical structure

Human Translation: The Attention Analogy

How do humans really translate long sentences?

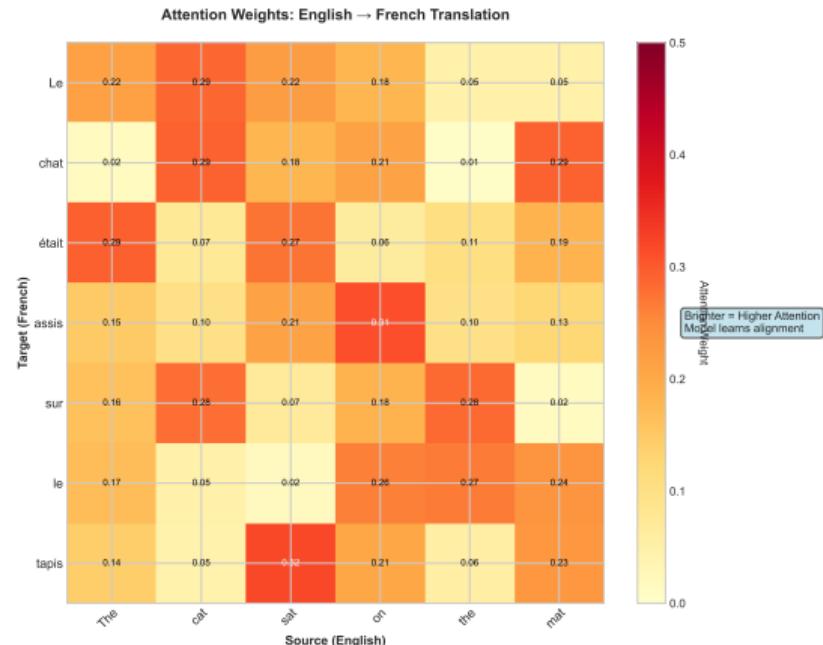
Translating Word by Word:

"The black cat that I saw yesterday sat"

When translating "sat":

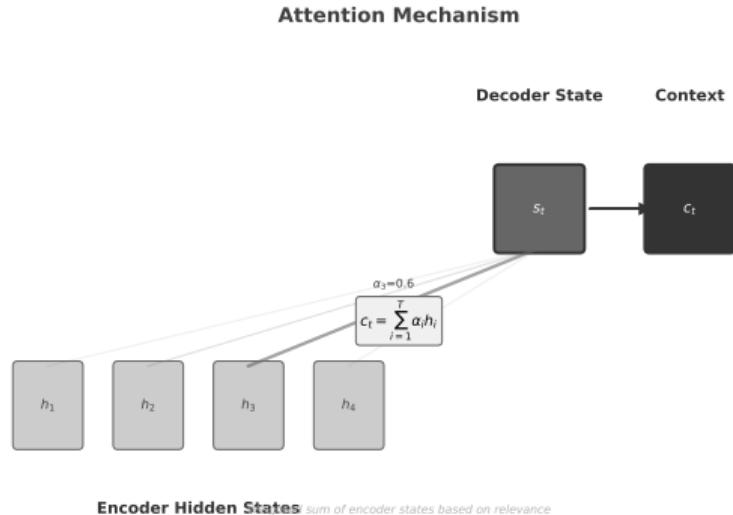
1. Look back at "cat" (subject)
2. Check tense markers
3. Verify agreement
4. Generate appropriate form

Key: We don't memorize everything! We look back as needed.



Attention Idea:

The Attention Mechanism: Dynamic Context



Old Way (Seq2Seq):

- Fixed context $c = h_n$
- Same for all decoder steps
- Information bottleneck
- Encodes each word

New Way (Attention):

- Dynamic context c_t
- Different for each word
- Weighted sum of all states
- Decodes each word

Attention Mathematics: Computing Weights

How to calculate attention weights:

Step 1: Score How relevant is each encoder state?

$$e_{ti} = \text{score}(s_{t-1}, h_i)$$

Step 2: Normalize Convert scores to probabilities:

$$\alpha_{ti} = \frac{\exp(e_{ti})}{\sum_j \exp(e_{tj})}$$

Common scoring functions:

- **Dot:** $s_{t-1} \cdot h_i$
- **General:** $s_{t-1} W h_i$
- **Concat:** $v \tanh(W[s_{t-1}; h_i])$

Step 3: Weighted Sum Compute dynamic context:

$$c_t = \sum_{i=1}^n \alpha_{ti} h_i$$

Each decoder step gets its own custom-weighted view of the source!

Mathematical Core: Three simple steps that revolutionized machine translation

Understanding the QKV Framework:

Components:

- **Query (s_{t-1})**: What I'm looking for
- **Keys (h_i)**: What's available
- **Values (h_i)**: What to retrieve
- **Weights (α_{ti})**: Relevance scores

Analogy:

- Query = Search term
- Keys = Document titles
- Values = Document content
- Attention = Search relevance

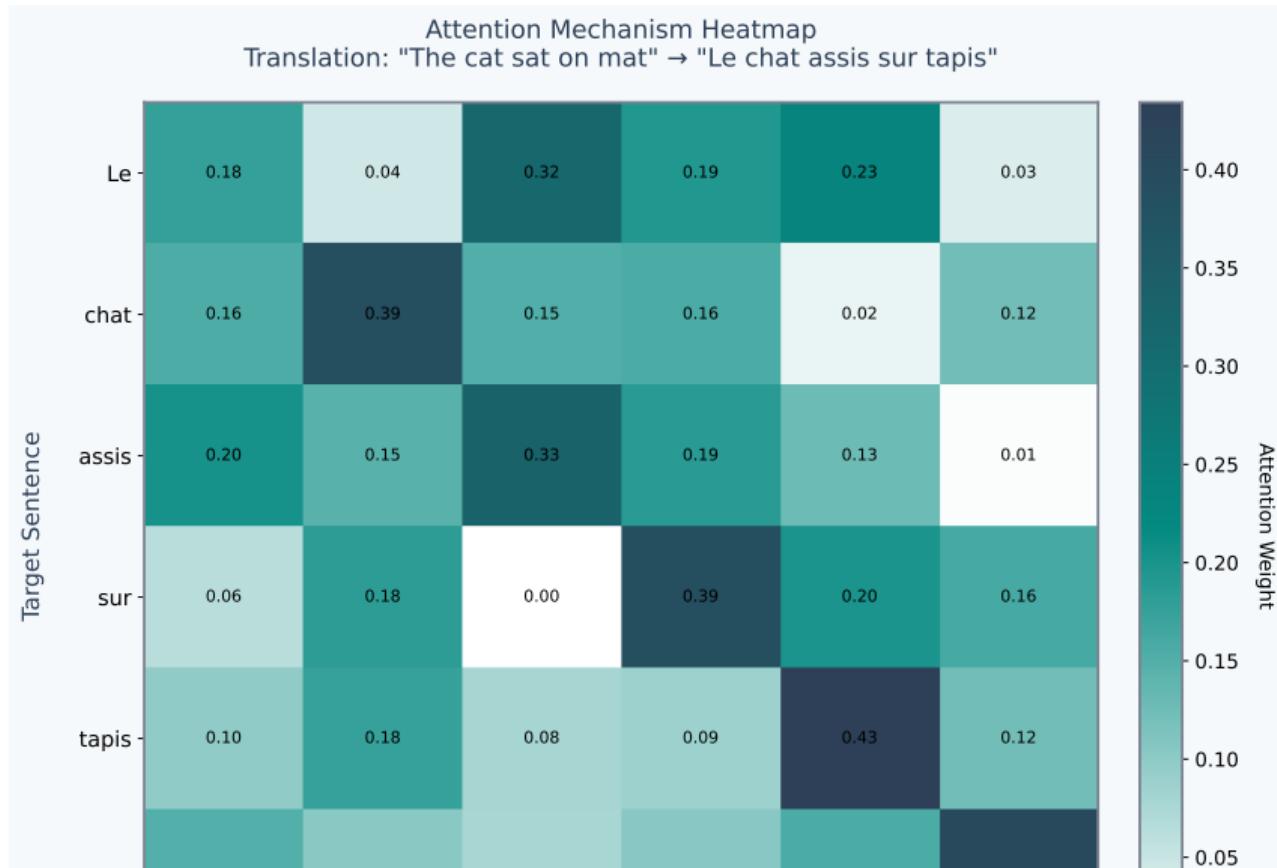
Critical Insight:

This simple mechanism
is the foundation of
ALL modern transformers!

GPT, BERT, T5, ChatGPT...
all use this QKV attention

Foundation: Attention as learned information retrieval - the key to modern AI

Visualizing Attention: What the Model Focuses On



Implementing Attention: Core Logic

```
1  class Attention(nn.Module):
2      def __init__(self, hidden_dim):
3          super().__init__()
4          # Score function parameters
5          self.attn = nn.Linear(hidden_dim * 2, hidden_dim)
6          self.v = nn.Linear(hidden_dim, 1, bias=False)
7
8      def forward(self, hidden, encoder_outputs):
9          # hidden: [1, batch, hidden_dim] - decoder state
10         # encoder_outputs: [batch, seq_len, hidden_dim] - all encoder states
11
12         batch_size = encoder_outputs.size(0)
13         seq_len = encoder_outputs.size(1)
14
15         # Repeat decoder hidden for each encoder position
16         hidden = hidden.squeeze(0).unsqueeze(1).repeat(1, seq_len, 1)
17
18         # Score each encoder state (concat scoring)
19         energy = torch.tanh(self.attn(
20             torch.cat((hidden, encoder_outputs), dim=2)
21         ))
22
23         # Convert to scalar scores
24         attention = self.v(energy).squeeze(2)
25
26         # Apply softmax to get weights
27         weights = F.softmax(attention, dim=1)
28
29         # Weighted sum of encoder states
30         context = torch.bmm(weights.unsqueeze(1), encoder_outputs)
31
```

Implementing Attention: Usage in Decoder

```
1  class DecoderWithAttention(nn.Module):
2      def __init__(self, vocab_size, embed_dim, hidden_dim):
3          super().__init__()
4          self.embedding = nn.Embedding(vocab_size, embed_dim)
5          self.attention = Attention(hidden_dim)
6          self.lstm = nn.LSTM(embed_dim + hidden_dim, hidden_dim)
7          self.output = nn.Linear(hidden_dim, vocab_size)
8
9      def forward(self, input_token, hidden, encoder_outputs):
10         # Embed current token
11         embedded = self.embedding(input_token) # [batch, 1, embed]
12
13         # Compute attention context
14         context, attn_weights = self.attention(hidden, encoder_outputs)
15         # context: [batch, 1, hidden_dim]
16
17         # Concatenate embedding with context
18         lstm_input = torch.cat((embedded, context), dim=2)
19
20         # LSTM step with attention-augmented input
21         output, hidden = self.lstm(lstm_input, hidden)
22
23         # Project to vocabulary
24         predictions = self.output(output)
25
26         return predictions, hidden, attn_weights
27
28 # Training loop usage
29 decoder = DecoderWithAttention(vocab_size=10000, embed_dim=256, hidden_dim=512)
```

Interactive Exercise: Calculate Attention Weights

Task: Compute attention for generating “noir” (black)

Given decoder state s_2 after generating “Le chat”:

Encoder states:

- h_1 : “The” = [0.1, 0.2]
- h_2 : “black” = [0.8, 0.9]
- h_3 : “cat” = [0.5, 0.4]
- h_4 : “sat” = [0.3, 0.1]

Decoder query:

- s_2 = [0.7, 0.8]

Your calculations:

1. Scores (dot product):

- $e_1 = s_2 \cdot h_1 = \underline{\hspace{2cm}}$
- $e_2 = s_2 \cdot h_2 = \underline{\hspace{2cm}}$
- $e_3 = s_2 \cdot h_3 = \underline{\hspace{2cm}}$
- $e_4 = s_2 \cdot h_4 = \underline{\hspace{2cm}}$

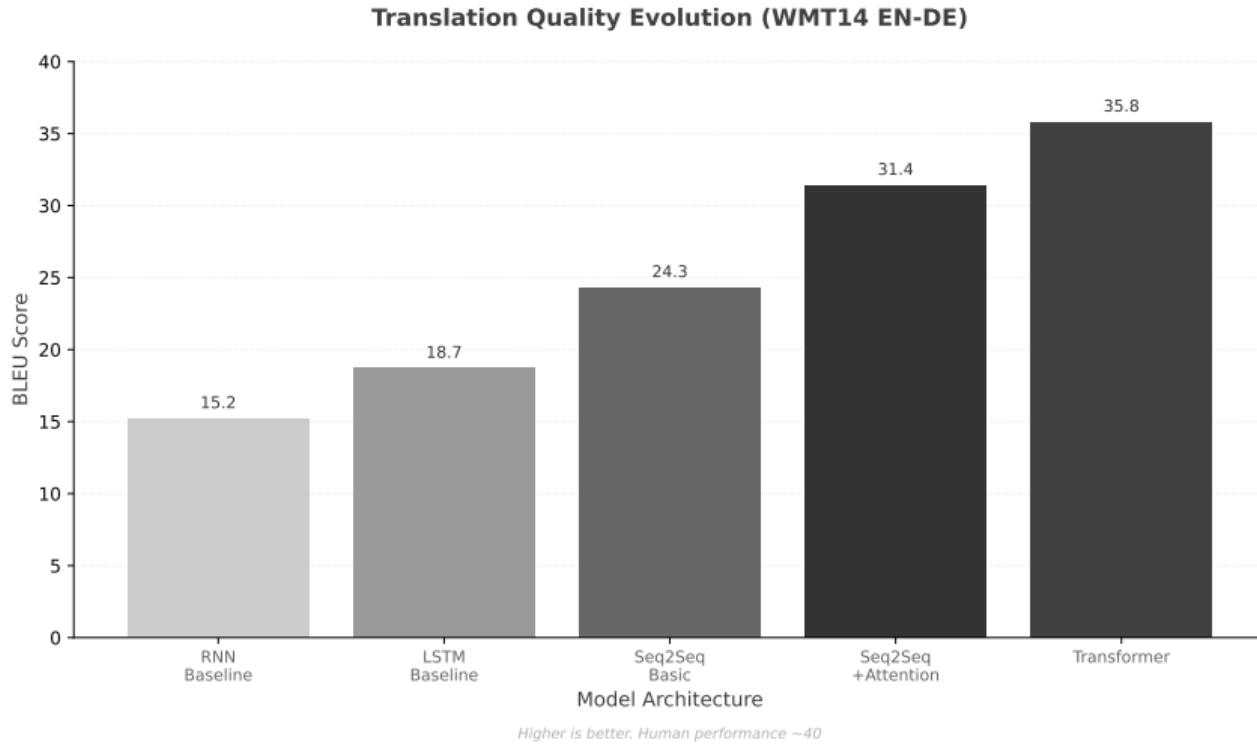
2. Softmax weights:

- $\alpha_2 = \underline{\hspace{2cm}}$ (highest!)

3. Context: weighted sum

Hands-On: Computing attention manually builds intuition for the mechanism

Impact of Attention: Dramatic Improvements



BLEU Score Improvements:

Why It Works:

Part 3 Summary: The Attention Revolution

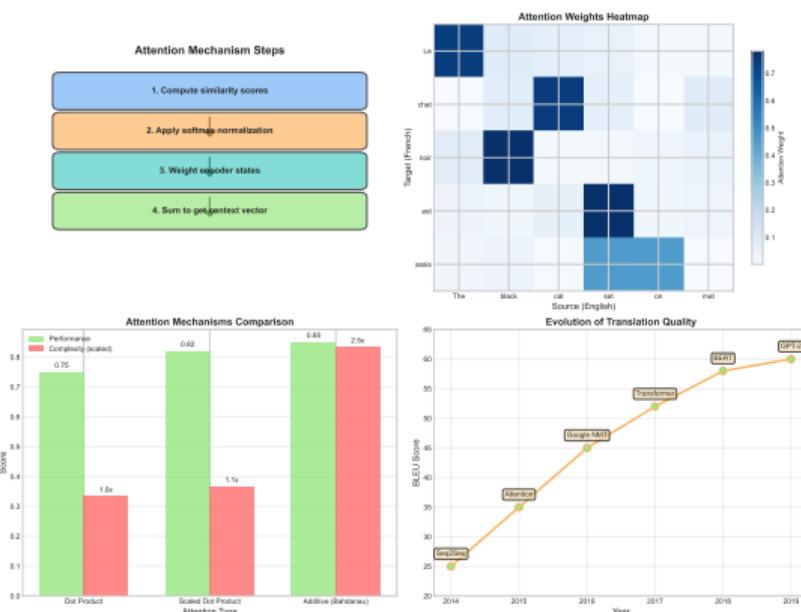
The Innovation:

- Dynamic context vectors
- Look at all encoder states
- Weighted by relevance
- Different for each word

Mathematical Core:

$$\alpha_{ti} = \text{softmax}(\text{score}(s_t, h_i)) \quad (11)$$

$$c_t = \sum_i \alpha_{ti} h_i \quad (12)$$



Impact:

Attention mechanism became foundation of all modern NLP

Historical Significance: Attention paper (2014) revolutionized entire field

Part 4: Implementation & Applications

From Research to Production

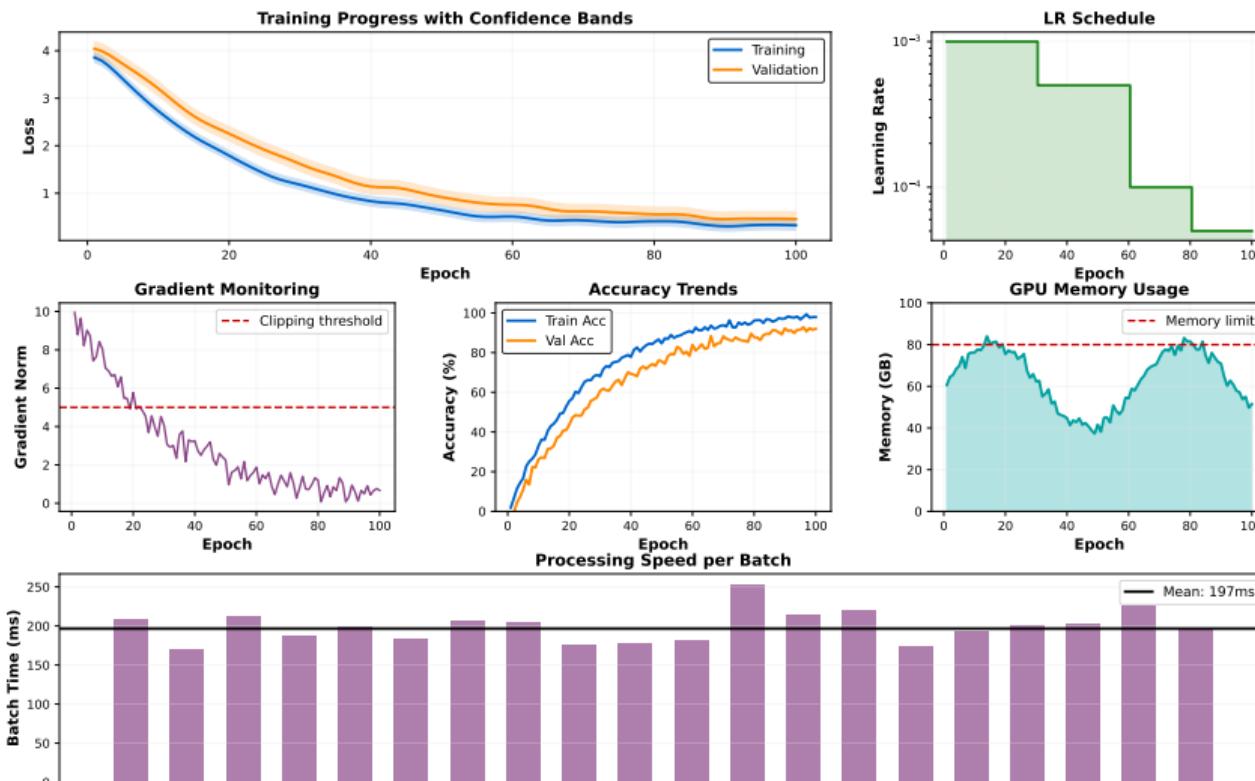
Complete Seq2Seq with Attention

```
1 class AttentionSeq2Seq(nn.Module):
2     def __init__(self, src_vocab,
3                  tgt_vocab, dim=512):
4         super().__init__()
5
5         # Components
6         self.encoder = Encoder(
7             src_vocab, dim
8         )
9         self.decoder = DecoderWithAttn(
10            tgt_vocab, dim
11        )
12         self.attention = Attention(dim)
13
14     def forward(self, src, tgt,
15                teacher_forcing=0.5):
16         # Encode all at once
17         enc_out, (h, c) = self.encoder(src)
18
19         batch = src.size(0)
20         max_len = tgt.size(1)
21         vocab = self.decoder.vocab_size
22
23         # Store outputs
24         outputs = torch.zeros(
25             batch, max_len, vocab
26         )
27
28         # First input
29         input = tgt[:, 0]
```

```
1         for t in range(1, max_len):
2             # Attention context
3             context, weights =
4                 self.attention(
5                     h, enc_out
6                 )
7
8             # Decode one step
9             output, (h, c) =
10                 self.decoder(
11                     input, (h, c),
12                     context
13                 )
14
15             outputs[:, t] = output
16
17             # Teacher forcing
18             use_teacher = random.random()
19                 < teacher_forcing
20
21             if use_teacher:
22                 input = tgt[:, t]
23             else:
24                 input = output.argmax(1)
25
26         return outputs
```

Training Dynamics: Learning to Translate

Training Monitoring Dashboard



Hyperparameters That Matter:

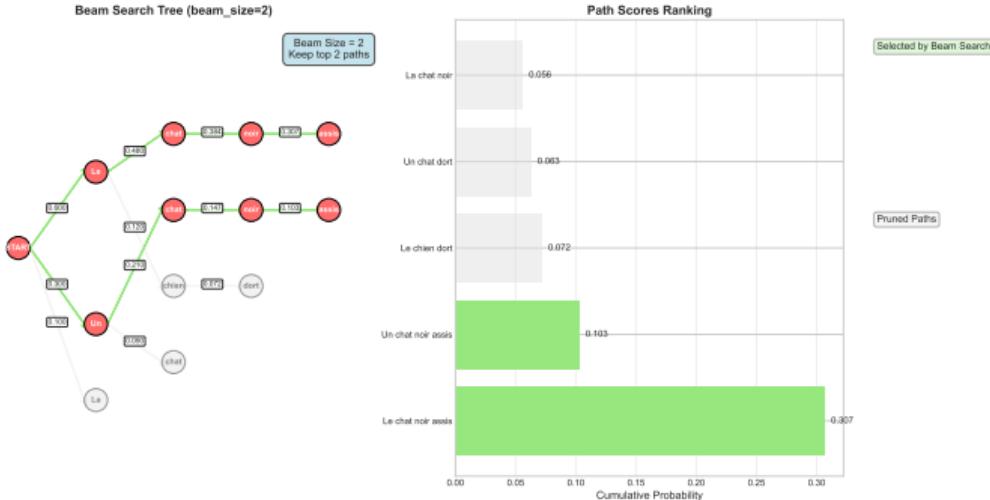
- **Learning rate:** Start 0.001, decay after epoch 10
- **Teacher forcing:** 0.5 → 0 over training
- **Gradient clip:** Essential (1.0 works well)
- **Batch size:** 32-64 optimal for GPU
- **Hidden size:** 512 is sweet spot
- **Layers:** 2-3 LSTM layers sufficient

Common Issues & Solutions:

- **Exploding loss:** → Reduce learning rate
- **Mode collapse:** → Add dropout (0.3)
- **Poor rare words:** → Increase min frequency
- **Slow training:** → Use GPU, reduce batch size
- **Overfitting:** → More data, regularization
- **Underfitting:** → Bigger model, longer training

Pro Tip: Start simple, add complexity gradually, monitor validation metrics closely

Beam Search: Better Decoding Strategy



Greedy vs Beam Search:

- Greedy: Pick best at each step
- Beam: Keep top-k hypotheses
- Explore multiple paths
- Better final translations

Beam size trade-off:

NLP Course 2025 (Professional Template Edition)

Example with beam=3:

Step 1: "Le", "Un", "Les"

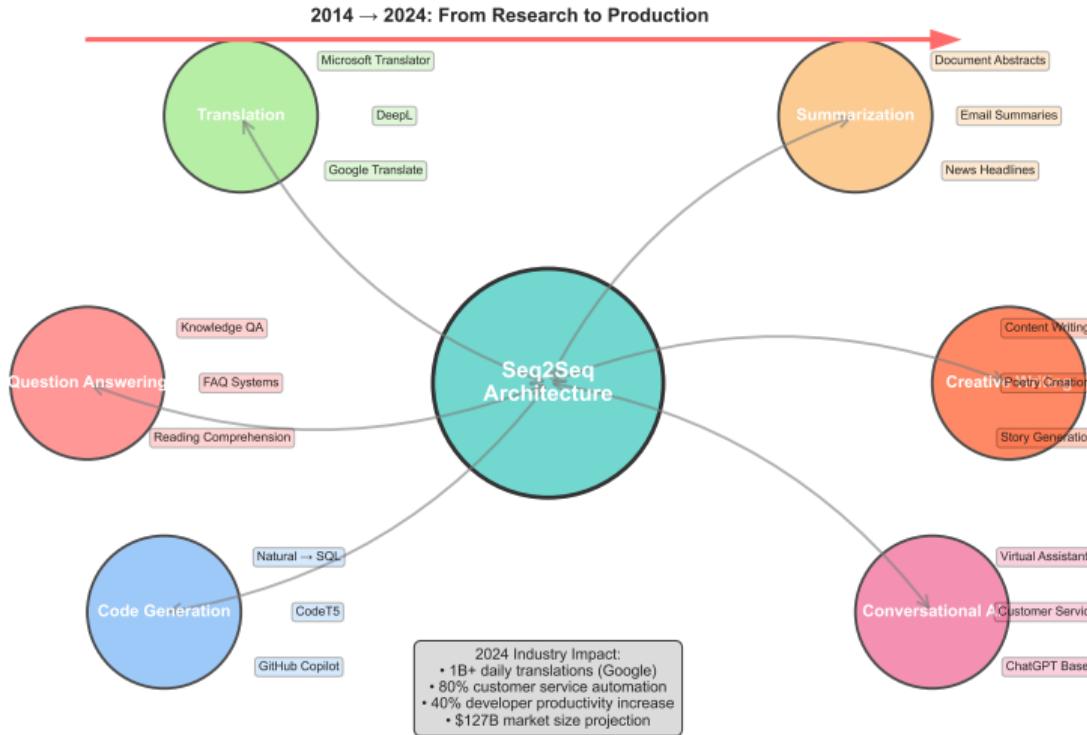
Step 2:

- "Le chat", "Le chien"
- "Un chat", "Les chats"

Step 3: Keep expanding top-3

Final: Pick highest scoring complete hypothesis

Seq2Seq Models: Modern Applications Ecosystem (2024)



Attention Everywhere:

- ChatGPT/Claude (attention-based)
- Image captioning
- Video understanding
- Code generation
- Music composition
- Speech recognition
- Medical diagnosis

Foundation Truth:

Seq2Seq + Attention
=
Modern AI Backbone

Every modern language model
builds on this architecture

Attention mechanism is the foundation of the entire AI revolution

Lab Preview: Build Your Own Translator

Week 4 Lab: English-French Neural Machine Translation

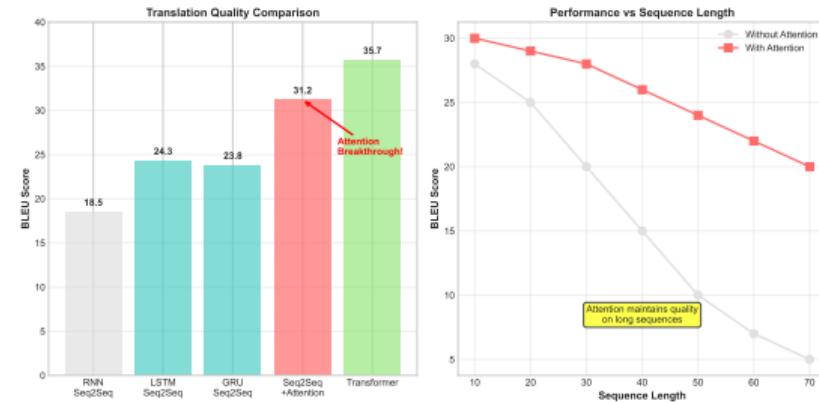
What You'll Build:

1. Load parallel corpus
2. Tokenize and preprocess
3. Implement encoder-decoder
4. Add attention mechanism
5. Train on GPU
6. Visualize attention weights
7. Compare with/without attention

Dataset:

- 10,000 sentence pairs
- English → French
- Average 15 words/sentence

Expected Results:



Bonus Challenges:

- Multi-head attention
- Bidirectional encoder
- Coverage mechanism
- Back-translation

Interactive Debugging: Common Training Issues

Your model isn't learning. Debug these issues:

Issue 1: Attention all uniform

Symptoms:

- All weights $\approx 1/n$
- Poor translation quality
- Not improving

Your fix: _____
Hint: Check score function

Common Fixes:

Fix 1: Initialize properly

- Use Xavier initialization
- Scale attention scores
- Add small epsilon to softmax

Issue 2: Mode collapse

Symptoms:

- Always generates "the the the"
- Loss plateaus high

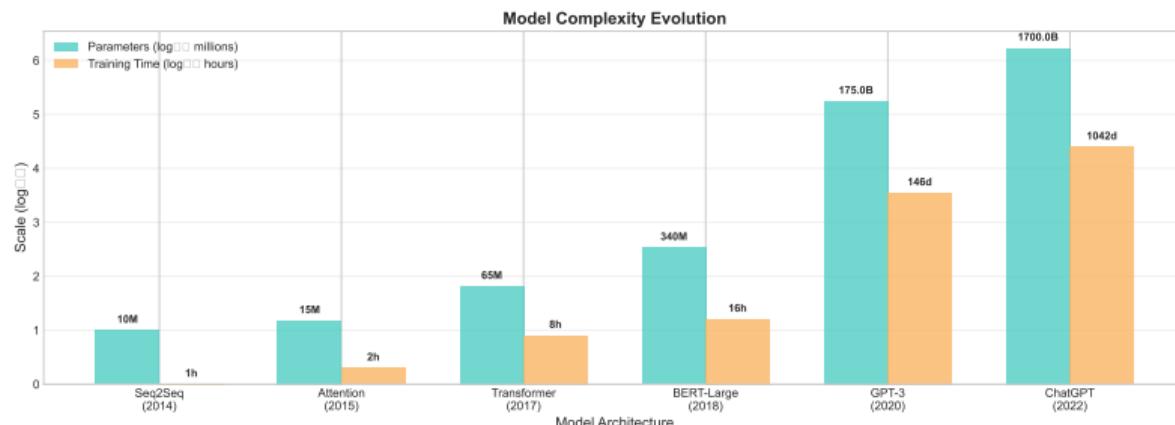
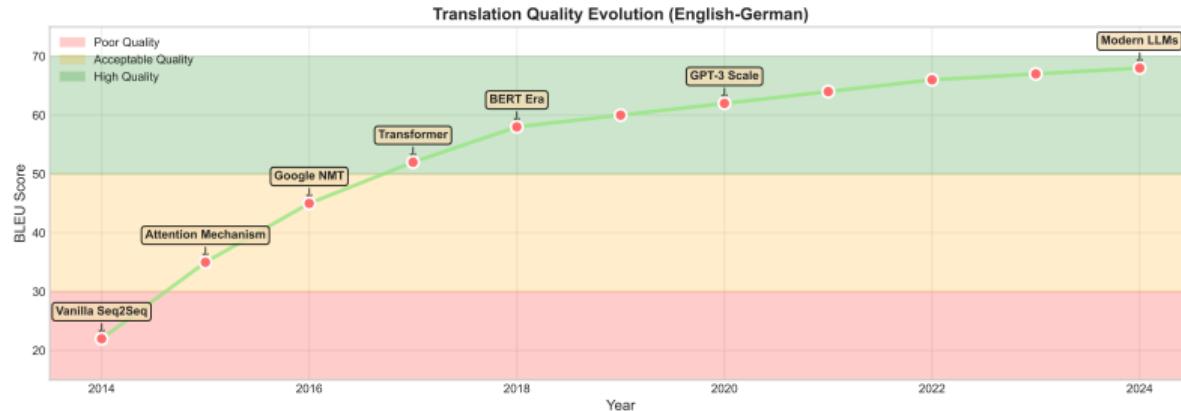
Fix 2: Teacher forcing

- Start with 100% teacher forcing
- Gradually reduce ratio
- Scheduled sampling

Debug systematically!

Debugging Skills: Most issues come from initialization or training schedule

Performance Evolution Timeline



2014: Seq2Seq

- BLEU: 20-25
- Simple, elegant
- Length problems
- 2-3 days training

Impact:

- First neural MT
- Proof of concept
- Beat phrase-based

2015: + Attention

- BLEU: 30-35
- Handles length
- Interpretable
- 4-5 days training

Impact:

- Production ready
- Google adoption
- 40% improvement

2017: Transformer

- BLEU: 40+
- All attention
- Parallel training
- 12 hours training!

Impact:

- New paradigm
- Enables GPT/BERT
- 10x faster training

Key Insight: Each innovation built on the previous - attention was THE breakthrough

Historical Progression: From RNN to attention to transformer architecture

The Bridge to Transformers (Week 5 Preview)

From Seq2Seq+Attention to Transformers:

What We Keep:

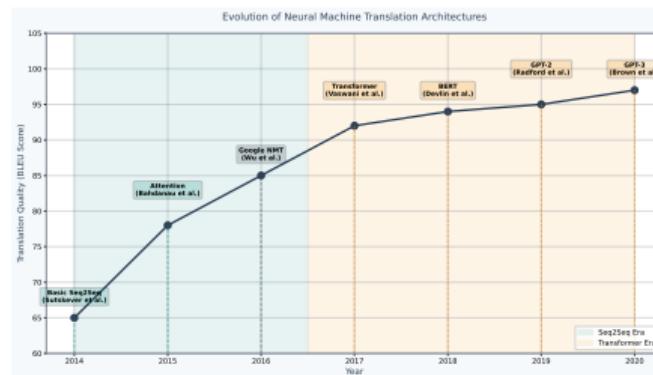
- Attention mechanism
- Query-Key-Value
- Position awareness
- Encoder-decoder structure

What We Add:

- Self-attention
- Multi-head attention
- Position encodings
- Layer normalization
- Parallel processing

What We Remove:

- RNN/LSTM cells
- Sequential processing
- Recurrent connections
- Hidden state passing



Part 1: Challenge

- Translation \neq word replacement
- Need meaning understanding
- Information bottleneck problem

Part 2: Seq2Seq

- Encoder-decoder architecture
- Fixed context vector
- Works but limited by bottleneck

Part 3: Attention

- Dynamic context vectors
- Look at all encoder states
- Massive performance improvement

Part 4: Applications

- Complete implementation
- Beam search decoding
- Powers modern translation
- Foundation for transformers

Key Takeaways:

1. Context vectors compress meaning
2. Attention removes bottleneck
3. Foundation of modern NLP
4. Bridge to transformers

Next Week: Transformers - Attention Without RNNs!

Achievement Unlocked: You understand the foundation of all modern language AI!