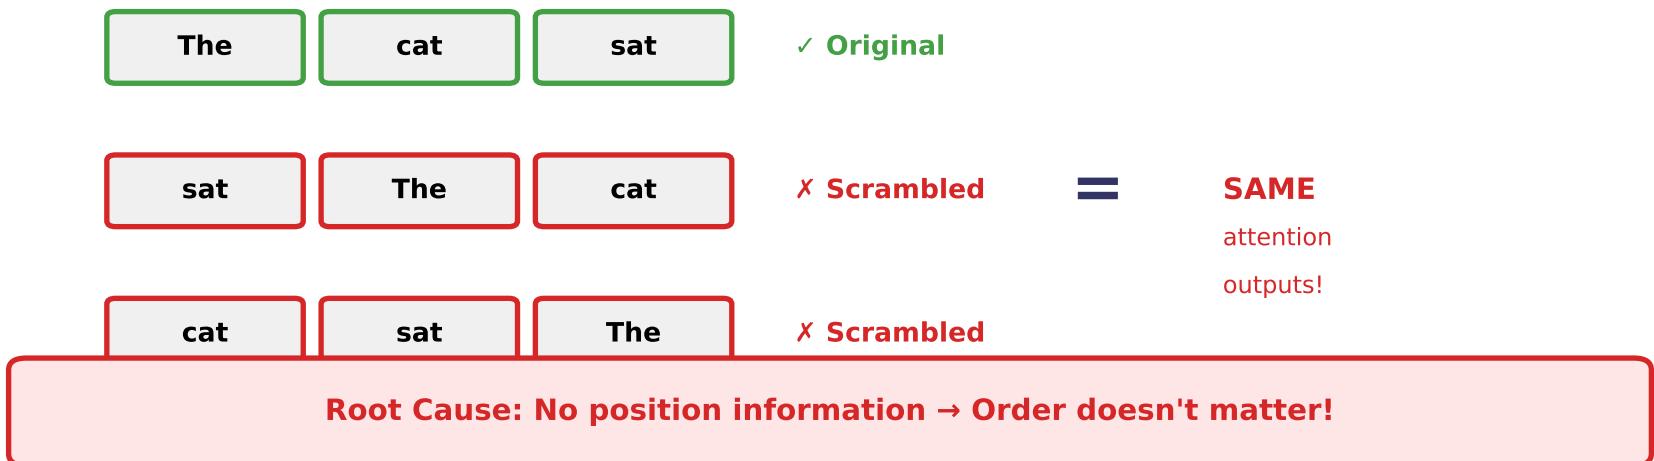


THE PROBLEM: Attention is Permutation Invariant



THE SOLUTION: Four Requirements for Position Encoding

1 Inject Position Info

Add unique pattern to each position

2 Parallel Computation

No sequential dependency

3 Any Sequence Length

Works for 10 or 1000 words

Position as added signal

4 Relative Positions

Preserves distance between words

Key Insight: Position can be a feature (added signal), not a process (sequential)