

Pre-trained Language Models

NLP Course 2025

September 30, 2025

A pedagogical journey from wasteful repetition to efficient transfer learning
45 slides • 4 sections • 2 checkpoints

What You Will Master

Conceptual Understanding

- Why pre-training revolutionized NLP
- Transfer learning principles
- BERT vs GPT paradigms
- Fine-tuning strategies
- Model selection criteria

Practical Skills

- Use pre-trained models
- Fine-tune for your task
- Choose the right model
- Implement with HuggingFace
- Avoid common pitfalls

By the end: You'll understand how to leverage billions of parameters trained on terabytes of text

Pre-training changed everything - from months to minutes, from millions to thousands

Section 1

The Challenge

Reinventing the Wheel

The Reading Revolution: An Analogy

Imagine two medical students:

Student A: Starting from Zero

- Never learned to read
- Must learn alphabet first
- Then words, grammar, sentences
- Finally can read medical texts
- Takes 10 years to become doctor

Student B: Transfer Learning

- Already knows how to read
- Starts directly with medical texts
- Focuses only on medical knowledge
- Becomes doctor in 4 years
- Better outcomes, less time

Warning: 90% of time learning to read, only 10% on medicine

100% of time on medicine, reading skill transfers

Intuition: Pre-training is like teaching AI to read before teaching it specific tasks!

This analogy captures why pre-training changed everything in NLP

The \$10 Million Problem No One Talked About

The wasteful reality of 2017-2018:

The Old Way:

- Company A: Sentiment analysis
 - Cost: \$500,000
 - Time: 2 weeks on 64 GPUs
 - Learns: Grammar from scratch
- Company B: Translation
 - Cost: \$500,000
 - Time: 2 weeks on 64 GPUs
 - Learns: Grammar from scratch
- Company C: Q&A
 - Cost: \$500,000
 - Time: 2 weeks on 64 GPUs
 - Learns: Grammar from scratch

Total waste: \$1.5 million!

The Smart Way (Today):

- Pre-train ONCE: \$1 million
- Fine-tune for each task:
 - Sentiment: \$10,000
 - Translation: \$10,000
 - Q&A: \$10,000

Total cost: \$1.03 million

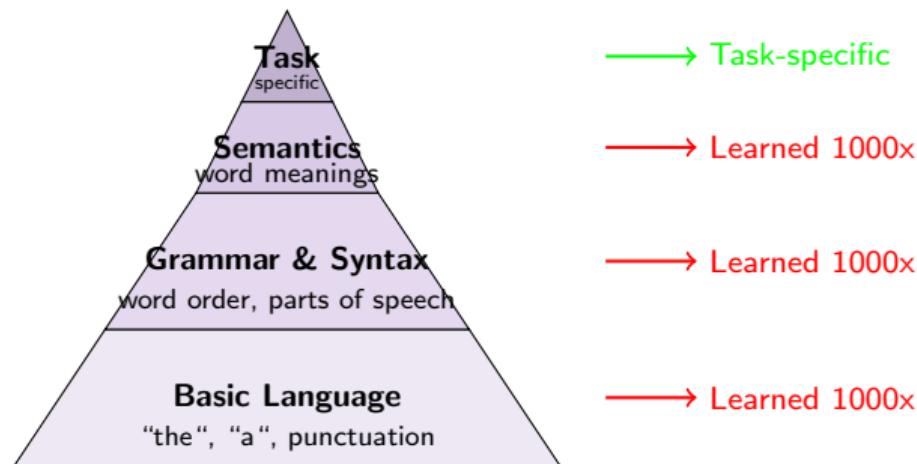
Savings: \$470,000!

Plus: Better performance!

Every team was re-learning what “the” means from scratch!

What Every Model Was Learning from Scratch

The redundant learning pyramid:

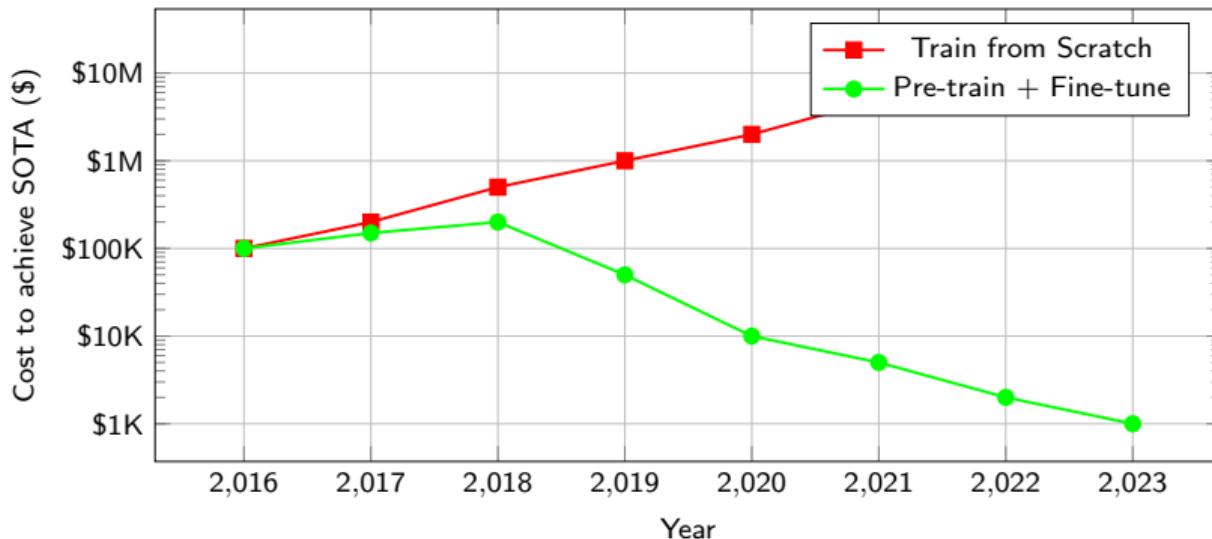


Common Misconception: “Each task needs its own model from scratch” - False! 95% of language understanding is general

The bottom 95% of the pyramid is identical for all NLP tasks

The Computational Reality

Training costs over time:



Pre-training inverted the cost curve - better models got cheaper!

The divergence started with BERT in 2018

The Carbon Footprint Reality

Environmental cost of training:

Training from Scratch (per model):

- Energy: 27,000 kWh
- CO2: 12 metric tons
- Equivalent to:
 - 5 cars for 1 year
 - 125 flights NYC-SF
 - Average home for 2.5 years

Pre-training + Fine-tuning:

- Pre-train once: 12 metric tons
- Fine-tune: 0.01 metric tons
- 1000 companies: 22 metric tons total

545x reduction in carbon footprint!

If 1000 companies do this:

- 12,000 metric tons CO2
- Small city's annual emissions

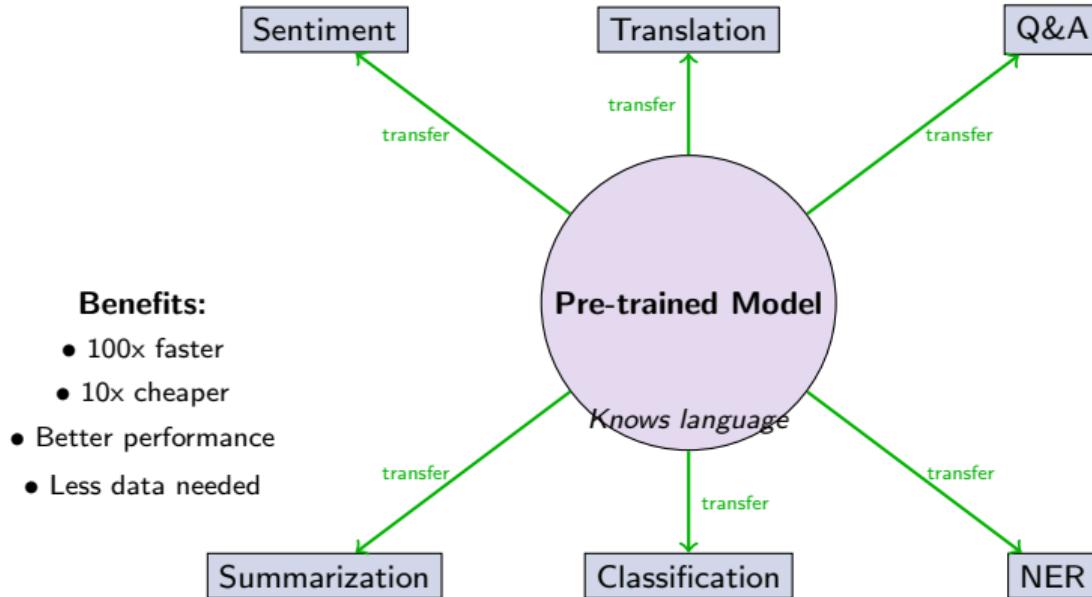
[Carbon footprint visualization]

Real World: Google's BERT pre-training saved an estimated 50,000 metric tons of CO2

Environmental responsibility drove adoption as much as cost savings

The Transfer Learning Opportunity

What if we could share knowledge?



One model's knowledge can bootstrap infinite downstream tasks

Section 2

The Foundation

Learning from Everything

What is Pre-training?

Teaching models to understand language itself:

Traditional Supervised Learning:

- Need: Labeled data
- Example: "Great movie!" → Positive
- Scale: Thousands of examples
- Cost: \$100K for annotation
- Coverage: One specific task

Pre-training (Self-Supervised):

- Need: Raw text only
- Example: "The [?] sat on the mat"
- Scale: Billions of examples
- Cost: Just compute time
- Coverage: All of language

Warning: Limited by human annotation speed

Unlimited data from the internet

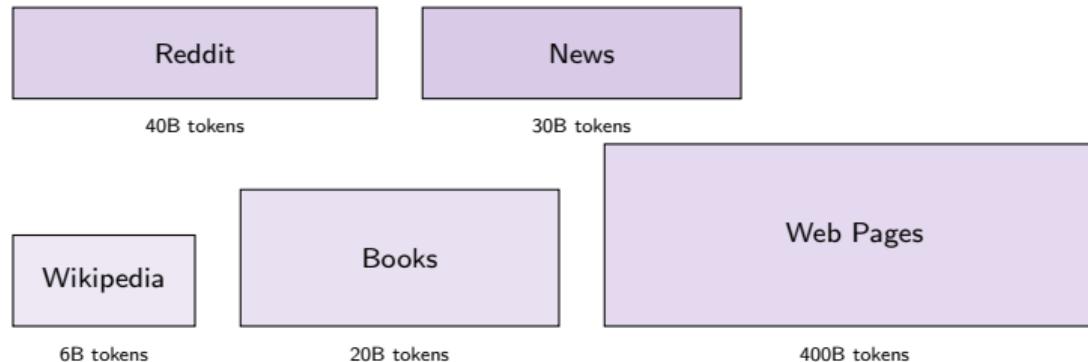
Intuition: Pre-training turns the internet into a free teacher - every sentence becomes a learning example

Self-supervision was the key unlock - no humans needed

The Scale of Pre-training Data

Where does the knowledge come from?

Total: 500+ Billion Tokens
(100,000 books worth of text)



Models read more in training than a human could in 10,000 lifetimes

GPT-3 trained on 45TB of compressed text

Self-Supervised Learning: The Secret Sauce

How to create infinite training examples:

Method 1: Masked Language (BERT)

- ① Take: "The cat sat on the mat"
- ② Mask: "The [MASK] sat on the mat"
- ③ Predict: "cat"
- ④ Learn: Context → Word

Method 2: Next Word (GPT)

- ① Take: "The cat sat on the "
- ② Predict: "mat"
- ③ Learn: Sequence → Next
- ④ Continue: Generate text

Advantages:

- Bidirectional context
- Natural for understanding
- Good for classification

Advantages:

- Natural generation
- Autoregressive
- Good for creation

Both methods turn unlabeled text into labeled training data automatically

This self-supervision eliminates the annotation bottleneck

Transfer Learning: The Theory

Why does knowledge transfer work?

[Transfer learning landscape visualization]

Key insights:

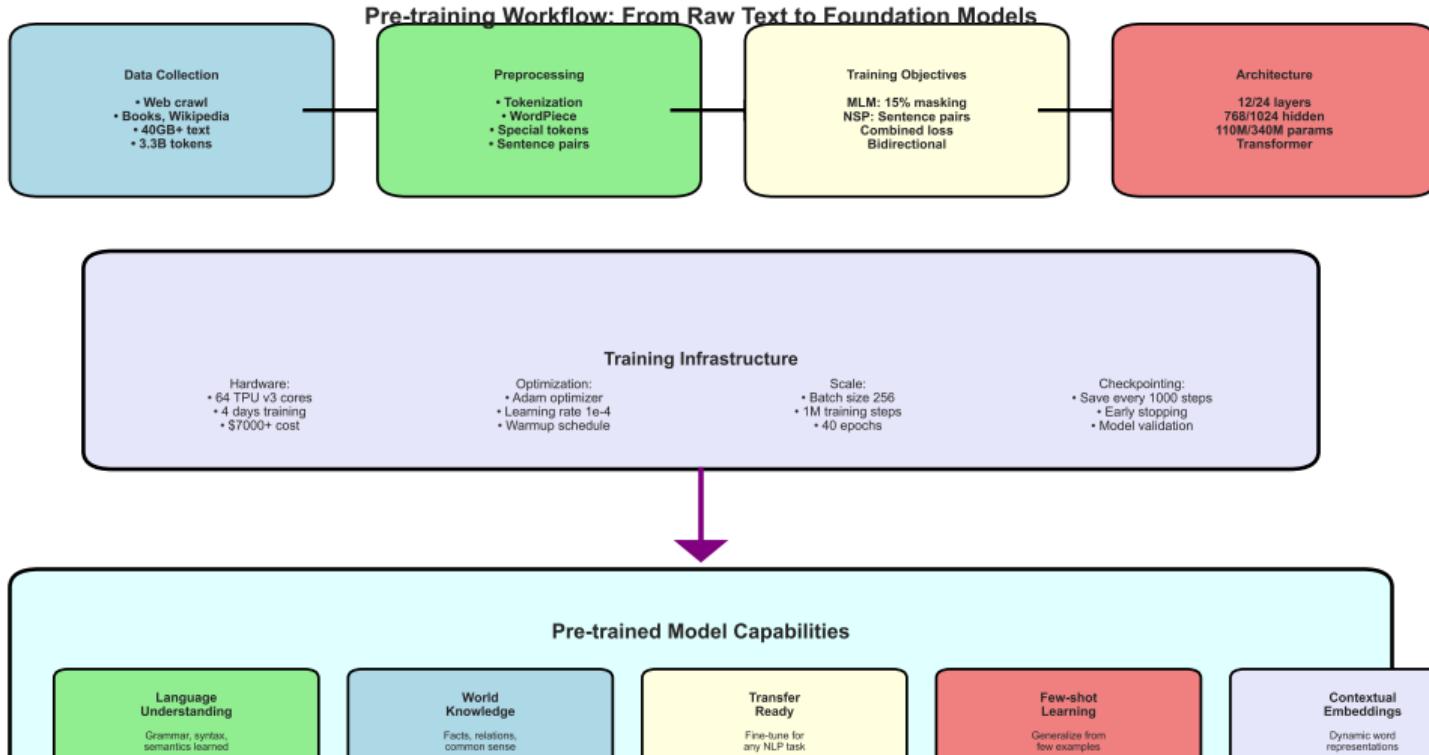
- **Lower layers:** Learn general features (word meanings, grammar)
- **Middle layers:** Learn combinations (phrases, concepts)
- **Upper layers:** Learn task-specific patterns
- **Transfer:** Reuse lower/middle, adapt upper

Intuition: Like learning piano helps with typing - the finger coordination transfers

Neural networks naturally learn hierarchical representations

The Pre-training Process Workflow

From raw text to capable model:



Layers of Understanding: What Models Learn

Knowledge emerges at different depths:

Early Layers (1-4):

- Character patterns
- Word boundaries
- Basic syntax
- Parts of speech

Example: "run" vs "runs"

Middle Layers (5-8):

- Semantic meaning
- Word relationships
- Phrase structure
- Dependencies

Example: "bank" (river/money)

Deep Layers (9-12):

- Abstract concepts
- Long-range patterns
- Task-specific features
- Reasoning steps

Example: Inference, causality

Each layer builds on the previous - from letters to logic

Probing studies reveal this consistent hierarchy across models

Test Your Understanding

Quick Quiz:

Q1: What's the main advantage of pre-training?

- A) Faster inference
- B) Transfer learning to any task
- C) Smaller models
- D) Less memory usage

Q2: How does self-supervision work?

- A) Humans label the data
- B) Models label each other
- C) Create labels from the text itself
- D) Random labeling

Q3: Why is pre-training expensive but fine-tuning cheap?

Answers:

A1: B - Transfer learning to any task

- Pre-trained knowledge applies everywhere
- One model bootstraps many applications
- Eliminates starting from scratch

A2: C - Create labels from the text itself

- Mask words → predict them
- Or predict next word
- No human annotation needed

A3: D - Both B and C

- General knowledge learned once
- Task-specific needs little data

Section 3

The Architecture

Two Paradigms: BERT and GPT

BERT: Bidirectional Encoder Representations

The fill-in-the-blank master:

Core Idea:

- Look at context from both sides
- Mask random words (15%)
- Predict what's hidden
- Learn deep bidirectional representations

[BERT masking visualization]

Example:

Input: "The [MASK] sat on the [MASK]"
Output: "cat", "mat" (high probability)

Architecture:

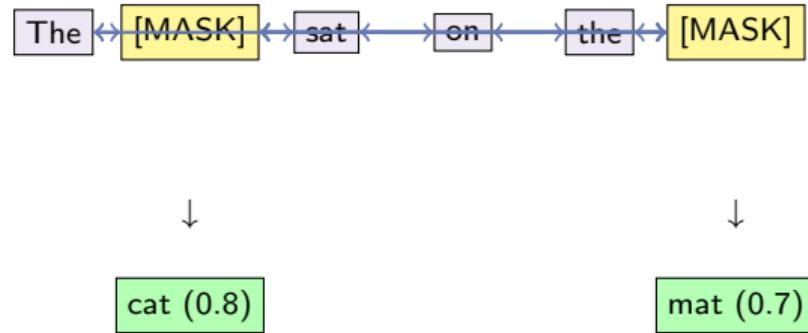
- Transformer encoder only
- 12/24 layers (Base/Large)
- 110M/340M parameters
- Attention in all directions

Intuition: BERT is like a detective examining all clues simultaneously to solve the mystery

BERT = Bidirectional Encoder Representations from Transformers

Masked Language Modeling (MLM) Explained

How BERT learns from context:



Training details:

- 80% of masks: Replace with [MASK]
- 10% of masks: Replace with random word
- 10% of masks: Keep original (but still predict)

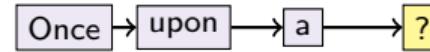
This prevents the model from only looking for [MASK] tokens

GPT: Generative Pre-trained Transformer

The story continuation expert:

Core Idea:

- Predict the next word
- Use only left context
- Autoregressive generation
- Learn to continue any text



time (0.85)
day (0.10)
night (0.05)

Architecture:

- Transformer decoder only
- 12/24/48 layers (GPT/GPT-2/GPT-3)
- 117M/1.5B/175B parameters
- Causal (left-to-right) attention

Example:

Input: "Once upon a"
Output: "time" (highest probability)

Intuition: GPT is like an author who can only see what's been written so far

How GPT learns to generate:

Training Process:

- ① Start with text sequence
- ② Predict each word from previous
- ③ Compare with actual next word
- ④ Update to improve predictions

Generation Process:

- ① Start with prompt
- ② Generate next token
- ③ Add to sequence
- ④ Repeat until done

Attention Mask:

Can attend				
	Can attend			
		Can attend		
			Can attend	
Cannot attend (future)				

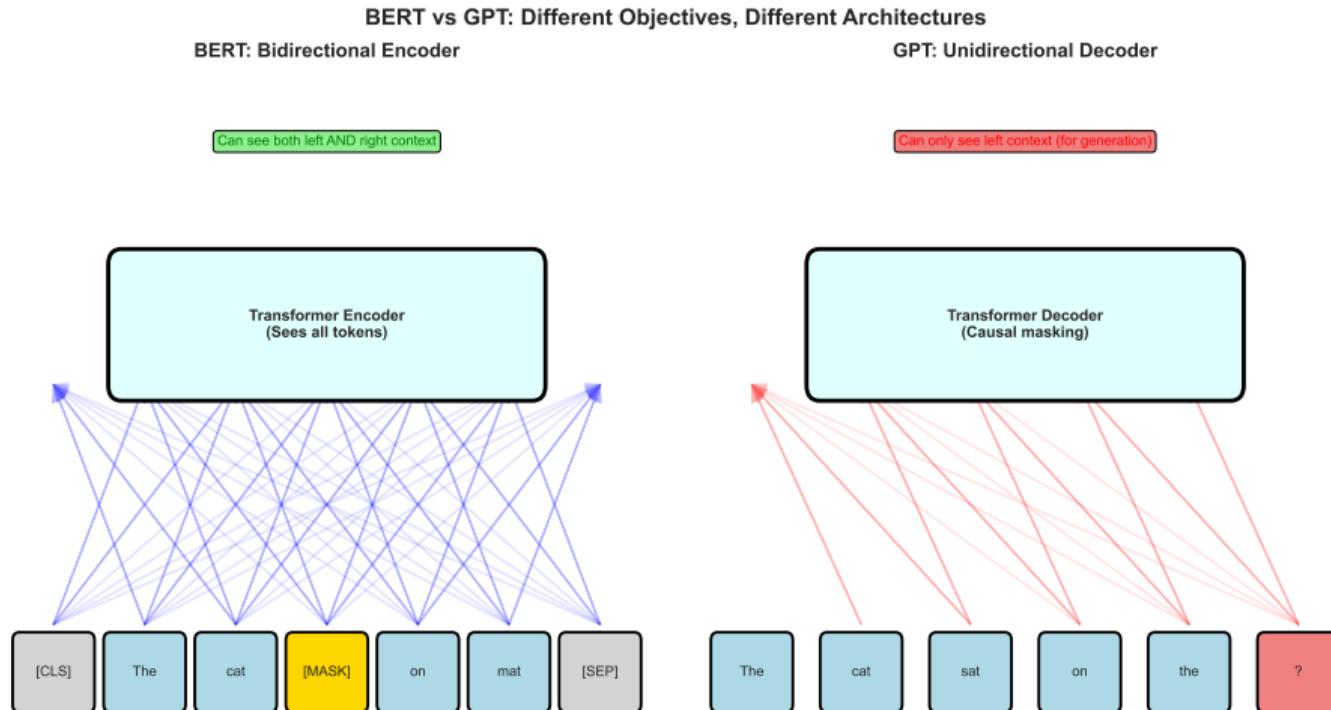
Can attend
Cannot attend (future)

Example Generation:

- Prompt: “The weather“
- Step 1: “The weather **is**“
- Step 2: “The weather is **nice**“
- Step 3: “The weather is nice **today**“

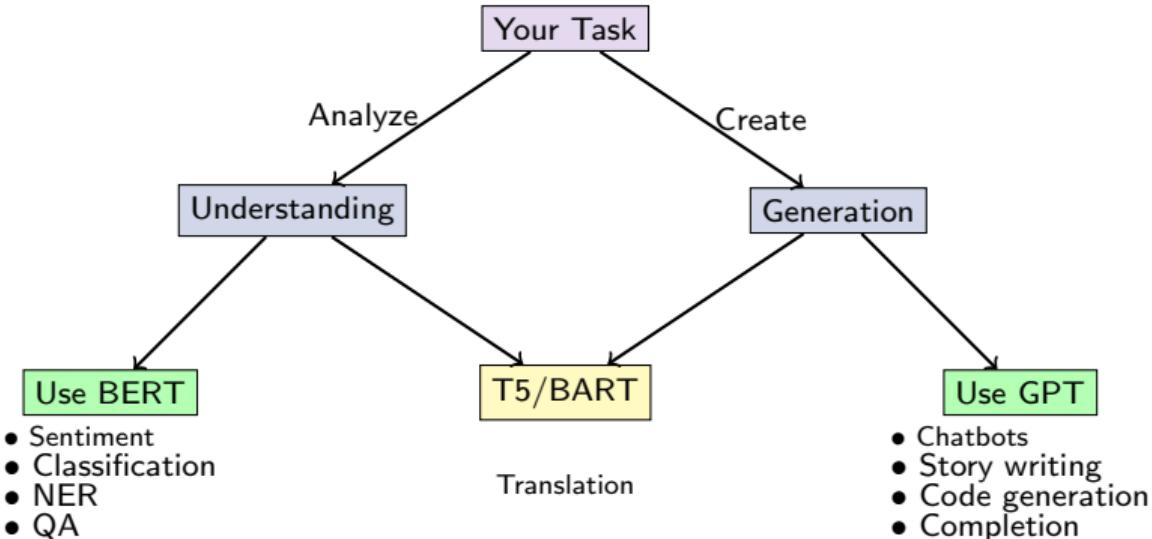
BERT vs GPT: Head-to-Head Comparison

Two paradigms, different strengths:



Choosing the Right Approach

Decision tree for model selection:



Real World: Google uses BERT for search understanding, OpenAI uses GPT for ChatGPT

Test Your Architecture Knowledge

Architecture Quiz:

Q1: Why can't BERT generate text well?

- A) Too small
- B) Bidirectional training
- C) Wrong architecture
- D) Not enough data

Q2: What makes GPT good at generation?

- A) Larger size
- B) Autoregressive training
- C) Better data
- D) Faster inference

Q3: For sentiment analysis, which is better?

Answers:

A1: B - Bidirectional training

- BERT sees future context during training
- Can't generate sequentially
- Designed for understanding, not creation

A2: B - Autoregressive training

- Learns to predict next token
- Natural for sequential generation
- Each output becomes next input

A3: B - BERT understands full context

- Sees entire sentence at once
- Better for classification tasks

Section 4

The Revolution

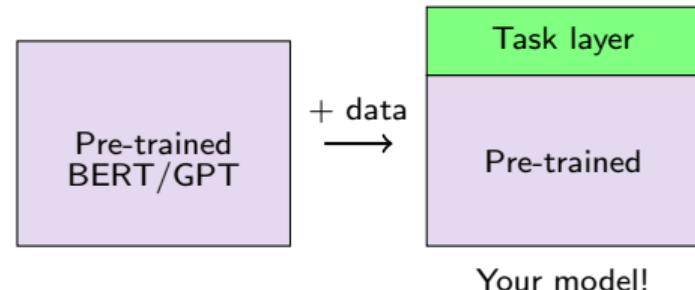
From Pre-training to Your Task

Fine-tuning: Teaching New Tricks

Adapting general knowledge to specific tasks:

The Process:

- ➊ Start with pre-trained model
- ➋ Add task-specific head
- ➌ Train on your labeled data
- ➍ Much less data needed (100s-1000s)
- ➎ 10-100x faster than from scratch



What changes:

- Top layers adapt most
- Lower layers change little
- Task head learns mapping

Example - Sentiment Analysis:

- Pre-trained BERT: Knows language
- Your data: 1000 movie reviews
- Fine-tune: 1 hour on 1 GPU
- Result: 95% accuracy

The Model Zoo: Choose Your Fighter!

Pre-trained models available today:

Model	Size	Type	Best For	Company
BERT-base	110M	Encoder	Classification	Google
BERT-large	340M	Encoder	Understanding	Google
RoBERTa	355M	Encoder	Robust tasks	Facebook
GPT-2	1.5B	Decoder	Generation	OpenAI
GPT-3	175B	Decoder	Few-shot	OpenAI
GPT-4	1T	Decoder	Everything	OpenAI
T5	11B	Enc-Dec	Text-to-text	Google
BART	400M	Enc-Dec	Summarization	Facebook
mT5	13B	Enc-Dec	Multilingual	Google

Real World: HuggingFace hosts 100,000+ pre-trained models ready to use

The model zoo grows daily - new models appear weekly

Using Pre-trained Models in Practice

5 lines of code to state-of-the-art:

```
from transformers import pipeline

# Load pre-trained model
classifier = pipeline('sentiment-analysis')

# Use immediately
result = classifier("I love this movie!")
# Output: [{'label': 'POSITIVE', 'score': 0.999}]

# Or fine-tune on your data
from transformers import AutoModelForSequenceClassification, Trainer

model = AutoModelForSequenceClassification.from_pretrained("bert-base")
trainer = Trainer(model=model, train_dataset=your_data)
trainer.train() # Fine-tunes in hours, not weeks
```

Modern libraries make pre-trained models as easy as importing a package

HuggingFace Transformers is the standard library for pre-trained models

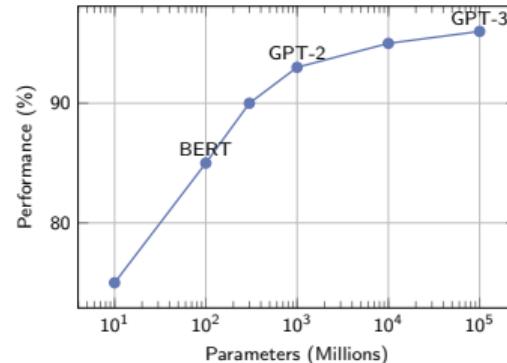
Which Model Should You Use?

Practical decision guide:

Start with these questions:

- ① What's your task?
 - Classification → BERT/RoBERTa
 - Generation → GPT-2/GPT-3
 - Translation → T5/mT5
- ② How much data do you have?
 - <100 examples → GPT-3 few-shot
 - 100-1000 → Fine-tune small model
 - >1000 → Fine-tune large model
- ③ What's your budget?
 - Free → BERT-base
 - \$100s → GPT-2
 - \$1000s → GPT-3 API

Model size vs performance:



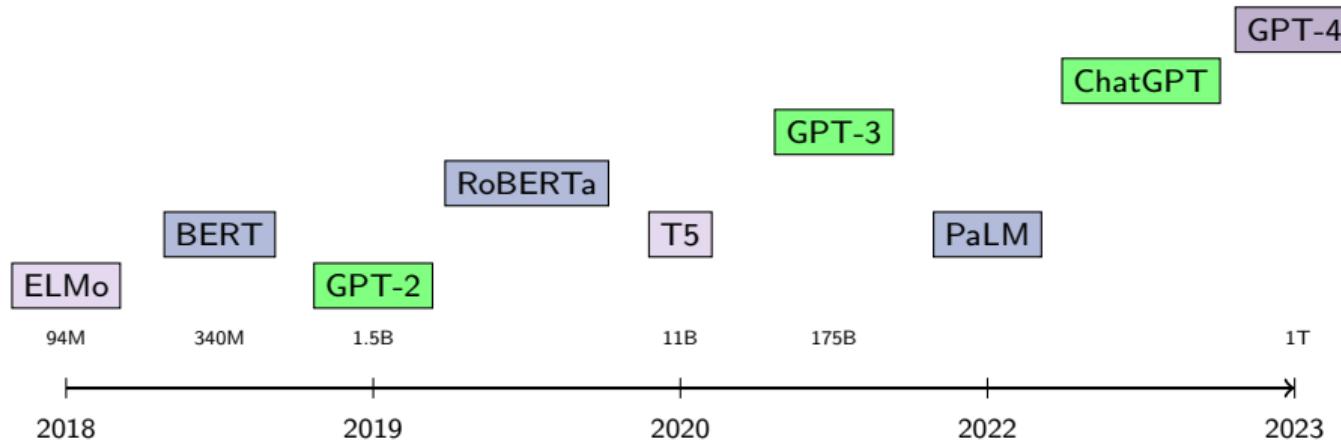
Diminishing returns:

10x size → 2-3% improvement

Common Misconception: “Bigger is always better” - False! BERT-base beats GPT-3 on many specific tasks

The Evolution: From ELMo to GPT-4

How we got here:



Key milestones:

- **2018:** Pre-training idea proven (ELMo, BERT)
- **2019:** Scale begins (GPT-2 “too dangerous to release”)
- **2020:** Few-shot learning (GPT-3)
- **2022:** Instruction following (ChatGPT)

Pre-trained models in production:

Search & QA:

- Google Search (BERT)
- Bing (GPT-4)
- Stack Overflow
- Customer support

10B+ queries daily

Content Creation:

- GitHub Copilot
- Copy.ai
- Jasper
- ChatGPT

\$1B+ market

Understanding:

- Gmail Smart Reply
- Grammarly
- Translation
- Medical diagnosis

1B+ users

Impact: Every text you read online has likely been processed by a pre-trained model

Pre-trained models are the foundation of modern AI applications

The Future: What's Next?

Where pre-training is heading:

Current Trends:

- **Multimodal:** Text + Images + Audio
- **Longer context:** 100K+ tokens
- **Efficiency:** Same performance, 10x smaller
- **Specialization:** Domain-specific models
- **Personalization:** Models that know you

Challenges:

- Computational cost
- Data quality
- Hallucination
- Bias and fairness

Emerging Capabilities:

- Reasoning and planning
- Tool use and APIs
- Self-correction
- Continuous learning
- Uncertainty awareness

Next 5 Years:

- Personal AI assistants
- Scientific discovery
- Creative partners
- Educational tutors
- Medical advisors

The Pre-training Revolution

① **The Problem:** Everyone was learning language from scratch

② **The Solution:** Pre-train once, fine-tune many times

③ **The Methods:**

- BERT: Bidirectional understanding (fill-in-blanks)
- GPT: Autoregressive generation (predict next)

④ **The Impact:**

- 100x faster development
- 10x cost reduction
- Better performance

⑤ **The Practice:** Use HuggingFace, start small, scale if needed

⑥ **The Future:** Multimodal, efficient, personalized

Your journey today:

The Challenge:

- Wasteful repetition
- \$10M problem
- Environmental impact
- Starting from scratch

The Architecture:

- BERT for understanding
- GPT for generation
- Choose based on task
- Both have strengths

The Foundation:

- Self-supervised learning
- Transfer learning theory
- Learning from internet
- Hierarchical knowledge

The Revolution:

- Fine-tuning magic
- Model zoo available
- 5 lines to SOTA
- Changed everything

Thank You!

Questions?

Next: Lab Session - Fine-tuning BERT

We'll fine-tune BERT for sentiment analysis in 30 minutes