

# Natural Language Processing Course

Week 6: Pre-trained Language Models

Joerg R. Osterrieder

[www.joergosterrieder.com](http://www.joergosterrieder.com)

## Week 6

# Pre-trained Language Models

Learning from All of Human Knowledge

# The \$10 Million Problem No One Talked About

**Training a language model from scratch (2018):** Company A trains model for sentiment analysis:

- Cost: \$500,000 in compute<sup>1</sup>
- Time: 2 weeks on 64 GPUs
- Learns: Grammar, syntax, word meanings, sentiment

Company B trains model for question answering:

- Cost: \$500,000 in compute
- Time: 2 weeks on 64 GPUs
- Learns: Grammar, syntax, word meanings... again!

Every team re-learning what "the" means from scratch!

**The waste:** 90% of training teaches same basic language understanding

---

<sup>1</sup>Based on 2018 cloud GPU pricing and typical training times

# The Revolution: Learn Language Once, Use Everywhere

## The breakthrough idea (2018):<sup>2</sup>

What if we:

- ① Train ONE model on massive text (Wikipedia, books, web)
- ② Learn general language understanding
- ③ Share this pre-trained model with everyone
- ④ Fine-tune for specific tasks with little data

## The impact:

- Training cost: \$10M → \$100<sup>3</sup>
- Training time: Weeks → Hours
- Data needed: 1M examples → 1K examples
- Performance: 70% → 95% accuracy

Transfer Learning: Don't start from scratch, start from knowledge!

<sup>1</sup>Howard & Ruder (2018) ULMFiT; Radford et al. (2018) GPT; Devlin et al. (2019) BERT

<sup>2</sup>Fine-tuning costs vs pre-training from scratch

# Pre-trained Models Power Everything (2024)

## Search & Understanding:

- Google Search: BERT since 2019<sup>4</sup>
- Affects 10% of all queries
- Bing: Multiple BERT variants
- Every modern search engine

## Writing Assistants:

- Grammarly: BERT-based
- Google Docs: Smart compose
- Microsoft Editor
- All use pre-trained models

## Business Applications:

- Customer service: 80% automated<sup>5</sup>
- Document analysis
- Email classification
- Resume screening

## Key Models:

- BERT: 110M-340M parameters
- GPT-2: 1.5B parameters
- RoBERTa: Better BERT training
- T5: Unified text-to-text

2024: You never train from scratch - always start from pre-trained

---

<sup>1</sup>Google blog: "Understanding searches better than ever before"

<sup>2</sup>Gartner report on AI automation

## Week 6: What You'll Master

**By the end of this week, you will:**

- **Understand** why pre-training changes everything
- **Master** BERT's bidirectional approach
- **Implement** masked language modeling
- **Compare** BERT vs GPT architectures
- **Fine-tune** models for your own tasks

**Core Insight:** Language understanding is a reusable skill

## Pre-training: Learning from Raw Text

### Traditional supervised learning:

- Need: Labeled data (expensive!)
- Example: 100K sentiment labels
- Problem: Starts from random weights

### Self-supervised pre-training:

- Need: Just text (free and abundant!)
- Example: All of Wikipedia (6B tokens)
- Advantage: Learns language before task

### The clever trick - Create labels from text itself:

- ① Take: "The cat sat on the [?]"
- ② Model predicts: "mat"
- ③ No human labeling needed!
- ④ Can use internet-scale data

Pre-training = Teaching models to read before teaching them tasks

# BERT: Bidirectional Understanding

## The limitation of GPT (left-to-right):

"The man went to the [MASK] to buy milk"

- GPT only sees: "The man went to the"
- Can't use "to buy milk" as context!

## BERT's innovation: Look both ways!<sup>6</sup>

- Sees entire sentence: "The man went to the [MASK] to buy milk"
- Can use both left AND right context
- Predicts: "store" (using "buy milk" as clue)

## Masked Language Model (MLM):


- ① Randomly mask 15% of tokens
- ② Predict masked tokens using all context
- ③ Learn deep bidirectional representations

BERT = Transformer encoder + Bidirectional context

---

<sup>6</sup>Devlin et al. (2019). "BERT: Pre-training of Deep Bidirectional Transformers"

# BERT's Training Process Visualized



## Next Sentence Prediction (NSP)

Sentence A: The man went to the store.

Sentence B: He bought milk and bread.

Label: IsNext

Sentence A: The cat slept all day.

Sentence B: Quantum physics is complex.

Label: NotNext

## Training Data

- BookCorpus: 800M words
- Wikipedia: 2.5B words
- Total: 3.3B words

- Training time: 4 days
- Hardware: 4-16 TPUs
- Parameters: 110M-340M

# GPT: Generative Pre-training

**Different philosophy: Predict the future<sup>7</sup>**

Given: "The cat sat on the"

Predict: "mat" (then "in", then "the", then "sun" ...)

**Autoregressive training:**

- ① Process text left-to-right
- ② Predict next token at each step
- ③ Can generate coherent text!
- ④ Same objective as traditional LM

**Key differences from BERT:**

- Unidirectional (left-to-right only)
- Can generate text naturally
- Simpler training (no masking tricks)
- Transformer decoder architecture

GPT = Transformer decoder + Autoregressive training

---

<sup>7</sup>Radford et al. (2018). "Improving Language Understanding by Generative Pre-Training"

# Implementing BERT's Masked Language Model

```
1 import torch
2 import torch.nn as nn
3 import random
4
5 class BERTMasking:
6     def __init__(self, vocab_size, mask_token_id, mask_prob=0.15):
7         """BERT-style masking for MLM"""
8         self.vocab_size = vocab_size
9         self.mask_token_id = mask_token_id
10        self.mask_prob = mask_prob
11
12    def mask_tokens(self, inputs, special_tokens_mask=None):
13        """Mask tokens for BERT training"""
14        labels = inputs.clone()
15
16        # Create probability matrix
17        probability_matrix = torch.full(labels.shape, self.mask_prob)
18        )
19        if special_tokens_mask is not None:
20            probability_matrix.masked_fill_(special_tokens_mask,
21                value=0.0)
22
23        masked_indices = torch.bernoulli(probability_matrix).bool()
24        labels[~masked_indices] = -100 # Only compute loss on
25        masked
26
27        # 80% of time, replace with [MASK]
28        indices_replaced = torch.bernoulli(
29            torch.full(labels.shape, 0.8)).bool() & masked_indices
30        inputs[indices_replaced] = self.mask_token_id
31
32        # 10% of time, replace with random word
33        indices_random = torch.bernoulli(
34            torch.full(labels.shape, 0.5)).bool() & masked_indices &
35            ~indices_replaced
```

## BERT's 80-10-10 Rule:

- 80%: Replace with [MASK]
- 10%: Replace with random token
- 10%: Keep original token

## Why this complexity?

MASK never seen in fine-tuning

- Random replacement adds noise
- Original tokens prevent mismatch

## Design validates through:

- Ablation studies in paper
- 15% masking is optimal
- Too much masking hurts learning

# BERT Model Architecture

```
1 class BERTModel(nn.Module):
2     def __init__(self, vocab_size, hidden_size=768, num_layers=12,
3                  num_heads=12, max_length=512, dropout=0.1):
4         """BERT model for masked language modeling"""
5         super().__init__()
6
6         # Token embeddings + position + segment
7         self.token_embeddings = nn.Embedding(vocab_size, hidden_size)
8
9         self.position_embeddings = nn.Embedding(max_length,
10                                               hidden_size)
11        self.segment_embeddings = nn.Embedding(2, hidden_size)
12
13        self.layer_norm = nn.LayerNorm(hidden_size)
14        self.dropout = nn.Dropout(dropout)
15
16        # Transformer encoder layers
17        encoder_layer = nn.TransformerEncoderLayer(
18            hidden_size, num_heads, dim_feedforward=4*hidden_size,
19            dropout=dropout, activation='gelu')
20
21        self.transformer = nn.TransformerEncoder(encoder_layer,
22                                              num_layers)
23
24        # MLM head
25        self.mlm_head = nn.Sequential(
26            nn.Linear(hidden_size, hidden_size),
27            nn.GELU(),
28            nn.LayerNorm(hidden_size),
29            nn.Linear(hidden_size, vocab_size)
30        )
31
32
33    def forward(self, input_ids, segment_ids=None, attention_mask=
34               None):
35        """Forward pass for MLM"""
36
```

## BERT Sizes:

- BASE: 12 layers, 768 hidden, 110M params<sup>8</sup>
- LARGE: 24 layers, 1024 hidden, 340M params

## Key Components:

- Segment embeddings for sentence pairs
- GELU activation (smoother than ReLU)
- Layer norm before attention

---


Most applications use BERT-BASE

# Fine-tuning: From General to Specific


## The fine-tuning recipe:

- ① Start with pre-trained BERT/GPT
- ② Add task-specific head (1-2 layers)
- ③ Train on your labeled data
- ④ 100x less data needed!

BERT Fine-tuning: Transfer Learning in Action




# Pre-training Impact: BERT Dominates Every Task



## Key Insights

- GLUE benchmark: 9 language understanding tasks
- BERT improved SOTA on ALL tasks simultaneously
- Average score: 82.1 (previous best: 75.1)
- Some tasks: 10+ point improvements

# The Power of Transfer Learning



## Key insights:

- With 100 examples: BERT matches 10K from-scratch training
- Especially powerful for low-resource languages
- Works across very different tasks
- The more diverse pre-training, the better transfer

# The Pre-trained Model Zoo (2024)

## General Purpose:

- BERT: Understanding tasks
- GPT-2/3: Generation tasks
- T5: Any text-to-text task
- ELECTRA: Efficient BERT alternative

## Domain Specific:

- BioBERT: Biomedical text
- SciBERT: Scientific papers
- FinBERT: Financial documents
- CodeBERT: Programming code

## Multilingual:

- mBERT: 104 languages
- XLM-R: 100 languages
- mT5: Text-to-text multilingual

## Where to find:

- Hugging Face: 500K+ models<sup>9</sup>
- Google TF Hub
- PyTorch Hub
- Company model hubs

2024 Reality: Never train from scratch - there's a model for that!

<sup>9</sup>Hugging Face hosts majority of open models

# Pre-trained Model Pitfalls

## 1. Domain Mismatch

- Problem: BERT trained on books/wiki, you have tweets
- Solution: Domain-adaptive pre-training
- Continue pre-training on your domain

## 2. Catastrophic Forgetting

- Problem: Fine-tuning destroys pre-trained knowledge
- Solution: Lower learning rates ( $2e-5$ ), gradual unfreezing

## 3. Overfitting on Small Data

- Problem: Large model, small dataset
- Solution: Freeze lower layers, only train top layers

### Real Example - Customer Service Bot:

- Started with BERT: 72% accuracy
- Continued pre-training on support tickets: 85%
- Fine-tuned with careful LR: 94%

## Week 6 Exercise: Fine-tune Your Own BERT

**Your Mission:** Take pre-trained BERT and make it expert at your task

### Part 1: Understand Pre-training

- Load pre-trained BERT from Hugging Face
- Examine what it already knows (probe tasks)
- Visualize attention patterns

### Part 2: Fine-tune for Classification

- Choose: Sentiment, spam, or topic classification
- Add classification head to BERT
- Compare: 100 vs 1000 vs 10000 examples
- Track how quickly it learns

### Part 3: Advanced Experiments

- Try different learning rates
- Freeze vs unfreeze layers
- Compare BERT vs training from scratch
- Measure training time savings

**You'll discover:** Why no one trains from scratch anymore!

## Key Takeaways: The Pre-training Revolution

### What we learned:

- Training from scratch wastes resources
- Pre-training learns reusable language understanding
- BERT: Bidirectional with masking
- GPT: Autoregressive generation
- Fine-tuning needs 100x less data

### The paradigm shift:

Task-specific training → Pre-train then fine-tune

### Why it matters:

- Democratized NLP (anyone can use BERT)
- Enabled rapid prototyping
- Made NLP practical for businesses

### Next week: Advanced Transformers

How do we scale to GPT-3's 175B parameters and beyond?

# References and Further Reading

## Foundational Papers:

- Devlin et al. (2019). "BERT: Pre-training of Deep Bidirectional Transformers"
- Radford et al. (2018). "Improving Language Understanding by Generative Pre-Training"
- Liu et al. (2019). "RoBERTa: A Robustly Optimized BERT"

## Practical Resources:

- Hugging Face Transformers library
- "The Illustrated BERT" by Jay Alammar
- Google's BERT GitHub repository

## Recent Advances:

- ELECTRA: More efficient pre-training
- DeBERTa: Disentangled attention
- Transfer learning survey (Qiu et al., 2020)