


The Alignment Problem

Intent ≠ Specification ≠ Behavior



- Helpful but not harmful

- Truthful but not overly verbose

- Creative but not hallucinating