

Natural Language Processing

Week 8: Breaking Words into Pieces

The Tokenization Revolution

NLP Course 2025

Learning Objectives

By the end of this lecture, you will understand:

- ① Why simple word-level tokenization fails (OOV problem)
- ② How subword tokenization preserves meaning for rare words
- ③ The BPE algorithm and why it learns from frequency
- ④ Why WordPiece and SentencePiece improve on BPE
- ⑤ How tokenization impacts multilingual models

Prerequisite

From previous weeks:

- Word embeddings represent words as vectors (Week 2)
- Transformers process sequences of tokens (Week 5)
- Vocabulary size affects model size and training

Table of Contents

- 1 The Word Tokenization Problem
- 2 Byte Pair Encoding (BPE)
- 3 Rare Word Handling
- 4 Vocabulary Size and OOV
- 5 WordPiece and SentencePiece
- 6 Multilingual Tokenization
- 7 Impact on Model Performance
- 8 Practical Considerations
- 9 Summary and Looking Ahead

Act 1: The Out-of-Vocabulary Crisis

Imagine you're building a translator...

You train on common English words: cat, dog, run, happy, etc.

Then a user types:

"I feel **unhappiness** about this situation"

Your model's vocabulary:

Known words:

- happy
- happiness
- sad
- unhappy

Unknown word:

- **unhappiness** → <UNK>
- **All meaning lost!**
- Model has no idea what user meant

The Dilemma: Can't memorize every possible word. English has millions!

Why Word-Level Tokenization Fails

Real-World Application

Real scenario from GPT-2 (2019):

Vocabulary size: 50,257 words

Reddit training data: Millions of rare words

Problem: 5-15% of test words were out-of-vocabulary!

Three approaches to consider:

1. Character-Level

- + Never OOV
- Too many tokens
- "the" = 3 tokens

2. Word-Level

- + Natural units
- High OOV rate
- Huge vocabulary

3. Subword-Level?

- + Best of both?
- + Balanced tokens
- + Rare words OK

Key Question: How do we automatically find meaningful subword units?

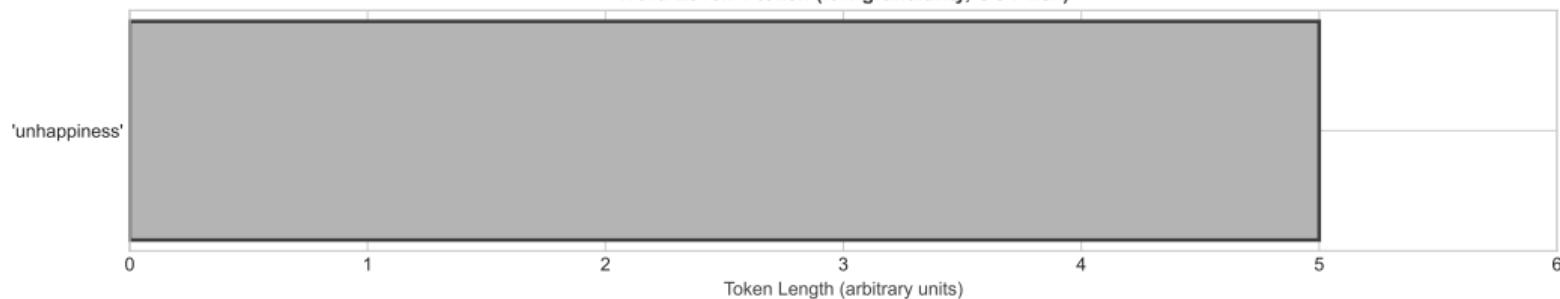
Visual: The Tokenization Spectrum

Tokenization Approaches: Visual Comparison

Character-Level: 11 tokens (high granularity)



Word-Level: 1 token (low granularity, OOV risk)



Subword (BPE): 3 tokens (balanced granularity, meaning preserved)



Act 2: The BPE Solution

Core Insight: Learn subword units from frequency patterns

The Algorithm (simplified):

- ① Start with character vocabulary: {a, b, c, ..., z}
- ② Count all adjacent character pairs in corpus
- ③ Merge the most frequent pair into a new token
- ④ Repeat until vocabulary reaches target size

Common Misconception

"BPE needs linguistic knowledge to find morphemes"

FALSE! BPE is purely statistical. It discovers that "un-", "-ing", "-ness" are common patterns automatically, without knowing they're prefixes/suffixes.

Key Properties:

- Language-agnostic (works for any language)
- Data-driven (learns from your corpus)
- Vocabulary size is a hyperparameter (typically 30K-50K)

BPE Example: Step-by-Step

Let's walk through a tiny example...

Training corpus:

```
low low low  
lower lower  
newest newest newest
```

Initial representation (character + word boundary):

```
l o w </w>  (appears 3 times)  
l o w e r </w>  (appears 2 times)  
n e w e s t </w>  (appears 4 times)
```

Count all pairs:

Pair	Frequency
(e, s)	4
(l, o)	5
(o, w)	5
(n, e)	4

Merge most frequent: (l, o) → lo OR (o, w) → ow

Visual: BPE Merge Detail

BPE Merge Operation: Step-by-Step Example

Step 1: Count all adjacent pairs

low
lower
newest

Step 2: Most frequent pair

'e s': 2
'l o': 2
'o w': 2

Step 3: Merge chosen pair (es)

e s → es

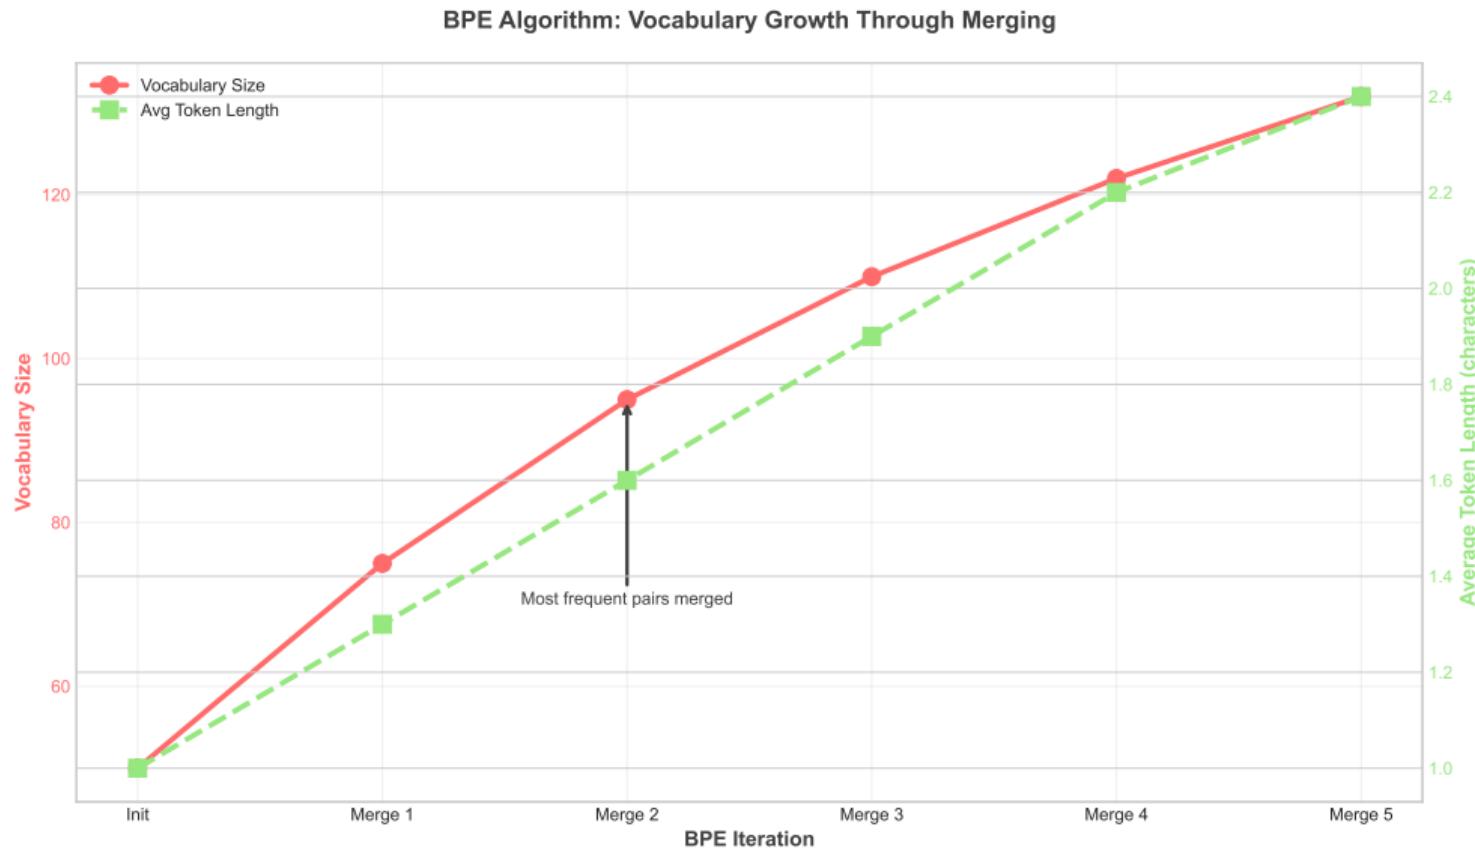
Step 4: Update corpus with new token

low
lower

Key Insight:

BPE learns subword units
from frequency patterns

Visual: BPE Algorithm Progression



Checkpoint 1: Understanding BPE

Checkpoint

Test your understanding:

Question 1: What does BPE merge at each iteration?

- A) Random character pairs
- B) Linguistically meaningful morphemes
- C) Most frequent adjacent pairs
- D) Longest possible substrings

Answer: C - Most frequent adjacent pairs

BPE is purely frequency-based. It doesn't know about linguistics!

Question 2: Why does BPE never produce out-of-vocabulary tokens?

- A) It memorizes all possible words
- B) It can always fall back to character-level
- C) It uses a special UNK token

Answer: B - Falls back to character-level

Since characters are always in vocabulary, any word can be decomposed.

Act 3: The Rare Word Advantage

Remember our “unhappiness” problem?

Let's see how BPE handles it...

After BPE training, vocabulary contains:

- “un” (common prefix in: unable, unhappy, unknown, etc.)
- “happi” (appears in: happy, happiness, unhappy, etc.)
- “ness” (common suffix in: happiness, sadness, kindness, etc.)

At test time:

“unhappiness” → [un] [happi] [ness]

Model can now understand:

- “un” = negation prefix
- “happi” = emotional state (positive)
- “ness” = quality/state suffix
- Combined: negated positive emotional state = negative feeling

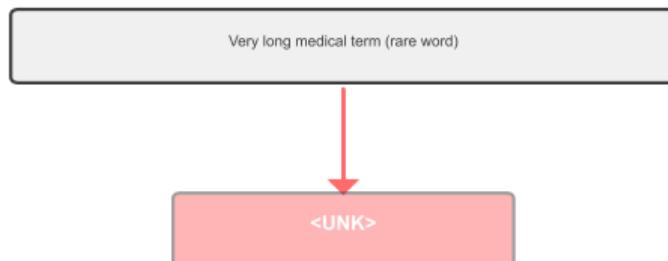
Visual: Rare Word Handling Comparison

Rare Word Handling: Word-Level vs Subword

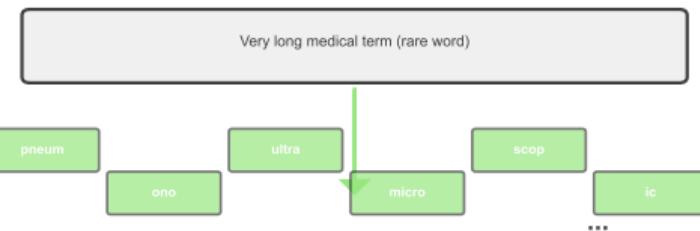
Word-Level: Information Loss

Subword (BPE): Meaning Preserved

Input: "pneumonultramicroscopicsilicovolcanoconiosis"



Input: "pneumonultramicroscopicsilicovolcanoconiosis"



Real-World Application

Medical domain example:

Word: "pneumonultramicroscopicsilicovolcanoconiosis" (lung disease)

The Vocabulary Size Trade-off

How large should our vocabulary be?

Small vocabulary (5K-10K):

- + Fast training
- + Memory efficient
- Many subword splits
- Longer sequences

Large vocabulary (50K-100K):

- + Fewer splits
- + More “whole words”
- Slower training
- More parameters

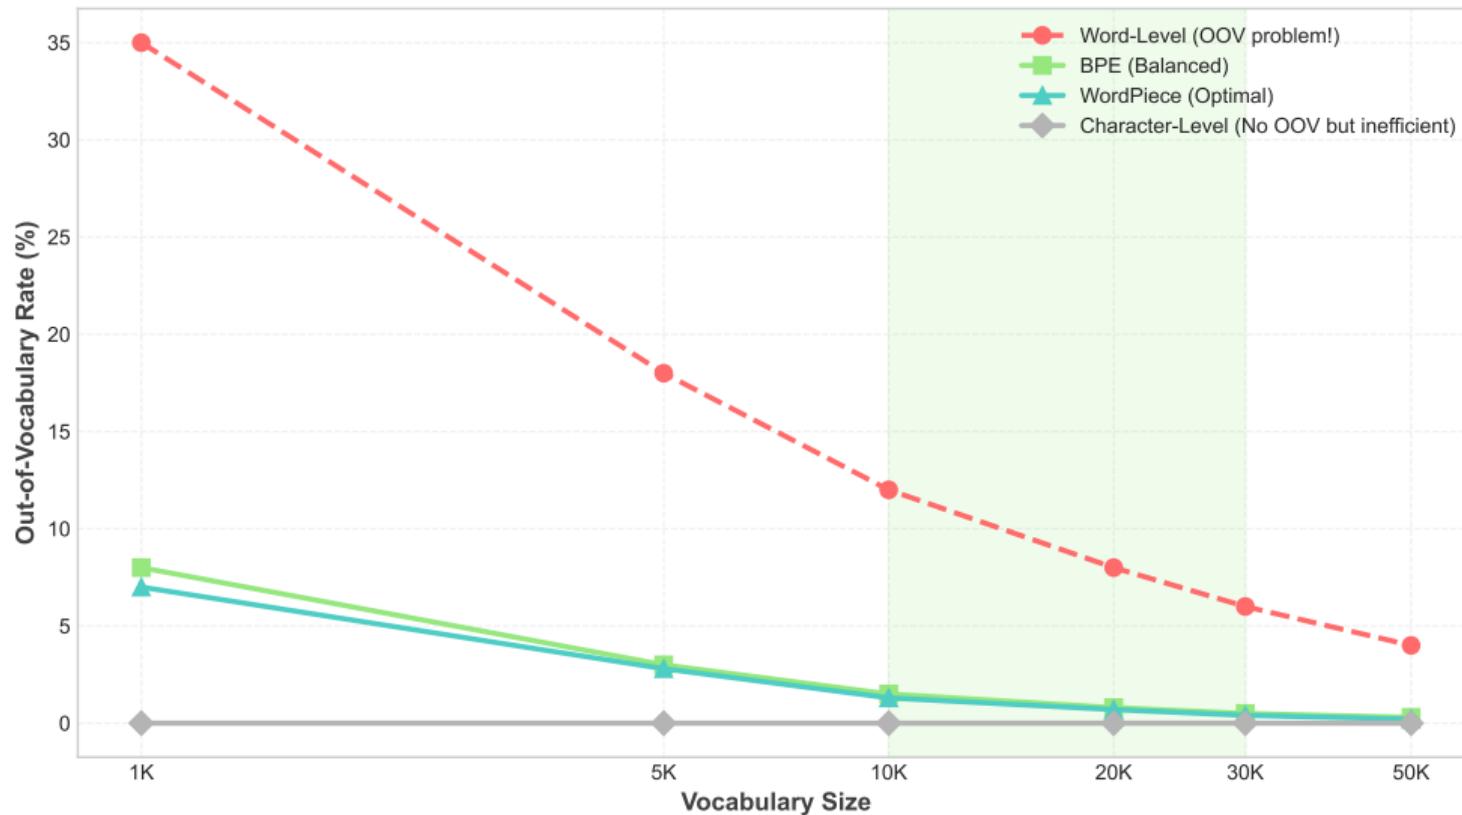
Industry standard: 30K-50K tokens (sweet spot)

Intuition

Think of it like image compression: Higher compression ratio = faster but lower quality. Lower compression = slower but better quality. Subword tokenization finds the optimal balance.

Visual: Vocabulary Size vs OOV Rate

Vocabulary Size vs OOV Rate: Why Subword Tokenization Wins



Act 4: Improvements on BPE

BPE was just the beginning...

1. WordPiece (Google, 2016):

- Used in BERT, T5, many Google models
- **Key difference:** Merges based on likelihood increase (not just frequency)
- Formula: Choose merge that maximizes $P(\text{corpus})$
- **Result:** Slightly better language model perplexity

2. SentencePiece (Google, 2018):

- Works on raw text (no pre-tokenization needed)
- Language-agnostic (handles Chinese, Japanese, Arabic, etc.)
- Treats spaces as tokens (e.g., “_hello”)
- Used in: XLNet, ALBERT, T5, many multilingual models

Common Misconception

“BPE is outdated, everyone uses WordPiece now”

FALSE! GPT-2, GPT-3, GPT-4, RoBERTa, and many others still use BPE. Choice depends on specific use case.

Comparison: BPE vs WordPiece vs SentencePiece

Property	BPE	WordPiece	SentencePiece
Merge criterion	Frequency	Likelihood	Both supported
Pre-tokenization	Required	Required	Not required
Space handling	Special char	Special char	Treated as token
Multilingual	Good	Good	Excellent
Used in	GPT-2/3/4	BERT	XLNet, T5

Real-World Application

Why does GPT-4 use BPE while BERT uses WordPiece?

BPE: Simpler, faster training, good for autoregressive models

WordPiece: Better perplexity, good for masked language modeling

Both work well - choice is often historical/engineering preference

The Multilingual Challenge

Different languages have different characteristics...

English:

- Space-separated words
- Moderate morphology
- “running” = run + ing

German:

- Compound words
- Rich morphology
- “Donaudampfschifffahrtsgesellschaftskapitan”

Chinese:

- No spaces between words
- Character-based writing
- Each character has meaning

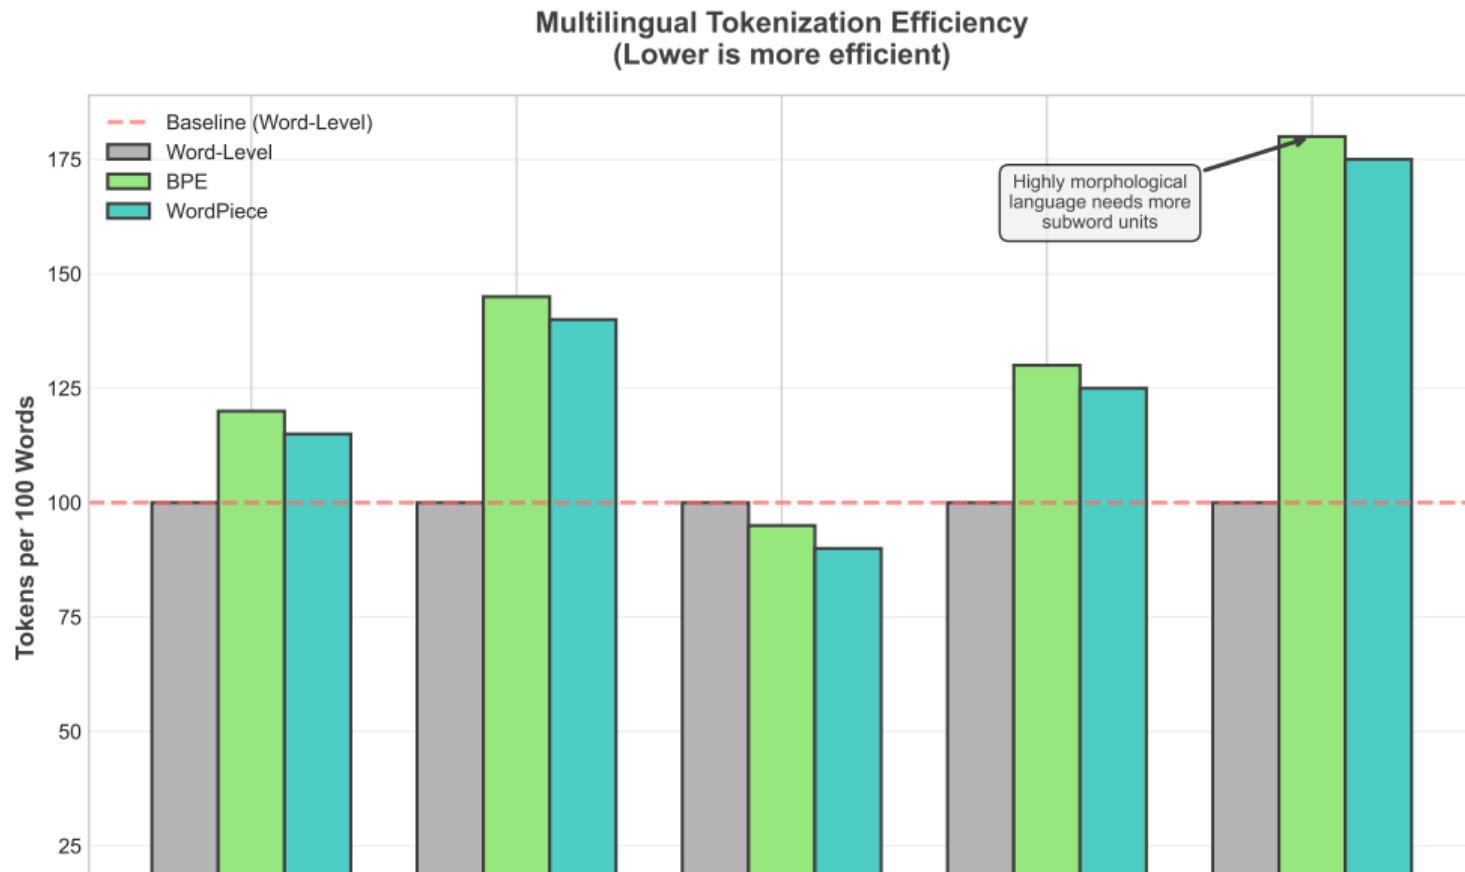
Finnish:

- Extremely complex morphology
- 15 grammatical cases
- “talossanikinko” = “in my house too?”

Challenge: How can one tokenization scheme work for all languages?

Answer: Subword tokenization adapts automatically! Languages with complex morphology get more subword splits, languages with simpler structure get fewer splits.

Visual: Multilingual Tokenization Efficiency



Checkpoint 2: Multilingual Understanding

Checkpoint

Test your understanding:

Question 1: Why does Finnish need more BPE tokens than English for the same text?

- A) Finnish words are longer
- B) Finnish has complex morphology requiring more subword units
- C) BPE doesn't work well for Finnish
- D) Finnish has a larger alphabet

Answer: B - Complex morphology

Finnish words encode grammatical information through suffixes, so need more subword splits. This is a feature, not a bug - the model learns morphological patterns!

Question 2: What's the main advantage of SentencePiece over BPE for Chinese?

- A) It understands Chinese grammar
- B) No pre-tokenization needed (no spaces in Chinese)
- C) It produces shorter sequences

Answer: B - No pre-tokenization needed

Chinese doesn't use spaces between words, so SentencePiece's raw text approach is ideal.

Does Tokenization Really Matter?

Yes! Tokenization significantly impacts model performance...

Real-World Application

Case study: GPT-2 (2019)

OpenAI experimented with different tokenization schemes:

Character-level: Perplexity = 85, Training time = 2.5x

Word-level: Perplexity = 120 (OOV problems!)

BPE (50K vocab): Perplexity = 45, Training time = 1.3x

Result: BPE became standard for GPT series

Three key metrics affected:

- ① **Model quality:** Better tokenization = lower perplexity
- ② **Training efficiency:** Balanced sequence length = faster training
- ③ **Memory usage:** Smaller vocabulary = fewer parameters

Visual: Tokenization Impact on Performance



Key Takeaway: WordPiece and BPE achieve best balance across all three metrics

Implementing BPE in Practice

Good news: You don't need to implement BPE yourself!

Popular libraries:

1. Hugging Face Tokenizers (Fast Rust implementation):

```
from tokenizers import Tokenizer
from tokenizers.models import BPE
from tokenizers.trainers import BpeTrainer

tokenizer = Tokenizer(BPE())
trainer = BpeTrainer(vocab_size=30000, special_tokens=["<PAD>", "<UNK>"])
tokenizer.train(files=["corpus.txt"], trainer=trainer)

# Use it
output = tokenizer.encode("unhappiness")
print(output.tokens) # ['un', 'happi', 'ness']
```

Implementing BPE in Practice (continued)

2. SentencePiece (Google):

```
import sentencepiece as spm

# Train SentencePiece model
spm.SentencePieceTrainer.train(
    input='corpus.txt',
    model_prefix='tokenizer',
    vocab_size=30000,
    model_type='bpe' # or 'unigram'
)

# Load and use
sp = spm.SentencePieceProcessor()
sp.load('tokenizer.model')
tokens = sp.encode_as_pieces("unhappiness")
print(tokens) # ['un', 'happi', 'ness']
```

3. tiktoken (OpenAI - for GPT models):

```
import tiktoken

enc = tiktoken.get_encoding("cl100k_base") # GPT-4 encoding
tokens = enc.encode("unhappiness")
print([enc.decode_single_token_bytes(t) for t in tokens])
```

Common Pitfalls and Best Practices

Common Misconception

"I should train a new tokenizer for every task"

Usually FALSE! Use pre-trained tokenizers when fine-tuning models. Only train new tokenizer if: (1) New language not in pre-trained vocab, or (2) Highly specialized domain (medical, legal) with unique terminology

Best Practices:

- ① **Vocabulary size:** Start with 30K-50K (industry standard)
- ② **Special tokens:** Always include PAD, UNK, BOS, EOS tokens
- ③ **Normalization:** Decide on casing (lowercase vs mixed case)
- ④ **Pre-tokenization:** For English/European languages, split on spaces first
- ⑤ **Corpus representativeness:** Train on data similar to inference data

Real-World Application

Real mistake: A company trained a tokenizer on formal business emails, then deployed it for social media text.
Result: Poor performance due to emoji, slang, abbreviations being split into too many tokens. Lesson: Match training data to use case!

Summary: The Tokenization Journey

What we learned today:

- ① **The Problem:** Word-level tokenization fails due to OOV (out-of-vocabulary)
- ② **The Solution:** Subword tokenization (BPE, WordPiece, SentencePiece)
- ③ **The Algorithm:** BPE iteratively merges most frequent adjacent pairs
- ④ **The Advantage:** Rare words decompose into meaningful subword units
- ⑤ **The Trade-off:** Vocabulary size balances efficiency vs granularity
- ⑥ **The Impact:** Tokenization significantly affects model performance

Key Insight: Tokenization is not a preprocessing afterthought - it's a fundamental design choice that affects model architecture, training, and performance!

Final Checkpoint: Complete Understanding

Checkpoint

Final self-assessment:

Can you explain these to a friend?

- Why does GPT-4 split “unhappiness” into [un][happi][ness]?
- What’s the difference between BPE and WordPiece?
- Why do multilingual models need larger vocabularies?
- How does tokenization affect training speed?

If you can answer these, you understand tokenization!

Challenge: Try encoding “supercalifragilisticexpialidocious” with BPE. How many subword units do you think it would produce? (Answer: Depends on vocabulary, but typically 8-12 units that capture phonetic patterns)

Looking Ahead: Week 9

Now that we have tokens, how do we generate text?

Next week: Decoding Strategies

- Greedy decoding (simple but flawed)
- Beam search (better but still limited)
- Sampling methods (temperature, top-k, top-p)
- Why do chatbots sometimes repeat themselves?
- How to control creativity vs coherence

Real-World Application

Teaser: Have you noticed ChatGPT sometimes gives boring responses and sometimes creative ones? That's decoding strategy at work! Next week we'll learn how to control this.

See you next week!