

## Natural Language Processing Course

Week 8: Tokenization and Subword Models

Joerg R. Osterrieder

[www.joergosterrieder.com](http://www.joergosterrieder.com)

# **Week 8**

# **Tokenization**

The Hidden Foundation of Every LLM

# The Word That Broke Google Translate

2016: A user typed "Pneumonoultramicroscopicsilicovolcanoconiosis"

Result: System crashed.<sup>1</sup>

## The problem:


- English: 170,000 words in current use
  - Medical terms: 100,000+ additional
  - New words daily: "COVID-19", "cryptocurrency", "mansplaining"
  - Misspellings: "recieve", "definately", "occured"
  - Other languages: 7,000+ languages worldwide!

Can't have a token for every possible word - memory explosion

**The dilemma:** How do you handle infinite vocabulary with finite memory

<sup>1</sup>Simplified example - real systems had various OOV issues

# The Tokenization Spectrum: Characters vs Words



# Tokenization Powers Every Modern LLM (2024)

## Model Vocabularies:

- GPT-2: 50,257 tokens
- GPT-3/4: 50,257 tokens
- BERT: 30,522 tokens
- T5: 32,000 tokens
- LLaMA: 32,000 tokens

## Why These Numbers?

- Power of 2 for efficiency
- Covers 99.9% of text
- Balances sequence length
- Works across languages

## Tokenization Methods:

- BPE: GPT family<sup>2</sup>
- WordPiece: BERT family
- SentencePiece: T5, mT5
- Unigram: Some Japanese models

## Critical Impact:

- Wrong tokenization → 50% performance drop
- Affects prompt cost (GPT-4: \$0.01/1K tokens)
- Determines max context length
- Controls multilingual ability

Tokenization is the most important choice you've never heard of

<sup>2</sup>Byte-level BPE specifically for GPT-2 onwards

## Week 8: What You'll Master

**By the end of this week, you will**

- **Understand** why subword tokenization dominates
  - **Implement** Byte Pair Encoding from scratch
  - **Master** the trade-offs in vocabulary design
  - **Analyze** tokenization's impact on different languages
  - **Build** your own tokenizer for any domain

**Core Insight:** Break words into learnable pieces

# The Two Extremes: Characters vs Words

**Example sentence:** "The quick brown fox jumps"

## Character tokenization:

- Tokens: [T, h, e, \_, q, u, i, c, k, \_, b, r, o, w, n, \_, f, o, x, \_, j, u, m, p, s]
- Length: 25 tokens
- Vocabulary: 100 (all ASCII)
- X Handles any text
- X Sequences too long
- X Model must learn spelling

## Word tokenization:

- Tokens: [The, quick, brown, fox, jumps]
- Length: 5 tokens
- Vocabulary: 170,000+ (English)
- X Efficient sequences
- X Out-of-vocabulary words
- X Can't handle typos

Neither extreme works well - we need a middle ground

# Byte Pair Encoding: Learning to Compress

**The brilliant idea (1994):**<sup>3</sup> Compress text by merging frequent pairs

**Example: "low lower lowest"**

- ① Start with characters: l o w \_ l o w e r \_ l o w e s t
- ② Count pairs: (l,o)=3, (o,w)=3, (w,e)=2, (e,r)=1...
- ③ Merge most frequent: "lo"
- ④ New: lo w \_ lo w e r \_ lo w e s t
- ⑤ Repeat: merge "low"
- ⑥ Final: low \_ low er \_ low est

**Result:** Learned meaningful subwords!

- "low" = common root
- "er" = comparative suffix
- "est" = superlative suffix


BPE discovers linguistic structure automatically!

<sup>3</sup>Gage (1994) for compression; Sennrich et al. (2016) for NMT

## BPE in Action: Building a Vocabulary

### Byte Pair Encoding: Learning Subwords


Corpus: "low lower lowest"



BPE discovers:  
• "low" = root  
• "er" = comparative  
• "est" = superlative



Final



# Implementing BPE: The Core Algorithm

```

1 import re
2 from collections import defaultdict, Counter
3
4 class BytePairEncoding:
5     def __init__(self, vocab_size):
6         """Initialize BPE tokenizer"""
7         self.vocab_size = vocab_size
8         self.word_tokenizer = re.compile(r'\w+|[^\\w\\s]')
9         self.vocab = {}
10        self.merges = []
11
12    def get_stats(self, words):
13        """Count frequency of adjacent pairs"""
14        pairs = defaultdict(int)
15        for word, freq in words.items():
16            symbols = word.split()
17            for i in range(len(symbols) - 1):
18                pairs[symbols[i], symbols[i + 1]] += freq
19        return pairs
20
21    def merge_vocab(self, pair, words):
22        """Merge most frequent pair"""
23        bigram = ' '.join(pair)
24        replacement = ','.join(pair)
25        new_words = {}
26        for word in words:
27            new_word = word.replace(bigram, replacement)
28            new_words[new_word] = words[word]
29        return new_words
30
31    def train(self, text, num_merges):
32        """Train BPE on text corpus"""
33        # Split into words
34        words = self.word_tokenizer.findall(text.lower())
35

```

## Algorithm Steps:

- Start with character vocabulary
- Count all adjacent pairs
- Merge most frequent pair
- Repeat until vocab size reached

## Design Choices:

- End-of-word token  $i/w_i$
- Preserves word boundaries
- Case normalization optional
- Typically 10K-50K merges

## Complexity:

- Training:  $O(NM)$  for  $N$  words,  $M$  merges
- Encoding:  $O(L^2)$  for length  $L$
- Memory:  $O(V)$  for vocabulary  $V$

## Tokenizer Comparison: Real Examples

**Text:** "Tokenization is surprisingly important for performance"

**GPT-2 (BPE):** ["Token", "ization", " is", " surprisingly", " important", " for", " performance"]  
7 tokens

**BERT (WordPiece):** ["token", "#**ization**", "is", "surprisingly", "important", "for", "performance"]  
7 tokens (note: ## indicates continuation)


**Character-level:** ["T", "o", "k", "e", "n", "i", "z", "a", "t", "i", "o", "n", " ", ...]  
47 tokens!

### Impact on rare words:

"Pneumonoultramicroscopicsilicovolcanoconiosis" →

- BPE: ["P", "neum", "ono", "ult", "ram", "ic", "ros", "cop", "ic", "sil", "ico", "vol", "can", "oc", "on", "ios", "is"]
- Still handles it! (17 subwords vs 45 characters)


# Tokenization Across Languages



## Key challenges:

- Chinese/Japanese: No spaces between words
- Arabic/Hebrew: Right-to-left scripts
- Korean: Syllable blocks (hangul)

# Tokenization's Hidden Impact



## Key Insights

- Wrong tokenizer: 50% performance drop

# State-of-the-Art Tokenizers (2024)

## Byte-level BPE (GPT family):

- Handles ANY text (even emojis)
- No unknown tokens
- Consistent across languages
- Used by: GPT-2/3/4, RoBERTa

## SentencePiece (T5, mT5):

- Language independent
- Reversible tokenization
- BPE or Unigram options
- Handles 100+ languages

## Recent Innovations:

- **CANINE**: Character-level (2021)
- **ByT5**: Byte-level tokens
- **PIXEL**: Image-based text
- **Charformer**: Learnable tokenization

## Tokenizer Selection:

- English only: BPE (30-50K)
- Multilingual: SentencePiece
- Code: Include special tokens
- Domain-specific: Custom training

2024 trend: Larger vocabularies (100K+) for better multilingual support

# Tokenization Best Practices

## 1. Vocabulary Size Trade-offs:

- Small (8K): Long sequences, better generalization
- Medium (32K): Standard choice, balanced
- Large (100K+): Shorter sequences, more memory

## 2. Special Tokens:

**PAD**, [UNK], [CLS], [SEP], [MASK]

- Domain tokens: [CODE], [MATH], [URL]
- Control tokens: [INST], [/INST]

## 3. Common Pitfalls:

- Tokenizer/model mismatch (50% accuracy drop!)
- Forgetting to handle special characters
- Not preserving whitespace information
- Case sensitivity mismatches

## Real Example - GPT costs:

- "Hello world" = 2 tokens = \$0.00002
- "Chinese text" = 4 tokens = \$0.00004 (2x cost!)
- Emoji = 1-3 tokens depending on tokenizer

# Week 8 Exercise: Build Your Domain Tokenizer

**Your Mission:** Create optimal tokenizer for your domain

## Part 1: Analyze Tokenization Impact

- Compare GPT-2, BERT, T5 tokenizers
- Tokenize: English, code, multilingual text
- Measure compression rates
- Calculate cost differences

## Part 2: Train Custom BPE

- Choose domain: medical, legal, code, etc.
- Collect domain corpus (10MB+)
- Train BPE with different vocab sizes
- Compare with general tokenizers

## Part 3: Optimization Experiments

- Test character vs subword vs word
- Measure model performance impact
- Analyze out-of-vocabulary rates
- Optimize for your use case

**You'll discover:** Why tokenization is make-or-break for LLMs!

## Key Takeaways: The Foundation of Language Models

### What we learned:

- Tokenization solves the infinite vocabulary problem
- Subwords balance efficiency and coverage
- BPE learns meaningful units automatically
- Different tokenizers for different needs
- Critical impact on cost and performance

### The evolution:

Characters → Words → Subwords → Learned tokenization

### Why it matters:

- Determines model capabilities
- Affects all downstream tasks
- Can't fix after training!

### Next week: Decoding Strategies

How do we generate coherent text from token probabilities?

# References and Further Reading

## Foundational Papers:

- Sennrich et al. (2016). "Neural Machine Translation of Rare Words with Subword Units"
- Kudo & Richardson (2018). "SentencePiece: Language-independent subword tokenizer"
- Schuster & Nakajima (2012). "Japanese and Korean voice search" (WordPiece)

## Recent Advances:

- Clark et al. (2021). "CANINE: Character-level transformers"
- Xue et al. (2021). "ByT5: Token-free transformers"
- Tay et al. (2021). "Charformer: Fast character transformers"

## Practical Resources:

- Hugging Face Tokenizers library
- Google SentencePiece
- OpenAI tiktoken library