

Contrastive Search: How It Works

Step 1: Get top-

- city: 0.45
- town: 0.18
- area: 0.15
- place: 0.12

Step 2: Adjust by blurring out probability ($\alpha=0.6$)

$$\text{score} = (1-\alpha) \times P(\text{token}) - \alpha \times \text{similarity}$$

$$\begin{aligned} \text{city: } & 0.4 \times 0.45 - 0.6 \times 0.92 = 0.372 \\ \text{town: } & 0.4 \times 0.18 - 0.6 \times 0.75 = 0.0378 \\ \text{area: } & 0.4 \times 0.15 - 0.6 \times 0.65 = 0.3302 \\ \text{place: } & 0.4 \times 0.12 - 0.6 \times 0.50 = 0.312 \end{aligned}$$

Step 2: Compute similarity to recent context

- Context: "the Winner"
- city: 0.92 (cosine similarity)
 - town: 0.75 (cosine similarity)
 - area: 0.65 (cosine similarity)
 - place: 0.60 (cosine similarity)

Winner: "town" (high prob, lower similarity)

Key Insight: Balance probability (coherence) with diversity (novelty)
 $\alpha=0$: Pure greedy | $\alpha=0.6$: Balanced | $\alpha=1.0$: Maximum diversity



Quantlet
Where Knowledge Meets Code
fostering transparency, replicability, and
integrity in scientific exploration