

# Decoding Strategies

Week 9: From Prediction to Generation

NLP Course 2025


Enhanced with Contrastive Search (2025)

November 9, 2025

**Two-Tier BSc Discovery**

Main: 25 slides — Appendix: 19 slides

## From Prediction to Generation: The Decoding Challenge



### Our Journey:

1. We trained models (Weeks 3-7: RNN → Transformers → BERT/GPT)
2. They learned to predict:  $P(\text{word}|\text{context})$
3. They output probability distributions over 50,000+ words
4. **Today:** How do we convert these probabilities into actual text?

---

Models predict probabilities. Decoding converts probabilities to text.

# Today's Challenge: From Probabilities to Text

**The Setup:** Model gives us probabilities for next word

Example: "The cat ..."

$$P(\text{sat}) = 0.45, \quad P(\text{is}) = 0.30, \quad P(\text{jumped}) = 0.25, \quad \dots \quad (50,000 \text{ words})$$

**Naive Approach 1:** Pick highest  
→ Greedy Decoding

**Result:**

"The city is a major city in the city..."

**Problem:** Repetitive, boring

**Naive Approach 2:** Pick random  
→ Pure Sampling

**Result:**

"The purple flying mathematics..."


**Problem:** Nonsense, incoherent

**Core Challenge:** Need text that is *coherent* AND *creative*

---

Today: 6 strategies to solve this - from simple to sophisticated

## The Decoding Paradox: Quality vs Diversity Tradeoff



**Discovery Question:** Why is best text boring and creative text nonsense?

The central challenge: How to balance coherence with creativity

# Three Decoding Families

## Deterministic

### Methods:

- Greedy
- Beam search

### Traits:

- Same output
- High quality
- No creativity

## Stochastic

### Methods:

- Temperature
- Top-k
- Nucleus

### Traits:

- Random
- Creative
- Variable quality

## Balanced

### Methods:

- Contrastive
- Hybrid

### Traits:

- Best of both
- No repetition
- Modern std

Different tasks need different strategies - no single best method

**6 methods total: Each optimizes different quality-diversity tradeoff**

# Greedy Decoding: The Baseline

## How It Works:

1. Compute probabilities
2. Pick highest probability
3. Add to sequence
4. Repeat until done

## Example:

$P(\text{cat}) = 0.45 \leftarrow \text{Pick!}$

$P(\text{dog}) = 0.30$

$P(\text{bird}) = 0.25$

## When to Use:

- Code generation
- Short responses
- Speed critical
- Reproducibility

## Problem: Degeneration


"The city is a city in a city..."

Model gets stuck in loops

---

Simplest but prone to repetition

## Beam Search: Exploring Multiple Paths




**Key Insight:** Keep top-k hypotheses at each step to find better sequences

Beam width = 3-5 typical. Balance greedy (1) vs exhaustive ( $\infty$ )

## Beam Search Example: "The cat..." (Width=3)

START: "The cat"



Worked example: "The cat..." with width=3 shows pruning in action

# Beam Search: Detail

## Algorithm:

1. Start: Keep top-k tokens
2. Expand: All continuations
3. Score: Multiply probs
4. Prune: Keep top-k sequences
5. Repeat until END

## Scoring:

$$\text{score} = \prod_{i=1}^t P(y_i | y_{<i})$$

With length normalization:

$$\text{score} = \frac{1}{t} \sum_{i=1}^t \log P(y_i | y_{<i})$$

## Best For:

- Machine translation
- Summarization
- Question answering
- "Correct" answer tasks

## Parameters:

Width = 3-5 (translation)


Width = 10 (more diverse)

- + Better than greedy
- + Diverse hypotheses
- Still deterministic
- 4-5× slower

---

Beam search: workhorse for deterministic tasks

### Temperature Effects on Probability Distribution



**Key Insight:** Temperature reshapes probability distribution

$T < 1$ : focused.  $T = 1$ : unchanged.  $T > 1$ : random

## Temperature Calculation: Step-by-Step

Given: Logits = [2.0, 1.0, 0.5, 0.2]  
 Tokens = ["cat", "dog", "bird", "fish"]

**Step 1: Scale by T=0.5 (MORE FOCUSED)**  
 Scaled =  $\frac{[2.0, 1.0, 0.5, 0.2]}{0.5} = [4.0, 2.0, 1.0, 0.4]$   
 Softmax =  $\frac{e^{[4.0, 2.0, 1.0, 0.4]}}{\sum e^{[4.0, 2.0, 1.0, 0.4]}} = [0.537, 0.19, 0.12, 0.10]$   
 → 73% on "cat" (VERY FOCUSED)

**Step 2: Scale by T=2.0 (FLATTER)**  
 Scaled =  $\frac{[4.0, 2.0, 1.0, 0.4]}{2.0} = [2.0, 1.0, 0.5, 0.2]$   
 Softmax =  $\frac{e^{[2.0, 1.0, 0.5, 0.2]}}{\sum e^{[2.0, 1.0, 0.5, 0.2]}} = [0.5, 0.17, 0.15]$   
 → 50% on "cat" (BALANCED)  
 → 25% on each token (FLATTER)

Lower  $T \rightarrow$  More confident (peaky)  
 Higher  $T \rightarrow$  More random (flat)

Concrete numbers show how temperature scaling works

# Temperature: Detail

## How It Works:

Given logits  $z_1, \dots, z_V$

Scale by temperature  $T$ :

$$p_i = \frac{\exp(z_i/T)}{\sum_j \exp(z_j/T)}$$

Sample from  $p$

## Effect:

$T \rightarrow 0$ : argmax

$T = 1$ : standard softmax

$T \rightarrow \infty$ : uniform

## Practical Settings:

- $T = 0.1\text{-}0.3$ : Factual Q&A
- $T = 0.7$ : Chatbots
- $T = 0.9\text{-}1.2$ : Creative
- $T = 1.5+$ : Experimental

+ Simple to implement


+ Continuous control

- No quality guarantee

- Can be nonsense

---

Temperature: simplest randomness control



**Key Insight:** Only sample from top-k most likely tokens

Prevents sampling from long tail of unlikely words

### Top-k Example: k=3

~~12.000000000000001~~ ~~Sampling Probabilities:~~  
~~cat: 0.45~~  
~~dog: 0.23~~  
~~bird: 0.19~~  
mouse: 0.08  
.... 0.04

$$0.45 + 0.23 + 0.19 = 0.87$$
$$0.87 / 3 = 0.29$$
$$0.29 \times 100 = 29\%$$

**Result: Sample from {cat: 58%, dog: 23%, bird: 19%}**

*Prevents sampling from long tail ("mouse" eliminated)*

---

Filtering + renormalization prevents tail sampling

# Top-k: Detail

## Algorithm:

1. Compute  $P(w_i)$  for all
2. Sort by probability
3. Keep only top-k
4. Renormalize
5. Sample from  $p'$

## Example (k=3):

Original: [0.45, 0.18, 0.15, ...]  
Keep: [0.45, 0.18, 0.15]  
Renorm: [0.58, 0.23, 0.19]

## Typical Values:

- k = 40-50 (balanced)
- k = 10-20 (focused)
- k = 100+ (very diverse)

## Limitation:


- Fixed cutoff regardless of distribution!
- Peaked: Wastes mass
- Flat: Too many bad tokens

## Solution: Dynamic cutoff

---

Top-k improves over temperature but inflexible



## Nucleus (Top-p) Sampling: Dynamic Cutoff



Key Insight: Adapt vocabulary size to distribution shape

Nucleus size grows/shrinks based on probability spread

## Nucleus Adapts to Distribution Shape ( $p=0.85$ )



Same p value → different vocabulary for peaked vs flat

# Nucleus (Top-p): Detail

## Algorithm:

1. Sort tokens by  $P(w_i)$
2. Compute cumulative sum
3. Find smallest set where  $\text{cum} \geq p$
4. Sample from nucleus

Example ( $p=0.85$ ):

Cumsum: [0.40, 0.60, 0.75, 0.85]

Nucleus: First 4 tokens

## Recommended:

- $p = 0.85\text{-}0.90$ : Dialogue
- $p = 0.90\text{-}0.95$ : Creative
- $p = 0.95\text{-}0.99$ : Very diverse

## Why Better:

Peaked → small nucleus

Flat → large nucleus

Adapts automatically!

## Current Standard

ChatGPT, Claude, GPT-4

---

Nucleus: modern standard for high-quality generation

## The Degeneration Problem: Model Repetition

### Real Output from Greedy Decoding:

"The city of New York is a major city in the United States. The city is known for its diverse culture and the city has many tourist attractions.

The city is also home to the city's financial district..."

**Problem: "the city" appears 6 times in 4 sentences!**

Why? Always picking argmax → same patterns repeated

**Solution: Penalize tokens similar to recent context (Contrastive Search)**

**Discovery:** Why do models repeat themselves?

---

Greedy/beam maximize probability - but high prob = repeating recent context

## Contrastive Search: How It Works

### Step 1: Get top

• city: 0.45  
• town: 0.18  
• place: 0.12

score =  $(1-\alpha) \times P(\text{token}) - \alpha \times \text{similarity}$

city:  $0.4 \times 0.45 - 0.6 \times 0.92 = -0.372$   
town:  $0.2 \times 0.18 - 0.6 \times 0.75 = -0.578$   
place:  $0.1 \times 0.12 - 0.6 \times 0.65 = -0.312$

### Step 2: Compute similarity to recent context

Context: "the **winner** is **town**" (high prob, lower similarity)

• city: 0.92 (cosine similarity)  
• town: 0.75 (cosine similarity)  
• place: 0.60 (cosine similarity)

*Key Insight: Balance probability (coherence) with diversity (novelty)*

$\alpha=0$ : Pure greedy |  $\alpha=0.6$ : Balanced |  $\alpha=1.0$ : Maximum diversity

**Key Insight:** Balance probability with diversity penalty

---

Explicitly avoid copying recent context - prevents degeneration

## Contrastive vs Nucleus: Comparison

### Same Prompt, Different Methods

Prompt: "The future of artificial intelligence is"

"...is pro  
many in  
significa  
educa


"...is rapidly evolving, bringing  
unprecedented opportunities across  
sectors ranging from medicine to  
climate science, while raising  
important ethical questions."

+ Diverse  
+ Creative  
+ No repetition

Contrastive Search explicitly penalizes copying recent context

---

Contrastive prevents repetition better for long generation



Pareto Frontier: No method dominates all others

Choose based on task: deterministic (left), creative (right)

### Choosing the Right Decoding Method

START: What kind of task?

Newspaper  
Translation  
Q&A  
News  
Translation  
Q&A  
Newspaper  
Translation  
Q&A

Creative Writing  
Dialogue  
Storytelling  
Short response generation?  
Creative Writing  
Dialogue  
Storytelling  
Short response generation?

→ BEAM SEARCH  
(width=3-5)  
*Or Temp=0.3-0.6*

*for consistency*

→ N → CONTRASTIVE  
( $p=0.5$ ) ( $\alpha=0.5-0.7$ )

**Special: Code Generation**  
→ Greedy or Beam (correctness critical)  
→ *Then verify syntax/semantics*

Start with task requirements, follow tree to recommended method


# Task-Specific Decoding Recommendations (2025)

Task	Recommended Decoder	Parameters	Why?
Machine Translation	Beam Search	width=3-5	Deterministic, quality critical
Factual Q&A	Greedy / Low Temp	T=0.1-0.3	Single correct answer needed
Summarization	Beam Search	width=4	Balance coverage + conciseness
Code Generation	Greedy	T=0	Syntax errors costly
Creative Writing	Nucleus / Contrastive	p=0.9, $\alpha$ =0.6	Diverse but coherent
Dialogue Systems	Nucleus	p=0.85-0.95	Natural variation needed
Story Generation	Contrastive	$\alpha$ =0.5-0.7	Avoid repetition in long text
Long-form Articles	Contrastive	$\alpha$ =0.6, p=0.9	Degeneration prevention

Task-Specific Tasks

Comprehensive mapping: 8 tasks → optimal strategies

# Computational Costs Matter



Contrastive: best quality-diversity but 12x slower. **Nucleus: best balanced**

Production tradeoff: Quality vs speed vs diversity

## Key Takeaways

1. **Deterministic** (Greedy, Beam): High quality, no diversity → factual tasks
2. **Temperature**: Simple randomness control → universal but crude
3. **Top-k**: Fixed vocabulary filter → prevents tail sampling
4. **Nucleus (Top-p)**: Dynamic cutoff → modern standard, adapts
5. **Contrastive (NEW)**: Degeneration prevention → long creative text
6. **Task matters**: Translation → Beam — Dialogue → Nucleus — Stories → Contrastive

**Next:** Lab - Implement all 6 methods, measure quality-diversity tradeoffs

---

Decoding strategy matters as much as model architecture

# Technical Appendix

19 slides: Complete mathematical treatment

A1-A5: Beam Search Mathematics

A6-A10: Sampling Mathematics

A11-A14: Contrastive Search & Degeneration

A15-A19: Advanced Topics & Production

## A1: Beam Search Formulation

**Objective:** Find  $y^* = \operatorname{argmax} P(y|x)$

**Decomposition:**

$$P(y|x) = \prod_{t=1}^T P(y_t|y_{<t}, x)$$

**Log-probability:**

$$\log P(y|x) = \sum_{t=1}^T \log P(y_t|y_{<t}, x)$$

**Beam Search Approximation:**

Instead of  $V^T$  sequences, maintain top-k hypotheses

**Complexity:**


Time:  $O(k \cdot V \cdot T)$

Space:  $O(k \cdot T)$

---

Tractable approximation to exact search

## A2: Length Normalization



**Solution:** Normalize by length

$$\text{score}(y) = \frac{1}{|y|^\alpha} \log P(y) \quad \text{where } \alpha \in [0.5, 1.0]$$

---

Length normalization essential for fair comparison

## A3: Beam Search Variants

### Diverse Beam:

Partition beams into groups  
Penalize within-group similarity

### Constrained Beam:

Force certain tokens  
Keywords, entities

### Stochastic Beam:

Sample beams (not argmax)  
More diverse

### Block n-gram:

Penalize n-gram repetition  
"city is a city" prevention

---

Many variants for specific requirements

## A4: Beam Search Stopping

### When to stop?

1. **Fixed length:** Stop at  $T_{\max}$  (simple but rigid)
2. **END token:** Special token (most common)
3. **Score threshold:** When best cannot improve
4. **Timeout:** Computational budget

---

Stopping criterion affects output length

## A5: Beam Search Limitations

### Fundamental Issues:

1. **Exposure bias:** Train with teacher forcing, test with own outputs
2. **Label bias:** Cannot compare different prefixes fairly
3. **Repetition:** Still can loop
4. **Bland outputs:** Maximizes probability  $\neq$  interestingness
5. **Search errors:** May miss better sequences

### When Beam Fails:

Open-ended generation, long-form text, creative tasks

→ Need sampling-based methods

---

Beam optimizes wrong objective for creative tasks

## A6: Sampling as Inference

**Goal:** Sample  $y \sim P(y|x)$  not  $\text{argmax } P(y|x)$

**Ancestral Sampling:**

For  $t = 1$  to  $T$ :

    Compute  $P(y_t|y_{<t}, x)$

    Sample  $y_t \sim P(\cdot|y_{<t}, x)$

**Variants:**

Temperature: Reshape before sample

Top-k: Truncate before sample

Nucleus: Dynamic truncation

---

**Sampling enables diversity but loses quality guarantees**

## A7: Temperature Mathematics

**Softmax with Temperature:**

$$p_i(T) = \frac{\exp(z_i/T)}{\sum_{j=1}^V \exp(z_j/T)}$$

**Limiting Cases:**

$T \rightarrow 0$ : deterministic (greedy)

$T \rightarrow \infty$ : uniform ( $1/V$ )

**Entropy Analysis:**

$H(p) = -\sum p_i \log p_i$  measures randomness

$H$  increases with  $T$

Low  $T$  (<0.5):  $H \approx 0$

High  $T$  (>2.0):  $H \approx \log V$

---

Temperature controls distribution entropy

## A8: Top-k Mathematics

### Formal Definition:

$$V_k = \{w_{\sigma(1)}, \dots, w_{\sigma(k)}\}$$

(top k by prob)

Truncated distribution:

$$p'(w) = \begin{cases} \frac{p(w)}{\sum_{w' \in V_k} p(w')} & w \in V_k \\ 0 & \text{otherwise} \end{cases}$$

### Information Loss:

$H(p') < H(p)$  (entropy decreases)

Loss  $\approx$  tail information

---

Top-k sacrifices tail probability for quality

## Formal Definition:

$$V_p = \min \left\{ V' : \sum_{w \in V'} p(w) \geq p \right\}$$

Smallest set with cumulative mass  $\geq p$

## Dynamic Size:

$$|V_p| = \min \{k : \sum_{i=1}^k p_{\sigma(i)} \geq p\}$$

Peaked: Small  $|V_p|$  (2-5)

Flat: Large  $|V_p|$  (50+)

**Why Nucleus > Top-k:**

Top-k: Fixed  $k$

Nucleus: Adapts to distribution

---

Nucleus adjusts vocabulary to distribution

# How to Measure Decoding Quality

## Quality Me Diversit Metrics

- Perplexity (Distinctiveness of unique n-grams)
- BLEU score (Success rate = more diverse)
- ROUGE score (summarization rate)
- Human evaluation (vocabulary richness)
- Grammaticality scope of outputs

## Combined Quality-Diversity Metrics

- Task success rate (does it solve the task?)
- Human preference (A/B testing)
- Pareto frontier analysis (multi-objective)

Need both quality AND diversity metrics

# A11: Degeneration Problem (Formal)

**Definition:** Unnatural repetitions in generated text

**Why It Happens:**

1. Model trained on natural text (low repetition)
2. Generation maximizes  $P(y_t|y_{<t})$
3. Recent context influences  $P$
4. Feedback loop: high prob → context → same high prob

**Quantifying:**

Greedy repetition: 15-30%

Human text: 2-5%

Gap = degeneration

**Examples:** “*I think that I think that I think...*”

---

Maximizing probability  $\neq$  natural text

## A12: Contrastive Search Objective

### Scoring Function:

$$\text{score}(w_t) = (1 - \alpha) \times P(w_t | y_{<t}) - \alpha \times \max_{w_i \in y_{<t}} \text{sim}(w_t, w_i)$$

where  $\alpha \in [0, 1]$  controls tradeoff

**Similarity:** Cosine similarity using token embeddings

$$\text{sim}(w_i, w_j) = \frac{h_i \cdot h_j}{\|h_i\| \cdot \|h_j\|}$$

### Algorithm:

1. Get top-k by probability
2. Compute similarity to all  $y_{<t}$
3. Apply penalty:  $\text{score} = \text{prob} - \alpha \times \text{max\_sim}$
4. Select highest score

---

### Explicit penalty for copying context

## A13: Contrastive Parameters

### Alpha ( $\alpha$ ):

- $\alpha = 0$ : Pure greedy
- $\alpha = 0.6$ : Balanced (default)
- $\alpha = 1.0$ : Max diversity

### Typical Settings:

- Short (<100):  $\alpha = 0.4 - 0.5$
- Medium (<500):  $\alpha = 0.5 - 0.6$
- Long (500+):  $\alpha = 0.6 - 0.7$

### Top-k (candidates):

- $k = 4$ : Fast, focused
- $k = 6$ : Balanced (default)
- $k = 10$ : Diverse

**Cost:**  $O(k \times T^2)$

12× slower than greedy  
Only for quality-critical

---

Hugging Face default:  $\alpha=0.6$ ,  $k=4$

### Research Findings:

- Greedy: 18-25% repetition (GPT-2), 12-18% (GPT-3)
- Nucleus: 8-12% (still above human 3-5%)
- Contrastive: 4-7% (closest to human)

### Why Probability Fails:

Training: Next token prediction

Generation: Global coherence

Mismatch: Local optimum  $\neq$  global quality

### Solutions Hierarchy:

1. Temp/Top-k/Nucleus: Reduce greedy determinism
2. Contrastive: Explicit penalty
3. RLHF/DPO: Align with humans (Week 10)

---

Contrastive addresses fundamental limitation

## A15: Hybrid Methods

**Combining Strategies:**

**Nucleus + Temperature:**

$$p_i(T) = \text{softmax}(z/T), \text{ then } V_p \leftarrow \text{nucleus}(p_i)$$

Used by GPT-3, ChatGPT

**Beam + Sampling:**

Beam search + stochastic selection

**Contrastive + Nucleus:**

Best of both worlds

---

Hybrid methods leverage complementary strengths

# A16: Constrained Decoding (2025)

**Goal:** Force tokens/patterns

**Lexically Constrained:**

Must include keywords

Beam variant tracks satisfaction

**Format Constraints:**

JSON output, code syntax

**Use Cases:**

Structured extraction, controllable summarization, code generation

---

**Constrained decoding enables controllable generation**

## A17: Computational Complexity

Method	Time/token	Total	Speed
Greedy	$O(V)$	$O(VT)$	1.0×
Temperature	$O(V)$	$O(VT)$	1.1×
Top-k	$O(V)$	$O(VT)$	1.2×
Nucleus	$O(V \log V)$	$O(V \log V \cdot T)$	1.3×
Beam (k=5)	$O(kV)$	$O(kVT)$	4.5×
Contrastive	$O(kT)$	$O(kT^2)$	12×

**Practical (1000 tokens):**

Greedy: 2.5s — Nucleus: 3.2s — Beam: 11s — Contrastive: 30s

---

**Computational cost matters for production**

## Production Decoding Settings (Real Systems 2024-2025)

System (2024-2 Method)	Parameters	Goal
GPT-3 API (2024)	Nucleus	T=0.7, p=1.0 Balanced default
ChatGPT	Nucleus + Temp	T=0.8, p=0.95 Creative but controlled
Google Translate	Beam Search	width=4 Quality critical
GitHub Copilot	Greedy	T=0 Code correctness
Claude	Nucleus	T=1.0, p=0.9 High quality generation
Hugging Face Diffa	Greedy	T=1.0 Deterministic baseline

---

Real settings from major production systems

## A19: Future Directions (2025)

### Active Research:

1. Quality-diversity optimization
2. Learned decoding (RLHF, DPO)
3. Speculative decoding (4–8× faster)
4. Adaptive methods
5. Energy-based decoding

### Open Problems:

- Auto-select best parameters for new task?
- Balance fluency + factuality + creativity?
- Efficient 100K+ token decoding?

Trend: Hand-tuned → learned decoding strategies

---

Active research area with many open questions