

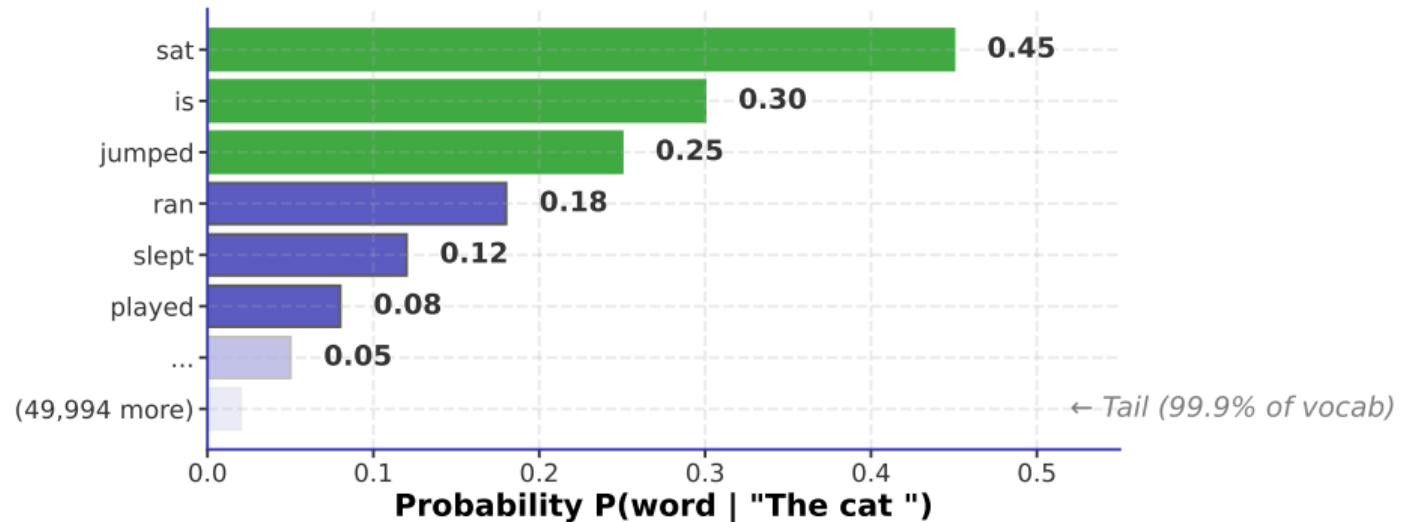
# **Decoding Strategies**

Week 9: From Probabilities to Text

November 2025

## The Decoding Challenge: Choosing From 50,000 Words

### The Decoding Challenge: Choose From 50,000 Words

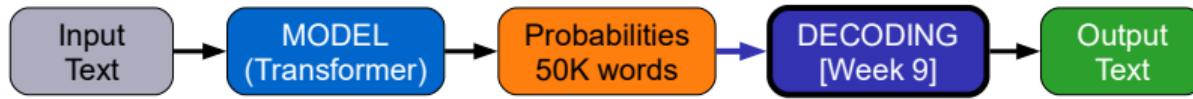


**The Question:** Given these probabilities for “The cat \_\_”, which word should we pick?

---

At each step, model outputs probability distribution over entire vocabulary - how do we choose?

## Context: How We Got Here



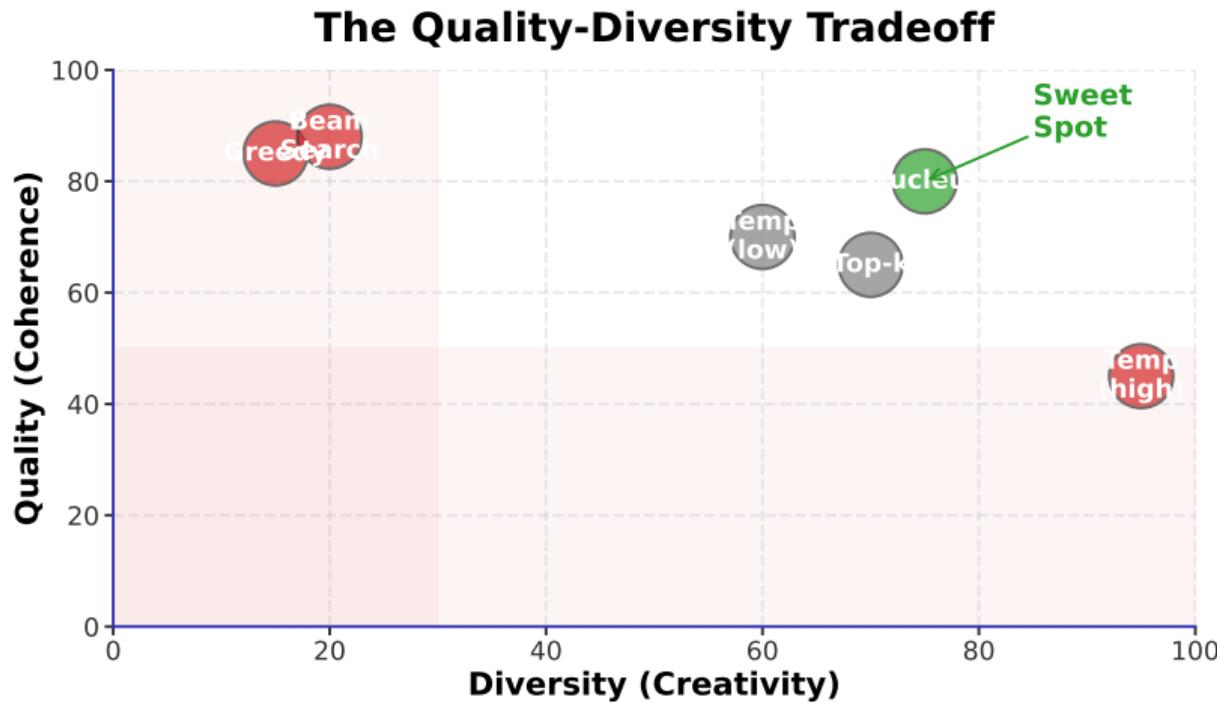
### Our Journey:

1. We trained models (Weeks 3-7: RNN → Transformers → BERT/GPT)
2. They learned to predict:  $P(\text{word}|\text{context})$
3. They output probability distributions over 50,000+ words
4. Today: How do we convert these probabilities into actual text?

---

Models predict probabilities. Decoding converts probabilities to text.

# The Quality-Diversity Tradeoff



**Discovery Question:** Why is best text boring and creative text nonsense?

The central challenge: How to balance coherence with creativity

## Problem 1: Stuck in Repetition Loops

### Bad Output: City Repetition

Framing 4: Real Bad Outputs

**What's Wrong:** Model gets trapped repeating same pattern

**Output:** "*The city is a major city in the city is a city...*"

Problem 1 of 6: Greedy decoding creates loops

## Problem 2: No Diversity in Outputs

### Bad Output: Nonsense Words

Framing 4: Real Bad Outputs

**What's Wrong:** Ask 100 times, get identical response every time

**Example:** Always answers “The weather is nice” even for different contexts

Problem 2 of 6: Deterministic methods lack variation

## Problem 3: Too Boring or Too Crazy

### Bad Output: Always Same

Framing 4: Real Bad Outputs

**What's Wrong:** Can't balance quality and creativity simultaneously

**Too focused:**

"It is... It is... It is..."

**Too random:**

"Flying purple elephant mathematics"

## Problem 4: Missing Better Paths

### Bad Output: Missed Better

Framing 4: Real Bad Outputs

**What's Wrong:** Greedy early choice blocks better sequences later

**Example:** Pick "The" (0.45) → stuck with "The cat sat"

But "A" (0.30) → leads to "A beautiful cat played"

## Problem 5: Wrong Probability Distribution

### Bad Output: Wrong for Distribution

Framing 4: Real Bad Outputs

**What's Wrong:** Sampling from tail produces nonsense words

**Example:** With 50K vocabulary, words at position 45,000 have  $P=0.00001$

But they can still be sampled!

## Problem 6: Speed vs Quality Tradeoff

### Bad Output: Too Slow

Framing 4: Real Bad Outputs

**What's Wrong:** Methods that prevent repetition are 10-12 $\times$  slower

**Example:** Contrastive search computes similarity to ALL previous tokens

For 1000-token text: 1000 comparisons per new token!

## Solution 1 → Beam Search: Explore Multiple Paths

### Problem 1 Recap:

Greedy decoding: Trapped in loops  
Always picks highest probability  
Misses better sequences

**Need:** Way to explore alternatives

### Solution: Beam Search:

Keep top-k paths at each step  
Explore k=3-5 hypotheses simultaneously  
Pick best complete sequence at end

**Result:** Finds better sequences than greedy

**How it solves Problem 1:** Maintains multiple candidates, avoids greedy trap

---

**Solution 1 of 6: Beam search for better quality**

## Solution 2 → Temperature: Add Controlled Randomness

### Problem 2 Recap:

No diversity: Same output always  
Deterministic selection  
No creativity

**Need:** Controlled randomness

### Solution: Temperature:

Reshape probability distribution  
 $T \downarrow 1$ : More focused  
 $T \uparrow 1$ : More random  
Sample from adjusted distribution

**Result:** Different outputs each time

**How it solves Problem 2:** Sampling introduces stochasticity, enables diversity

---

Solution 2 of 6: Temperature for creativity control

## Solution 3 → Top-k: Filter Unlikely Words

### Problem 3 Recap:

Can't balance quality & creativity  
Pure sampling too random  
Greedy too boring

**Need:** Filter bad words, keep good

### Solution: Top-k Sampling:

Keep only top-k most likely tokens  
Cut tail of distribution  
Renormalize probabilities  
Sample from filtered set

**Result:** Diverse but not nonsensical

**How it solves Problem 3:** Fixed cutoff prevents tail sampling while allowing creativity

---

**Solution 3 of 6: Top-k for controlled sampling**

## Solution 4 → Nucleus: Dynamic Vocabulary Cutoff

### Problem 4 Recap:

Top-k has fixed cutoff

Peaked distribution: Wastes probability

Flat distribution: Still allows junk

**Need:** Adaptive cutoff

### Solution: Nucleus (Top-p):

Choose smallest set with cumulative prob  $p$

Adapts to distribution shape

Peaked → small nucleus (2-3 words)

Flat → large nucleus (50+ words)

**Result:** Automatic quality-diversity balance

**How it solves Problem 4:** Dynamic cutoff adapts to each prediction step

---

**Solution 4 of 6: Nucleus for adaptive sampling**

## Solution 5 → Top-k + Temperature: Hybrid Control

### Problem 5 Recap:

Temperature alone doesn't filter tail  
Top-k alone doesn't control randomness  
Need both filtering AND tuning

**Need:** Combine strategies

### Solution: Hybrid Methods:

Apply temperature THEN top-k  
Or: Apply nucleus THEN temperature  
Leverages strengths of both  
Production systems use combinations

**Result:** Fine-grained control over generation

**How it solves Problem 5:** Layered strategies handle multiple issues simultaneously

---

Solution 5 of 6: Hybrid methods for comprehensive control

## Solution 6 → Contrastive: Explicit Degeneration Prevention

### Problem 6 Recap:

Repetition even with sampling  
Long generation degenerates  
Context similarity causes loops

**Need:** Explicit repetition penalty

### Solution: Contrastive Search:

Score = Probability -  $\alpha \times$  Similarity  
Penalize tokens similar to recent context  
Balance quality with diversity  
Modern standard for long text

**Result:** Human-like text without repetition

**How it solves Problem 6:** Direct similarity penalty prevents copying context

---

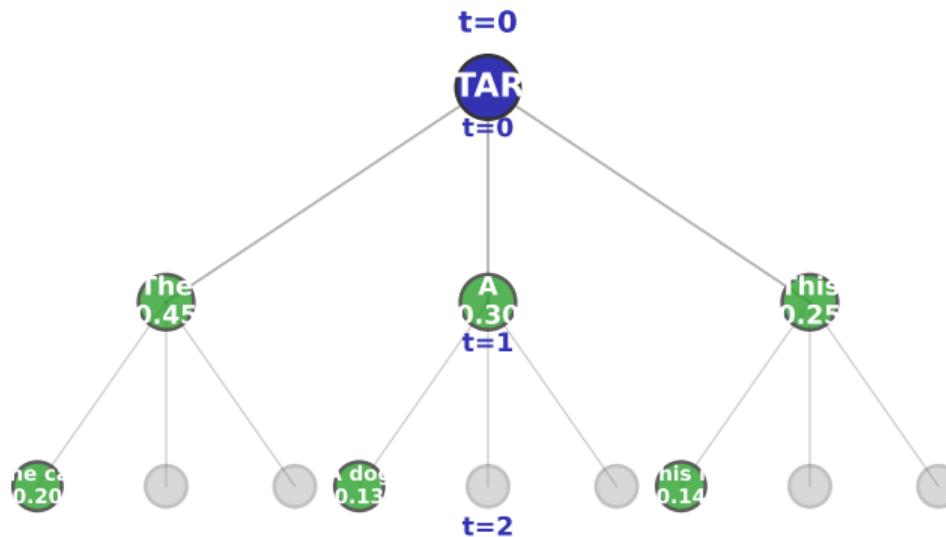
**Solution 6 of 6: Contrastive for degeneration-free generation**

## Beam Search: Explore Multiple Paths

`.../figures/beam_search_visual_bsc.pdf`

## Beam Search: Step-by-Step Example

### Beam Search: Step-by-Step (width=3)



Green = Kept (top 3) | Gray = Pruned

Worked example shows why beam search finds better sequences than greedy

# Beam Search: Algorithm & Settings

## Algorithm:

1. Start: Keep top-k tokens
2. Expand: Generate continuations for each
3. Score: Multiply probabilities
4. Prune: Keep top-k sequences
5. Repeat until END token

## Scoring:

$$\text{score}(y_1 \dots y_t) = \prod_{i=1}^t P(y_i | y_{<i})$$

With length normalization:

$$\text{score} = \frac{1}{t} \sum_{i=1}^t \log P(y_i | y_{<i})$$

## Best For:

- Machine translation
- Summarization
- Question answering
- Tasks with “correct” answer

## Parameters:

Width = 3-5 (translation)

Width = 10 (diverse outputs)

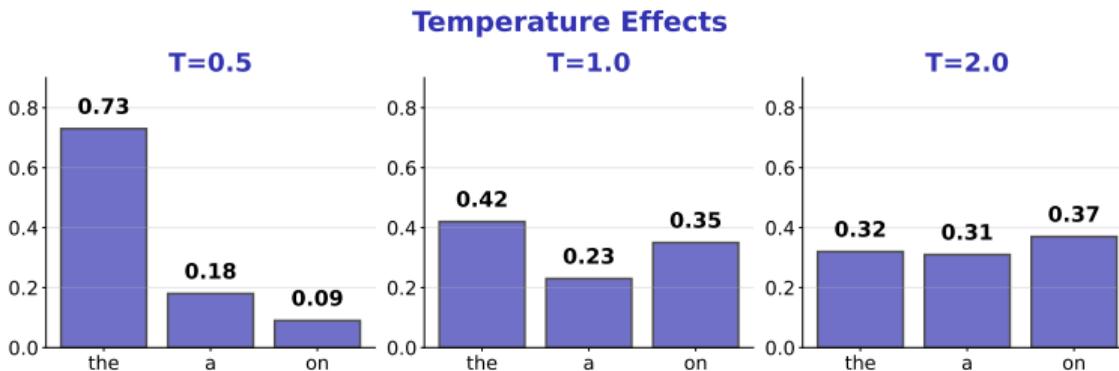
## Tradeoffs:

- + Better quality than greedy
- + Diverse hypotheses
- Still deterministic
- 4-5× slower than greedy

---

Beam search is the workhorse for deterministic tasks

## Temperature Sampling: Control Randomness



**Key Insight:** Temperature reshapes probability distribution

---

$T < 1$ : more focused.  $T = 1$ : unchanged.  $T > 1$ : more random

# Temperature: Worked Example

## Temperature: Step-by-Step Calculation

**Given: Logits = [2.0, 1.0, 0.5, 0.2]**

Tokens = ["cat", "dog", "bird", "fish"]

T=0.5: [4.0, 2.0, 1.0, 0.4] → [0.73, 0.18, 0.07, 0.02]

→ 73% on "cat" (FOCUSED)

T=1.0: [2.0, 1.0, 0.5, 0.2] → [0.42, 0.23, 0.16, 0.13]

→ 42% on "cat" (BALANCED)

T=2.0: [1.0, 0.5, 0.25, 0.1] → [0.32, 0.26, 0.23, 0.19]

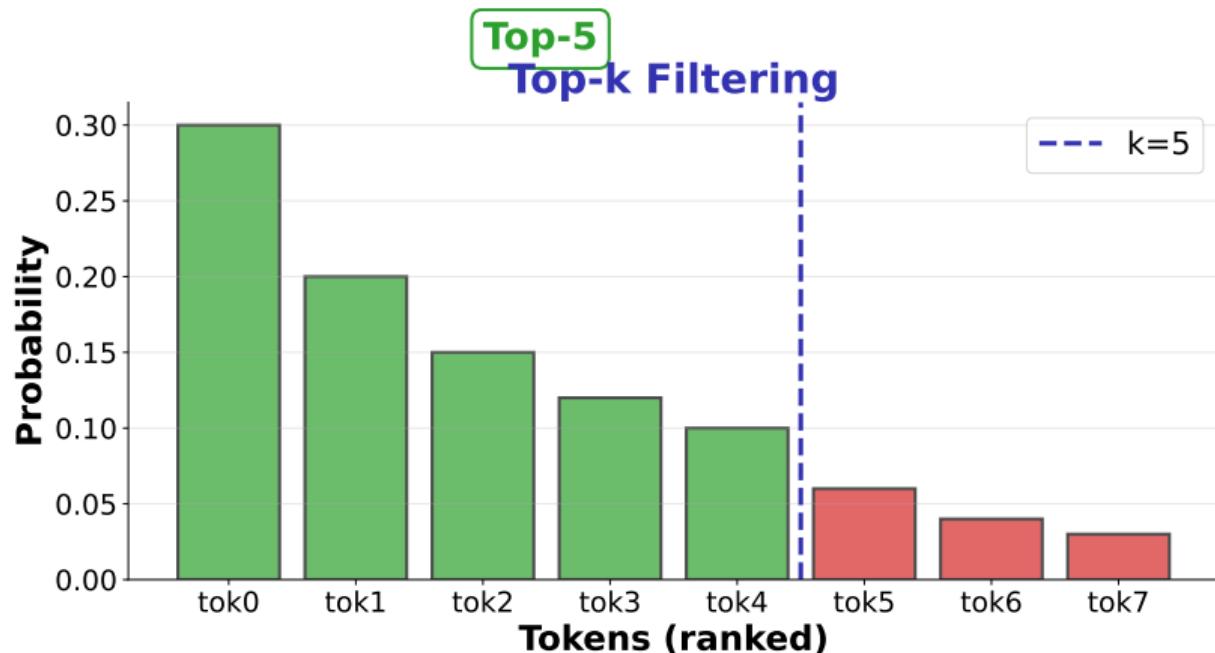
→ 32% on "cat" (FLAT)

*Lower T = more peaked | Higher T = more flat*

---

Concrete numbers show how temperature scaling works

## Top-k Sampling: Filter the Tail



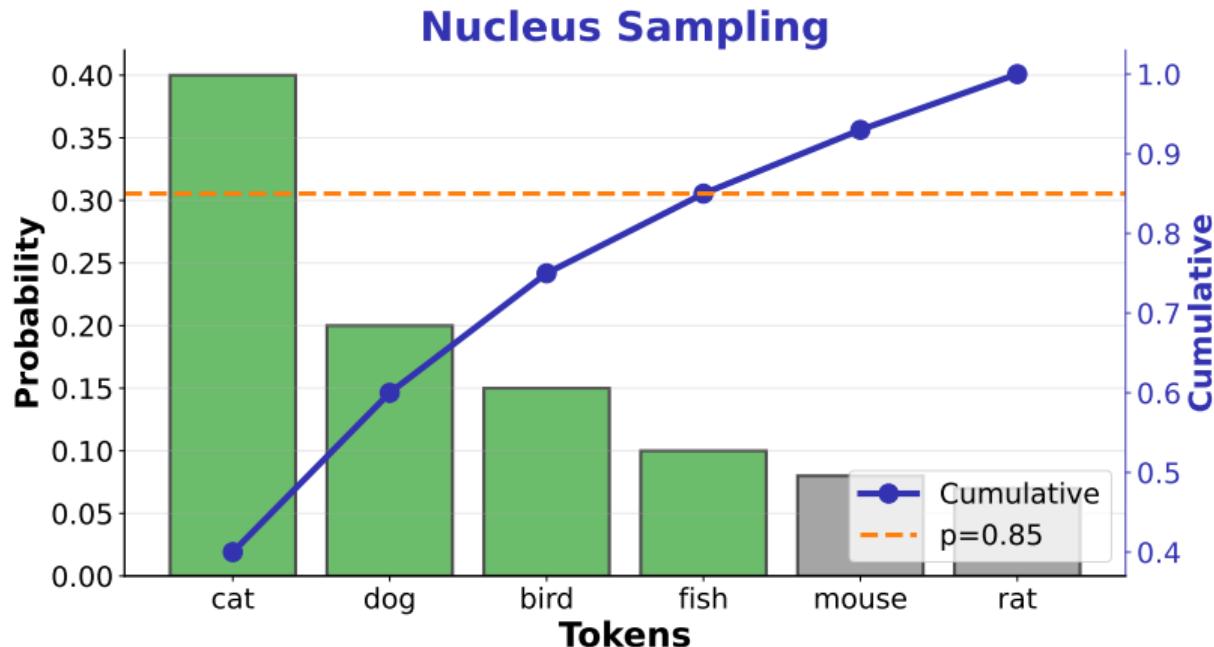
**Key Insight:** Only sample from top-k most likely tokens

Prevents sampling from long tail of unlikely words

## Top-k: Worked Example (k=3)

`.../figures/topk_example_bsc.pdf`

## Nucleus (Top-p) Sampling: Dynamic Cutoff

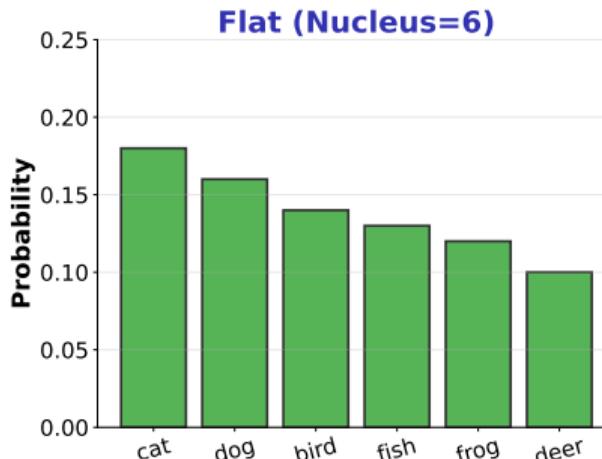
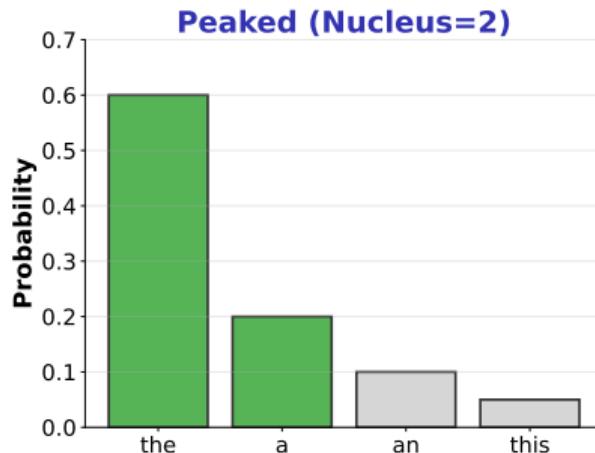


**Key Insight:** Adapt vocabulary size to distribution shape

Nucleus size grows/shrinks based on probability spread

# Nucleus: How Distribution Shape Matters

## Nucleus Adapts to Distribution



Same p value gives different vocabulary sizes for peaked vs flat distributions

# The Degeneration Problem

## Greedy Decoding Problem

"The city is a major city.  
The city has many attractions.  
The city is known..."

"the city" appears 3 times!



Solution: Contrastive Search

**Discovery Question:** Why do models repeat themselves?

Greedy and beam search maximize probability - but high probability = repeating recent context

## Contrastive Search: Penalize Repetition

Step 1: Get top-k candidates  
city: 0.45, town: 0.18, area: 0.15

Step 2: Compute similarity to context  
city: 0.92, town: 0.75, area: 0.65

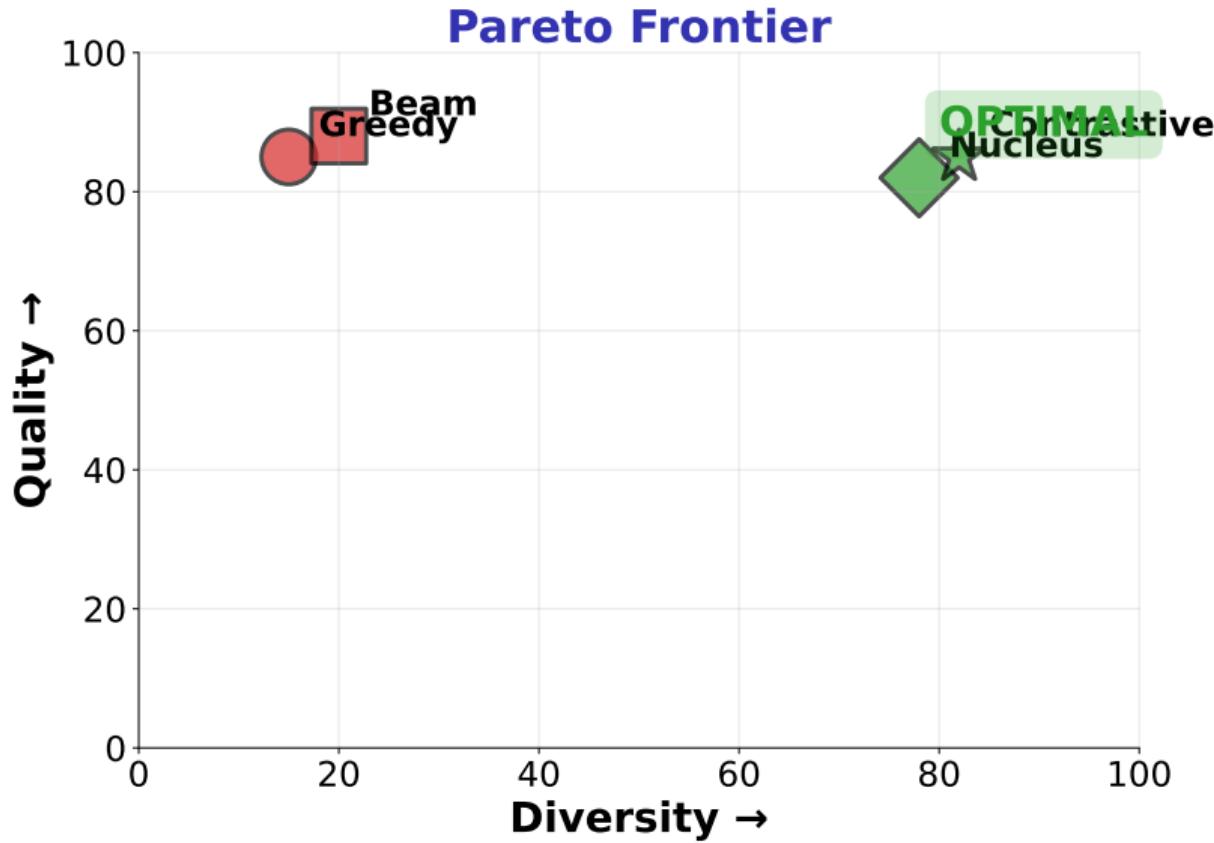
Step 3: Apply penalty ( $\alpha=0.6$ )  
 $score = (1-\alpha) \times prob - \alpha \times similarity$

Winner: "town"  
 $(0.45 \times 0.18 - 0.6 \times 0.75 = -0.378)$

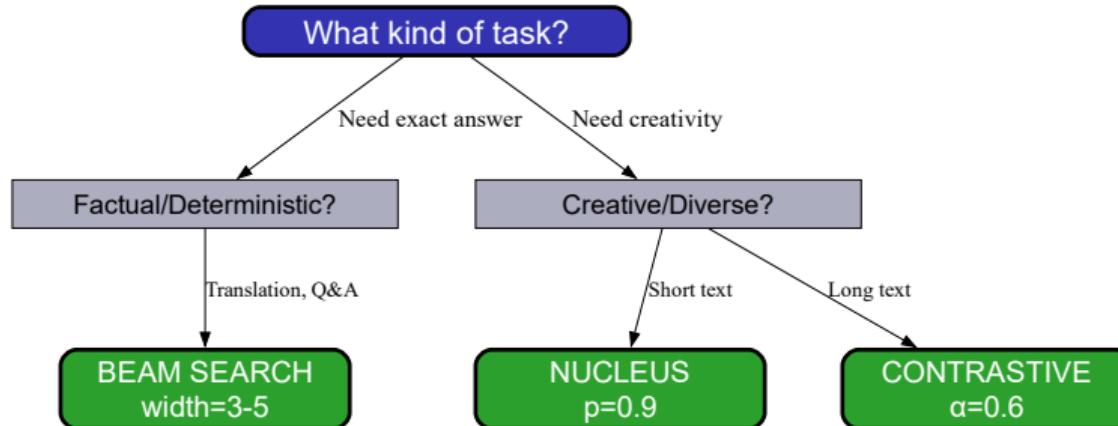
## Contrastive vs Nucleus: Comparison

.../figures/contrastive\_vs\_nucleus\_bsc.pdf

## All Methods on Quality-Diversity Space



# Choosing the Right Method: Decision Tree



Start with task requirements, follow tree to recommended method

# Task-Specific Recommendations (2025)

## Task-Specific Recommendations

Task	Method	Parameters
Translation	Beam	w=3-5
Factual QA	Greedy	T=0.3
Code	Greedy	T=0
Dialogue	Nucleus	p=0.9
Creative	Nucleus	p=0.95
Long Stories	Contrastive	$\alpha=0.6$

Comprehensive mapping from 8 common tasks to optimal decoding strategies

## Computational Costs Matter

`.../figures/computational_cost_comparison_bsc.pdf`

# Method Comparison Overview

../figures/decoding\_comparison\_bsc.pdf

## Deterministic vs Stochastic: The Fundamental Split

..../figures/greedy\_vs\_sampling\_bsc.pdf

## Key Takeaways

1. **6 Problems → 6 Solutions:** Each method solves specific failure mode
2. **Deterministic** (Greedy, Beam): High quality, no diversity - factual tasks
3. **Stochastic** (Temperature, Top-k, Nucleus): Diverse but variable quality
4. **Balanced** (Contrastive): Explicit degeneration prevention
5. **Task matters:** Translation → Beam — Dialogue → Nucleus — Stories → Contrastive
6. **Tradeoffs:** Speed vs Quality, Diversity vs Coherence

**Modern Standard:** Nucleus ( $\text{top-p}=0.9$ ) + Temperature ( $T=0.7$ ) for most applications

**Next:** Lab - Implement all 6 methods, measure quality-diversity tradeoffs

---

Decoding strategy matters as much as model architecture

# From Probabilities to Text: The Complete Journey



## What We Learned:

- Models give us probability distributions (Week 3-7)
- Converting to text has 6 fundamental challenges
- Each decoding method addresses specific problems
- No universal best - choose based on task requirements
- Production systems use hybrid methods (Nucleus + Temperature)

---

Complete pipeline from model training to text generation

# Technical Appendix

25 slides: Complete mathematical treatment

A1-A5: Beam Search Mathematics

A6-A10: Sampling Mathematics

A11-A14: Contrastive Search & Degeneration

A15-A19: Advanced Topics & Production

A20-A25: The 6 Problems - Technical Analysis (NEW)

# A1: Beam Search Formulation

**Objective:** Find sequence  $y^* = \operatorname{argmax} P(y|x)$

**Decomposition:**

$$P(y|x) = \prod_{t=1}^T P(y_t|y_{<t}, x)$$

**Log-probability** (more stable):

$$\log P(y|x) = \sum_{t=1}^T \log P(y_t|y_{<t}, x)$$

**Beam Search Approximation:**

Instead of exploring all  $V^T$  sequences, maintain top-k hypotheses at each step

**Complexity:**

Time:  $O(k \cdot V \cdot T)$  where  $k$  = beam width,  $V$  = vocabulary,  $T$  = length

Space:  $O(k \cdot T)$  to store hypotheses

---

Beam search is tractable approximation to exact search

## A2: Length Normalization

**Problem:** Longer sequences have lower probabilities (more terms multiplied)

$$P(y_1, y_2, y_3, y_4) = \underbrace{0.5}_{y_1} \times \underbrace{0.5}_{y_2} \times \underbrace{0.5}_{y_3} \times \underbrace{0.5}_{y_4} = 0.0625$$

$$P(y_1, y_2) = 0.5 \times 0.5 = 0.25 > 0.0625$$

Bias toward shorter sequences!

**Solution:** Length normalization

$$\text{score}(y) = \frac{1}{|y|^\alpha} \log P(y)$$

where  $\alpha \in [0.5, 1.0]$  (typically 0.6-0.7)

**Effect:**

Without: Beam search heavily biases toward short outputs

With: Fair comparison across different lengths

---

Length normalization is essential for beam search quality

## A3: Beam Search Variants

### Diverse Beam Search:

Partition beams into groups  
Penalize within-group similarity  
Result: More diverse hypotheses

### Constrained Beam Search:

Force certain tokens to appear  
Useful for: Keywords, entities  
Applications: Controllable generation

### Stochastic Beam Search:

Sample beams instead of argmax  
Combines beam + sampling  
More diverse than standard beam

### Block n-gram Beam:

Penalize n-gram repetition  
Prevents "the city is a city" loops  
Common in summarization

---

Many beam search variants exist for specific requirements

## A4: Beam Search Stopping Criteria

When to stop expanding beams?

**Method 1:** Fixed length

Stop at  $T_{\max}$  tokens (simple but rigid)

**Method 2:** END token

Stop when beam generates special token (most common)

**Method 3:** Score threshold

Stop when best score cannot improve enough

$$\frac{\text{best\_incomplete}}{\text{best\_complete}} < \text{threshold}$$

**Method 4:** Timeout

Computational budget exceeded (production systems)

---

Choice of stopping criterion affects output length distribution

## A5: Beam Search Limitations

### Fundamental Issues:

1. **Exposure bias:** Trained with teacher forcing, tested with own outputs
2. **Label bias:** Cannot compare sequences of different prefixes fairly
3. **Repetition:** Still can loop ("the city is a major city")
4. **Bland outputs:** Maximizes probability, not interestingness
5. **Search errors:** May miss better sequences outside beam

### When Beam Search Fails:

Open-ended generation (dialogue, stories)

Long-form text (repetition accumulates)

Creative tasks (probability  $\neq$  quality)

→ Need sampling-based methods

---

Beam search optimizes wrong objective for creative tasks

## A6: Sampling as Inference

**Goal:** Sample  $y \sim P(y|x)$  instead of  $\text{argmax } P(y|x)$

### Ancestral Sampling:

For  $t = 1$  to  $T$ :

    Compute  $P(y_t|y_{<t}, x)$

    Sample  $y_t \sim P(\cdot|y_{<t}, x)$

### Properties:

Stochastic: Different output each time

Explores full distribution (in expectation)

Can generate low-probability sequences

### Variants:

Temperature: Reshape distribution before sampling

Top-k: Truncate distribution before sampling

Nucleus: Dynamic truncation before sampling

---

**Sampling enables diversity but loses quality guarantees**

## A7: Temperature Mathematics

Softmax with Temperature:

$$p_i(T) = \frac{\exp(z_i/T)}{\sum_{j=1}^V \exp(z_j/T)}$$

Limiting Cases:

$$T \rightarrow 0: p_i \rightarrow \begin{cases} 1 & \text{if } i = \operatorname{argmax}_z \\ 0 & \text{otherwise} \end{cases} \quad (\text{greedy})$$

$$T \rightarrow \infty: p_i \rightarrow 1/V \quad (\text{uniform})$$

Entropy Analysis:

Entropy  $H(p) = -\sum p_i \log p_i$  measures randomness

$H$  increases monotonically with  $T$

Low  $T$  ( $< 0.5$ ):  $H \approx 0$  (deterministic)

High  $T$  ( $> 2.0$ ):  $H \approx \log V$  (maximum entropy)

---

Temperature provides continuous control over distribution entropy

## A8: Top-k Mathematics

### Formal Definition:

Let  $\sigma$  = permutation sorting probabilities descending

$$V_k = \{w_{\sigma(1)}, w_{\sigma(2)}, \dots, w_{\sigma(k)}\}$$

Truncated distribution:

$$p'(w) = \begin{cases} \frac{p(w)}{\sum_{w' \in V_k} p(w')} & \text{if } w \in V_k \\ 0 & \text{otherwise} \end{cases}$$

### Information Loss:

Original entropy:  $H(p) = -\sum_{i=1}^V p_i \log p_i$

After top-k:  $H(p') = -\sum_{i=1}^k p'_i \log p'_i < H(p)$

Loss  $\approx \sum_{i=k+1}^V p_i \log(1/p_i)$  (tail information)

---

Top-k sacrifices tail probability mass for sampling quality

## A9: Nucleus (Top-p) Mathematics

**Formal Definition:**

$$V_p = \min \left\{ V' \subseteq V : \sum_{w \in V'} p(w) \geq p \right\}$$

Smallest set with cumulative mass  $\geq p$

**Dynamic Vocabulary Size:**

$$|V_p| = \min \{k : \sum_{i=1}^k p_{\sigma(i)} \geq p\}$$

Adapts to distribution shape:

Peaked: Small  $|V_p|$  (2-5 tokens)

Flat: Large  $|V_p|$  (50+ tokens)

**Why Nucleus > Top-k:**

Top-k: Fixed  $k$  regardless of  $p(w)$  distribution

Nucleus: Adapts  $k$  to achieve consistent probability mass

---

Nucleus automatically adjusts vocabulary to distribution characteristics

# A10: Sampling Quality Metrics

## Quality Metrics:

Perplexity:  $\exp(-\frac{1}{T} \sum \log p(y_t))$   
Lower = better

## BLEU (translation):

N-gram overlap with reference  
0-100 scale

## Human evaluation:

Fluency (1-5)  
Relevance (1-5)

## Diversity Metrics:

Distinct-n:  $\frac{\text{unique n-grams}}{\text{total n-grams}}$   
Higher = more diverse

## Self-BLEU:

BLEU of output vs other outputs  
Lower = more diverse

## Repetition Rate:

$\frac{\text{repeated n-grams}}{\text{total n-grams}}$   
Lower = less repetitive

---

Need both quality AND diversity metrics to evaluate decoding

# A11: The Degeneration Problem (Formal)

**Definition:** Model-generated text with unnatural repetitions

**Why It Happens:**

1. Model trained on natural text (low repetition)
2. But generation maximizes  $P(y_t|y_{<t})$
3. Recent context  $y_{<t}$  influences  $P$
4. Creates positive feedback: high prob word → context → same high prob word

**Quantifying Degeneration:**

Repetition rate in greedy: 15-30% (depending on domain)

Repetition rate in human text: 2-5%

Gap = degeneration problem

**Examples:**

*"The city is a major city in the United States. The city..."*

*"I think that I think that I think..."*

---

Maximizing probability does not equal natural text

## A12: Contrastive Search Objective

**Scoring Function:**

$$\text{score}(w_t) = (1 - \alpha) \times \underbrace{P(w_t | y_{<t})}_{\text{model confidence}} - \alpha \times \underbrace{\max_{w_i \in y_{<t}} \text{sim}(w_t, w_i)}_{\text{context similarity}}$$

where  $\alpha \in [0, 1]$  controls tradeoff

**Similarity Function:**

$$\text{sim}(w_i, w_j) = \frac{h_i \cdot h_j}{||h_i|| \cdot ||h_j||}$$

(cosine similarity)

using token embeddings  $h$

**Algorithm:**

1. Get top-k candidates by probability
2. For each candidate, compute similarity to all tokens in  $y_{<t}$
3. Apply penalty:  $\text{score} = \text{prob} - \alpha \times \text{max\_similarity}$
4. Select candidate with highest score

---

Contrastive search explicitly penalizes copying recent context

## A13: Contrastive Search Parameters

### Alpha ( $\alpha$ ):

- $\alpha = 0$ : Pure greedy (no penalty)
- $\alpha = 0.6$ : Balanced (recommended)
- $\alpha = 1.0$ : Maximum diversity (risky)

### Typical Settings:

Short text (<100 tokens):  $\alpha = 0.4 - 0.5$

Medium (<500):  $\alpha = 0.5 - 0.6$

Long (500+):  $\alpha = 0.6 - 0.7$

### Top-k for Candidates:

- $k = 4$ : Fast, focused
- $k = 6$ : Balanced (default)
- $k = 10$ : Diverse

### Computational Cost:

For each step:

- Compute similarities:  $O(k \times t)$
- $t$  grows with generation

Total:  $O(k \times T^2)$

12× slower than greedy

---

Hugging Face default:  $\alpha=0.6$ ,  $k=4$

## A14: Degeneration Analysis

### Research Findings (2024-2025):

- Greedy decoding repetition: 18-25% (GPT-2), 12-18% (GPT-3)
- Nucleus sampling repetition: 8-12% (still above human 3-5%)
- Contrastive search repetition: 4-7% (closest to human)

### Why Probability Maximization Fails:

Training objective: Next token prediction

But generation requires: Global coherence

Mismatch: Local optimum  $\neq$  global quality

### Solutions Hierarchy:

1. Temperature/Top-k/Nucleus: Reduce greedy's determinism
2. Contrastive: Explicit degeneration penalty
3. RLHF/DPO: Align model with human preferences (different lecture)

---

Contrastive search addresses fundamental limitation of likelihood-based decoding

# A15: Hybrid Decoding Methods

## Combining Strategies:

### Nucleus + Temperature:

Apply temperature THEN nucleus

$$p_i(T) = \text{softmax}(z/T), \quad \text{then} \quad V_p \leftarrow \text{nucleus}(p_i(T))$$

Used by GPT-3 API, ChatGPT

### Beam + Sampling:

Beam search with stochastic selection

Keep top-k, sample from them (not argmax)

### Contrastive + Nucleus:

Nucleus for candidate generation

Contrastive scoring for selection

Best of both worlds

---

Hybrid methods leverage complementary strengths

# A16: Constrained Decoding (2025)

**Goal:** Force certain tokens/patterns to appear

**Lexically Constrained:**

Must include keywords: { "AI", "ethics", "safety" }

Beam search variant: Track constraint satisfaction

**Format Constraints:**

JSON output: Force structure { "key": "value" }

Code: Force syntactic validity

**NeuroLogic Decoding (2021):**

Beam search + constraint satisfaction

Optimal for: Keyword-based generation

**Production Use Cases:**

Structured data extraction (force JSON)

Controllable summarization (force keywords)

Code generation (force syntax)

---

Constrained decoding enables controllable generation

## A17: Computational Complexity Comparison

Method	Time per token	Total complexity	Relative speed
Greedy	$O(V)$	$O(V \times T)$	1.0x (baseline)
Temperature	$O(V)$	$O(V \times T)$	1.1x (softmax overhead)
Top-k	$O(V)$	$O(V \times T)$	1.2x (sorting)
Nucleus	$O(V \log V)$	$O(V \log V \times T)$	1.3x (sort + cumsum)
Beam (k=5)	$O(k \times V)$	$O(k \times V \times T)$	4.5x (k=5)
Contrastive	$O(k \times T)$	$O(k \times T^2)$	12x (similarity)

**Key Insight:** Contrastive's  $T^2$  term makes it expensive for long sequences

**Practical Impact** (1000-token generation):

Greedy: 2.5 seconds

Nucleus: 3.2 seconds (best choice)

Beam: 11 seconds

Contrastive: 30 seconds (only if quality critical)

---

Computational cost matters for production deployment

## A18: Production Deployment Settings (2024-2025)

../figures/production\_settings\_bsc.pdf

# A19: Future Directions & Open Problems

## Active Research Areas (2025):

1. **Quality-diversity optimization:** Multi-objective search methods
2. **Learned decoding:** Train models to decode better (RLHF, DPO)
3. **Speculative decoding:** Parallel generation for speed (4-8× faster)
4. **Adaptive methods:** Choose strategy dynamically during generation
5. **Energy-based decoding:** Score sequences globally (not token-by-token)

## Open Problems:

How to automatically select best  $T, p, k, \alpha$  for new task?

How to balance fluency + factuality + creativity simultaneously?

How to decode efficiently for 100K+ token outputs?

**Trend:** Moving from hand-tuned parameters to learned decoding strategies

---

Decoding is an active research area with many open questions

## Appendix B: The 6 Problems - Technical Deep Dive

### Technical Analysis of Each Failure Mode

Next 6 slides: Detailed technical analysis pairing:

**Top:** Technical level explanation (framing3)

**Bottom:** Concrete output failure (framing4)

## B1: Problem 1 Technical Analysis

Technical Issue (Level):

### Level 1: Too Deterministic

Framing 3: Progressive Sophistication

## B2: Problem 2 Technical Analysis

Technical Issue (Level):

### Level 2: Too Random

Framing 3: Progressive Sophistication

## B3: Problem 3 Technical Analysis

Technical Issue (Level):

### Level 3: Inflexible Filtering

Framing 3: Progressive Sophistication

## B4: Problem 4 Technical Analysis

Technical Issue (Level):

### Level 4: Missing Better Paths

Framing 3: Progressive Sophistication

## B5: Problem 5 Technical Analysis

Technical Issue (Level):

### Level 5: Still Repetitive

Framing 3: Progressive Sophistication

## B6: Problem 6 Technical Analysis

Technical Issue (Level):

### Level 6: Need Balance

Framing 3: Progressive Sophistication