

Efficiency & Optimization

Week 11 - NLP Course 2025

NLP Course 2025

September 28, 2025

Efficiency & Optimization

Week 11: Overview

Learning Objectives

- Understand Quantization
- Understand Pruning
- Apply concepts in practice
- Evaluate different approaches

Topics Covered

- Quantization
- Pruning
- Distillation
- Efficient architectures

Introduction

Efficiency & Optimization

Content for Efficiency & Optimization...

Key Concepts

- Concept 1: Quantization
- Concept 2: Pruning
- Practical applications
- Current research directions

Summary

Key Takeaways

- Main learning point 1
- Main learning point 2
- Main learning point 3

Next week: Continue to next topic