

Natural Language Processing Course

Week 11: Efficiency and Deployment

Joerg R. Osterrieder

www.joergosterrieder.com

Week 11

Efficiency & Deployment

From Cloud Giants to Pocket-Sized Models

Why Your Phone Can't Run GPT-4 (Yet)

The shocking reality of modern models

- GPT-3: 175B parameters = 350GB in FP16
 - Your phone: 6GB RAM
 - Inference cost: \$0.06 per 1K tokens
 - Latency: 100ms+ per token
 - Energy: 0.1 kWh per conversation

One ChatGPT query = 10x the energy of a Google search

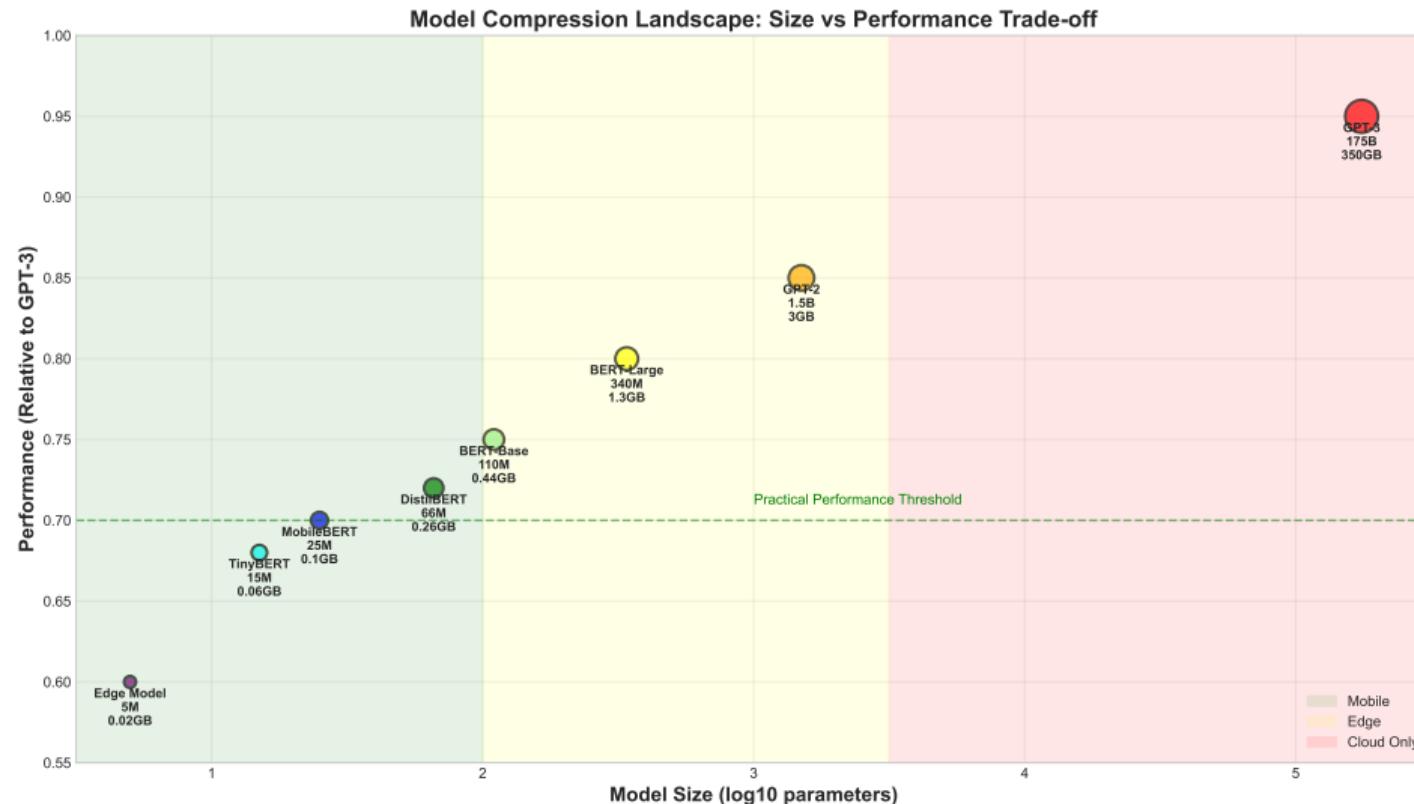
The challenge:

- Users want instant responses
 - Privacy requires on-device processing
 - Costs are unsustainable at scale
 - Edge devices have limited resources

This week: How to shrink giants into pocket-sized assistants

¹Brown et al. (2020); Strubell et al. (2019) on model carbon footprint

From Supercomputers to Smartphones: The Deployment Journey



Efficient Models in Production (2024)

Success Stories:

- iPhone: On-device Siri (5B → 200M)²
- Google: Pixel voice typing
- Microsoft: Excel formula suggestions
- WhatsApp: Real-time translation
- Tesla: In-car voice commands

Business Impact:

- 100x cost reduction
- 10ms latency (from 1s)
- Works offline
- Privacy preserved
- Scales to billions

Efficiency Techniques:

- INT8 quantization: Standard
- INT4/Binary: Emerging
- Distillation: 95% performance
- Structured pruning: Hardware-friendly
- Flash attention: Memory efficient

Deployment Targets:

- Mobile phones: 2-4GB limit
- Browsers: WebAssembly
- IoT devices: MCUs
- Edge servers: Local processing
- Specialized chips: NPUs/TPUs

2024: Every device runs neural networks - efficiency made it possible

²Apple ML Research Blog (2023); Google I/O presentations

Week 11: What You'll Master

By the end of this week, you will:

- **Understand** why models are so large
- **Implement** quantization techniques
- **Master** knowledge distillation
- **Apply** pruning strategies
- **Deploy** models to edge devices

Core Insight: 90% of weights do 10% of the work

Why Are Models So Large? The Overparameterization Mystery

The paradox:

Models have way more parameters than necessary!

Evidence:

- Lottery Ticket Hypothesis: Small subnetworks work just as well³
- Magnitude pruning: Remove 90% weights, maintain accuracy
- Low-rank decomposition: Matrices are redundant
- Quantization: 32 bits → 4 bits still works

Why overparameterize?

- Easier optimization landscape
- Better generalization (surprisingly!)
- Redundancy provides robustness
- Training dynamics require it

Large models are like rough marble blocks - we can carve out efficient versions

³Frankle & Carbin (2019) "The Lottery Ticket Hypothesis"

Quantization: From Float32 to Int8 and Beyond

The quantization spectrum:

- FP32 (32 bits): Training precision
- FP16 (16 bits): Mixed precision training
- INT8 (8 bits): Standard deployment
- INT4 (4 bits): Aggressive compression
- Binary (1 bit): Research frontier

Implementing Post-Training Quantization

```
1 import torch
2 import torch.nn as nn
3
4 def quantize_tensor(x, num_bits=8):
5     """Quantize tensor to n bits"""
6     qmin = -(2**((num_bits-1)))
7     qmax = 2**((num_bits-1)) - 1
8
9     min_val = x.min()
10    max_val = x.max()
11
12    scale = (max_val - min_val) / (qmax - qmin)
13    zero_point = qmin - min_val / scale
14
15    q_x = torch.round(x / scale + zero_point)
16    q_x = torch.clamp(q_x, qmin, qmax)
17
18    return q_x, scale, zero_point
19
20 def dequantize_tensor(q_x, scale, zero_point):
21     """Dequantize back to float"""
22     return scale * (q_x - zero_point)
23
24 class QuantizedLinear(nn.Module):
25     def __init__(self, weight, bias, num_bits=8):
26         super().__init__()
27
28         self.q_weight, self.w_scale, self.w_zp = quantize_tensor(
29             weight, num_bits)
30         self.q_weight = self.q_weight.to(torch.int8)
31
32         if bias is not None:
33             self.q_bias, self.b_scale, self.b_zp = quantize_tensor(
34                 bias, num_bits)
35         else:
```

Quantization Impact:

- Memory: 4x reduction (INT8)
- Speed: 2-4x on CPUs
- Accuracy: -1% typical
- Energy: 10x savings

Advanced Techniques:

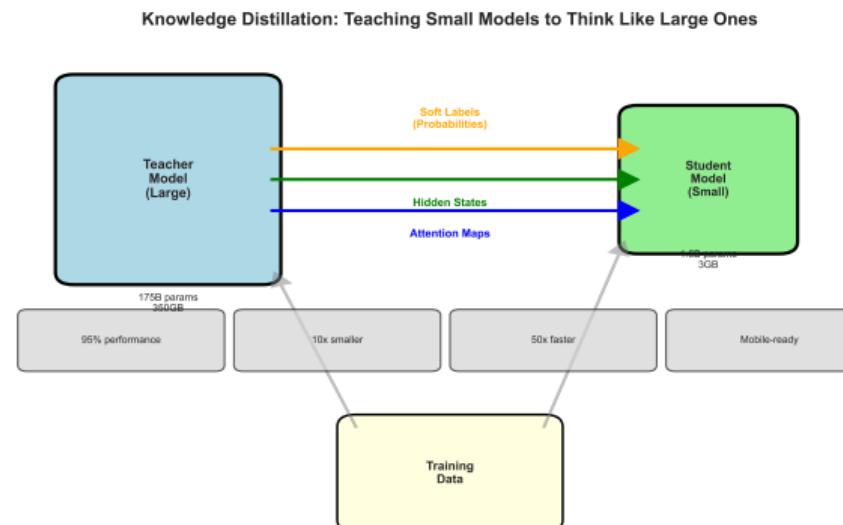
- Symmetric vs asymmetric
- Per-channel quantization
- Quantization-aware training
- Mixed bit-width

Hardware Support:

- INT8: All modern chips
- INT4: Newer GPUs/NPUs
- Binary: Research only

Knowledge Distillation: Teaching Small Models to Think Big

The teacher-student paradigm:⁴



Key insight: Soft labels contain more information

- Hard label: "cat" (probability = 1.0)
- Soft labels: cat=0.9, tiger=0.05, dog=0.03, ...
- Student learns relationships between classes

Implementing Knowledge Distillation

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4
5 class DistillationLoss(nn.Module):
6     def __init__(self, temperature=3.0, alpha=0.7):
7         """
8             temperature: Softens probability distributions
9             alpha: Weight between distillation and true label loss
10            """
11            super().__init__()
12            self.temperature = temperature
13            self.alpha = alpha
14            self.ce_loss = nn.CrossEntropyLoss()
15
16    def forward(self, student_logits, teacher_logits, labels):
17        soft_targets = F.softmax(teacher_logits / self.temperature, dim=-1)
18        soft_prob = F.log_softmax(student_logits / self.temperature, dim=-1)
19
20        distillation_loss = F.kl_div(soft_prob, soft_targets,
21                                     reduction='batchmean') * (self.temperature
22                                     ** 2)
23
24        student_loss = self.ce_loss(student_logits, labels)
25
26        loss = self.alpha * distillation_loss + (1 - self.alpha) * student_loss
27        return loss
28
29    def train_student(student_model, teacher_model, dataloader,
30                      epochs=10, temperature=3.0):
31        """
32            Train student to mimic teacher
33        """
34        teacher_model.eval()
35        optimizer = torch.optim.Adam(student_model.parameters(), lr=1e-3)
36        criterion = DistillationLoss(temperature=temperature)
```

Distillation Results:

- Size: 10-100x smaller
- Speed: 5-50x faster
- Accuracy: 95-98% retained
- Examples: DistilBERT, TinyBERT

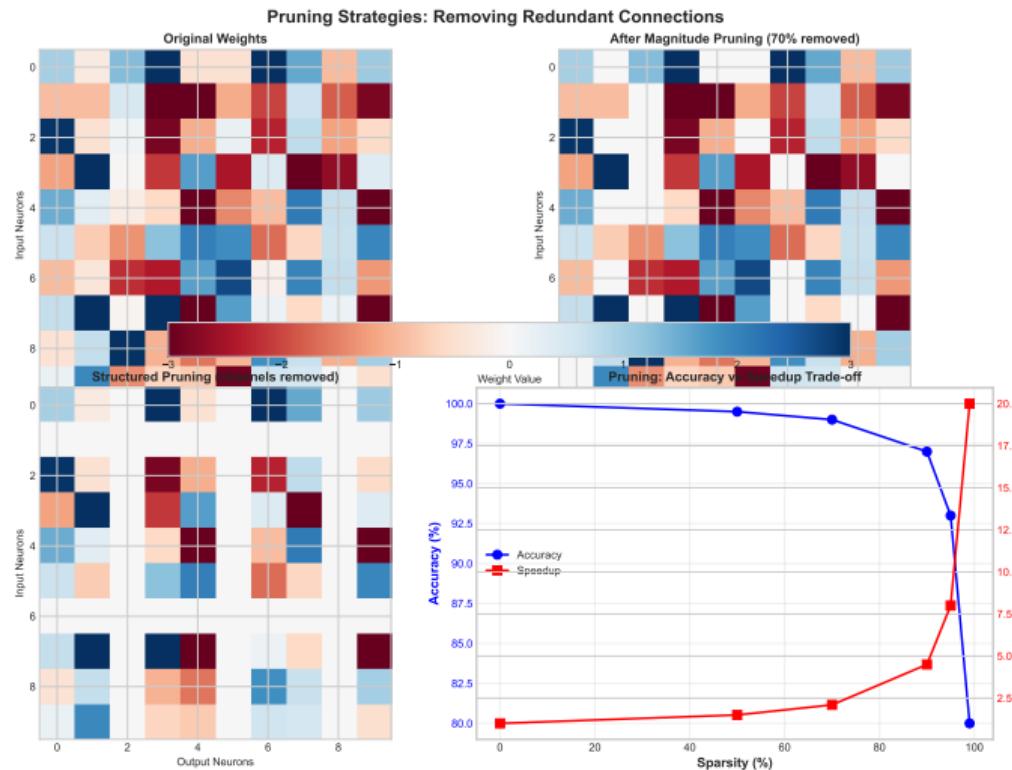
Temperature Effect:

- T=1: Normal softmax
- T=3-5: Typical for NLP
- T>10: Very soft labels
- Higher T = more dark knowledge

Advanced Methods:

- Feature distillation
- Attention transfer
- Progressive distillation

Pruning: Finding the Essential Subnetwork



Pruning approaches:

- Magnitude pruning: Remove small weights

Mobile-First Architectures: Designed for Efficiency

MobileBERT:⁵

- Bottleneck structure
- 4.3x smaller, 5.5x faster
- Depth-wise convolutions
- Progressive knowledge transfer

ALBERT:

- Parameter sharing
- Factorized embeddings
- 18x fewer parameters
- Same performance as BERT

EdgeBERT:

- Hardware-aware design
- Adaptive computation
- Early exit mechanisms
- Dynamic depth/width

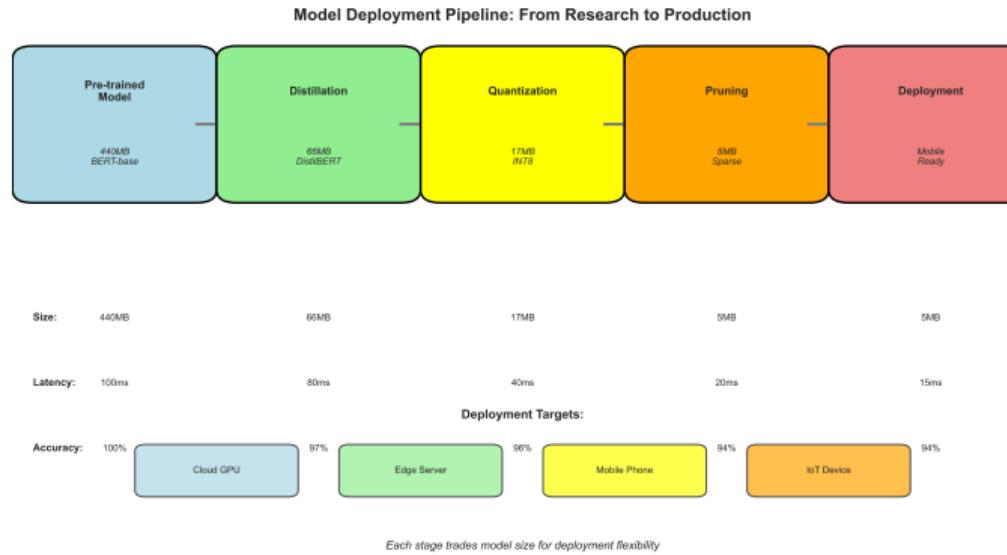
Design Principles:

- Prefer depth over width
- Share parameters
- Use separable operations
- Minimize memory access

Mobile architectures: Not compressed models, but efficient by design

⁵Sun et al. (2020) "MobileBERT"; Lan et al. (2020) "ALBERT"

Deployment Optimization Pipeline



Key Insights

- Start with pretrained model
- Apply compression techniques
- Optimize for target hardware
- Profile and iterate

Hardware Acceleration: From GPUs to Edge TPUs

Compute Platforms:

- GPUs: Training & cloud inference
- TPUs: Optimized for transformers
- NPUs: Mobile inference
- DSPs: Ultra-low power
- FPGAs: Custom acceleration

Optimization Techniques:

- Kernel fusion
- Memory pooling
- Graph optimization
- Operator scheduling
- Cache optimization

Framework Support:

- TensorFlow Lite: Mobile/edge
- ONNX Runtime: Cross-platform
- Core ML: iOS devices
- TensorRT: NVIDIA optimization
- OpenVINO: Intel hardware

Performance Gains:

- Quantization: 2-4x speedup
- Pruning: 2-10x speedup
- Hardware opt: 5-50x
- Combined: 100x+ possible

2024: Every major chip has AI acceleration built-in

Real-World Deployment: BERT on Mobile

```
1 import torch
2 import torch.quantization as quant
3 from transformers import BertModel
4
5 def prepare_mobile_bert(model_name='bert-base-uncased'):
6     """Complete pipeline for mobile deployment"""
7
8     model = BertModel.from_pretrained(model_name)
9     model.eval()
10
11    print(f"Original size: {get_model_size(model):.1f} MB")
12
13    distilled_model = distill_bert(model, num_layers=6)
14    print(f"After distillation: {get_model_size(distilled_model):.1f} MB")
15
16    model.qconfig = quant.get_default_qconfig('qnnpack')
17    quant.prepare(model, inplace=True)
18    quant.convert(model, inplace=True)
19    print(f"After INT8 quantization: {get_model_size(model):.1f} MB")
20
21    pruned_model = magnitude_prune(model, sparsity=0.9)
22    print(f"After pruning: {get_model_size(pruned_model):.1f} MB")
23
24    optimized_model = torch.jit.script(pruned_model)
25    optimized_model.save("mobile_bert.pt")
26
27    print("\nDeployment stats:")
28    print(f"Final size: {get_model_size(optimized_model):.1f} MB")
29    print(f"Inference time: {benchmark_model(optimized_model):.1f} ms")
30    print(f"Memory usage: {get_memory_usage(optimized_model):.1f} MB")
31
32    return optimized_model
33
34 def export_to_mobile(model, example_input):
35     """Export for iOS/Android"""
```

Compression Results:

- BERT-base: 440MB → 13MB
- 33x size reduction
- 20x speedup on mobile
- 95% accuracy retained

Deployment Checklist:

- Profile target device
- Choose compression mix
- Validate accuracy
- Optimize runtime
- Monitor in production

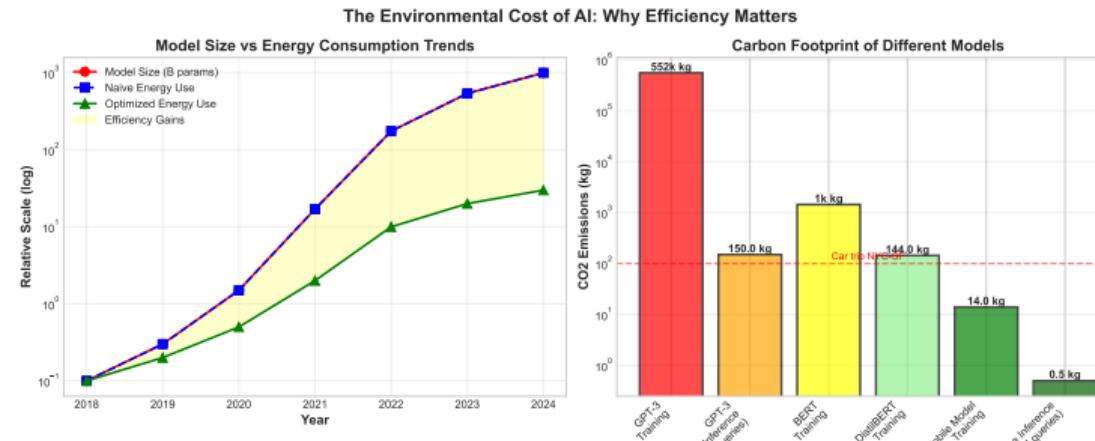
The Green AI Movement: Energy-Efficient Models

The environmental cost of AI:⁶

- Training GPT-3: 1,287 MWh (123 cars for a year)
- One ChatGPT query: 0.0003 kWh
- Global AI energy: Doubling every 3.4 months

Efficiency improvements:

- Sparse models: 10x less energy
- Quantization: 4x less energy
- Better algorithms: 100x over 10 years
- Hardware efficiency: 1000x since 2012



The Future of Efficient AI (2024 and Beyond)

Emerging Techniques:

- 1-bit models (BitNet)
- Mixture of Experts (MoE)
- Dynamic neural networks
- Neuromorphic computing
- Photonic processors

Research Frontiers:

- Sub-1-bit quantization
- Learned compression
- Architecture search for efficiency
- Federated model compression

Industry Trends:

- Every device runs AI (AIoT)
- Privacy-preserving inference
- Real-time translation everywhere
- Personalized on-device models
- Zero-latency assistants

Hardware Evolution:

- Analog AI chips
- In-memory computing
- Quantum advantage?
- Brain-inspired chips

The future: AI everywhere, invisible, instant, and sustainable

Week 11 Exercise: Deploy Your Own Mobile Model

Your Mission: Take a large model and make it mobile-ready

Part 1: Baseline and Profiling

- Choose a pretrained model (BERT, GPT-2, etc.)
- Profile: size, latency, memory, energy
- Identify bottlenecks
- Set target constraints (e.g., 50MB, 100ms)

Part 2: Apply Compression

- Implement INT8 quantization
- Distill to smaller student
- Apply magnitude pruning
- Combine techniques

Part 3: Deploy and Benchmark

- Export to ONNX/TFLite
- Run on real device (phone/Raspberry Pi)
- Measure real-world performance
- Compare accuracy vs efficiency
- Create deployment package

Bonus: Build demo app showcasing your efficient model!

Key Takeaways: Efficiency Enables Everything

What we learned:

- Models are vastly overparameterized
- 90% compression with minimal loss
- Quantization: Smaller and faster
- Distillation: Transfer capabilities
- Hardware matters immensely

The efficiency toolkit:

Quantization → Distillation → Pruning → Hardware optimization

Why it matters:

- Enables edge deployment
- Reduces costs dramatically
- Improves user experience
- Environmental sustainability

Next week: Ethics and Future Directions

With great power comes great responsibility - what should we build?

References and Further Reading

Foundational Papers:

- Hinton et al. (2015). "Distilling the Knowledge in a Neural Network"
- Han et al. (2016). "Deep Compression: Compressing DNNs"
- Frankle & Carbin (2019). "The Lottery Ticket Hypothesis"

Quantization:

- Jacob et al. (2018). "Quantization and Training of Neural Networks"
- Gholami et al. (2021). "A Survey of Quantization Methods"
- Dettmers et al. (2022). "LLM.int8(): 8-bit Matrix Multiplication"

Practical Resources:

- PyTorch Quantization Documentation
- TensorFlow Lite Guide
- ONNX Runtime Optimization
- Edge Impulse Platform