

Retrieval-Augmented Generation

Grounding LLMs in External Knowledge

NLP Course – Lecture 1

Advanced Topics in Natural Language Processing

Why Do LLMs Hallucinate?

The Problem

- LLMs predict the most likely next token
- They have no access to real-time information
- Knowledge is frozen at training time
- No mechanism to verify facts

The Solution: RAG

- Retrieve relevant documents first
- Augment the prompt with facts
- Generate grounded responses
- Cite sources for verification

This lecture: How to ground LLMs in external knowledge

RAG is the most widely deployed technique for making LLMs factually accurate.

Act I: RAG & AI Agents

Making LLMs Useful in the Real World

The Problem

LLMs confidently state wrong facts:

- “The current CEO of OpenAI is...” (outdated)
- “The 2024 Olympic gold medalist was...” (unknown)
- “Your company’s Q3 revenue was...” (not in training data)

Root Causes

- Knowledge frozen at training time
- No access to private/recent information
- Model “fills in gaps” with plausible text

Why This Matters

For real applications, we need:

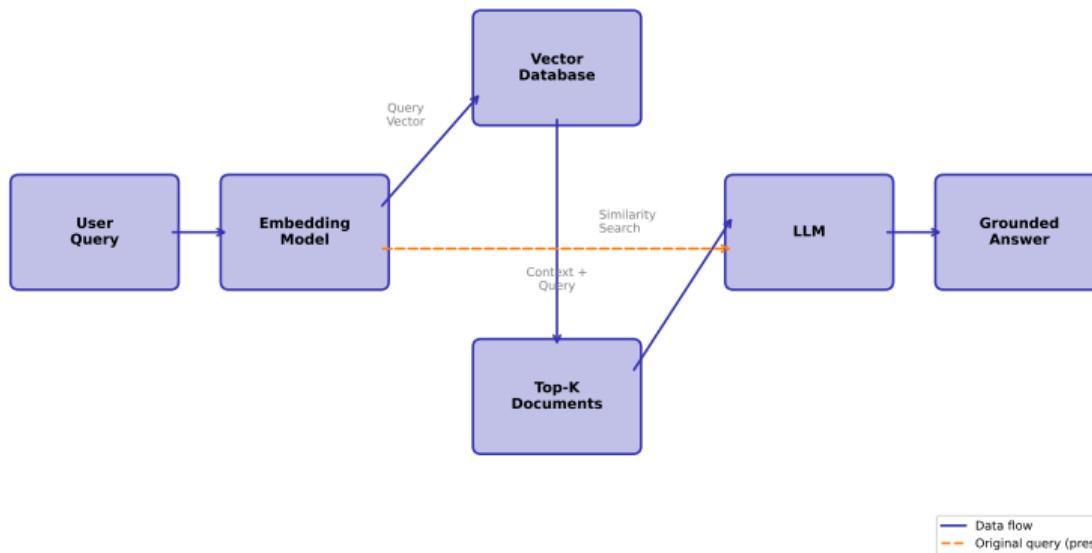
- Access to current information
- Grounding in verifiable sources
- Ability to say “I don’t know”

Connection to Ethics Week

Remember: LLMs don’t “know” anything – they predict tokens. Without grounding, this is dangerous.

Solution: Don’t try to store everything in parameters. Retrieve at inference time.

RAG (Retrieval-Augmented Generation) Architecture



Key insight: Separation of concerns – parametric knowledge (the model) vs. retrieved knowledge (the database)

Core Idea

Instead of: $p(y|x)$ (generate from query alone)
RAG marginalizes over retrieved documents:

$$p(y|x) = \sum_{z \in \text{top-}k} p(z|x) \cdot p(y|z)$$

Why no z on left? We sum over all z (marginalization) –
the result depends only on x .

Where: x = query, z = retrieved doc, y = response

Key Equation: Dense Retrieval

$$\text{sim}(q, d) = \frac{E_q(q)^T \cdot E_d(d)}{\|E_q(q)\| \cdot \|E_d(d)\|}$$

Retrieval probability (softmax):

$$p(z_i|x) = \frac{\exp(\text{sim}(x, z_i)/\tau)}{\sum_{j=1}^k \exp(\text{sim}(x, z_j)/\tau)}$$

You Already Know This!

This is just attention over an external memory.

Lewis et al. (2020): “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks”

Query x : "What is the capital of France?"

Retrieved Documents (with similarity scores):

z_1 : "Paris is the capital and largest city of France..." 0.92

z_2 : "France is a country in Western Europe..." 0.71

z_3 : "The Eiffel Tower is located in Paris..." 0.65

Generation Probabilities $p(y|x, z_i)$:

For answer y = "Paris":

- $p(y|x, z_1) = 0.95$ – directly states "Paris is capital"
- $p(y|x, z_2) = 0.40$ – mentions France, not Paris
- $p(y|x, z_3) = 0.70$ – mentions Paris, not as capital

Step 1: Retrieval Probabilities

$$\text{Softmax: } p(z_i|x) = \frac{e^{\text{sim}_i}}{\sum_j e^{\text{sim}_j}}$$

$$\begin{aligned} p(z_1|x) &= 0.52 && (\text{most relevant}) \\ p(z_2|x) &= 0.27 \\ p(z_3|x) &= 0.21 \end{aligned}$$

Step 2: Marginalization

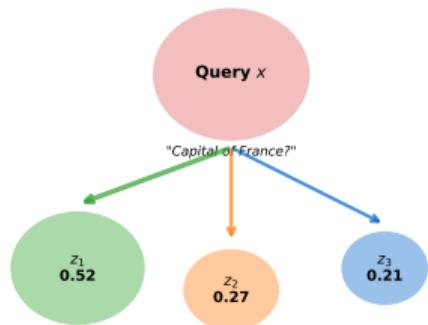
$$\begin{aligned} p(y|x) &= \sum_{i=1}^3 p(z_i|x) \cdot p(y|x, z_i) \\ &= 0.52 \times 0.95 && (\text{from } z_1) \\ &+ 0.27 \times 0.40 && (\text{from } z_2) \\ &+ 0.21 \times 0.70 && (\text{from } z_3) \\ &= 0.494 + 0.108 + 0.147 \\ &= \mathbf{0.75} \end{aligned}$$

Key: $p(z|x)$ = how relevant is doc? $p(y|x, z)$ = given this doc, how likely is answer?

RAG Conditional Probabilities: Visual Intuition

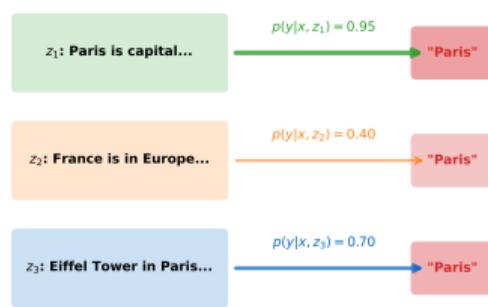
RAG Conditional Probabilities: Visual Intuition

$p(z_i|x)$: Retrieval Probability



Size/arrow = retrieval probability
(how relevant is doc to query?)

$p(y|x, z_i)$: Generation Probability



Arrow thickness = generation probability
(given doc, how likely is answer?)

$$p(y|x) = \sum_i p(z_i|x) \cdot p(y|x, z_i)$$

$$p(y|x) = 0.75$$

$$\begin{aligned} z_3: & 0.21 \times 0.70 \\ & = 0.147 \\ z_2: & 0.27 \times 0.40 \\ & = 0.108 \\ z_1: & 0.52 \times 0.95 \\ & = 0.494 \end{aligned}$$

$$\begin{aligned} &= 0.52 \times 0.95 \\ &+ 0.27 \times 0.40 \\ &+ 0.21 \times 0.70 \\ &= 0.75 \end{aligned}$$

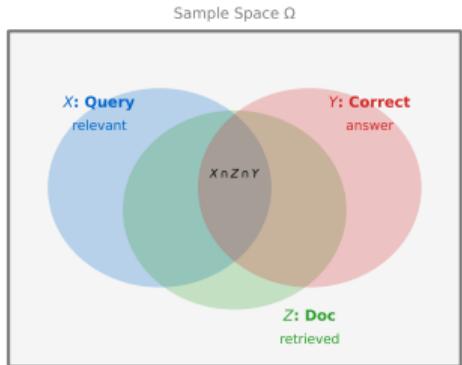
Each document contributes to final answer
weighted by its retrieval probability

Marginalization: Sum over all docs, each weighted by retrieval probability times generation probability

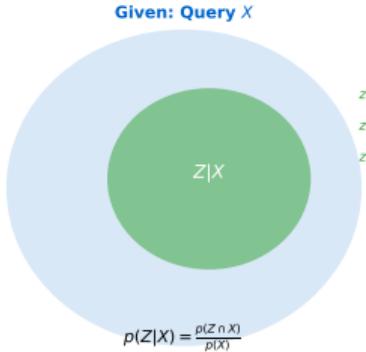
RAG Probabilities: Venn Diagram Interpretation

RAG Probabilities: Venn Diagram Interpretation

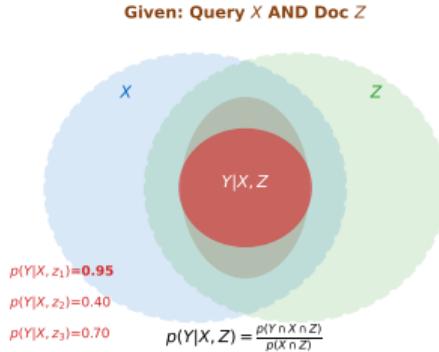
Sample Space: All Possible Outcomes



$p(Z|X)$: Retrieval Probability



$p(Y|X, Z)$: Generation Probability



X = query context | Z = retrieved doc | Y = correct answer "Paris"

How likely is doc Z retrieved given query X?

Given query AND doc, how likely is correct answer?

Conditional probability: We restrict the sample space to the given event, then measure probability within it

Query and Response

- x – User query (input question)
- y – Generated response (output)
- q – Query after embedding

Documents and Retrieval

- z – Retrieved document(s)
- z_i – The i -th retrieved document
- \mathcal{Z} – Full document corpus
- d – Single document in corpus
- k – Number of documents retrieved (top- k)

Embedding Functions

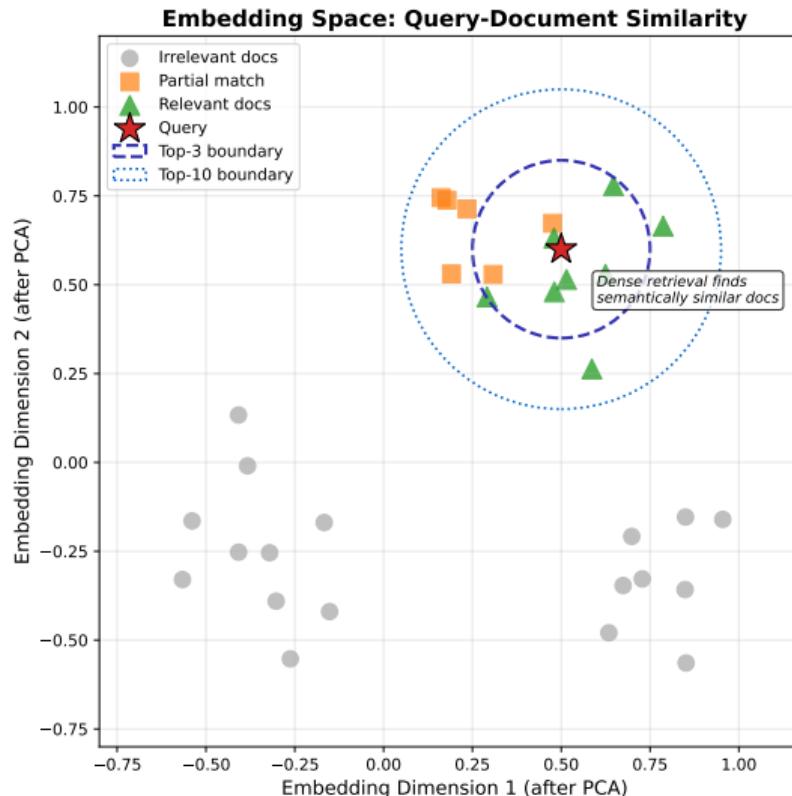
- $E_q(\cdot)$ – Query encoder (embeds queries)
- $E_d(\cdot)$ – Document encoder (embeds documents)
- Often $E_q = E_d$ (same encoder for both)

Similarity and Probability

- $\text{sim}(q, d)$ – Cosine similarity between query and document vectors
- τ – Temperature parameter (controls softmax sharpness)
- $p(z|x)$ – Probability of retrieving document z given query x
- $p(y|x, z)$ – Generation probability given query and retrieved docs

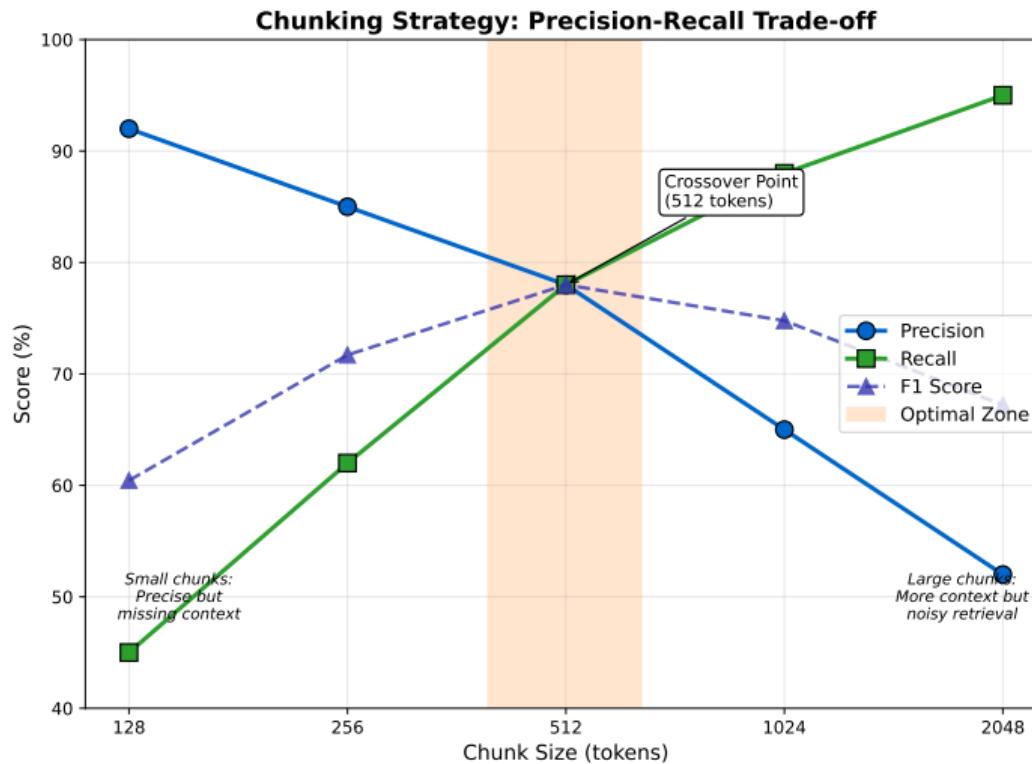
Understanding notation: Embedding similarity drives retrieval, retrieval augments generation

Visualizing the Embedding Space



Dense retrieval works by finding documents whose embeddings are closest to the query embedding

Chunking Trade-offs: Precision vs Recall



Rule of thumb: Start with 512 tokens, adjust based on your retrieval quality metrics

Embedding Models

- Sentence transformers
- OpenAI embeddings
- Cohere, Voyage, etc.

Output

Dense vectors (e.g., 1536-dim)

Vector Databases

- FAISS (Facebook)
- Pinecone (managed)
- ChromaDB (local)
- Weaviate, Milvus

Key Operation

Approximate nearest neighbor search

Chunking Strategies

- Fixed-size (512 tokens)
- Semantic (by paragraph)
- Hierarchical (nested)
- Sliding window

Trade-off

Small chunks = precise retrieval
Large chunks = more context

The choice of chunking strategy significantly impacts retrieval quality

What Is a Vector Database?

Specialized database for storing and querying high-dimensional vectors (embeddings).

Key Operation: ANN Search

Approximate Nearest Neighbor (ANN):

- Exact search is $O(n)$ – too slow
- ANN trades accuracy for speed
- Typical: 95%+ recall at 10-100x speedup

Index Structures

- HNSW (Hierarchical Navigable Small World)
- IVF (Inverted File Index)
- LSH (Locality Sensitive Hashing)

Popular Vector Databases

Open Source:

- FAISS (Meta) – In-memory, very fast
- ChromaDB – Simple, Python-native
- Milvus – Distributed, scalable
- Weaviate – GraphQL interface

Managed Services:

- Pinecone – Fully managed
- Qdrant – Self-hosted or cloud

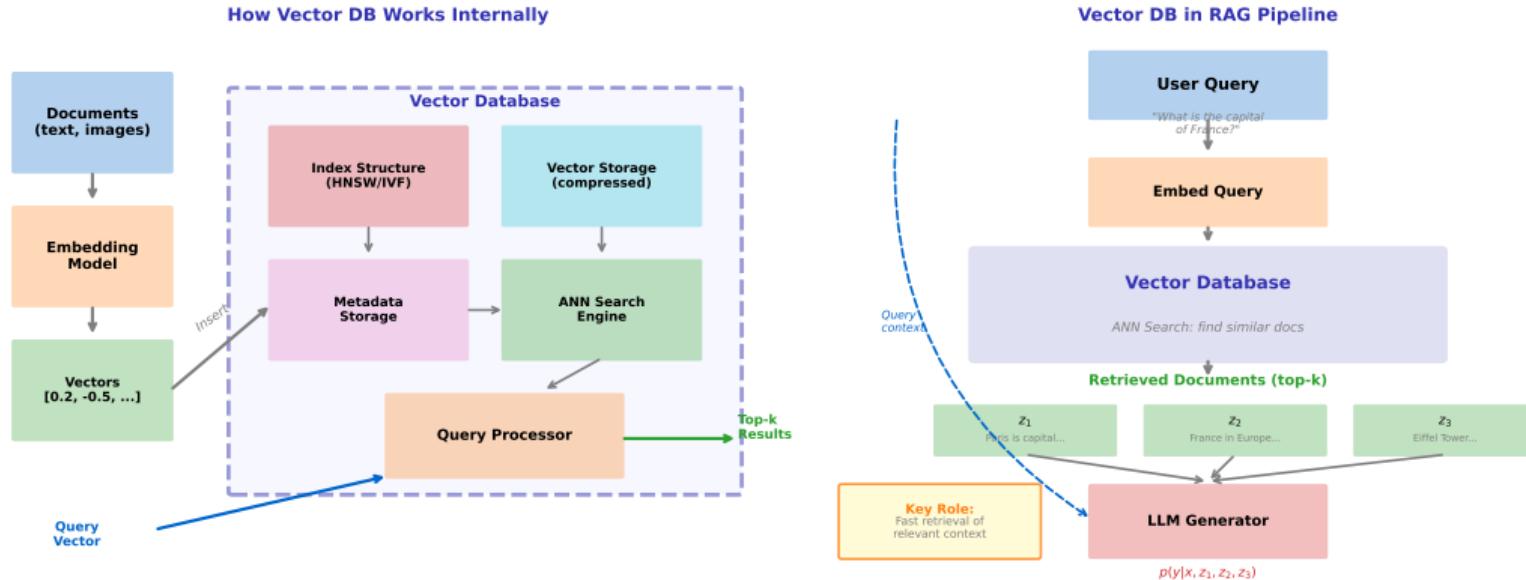
Typical Workflow

1. Embed documents → vectors
2. Store vectors with metadata
3. Query: embed query → find top- k similar

Vector databases are the “memory” that makes RAG possible at scale

Vector Database: Architecture and Role in RAG

Vector Database: Architecture and Role in RAG



Vector DBs enable fast retrieval: embed documents once, search in milliseconds at query time

Approximate Nearest Neighbor: Why and How

Approximate Nearest Neighbor (ANN): The Core Idea

The Problem

Given: Database of n vectors
 $D = \{d_1, d_2, \dots, d_n\}$

Query: Find k vectors closest to q

Exact solution requires:

- Compute distance to ALL n vectors
- Sort and return top- k
- Time: $O(n)$ per query
 $n = 1 \text{ billion? That is } 1 \text{ billion distance calculations per query!}$

The Mathematics

Exact k-NN:

$$N_k(q) = \operatorname{argmin}_{S=k} \max_{d \in S} \|q - d\|$$

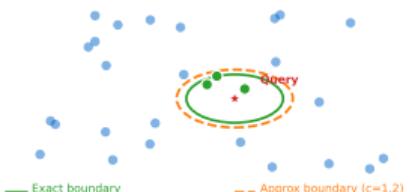
c -Approximate k-NN:

$$\text{For all } d \in ANN_k(q): \\ \|q - d\| \leq c * \|q - d^*\|$$

where d^* is the true k -th neighbor
and $c > 1$ is the approximation factor.

^{1.05} means we accept neighbors up to 5% farther than optimal

Visual Intuition



The Trade-off

Method	Time	Recall	Use Case
Exact (brute)	$O(n)$	100%	Small datasets
IVF	$O(\sqrt{n})$	~95%	Medium scale
HNW	$O(\log n)$	~99%	Production
LSH	$O(1)^*$	~90%	Massive scale

Key: Accept 1-5% accuracy loss for 100-1000x speedup

* LSH: $O(1)$ query but $O(n)$ space for hash tables

ANN is the key enabler for billion-scale vector search: trade small accuracy for massive speedup

Exact k -NN Problem

Given query q and database $D = \{d_1, \dots, d_n\}$, find:

$$N_k(q) = \arg \min_{S \subseteq D, |S|=k} \max_{d \in S} \|q - d\|$$

c -Approximate k -NN

An algorithm returns $\text{ANN}_k(q)$ such that:

$$\forall d \in \text{ANN}_k(q) : \|q - d\| \leq c \cdot \|q - d^*\|$$

where d^* is the **true k -th nearest neighbor** and $c \geq 1$ is the **approximation factor**.

What This Means

- $c = 1.0$: Exact (no approximation)
- $c = 1.05$: At most 5% farther
- $c = 1.10$: At most 10% farther

The Trade-off

- $c \rightarrow 1$ Slower, exact
 $c > 1$ Faster, approximate

In Practice

Most systems achieve $c \approx 1.01$ to 1.05 with $100\text{--}1000\times$ speedup.

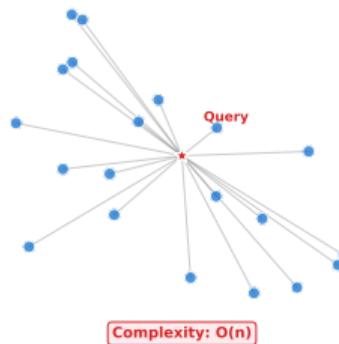
The c -approximation guarantee means returned neighbors are at most c times farther than the true nearest

HNSW: The Most Popular ANN Algorithm

Exact Search vs HNSW: Why Approximate is Faster

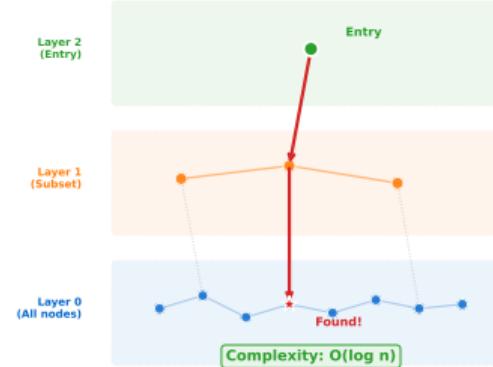
Exact Search (Brute Force)

Compare query to ALL documents



HNSW (Hierarchical Navigable Small World)

Navigate graph: sparse top -> dense bottom



Trade-off: HNSW achieves 95-99% recall with 100-1000x speedup over exact search

HNSW builds a navigable graph: start at sparse top layers, greedily descend to find nearest neighbors

The Key Idea

Combine two concepts:

1. **Skip Lists:** Hierarchical layers for $O(\log n)$ traversal
2. **Navigable Small World:** Each node connected to “nearby” nodes

Layer Structure

- Layer 0: All n nodes (dense)
- Layer 1: $\sim n/m_L$ nodes
- Layer 2: $\sim n/m_L^2$ nodes
- Top: Few entry points

How are nodes assigned?

Each node's max layer is **random**:

$$\ell = \lfloor -\ln(\text{uniform}(0, 1)) \cdot m_L \rfloor$$

Most nodes: layer 0 only. Few “lucky” nodes reach higher layers (like express stops).

The hierarchical structure enables logarithmic search: coarse navigation at top, fine-grained at bottom

Construction Algorithm

For each new vector v :

1. Sample max layer ℓ (formula on left)
2. Insert v into layers $0, 1, \dots, \ell$
3. At each layer, connect to M nearest neighbors

Key Parameters

M Max connections/node

ef Search beam width

m_L Level multiplier

Typical: $M = 16$, $ef = 100$, $m_L = 1/\ln(M)$

Intuition: Like a subway system – express lines (top layers) connect major hubs, local lines (layer 0) reach everywhere.

Greedy Search Procedure

1. Start at entry point (top layer)
2. At each layer:
 - Greedily move to nearest neighbor
 - Repeat until no closer neighbor exists
3. Descend to next layer
4. At layer 0: expand search with beam width ef
5. Return top- k from candidates

Complexity

Search: $O(\log n)$
Insert: $O(\log n)$
Space: $O(n \cdot M)$

Why It Works

Small World Property: Any two nodes connected by short path ($\sim \log n$ hops).

Hierarchical Speedup: Top layers skip large distances; bottom layers refine.

Pseudocode

```
search(q, k, ef):  
    ep = entry_point  
    for layer in top...1:  
        ep = greedy(q, ep, layer)  
    cands = beam(q, ep, L0, ef)  
    return top_k(cands, k)
```

ef controls accuracy/speed trade-off.

HNSW achieves >99% recall with 10–100× speedup; used in FAISS, Pinecone, Weaviate, Qdrant

HNSW: A Simple Example

Setup: 8 cities, find nearest to query “Berlin”

Layer 2 (Top) – 2 nodes

Entry points: Paris, Tokyo

Query: Berlin → Check Paris, Tokyo
→ Paris closer → **go to Paris**

Layer 1 – 4 nodes

Paris, Tokyo, London, Sydney

From Paris → Check neighbors
→ London closer → **go to London**

Layer 0 (Bottom) – all 8 nodes

From London → Check all neighbors
→ **Found: Amsterdam** (nearest!)

What Happened

Layer 2: 2 comparisons

Layer 1: 3 comparisons

Layer 0: 4 comparisons

Total: **9 comparisons**

Brute Force

8 comparisons (check all)

With 1 Billion Nodes

Brute: 1,000,000,000

HNSW: ~30 (log scale!)

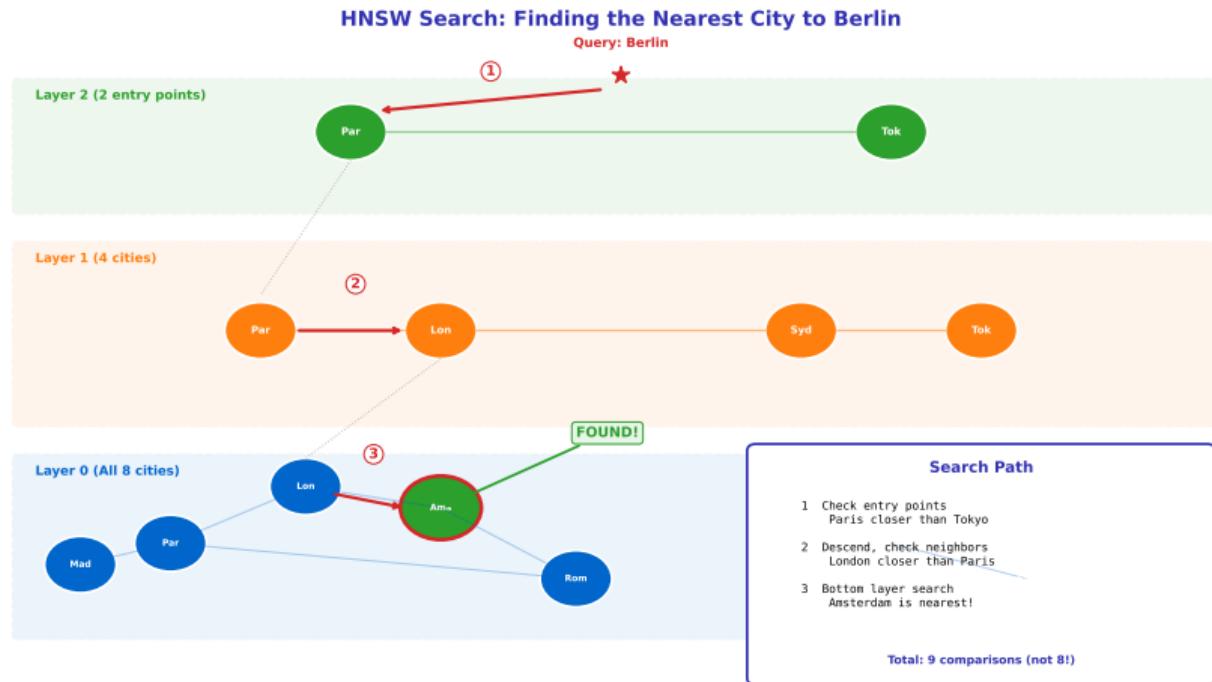
Key Insight

Top layers = “highways”

Bottom layer = “local streets”

HNSW is like using a map: zoom out to find the region, then zoom in to find the exact location

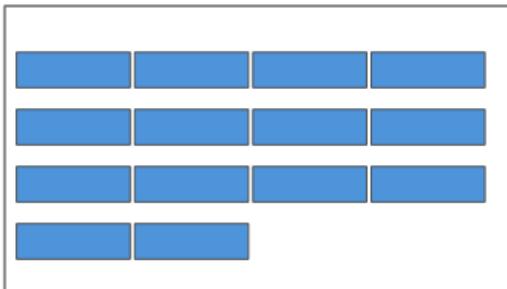
HNSW: Visual Walkthrough



Each layer narrows the search: start broad at the top, refine at the bottom

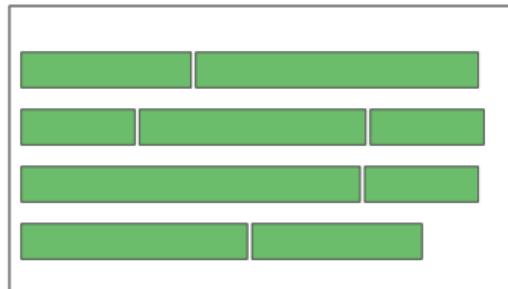
Chunking Strategies Deep Dive

Fixed-Size Chunking



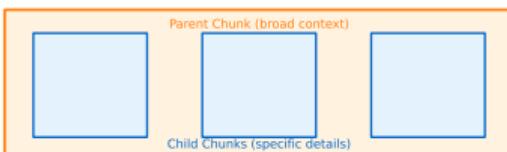
Simple: Split every N tokens (e.g., 512)

Semantic Chunking



Split at paragraph/section boundaries

Hierarchical Chunking



Multi-level: Query routes to appropriate granularity

Sliding Window



Overlapping windows: No info lost at boundaries

Chunking is often the difference between RAG that works and RAG that fails – start with 512 tokens, 10% overlap

Naive RAG

- Simple retrieve-then-generate
- Fixed number of chunks
- No query preprocessing

Advanced RAG

- Query rewriting
- Re-ranking retrieved documents
- Iterative retrieval
- Multi-stage retrieval

Modular RAG

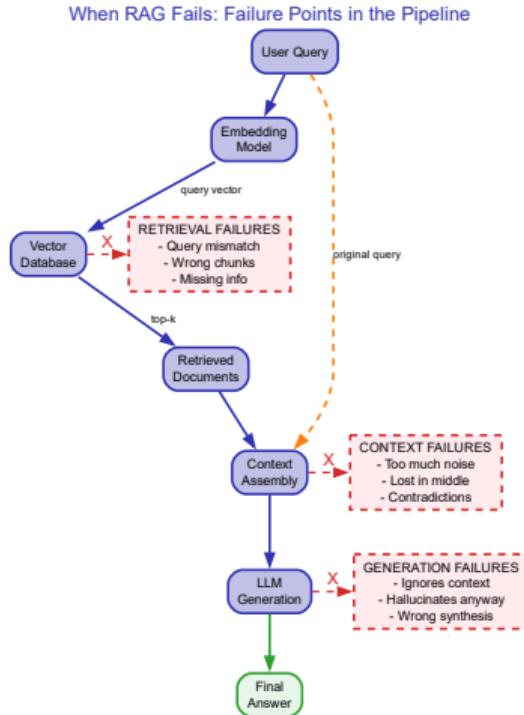
- Self-RAG: decide *when* to retrieve
- CRAG: correct retrieval errors
- Adaptive: retrieve more if needed

Agentic RAG (2024+)

- Agent decides retrieval strategy
- Multiple retrieval sources
- Tool use for specialized queries

Trend: More intelligence in the retrieval process, not just generation

When RAG Fails: Failure Points in the Pipeline



RAG requires careful engineering at every pipeline stage

Retrieval Fixes

- Query expansion/rewriting
- Multi-stage retrieval
- Better chunking strategies
- Cross-encoder re-ranking
- Multiple retrieval passes

Context Fixes

- Smart chunk ordering
- Compression/summarization
- Relevance filtering
- Hierarchical retrieval
- Attention to chunk boundaries

Generation Fixes

- Instruction tuning for RAG
- Citation requirements
- Self-consistency checks
- Confidence calibration
- Fallback to “I don’t know”

“Lost in the middle” problem: LLMs often ignore content in the middle of long contexts.

Solution: Place most relevant chunks at beginning and end.

Each failure mode has specific mitigations – production RAG requires all of them

Key Takeaways: RAG

1. **RAG solves hallucination** by grounding LLMs in external documents
2. **Vector search** enables millisecond retrieval from billions of documents
3. **HNSW** provides $O(\log n)$ approximate nearest neighbor search
4. **Chunking strategy** critically affects retrieval quality
5. **RAG can fail** at retrieval, ranking, or generation stages

Key Equations:

- Dense retrieval: $\text{sim}(q, d) = \cos(E_q(q), E_d(d))$
- RAG probability: $p(y|x) = \sum_z p(z|x) \cdot p(y|z, x)$

RAG is the foundation of most production LLM applications today.

Further Reading: RAG

Foundational Papers:

- Lewis et al. (2020) - “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks”
- Karpukhin et al. (2020) - “Dense Passage Retrieval”
- Malkov & Yashunin (2018) - “HNSW: Hierarchical Navigable Small World Graphs”

Tools & Frameworks:

- Vector DBs: Pinecone, Weaviate, ChromaDB, FAISS
- Frameworks: LangChain, LlamaIndex

Repository: github.com/Digital-AI-Finance/Natural-Language-Processing