

# LLM Reasoning

From Chain-of-Thought to Test-Time Compute

NLP Course – Lecture 3

Advanced Topics in Natural Language Processing

# Why Do LLMs Struggle with Reasoning?

## The Problem

- LLMs give instant responses
- No “working memory” for computation
- Multi-step problems require planning
- Direct answers often wrong

## The Solution

- Chain-of-Thought: think step by step
- Test-time compute: think longer
- Process reward models: verify steps
- DeepSeek-R1: trained to reason

This lecture: Teaching LLMs to think before answering

---

Reasoning capabilities have dramatically improved since 2022.

## Reliability Issues

- Agents get stuck in loops
- Wrong tool selection
- Hallucinated tool parameters
- Failure to know when to stop

## Cost Accumulation

- Each step = API call
- Complex tasks = many calls
- Costs can spiral quickly

## Security Concerns

- Tool access = system access
- Prompt injection attacks
- Unintended actions

## What Works Today

- Well-defined, bounded tasks
- Human oversight/approval
- Retrieval-heavy workflows
- Single-domain expertise

*"Agents are promising but not production-ready for autonomous operation." – 2024 consensus*

---

Connection to reasoning: Better reasoning = more reliable agents. This leads us to Act II...

## Act II: Reasoning in LLMs

Chain-of-Thought, Test-Time Compute, and DeepSeek-R1

# The Surprising Discovery (2022)

## The Experiment

Google researchers found something remarkable:  
Simply adding “*Let’s think step by step*” to a prompt  
improved math accuracy by **40%+**

No model changes. No fine-tuning. Just a prompt.

## Why?

- Creates “scratchpad” for computation
- Forces sequential reasoning
- Mirrors human problem-solving

## Before CoT

Q: Roger has 5 tennis balls. He buys 2 cans of 3 balls each. How many does he have now?  
A: **11** (*direct answer, sometimes wrong*)

## After CoT

Q: [same question] *Let’s think step by step.*  
A: Roger starts with 5 balls.  
He buys  $2 \text{ cans} \times 3 \text{ balls} = 6 \text{ balls}$ .  
**Total =  $5 + 6 = 11 \text{ balls}$**   
(*reasoning chain makes answer verifiable*)

---

Wei et al. (2022): “Chain-of-Thought Prompting Elicits Reasoning in Large Language Models”

## Standard Prompting

Direct answer generation:

$$p(\text{answer}|\text{question})$$

The model jumps straight to the answer in a single forward pass.

### Problem

Complex reasoning requires multiple “steps” – but each token is generated independently.

## Chain-of-Thought Prompting

Decompose into two stages:

$$p(r|q) \cdot p(a|r, q)$$

Where:

- $q$  = question
- $r$  = reasoning chain
- $a$  = final answer


### Key Insight

The reasoning tokens  $r$  create intermediate computation space that the model can “use” to solve harder problems.

---

Connection to Week 9 (Decoding): CoT changes what the model generates, not how it decodes

# Intermediate Computation Space: Why CoT Works



**Key insight:** Reasoning tokens create a “scratchpad” that enables multi-step computation within a single generation

# Chain-of-Thought Variants

## Zero-Shot CoT

Just add: "Let's think step by step"  
No examples needed. Works surprisingly well.

## Few-Shot CoT

Provide examples with reasoning chains:

*Example 1: [problem] [reasoning] [answer]*

*Example 2: [problem] [reasoning] [answer]*

*Your turn: [problem]*

More reliable but requires good examples.

## Self-Consistency

Generate  $N$  reasoning chains (with temperature  $> 0$ ).  
Take majority vote on final answer:

$$\hat{a} = \arg \max_a \sum_{i=1}^N \mathbf{1}[a_i = a]$$

## Tree of Thoughts

Explore multiple reasoning *paths*, not just one chain.  
Allows backtracking on dead ends.

## Least-to-Most


Decompose into subproblems first, then solve.

---

CoT is the most powerful prompt engineering technique discovered so far

# Self-Consistency: Sample and Vote

## Self-Consistency: Multiple Reasoning Paths with Majority Voting



Self-consistency adds 5-10% accuracy on top of CoT by marginalizing over reasoning paths

## Example 1: Math Problem

*Q: A farmer has 17 sheep. All but 9 die. How many are left?*

**Let's think step by step.** "All but 9 die" means 9 survive. Answer: **9 sheep**

---

## Example 2: Logic Puzzle

*Q: If all roses are flowers and some flowers fade quickly, can we conclude that some roses fade quickly?*

**Let's analyze:** (1) All roses ⊂ flowers. (2) Some flowers fade quickly – but which ones? Could be non-rose flowers. (3) We cannot conclude roses fade quickly. Answer: **No, invalid inference**

---

## Example 3: Code Debugging

*Q: Why does `sum([1,2,3][:2])` return 3, not 6?*

**Let's trace:** (1) `[1,2,3]` creates list. (2) `[:2]` slices indices 0,1 → `[1,2]`. (3) `sum([1,2]) = 3`. The slice excludes index 2.

---

CoT works across domains: math, logic, code – wherever step-by-step reasoning helps

# The Paradigm Shift: Test-Time Compute

## Old Paradigm: Scale Training

$$\text{Performance} \propto \log(\text{Parameters})$$

Bigger models = Better performance

GPT-2 → GPT-3 → GPT-4

### Problem

Training cost grows exponentially.

Diminishing returns at scale.

One-size-fits-all computation.

## New Paradigm: Scale Inference

$$\text{Performance} \propto \log(\text{Test-Time Compute})$$

Same model, more “thinking” = Better answers

### Key Insight

Not all questions need the same compute.

Hard problems deserve more thinking time.


Let the model allocate compute adaptively.

**This is revolutionary!**

---


Snell et al. (2024): “Scaling LLM Test-Time Compute Optimally can be More Effective than Scaling Model Parameters”

# Test-Time vs Pre-Training Scaling



For hard problems, test-time compute scaling can outperform pre-training scaling at equivalent FLOPs

# Inference Token Scaling: More Thinking = Better Answers



**Key insight:** Reasoning models show log-linear improvement with inference tokens; standard models plateau

## 1. Best-of-N with Verifiers

Generate  $N$  candidate solutions.

Score each with a verifier (PRM).

Select the best one.

$$\hat{y} = \arg \max_{y \in \{y_1, \dots, y_N\}} r_{\text{PRM}}(y)$$

## Process Reward Models (PRMs)

Score each step of reasoning:

$$r_{\text{PRM}}(s_1, \dots, s_T) = \prod_{t=1}^T p(\text{correct} | s_1, \dots, s_t)$$

More compute = more candidates = better selection.

## 2. Extended Reasoning

Let the model think for more tokens.

Longer reasoning = better answers.

### How o1 Does It

Hidden “thinking” tokens before answering.

Model trained to use this space productively.

### Adaptive Allocation

Easy questions: Short reasoning

Hard questions: Long reasoning

The model learns to allocate compute based on difficulty.

---

Both mechanisms: more compute at inference = better results (with diminishing returns)

# The Cost-Quality Tradeoff

Model	Tokens Generated	Accuracy	Relative Cost
GPT-4 (direct)	~50	78%	1x
GPT-4 (CoT prompt)	~150	89%	3x
o1-mini	~500	95%	10x
o1	~2000	97%	40x
o1-pro	~5000+	99%	100x+

## Practical Implication

You can choose your accuracy/cost tradeoff:

- Simple queries: Use fast, cheap model
- Complex reasoning: Invest in more compute
- Critical decisions: Use maximum reasoning

---

This is like choosing car vs. plane – different tools for different journeys

## The Announcement

DeepSeek (Chinese lab) releases R1:

- Matches OpenAI o1 performance
- Fully open-source (weights + paper)
- Fraction of training cost
- Multiple distilled sizes available

## Why It Matters

Demonstrated that reasoning can be achieved with:

- Open research
- Smaller budgets
- Novel training approaches

## Key Results

AIME 2024 (math olympiad):

**15.6% → 71.0%** (pass@1)

Matches o1-1217 on most benchmarks.

## Available Models

DeepSeek-R1-Distill-Qwen:

1.5B, 7B, 14B, 32B

DeepSeek-R1-Distill-Llama:

8B, 70B

All on HuggingFace, open weights.

---

DeepSeek-AI (2025): "DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning"

## The Experiment

What if we train reasoning with *pure RL*, no supervised fine-tuning?

## DeepSeek-R1-Zero

- Start from base model
- Apply RL directly
- Reward only final answer correctness
- No human demonstrations of reasoning

## Result

The model *spontaneously* learned to:

- Generate reasoning chains
- Self-verify answers
- Reflect on mistakes
- Allocate more tokens to hard problems

## Why This Is Shocking

“Reasoning” emerged from the objective alone.  
No one told the model *how* to reason – just rewarded correct answers.

## Emergent Behaviors

*Self-verification:*

“Let me check:  $3 \times 5 = 15$ , correct.”

*Reflection:*

“Wait, I made an error. Let me reconsider...”

*Extended thinking:*

Hard problems → longer reasoning traces

## Implication

Reasoning might be more fundamental than we thought – it emerges when you optimize for correctness.

---

This suggests reasoning is an “attractor” in the optimization landscape, not a special trick

## Standard RL (PPO)

Requires:

- Critic network (value function)
- Reward model
- Complex optimization

## GRPO Simplification

No critic network needed!

Compute advantage relative to group:

$$A(x, y) = r(y) - \frac{1}{|G|} \sum_{y' \in G} r(y')$$

For each prompt, generate multiple outputs, compare to each other.

## Rule-Based Rewards

No neural reward model either!

*Accuracy reward:*

$$r_{\text{acc}} = \mathbf{1}[\text{answer correct}]$$

*Format reward:*

$$r_{\text{fmt}} = \mathbf{1}[\text{ithink}_L \dots \text{i/think}_L \text{ tags present}]$$

## Why This Works

For math/code: correctness is verifiable.

No need to learn “what humans prefer.”

---

**GRPO: Simpler than PPO, no reward model, no critic – yet achieves state-of-the-art reasoning**

## Stage 1: Cold Start (Optional)

Small amount of SFT on reasoning examples.  
Teaches the format: <think>...</think>  
Not strictly necessary (R1-Zero skips this).

## Stage 2: Reasoning RL

Pure RL with GRPO.  
Reward: correctness + format.  
Model learns to reason.

## Stage 3: Rejection Sampling

Generate many responses from RL model.  
Filter for correct + well-formatted.  
Creates high-quality reasoning dataset.

## Stage 4: Final SFT


Fine-tune on curated reasoning data.  
Adds general capabilities back.  
Balances reasoning with helpfulness.

## Result: DeepSeek-R1

---

Key insight: RL discovers reasoning, then SFT polishes and generalizes it

# DeepSeek-R1: The Four-Stage Training Pipeline



DeepSeek-R1 shows that open-source reasoning can match proprietary models with clever training

## OpenAI o1

- Closed source, proprietary
- Hidden “thinking” tokens (not shown to user)
- Likely uses process supervision
- Rumored to use search/planning
- Available via API only

### Strengths

Polish, reliability, integration with OpenAI ecosystem.

## DeepSeek-R1

- Open source (weights + paper)
- Visible reasoning traces
- Pure RL approach documented
- Distilled to many sizes
- Run locally or via API

### Strengths

Transparency, customizability, research value.

### Performance

Comparable on most benchmarks.

---

The gap between closed and open reasoning models is narrowing rapidly

## Key Takeaways: LLM Reasoning

1. **Chain-of-Thought** dramatically improves reasoning (+40% on math)
2. **Intermediate tokens** serve as computational scratchpad
3. **Test-time compute** is the new scaling paradigm
4. DeepSeek-R1 showed pure RL can develop reasoning
5. **Process reward models** enable verification of reasoning steps

**Key Insight:** “Let the model think longer” is often more effective than making models bigger.

---

Reasoning capabilities define the frontier of AI capabilities in 2025.

## Foundational Papers:

- Wei et al. (2022) - "Chain-of-Thought Prompting"
- Wang et al. (2023) - "Self-Consistency"
- DeepSeek (2025) - "DeepSeek-R1"
- OpenAI (2024) - "o1 System Card"

## Key Concepts:

- Test-time compute scaling
- Process Reward Models (PRMs)
- GRPO (Group Relative Policy Optimization)

---

Repository: [github.com/Digital-AI-Finance/Natural-Language-Processing](https://github.com/Digital-AI-Finance/Natural-Language-Processing)